

OpenAI's 719 AI-generated math papers, explained from zero

On 2026-10-06 OpenAI published **722 math manuscripts (372 result families)** generated by an unreleased internal model. Among the claims: the quasi-Riemann hypothesis, the Unique Games Conjecture, matrix-multiplication exponent $\omega \leq 9/4$, non-amenability of Thompson's group F, and the plane not being 5-colorable — a run of problems open for decades. The next day three manuscripts were withdrawn over a sign error. Mathematicians dubbed it the "Mathocalypse." The goal of this deck: using **only linear algebra, calculus and high-school math**, take 11 representative results far enough that you can judge how big each is, how hard it was, why it matters, and how far it has been verified.

First edition 2026-10-08; revised 10-09 (each result now has one from-zero version; new §14 on the Kakeya problem and Fourier restriction; Part II outlook rewritten as §18-§24) · repo snapshot openai/math@fd4aeeb (includes the Oct 7 withdrawals and fixes) · every fact is sourced; [P] = primary material read (paper / repo / Lean source), [R] = secondary reporting, Δ = unverified or conflicting · English edition of notes.html (Chinese)

ORIENTATION

0. What this deck will and won't give you

YOU WILL GET

A "ladder" understanding of each of 11 headline results: start from a toy example you can compute by hand and climb level by level to the real statement; learn how long the problem was stuck and where; see exactly what OpenAI claims, how much the Lean formalization covers, and what human experts have said so far; and what it means for mathematics, for AI, and for AI safety.

YOU WON'T GET

The ability to judge whether any proof is correct. As of 2026-10-09, **none** of these 11 results has a human domain expert publicly saying "I have read and checked it." "Verification" in this deck means two things only: (1) a static review of the Lean source and statements, and (2) a summary of public reactions. Both are labeled with sources and confidence.

HOW THE RESULTS WERE CHOSEN

The deck covers 11 of the 372 result families. Selection criteria: (1) the problem is famous and recognized as hard by experts (those named in first-day coverage and by mathematicians came first); (2) a Lean formalization exists where possible, so the deck can say how far each claim is verified; (3) it can be built up from high-school math and connects

to computing or AI; (4) families with a published reasoning summary were preferred. The first edition (10-08) was in practice picked from four fields (number theory, theoretical computer science, combinatorics and group theory, physics and PDE) without a systematic pass over all 17 subject areas. The whole harmonic-analysis cluster (Takeya, Fourier restriction, Bochner–Riesz, local smoothing, Falconer distance) was missed that way and was added on 10-09 as §14. **Still not covered, though at least comparable by reputation:** Hilbert's tenth problem over $\mathbb{Q}$ (004), counterexamples to Hadwiger's conjecture (157), a counterexample to Kaplansky's zero-divisor conjecture (196), isomorphism of the free group factors (287), counterexamples to Baum–Connes and Kadison–Kaplansky (285), the Hilbert–Smith conjecture (304), uniform bounds in Hilbert's sixteenth problem (143), and Erdős's reciprocal-sum conjecture with quasipolynomial Szemerédi bounds (159) [speculation: judged from titles only, papers not read]. The full catalogue is in [overview.pdf](#).

0.1 Reading path

§1 fills in three basics (what mathematicians do, what an "open problem" is, how to judge how big a result is). §2–§3 cover the event and what Lean actually guarantees — **these two sections determine how you should read every "claim" that follows**. §4–§14 are the 11 results. Each is built up from the most basic concepts, with hand-checkable examples and [Check] exercises, and follows the same structure: *bottom line* → *concepts from zero* → *what OpenAI claims* → *how far it is verified* → *what it would mean if true*; read them in any order. §15 is the withdrawal case study, §16–§17 give cross-cutting patterns and implications for AI. **Part II (§18–§24) looks ahead:** if AI solves mathematics at scale, or its other abilities break through too, what happens in computing and cryptography, physics and chemistry, biology and medicine, AI itself, and the economy and governance? §18 calibrates with history and three capability curves; §19–§23 each pick the problems that truly matter in one field and explain whether math is the bottleneck, what AI changes, the deeper impacts and rough timing; §24 gives the big picture, a top 10 and a timeline. §25 collects reactions, §26 lists open questions and a learning path, §27 is the glossary and sources.

0.2 Cheat sheet

§	RESULT (FAMILY NO.)	OPEN FOR	LEAN	ONE LINE
4	Quasi-Riemann hypothesis (003)	~125 yrs	yes	ζ has no zeros at all in $\text{Re } s > 7/8$
5	Irrationality exponent of $\pi = 2$ (017)	~70 yrs	yes Δ	π cannot be approximated "unusually well" by fractions
6	Matrix multiplication $\omega \leq 9/4$ (107)	57 yrs	yes	a drop of 0.12 at once; the previous 12 years gave 0.0017
7	Unique Games Conjecture (102)	24 yrs	yes	optimality of many approximation algorithms becomes theorem
8	$L = RL = BPL$ (103)	~30 yrs	yes	randomness is useless for small-memory computation
9	The plane is not 5-colorable (158)	76 yrs (last lower-bound step 2018)	yes	chromatic number of the plane $\in \{6, 7\}$
10	Thompson's group F non-amenable (248)	47 yrs	yes	famous for "proofs" in both directions later withdrawn
11	Navier–Stokes blowup (separate Sept release)	Millennium Prize problem	yes	only the "with forcing" alternative; credit disputed
12	Relativistic Vlasov–Maxwell (362)	40 yrs	yes	large-data global smoothness for a plasma model (one species)
13	Mézard–Parisi formula (221)	25 yrs	yes	free-energy formula for sparse spin glasses (even-arity case)
14	Keakeya: 4D set + 3D maximal (074)	4D: decades; the 3D set case was only solved in 2025	no	every 4D Keakeya set has full dimension 4 (previously about 3.059); the same release also claims the 3D Fourier restriction conjecture (077)
15	Hodge-related withdrawals (032 etc.)	—	no	one sign error, three papers withdrawn: the risk of no formalization

FUNDAMENTALS · PRIMER

1. What do mathematicians actually do?

In one sentence: **they study the structure of abstract objects and establish truths about them by proof.** What separates mathematics from science and engineering is *how you know something is right*: mathematics by logical proof (once proven, true forever), science by experiment (can be overturned), engineering by "it works and measures well enough."

1.1 Four words

WORD	MEANING	EXAMPLE FROM THIS EVENT
conjecture	a statement widely believed but unproven	the Unique Games Conjecture (Khot 2002)
theorem	a statement that has been proven	the Goemans–Williamson 0.878 algorithm
counterexample	a concrete object showing that a "for all ..." conjecture is false	the 9-dimensional Borsuk counterexample; about 20% of this batch are counterexamples/disproofs
formalization	writing a proof as code a computer can check step by step (here Lean 4)	about 300/719 top-line results have Lean proofs

1.2 Two cultures of mathematics

In *The Two Cultures of Mathematics* (2000), Fields medalist Tim Gowers splits mathematicians into **problem-solvers** (who crack specific hard problems, like Erdős problems) and **theory-builders** (who invent new concepts and frameworks, like "group" or "manifold"). Thurston's *On Proof and Progress in Mathematics* (1994) goes further: the point of mathematics is not to produce theorems but to **advance human understanding**. These two essays are the key to this controversy — what AI has broken open at scale is the problem-solving culture, and that is what Tao means by "the end of Math 1.0."

1.3 Five questions for judging how big a result is

1. **How long was it open, and who tried?** A problem open for 50 years with repeated failures by top experts weighs very differently from one posed two years ago.
2. **Last step or giant leap?** Moving ω from 2.3713 to 2.3711 is the former; to 2.25 is the latter.
3. **How much does it unlock?** If UGC holds, hundreds of "if UGC then ..." papers become theorems at once.
4. **Is the method new?** New methods often matter more than the result, because they can be reused.
5. **Who chose the problem, and how many attempts failed?** The most overlooked question in this event; see §2.3 and §16.

2. What happened

DATE	EVENT
2025-07	An OpenAI experimental model and Google DeepMind's Gemini Deep Think reach gold-medal scores at IMO 2025 (35/42) — competition problems with known answers fall. [R]
2026-09-08	OpenAI announces that ~10,000 agents over 88 hours proved finite-time blowup "with forcing" for 3D Navier–Stokes (166 pages + Lean), sparking a credit dispute (§11). [P] [R]
late 2026-09	The Advisory Group on Mathematics and AI hosted by the Institute for Advanced Study (members include Terence Tao) publishes recommendations for releasing AI-generated mathematics. [R]
2026-10-06	openai/math released: 722 manuscripts / 372 result families, Apache-2.0, with a Lean library and 10 abridged reasoning summaries. [P]
2026-10-07	3 manuscripts withdrawn (one sign error, cascading), 14 repaired, 13 re-cited, 6 new formalizations → 719 manuscripts , "300/719 $\approx$ 42% of top-line results formalized." A statement from the Association for Human Mathematics (AHM) is reposted on Tao's blog; Scott Aaronson posts "The Mathocalypse". [P] [R]
2026-10-08	Tao's blog runs a guest post by Lozano-Robledo, What should we tell our students? [P]

2.1 How the results were produced


From the README [P]: "The vast majority of results were obtained with the same procedure using an unreleased internal OpenAI model. On average, each result used three hours of ChatGPT Pro thinking compute... the model was posed approximately 4,000 problems... requiring an appropriate level of significance led to the catalog." In short: **~4,000 open problems → filtered for significance → 372 families**. An OpenAI spokesperson told Scientific American that almost everything came from "a single prompt handed to a single AI agent" [R]; MIT's Andrew Sutherland responded that until the model can be replicated, "one-shotting" claims should be treated as unverified — "We should ask for receipts."

Exceptions: the README concedes that the zero-free region for ζ (§4) and the Hodge conjecture for CM abelian varieties (§15) were *not* produced by the fixed procedure, and that the $\text{Re } s > 11/12$ write-up was human-edited. What exactly differed (more compute? human guidance?) is not disclosed $\triangle$.

2.2 Reconciling the numbers

CLAIM	WHAT IT ACTUALLY MEANS
"722" vs "719"	722 at release; 719 after the three Oct 7 withdrawals. Early news used 722.
"372 families" vs "377"	The catalog is numbered 001–377 with 045/061/070/123/163 unused: 372 actual families.
"~42% formalized"	$300 \text{ top-line results} \div 719 \text{ manuscripts}$ — numerator and denominator in different units. By family, $242/372 \approx 65\%$ have a Lean scope document.
"162 fully + 73 partially formalized"	Reported by some outlets; source not found △ .

372 result families by field (orange = this deck covers at least one result in depth)



Source: the 17 subject sections of [openai/math overview.pdf](#), counted entry by entry (total 372). Catalog numbers 045/061/070/123/163 are unused; no reason given.

2.3 The one thing to keep in mind: selection effects

Epoch AI's FrontierMath Erdős (2026-09) took 68 Erdős problems **pre-registered** as open and required Lean proofs; GPT-6 Astra solved 2 ($\approx 3\%$), all other models 0. OpenAI's release is **post-hoc selection**: of $\sim 4,000$ problems it kept the "significant" ones and published no failure list. A rough hit rate is $\leq 9\%$ (by family) — but problem difficulty is not comparable.

Both numbers are right; they answer different questions: the first asks "given an arbitrary open problem, can it solve it?", the second "given 4,000, how many good results can it pick out?"

3. What "has a Lean proof" actually guarantees

Lean is a proof assistant: you write a theorem and its proof as code, and a small **kernel** checks every inference step. If the kernel has no bugs and only accepted axioms are used, passing the check means "this Lean statement is true." The difference from human refereeing is like the difference between a compiler with a type system and code review — the latter misses things; the former doesn't miss *what it checks*.

What "has a Lean proof" guarantees: the trust stack, bottom up

① **Lean kernel checks every proof step**

this repo: sorry = 0, added axioms = 0, native_decide = 0

strong

② **Only the 3 standard axioms**

all 416 Comparator configs allow only propext / Quot.sound / Classical.choice

strong

③ **Lean statement = what the paper claims?**

challenge files written by OpenAI (agents); a human must read the definitions

human

④ **Are definitions in deps and patches right?**

24 patches (largest ~2 MB) + 11 lana-agents/* dependencies of unknown origin

unaudited

⑤ **The ~58% with no formalization**

no machine check at all; the 3 Hodge-related papers withdrawn on 10-07 were here

none

①② are guaranteed by the machine; ③④⑤ still need humans. Audit data: source grep of `lean/OAI/` (122,458 .lean files, ~26M lines) plus Comparator config counts (2026-10-08, commit fd4aeeb); not compiled.

3.1 Our static audit of the Lean library [P]

CHECK	RESULT	MEANING
sorry (unfinished proof)	0	no "fill in later" steps
axiom declarations	0 (all hits are in comments)	no smuggled-in axioms
native_decide / implemented_by / extern	0	no "trust the compiler" backdoors
decide +kernel	≈1.31M lines, 5,848 files	large finite computations done by the kernel — still trusted, but the proofs are enormous
Comparator configs	416, all allowing only the three standard axioms; only 2 enable the second independent kernel nanoda	see below
lean/patches/	24 patches to third-party dependencies (not the Lean kernel, not Mathlib), largest ~2 MB	patches change definitions and proofs in dependencies △
dependencies	30, of which 11 are <code>lana-agents/*</code> repos of unknown origin	an open point in the trust chain △
formalization.yaml	<code>review: status: unchecked</code>	OpenAI's own label: not reviewed

Scope: lean/OAI/ has 122,458 .lean files, ~26M lines; we only grepped the source and counted configs, without compiling (a full build needs many GB of memory and a Linux sandbox). Independent re-checks: Dave Goldblatt ran the ζ 7/8 proof through Comparator and the nanoda kernel (~63 minutes) [repo]; a third party reviewed the statement of the 9/4 matrix-multiplication theorem [repo].

3.2 Comparator: what it guards against, and what it doesn't

The Lean FRO's Comparator takes a "challenge file" (only the theorem statement, with `sorry` as the proof, importing only trusted libraries such as Mathlib) and a "solution module": the theorem proven in the solution must be **the same type** as the challenge statement, using only allowed axioms. It guards against swapping the statement, changing definitions or adding axioms inside the proof. It does **not** guard against the challenge statement itself failing to say what the paper claims — and the challenge files were written by OpenAI (its agents). That is why the "verification status" of each section below focuses on **reading the statement line by line**.

A LIVE COUNTEREXAMPLE

A footnote in OpenAI's plane-coloring paper mentions a May 2026 manuscript that claimed the chromatic number of the plane is 7 with a "formalized theorem" — but that theorem **took a false density bound as a hypothesis**. The formalization was fine; the proven statement was the problem. [P, as reported in the OpenAI paper; original not read △] This is the

same kind of problem as specification gaming in AI: the system perfectly satisfies the written spec, but the spec is not what you wanted.

NUMBER THEORY · FAMILY 003

4. The quasi-Riemann hypothesis: a ceiling on the noise in the primes

Number theory · OpenAI family 003. All you need going in: you can work with powers and logarithms, you know the summation sign Σ , and you know what a function is. Every new concept is explained before it is used. Every number in this chapter was either computed by script or has a source. Tags: [P] = a primary source we read directly (paper, repo file, Lean source), [R] = secondary reporting, $\triangle$ = unverified or doubtful. The **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. Estimated reading time: 60–90 minutes. You can read it in three sittings: §4.1–4.3 (primes and the error), §4.4–4.6 (the ζ function and its zeros), §4.7–4.10 (OpenAI's result and what it means).

4.0 The bottom line: the one thing this chapter explains

Prime numbers (2, 3, 5, 7, 11, ...) look randomly scattered, but taken as a whole they are very regular. Mathematicians can write down a formula that predicts, accurately, "roughly how many primes there are up to x ." The truly hard question is: **how wrong can that prediction be, at most?**

This chapter explains one causal chain:

How big the prediction error is $\Leftarrow$ is determined by a set of "waves" $\Leftarrow$ each wave corresponds to a "zero" of the ζ function $\Leftarrow$ the farther a zero sits from a certain line, the bigger the error.

The Riemann hypothesis says that all the zeros line up neatly on one line, so the error is as small as it can be. Nobody can prove it.

The **quasi-Riemann hypothesis** is a much weaker version. It does not require the zeros to lie on that line; it only requires that they **not sit too far to the right**. Nobody has proved even this weak version in 125 years. OpenAI claims to have proved it.

After reading this chapter, you should be able to explain in your own words:

- what exactly the relationship is between "zeros" and "the error in counting primes";

- why there is a qualitative difference between "a zero-free region that gets narrower and narrower" and "a zero-free region of fixed width";
- what OpenAI's claimed result, if true, concretely means, and what it does not mean.

4.1 Primes: count them, look for a pattern

4.1.1 What primes are and why they matter

A prime is an integer greater than 1 that is divisible only by 1 and itself: 2, 3, 5, 7, 11, 13, 17, 19, 23, 29, ...

They matter because **every integer can be split into a product of primes in exactly one way** (the fundamental theorem of arithmetic):

- $12 = 2 \times 2 \times 3$
- $360 = 2 \times 2 \times 2 \times 3 \times 3 \times 5$

So primes are the "atoms" of the integers. The HTTPS encryption (RSA) you use every day rests on one fact: multiplying two large primes is easy, but splitting the product back apart is extremely hard.

4.1.2 Counting

Write $\pi(x)$ for "the number of primes that are at most x ." This π is just a name. It has nothing to do with the circle constant; it is a historical notation.

x	$\pi(x)$ (actual count)
10	4 (2, 3, 5, 7)
100	25
1,000	168
1,000,000	78,498
1,000,000,000	50,847,534

Primes get sparser as you go further out. Below 100, one number in four is prime. Near 1 million, only about 72 in every 1000 numbers are prime (that is the estimate from the $1/\ln x$ rule in the next section; counted by script, the 1000 numbers from 999,001 to 1,000,000 actually contain 65 primes).

4.1.3 The first pattern: the density is about $1/\ln x$

As a teenager, Gauss stared at prime tables and noticed: **near x , about one number in every $\ln x$ is prime.**

Here $\ln$ is the natural logarithm, with base $e \approx 2.718$. $\ln(1,000,000) \approx 13.8$, which means that near 1 million about one number in 14 is prime. The script count above was 65 per 1000, or about one in 15. Very close.

[Check 1] $\ln(1000) \approx 6.9$. By this rule, about how many primes are there per 1000 numbers near 1000?

4.1.4 The prime number theorem

Integrate the "density $1/\ln t$ " from 2 to x and you get a prediction for the number of primes up to x :

$$\pi(x) \approx \text{Li}(x) = \int_2^x \frac{dt}{\ln t}$$

$\text{Li}(x)$ is called the "logarithmic integral." If you are not comfortable with integrals, that is fine: think of it as "add up the density $1/\ln t$ at every position."

x	actual $\pi(x)$	rough prediction $x/\ln x$	precise prediction $\text{Li}(x)$	error of Li
1,000	168	145	177	+9
1,000,000	78,498	72,382	78,627	+129
1,000,000,000	50,847,534	48,254,942	50,849,235	+1,701

($\text{Li}(10^9)$ uses the standard value from the literature; the rest were computed by script.)

The **prime number theorem** (proved independently by Hadamard and de la Vallée Poussin in 1896) says that the **ratio** of $\pi(x)$ to $\text{Li}(x)$ tends to 1. You can see this in the table: the absolute error is growing ($9 \rightarrow 129 \rightarrow 1,701$), but relative to the total it becomes more and more negligible (at 10^9 it is off by only 0.3 parts in 10,000).

4.2 The real question: how big is the error?

4.2.1 Two ways of saying "the error is getting smaller," with very different strength

The prime number theorem only says "error / $x \rightarrow 0$," that is, the relative error goes to 0. It does not say **how fast** it goes to 0.

An intuitive way to measure this: roughly what power of x is the error?

- If the error $\approx x^{0.5} = \sqrt{x}$: at 10^9 , $\sqrt{x} \approx 31,623$. The actual error in the table above, 1,701, is even smaller than that.
- If the error $\approx x^{0.875} = x^{(7/8)}$: at 10^9 this is about 75 million, far bigger than the actual error.

Note: the **actual error** and the **error bound you can prove** are two different things. Numerically, the actual error is very small (roughly on the order of $\sqrt{x}$). But mathematicians need a **proof**: "for all x , the error is definitely no more than such-and-such." Everything in this chapter is about the **bound you can prove**.

4.2.2 "Power saving": why $x^{(7/8)}$ is a qualitative change

Since 1958, the best provable bound has roughly had the form:

$$\text{error} \leq x \cdot \exp(-c \cdot (\ln x)^{0.6}) \quad (c \text{ is some small constant})$$

Compare this with $x^{(7/8)}$ by looking at how "error / x " (the relative error) changes with x . To compare the trends, take $c = 1$ below (the real c is smaller, so the old bound is actually even worse):

x	relative error of the old bound $\exp(-(\ln x)^{0.6})$	relative error of the new bound $x^{(-1/8)}$
10^{100}	$10^{(-11.4)}$	$10^{(-12.5)}$
10^{1000}	$10^{(-45.2)}$	$10^{(-125)}$
10^{10000}	$10^{(-180)}$	$10^{(-1250)}$

(Computed by script.)

The old bound also shrinks, but **more slowly than any "negative power of x ."** The new bound is " x to the power $-1/8$," and the gap widens the further out you go. Mathematicians call this a **power saving**: the error is smaller than x by a fixed power.

Why this matters: many harder problems in number theory plug the error term of the prime number theorem into further calculations. If the error is only "a little" smaller than x (like the old bound), it often gets swamped by other error terms once you plug it in. If it is smaller by a fixed power, you have room to spare for all kinds of consequences.

4.3 The key turn: the error is made of "waves"

This was Riemann's stroke of genius in 1859, and it is the single most important step in the chapter.

4.3.1 First, an analogy with sound

A sound (say, a chord) can be broken down into a sum of pure tones (sine waves). Each pure tone has two properties:

- **frequency:** how fast it vibrates

- **amplitude:** how loud it is

This is called **Fourier decomposition**. The EEG spectra you saw in neuroscience (alpha waves at 8–12 Hz, beta waves, ...) are the same idea: break a complicated signal into waves of different frequencies.

4.3.2 Riemann's discovery

Riemann discovered that the error in counting primes can be broken down the same way:

$$(\text{weighted count of primes}) = (\text{smooth main term}) - \Sigma(\text{individual "waves"}) - (\text{some small corrections})$$

(To keep the formula clean, mathematicians usually use a weighted version of the prime count, $\psi(x)$, in which each prime p gets weight $\ln p$. It is essentially the same thing as $\pi(x)$.)

The main term is smooth, roughly just x . All of the error comes from those "waves," and:

- **Each wave corresponds to one zero of the ζ function** (§4.5 explains what a zero is).
- A zero is written in the form $\rho = \beta + i\gamma$, with two coordinates β and γ .
- **γ determines the wave's frequency** (how fast it oscillates).
- **β determines how fast the wave's amplitude grows with x :** the size of this wave is about x^β .

So:

- If every zero has β equal to $1/2$, each wave has size about $\sqrt{x}$, and the error is about $\sqrt{x}$. This is exactly the **Riemann hypothesis**.
- If some zero has β close to 1, its wave has size close to x , almost as big as the main term, and the error will be huge.
- If you can prove that **every zero has β at most θ** , then the error is at most about x^θ .

In one sentence: the β coordinate of the zeros (how far "to the right" they sit) directly determines how big the error in counting primes is. That is why number theorists care so much about "where the zeros are."

4.3.3 Why complex numbers show up

To describe these "waves" we need complex numbers. The next section explains them first.

4.4 Complex numbers: the 5-minute version

4.4.1 A complex number is a point in the plane

Real numbers sit on a line. **Complex numbers are points in a plane**: the horizontal coordinate is called the real part, the vertical coordinate the imaginary part.

$s = \sigma + i \cdot t$ σ is the horizontal coordinate (real part, written $\text{Re } s$), t is the vertical coordinate (imaginary part, written $\text{Im } s$)

i is a symbol satisfying $i^2 = -1$. Don't worry about whether it "exists"; treat it as a marker for "one step in the vertical direction" and that is enough.

4.4.2 Multiplying by a complex number = scaling + rotation

The most useful property of complex numbers: **multiplying by a complex number rotates a point around the origin and scales it**. For example, multiplying by i rotates 90 degrees counterclockwise: $1 \rightarrow i \rightarrow -1 \rightarrow -i \rightarrow 1$, one full turn.

4.4.3 x to a complex power = "size part" $\times$ "rotation part"

For a positive real number x and a complex number $\beta + i\gamma$:

$$x^{(\beta+i\gamma)} = x^\beta \times (\cos(\gamma \cdot \ln x) + i \cdot \sin(\gamma \cdot \ln x))$$

- x^β is the size part: it grows with x , and how fast is set by β .
- $\cos(\gamma \cdot \ln x) + i \cdot \sin(\gamma \cdot \ln x)$ is the rotation part: its length is always 1; it just goes around the unit circle. How fast it turns is set by γ .

These are the "waves" from §4.3: **β controls the amplitude (x^β), γ controls the frequency (oscillation on the scale of $\ln x$).**

4.5 The ζ function and its zeros

4.5.1 Definition of the ζ function

$$\zeta(s) = 1/1^s + 1/2^s + 1/3^s + 1/4^s + \dots$$

At $s = 2$: $\zeta(2) = 1 + 1/4 + 1/9 + 1/16 + \dots = \pi^2/6 \approx 1.6449$ (computed by Euler in 1734; here π is the circle constant).

At $s = 1$: $\zeta(1) = 1 + 1/2 + 1/3 + 1/4 + \dots$, the "harmonic series," which **diverges to infinity**.
Computed by script: the sum of the first 1000 terms ≈ 7.49 , the first 1 million terms ≈ 14.39 .
It grows very slowly but never stops.

4.5.2 The Euler product: how ζ "knows" about primes

Euler discovered:

$$\zeta(s) = 1/(1-2^{-s}) \times 1/(1-3^{-s}) \times 1/(1-5^{-s}) \times 1/(1-7^{-s}) \times \dots \quad (\text{product over all primes})$$

Why it holds: each factor $1/(1-p^{-s})$ expands into $1 + 1/p^s + 1/p^{2s} + \dots$ (a geometric series). Multiply these expansions together over all primes, and every term has the form $1/(p_1^a \cdot p_2^b \cdot \dots)^s$. Because every integer has a **unique** prime factorization, each $1/n^s$ appears exactly once, and the product is exactly $\zeta(s)$.

By hand for $s = 2$ (checked by script):

- Multiplying only the three factors for $p = 2, 3, 5$: $(4/3)(9/8)(25/24) = 1.5625$
- Multiplying the factors for all primes below 100: 1.6419
- True value $\pi^2/6 = 1.6449$

So **the ζ function ties "all integers" to "all primes"**, and this is the bridge for the whole story.

A beautiful consequence (Euler's proof that "there are infinitely many primes"): at $s = 1$ the left side is the harmonic series, which diverges to infinity. If there were only finitely many primes, the right side would be a product of finitely many finite numbers and could not be infinite. So there must be infinitely many primes.

4.5.3 "Analytic continuation": extending the definition to the whole plane

The sum above only converges (adds up to a finite number) when the real part of s is > 1 . But Riemann needed to talk about $\zeta(s)$ on the whole complex plane.

An analogy for how to extend it:

$$1 + x + x^2 + x^3 + \dots = 1/(1-x)$$

The left side converges only when $|x| < 1$ (for example, at $x = 2$ the left side is $1+2+4+8+\dots$, which diverges), but the right side $1/(1-x)$ makes sense everywhere except $x = 1$ (at $x = 2$ it equals -1). We can **use the right side to define** something that originally only made sense for $|x| < 1$. This is called **analytic continuation**.

One can prove mathematically that such a continuation, if it exists, is unique. Riemann used a similar but more elaborate technique to continue $\zeta(s)$ to the whole complex plane (except the single point $s = 1$).

4.5.4 Zeros

A **zero** is a complex number s that makes $\zeta(s) = 0$.

- **Trivial zeros:** $s = -2, -4, -6, \dots$, which come from a simple factor in the continuation formula. Nothing mysterious.
- **Non-trivial zeros:** all the other zeros. They all lie in the vertical strip $0 < \operatorname{Re} s < 1$, which is called the **critical strip**.

The first three non-trivial zeros (in the upper half-plane) are approximately:

$$\frac{1}{2} + 14.13i, \quad \frac{1}{2} + 21.02i, \quad \frac{1}{2} + 25.01i$$

Their real parts are all exactly $1/2$.

4.5.5 The Riemann hypothesis

Riemann hypothesis (1859): every non-trivial zero has real part equal to $1/2$.

Numerical checks have examined **more than ten trillion zeros**, all on the line $\operatorname{Re} s = 1/2$ (Gourdon in 2004 verified the first 10^{13} zeros $\triangle$ defer to the original paper for the exact figure). But no amount of checking is a proof. It is one of the seven Clay Millennium Prize Problems, with a \$1 million prize.

Connecting §4.3 to here: **Riemann hypothesis $\Leftrightarrow$ every wave has amplitude $\sqrt{x} \Leftrightarrow$ the error in counting primes is about $\sqrt{x}$.**

4.6 Zero-free regions: what can be proved when the Riemann hypothesis can't

Since nobody can prove "all zeros are on the line $1/2$," mathematicians take a step back and ask: **can we at least prove where the zeros are not?** A region proved to contain no zeros is called a **zero-free region**.

4.6.1 Step one: no zeros on the line $\operatorname{Re} s = 1$ (1896)

The proof of the prime number theorem essentially amounts to proving that **there are no zeros on the vertical line $\operatorname{Re} s = 1$** .

The key trick is a trigonometric identity:

$$3 + 4\cos \theta + \cos 2\theta \geq 0 \quad (\text{holds for all } \theta)$$

[Check 2] Use $\cos 2\theta = 2\cos^2\theta - 1$ to prove that $3 + 4\cos \theta + \cos 2\theta = 2(1 + \cos \theta)^2$, so it is always ≥ 0 . (The script swept 6,284 points on $[0, 2\pi]$; the minimum was 1.4×10^{-14} , i.e. 0 plus

rounding error.)

Combining this identity with the Euler product, one can show that if $1 + iy$ were a zero, it would lead to a contradiction.

4.6.2 Step two: a strip to the left of $\text{Re } s = 1$ that "keeps getting narrower" (1899, 1958)

de la Vallée Poussin (1899) pushed the result a little further: not only are there no zeros on $\text{Re } s = 1$, there are none in **a narrow strip to its left** either:

zero-free region: $\text{Re } s > 1 - c / \ln(t)$ (t is the height, i.e. the imaginary part)

Note the $\ln t$ in the denominator: **the bigger the height t , the narrower this strip.**

- At $t = 100$, the width is about $c / 4.6$
- At $t = 10^{10}$, the width is about $c / 23$
- As $t \rightarrow \infty$, the width $\rightarrow 0$

Vinogradov–Korobov (1958) widened the strip a bit (the denominator changes from $\ln t$ to $(\ln t)^{2/3}$ times an even smaller factor), but **it still shrinks toward 0 with height**. This was the best result for the following 60-plus years.

4.6.3 Why "keeps getting narrower" and "fixed width" are qualitatively different

Back to §4.3: the error is the sum of the waves from all zeros, and **the rightmost zero** sets the size of the error.

- If the zero-free region keeps getting narrower: in principle, higher and higher up there could be zeros whose β gets closer and closer to 1. You cannot rule them out, so the error bound can only be "a little" smaller than x . That is the $\exp(-c(\ln x)^{0.6})$ from §4.2.
- If there is a zero-free region of **fixed width**, say no zeros at all in $\text{Re } s > 7/8$: then every zero has $\beta \leq 7/8$, every wave is $\leq x^{7/8}$, and the error bound becomes of order $x^{7/8}$. That is the **power saving** from §4.2.


The **quasi-Riemann hypothesis** is exactly such a statement:

There is some $\theta < 1$ such that ζ has no zeros anywhere in the region $\text{Re } s > \theta$.

It is much weaker than the Riemann hypothesis ($\theta = 1/2$), but before 2026 **nobody could prove it for any $\theta < 1$.**

The figure below puts §4.5 and this section on one picture of the complex plane: the known zeros, the classical zero-free region that narrows as you go up, and the fixed-width zero-

free region OpenAI claims.



Blue dots: the first three non-trivial zeros of ζ (all on $\text{Re } s = \frac{1}{2}$; the Riemann hypothesis says every one is). Grey band: schematic de la Vallée Poussin-type zero-free region (its width shrinks to 0 as the height t grows; curve shape illustrative, constant 0.1). Red region $\text{Re } s > 7/8$: the fixed-width zero-free region claimed by OpenAI family 003. Nobody had previously proven any fixed-width zero-free strip.

4.6.4 Why the classical methods can't do it

The " $3 + 4\cos \theta + \cos 2\theta \geq 0$ "-type argument from §4.6.1 starts from properties of the Euler product in $\text{Re } s > 1$ and "seeps" a little way to the left. But this seeping power decays with the height t , so all you can get is a strip that keeps getting narrower. To get a fixed width, you need **a fundamentally different mechanism**.

Another line of attack is called "zero-density estimates" (for example, the important breakthrough of Guth–Maynard 2024, arXiv:2405.20552). These can prove that "there are **few** zeros with $\text{Re } s > \sigma$," but not that "there are **none at all**." The introduction of OpenAI's paper makes the same distinction explicitly [P]. And for the error bound, even a single zero with β close to 1 is enough to make the error large.

4.7 Dirichlet L-functions and Siegel zeros

OpenAI's result is not only about ζ ; it also covers all **Dirichlet L-functions**.

4.7.1 Primes in arithmetic progressions

Apart from 2, all primes are odd, so each has the form $4k+1$ (5, 13, 17, ...) or $4k+3$ (3, 7, 11, 19, ...).

Computed by script:

- Below 100: 11 of type $4k+1$, 13 of type $4k+3$
- Below 1000: 80 of type $4k+1$, 87 of type $4k+3$

The two types are roughly half and half (type $4k+3$ is usually slightly ahead, but not always: at $x = 26,861$ type $4k+1$ takes the lead for the first time; this phenomenon is called Chebyshev's bias). In 1837 Dirichlet proved that for any modulus q and any a coprime to q , there are infinitely many primes of the form $qk + a$, and that in the long run the classes split evenly.

4.7.2 Dirichlet L-functions

To study each class separately, Dirichlet built variants of ζ :

$$L(s, \chi) = \chi(1)/1^s + \chi(2)/2^s + \chi(3)/3^s + \dots$$

χ (pronounced "chi") is called a "character." It is a function that maps integers to certain complex numbers, and its job is to "pick out" and "tell apart" the different residue classes. For example, the non-trivial character mod 4: $\chi(\text{odd numbers } 4k+1) = 1$, $\chi(\text{odd numbers } 4k+3) = -1$, $\chi(\text{even numbers}) = 0$.

Just like ζ , the zeros of L-functions control the error in the distribution of primes across the classes mod q .

4.7.3 Siegel zeros: a ghost that has haunted the field for 90 years

For certain special L-functions, classical methods cannot rule out one possibility: **that it has a real zero very, very close to 1**. Such a hypothetical zero is called a **Siegel zero** (or Landau–Siegel zero).

If one exists, the consequences are strange: primes for the corresponding modulus would be heavily skewed toward certain residue classes. Although people widely believe it does not exist, nobody can rule it out.

Worse, Siegel's related theorem from 1935 is "ineffective": it proves that a certain constant **exists** but tells you nothing at all about **how big** it is. As a result, a whole set of problems, such as "how large is the class number of an imaginary quadratic field at least," can only get "there exists a bound," not an actual number.

If the quasi-Riemann hypothesis holds for all Dirichlet L-functions, Siegel zeros are automatically ruled out: there are no zeros anywhere in $\text{Re } s > 7/8$, so of course there is no zero close to 1.

4.8 What OpenAI claims

4.8.1 The main result

Abstract of the paper *The Quasi-Riemann Hypothesis: A Zero-Free Half-Plane $\Re s > 7/8$* (2026-09-30, 199 pages) [P]:

"We prove that all finite-order Hecke L-functions over $\mathbb{Q}(\sqrt{-3})$ and all Dirichlet L-functions are zero-free in the half-plane $\Re s > 7/8$... In particular, the Riemann zeta function is zero-free in this half-plane, proving the quasi-Riemann hypothesis."

In plain terms: **ζ and all Dirichlet L-functions have no zeros anywhere in the half-plane $\Re s > 7/8$.**

The paper itself makes a point of stressing: "The Riemann hypothesis remains open." The Riemann hypothesis says the zeros are all at $1/2$; this only proves they are all $\leq 7/8$. Between $7/8$ and $1/2$ there is still a large unknown territory.

There are two more papers in the same family (family 003): one proves an $11/12$ version by a different method (49 pages; the README says its writing was edited by humans), and a 9-page paper separately proves that Siegel zeros are uniformly ruled out.

4.8.2 Intuition for the method (we only read the introduction; this is a simplified version)

- 1. Switch number systems.** Instead of working in the ordinary integers, work in the "Eisenstein integers" $\mathbb{Z}[\omega]$, where ω is a cube root of unity: $\omega^3 = 1$, $\omega \neq 1$. You can picture them as a triangular lattice of points in the plane (the ordinary integers are points on a line, the Gaussian integers a square lattice, the Eisenstein integers a triangular lattice). The field corresponding to this number system is written $\mathbb{Q}(\sqrt{-3})$.
- 2. Build one quantity and compute it two ways.** The authors build a sum from the Fourier coefficients of a "cubic theta series" (the paper calls it a "completed sum"), which can be computed along two independent routes: one uses a "reflection formula," the other uses "Poisson summation." The result of the second route contains terms like $1/L(s, \eta)$, i.e. the reciprocal of an L-function.
- 3. Contradiction.** If some L-function had a zero in $\Re s > 7/8$, $1/L$ would blow up there, and the two computations could not agree. Think of it as weighing the same object on two independent scales: if an assumption would make the two scales read differently, the assumption is wrong.
- 4. Transfer to ordinary Dirichlet L-functions.** Using a factorization formula $L_F(s, \chi \circ N) = L(s, \chi) \cdot L(s, \chi\chi_{-3})$, the conclusion over $\mathbb{Q}(\sqrt{-3})$ can be "split" back into L-functions over the ordinary integers, including ζ .

Technically there are two stages: first a simpler version gives 11/12, then extra correction terms push it to 7/8. The tools used include the sextic large sieve, Heath-Brown's estimates for cubic sums, and Dunn–Radziwiłł's 2024 Annals result on the Patterson conjecture.

4.9 How far it has been verified

4.9.1 The Lean statement: one line, and you can read it yourself

The "challenge statement" file in the repo,

`lean/ComparatorChallenges/QuasiRiemannHypothesis.lean` (it imports only Mathlib) [P]:

```
theorem riemannZeta_ne_zero_of_seven_eighths_lt_re {s : ℂ} (hs : (7/8 : ℝ) < s.re) :  
  riemannZeta s ≠ 0
```

Word by word:

- `theorem ...`: the name of the theorem, which literally means "ζ is nonzero when the real part is greater than 7/8."
- `{s : ℂ}`: s is a complex number (ℂ is the set of complex numbers).
- `(hs : (7/8 : ℝ) < s.re)`: the hypothesis hs is "the real part of s > 7/8."
- `riemannZeta s ≠ 0`: the conclusion is "ζ(s) is not equal to 0."

So this one line says: **for any complex number s, if its real part is greater than 7/8, then ζ(s) ≠ 0**. That is exactly the quasi-Riemann hypothesis.

The `riemannZeta` here is the **standard definition** that has long been maintained in Mathlib (Lean's community mathematics library), not something OpenAI wrote itself. This matters: if the definition were home-made, it could quietly "change the subject."

4.9.2 The proof itself

- The proof is about 2,926 Lean files and 487,000 lines [P].
- We grepped it: no `sorry` ("leave this step blank for now"), no added axioms, no `native_decide` (a backdoor that trusts the compiler) [P].
- Independent re-check: Dave Goldblatt used the Lean FRO's Comparator tool, plus an independently implemented checking kernel (nanoda), and got the proof to go through, in about 63 minutes; `#print axioms` lists only the three standard axioms (`propext`, `Quot.sound`, `Classical.choice`) [P]. Re-check repo: davegoldblatt/openai-zeta-proof-check.

- **⚠ Remaining risk:** the dependency patch `lean/patches/PrimeNumberTheoremAnd-lean4341.patch` is about 40,000 lines, and in it two upstream `sorry`s were removed and replaced with proofs. In other words, part of the so-called "dependency library" is OpenAI's own code. (PrimeNumberTheoremAnd is the prime-number-theorem formalization project that Kontorovich, quoted below, started.) The Lean kernel checks these proofs, but if definitions were changed, that could affect what the statement means. Nobody has reviewed this line by line yet.
- **The formalization covers only the zero-free regions themselves.** Comparator checks four statements: the $7/8$ zero-free region for ζ , for Dirichlet L-functions and for Hecke L-functions over $\mathbb{Q}(\sqrt{-3})$, plus a uniform gap for Siegel zeros. The scope note `lean/docs/003.md` says, verbatim, "The paper's later applications are not included" [P]. So the least-quadratic-non-residue bound and the square-root algorithm in §4.10.2 have only the paper's human-readable proof, not a Lean proof.

4.9.3 Human experts

- Alex Kontorovich (Rutgers; he himself started PrimeNumberTheoremAnd, the Lean formalization project for the prime number theorem): "Quasi-RH?!?!?! Are you kidding me? If a human did this, it would be an instant Fields Medal..." [R] (X post; we quote it via officechai, as the original post could not be fetched)
- Levent Alpöge (Anthropic): "Big, big, big, big props for quasiriemann and no Siegel zeroes." [R] (X post, quoted via Latent Space)
- **⚠ As of 2026-10-08, no human expert has publicly said they have read and verified the 199-page paper.** Both comments above are reactions, not referee reports.

4.9.4 How to think about "proved by machine, but not understood by people"

If you believe two things, (1) Mathlib's definition of ζ is correct and (2) the Lean kernel has no bugs, then this theorem has been proved in the machine sense, and it is an order of magnitude more trustworthy than a paper nobody has refereed. The remaining risk sits in the build chain: the patches and dependency libraries all come from OpenAI (see 9.2).

But the mathematical community usually also asks for **understanding**: why is it true? Does the method generalize? This is what Tao means by "the end of Math 1.0": a conclusion can be confirmed by machine first, and understanding may only catch up much later.

4.10 If true, what it means (and what it does not mean)

4.10.1 For mathematics

1. **For the first time, the prime number theorem has an error bound with a power saving:**
 $\pi(x) - \text{Li}(x) = O(x^{7/8} \cdot (\ln x)^2)$. This is one of the biggest single steps forward in the

theory of the distribution of primes in 125 years.

2. **Siegel zeros are ruled out:** an obstacle that since 1935 has blocked a set of "can this constant be computed?" problems disappears, and a large batch of constants that "exist but can't be computed" become computable, e.g. effective lower bounds for class numbers of imaginary quadratic fields (§4.7.3).
3. **Pushing from 1 to 7/8 is a qualitative change, but it is still far from 1/2.** The Riemann hypothesis remains open.

4.10.2 For computer science: square roots mod p

One of the paper's corollaries is an upper bound on the "least quadratic non-residue." First, what that means.

Quadratic residue: a number that "is the square of something" mod p . Take $p = 7$ as an example:

- $1^2 = 1$, $2^2 = 4$, $3^2 = 9 \equiv 2$, $4^2 = 16 \equiv 2$, $5^2 = 25 \equiv 4$, $6^2 = 36 \equiv 1 \pmod{7}$
- So the squares mod 7 are $\{1, 2, 4\}$; these are called quadratic residues. The rest, $\{3, 5, 6\}$, are called **quadratic non-residues**.
- The **least quadratic non-residue** is 3.

[Check 3] Which numbers are quadratic residues mod 23? What is the least quadratic non-residue? (Answer at the end of the chapter.)

Why it matters: to take a square root mod p (given a , find x with $x^2 \equiv a$), there is a classical algorithm (Tonelli–Shanks) that first needs to find **any one** quadratic non-residue. Pick a number at random and there is a one-in-two chance it is a non-residue, so the randomized algorithm is fast. But a **deterministic** algorithm (no dice) has to guarantee that "if you try numbers upward from 2, you will hit a non-residue quickly," and for that you need to know how small the least non-residue is.

The paper's Corollary 1.2 derives: least quadratic non-residue $\leq C \cdot (\ln p)^{32}$ (C is an absolute constant; the paper says 32 is just one workable value and was not optimized) [P]. That is, trying upward from 2, you are guaranteed to find one within on the order of $(\ln p)^{32}$ tries, so taking square roots gets an **unconditional** deterministic polynomial-time algorithm. Previously this conclusion required assuming the "generalized Riemann hypothesis."

The same fixed-width zero-free region also makes the **Miller primality test** (Miller 1976, which likewise relied on the generalized Riemann hypothesis) an unconditional deterministic polynomial-time algorithm. This one is stated in the introduction of the companion 11/12

paper in the same family [P]; Corollary 1.2 of the 7/8 paper itself only states the non-residue bound and square roots. Since 7/8 is stronger than 11/12, the conclusion carries over.

A reminder (see 9.2): none of these corollaries is inside the Lean formalization.

No impact on practical cryptography:

- In practice people have always used the randomized algorithm, which is already fast.
- A deterministic polynomial-time algorithm for primality testing has existed since 2002 (AKS).
- The security of RSA and elliptic-curve cryptography rests on "factoring large numbers is hard" and "discrete logarithms are hard," which have nothing to do with this result.

4.10.3 For AI

This is one of the strongest cases of "AI gets a result on an open problem that humans agree is hard-core, and the result is machine-verified." It also creates a new situation: **a 199-page paper + 480,000 lines of Lean**, where trust rests mainly on "are the kernel and the definitions correct," not on "has anyone understood it." Human reviewing bandwidth has become the new bottleneck.

4.11 Self-test

Try to answer these after reading. If you can't, go back and reread the corresponding section.

1. In one sentence: why does the "real part" of the zeros affect the error in counting primes? (§4.3)
2. What is the fundamental difference in shape between de la Vallée Poussin's zero-free region and the quasi-Riemann hypothesis's zero-free region? Why does that difference matter? (§4.6.2–4.6.3)
3. Why can't zero-density estimates, however good, replace a zero-free region? (§4.6.4)
4. What is a Siegel zero? Why does the quasi-Riemann hypothesis rule it out? (§4.7.3)
5. In the Lean statement, what does the hypothesis `hs` in front of `riemannZeta s ≠ 0` say? (§4.9.1)
6. Why does this result not affect the security of RSA? (§4.10.2)

Appendix: answers to the checks

- **Check 1:** about $1000/6.9 \approx 145$ primes per 1000 numbers. In reality there are 168 primes below 1000, so the order of magnitude matches.

- **Check 2:** $\cos 2\theta = 2\cos^2\theta - 1$; substituting gives $3 + 4\cos \theta + 2\cos^2\theta - 1 = 2\cos^2\theta + 4\cos \theta + 2 = 2(\cos \theta + 1)^2 \geq 0$.
- **Check 3:** the quadratic residues mod 23 are {1, 2, 3, 4, 6, 8, 9, 12, 13, 16, 18} (11 of them, exactly half of 22), and the least quadratic non-residue is **5** (checked by script).

Appendix: Glossary

Term	One-line explanation
Prime number theorem	The number of primes up to x is about $\text{Li}(x)$; the ratio tends to 1
Power saving	The error is smaller than x by a fixed power, e.g. $x^{7/8}$
Complex number	A point in the plane; multiplication = scaling + rotation
ζ function	$1/1^s + 1/2^s + \dots$, connected to the primes through the Euler product
Analytic continuation	Uniquely extending a function defined only on part of a region to a larger region
Zero	A point where the function's value is 0
Critical strip	The vertical strip $0 < \text{Re } s < 1$, where all non-trivial zeros lie
Riemann hypothesis	Every non-trivial zero has real part $1/2$
Zero-free region	A region proved to contain no zeros
Quasi-Riemann hypothesis	There is a fixed $\theta < 1$ such that there are no zeros where $\text{Re } s > \theta$
Dirichlet L-function	A variant of ζ used to study primes in arithmetic progressions
Siegel zero	A hypothetical real zero of certain L-functions, very close to 1
Quadratic residue	A number that is the square of something mod p

Appendix: Further reading

- **3Blue1Brown**, *Visualizing the Riemann zeta function and analytic continuation* (YouTube): an animated version of §4.5.3, "analytic continuation." Highly recommended.
- **3Blue1Brown**, *But what is the Fourier Transform?: the intuition behind §4.3*, "breaking a signal into waves."
- **Barry Mazur & William Stein**, *Prime Numbers and the Riemann Hypothesis* (Cambridge University Press, 2016): a book written specifically for readers who are not math specialists; the first half needs only high-school math.
- Original paper: <https://github.com/openai/math/blob/main/preprints/The-Quasi-Riemann-Hypothesis-September-30-2026/paper.pdf>
- Companion papers: the 11/12 version <https://github.com/openai/math/blob/main/preprints/The-Quasi-Riemann-Hypothesis-October-5-2026/paper2.pdf> ; Siegel zeros

<https://github.com/openai/math/blob/main/preprints/Uniform-exclusion-of-Landau-Siegel-zeros-October-1-2026/paper.pdf>

- Lean scope note: <https://github.com/openai/math/blob/main/lean/docs/003.md> ; challenge statement:
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/QuasiRiemannHypothesis.lean>
- Dave Goldblatt's independent re-check: <https://github.com/davegoldblatt/openai-zeta-proof-check>

NUMBER THEORY · FAMILY 017 (+ 005)

5. How well can fractions approximate π : the irrationality exponent $\mu(\pi) = 2$

Number theory · OpenAI family 017 (§5.7.4 also covers family 005 from the same release: Catalan's constant). All you need going in: you can work with fractions, powers and logarithms, and you know what it means for an infinite sum to "converge." Every new concept is explained before it is used. Every number in this chapter was either computed by script (Python + mpmath, with π to 3000 digits of precision) or has a source. Tags: [P] = a primary source we read directly (paper, repo file, Lean source), [R] = secondary reporting, $\triangle$ = unverified or doubtful. The **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. Estimated reading time: 60–80 minutes. You can read it in three sittings: §5.1–5.3 (approximating by fractions, continued fractions, Dirichlet), §5.4–5.6 (the irrationality exponent, the history for π , the Flint Hills series), §5.7–5.9 (OpenAI's result, its verification and what it means).

5.0 The bottom line: the one thing this chapter explains

π is irrational: it cannot be written as a fraction. But it can be **approximated** by fractions: $22/7$, $355/113$, ... The bigger the denominator, the more accurate the approximation, which is nothing special. The real question is:

Relative to the size of the denominator, how "well" can π be approximated?

This chapter explains one chain:

Any irrational number can be approximated to within "error $\approx 1/q^2$ " (Dirichlet; everyone gets this) $\Rightarrow$ the question becomes "can you do **much better than $1/q^2$** , and do so infinitely often" $\Rightarrow$ a number μ (the irrationality exponent) measures this capacity for "exceptional approximation" $\Rightarrow$ almost every real number has $\mu = 2$, but for the **specific** number π , for 70 years all anyone could prove was $\mu \leq 7.1$.

OpenAI claims to have proved $\mu(\pi) = 2$. That is: in this respect π is as ordinary as "a real number picked at random"; there are not infinitely many "absurdly good" fractional approximations. As a side effect, it settles a popular puzzle: the Flint Hills series $\sum 1/(n^3 \sin^2 n)$ converges.

After reading this chapter, you should be able to explain in your own words:

- how to measure precisely "how well a number can be approximated by fractions," and why the baseline is $1/q^2$;
- why "almost every number has $\mu = 2$ " is easy to prove, while " π has $\mu = 2$ " is extremely hard;
- what playbook OpenAI's proof uses (in one sentence: "a nonzero integer is at least 1"), and how far it has been verified;
- what this result does **not** mean (it is not transcendence, not "normality," not "bounded continued fraction").

5.1 Approximating π by fractions: from 3 to 355/113

5.1.1 A table

$\pi = 3.14159265358979\dots$ Some historically famous fractions:

Fraction	Origin	Error $ \pi - p/q $	$1/q^2$	$q^2 \times \text{error}$
3/1	Antiquity	1.42×10^{-1}	1	0.142
314/100	Decimal truncation	1.59×10^{-3}	1×10^{-4}	15.9
22/7	Archimedes' "yuelü" ("approximate ratio")	1.26×10^{-3}	2.04×10^{-2}	0.062
333/106		8.32×10^{-5}	8.90×10^{-5}	0.935
355/113	Zu Chongzhi's "milü" ("close ratio," 5th century)	2.67×10^{-7}	7.83×10^{-5}	0.0034
103993/33102		5.78×10^{-10}	9.13×10^{-10}	0.633

(Computed by script.)

5.1.2 Why divide by the "cost of the denominator"

A smaller error is not remarkable in itself: if you allow denominators up to 10^6 , you can just take 3141593/1000000 and push the error down to around 10^{-7} . So looking only at the error is unfair; you have to look at the error **relative to the size of the denominator**.

Compare 314/100 and 22/7: their errors are about the same (both around 10^{-3}), but 22/7 uses only the denominator 7, while 314/100 uses the denominator 100. 22/7 is much better "value for money."

A natural measure: **the error is one over what power of the denominator?** If $|\pi - p/q| \approx q^{-\nu}$, we say this approximation has "effective exponent" ν . The bigger ν , the more "exceptional" the approximation.

- $314/100$: error $\approx 100^{-1.4}$, $\nu \approx 1.4$. Decimal truncation only reaches about the $\nu \approx 1$ level (error $\sim 1/q$).
- $22/7$: $\nu \approx 3.43$ (computed by script).
- $355/113$: $\nu \approx 3.20$ (computed by script). This is what is magical about the "milü": a three-digit denominator gives accuracy of about 10^{-7} (the first 6 decimal places are all correct).

The last column of the table, $q^2 \times \text{error}$, is another angle: if it is < 1 , the error is smaller than $1/q^2$. For $355/113$ this column is 0.0034 , i.e. about **290 times** better than $1/q^2$.

[Check 1] Verify with a calculator: $355/113 = 3.14159292\dots$; how far is it from π ? What is 113^2 ? Roughly what is the error times 113^2 ?

5.2 Continued fractions: a machine for finding "best approximations"

The fractions 3 , $22/7$, $333/106$, $355/113$ in the table were not cobbled together; they are the output of a machine. The machine is called a **continued fraction**.

5.2.1 The algorithm: take the integer part, take the reciprocal, repeat

For a number x :

1. Write down the integer part a_0 ; what remains is the fractional part r .
2. Take the reciprocal $1/r$, write down its integer part a_1 , and keep the new fractional part.
3. Repeat.

The first two steps for π by hand:

- $\pi = 3 + 0.14159\dots$, so $a_0 = 3$.
- $1/0.14159\dots = 7.0625\dots$, so $a_1 = 7$.
- $1/0.0625\dots = 15.996\dots$, so $a_2 = 15$.

Continuing (computed by script):

$$\pi = [3; 7, 15, 1, 292, 1, 1, 1, 2, 1, 3, 1, 14, 2, 1, 1, 2, 2, 2, 2, \dots]$$

This means $\pi = 3 + 1/(7 + 1/(15 + 1/(1 + 1/(292 + \dots))))$.

5.2.2 Truncation = good fractions

Cut it off at some position and you get a fraction, called a **convergent**:

- Cut at [3] $\rightarrow 3$
- Cut at [3; 7] $\rightarrow 3 + 1/7 = 22/7$
- Cut at [3; 7, 15] $\rightarrow 333/106$
- Cut at [3; 7, 15, 1] $\rightarrow 355/113$
- Cut at [3; 7, 15, 1, 292] $\rightarrow 103993/33102$

A classical theorem (Lagrange and others, 18th century): convergents are "best approximations." Among all fractions with denominator at most q , none is closer to x .

5.2.3 The bigger the next number, the more accurate the current fraction

This is the most useful rule of thumb in the chapter:

The error of a convergent p/q is about $1 / (a_{\text{next}} \times q^2)$, where a_{next} is the number in the continued fraction that **comes right after the cut** (called a "partial quotient").

- $22/7$ is followed by 15: error $\approx 1/(15 \times 49) \approx 1.4 \times 10^{-3}$; actual 1.26×10^{-3} .
- $355/113$ is followed by **292**: error $\approx 1/(292 \times 12769) \approx 2.7 \times 10^{-7}$; actual 2.67×10^{-7} .
- $333/106$ is followed by 1: error $\approx 1/(1 \times 106^2)$, the same order as $1/q^2$, not "exceptional" at all ($q^2 \times \text{error} = 0.935$).

So asking "can π be approximated exceptionally well" is equivalent to asking "do especially huge partial quotients show up in π 's continued fraction." 292 is one. Will there be even more outrageous ones later?

5.2.4 Control group: some "well-behaved-looking" numbers

Number	Continued fraction (computed by script)	Pattern
Golden ratio $\phi = (1+\sqrt{5})/2$	[1; 1, 1, 1, 1, ...]	All 1s; the hardest number to approximate
$\sqrt{2}$	[1; 2, 2, 2, 2, ...]	Always 2
e	[2; 1, 2, 1, 1, 4, 1, 1, 6, 1, 1, 8, ...]	1, 1, 2k pattern; grows slowly
π	[3; 7, 15, 1, 292, 1, 1, 1, 2, ...]	No known pattern

For patterned numbers like $\sqrt{2}$ and e , the size of the partial quotients is obvious at a glance, so "how well it can be approximated" can be computed directly. π 's continued fraction has **no known pattern at all**, and that is the root of the difficulty: you can't just stare at its continued fraction and prove "a super-large number will never appear later."

[Check 2] Compute the continued fraction of $355/113$ by hand. (Hint: $355 = 3 \times 113 + 16$; then continue with $113/16$.)

5.3 Dirichlet's theorem: exponent 2 comes free for everyone

5.3.1 The statement

Dirichlet's approximation theorem (1842): for any irrational number x , there are infinitely many fractions p/q with $|x - p/q| < 1/q^2$.

In other words, effective exponent 2 is "free": every irrational number can achieve it. That is why the baseline is $1/q^2$ and not something else.

5.3.2 Proof: the pigeonhole principle, doable by hand

Take an integer Q , say $Q = 10$. Look at the **fractional parts** of the 11 numbers $0 \cdot \pi, 1 \cdot \pi, 2 \cdot \pi, \dots, 10 \cdot \pi$ (computed by script):

q	fractional part of $q\pi$
0	0.000
1	0.142
2	0.283
3	0.425
4	0.566
5	0.708
6	0.850
7	0.991
8	0.133
9	0.274
10	0.416

Cut $[0, 1)$ into 10 drawers of length 0.1. Put 11 numbers into 10 drawers and **at least two land in the same drawer**. From the table: $q = 1$ (0.142) and $q = 8$ (0.133) are both in $[0.1, 0.2)$.

Subtract: $8\pi - 1\pi = 7\pi$, and the fractional parts differ by only 0.009. This says 7π is very close to an integer: $7\pi \approx 21.991$, only 0.0089 away from 22. So

$$|7\pi - 22| < 1/10 \quad \Rightarrow \quad |\pi - 22/7| < 1/(7 \times 10) \leq 1/7^2$$

In general, with Q drawers and $Q+1$ numbers, you can always find $q \leq Q$ with $|qx - p| < 1/Q$, so $|x - p/q| < 1/(qQ) \leq 1/q^2$. Letting Q grow larger and larger gives infinitely many such fractions. QED.

[Check 3] Find another pair of q values in the table that land in the same drawer. What is their difference? Which fraction does it give?

5.3.3 Can you do better than $1/q^2$?

Computed by script: among all reduced fractions with denominators $2 \leq q \leq 100,000$, count how many satisfy $|\pi - p/q| < q^{-(v)}$:

v	Number of fractions that satisfy it	Their denominators
2	8	6, 7, 106, 113, 33102, 33215, 66317, 99532
2.2	2	7, 113
2.5	2	7, 113
3	2	7, 113

At $v = 2$ the count keeps growing (Dirichlet guarantees infinitely many). Just above 2 a few remain (computed by script: denominators 6, 106, 33102, 33215, 66317 and 99532 have effective exponents between 2.01 and 2.11); once $v \geq 2.2$, within 100,000 only two "early miracles" remain: 22/7 and 355/113.

$\mu(\pi) = 2$ says exactly this: for any $v > 2$ (even 2.0001), there are only **finitely many** such fractions. There can be a few miracles early on, but miracles cannot keep appearing forever.

5.4 The irrationality exponent μ : definition and a "zoo"

5.4.1 Definition

Irrationality exponent (also called the irrationality measure) $\mu(x)$ = the supremum of all v such that " $|x - p/q| < q^{-(v)}$ has infinitely many solutions."

In plain words: $\mu(x)$ is the highest exponent at which x can be "exceptionally approximated infinitely many times." The bigger μ is, the easier it is to approximate x very well by fractions, and in that sense the "closer to rational" x is.

Continued-fraction version (a standard result): $\mu(x) = 2 + \limsup (\ln a_{n+1} / \ln q_n)$. That is, **$\mu = 2$ if and only if the partial quotients a_{n+1} grow more slowly than any positive power of q_n .**

An example computed by script: among the first 1000 terms of π 's continued fraction, the largest partial quotient is 20776 (term 431), and the denominator of the convergent just before it has 216 digits. The effective exponent of this approximation is only **2.020**. With a denominator that big, even a partial quotient of more than 20,000 contributes almost nothing to the exponent. By contrast, 355/113 has effective exponent 3.20; that was "luck while the denominator was small."

5.4.2 The zoo

Number	μ	Who proved it
Rational numbers	1	Trivial
Any irrational number	≥ 2	Dirichlet 1842 (§5.3)
Quadratic irrationals such as $\sqrt{2}$	2	Liouville-type argument, see 4.3
Any algebraic irrational (e.g. $\sqrt[3]{2}$)	2	Roth 1955 , Fields Medal 1958
e	2	Computed directly from its patterned continued fraction
Almost every real number (in the Lebesgue-measure sense)	2	Khinchin (1920s), see 4.5
Liouville number $\sum 10^{-(k!)}$	∞	Liouville, mid-19th century
π	Known ≤ 7.10 ; OpenAI claims = 2	See §5.5 and 7

5.4.3 A proof you can do by hand: why $\sqrt{2}$ can't be approximated exceptionally well

This small proof contains the core playbook of the whole chapter (including OpenAI's proof), so it is worth reading slowly.

For any fraction p/q , look at the integer $p^2 - 2q^2$:

- It is **not 0** (because $\sqrt{2}$ is irrational).
- It is an **integer**.
- So $|p^2 - 2q^2| \geq 1$.

On the other hand, $p^2 - 2q^2 = q^2 (p/q - \sqrt{2})(p/q + \sqrt{2})$. So

$$|\sqrt{2} - p/q| = |p^2 - 2q^2| / (q^2 \cdot (p/q + \sqrt{2})) \geq 1 / (q^2 \cdot (p/q + \sqrt{2})) \approx 1/(2.83 q^2)$$

The error is never smaller than about $0.35/q^2$, so the effective exponent can't exceed 2, and $\mu(\sqrt{2}) = 2$.

By hand for $p/q = 7/5$: $p^2 - 2q^2 = 49 - 50 = -1$, absolute value exactly 1. Predicted error = $1/(25 \times (1.4 + 1.41421)) = 0.01421$; actual $|1.41421 - 1.4| = 0.01421$. (The script also checked that for the convergents of $\sqrt{2}$, $q^2 \times \text{error}$ stays steady at around 0.35.)

The core playbook:

Build a quantity tied to the approximation. **Arithmetically**: it is a nonzero integer, so it is ≥ 1 . **Analytically**: if the approximation were too good, it would be < 1 . The two contradict each other.

Remember this sentence; you will meet it again in §5.7.

5.4.4 Liouville numbers: the monsters with $\mu = \infty$

$L = \sum 10^{-(k!)} = 0.1100010000000000000000001000\dots$ (the 1s are at decimal places 1, 2, 6, 24, 120, ...).

Cut it off after one of the 1s, and the resulting fraction has a tiny error (computed by script):

- Cut at place 2: $q = 10^2$, error $\approx 10^{-6} = q^{-3}$
- Cut at place 6: $q = 10^6$, error $\approx 10^{-24} = q^{-4}$
- Cut at place 24: $q = 10^{24}$, error $\approx 10^{-120} = q^{-5}$

The effective exponents 3, 4, 5, ... grow without bound, so $\mu(L) = \infty$. Liouville used numbers of this kind to construct the first numbers in history that were **proved to be transcendental** (since algebraic numbers have finite μ , a number with $\mu = \infty$ cannot be algebraic).

This point is crucial for understanding the difficulty: Liouville numbers are transcendental, with $\mu = \infty$. π is also transcendental (Lindemann 1882). So **the fact that " π is transcendental" is of no help at all for $\mu(\pi)$:** any argument that uses only " π is transcendental" would be refuted by Liouville numbers. To prove $\mu(\pi) = 2$, you have to use structure that only π has.

5.4.5 Why "almost every number has $\mu = 2$ " is easy to prove

Throw a random point x into $[0, 1]$. For a fixed q , the region where " x is within $q^{-(v)}$ of some p/q " has total length about q small intervals $\times$ length $2q^{-(v)}$ each $= 2q^{(1-v)}$.

Add this up over all q : $\sum 2q^{(1-v)}$. When $v > 2$, the exponent $1 - v < -1$, and the sum **converges** (it is just 2 times $\zeta(v-1)$ from §4; $\zeta(s)$ is finite for $s > 1$). Probability theory has a result called the Borel–Cantelli lemma: if the probabilities of a sequence of events have a finite sum, then almost surely only finitely many of the events happen. So almost every x has only finitely many exceptional approximations with exponent v , and $\mu(x) \leq v$. v can be taken arbitrarily close to 2, so almost every x has $\mu = 2$.

This argument cannot be applied to π at all: it is about "a random number," and π is one specific number. It's like how "almost nobody is left-handed" can't tell you whether a particular person is left-handed. This is the typical shape of problems like this: **typical behavior is obvious at a glance; a specific instance is extremely hard to prove** (" π is normal" has the same shape, and is still unsolved).

5.5 The history of upper bounds for π : 70 years from 42 to 7.1

5.5.1 Timeline

Year	Author	Proved upper bound on $\mu(\pi)$	Source
1953	Mahler	42 (and said it could eventually be brought to 30)	<i>On the approximation of π</i> , Indag. Math. 15
1974	Mignotte	21 (all denominators) / 20 (eventually)	Bull. SMF Mém. 37
1993	Hata	8.01604539...	Acta Arith. 63
2008	Salikhov	7.606308...	Russ. Math. Surveys 63(3)
2020	Zeilberger–Zudilin	7.103205334137...	Moscow J. Comb. Number Theory 9(4); arXiv:1912.06345
2026-09	Bai (preprint)	7.101862832357	arXiv:2609.11276 $\triangle$ not peer-reviewed
2026-09-24	OpenAI (claimed)	= 2	§5.7 of this chapter

(Values in the table follow the survey in §1.1 of the OpenAI paper, cross-checked against the deck author's research notes.)

Note the size of the last jump: the first 70 years went $42 \rightarrow 7.10$, each step a large amount of delicate work; then a single jump to 2, and it is an **exact value**, not yet another upper bound.

5.5.2 The ceiling of the old methods

All previous upper bounds came from the same framework (Zudilin's survey arXiv:math/0404523 explains it clearly):

1. Build a sequence of "linear forms" with integer coefficients, $a_n\pi - b_n$, that are **very small but not 0**. Usually these are built from carefully designed integrals.
2. Two quantities pull against each other:
 - how fast the linear form decays (faster is better);
 - how fast the coefficients a_n grow (slower is better). To make the coefficients integers, you usually have to multiply by a large number like $\text{lcm}(1, 2, \dots, n)$ to "clear denominators," and this cost is high.
3. A standard lemma: if the decay rate is $e^{(-\beta n)}$ and the coefficient growth is $e^{(\alpha n)}$, you get $\mu \leq 1 + \alpha/\beta$.

To get $\mu = 2$ you need α and β to be almost equal, i.e. the "cost of clearing denominators" has to be almost zero. For every known integral construction for π , this ratio is far from 1. The Zeilberger–Zudilin generation of work was tightening screws inside this framework (optimizing the parameters of the integrals), and each time could only move the bound a little.

5.5.3 Why Roth's method can't be used directly

Roth's theorem gives $\mu = 2$ for all algebraic numbers, and its method is **multivariable**: take several approximations $p_1/q_1, p_2/q_2, \dots$ at once, with wildly different denominators, build a multivariable polynomial, and derive a contradiction using the "nonzero integer ≥ 1 " playbook from §5.4.3.

But Roth's proof fundamentally uses the fact that x satisfies a polynomial equation with integer coefficients (x is algebraic). π is not algebraic, so there is no such equation to use. And as §5.4.4 said, you also can't use only " π is transcendental." OpenAI's reasoning summary happens to record the moment the model ran into this wall (verbatim):

"Wait this would work for any transcend, contradiction with Liouville. Examine."

Meaning: it noticed that one of its arguments, if valid, would work for **all transcendental numbers**, but Liouville numbers are a counterexample, so the argument must have a hole.

5.6 A concrete consequence: the Flint Hills series

5.6.1 What the series is

$$\sum 1/(n^3 \sin^2 n) = 1/(1^3 \sin^2 1) + 1/(2^3 \sin^2 2) + 1/(3^3 \sin^2 3) + \dots \quad (n \text{ in radians})$$

The question: does it converge? Pickover popularized the problem, and it stayed a popular puzzle for more than two decades; Alekseyev 2011 (arXiv:1104.5100) discussed it systematically, and it remained unsolved.

5.6.2 Why it is related to approximating π


The denominator contains $\sin n$. When an integer n is very close to some integer multiple $k\pi$ of π , $\sin n$ is very small, and that term blows up. And " $n \approx k\pi$ " means " $n/k \approx \pi$," i.e. **n/k is a good fractional approximation of π .**

Partial sums computed by script:

n	This term	Partial sum
1	1.412	1.412
3	1.860	3.423
21	0.00015	3.555
22 (22/7)	1.199	4.754
354	3.2×10^{-8}	4.807
355 (355/113)	24.60	29.406

n	This term	Partial sum
1,000	1.5×10^{-9}	30.175
103,993 (103993/33102)	2.4×10^{-6}	30.3145
1,000,000	8.2×10^{-18}	30.3145

Here are the partial sums for $n \leq 400$, plotted:



Each jump happens when n is the numerator of an unusually good fraction approximating π (so $\sin n$ is tiny). Whether the series converges depends on how often and how large these jumps are, which is governed by the irrationality exponent of π . Values computed by the build script of this deck.

The whole curve is a staircase: most of the time it barely moves, and every time n hits the numerator of an exceptionally good approximation of π , it jumps. The $n = 355$ term alone contributes 24.6, which is 81% of the sum of the first 1 million terms.

[Check 4] Explain why the $n = 355$ term is about 24.6. (Hint: $\sin(355) = \sin(355 - 113\pi)$, and $355 - 113\pi \approx 3.01 \times 10^{-5}$; the sine of a very small angle $\approx$ the angle itself.)

5.6.3 Whether it converges is decided entirely by $\mu(\pi)$

- If π has infinitely many "too good" approximations, there will be infinitely many large enough jumps, and the series diverges.
- If exceptional approximations are rare enough, the jumps add up to something finite, and the series converges.

Two published criteria tie this precisely to $\mu(\pi)$:

- Alekseyev 2011: if the series converges, then $\mu(\pi) \leq 5/2$.
- Meiburg 2022 (arXiv:2208.13356): if $\mu(\pi) < 5/2$, then the series converges.

So you only need to prove $\mu(\pi) < 2.5$ to settle the problem. But the best known bound was 7.10, far away. OpenAI's reasoning summary shows that the model's initial goal was exactly $\mu < 5/2$, and only later did it find it could push all the way to 2 (see §5.7.3).

The Flint Hills corollary in OpenAI's paper (Corollary 1.2) takes exactly this route: the main theorem gives $\mu(\pi) = 2 < 5/2$, and Meiburg's published human result then gives convergence [P]. §5 of the paper also gives its own short proof (group the integers n by the nearest $q\pi$, and control each group with $|\sin n| \geq (2/\pi) \cdot |n - q\pi|$), and derives a more general result: $\sum 1/(n^a |\sin n|^b)$ converges if and only if $a > \max\{1, b\}$.

5.7 What OpenAI claims

5.7.1 The main result

Abstract of the paper *The irrationality exponent of π is 2* (OpenAI, 2026-09-24, 23 pages) [P]:

"We prove the conjecture that the irrationality exponent of π is 2. As a consequence, the classical Flint–Hills series $\sum_{n \geq 1} 1/(n^3 \sin^2 n)$ converges, with angles in radians."

The precise statement of the main theorem (Theorem 1.1):

For every real number $\nu > 2$, there exists an integer $Q(\nu)$ such that for all integers p and all $q \geq Q(\nu)$, $|\pi - p/q| \geq q^{-(\nu)}$.

The paper itself spells out three points (all important):

1. p/q is not required to be in lowest terms.
2. **$Q(\nu)$ is not effectively computable:** "the argument does not give it effectively". That is, it proves "from some denominator on, there are no more exceptional approximations with exponent ν ," but does not tell you where that starts.
3. **$\mu = 2$ is not the same as "bounded partial quotients":** "an exponent of 2 is different from a uniform positive lower bound of the form c/q^2 , or equivalently from bounded continued-fraction partial quotients". The latter (whether the partial quotients of π 's continued fraction are bounded) is still unsolved.

The paper also records other claims in the same area: Mantzakouras and López Zapata's Flint Hills convergence claim was withdrawn in 2026-09; Carella claims $\mu(\pi) = 2$ and the stronger "bounded partial quotients," and the OpenAI paper points out in a footnote that his proof of the stronger claim has a sign error (this is OpenAI's side of the story $\triangle$, with no third-party confirmation).

The paper says that Waldschmidt recorded this statement as a conjecture in *Open Diophantine problems* (Moscow Math. J. 4(1), 2004, p. 265).

5.7.2 Intuition for the method (we only read the introduction and the paper's §5; this is a simplified version)

Recall the playbook from §5.4.3: "a nonzero integer is ≥ 1 , but an approximation that is too good would make it < 1 ." OpenAI's proof is a very heavy-duty version of this playbook.

1. **Proof by contradiction.** Fix some $v > 2$, and assume π has infinitely many exceptional approximations with exponent v .
2. **Pick a few, with widely separated denominators.** From these infinitely many, pick m of them, each denominator much bigger than the last: $q_1 \ll q_2 \ll \dots \ll q_m$. This step is the same as Roth (multivariable, separated denominators).
3. **Use structure only π has: it is a period of the exponential function.** $e^{2\pi i} = 1$, so $2\pi i \cdot j$ is a "period" of the exponential function, and also the source of the multivaluedness of the logarithm. The authors turn the exceptional approximations p_i/q_i into $2i \cdot j \cdot p_i/q_i$, which approximate the true periods $2\pi i \cdot j$. This step answers the question from §5.4.4: the proof does not use " π is transcendental," but "this specific relationship between π and $\exp/\log$," which Liouville numbers do not have.
4. **Build a big determinant.** Take a multivariable polynomial $P(Y, X_1, \dots, X_m)$, expand it along the "logarithmic curve" $Y = 1 + t$, $X_i = c_i + \log(1 + t)$, take the Taylor coefficients, and arrange them into a big matrix. A new theorem in Section 2 of the paper ("separated-weight interpolation," using tools from algebraic geometry such as blowups and ampleness) guarantees that this matrix has full rank, so a **nonzero** determinant can be picked out.
5. **Arithmetic lower bound.** The matrix entries are all in $\mathbb{Q}(i)$ (of the form $a + bi$, with a and b fractions). After multiplying out the denominators, the determinant is a nonzero "Gaussian integer," with absolute value ≥ 1 .
6. **Analytic upper bound.** Swap the approximations $2i \cdot j \cdot p_i/q_i$ back for the true periods $2\pi i \cdot j$, and many rows become Taylor coefficients of the same functions, nearly dependent on one another, so the determinant cancels heavily (the paper calls this "quadratic saving"); the remaining terms carry the tiny approximation errors $|\pi - p_i/q_i|$ as factors. Result: the determinant has absolute value < 1 .
7. **Contradiction.** ≥ 1 and < 1 at the same time is impossible, so the assumption is false, and there are only finitely many exceptional approximations with exponent v .

Why it reaches any $v > 2$: the dimension m (how many approximations you pick) can grow with v . The higher the dimension, the faster the "cancellation saving" grows, until it eventually beats the cost of any $v > 2$. Why $Q(v)$ can't be computed: the argument starts

from "assume there are infinitely many," and the approximations it picks have no concrete location. This is the same source as the "ineffectivity" of Roth's theorem.

5.7.3 The reasoning summary: how the AI got here

OpenAI released a 42-page "abridged chain-of-thought summary" (0) [P]. The script read its full table of contents; the structure is:

- **The summary's first part (sections 1–44):** the goal is only the $\mu < 5/2$ needed for Flint Hills. The model tried dozens of routes, including the BBP formula, modular forms, Padé approximation, the Chudnovsky series, p-adic logarithms and elliptic functions, and kept running into the wall it repeatedly called "height vs denominator cost" (height barrier), i.e. the "cost of clearing denominators" from §5.5.2. Along the way it repeatedly drew on the determinant idea from OpenAI's own Catalan's-constant manuscript (the phrase "OpenAI's Catalan determinants" recurs in the summary; the manuscript is covered in §5.7.4). In the end it used a weighted interpolation determinant to get $\mu \leq 62/25 = 2.48$, and hence Flint Hills convergence. Almost all the titles of sections 32–44 are "audit" or "stress-testing," i.e. repeated self-checking of this step.
- **The summary's second part (section 45):** "takes the earlier report's interpolation-and-determinant technique as a starting point and claims exponent two, explicitly without assuming the earlier approximation bound". On the same technique, it tuned parameters and raised the dimension, pushing to $\mu = 2$.

This document is itself a good case study of the research process: the goal was set conservatively (2.5), and the audit revealed that the technique's actual limit was 2.

5.7.4 Another result from the same release: Catalan's constant is irrational (family 005)

The "Catalan manuscript" that the reasoning summary keeps borrowing from is another paper in the same release: *Catalan's constant is irrational* (2026-09-24, 44 pages) [P]. It pairs with the π result: on method, the π proof borrowed its determinants; on the problem, it was stuck at the same wall as §5.5.2.

What it is. Catalan's constant $G = 1 - 1/3^2 + 1/5^2 - 1/7^2 + \dots \approx 0.91597$. It is the value at $s = 2$ of the Dirichlet beta function $\beta(s) = 1 - 1/3^s + 1/5^s - \dots$ (here β is a function name, unrelated to the decay rate β in §5.5.2). Put another way, it is the value at $s = 2$ of the Dirichlet L-function for the mod-4 character in §5.7.2 of §4.

Why it was hard. In the language of §5.5.2: to prove a constant irrational you want a sequence of integer linear forms $a_n G - b_n$ that are nonzero but tend to 0. In 1978 Apéry proved $\zeta(3)$ irrational this way (van der Poorten's article telling the story is titled *A proof that Euler missed*). The same idea fails for G : Zudilin built Apéry-like recurrences whose

fractions b_n/a_n converge to G quickly, but after clearing denominators the linear forms **do not tend to 0**, so the "cost of clearing denominators" wins [P, as summarized in §1.1 of the paper]. The best previous results were family results: at least one of the five numbers $\beta(2)$, $\beta(4)$, ..., $\beta(10)$ is irrational (Rivoal–Zudilin 2003 first proved it up to $\beta(14)$, Zudilin 2019 narrowed it to $\beta(12)$, Lai–Zhou 2022 to $\beta(10)$), with no way to say which one.

What OpenAI claims: G is irrational. On the Lean side, `lean/docs/005.md` says the claim is formalized, and the challenge statement is `Irrational ( $\sum' j : \mathbb{N}, (-1 : \mathbb{R}) ^ j / ((2 * j + 1 : \mathbb{N}) : \mathbb{R}) ^ 2$ )`, again using only standard Mathlib concepts [P]. The solution is about 981 files and 574,000 lines, and grep finds no `sorry` [P]. $\triangle$ It is likewise **not registered** in `formalization.yaml`; and unlike the π proof, it depends on Catalan modules that OpenAI added to the PrimeNumberTheoremAnd library by patch (the same ~40,000-line patch as in §5.9.2 of §4; we did not grep that part or run the comparator).

If it holds: it is one of the most important results of the "this single classical constant is irrational" kind in the nearly half-century since Apéry.

5.8 How far it has been verified

5.8.1 The Lean statement: you can read it yourself

The "challenge statement" file in the repo,

`lean/ComparatorChallenges/PiExponent.lean` (it imports only Mathlib; the scope note is `lean/docs/017.md`) [P]:

```
theorem main :
  (∀ v : ℝ, 2 < v → ∃ Q : ℤ, 2 ≤ Q ∧
    ∀ p q : ℤ, Q ≤ q →
      (q : ℝ) ^ (-v) ≤ |Real.pi - (p : ℝ) / (q : ℝ)|) ∧
  sSup {v : ℝ | 0 < v ∧
    Set.Infinite {r : ℚ | 2 ≤ r.den ∧
      0 < |Real.pi - (r : ℝ)| ∧
      |Real.pi - (r : ℝ)| < (r.den : ℝ) ^ (-v)}} = 2
```

Piece by piece:

- **First part** (before the `∧`):
 - `∀ v : ℝ, 2 < v →`: for any real number v greater than 2,
 - `∃ Q : ℤ, 2 ≤ Q ∧`: there exists an integer threshold $Q \geq 2$,
 - `∀ p q : ℤ, Q ≤ q →`: such that for all integers p and all integers $q \geq Q$,
 - `(q : ℝ) ^ (-v) ≤ |Real.pi - p/q|`: the distance between π and p/q is at least q^{-v} .

- This is exactly Theorem 1.1. Note that it only says "Q exists"; it does not give Q, consistent with the paper's "not effective."
- **Second part** (after the $\wedge$):
 - $\{r : \mathbb{Q} \mid 2 \leq r.\text{den} \wedge 0 < |\pi - r| \wedge |\pi - r| < r.\text{den}^{(-v)}\}$: the set of all rationals with denominator ≥ 2 that are exceptional approximations with exponent v (Lean's $\mathbb{Q}$ is automatically in lowest terms; `r.den` is the denominator).
 - `Set.Infinite ...`: this set is infinite.
 - `sSup {v | 0 < v  $\wedge$  ...} = 2`: the supremum of all v that "have infinitely many solutions" is 2. This is the definition of μ from §5.4.1, matched word for word.

Two details worth noting:

- The statement **uses only Mathlib's standard definitions** (`Real.pi`, `sSup`, `Set.Infinite`, `ℚ`), and no concepts defined by OpenAI. This is much more trustworthy than "writing the statement with custom definitions."
- Lean's `sSup` has a "junk value" convention: for an empty set or a set with no upper bound, it returns 0. Here the conclusion is $= 2$, not $= 0$, so it does not fall into this trap.

5.8.2 The proof itself (counted by script on a local sparse clone, snapshot fd4aeeb) [P]

- Solution directory `lean/OAI/NumberTheory/PiExponent/`: **869 .lean files, 94,948 lines** (about 95,000).
- `grep`: 0 occurrences of `sorry`, 0 of `native_decide`, 0 custom `axiom`s.
- The comparator config `PiExponent.json` allows only the three standard axioms (`propext`, `Quot.sound`, `Classical.choice`).
- The solution imports only Mathlib and files in its own directory, and **does not depend on any patched external library in `lean/patches/`** (the script checked every import). Compare §4: the proof of the quasi-Riemann hypothesis depends on a patch library of about 40,000 lines, which is one of its main remaining risks; the π proof does not have this problem.
- The entry file `Main.lean` also contains a `flint_hills_summable`: it uses Mathlib's `Summable` to state that $\sum 1/(n^3 \sin^2 n)$ converges, and gives a proof. But the official scope note `lean/docs/017.md` says Flint Hills is "outside this selected statement," meaning it is **not in the statement the comparator checks**. The accurate way to put it: a Lean proof of Flint Hills exists in the repo, but it has not gone through the same "statement-proof" comparison process as the main theorem.

5.8.3 What has not been done ($\triangle$)

- ⚠ This formalization is **not registered** in `lean/formalization.yaml` (we grepped for PiExponent and 017 and found neither). It may have been missed in the 10-07 update.
- ⚠ We **did not compile and run** the comparator ourselves, and we **found no third-party re-check** (Dave Goldblatt's independent check in §4 covered only the ζ proof).
- On 2026-10-07, an issue appeared in Google DeepMind's formal-conjectures repository (#6909, by mo271) proposing to mark its `flint_hills_series_converges` as solved, but the issue itself notes it was "prepared with AI assistance and may contain mistakes" and asks people to check it first. When we looked, it had no comments.
- ⚠ **As of 2026-10-08, we found no named expert in Diophantine approximation / transcendental number theory publicly commenting on the correctness of this paper.** John D. Cook's 2026-10-07 blog post is only a popular explainer and does not assess the proof [R].

5.8.4 How to think about it

Compared with the quasi-Riemann hypothesis, this one is cleaner on the "machine verification" dimension: the statement uses only standard Mathlib concepts and reads like the textbook definition; the proof does not depend on a patch library; and the size (about 95,000 lines) is much smaller. If you trust that Mathlib's definitions of π , real powers and suprema are sound, and that the Lean kernel has no bugs, then " $\mu(\pi) = 2$ " is very credible in the machine sense. The remaining gaps are procedural (nobody has run a third-party comparator) and about understanding (no human expert has publicly said they understand the proof).

5.9 If true, what it means (and what it does not mean)

5.9.1 For mathematics: what it means

1. **A formally recorded conjecture is solved:** π is "typical" in its approximation properties, like almost every real number. The 70-year race of upper bounds ($42 \rightarrow 7.10$) ends with an "exact value."
2. **The Flint Hills series converges:** a puzzle popular for more than two decades is settled, along with a complete convergence criterion for the more general $\sum 1/(n^a |\sin n|^b)$.
3. **The method may matter more than the result.** If "separated-weight interpolation along the logarithmic curve + determinants" generalizes, the next targets would be the irrationality exponents of $\log 2$, $\zeta(3)$ and other logarithm values (section 43 of the reasoning summary mentions the model "considered an extension to algebraic logarithms," but this is a speculative direction with no conclusion ⚠). The same family of determinant techniques was also used in the Catalan's-constant manuscript from the same release (§5.7.4).

5.9.2 For mathematics: what it does not mean

Commonly confused claim	Relation to $\mu(\pi) = 2$
π is irrational	Long known (Lambert 1761). $\mu = 2$ is far stronger than "irrational"
π is transcendental	Long known (Lindemann 1882). And transcendental numbers can have $\mu = \infty$ (Liouville numbers), so neither implies the other
The partial quotients of π 's continued fraction are bounded	Still unsolved. The paper says explicitly that $\mu = 2$ is not the same thing
π is normal (every digit string appears with the same frequency)	Still unsolved; a completely different problem
You can compute "from which denominator on there are no more exceptional approximations"	No; $Q(v)$ is not effectively computable
Good approximations like $22/7$ and $355/113$ don't exist	Of course they exist. $\mu = 2$ only says there are finitely many "miracles with exponent > 2 "

5.9.3 For computer science

Almost no direct impact.

- Floating-point trig functions need "argument reduction" (reducing a large x into the interval $[-\pi/4, \pi/4]$), and the worst case depends on how close some floating-point number gets to a multiple of $\pi/2$. But this is a **finite-range** problem (the range of a double is finite), and engineers long ago computed the worst cases directly with continued fractions (see J.-M. Muller, *Elementary Functions*, the chapter on argument reduction [△](#) check the original book for the exact chapter). An "asymptotic theorem that gives no threshold" will not change these concrete numbers.
- No connection to cryptography or complexity theory.

5.9.4 For AI

- This is a relatively **easy-to-check** case of "AI solves a classical conjecture formally recorded in the literature": the statement is short, uses only standard definitions, and the Lean proof does not depend on a patch library. If you want to show someone "which results in this batch are the most trustworthy," it is a better example than the quasi-Riemann hypothesis (§4).
- The reasoning summary shows a process feature worth studying: a conservative goal (2.5) + a long stretch of failed exploration (40-plus sections) + extensive self-auditing, ending with the discovery of the technique's true limit. This is useful for studying the research process of AI itself (and for eval design).
- At the ecosystem level: in the same month there is Bai's 7.1018 preprint (the deck author's research notes record that Bai's acknowledgments mention help from OpenAI models [△](#) not checked against the original), a withdrawn Flint Hills claim, and an AI-drafted issue on

formal-conjectures. The number of "AI-assisted math claims" is surging, and human expert reviewing bandwidth is still the bottleneck, consistent with the conclusion of §4.

5.10 Self-test

Try to answer these after reading. If you can't, go back and reread the corresponding section.

1. Why can't you judge a fractional approximation by the size of its error alone? How is the "effective exponent v " defined? (§5.1.2)
2. Why is $355/113$ especially accurate? Explain in the language of continued fractions. (§5.2.3)
3. Use the pigeonhole principle to explain why every irrational number has $\mu \geq 2$. (§5.3.2)
4. In the "nonzero integer ≥ 1 " playbook, what quantity plays that role in the $\sqrt{2}$ proof, and in OpenAI's proof? (§5.4.3, 7.2)
5. Why doesn't " π is transcendental" help prove $\mu(\pi) = 2$? What structure of π does OpenAI's proof use? (§5.4.4, 7.2)
6. Why are the partial sums of the Flint Hills series a staircase? Why is $\mu(\pi) < 5/2$ enough to make it converge? (§5.6)
7. Which words in the definition of μ does the second part of the Lean statement, `sSup {...} = 2`, correspond to? (§5.8.1)
8. Does $\mu(\pi) = 2$ imply "the partial quotients of π 's continued fraction are bounded"? Does it imply " π is normal"? (§5.9.2)

Appendix: answers to the checks

- **Check 1:** $355/113 = 3.14159292\dots$, which differs from π by about 2.67×10^{-7} ; $113^2 = 12769$; the product is about 0.0034, i.e. about 290 times better than $1/q^2$.
- **Check 2:** $355 = 3 \times 113 + 16$, so $355/113 = 3 + 16/113 = 3 + 1/(113/16)$; $113 = 7 \times 16 + 1$, so $113/16 = 7 + 1/16$. Putting it together, $355/113 = [3; 7, 16]$. (π itself is $[3; 7, 15, 1, \dots]$, and $[3; 7, 15, 1] = [3; 7, 16]$; the two forms are the same fraction.)
- **Check 3:** $q = 2$ (0.283) and $q = 9$ (0.274) are both in $[0.2, 0.3]$; the difference is 7, giving $22/7$ again. ($q = 3$ and $q = 10$ are both in $[0.4, 0.5]$, also with difference 7. With $Q = 10$ the pigeonhole principle can only "see" fractions with denominator ≤ 10 , and the best of those is $22/7$.)
- **Check 4:** $\sin(355) = \sin(355 - 113\pi) \approx 3.01 \times 10^{-5}$ (computed by script: $355 - 113\pi = 3.0144 \times 10^{-5}$). $\sin^2 \approx 9.09 \times 10^{-10}$; $355^3 \approx 4.47 \times 10^7$; the product ≈ 0.0407 , and its reciprocal ≈ 24.6 .

Appendix: Glossary

Term	One-line explanation
Effective exponent ν	If $ x - p/q \approx q^{(-\nu)}$, then ν is how "exceptional" this approximation is
Continued fraction	The expansion $[a_0; a_1, a_2, \dots]$ obtained by repeatedly "taking the integer part, taking the reciprocal"
Convergent	A fraction obtained by truncating a continued fraction; the best approximation among fractions with comparable denominators
Partial quotient	Each a_i in a continued fraction; the bigger it is, the more accurate the convergent just before it
Dirichlet's approximation theorem	Every irrational number has infinitely many approximations with $ x - p/q < 1/q^2$
Irrationality exponent $\mu(x)$	The largest ν at which x can be approximated "infinitely many times" to accuracy $q^{(-\nu)}$
Algebraic number	A root of some polynomial with integer coefficients, such as $\sqrt{2}$ or $\sqrt[3]{2}$
Transcendental number	A number that is not algebraic, such as π , e , or Liouville numbers
Roth's theorem	Every algebraic irrational has $\mu = 2$ (1955, Fields Medal)
Liouville number	A number with $\mu = \infty$, e.g. $\sum 10^{(-k!)}$
Linear form $a\pi - b$	The "very small but nonzero" quantity built by the old methods to put an upper bound on μ
Ineffective	A proof shows a constant exists but cannot compute its value
Flint Hills series	$\sum 1/(n^3 \sin^2 n)$; whether it converges is decided by $\mu(\pi)$
Bounded partial quotients (badly approximable)	$ x - p/q \geq c/q^2$ holds for all fractions; stronger than $\mu = 2$
Normal number	A number in whose expansion every digit string appears with the same frequency; unrelated to μ
Catalan's constant G	$1 - 1/3^2 + 1/5^2 - \dots$; OpenAI claims, in the same release, to prove it irrational (§5.7.4)

Appendix: Further reading

- **Numberphile**, *The Golden Ratio (why it is so irrational)* (YouTube, Ben Sparks): an intuitive version of §5.2.4's "the golden ratio is the hardest number to approximate," which is about its continued fraction being all 1s.
- **Mathologer**, *Infinite fractions and the most irrational number* (YouTube): continued fractions and approximation, with clear animations.
- **A. Ya. Khinchin**, *Continued Fractions* (Dover reprint): a slim book, the classical source for §5.2–5.4 and §5.4.5; readable with undergraduate math.
- **Wadim Zudilin**, *An essay on irrationality measures of π and other logarithms*, arXiv:math/0404523: the authoritative survey of §5.5.2's "linear forms framework."

- **Michel Waldschmidt**, *Open Diophantine problems*, Moscow Math. J. 4(1), 2004:
<https://webusers.imj-prg.fr/~michel.waldschmidt/articles/pdf/odp.pdf> (p. 265 records this conjecture)
 - Zeilberger–Zudilin, arXiv:1912.06345; Alekseyev, arXiv:1104.5100; Meiburg, arXiv:2208.13356.
 - Original paper: <https://github.com/openai/math/blob/main/preprints/The-irrationality-exponent-of-pi-is-2-September-24-2026/paper.pdf>
 - Reasoning summary:
https://github.com/openai/math/blob/main/reasoning_traces/irrationality-exponent-of-pi.pdf
 - Lean scope note: <https://github.com/openai/math/blob/main/lean/docs/017.md> ; challenge statement:
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/PiExponent.lean>
 - Catalan's-constant paper (§5.7.4):
<https://github.com/openai/math/blob/main/preprints/Catalans-constant-is-irrational-September-24-2026/paper.pdf> ; Lean scope note:
<https://github.com/openai/math/blob/main/lean/docs/005.md>
 - formal-conjectures issue #6909: <https://github.com/google-deepmind/formal-conjectures/issues/6909>
 - John D. Cook popular explainer: <https://www.johndcook.com/blog/2026/10/07/irrationality-exponent-of-pi/>
-

THEORETICAL CS · FAMILY 107

6. How fast can matrices be multiplied? $\omega \leq 9/4$

You only need to know matrix multiplication, how to take a logarithm (things like $\log_2 8 = 3$), and what "complexity $O(n^3)$ " means. New ideas such as tensors, rank and "tensor characters" are all explained before they are used. Every number here is either computed by script (script: `精讲/scripts/03_matmul.py`, with symbolic checks in sympy) or has a source. **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. $\triangle$ = I could not fully verify this. Estimated reading time: 60–90 minutes. You can split it into three sittings: §6.1–6.3 (Strassen and ω), §6.4–6.7 (tensors, approximation, the asymptotic spectrum), §6.8–6.10 (OpenAI's result and what it means). One difference from the quasi-Riemann hypothesis (§4): this paper is **only 13 pages, and I read all of it** (not just the introduction). So §6.8 can lay out the full skeleton of the proof.

6.0 The bottom line: the one thing this chapter explains

Multiplying two $n \times n$ matrices the textbook way takes n^3 multiplications. In 1969 Strassen found that you can do it with fewer. Ever since, theoretical computer science has been asking: **what is the smallest power of n you need?** That power is called ω (the matrix multiplication exponent).

This chapter explains one chain of cause and effect:

How small ω is $\iff$ how fast the rank of the matrix multiplication "tensor" grows with $n \iff$ (Strassen's duality theorem) how large every "tensor character" is on matrix multiplication $\iff$ OpenAI: using a simpler tensor, **polynomial multiplication**, they pinned down the value of every character.

Result: $\omega \leq 9/4 = 2.25$. The previous best human result was about 2.3712 (2026-08), and for 35 years progress had only moved the third decimal place.

After this chapter you should be able to explain in your own words:

- why "one fewer multiplication for 2×2 " can rewrite the whole n^3 ;
- what ω means exactly, and why it is the same thing as "the rank of a tensor";
- roughly what steps OpenAI's 13-page proof takes, and why it is not the same road as the "laser method" of the past 35 years;
- why this result **gives no direct speedup at all to GPUs or LLM training**.

6.1 Matrix multiplication: counting the multiplications

6.1.1 The textbook algorithm

$C = A \cdot B$. Each entry of C is a row times a column:

$$c_{ij} = a_{i1} \cdot b_{1j} + a_{i2} \cdot b_{2j} + \dots + a_{in} \cdot b_{nj}$$

C has n^2 entries, and each takes n multiplications and $n-1$ additions. In total: **n^3 multiplications + $(n^3 - n^2)$ additions**. For 2×2 : 8 multiplications, 4 additions.

6.1.2 Why theorists only count "multiplications"

You may find this odd: additions also take time, so why count only multiplications? The reason is recursion (you'll see it in §6.2 shortly). When a and b are themselves matrix blocks, "adding blocks" takes only $(n/2)^2$ operations, while "multiplying blocks" has to recurse, which costs far more. **In a recursion, the number of multiplications sets the exponent; additions only affect the constant.**

6.1.3 How this connects to your daily work

Most of the FLOPs in LLM training are GEMMs (general matrix multiplies). The GEMMs in cuBLAS and CUTLASS use the n^3 algorithm, just cut into small tiles that fit Tensor Cores and shared memory (tiling). That will still be true at the end of this chapter; §6.10.2 explains why.

6.2 Strassen 1969: 7 is enough

6.2.1 Seven products

Let $A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$ and $B = \begin{bmatrix} e & f \\ g & h \end{bmatrix}$. Use only 7 multiplications:

$$\begin{array}{llll} M1 = (a+d)(e+h) & M2 = (c+d)e & M3 = a(f-h) & M4 = d(g-e) \\ M5 = (a+b)h & M6 = (c-a)(e+f) & & M7 = (b-d)(g+h) \\ \\ C11 = M1 + M4 - M5 + M7 & C12 = M3 + M5 & & \\ C21 = M2 + M4 & C22 = M1 - M2 + M3 + M6 & & \end{array}$$

Check C12: $M3 + M5 = af - ah + ah + bh = af + bh \checkmark$. (The script expanded all four entries with sympy and checked them; all four are correct.)

[Check 1] Verify by hand that $C21 = M2 + M4 = ce + dg$.

The cost: multiplications go down from 8 to 7, but additions/subtractions go up from 4 to 18. Doesn't look like a good deal, does it?

6.2.2 The key: the formulas never use $ab = ba$

Notice that in these 7 products, every multiplication has "only entries of A on the left, only entries of B on the right", and the order of multiplication is never swapped. So a...h can **themselves be matrix blocks** (matrix multiplication is not commutative, but we never need it to be). So:

- cut the $n \times n$ matrix into 4 blocks of size $(n/2) \times (n/2)$;
- do 7 "block multiplications" (each one a recursive call) + 18 "block additions";
- recurrence: $T(n) = 7 \cdot T(n/2) + 18 \cdot (n/2)^2$.

Unroll the recursion: at each level down, the number of multiplications goes $\times 7$ and the size goes $\div 2$. There are $\log_2 n$ levels, so the total number of multiplications $= 7^{(\log_2 n)} = n^{(\log_2 7)}$.

$$\log_2 7 = 2.807354... \text{ (computed by script)}$$

The one saved multiplication compounds at every level of the recursion. The table below compares exact counts for $n = 2^k$, recursing all the way down to 1×1 :

n	Naive multiplications n^3	Strassen multiplications 7^k	Ratio	Naive total ops	Strassen total ops (incl. 18 block additions)	Ratio
2	8	7	0.875	12	25	2.08
16	4,096	2,401	0.586	7,936	15,271	1.92
128	2,097,152	823,543	0.393	4,177,920	5,666,497	1.36
1024	1,073,741,824	282,475,249	0.263	2,146,435,072	1,971,035,287	0.918
4096	6.87×10^{10}	1.38×10^{10}	0.201	1.37×10^{11}	9.68×10^{10}	0.704

(Computed by script.) Two things are worth noticing:

1. Counting only multiplications, Strassen wins right away. Counting additions too, **it only starts winning at $n = 1024$** (if you recurse all the way to 1×1). Real implementations switch to the naive algorithm once blocks get small enough (say 64), which moves the crossover earlier. But the pattern "asymptotically faster $\neq$ faster at small sizes" will come up again and again in this chapter.
2. This is only the operation count. Real hardware also has memory traffic and numerical stability; more on that in §6.10.2.

6.3 ω : the matrix multiplication exponent

6.3.1 Definition

ω = the infimum of all τ such that, for every $\varepsilon > 0$, $n \times n$ matrix multiplication can be done with $O(n^{(\tau+\varepsilon)})$ arithmetic operations (+, −, ×).

Unpacking this:

- "Infimum": ω itself might not be achievable, but you can get arbitrarily close to it.
- " $+\varepsilon$ ": allows a tiny bit of slack. That way something like $n^2 \cdot \log n$ still counts as exponent 2.
- The constant inside $O(\dots)$ **can grow arbitrarily large as ε shrinks**. This is the root of the "galactic algorithms" in §6.10.
- The computation is exact, over the complex numbers (or any fixed field). Floating-point error and the bit length of numbers are ignored.

6.3.2 Why $2 \leq \omega \leq 3$

- $\omega \leq 3$: the textbook algorithm.
- $\omega \geq 2$: the output has n^2 numbers, and each has to be computed at least once. (Reading in the $2n^2$ inputs also takes that many steps.)

- Strassen proved $\omega \leq 2.807$.

The mainstream conjecture: $\omega = 2$. That is, matrix multiplication can be almost as fast as "reading the data once". Nobody can prove it, and nobody can prove $\omega > 2$ either.

6.3.3 Why it matters

For many basic problems, the best known algorithm's complexity can be written directly as n^ω : solving linear systems, matrix inversion, determinants, LU decomposition, transitive closure of a graph, maximum matching in bipartite graphs, and more [P] (standard results; see the Bürgisser–Clausen–Shokrollahi textbook in "Further reading"). Linear-programming solvers are on the list too: the interior-point method of Cohen–Lee–Song (*Solving Linear Programs in the Current Matrix Multiplication Time*, STOC 2019, arXiv:1810.07896) solves a linear program with n variables in about n^ω time at the current ω (ignoring log factors) [P]. When ω goes down, their theoretical complexity goes down with it.

6.4 Tensors: turning "an algorithm" into "a number"

This section is the foundation of the whole chapter. It replaces the task of "designing an algorithm" with the task of "finding the rank of a three-dimensional array".

6.4.1 Starting from the rank of a matrix

The rank of a matrix = the smallest number of "outer products" (column vector $\times$ row vector) it can be written as a sum of. For example, $[[1,2],[2,4]] = [1,2]^T \cdot [1,2]$, which has rank 1.

6.4.2 The matrix multiplication tensor

Write matrix multiplication as a big "polynomial" in three groups of variables (a trilinear form):

$$T_n = \sum_{\{i,j,k\}} x_{ij} \cdot y_{jk} \cdot z_{ki}$$

x stands for the entries of A , y for the entries of B , and z is a placeholder for "position (i,k) of C ". The expression means: "entry (i,j) of A times entry (j,k) of B contributes to position (i,k) of C ." Arrange the coefficients into a three-dimensional array and you have a **tensor** (think of it as a three-dimensional matrix). Its three groups of variables x, y, z are called its three **legs**.

For 2×2 , T_2 has $2^3 = 8$ terms, matching the 8 multiplications of the textbook algorithm.

6.4.3 Tensor rank = the minimum number of multiplications

The **rank of a tensor $R(T)$** = the smallest number of "rank-1 terms" it can be written as a sum of, where each rank-1 term looks like:

(some linear combination of x) $\cdot$ (some linear combination of y) $\cdot$ (some linear combination of z)

Compare with Strassen's $M1 = (a+d)(e+h)$: the first two brackets are exactly "a linear combination of x " and "a linear combination of y ", and the third bracket records which positions of C $M1$ gets added to ($C11$ and $C22$). **A rank- r decomposition = an algorithm that uses only r multiplications**, and vice versa (such algorithms are called "bilinear algorithms"). So:

Strassen 1969 $\Leftrightarrow R(T_2) \leq 7$. (Winograd 1971 proved that 7 is optimal: $R(T_2) = 7$ [P].)

6.4.4 Tensor product = recursion

"Multiply" T_n and T_m (the tensor product, which pairs up the indices on each leg) and you get exactly $T_{\{nm\}}$. This is the algebraic version of the "blocks inside blocks" recursion from §6.2.2. So $R(T_{\{nm\}}) \leq R(T_n) \cdot R(T_m)$, and:

As soon as you find a $k \times k$ decomposition using r multiplications, you get $\omega \leq \log_k r$.

Size k	Multiplications r	Source	$\log_k r$ (computed by script)
2	8	Naive	3.0000
2	7	Strassen 1969	2.8074
3	23	Laderman 1976	2.8540
4	49	Strassen applied twice	2.8074
4	48	AlphaEvolve 2025 (over the complex numbers)	2.7925
4	47	AlphaTensor 2022 (GF(2) only, i.e. arithmetic mod 2)	2.7773

Note: by "finding one good decomposition of a fixed small size", the best anyone has reached is about **2.77**. Every upper bound on ω below that number comes from the "asymptotic" methods of §6.5–6.7.

[Check 2] If someone found a 4×4 decomposition with 46 multiplications, what bound on ω would follow?

6.4.5 Another tensor: polynomial multiplication (the star of OpenAI's proof)

Multiply two polynomials: $(x_0 + x_1t) \cdot (y_0 + y_1t) = x_0y_0 + (x_0y_1 + x_1y_0)t + x_1y_1t^2$. Written as a tensor:

$$C(a, b) = \sum_{\{i < a, j < b\}} x_i \cdot y_j \cdot z_{\{i+j\}}$$

a and b are the numbers of coefficients of the two input polynomials, and the output has $a + b - 1$ coefficients. The naive method takes $a \cdot b$ multiplications. But its rank is only **$a + b - 1$** :

- evaluate both polynomials at $a + b - 1$ different points and multiply the values ($a + b - 1$ multiplications);
- then recover the coefficients of the product polynomial by interpolation (interpolation uses only additions and multiplication by constants, which don't count as "variable $\times$ variable" multiplications).

The case $a = b = 2$ is the **Karatsuba trick**: 3 multiplications $P_0 = x_0 y_0$, $P_1 = (x_0 + x_1)(y_0 + y_1)$, $P_2 = x_1 y_1$, and the middle coefficient $= P_1 - P_0 - P_2$ (checked by script: it expands to $x_0 y_1 + x_1 y_0$ ✓). This is also where the idea behind fast FFT multiplication comes from.

Key point: the rank of polynomial multiplication is known exactly ($a + b - 1$); the rank of matrix multiplication is unknown. OpenAI's entire proof is "use the known thing to pin down the unknown thing".

[Check 3] Using "evaluate at the three points $t = 0, 1, -1$ ", design a different 3-multiplication algorithm for $(x_0 + x_1 t)(y_0 + y_1 t)$, and write down how to recover the middle coefficient.

6.5 Approximations count too: border rank and "degeneration"

6.5.1 A strange example

Look at this small tensor (it has nothing to do with matrix multiplication; it's just a demonstration):

$$W = x_0 y_0 z_1 + x_0 y_1 z_0 + x_1 y_0 z_0$$

One can prove its rank is 3. But look at this expression (ϵ is a small parameter):

$$\begin{aligned} & [(x_0 + \epsilon x_1)(y_0 + \epsilon y_1)(z_0 + \epsilon z_1) - x_0 y_0 z_0] / \epsilon \\ & = W + \epsilon \cdot (x_0 y_1 z_1 + x_1 y_0 z_1 + x_1 y_1 z_0) + \epsilon^2 \cdot x_1 y_1 z_1 \end{aligned}$$

(Expanded and checked by script with sympy.) The left side has only **2** rank-1 terms, and as $\epsilon \rightarrow 0$ it tends to W . So W can be approximated arbitrarily well by "something of rank 2". We say W has **border rank 2**.

6.5.2 Bini's interpolation: approximate algorithms can become exact ones

Bini et al. found in 1979 that a variant of 2×2 matrix multiplication shows this kind of "approximately save a multiplication" behavior, and Bini 1980 proved a key fact [P]:

asymptotically, approximate algorithms are as good as exact ones. The method: treat the expression containing ε as a polynomial in ε , compute it at several different values of ε , then interpolate out the ε^0 term. The cost is only a factor independent of n , which disappears in the exponent.

6.5.3 "Degeneration": giving variables weights

The ε trick above has a more systematic version, which OpenAI's proof uses over and over:

1. give each variable an integer **weight** w ;
2. multiply each variable by ε^w ;
3. the total weight of a term = the sum of the weights of its three variables; **keep only the terms with the smallest total weight**. All other terms carry higher powers of ε and vanish in the limit.

Getting B from A this way is written $A \rightsquigarrow B$. Bini's interpolation guarantees that B is no "harder" than A . This is a powerful tool: **by choosing weights cleverly, you can "cut" the small tensor you want out of a big tensor, while throwing away every term you don't want.**

6.6 Fifty years of history: a war over the third decimal place

6.6.1 The asymptotic sum inequality and the laser method

- **Schönhage 1981:** if you can compute several **mutually independent** matrix multiplications **at the same time** more cheaply than computing them separately, you also get an upper bound on ω (the "asymptotic sum inequality"). Note that "independent" means no variables are shared on any of the three legs. This is called a **full direct sum**.
- **Strassen 1986–87: the laser method.** Find a tensor that is not itself matrix multiplication but has very low rank, take a very high power of it (tensor it with itself many times), and then "cut" a large number of mutually independent matrix multiplications out of it. Like a laser, it filters the in-phase part out of messy light.
- **Coppersmith–Winograd 1990:** found a particularly good starting tensor (the CW tensor) and, combined with the combinatorics of "sets with no three-term arithmetic progression" for the cutting, got 2.3755.

From 1986 to August 2026, nearly 40 years, almost all progress on the upper bound for ω came from refining the laser method (every row after 1986 in the table below).

6.6.2 Timeline

Year	Authors	Upper bound on ω	Drop from the previous row (computed by script)
1969	Strassen	2.8074	—
1978	Pan	2.796	0.011
1979	Bini et al. (border rank)	2.780	0.016
1981	Schönhage (asymptotic sum inequality)	2.522	0.258
1986	Strassen (laser method)	2.479	0.043
1990	Coppersmith–Winograd	2.3755	0.104
2014	Le Gall (preceded by Stothers 2010, Vassilevska Williams 2012)	2.3728639	0.0026
2021	Alman–Vassilevska Williams	2.3728596	0.0000043
2023	Duan–Wu–Zhou	2.371866	0.00099
2024	Vassilevska Williams–Xu–Xu–Zhou	2.371552	0.00031
2024	Alman–Duan–Vassilevska Williams–Xu–Xu–Zhou	2.371339	0.00021
2026-08	Dupont et al. (incl. AlphaEvolve)	2.371177	0.00016
2026-09-24	OpenAI claim (any field, still within the laser-method framework) [P]	2.371054886	0.00012
2026-09-24	OpenAI claim (complex numbers) [P]	2.258	0.113
2026-10-02	OpenAI claim (complex numbers, the subject of this chapter) [P]	2.25 = 9/4	0.008

(Values before 2026-08 are all peer-reviewed results [P]; Dupont et al.'s 2.371177 is the arXiv preprint 2608.16884, and OpenAI's 9/4 paper itself cites it as the previous record. The three OpenAI rows are taken from the papers in the openai/math repo and have not been peer-reviewed.)

A few comparisons computed by script:


- The 34 years from 1990 → 2024: a total drop of 0.0041.
- The 12 years from 2014 → 2026-08: a total drop of 0.0017.
- 2026-08 → 9/4: a single drop of **0.121**.

6.6.3 The wall the laser method hit

Ambainis–Filmus–Le Gall (STOC 2015) [P] proved that this kind of laser-method analysis, applied to powers of the CW tensor, **cannot go below about 2.3078**. Later, Alman–Vassilevska Williams and others gave more general barriers [P] $\triangle$ (I did not check the specific values one by one). So 2.25 cannot be "the laser method optimized a bit more"; it

has to come from a different road. This is exactly what Aaronson means by "via a completely different approach".

The figure below puts the timeline from §6.6.2 and this barrier together: OpenAI's 2.25 sits below the dashed line (about 2.3078).



6.7 Strassen's asymptotic spectrum: "measuring" instead of "constructing"

This section is the key to understanding OpenAI's proof, and also the most abstract part of the chapter. I'll use engineering analogies as much as I can, and flag where an analogy breaks down.

6.7.1 "Which tensor is stronger"

$A \geq B$ (B is a "restriction" of A) means: you can get B by applying a linear substitution to each of A 's three legs. Intuitively, **if you can compute A , you can compute B** . For example:

- $R(T) \leq r$ is equivalent to " r independent scalar multiplications $\geq T$ " (where r denotes "a sum of r mutually unrelated xyz terms").
- Drop all terms of T_n with $i \neq j$, and what's left is n independent scalar multiplications, so $T_n \geq n$.

6.7.2 Tensor characters: a "consistent measuring device"

A **tensor character** (also called a "point" of the asymptotic spectrum) is a function λ that maps tensors to non-negative real numbers and satisfies four properties:

Property	Meaning	Engineering analogy
$\lambda(1) = 1$	A single scalar multiplication has value 1	Unit calibration
$\lambda(A \oplus B) = \lambda(A) + \lambda(B)$	Independent tasks run in parallel; values add	Throughput adds up
$\lambda(A \otimes B) = \lambda(A) \cdot \lambda(B)$	Nesting (recursion); values multiply	Recursive costs multiply
$A \geq B \Rightarrow \lambda(A) \geq \lambda(B)$	If you can compute A you can compute B, so A's value is no smaller	Monotone

These four immediately imply: **for any nonzero tensor, $1 \leq \lambda(A) \leq R(A)$** . So every character gives a **lower bound** on rank.

Example: flattening rank. Stand one leg of the three-dimensional array up vertically and lay the other two legs out together as a row. You get an ordinary matrix; take its matrix rank. This satisfies all four properties, and these are the three simplest characters. For T_n , the flattening rank is n^2 .

6.7.3 Every character's value on matrix multiplication is n^{3t}

§2.1 of the paper uses a small trick to prove that for any character λ there is a number $t > 0$ such that

$$\lambda(T_n) = n^{3t} \quad \text{for all } n$$

(Proof idea: T_n can be split into a tensor product of three "dot product" tensors, and each dot product has value some power of n under λ ; multiplying gives $n^{(p_X + p_Y + p_Z)}$, and we write $3t = p_X + p_Y + p_Z$.)

For example, flattening rank: $\lambda(T_n) = n^2 \Rightarrow 3t = 2, t = 2/3$.

6.7.4 Strassen's duality theorem: measurements decide everything

Strassen 1986–88 [P] proved a deep theorem (Appendix A of the OpenAI paper re-proves the special case it needs):

The rank exponent of matrix multiplication $\omega \leq 3 \cdot (\text{the largest } t \text{ over all characters})$. More precisely: if $R(T_d)$ is large, then there must be some character λ that "detects" it, i.e. $\lambda(T_d)$ is also large.

Analogy: this is a lot like duality in linear programming, or Farkas' lemma: either a good algorithm exists, or there is a "certificate" (a character) proving it doesn't. So to prove "a good algorithm exists" (ω is small), you **don't need to construct the algorithm**. You only need to prove that "**no** certificate can show it doesn't exist", i.e. **every character has $t \leq 3/4$** .

This is the most important turn in the chapter:

- Laser method: **construct** one specific large tensor decomposition.
- Asymptotic spectrum: **bound every possible measurement**, then use the duality theorem to conclude that an algorithm exists.

An honest note: Strassen proposed the asymptotic-spectrum route back in the 1980s. The hard part is that nobody knows what all the characters look like. Over the complex numbers, the only known infinite family is the "quantum functionals" of Christandl–Vrana–Zuiddam (JAMS 2023) [P]. OpenAI's proof **sidesteps** this difficulty: it doesn't need to know what the characters look like, and derives its inequalities from the four properties alone.

If the $\omega = 2$ conjecture is true, then every character has $t = 2/3$ on matrix multiplication (the same as flattening rank). The $9/4$ result says: $t \leq 3/4$.

6.8 What OpenAI claims

6.8.1 The main result

The abstract of the paper *An Upper Bound of $9/4$ for the Matrix Multiplication Exponent* (2026-10-02, 13 pages) is a single sentence:

"We prove that the exponent of matrix multiplication over the complex numbers is at most $9/4$." [P]

Theorem 1.1: "For every $\varepsilon > 0$, two $n \times n$ complex matrices can be multiplied using $O_\varepsilon(n^{(9/4+\varepsilon)})$ arithmetic operations. In particular, $\omega \leq 9/4$."

The paper itself also says: "Theorem 1.1 concerns the asymptotic arithmetic exponent; its proof does not specify a competitive finite matrix size."

The same family (number 107) has two more manuscripts:

- *Complex Matrix Multiplication Below 2.258 and Rectangular Bounds* (2026-09-24, 83 pages): over fields of characteristic 0, $\omega < 2.258$, the "dual exponent" $\alpha > 0.465$, and $\omega(1, 0.709, 1) < 2.092$. Once $9/4$ came out, the 2.258 was superseded by it.
- *Staggered extraction for exact matrix multiplication over every field* (2026-09-24, 35 pages): over **every field** (including finite fields and positive characteristic), $\omega <$

2.371054886006746. Still within the laser-method framework; about 1.2×10^{-4} better than Dupont et al.

6.8.2 The skeleton of the proof (I read the whole paper; this is a faithful simplification)

The whole proof uses only the tools from §6.4–6.7. It has five steps.

Step 1: separate blocks that "share a leg" (paper §3, the separation lemma).

An old problem for the laser method is that the small blocks you cut out often share variables, so they are not fully independent. OpenAI's core construction handles this case: M blocks that are independent of each other on the y and z legs but **share** the x leg. The method:

- make $5M$ copies, and give each one a phase from a different power of an $L = 5M$ -th root of unity ζ (this is a finite Fourier transform), tagging each x variable with a "tentative label" g ;
- then use the "weighting" from §6.5.3: x gets weight g^2 , y gets weight $h_u - h^2$, z gets weight $-h_v$. On the terms that survive the Fourier transform, the total weight is exactly $(g - h)^2$;
- keep only the terms with weight 0, i.e. $g = h$: each block gets its own private copy of x , and the blocks become fully independent. As a side effect, each block also gains a "dot product" of length M .

The intuition for this step: **use the square $(g-h)^2 \geq 0$ as a "penalty", so that every term whose labels don't match vanishes as $\epsilon \rightarrow 0$** . Combined with the standard "count by frequency" technique, this gives an entropy inequality (Corollary 3.2).

Step 2: look only at polynomial multiplication. For any character λ , define the "symmetrized value" $P(a, b)$ of the polynomial multiplication tensor $C(a, b)$ (evaluate all 6 permutations of the three legs, multiply the 6 values together, then take the $1/(6t)$ -th power; see the paper's eq. 4.4). It satisfies three simple properties (Eq. 4.5):

- $P(a, b) = P(b, a)$;
- $P(1, b) = b$ ($C(1, b)$ is just a dot product);
- **$P(a, b) \leq (a + b - 1)^{1/t}$** : because the rank of polynomial multiplication is $a + b - 1$ (§6.4.5), and a character is $\leq$ rank.

Here t is exactly the "exponent of λ on matrix multiplication" from §6.7.3.

Step 3: two new inequalities (paper §4).

- **Discrete concavity:** $2P(a, b) \geq P(a, b+1) + P(a, b-1)$. (A "determinantal filtration" degenerates $C(a, b) \otimes$ dot product into two blocks, $C(a, b+1)$ and $C(a, b-1)$; then apply Step 1.)

- **Shifted tripling:** $P(a, 3h + a - 1) \geq 3 \cdot P(a, h)$. (Cut the y and z indices of $C(a, 3h+a-1)$ into left, middle and right segments; after weighting, exactly three copies of $C(a, h)$ remain.)

You can work out the smallest case of shifted tripling by hand: $a = 2, h = 1$, i.e. $C(2, 4)$, which has 8 terms $x_i \cdot y_j \cdot z_{i+j}$. Cut y into left $\{0\}$, middle $\{1, 2\}$, right $\{3\}$; cut z into left $\{0, 1\}$, middle $\{2\}$, right $\{3, 4\}$. Middle y gets weight $+1$, middle z gets weight -1 , everything else 0:

Term (x, y, z)	Weight	Result
$(0,0,0), (1,0,1)$	0	Left block kept
$(0,2,2), (1,1,2)$	$+1-1 = 0$	Middle block kept
$(0,3,3), (1,3,4)$	0	Right block kept
$(0,1,1), (1,2,3)$	$+1$	Vanishes

(Checked term by term by script; matches Figure 1 of the paper.) The three blocks don't overlap on y or z , and they share x : exactly the shape Step 1 can handle.

[Check 4] Without looking at the table, work out why the term $(0,1,1)$ has weight $+1$.

Step 4: a pure sequence argument (paper §5, Lemma 5.1). Using only "symmetry + $P(1,b) = b + \text{concavity} + \text{shifted tripling}$ ", you can derive growth along the diagonal:

$$P(a, a) \geq a^{4/3}$$

This step is entirely elementary: concavity says the increments decrease, tripling says the increments can't decrease too fast, and together they squeeze out a recurrence $H_a \geq (1 + 1/(3(a-1))) \cdot H_{a-1}$, finished off with $(1 + 1/(3m))^3 \geq 1 + 1/m$. The script checked this chain of lower bounds (for example, at $a = 1000$ the paper's intermediate lower bound gives $P(a,a) \geq 16,790$, while $a^{4/3} = 10,000$).

Step 5: squeeze from both sides. Step 2 gives an upper bound and Step 4 gives a lower bound:

$$a^{4/3} \leq P(a, a) \leq (2a - 1)^{(1/t)}$$

As $a \rightarrow \infty$, the left side behaves like $a^{1.333}$ and the right side like $a^{(1/t)}$. For this to hold for all a , we need $1/t \geq 4/3$, i.e. $t \leq 3/4$. Computed by script: if some character had $t = 0.8$, this inequality would already be violated at $a = 32,761$; with $t = 0.76$ it would take until $a \approx 3.8 \times 10^{22}$.

So **every** character has $\lambda(T_d) = d^{(3t)} \leq d^{(9/4)}$. Then by the duality theorem of §6.7.4, the rank exponent is $\leq 9/4$, so $\omega \leq 9/4$.

6.8.3 Why the final answer is a rational number

Laser-method bounds come from the optimal value of a numerical optimization problem, so they are "computed" irrational-looking numbers like 2.3728639. Here $9/4 = 3 \times 3/4$; the $3/4$ comes from "growth like the $4/3$ power", and the $4/3$ in turn comes from the 3 in the "tripling" construction. Hence Aaronson's line: "a *rational* exponent for once (!)" [R].

6.8.4 What "does not specify a size" actually means

The last paragraph of the paper turns " $\omega \leq 9/4$ " into an algorithm: by the definition of the rank exponent, there **exists** some fixed u and a decomposition of T_u with $r < u^{(9/4+\delta)}$ multiplications; then you recurse as Strassen did. But the proof **doesn't tell you what u is, and doesn't give the decomposition**, because the duality theorem of §6.7.4 uses compactness and a fixed-point theorem (Schauder–Tychonoff). It is essentially an existence proof.

What this means for engineering can be estimated with known lower bounds (computed by script):

- $u = 2$: $R(T_2) = 7 > 2^{2.25} \approx 4.76$. Impossible.
- $u = 4$: you'd need $r \leq 22$, but Bläser 2003 proved $R(T_n) \geq 2.5n^2 - 3n$ [P] $\triangle$ (details of the source not re-checked), which gives a lower bound of 28 at $n = 4$. Impossible.
- By this lower bound, **u must be at least 34** for a decomposition meeting the requirement to possibly exist. The real u may be far larger; nobody knows.

Remark 5.2 of the paper adds that the coefficients of this decomposition can be taken to be algebraic numbers (in some number field). So in principle it could be found by a finite search, but there is no practical feasibility at all. (This is my inference $\triangle$.)

A comment on Aaronson's blog (Bram Cohen, #42) says the result "isn't constructive". That judgment is right, but his description "ramsey type thing" is inaccurate: the non-constructiveness comes from the compactness + fixed-point argument in the duality theorem, not from a Ramsey-type combinatorial argument.

6.9 How far it has been verified

6.9.1 The Lean statement: read the model first, then the theorem

Unlike the quasi-Riemann hypothesis (§4), ω here is not a ready-made Mathlib definition; it is written by the challenge file `lean/ComparatorChallenges/MatrixMultiplication.lean`

itself. So **reviewing the definitions** matters more than reviewing the theorem. Here is my block-by-block translation from reading the Lean source [P]:

- `Gate`: there are only five kinds of "gate": `constant` (a constant), `input` (read one input), `add`, `sub`, `mul` (operate on two existing registers).
- `Gate.cost`: constants and inputs cost 0; **add, subtract and multiply each cost 1** (the source reads `.constant _ => 0`, `.input _ => 0`, and `=> 1` for each of `.add`, `.sub`, `.mul`). So this is the standard "number of arithmetic operations" model. (Multiplying by a constant also uses a `mul` gate, so it also costs 1.)
- `Program`: a sequence of gates, where each gate can only refer to earlier registers. This is a "straight-line program": no loops, no branches, no division.
- `MatrixAlgorithm F a b c`: a program, plus "which register each position of the output matrix takes".
- `Correct`: for **all** input matrices A, B, the output equals `A * B` (Mathlib's standard matrix multiplication).
- `AdmissibleExponent F τ`: for every $\varepsilon > 0$ there is a constant C such that for **every** $n \geq 1$ there is a correct program with cost $\leq C \cdot n^{(\tau+\varepsilon)}$.
- `omega F := sInf {τ | AdmissibleExponent F τ}`: the infimum.

Then the theorem itself:

```
theorem complex_omega_le_nine_quarters :  
  Arithmetic.omega C ≤ (9 : ℝ) / 4
```

Its full name is `OAI.MatrixMultiplication.complex_omega_le_nine_quarters`.

Translation: **over the complex numbers, the matrix multiplication exponent, defined by the model above, is $\leq 9/4$** . This matches the definition in §6.3.1. Note that it is "non-uniform" (each n may have a different program; there is no requirement for one uniform algorithm that generates them), which is also the standard definition of ω in the literature (non-uniform algebraic complexity). My judgment: the model is consistent with the literature, with no hidden assumptions.

A specification trap worth noting (especially interesting for AI safety people): in Lean/Mathlib, **the `sInf` of the empty set is defined to be 0**, and the `sInf` of a set with no lower bound is also defined to be 0. So if some definition were written wrong so that "no program is ever judged correct", then $\omega = 0$ and " $\omega \leq 9/4$ " would hold **trivially**. I checked this: the solution proves `omega_le_three` (the naive algorithm works, so the set is non-empty) and `admissibleExponent_two_le` (every admissible exponent is ≥ 2 , so the

set is bounded below), and the $9/4$ derivation goes through `csInf_le`, which requires actually exhibiting an admissible exponent. So this trap is ruled out.

The same challenge file has two more theorems: $\alpha > 93/200$ and $\omega(\mathbb{C}; 1, 0.709, 1) < 523/250$.

6.9.2 The proof itself

- The Lean proof of the $9/4$ theorem is in `OAI/LinearAlgebra/MatrixMultiplication/AuxiliarySeparation/` (73 files, 10,895 lines). Computing the import closure from its `Main.lean`, the project-internal total is **129 files, about 20,700 lines** (computed by script, including shared directories such as Arithmetic, Tensor and Polynomial). The whole matrix multiplication directory (including 2.258 and the rectangular bounds) is 509 files, about 83,000 lines.
- I grepped: no `sorry`, no custom `axiom`, no `native_decide` (the three hits for `admit` are all the English word "admits" in comments).
- The Comparator config allows only three standard axioms: `propext`, `Quot.sound`, `Classical.choice`. Note that the config has `"enable_nanoda": false`, i.e. **this challenge did not turn on a second, independent kernel**.
- **External dependency**: the fixed-point theorem in Appendix A uses the third-party library `fixed-point-theorems-lean4` (harfe, the Brouwer fixed-point theorem). OpenAI applied a porting patch to it (about 77 lines added, 63 deleted). I looked at it: it mainly adapts imports and syntax to Lean 4.34, and I saw no `sorry` or new axioms.
- A small documentation oddity: the scope note `lean/docs/107.md` lists only the 2.258 and staggered papers, not the $9/4$ paper; but the $9/4$ theorem is indeed proved in `AuxiliarySeparation/Main.lean` (the file comment says "Theorem 1.1").

6.9.3 Independent re-checks

- `erenciracioglu-dotcom/openai-math-9-4-lean-check` [R]: compiled these three theorems with Lean v4.34.1 on Windows 11, used `pp.all` to compare the challenge statement and the solution statement and found them identical, and reports dependence on only the three standard axioms. But it says itself: "It does **not** replace Comparator: there was no sandbox and no second, independent kernel replay." Also, it used openai/math commit `adc7f12`, not the `fd4aeeb` used in these notes $\triangle$ (I did not compare whether this part changed between the two commits).
- I **did not compile** it myself.

6.9.4 Human experts

- Scott Aaronson ("The Mathocalypse", 2026-10-07): "Matrix multiplication in $n^{(9/4+o(1))}$ time—a rational exponent for once (!), and via a completely different approach than was

used for $O(n^{(2.373)})$ and so forth" [R]. This is a comment, not a verification.

- **As of 2026-10-08, I found no laser-method / tensor expert (Vassilevska Williams, Alman, Le Gall, Zuiddam, Christandl, etc.) publicly saying they have read and verified this paper** $\triangle$. The vibemathed.com entry also says "No independent mathematician has checked this yet" [R].
- One favorable factor (my judgment): the paper is only 13 pages, uses only standard tools, and Appendix A fully re-proves the special case of the duality theorem it needs. It is **one of the headline results of the whole openai/math release most likely to be fully understood by human experts within a few weeks.**

6.9.5 Remaining risks

1. The custom computational model: I read it and think it is faithful (§6.9.1). But that is the level of "one person read it once".
2. No second-kernel re-check.
3. Correctness of the third-party fixed-point library: the Lean kernel checks its proofs, so the risk is mainly at the level of definitions, and it is low.
4. Risks beyond correctness: the conclusion is right, but humans don't yet understand "why nobody made this road work before".

6.10 What it means if true (and what it doesn't)

6.10.1 For mathematics and theoretical computer science

1. **The biggest single step in 35 years:** from 2.3712 to 2.25, a drop about 70 times the total of the previous 12 years (computed by script: 0.121 vs 0.0017).
2. **It gets around the laser-method barrier** (§6.6.3), showing that Strassen's asymptotic-spectrum route is workable, and can be used in a quite "elementary" way.
3. The theoretical upper bounds of all algorithms whose complexity is written as n^ω (§6.3.3) go down together.
4. **Is $\omega = 2$ now closer?** This is speculation. Where this method's bottleneck lies (whether the $4/3$ growth rate is a limit of the method itself), and whether it can be pushed to 2, the paper doesn't discuss $\triangle$.
5. **Only for the complex numbers** (and, via the algebraic-coefficient argument in the paper's Remark 5.2). Over fields of positive characteristic, OpenAI only claims 2.371054886.

6.10.2 For GPUs and LLM training: no direct impact

This is where engineers most easily get it wrong, so let me be concrete:

1. **We barely use even Strassen.** The table in §6.2.2 shows that, counting additions, plain Strassen only wins on operation count at $n \approx 1024$. On real hardware there is also: (a) the extra block additions need extra memory reads and writes with an irregular access pattern, and GEMM is often bound by memory bandwidth and Tensor Core throughput; (b) Tensor Cores only accelerate "multiply-add", so block additions don't get accelerated; (c) numerical stability is worse than the naive algorithm (Higham's textbook analyzes this specifically). Some work has achieved practical Strassen speedups on CPUs (Huang et al., "Strassen's Algorithm Reloaded", SC16), but mainstream deep learning libraries don't use it.
2. **The 9/4 algorithm is a "galactic algorithm":** asymptotically faster, but only wins on absurdly large inputs. The paper gives no constant, so what follows is a **hypothetical** calculation (computed by script): if the 9/4 algorithm's constant is C and the naive algorithm's is $2n^3$, the crossover is around $n \approx (C/2)^{4/3}$. With $C = 10^6$, $n \approx 4 \times 10^7$, and a single such fp16 matrix takes about 3×10^{15} bytes (3 PB). And nobody knows the real constant; as §6.8.4 said, even the size u of the base decomposition is at least 34, and the actual decomposition is unknown.
3. **The 2.37 laser-method algorithms are galactic algorithms too.** So this is not a problem specific to 9/4: **none of the results with ω below about 2.77 has ever been usable in practice.**
4. **Compare AlphaTensor / AlphaEvolve:** AlphaTensor (Fawzi et al., *Nature* 610, 2022) [P] and AlphaEvolve (2025, 48 multiplications for 4×4 complex matrices) [R] look for low-rank decompositions of a **fixed small size**. That kind of thing can be written directly into a kernel and gives constant-factor speedups (the AlphaTensor paper even optimized decompositions for specific GPUs). ω is about the exponent as $n \rightarrow \infty$. The two are almost different problems. Dupont et al. (2026-08) used AlphaEvolve to optimize parameters inside the laser method; that was the first time AI helped set a new record bound on ω , but the improvement was only 1.6×10^{-4} .

6.10.3 For AI

- **This is a "new idea", not a "bigger optimization".** Compared with Dupont et al. "using AI to tune parameters and squeeze out 1.6×10^{-4} ", the 9/4 proof switched to a different route, and the final argument is short enough to write out by hand. If it holds up, it is one of the clearest examples of "AI producing a mathematical idea that humans didn't find in 35 years".
- **13 pages + about 20,000 lines of Lean make a good touchstone for human verification.** Unlike the 199 pages + 480,000 lines of Lean for the quasi-Riemann hypothesis (§4), here human experts can read the whole thing independently. How experts assess it over the next few weeks is a good signal for judging "the credibility of the whole openai/math release".

- **Spec review matters more than proof review** (§6.9.1): a Lean proof is only trustworthy if "the computational model is written correctly". This is the math version of the "spec gaming" problem in AI safety, and here the review only requires reading 100 lines of definitions.

6.11 Self-test

Try answering these after reading. If you can't, go back and reread the relevant section.

1. Why can Strassen's 7-multiplication formula be used recursively? Which property does it crucially not use? (§6.2.2)
2. What is the relationship between "a rank- r tensor decomposition" and "an algorithm using r multiplications"? (§6.4.3)
3. Why is the rank of the polynomial multiplication tensor $C(a, b)$ equal to $a + b - 1$? How does this relate to Karatsuba? (§6.4.5)
4. What four properties must a tensor character satisfy? Why is every character $\leq$ rank? (§6.7.2)
5. Between the laser method and the asymptotic-spectrum route, which is "constructing" and which is "measuring"? Why can the latter get around the known barriers? (§6.6.3, §6.7.4)
6. In the last step of OpenAI's proof, which two inequalities squeeze out $t \leq 3/4$? (§6.8.2, step 5)
7. In Lean, why does " $\text{slnf of the empty set} = 0$ " create a specification risk? How is it ruled out here? (§6.9.1)
8. Why won't $\omega \leq 9/4$ make your GPU training faster? Give at least three reasons. (§6.10.2)

Appendix: answers to the checks

- **Check 1:** $M_2 + M_4 = (c+d)e + d(g-e) = ce + de + dg - de = ce + dg \checkmark$.
- **Check 2:** $\omega \leq \log_4 46 \approx 2.7618$ (computed by script). Better than Strassen's 2.8074, but far from the laser method's 2.37.
- **Check 3:** Evaluate at $t = 0$: $P_0 = x_0 y_0$; $t = 1$: $P_1 = (x_0 + x_1)(y_0 + y_1)$; $t = -1$: $P_2 = (x_0 - x_1)(y_0 - y_1)$. The product polynomial is $c_0 + c_1 t + c_2 t^2$, so $P_1 = c_0 + c_1 + c_2$ and $P_2 = c_0 - c_1 + c_2$. Hence $c_0 = P_0$, $c_1 = (P_1 - P_2)/2$, $c_2 = (P_1 + P_2)/2 - P_0$. Dividing by 2 is multiplication by a constant, not a "variable $\times$ variable" multiplication.
- **Check 4:** In $(0, 1, 1)$, the y index 1 is in the middle segment, weight +1; the z index 1 is in the left segment $\{0, 1\}$, weight 0. The total weight is $+1 > 0$, so it vanishes as $\epsilon \rightarrow 0$.

Appendix: Glossary

Term	One-line explanation
ω (matrix multiplication exponent)	The optimal exponent for the number of arithmetic operations needed for $n \times n$ matrix multiplication; $2 \leq \omega \leq 3$
Bilinear algorithm	An algorithm in which every multiplication is "a linear combination of A's entries $\times$ a linear combination of B's entries"
Tensor	A three-dimensional array; here equivalent to a trilinear form $\sum c \cdot x \cdot y \cdot z$
Leg	One of the tensor's three groups of variables x, y, z
Tensor rank	The smallest number of rank-1 terms it can be written as a sum of; equals the minimum number of multiplications of a bilinear algorithm
Border rank	The smallest rank when approximation by expressions containing ϵ is allowed
Degeneration	Give variables weights as powers of ϵ and keep only the lowest-weight terms
Restriction	Get another tensor by linear substitutions on the three legs; $A \geq B$ means "if you can compute A you can compute B"
Full direct sum	A sum of several tensors that share no variables on any of the three legs
Laser method	Cutting many independent matrix multiplications out of a high power of some low-rank tensor (Strassen 1986)
Asymptotic spectrum / tensor character	The set of all "measurement functions" that are normalized, additive, multiplicative and monotone; Strassen proved they determine asymptotic complexity
Flattening rank	The rank of the tensor after laying it out as a matrix; the simplest tensor character
Galactic algorithm	An algorithm that is asymptotically faster but only wins on unrealistically large inputs
Straight-line program	A sequence of operations with no loops or branches; the Lean formalization uses it to model algorithms

Appendix: Further reading

- **The original paper** (13 pages, readable in full):
<https://github.com/openai/math/blob/main/preprints/Matrix-Multiplication-Nine-Fourths-October-2-2026/paper.pdf>
- **Lean statement** (about 100 lines; I recommend reading the definitions yourself):
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/MatrixMultiplication.lean>
- **Independent Lean re-check**: <https://github.com/erenciracioglu-dotcom/openai-math-9-4-lean-check>
- **Quanta Magazine**, *New Breakthrough Brings Matrix Multiplication Closer to Ideal* (2024-03-07): popular coverage of the Duan–Wu–Zhou step, with intuitive pictures of the laser

method. <https://www.quantamagazine.org/new-breakthrough-brings-matrix-multiplication-closer-to-ideal-20240307/>

- **Markus Bläser**, *Fast Matrix Multiplication* (Theory of Computing Graduate Surveys 5, 2013): the best introductory survey, free. <https://theoryofcomputing.org/articles/gs005/>
- **J. M. Landsberg**, *The complexity of matrix multiplication: developments since 2014* (arXiv 1811.11667): a survey from a more geometric point of view.
- **Bürgisser, Clausen, Shokrollahi**, *Algebraic Complexity Theory* (Springer 1997): Chapters 14–15 cover tensor rank, the laser method and the asymptotic spectrum; the standard textbook.
- **Christandl, Vrana, Zuiddam**, *Universal points in the asymptotic spectrum of tensors* (JAMS 36, 2023; arXiv 1709.07851): "what the known characters look like" from §6.7.
- **Ambainis, Filmus, Le Gall**, the laser-method barrier (STOC 2015): <https://arxiv.org/abs/1411.5414>
- **Alman et al.**, *More Asymmetry Yields Faster Matrix Multiplication* (SODA 2025): <https://arxiv.org/abs/2404.16349>
- **Dupont et al.**, *Improving the matrix multiplication exponent with modern optimization and AlphaEvolve* (2026-08): <https://arxiv.org/abs/2608.16884>
- **Fawzi et al.**, *Discovering faster matrix multiplication algorithms with reinforcement learning* (AlphaTensor, *Nature* 610, 2022)
- **Huang, Smith, Henry, van de Geijn**, *Strassen's Algorithm Reloaded* (SC16): when Strassen is worth using on real hardware.
- **Scott Aaronson**, *The Mathocalypse* (2026-10-07): <https://scottaaronson.blog/?p=10169>
- **3Blue1Brown**, *Linear transformations and matrices* (YouTube, Essence of Linear Algebra episodes 3–4): if you want to shore up the intuition that "matrix multiplication = composition of linear transformations" first.

THEORETICAL CS · FAMILY 102

7. The Unique Games Conjecture: where is the ceiling for approximation algorithms?

You only need to know what a graph is (vertices and edges), how to do modular arithmetic ($7 \bmod 3 = 1$), that "polynomial time" roughly means "grows mildly with input size", and to have a vague sense of P vs NP. Every new idea is explained before it is used. Every number here is either computed by script (script: `精讲/scripts/04_ugc.py`) or has a source.

[Check] items are small exercises you can work out yourself; answers are at the end of the

chapter. $\triangle$ = I could not fully verify this. Estimated reading time: 60–90 minutes. You can split it into three sittings: §7.1–7.4 (NP-hardness, approximation, 0.878, PCP), §7.5–7.7 (what a unique game is, why it matters, the 24-year tug-of-war), §7.8–7.10 (OpenAI's result, verification, what it means).

7.0 The bottom line: the one thing this chapter explains

Many important optimization problems (cutting a network into two halves, placing the fewest guards to watch every road...) are believed to have **no fast exact algorithm**. As a fallback, we ask: **what percentage of the optimum can a fast algorithm guarantee at best?**

For the most classic of these, the Max-Cut problem, someone gave an algorithm in 1995 that guarantees **87.8%**. On the other side, in 1997–2001 it was proved that guaranteeing more than **94.1%** is NP-hard. For 30 years, nobody knew where the truth lay in the 6 percentage points in between.

This chapter explains one chain of cause and effect:

How well you can approximate $\Leftarrow$ is decided by how hard it is to "tell nearly satisfiable apart from far from satisfiable" $\Leftarrow$ the most important such constraint system is called Unique Games $\Leftarrow$ the Unique Games Conjecture (UGC) says: telling these apart is NP-hard.

If UGC is true, then 87.8% is **exactly** the limit for Max-Cut, and the same argument settles the approximation limits of a whole class of problems at once. Khot proposed it in 2002, and for 24 years nobody could prove or refute it. **OpenAI claims to have proved it**, with a Lean formal proof attached.

After this chapter you should be able to explain in your own words:

- why a "fully satisfiable" unique game is easy, while a "nearly satisfiable" one is conjectured to be hard;
- exactly how UGC relates to the number 0.878;
- what OpenAI's result means if true, and what it does **not** mean (it has nothing to do with P vs NP).

7.1 P vs NP and NP-hardness: the 5-minute version

7.1.1 "Easy to check" is not "easy to find"

Sudoku: given a filled-in grid, checking whether it's right takes seconds; filling it in from an empty grid can take a long time of trial and error.

- **P**: problems where you can **find** the answer in polynomial time (say n^2 or n^3 steps).

- **NP**: problems where, given a candidate answer, you can **check** whether it is right in polynomial time.

P vs NP asks: is everything that is easy to check also easy to find? Almost everyone believes **P $\neq$ NP**, but nobody has proved it (it is one of the Clay Millennium Problems).

7.1.2 Reductions and NP-hardness

Reduction: quickly translate any instance of problem A into an instance of problem B, so that B's answer is A's answer. If you can do this, then "B has a fast algorithm" implies "A has a fast algorithm", so B is at least as hard as A.

3SAT: given a bunch of clauses like " $(x_1 \text{ OR NOT } x_2 \text{ OR } x_5)$ ", each with 3 variables, can you assign true/false to the variables so that all clauses hold at once? The Cook–Levin theorem (1971) says every problem in NP can be reduced to 3SAT.

NP-hard: if 3SAT can be reduced to problem B, B is called NP-hard. Meaning: **unless P = NP, B has no polynomial-time algorithm**.

Every "X is NP-hard" in this chapter should be read as "unless P = NP, X cannot be done". None of them **proves** P $\neq$ NP; they just tie X's fate to P vs NP.

7.2 The fallback: approximation algorithms

7.2.1 Max-Cut

Given a graph, split the vertices into two groups (say, color them black and white). An edge is "cut" if its two ends have different colors. **Max-Cut** asks: what is the largest number of edges you can cut? Deciding "can you cut $\geq k$ edges" is one of Karp's 21 NP-complete problems from 1972.

Computed by script (brute-force enumeration of all colorings):

Graph	Edges	Max cut	Expected cut of random coloring	Random / optimal
Triangle	3	2	1.5	0.750
5-cycle (pentagon)	5	4	2.5	0.625
Petersen graph	15	12	7.5	0.625

Why a triangle can cut at most 2: three vertices and only two colors means two vertices must share a color, and the edge between them can't be cut.

7.2.2 Approximation ratio, and "1/2 for free"

An **α -approximation algorithm**: runs in polynomial time and always returns a solution worth at least " $\alpha \times$ the optimum".

Random coloring: each vertex flips an independent coin to pick its color. The two ends of an edge get different colors with probability $1/2$, so the expected cut is "half the edges". Since the optimum can't exceed the number of edges, random coloring is at least a **$1/2$ -approximation**.

From 1976 (Sahni–Gonzalez) until 1995, $1/2$ was the best guarantee known.

[Check 1] What is the max cut of a 4-cycle (a square, 4 edges)? How many edges does a random coloring cut in expectation?

7.2.3 Another lead character: Vertex Cover

Pick the fewest vertices in a graph so that every edge has at least one endpoint picked (imagine posting guards at intersections to watch every road).

A simple algorithm known since the 1970s: repeatedly find an edge with neither end picked, and **pick both ends**. The optimal solution has to pick at least one end of that edge, and we picked both, so we use at most **2 times** the optimum. That is a **2-approximation**.

For 50 years, nobody has managed "1.99 times".

7.3 Goemans–Williamson: turning vertices into vectors

7.3.1 Relaxation: first solve a looser problem

Max-Cut is hard because each vertex can take only two values (black = $+1$, white = -1). An edge (i, j) is cut exactly when $x_i \cdot x_j = -1$, so the number of cut edges = $\sum (1 - x_i \cdot x_j)/2$.

The idea of Goemans–Williamson (1995): **replace each ± 1 with a unit vector v_i** , and replace the product $x_i \cdot x_j$ with the inner product $v_i \cdot v_j$:

maximize $\sum_{\text{edges}} (1 - v_i \cdot v_j)/2$ subject to: each v_i is a vector of length 1

This is called a **semidefinite program (SDP)**, and it can be solved (approximately) in polynomial time. Since ± 1 are "one-dimensional unit vectors", every solution of the original problem is also a solution of this one, so **the SDP optimum $\geq$ the true max cut**. Intuitively: the more "opposite" the two ends of an edge are pushed (the closer their angle is to 180°), the better.

7.3.2 Random hyperplane: turning vectors back into black and white

Once you have the vectors, pick a random plane through the origin, splitting space in two: color the points on one side black and the other side white.

Key fact: **two vectors at angle θ are separated by a random plane with probability exactly θ/π** . It's easiest to see in 2D: a random line through the origin falls "between" the two vectors with probability $\theta/180^\circ$.

Verified by script with 1,000,000 random planes:

Angle	Theory θ/π	Simulation
90°	0.5000	0.4994
120°	0.6667	0.6659
150°	0.8333	0.8337

7.3.3 Where 0.878 comes from

For each edge, compare two quantities:

- the "score" the SDP gives it: $(1 - \cos \theta)/2$;
- the probability it is actually cut after rounding: θ/π .

As long as, for **every** angle, the second is at least α times the first, the whole graph gets at least α times the SDP value, and so at least α times the optimum. So the guarantee = the minimum of this ratio over all angles:

$$\alpha_{\text{GW}} = \min \text{ over } \theta \text{ of } (\theta/\pi) / ((1 - \cos \theta)/2) = \min 2\theta / (\pi(1 - \cos \theta))$$

Computed by script (numerical minimization):

Angle θ	SDP score	Probability cut	Ratio
60°	0.250	0.333	1.333
90°	0.500	0.500	1.000
120°	0.750	0.667	0.889
133.56°	0.845	0.742	0.8786 (minimum)
150°	0.933	0.833	0.893
180°	1.000	1.000	1.000

$\alpha_{\text{GW}} = 0.878567\dots$, attained at $\theta \approx 133.56^\circ$ (inner product $\rho \approx -0.689$). This matches the 0.878567... given in Eq. (1.1) of OpenAI's Max-Cut paper.

7.3.4 An example: the pentagon

The optimal SDP solution for the 5-cycle places the 5 vectors in a plane, one every 144° ($= 4\pi/5$). Computed by script:

- $\text{SDP value} = 5 \times (1 - \cos 144^\circ)/2 = \mathbf{4.5225}$;
- $\text{true max cut} = 4$;
- $\text{expected cut of a random plane} = 5 \times 144/180 = \mathbf{4.000}$.

So on the pentagon the SDP value is noticeably higher than the true optimum ($4/4.5225 = 0.8845$). This kind of "SDP is too optimistic" gap is called the **integrality gap**. Feige–Schechtman (2002) built graphs where this ratio gets arbitrarily close to 0.878, so **this SDP plus any rounding method can never beat 0.878**. But that only says "this one algorithm" has hit its limit, not that "all algorithms" have.

7.4 Approximation can be hard too: the PCP theorem

7.4.1 The intuition behind the PCP theorem

The **PCP theorem** (Arora–Safra; Arora–Lund–Motwani–Sudan–Szegedy, 1992 conference, 1998 journal; Dinur 2007 gave a combinatorial proof) says something that sounds like magic:

Any mathematical proof can be rewritten in a special format such that a referee who **randomly reads only a constant number of bits** of it can tell, with high confidence, whether the proof is right: a correct proof always passes; a wrong proof, no matter how it is written, gets caught with probability at least one half.

7.4.2 Why it has to do with "approximation"

Restate PCP: there is a reduction starting from 3SAT that produces a constraint system such that

- formula satisfiable $\Rightarrow$ the constraint system can be **fully** satisfied;
- formula unsatisfiable $\Rightarrow$ the constraint system can satisfy **at most** a fixed fraction $s < 1$.

This is called a **gap**. Now suppose some algorithm guarantees an approximation ratio $\alpha > s$. Then you could use it to tell the two cases apart, and so solve 3SAT. Therefore: **a gap $\Rightarrow$ approximating above the gap is NP-hard**.

The game for the following decades was to make this gap "sharper and sharper", until it sits exactly at the position of a known algorithm.

7.4.3 The gap for Max-Cut


Håstad (2001, *JACM*), together with the "gadgets" of Trevisan–Sorkin–Sudan–Williamson, proved: **approximating Max-Cut above $16/17 \approx 0.941$ is NP-hard**.

So:



Who owns the region between 0.878 and 0.941? A better algorithm, or a stronger hardness proof? This is what UGC is meant to answer.

The top half of the figure below is the same number line; the bottom half is the "completeness ladder" for unique games that §7.5–7.7 explain (proved up to about $1/2$ in 2018; UGC needs $1 - \epsilon$). Skim it now and come back after §7.7.



Vertex cover has a similar gap: the algorithm achieves 2, and the NP-hardness lower bound was Dinur–Safra's **1.3606** (2005, *Annals*), raised in 2018 to $\sqrt{2} \approx 1.414$ (Khot–Minzer–Safra, see §7.7.2).

7.5 What Unique Games are

7.5.1 Labeling games

A **label constraint system** has:

- a graph;
- an "alphabet": each vertex picks a label from $\{0, 1, \dots, k-1\}$;
- a constraint on each edge, saying which labels at the two ends count as "compatible".

The goal: pick labels so that the fraction of satisfied constraints (called the instance's **value**) is as high as possible.

"Label Cover", the standard output of the PCP theorem, has exactly this form, except that its constraints are "many-to-one": once one end is fixed, the other end has several choices.

7.5.2 "Unique": fix one end and the other is determined

Unique Game: every edge constraint is a **permutation** (a one-to-one correspondence). Once one end picks a label, there is **exactly one** label at the other end that satisfies the edge.

The most common example, and the one this chapter uses: **difference equations mod k**

$$x_i - x_j \equiv c \pmod{k}$$

Given x_j , x_i is uniquely determined as $x_j + c$. (OpenAI's paper uses "translation constraints" $a(v) = a(u) + c$, which just replace numbers mod k with bitwise XOR on 0/1 vectors; it's essentially the same.)

7.5.3 By hand: a triangle, mod 3

Three variables a, b, c , labels $\{0, 1, 2\}$:

Instance A: $b - a \equiv 1, c - b \equiv 1, a - c \equiv 1 \pmod{3}$. Add them up around the loop: $0 \equiv 3 \equiv 0$, no contradiction. Take $a = 0$; propagating gives $b = 1, c = 2$; check $a - c = -2 \equiv 1 \checkmark$.

Everything is satisfied; value = 1.

Instance B: change the last one to $a - c \equiv 2$. Add them up around the loop: $0 \equiv 1 + 1 + 2 = 4 \equiv 1$, a contradiction. So the three can't all hold at once, but dropping any one of them makes the rest satisfiable. **Value = 2/3.** (Confirmed by the script by brute-force enumeration of all 27 labelings.)

The general rule: **on a cycle, add up the constraint constants around the loop; if the sum is not 0 (mod k), at least one constraint on that cycle must be violated.**

[Check 2] Four variables $x_0 \dots x_3$, mod 5, constraints: $x_1 - x_0 \equiv 2, x_2 - x_1 \equiv 3, x_3 - x_2 \equiv 1, x_0 - x_3 \equiv 4, x_2 - x_0 \equiv 0$. Can they all be satisfied? What if the last one is changed to $x_2 - x_0 \equiv 1$?

7.5.4 Why "fully satisfiable" is easy

If a unique game can be **fully** satisfied, there is a simple polynomial-time algorithm:

1. in each connected component, pick any vertex and try one of its k labels;
2. propagate along the edges: each edge uniquely determines the neighbor's label;
3. on hitting a contradiction, move to the next starting label; if all k fail, it can't be fully satisfied.

Each attempt is one pass over the graph, k attempts in total. Compare brute-force enumeration: 100 variables mod 3 means trying $3^{100} \approx 10^{47.7}$ possibilities (computed by script).

This is completely different from 3SAT: even when a 3SAT formula is "fully satisfiable", finding a solution is hard. **So UGC can't be about "satisfiable vs unsatisfiable"; it has to be about "nearly satisfiable vs far from satisfiable".** That is where the $1 - \epsilon$ in UGC comes from.

7.5.5 Why "nearly satisfiable" might be hard

Once a few constraints are "wrong", propagation carries the error a long way. The script ran an experiment: a random graph with 2000 vertices and on average 6 edges per vertex, mod 7; first plant a reference answer, then randomly tamper with a fraction of the constraints; then assign labels by propagation (breadth-first) from random starting points, trying 20 starting points:

Fraction tampered	Value of planted answer	Propagation, average value	Propagation, best value
0%	1.000	1.000	1.000
1%	0.990	0.958	0.982
5%	0.947	0.728	0.845
10%	0.901	0.610	0.709

(Computed by script; the random graph has a very small number of isolated vertices, whose label defaults to 0; this doesn't affect the conclusion.)

The reason: a vertex's label is "passed along" a path from the starting point, and as soon as the path crosses one bad edge, everything downstream is shifted.

To be honest: this experiment only shows that "propagation" fails, not that the problem itself is hard. On **random** noisy instances like these, smarter algorithms (such as SDP) actually do very well. What UGC says is: there **exist** carefully constructed instances where no polynomial-time algorithm can tell "it is 99% satisfiable" from "it is only 1% satisfiable".

7.5.6 The precise statement of UGC

Unique Games Conjecture (Khot 2002): for every $\epsilon > 0$ and $\delta > 0$, there is an alphabet size k such that, for unique games with alphabet k , it is NP-hard to tell apart the following two cases:

- YES: value $\geq 1 - \epsilon$ (nearly all satisfiable);
- NO: value $\leq \delta$ (nearly all unsatisfiable).

Note the **order of quantifiers**: first ϵ and δ are given, **then** the alphabet size k is chosen. If you flip it and let k grow with the instance, there are known algorithms that solve it (the SDP algorithm of Charikar–Makarychev–Makarychev 2006 gets about $1 - O(\sqrt{\epsilon \log k})$, which is very good when k is small). §1.2 of OpenAI's paper discusses this point specifically.

7.6 Why UGC matters: one conjecture governs a whole territory

The appeal of UGC: assuming it holds, the **exact** approximation limits of many problems can be computed, and they often **exactly equal some known algorithm**.

Problem	Best algorithm	Unconditional NP-hardness (before 2026)	Assuming UGC
Max-Cut	0.878 (GW 1995)	$16/17 \approx 0.941$ (Håstad 2001)	0.878 is the limit (KKMO 2004/07)
Vertex cover	2	$\sqrt{2} \approx 1.414$ (KMS 2018)	2 is the limit (Khot–Regev 2003/08)
All constraint satisfaction problems (CSPs)	Basic SDP + rounding	Varies	Basic SDP is the optimal algorithm (Raghavendra 2008)

A few key results:

- **KKMO** (Khot–Kindler–Mossel–O'Donnell, FOCS 2004, *SICOMP* 2007): UGC $\Rightarrow$ beating α_{GW} on Max-Cut is NP-hard. A probability statement they needed, "Majority Is Stablest", was later proved by Mossel–O'Donnell–Oleszkiewicz (2005, *Annals* 2010). The odd 133.56° angle in the GW algorithm corresponds exactly to the "noise" parameter in the hardness reduction: the algorithm and the hardness meet at the same number.
- **Khot–Regev** (CCC 2003, *JCSS* 2008): UGC $\Rightarrow$ vertex cover can't achieve $2 - \epsilon$. The clumsy 1970s "pick both ends" algorithm is optimal.
- **Raghavendra** (STOC 2008): UGC $\Rightarrow$ for **every** CSP, the same generic SDP algorithm is the optimal polynomial-time approximation algorithm. One algorithm plus one conjecture settles an entire class of problems.

This is why UGC is called "the map of the field of approximation algorithms". Khot received the 2014 Nevanlinna Prize for it.

7.7 A 24-year tug-of-war: support and doubt

7.7.1 Doubt: a subexponential-time algorithm (2010)

Arora–Barak–Steurer (FOCS 2010, *JACM* 2015) gave an algorithm that, for a unique game with alphabet k , n variables and value $\geq 1 - \epsilon^c$ (c a fixed constant), finds a labeling satisfying $1 - \epsilon$ in $\exp(k \cdot n^\epsilon)$ time.

Why did this make people doubt UGC? Because 3SAT is generally believed to need close to 2^n time, while UG needs only 2 to "a small power of n ". Comparison computed by script:

n	3^n (brute force)	$2^{(n^{0.1})}$	$2^{(n^{1/3})}$	n^3
100	$10^{47.7}$	3	$10^{1.4}$	10^6
1,000	10^{477}	4.0	$10^{3.0}$	10^9
1,000,000	10^{477121}	15.8	$10^{30.1}$	10^{18}

If UGC holds, the reduction from 3SAT to UG must **blow up** the instance to a very large polynomial size (otherwise 3SAT would also have a subexponential algorithm). That is not a contradiction, but it shows that even if UG is hard, it is "not thoroughly hard". Some researchers leaned toward believing UGC is false because of this.

7.7.2 Support: the 2-to-2 theorem (2018)

2-to-2 constraints: a bit looser than unique; once one end is fixed, the other end has 2 allowed labels. A series of works by Khot–Minzer–Safra and Dinur–Khot–Kindler–Minzer–Safra in 2017–2018 (the last piece being the KMS FOCS 2018 theorem on expansion in Grassmann graphs) proved: **nearly satisfiable 2-to-2 games are NP-hard**.

Corollaries:

- Unique games with **completeness about $1/2$** are NP-hard: telling " $\text{value} \geq 1/2 - \epsilon$ " from " $\text{value} \leq \delta$ ". This has been called "half of UGC".
- Vertex cover to within $\sqrt{2} - \epsilon$ is NP-hard.

What remained hard: pushing completeness from $1/2$ to $1 - \epsilon$. §1.3 of OpenAI's paper also says: "Obtaining Unique Games completeness arbitrarily close to one was a separate problem."

7.7.3 September 2026: a race

According to Quanta (2026-10-07), MIT's Dor Minzer and his students Yumou Fei and Shuo Wang, after hearing rumors that OpenAI was about to release a proof of UGC, rushed out a proof of a 4-to-1 variant of Khot's 2-to-1 conjecture on 2026-09-14. O'Donnell's comment: "They solved the problem in the old-fashioned way, with their minds, and wrote it with their own fingers." [R]

7.8 What OpenAI claims

7.8.1 The main result

The abstract of the paper *The Unique Games Theorem* (2026-09-23, 58 pages, author listed as "OpenAI"):

"We prove the Unique Games Conjecture. For every fixed $\epsilon, \delta \in (0, 1/2)$, we give a deterministic polynomial-time reduction from 3SAT to Unique Games over a fixed finite alphabet, with completeness at least $1 - \epsilon$ and soundness at most δ ." [P]

In plain words: **for any fixed ϵ, δ , there is a deterministic polynomial-time reduction that translates 3SAT into unique games with a fixed alphabet: satisfiable formulas become instances with value $\geq 1 - \epsilon$, and unsatisfiable ones become instances with value $\leq \delta$.**

This is exactly the statement of UGC from §7.5.6.

Theorem 1.1 also states a stronger form: the alphabet is the 0/1 vector space F_2^s , every constraint is a translation, and the graph is simple, bipartite and unweighted. The paper says this is a known equivalent form of UGC (citing KKM0), so "stronger" is not "swapping in a different problem".

7.8.2 Other papers in the same family (family 102)

Paper	Pages	Conclusion
A Direct Proof of Optimal Max-Cut Hardness	38	Beating α_{GW} on Max-Cut is NP-hard
The Factor-Two Hardness Threshold for Vertex Cover	25	Vertex cover can't be approximated to any constant factor < 2
Constant-factor hardness of Min-UnCut	—	Every constant factor is NP-hard
Constant-factor hardness of directed feedback vertex set	—	Every constant factor is NP-hard

The repo's CONTENTS.md says these 4 papers are **independent direct reductions** that "use established PCP and Label Cover hardness results" and don't depend on the UGC paper. In other words, even if the UGC paper has an error, its two most famous consequences, Max-Cut 0.878 and vertex cover 2, have a separate independent path. The complexity theorist Dana Moshkovitz put it this way, in commentary that Aaronson posted as comment #2 on "The Mathocalypse" (it opens "Additional commentary from Dana on the UGC proof"): "The Unique Games Conjecture is true, though it wasn't really needed for the proof of its most famous corollaries", with optimal hardness of approximating Max-Cut and other constraint satisfaction problems as her examples in parentheses [R]. These 4 direct reductions are what she means.

In addition, **family 105** (*Perfect completeness for 2-to-1 games*) claims a proof of Khot's 2-to-1 conjecture (the perfect-completeness version). Footnote 1 of the UGC paper says it is "not an input to this proof".

7.8.3 Intuition for the method (I read only the introduction; this is a simplified version)

1. **Starting point:** Håstad's hardness result for "nearly satisfiable systems of parity equations" (each equation $x + y + z \equiv b \pmod{2}$).
2. **Why the 2-to-2 road got stuck at 1/2:** in the "matrix short code" of Barak–Kothari–Steurer, the probability that an honest proof's answer stays unchanged under one random rank-one perturbation is $(1 + 2^{(-\ell)})/2$, i.e. about 1/2. So completeness could only reach 1/2.
3. **The key new idea: a different kind of noise.** The paper constructs a nonlinear "encoder" C and a special noise μ : honest answers barely change under this noise (change probability $\leq$ an arbitrarily small p), but any "cheating", high-rank linear observation still notices the noise with probability at least 1/8. An analogy: design a jitter that barely affects correct readings but exposes forged ones. This lets completeness be pushed to $1 - \epsilon$ while keeping the structural tools that the 2-to-2 approach uses for soundness.
4. **Soundness (unsatisfiable formulas have value $\leq \delta$):** reuses the Grassmann expansion theorem of KMS 2018 (in the form of an "inverse short code theorem") to extract a strategy, then uses parallel repetition (Dinur–Steurer 2014) to push the success rate down to $\exp(-c \cdot k^{(1/3)})$, giving a contradiction.
5. **Finishing up:** remove weights and repeated edges to get a simple, unweighted, bipartite instance.

7.9 How far it has been verified

7.9.1 The Lean statement: much longer than the quasi-Riemann hypothesis, but still readable

The Lean statement of the quasi-Riemann hypothesis (§4) is a single line, because Mathlib already has the ζ function. UGC is different: Mathlib has **no** ready-made definitions of "unique game", "3SAT encoding" or "polynomial-time reduction to a graph", so the challenge file `ComparatorChallenges/UniqueGamesTheorem.lean` (239 lines) has to define all of these itself. The core theorem (full name `OAI.UniqueGamesTheorem.theorem11`); the translation below is from my reading of the Lean source [P]):

```
theorem theorem11 (ε δ : ℝ)
  (hε : 0 < ε) (hεhalf : ε < 1 / 2)
  (hδ : 0 < δ) (hδhalf : δ < 1 / 2) :
  Nonempty (Explicit.MachineOutputContract.BinaryGapReduction ε δ)
```

Word by word:

- `(ε δ : ℝ)`: ϵ and δ are real numbers.
- `hε`, `hδ`: assume $0 < \epsilon < 1/2$ and $0 < \delta < 1/2$.

- `Nonempty (BinaryGapReduction  $\epsilon$   $\delta$ )`: the conclusion is "there **exists** a reduction satisfying all of the conditions below".

`BinaryGapReduction  $\epsilon$   $\delta$` is a structure that requires all of the following at once:

Field	Meaning
<code>alphabet</code> , <code>dimension</code> , <code>coordinates</code>	Alphabet size $k \geq 2$, identified with F_2^s (0/1 vectors)
<code>construct : List Bool $\rightarrow$ Instance alphabet</code>	Takes a string of bits (the encoding of a 3SAT formula) and outputs a unique game
<code>simpleBipartite</code>	The output graph is a simple bipartite graph
<code>translations</code>	Every constraint is a translation $a(v) = a(u) + c$
<code>computation : Turing.TM2ComputableInPolyTime ...</code>	<code>construct</code> can be computed in polynomial time by Mathlib's multi-tape Turing machine
<code>finiteAlphabet</code>	The Turing machine's tape alphabet is finite
<code>completeness</code>	Input is a satisfiable formula $\Rightarrow$ there is a labeling satisfying a fraction $\geq 1 - \epsilon$
<code>soundness</code>	Input is not a satisfiable formula $\Rightarrow$ every labeling satisfies a fraction $\leq \delta$

A few details worth noting:

- Permutations are represented as "forward table + inverse table + two proofs that they are inverses", so they really are one-to-one.
- Repeated edges are counted by multiplicity; there are no hidden weights.
- If the input string can't even be parsed as a formula, it counts as "unsatisfiable", and the reduction must still output an instance with value $\leq \delta$. This is a stricter requirement, not a loophole.
- Restricting $\epsilon, \delta < 1/2$ doesn't affect the conclusion: UGC is precisely about small ϵ, δ .
- Completeness is written as $\geq 1 - \epsilon$ rather than $= 1$, which is correct: as §7.5.4 said, the fully satisfiable case is easy.

My judgment: this is a faithful encoding of UGC, in a stronger form. One outside fact the reader has to supply: the theorem says "reduces from 3SAT"; "therefore NP-hard" also needs the Cook–Levin theorem (3SAT is NP-complete). That step is not in this statement, but it is a textbook result and uncontroversial.

7.9.2 The proof itself

- Solution directory `OAI/Computability/UniqueGames/`: **478 Lean files, 186,782 lines** (counted by script).

- grep for `sorry`, `axiom`, `admit`, `native_decide`, `implemented_by`, `unsafe`: 0 hits. The Comparator config allows only three standard axioms (`propext`, `Quot.sound`, `Classical.choice`).
- Apart from Mathlib, the proof imports nothing from other directories in the repo. Since Mathlib has no PCP theorem, **the PCP theorem, Håstad's result, KMS's Grassmann expansion theorem, parallel repetition, and the polynomial-time Turing machine implementation** should all have been formalized from scratch within these 478 files. The subdirectory names (`PCP/` with 104 files, `Inverse/` with 39 files including `KMS...`, `Repetition/`, `Machines/` with 144 files) are consistent with this. $\triangle$ This is inferred from the directory structure and import relations; I did not open each one to check.
- The direct proofs for Max-Cut and vertex cover also have Lean: `MaxCut/` has 372 files, 186,763 lines; `VertexCover/` has 354 files, 66,335 lines. In the Max-Cut statement, α_{GW} is defined directly as the `sInf` of $2 \cdot \arccos \rho / (\pi(1 - \rho))$, which is exactly the formula from §7.3.3 ($\arccos \rho = \theta$).
- **What wasn't done:** I did not compile it. The Comparator config for the UGC family has `enable_nanoda: false`, meaning no second, independent kernel re-check. Unlike the quasi-Riemann hypothesis (§4), I found no third party who has publicly re-run this proof $\triangle$. The repo's `formalization.yaml` itself says `review: status: unchecked`.

7.9.3 Human experts

- Lance Fortnow (blog post "Open No More", 2026-10-07): "if they hold up, we've seen more progress in TCS in the last 24 hours than in the previous three decades combined." [R]
- Mark Braverman (Princeton, Quanta 2026-10-07): "Math by press release is not that healthy for math." [R]
- Quanta's report says this proof "didn't undergo any human editing or review by independent experts". [R]
- Dana Moshkovitz (a complexity theorist, and Aaronson's wife): her views come from private texts and comments that Aaronson posted on his blog, not from a formal review [R]. In the post body: "Basically the paper is so horribly written that it's impossible to read it without AI help"; in comment #2: "the writeup is poorly written". Her description of the method: "The UGC proof invents a completely new bizarre code with a noise test" (this is the new encoder and noise of step 3 in §7.8.3). Later, in comment #94, Aaronson wrote: "Dana tells me she now mostly understands the proof of the UGC". Secondhand retellings in outlets such as Implicator word this differently $\triangle$; go by the original on Aaronson's blog.
- Khot himself: as of 2026-10-08 I found no public statement $\triangle$.
- **As of 2026-10-08, no human expert has publicly said they have read through and verified these 58 pages.**

7.9.4 Remaining risks

1. **Statement level:** I read it line by line and found no problems, but I am not a complexity theory expert. The things most worth an expert's re-check are the use of `TM2ComputableInPolyTime` (whether the input and output encodings are reasonable, and whether the polynomial really is independent of the input), and whether the 3SAT encoding has a loophole that makes the problem easier.
2. **Compilation level:** depends on third parties or re-running it yourself. No nanoda re-check.
3. **Understanding level:** this is the same situation as the quasi-Riemann hypothesis (§4). It may already be proved in the machine sense, but nobody can explain "why that nonlinear noise works". The paper is very hard to read, which will slow down human understanding.

7.10 What it means if true (and what it doesn't)

7.10.1 For theoretical computer science

1. **A large batch of conditional results become theorems:** Max-Cut's 0.878, vertex cover's 2, Raghavendra's "SDP is optimal for every CSP", and so on. The premise is removed from the hundreds of papers of the past 20 years that say "assuming UGC, then...".
2. **The map of approximation algorithms is basically complete:** for many problems, the gap between "achievable" and "not achievable" is closed. The 1995 GW algorithm is proved optimal.
3. **The subexponential algorithm is no longer a counterexample:** UGC being true and ABS's $\exp(n^\epsilon)$ algorithm coexist. This means the reduction from 3SAT to UG must produce an extremely large polynomial (the paper says the degree and constants of the polynomial depend only on ϵ , δ , but gives no concrete numbers).

7.10.2 What it does not mean

- **It does not resolve P vs NP.** Every conclusion is "X is NP-hard", i.e. "unless $P = NP$, X cannot be done". If someone proved $P = NP$ tomorrow, all these hardness results would become void. Fortnow says in the same blog post that these results don't bring us closer to resolving P vs NP [R].
- **It does not affect practical engineering.** Industry uses heuristics for graph partitioning and clustering, and on real data they often do far better than the worst-case 0.878. UGC is about **guarantees** in the **worst case**.
- **It does not overturn any known algorithm.** It only proves that "existing algorithms can't be improved further".

7.10.3 For AI

- This is the result in OpenAI's batch **closest to AI researchers' daily work:** its core objects (SDP relaxations, randomized rounding, PCP-style verification) are basic tools of computer

science.

- An association (**my speculation, not an established result**): the complexity-theoretic foundation of debate-style scalable oversight (Irving–Christiano–Amodei 2018, arXiv:1805.00899) is interactive proofs and PCP: "can a weak verifier, looking at only a little information, judge a strong prover's argument?" UGC characterizes the boundary where "a little noise makes efficient distinction impossible". For designing "weak overseer + strong model" protocols, it suggests there may be an essential computational-hardness difference between **perfect agreement** and **near agreement**.
- As with the quasi-Riemann hypothesis (§4), we again have the combination "a hard-to-read 58-page paper + 187,000 lines of Lean". The bottleneck of trust moves from "is the proof right" to "**was the statement encoded correctly**", which is exactly the spec review that §7.9.1 does.

7.11 Self-test

Try answering these after reading. If you can't, go back and reread the relevant section.

1. What exactly does "X is NP-hard" mean? Does it prove $P \neq NP$? (§7.1.2)
2. Why is random coloring at least a $1/2$ -approximation? (§7.2.2)
3. In the GW algorithm, 0.878 is the minimum of the ratio of which two quantities, and where is it attained? (§7.3.3)
4. Why does the "gap" in the PCP theorem imply hardness of approximation? (§7.4.2)
5. Why can a fully satisfiable unique game be solved in polynomial time, while 3SAT can't? (§7.5.4)
6. Why does the order "first ϵ , δ , then the alphabet size" in UGC matter? (§7.5.6)
7. What has the 2-to-2 theorem already proved, and what was still missing? (§7.7.2)
8. What does the **soundness** field in the Lean statement say? What happens if the input string isn't a valid formula encoding at all? (§7.9.1)
9. Why might the optimality of Max-Cut 0.878 still hold even if the UGC paper has an error? (§7.8.2)

Appendix: answers to the checks

- **Check 1:** A 4-cycle is bipartite; color it alternately black and white and all 4 edges are cut, so max cut = 4. A random coloring cuts $4 \times 1/2 = 2$ edges in expectation.
- **Check 2:** The first four constraints form the cycle $x_0 \rightarrow x_1 \rightarrow x_2 \rightarrow x_3 \rightarrow x_0$, with constant sum $2 + 3 + 1 + 4 = 10 \equiv 0 \pmod{5}$, so no contradiction. The fifth, $x_2 - x_0 \equiv 0$: along the cycle, $x_2 - x_0 = 2 + 3 = 5 \equiv 0 \checkmark$. So everything can be satisfied, e.g. $(x_0, x_1, x_2, x_3) = (0, 2, 0, 1)$. After changing it to $x_2 - x_0 \equiv 1$, the cycle $x_0 \rightarrow x_1 \rightarrow x_2 \rightarrow x_0$ sums to $2 + 3 - 1 = 4 \not\equiv 0$, so at least one

constraint is violated, and at most 4/5 can be satisfied (confirmed by the script by brute-force enumeration of all 625 labelings).

Appendix: Glossary

Term	One-line explanation
P / NP	Problems where you can quickly find the answer / problems where you can quickly check an answer
Reduction	Quickly translating problem A into problem B, so that a fast algorithm for B solves A
NP-hard	3SAT reduces to it; unless $P = NP$, it has no polynomial-time algorithm
3SAT	Boolean satisfiability with 3 variables per clause
Max-Cut	Split the vertices into two groups to maximize the edges crossing between groups
Vertex cover	Pick the fewest vertices that cover every edge
α -approximation	Always reaches α times the optimum in polynomial time
Semidefinite program (SDP)	An optimization problem that relaxes ± 1 variables to unit vectors; solvable in polynomial time
Random hyperplane rounding	Make one random cut to turn vectors back into black and white
α_{GW}	0.878567..., the guarantee of the GW algorithm
Integrality gap	The ratio gap between the relaxation's optimum and the true optimum
PCP theorem	Proofs can be rewritten in a form that can be verified by randomly spot-checking a constant number of bits
Gap	YES instances have value $\geq c$, NO instances have value $\leq s$, with $c > s$
Completeness / soundness	The fraction satisfiable in the YES case / the most that can be satisfied in the NO case
Label Cover	The standard output of the PCP theorem: a labeling problem with "many-to-one" constraints on each edge
Unique Game	A labeling problem where each edge constraint is a permutation (fix one end and the other is unique)
UGC	Telling apart unique games with value $\geq 1 - \epsilon$ from those with value $\leq \delta$ is NP-hard
2-to-2 theorem	"Half of UGC" (completeness about 1/2), proved in 2018
Subexponential time	Faster than any $2^{(n^c)}$ (for fixed $c > 0$), slower than polynomial

Appendix: Further reading

- **Ryan O'Donnell**, *Analysis of Boolean Functions* (Cambridge University Press, 2014; free online version <https://analysisofbooleanfunctions.net/>): Chapter 7 covers PCP and Håstad,

Chapter 11 covers "Majority Is Stablest" and KKM0. The best textbook for understanding the UGC hardness proofs. $\triangle$ Chapter numbers from memory; go by the table of contents.

- **Williamson & Shmoys**, *The Design of Approximation Algorithms* (Cambridge, 2011; free PDF <http://www.designofapproxalgs.com/>): Chapter 6 covers the GW algorithm, Chapter 16 covers hardness of approximation and UGC. $\triangle$ Chapter numbers from memory.
- **Subhash Khot**, "On the Unique Games Conjecture" (survey, CCC 2010): a survey by the person who proposed it. <https://cs.nyu.edu/~khot/papers/UGCSurvey.pdf>
- **Erica Klarreich**, "Approximately Hard: The Unique Games Conjecture" (Simons Foundation, 2011): popular science for a general audience.
- **Quanta**, "As AI Closed In on 'Unique Games' Proof, Researchers Raced to Beat the Machines" (2026-10-07): <https://www.quantamagazine.org/as-ai-closed-in-on-unique-games-proof-researchers-raced-to-beat-the-machines-20261007/>
- **Lance Fortnow**, "Open No More" (2026-10-07): <https://blog.computationalcomplexity.org/>
- **Scott Aaronson**, "The Mathocalypse" (2026-10-07): <https://scottaaronson.blog/?p=10169>
- Classic original references: Goemans–Williamson, *JACM* 42 (1995); Håstad, *JACM* 48 (2001); Khot, STOC 2002; KKM0, *SICOMP* 37 (2007); Khot–Regev, *JCSS* 74 (2008); Raghavendra, STOC 2008; Arora–Barak–Steurer, FOCS 2010 / *JACM* 62 (2015); Khot–Minzer–Safra, FOCS 2018.
- Original paper: <https://github.com/openai/math/blob/main/preprints/The-Unique-Games-Theorem-September-23-2026/paper.pdf>
- Direct Max-Cut proof: <https://github.com/openai/math/blob/main/preprints/A-Direct-Proof-of-Optimal-Max-Cut-Hardness-September-23-2026/paper.pdf>
- Lean statement:
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/UniqueGamesTheorem.lean>

THEORETICAL CS · FAMILY 103 (+ 109)

8. $L = RL = BPL$ (in small-memory computation, randomness doesn't help)

You only need these prerequisites: you know what $\log_2$ is, you know the phrase "an algorithm's running time grows with the input size n ", and you have seen matrix multiplication. Every new concept is explained before it is used. Every number in the text is either computed by script or has a source (scripts: `精讲/scripts/05_bpl_toy.py`, `05_bpl_spec.py`, implemented in numpy). The **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. Tags: **[Established]** = a published,

peer-reviewed result; **[OpenAI claim]** = a manuscript in the openai/math repository; **[Speculation]** = my own judgment; **[P]** = primary material read (paper, repository file or Lean source); **[R]** = secondary reporting (blogs, news); Δ = something I could not fully verify, or where sources conflict. Estimated reading time: 60–90 minutes. You can read it in three sittings: §8.1–8.3 (memory complexity and randomness), §8.4–8.6 (random walks, Reingold, the history of derandomization), §8.7–8.9 (OpenAI's result, its verification status, and what it means, ending with the integer-multiplication result from the same batch).

8.0 The bottom line: the one thing this chapter explains

Picture a machine with tiny memory: it can read the input as much as it likes, but its working memory only holds **a few pointers**. If this machine can flip coins, can it solve more problems?

This chapter explains the following causal chain:

The whole computation of a small-memory machine $\Leftrightarrow$ a random walk on a "configuration graph" $\Leftrightarrow$ acceptance probability = a sum of powers of a large matrix $\Rightarrow$ derandomization = computing this probability accurately with very little memory $\Rightarrow$ the difficulty is that the precision and the recursion depth needed to "compute it accurately" eat up the memory.

L is the set of problems that deterministic small-memory computation can solve. **BPL** is the set of problems solvable with "small memory + coin flips + a small allowed probability of error". Almost everyone believes $L = BPL$, but since 1999 the best unconditional result has only reached "BPL can be simulated deterministically with $(\log n)^{1.5}$ memory". That has always been a factor of $\sqrt{\log n}$ away from $\log n$.

OpenAI claims to have closed this gap completely: **$L = RL = BPL$** . Few people ever doubted the conclusion itself; the hard part is the proof. It comes with a Lean formalization, but as of 2026-10-08 no human expert has publicly said they have read and verified this 108-page paper.

After reading this chapter, you should be able to explain in your own words:

- why " $O(\log n)$ bits of memory" roughly equals "can store a constant number of pointers", and why such a machine must halt in polynomial time;
- why a random walk can solve undirected graph connectivity with tiny memory, and why the same method fails on directed graphs;
- why "seed length" is the central measure in this field, and what separates Nisan's $\log^2 n$ from the target $\log n$;
- what OpenAI's result would mean if true, and what it does **not** mean (it does not mean $P = BPP$, and it does not mean $L = NL$).

8.1 Space complexity: memory is a resource

8.1.1 Time and space

There are two rulers for measuring an algorithm:

- **Time:** how many steps it takes.
- **Space:** the maximum number of bits of **working memory** it uses at any one moment.

There is a key convention here: **the input does not count toward space**. The input sits on a "read-only tape"; the machine can read it back and forth but cannot write to it. Only the "work tape" (read-write scratch paper) counts as space. This is what lets us talk about "memory smaller than the input".

The engineering analogy is **streaming plus random access**: the data sits on a big read-only disk, you can seek to any position and read it, but your local RAM is only a few dozen bytes.

8.1.2 Logarithmic space L: room for a constant number of pointers

L (logspace) = the problems that a deterministic machine can decide with $O(\log n)$ bits of working memory.

Why $\log n$? Because a pointer that can point to any position in the input needs $\lceil \log_2 n \rceil$ bits. $O(\log n)$ bits means "a constant number of pointers or counters".

Input length n	$\log_2 n$ (one pointer)	$(\log_2 n)^{1.5}$	$(\log_2 n)^2$
10^3	10	31	99
10^6	20	89	397
10^9	30	163	894
10^{12}	40	252	1,589

(Computed by script; the unit is bits.) With an input of 10^9 bytes, L allows only tens to hundreds of bits of memory, about the size of a few registers. The $(\log n)^2$ and $(\log n)^{1.5}$ columns will come up again and again; they are the two milestones in the history of derandomization.

[Check 1] The input is a graph with 10^9 vertices. How many bits does it take to store one vertex ID? How many bits in total, roughly, to store "current vertex + target vertex + a step counter (counting up to at most n^3)"?

8.1.3 An important fact: a deterministic logspace machine must halt in polynomial time

The machine's complete state at a given moment is called its **configuration**: control state + the positions of all read/write heads + the contents of the work tape. The work tape has only

$c \cdot \log n$ bits, so it has at most $2^{(c \cdot \log n)} = n^c$ possible contents; the read head has n possible positions; the control state has a constant number of possibilities. So **the total number of configurations is polynomial in n** .

If a deterministic machine halts, it can never pass through the same configuration twice (otherwise it would loop forever). So its number of steps is at most the number of configurations, which means polynomial time.

This fact that "the number of configurations is polynomial" is the cornerstone of the whole chapter. Everything that follows can be seen as working on a polynomial-size "configuration graph".

8.2 What randomness can buy

8.2.1 Example: polynomial identity testing

You are given two polynomials written as formulas and asked whether they are identical. For example, $(x+y)^2$ and $x^2 + 2xy + y^2$.

You can settle this example by expanding by hand. The hard case is when the polynomial is written as a very compact formula or circuit that would have exponentially many terms once expanded (for example, an $n \times n$ determinant has $n!$ terms when expanded). In that case no fast **deterministic** algorithm is known to this day. This is the **PIT problem** (polynomial identity testing). [Established]

The randomized algorithm is simple: **plug in a few random numbers and check whether the two sides are equal.**

- If the polynomials are identical, any input gives equal values.
- If they are not identical, the difference of the two sides is a nonzero polynomial, and it has few roots. The Schwartz–Zippel lemma says: if you pick values at random from a set of size $|S|$, the probability that a nonzero polynomial of degree d evaluates to 0 is $\leq d/|S|$.

[Established]

Worked example: the false "identity" $(x+y)^2 = x^2 + y^2$. The difference of the two sides is $2xy$, a degree-2 polynomial. Pick x and y at random from $\{0, \dots, 99\}$: 10,000 combinations in total.

[Check 2] How many combinations make the two sides exactly equal (that is, "fool" the random test)?

Computed by script: 199, or 1.99%. The Schwartz–Zippel bound is $d/|S| = 2/100 = 2\%$, which matches. With one random substitution, the chance of being fooled is under 2%; with 10 independent substitutions, it is under $2\%^{10} \approx 10^{-17}$.

8.2.2 P and BPP

- **P**: problems solvable in deterministic polynomial time.
- **BPP**: problems solvable in polynomial time + random bits, with error probability $\leq 1/3$ on every input.

The number $1/3$ does not matter, because you can **amplify**: run independently k times and take a majority vote.

Repetitions k	Majority-vote error probability (single-run success $2/3$)	Bits needed for the counter
1	3.3×10^{-1}	1
11	1.2×10^{-1}	4
51	6.9×10^{-3}	6
101	2.7×10^{-4}	7
201	5.5×10^{-7}	8

(Computed by script, summing binomial tails.) Note the last column: amplification only needs a counter, **so it works for small-memory machines too**. That is why it doesn't matter whether the definition of BPL uses a threshold of $2/3$ or 0.51 .

The mainstream conjecture is $P = BPP$, meaning randomness gives no fundamental advantage. The evidence is strong: Impagliazzo–Wigderson 1997 proved that as long as certain problems require exponential-size circuits (a widely believed "hardness assumption"), $P = BPP$. [Established] But proving $P = BPP$ **unconditionally** is considered extremely hard: Kabanets–Impagliazzo 2004 proved that even just derandomizing PIT would imply certain circuit lower bounds that we currently have no idea how to prove. [Established]

This is the "hardness vs. randomness" framework: **derandomization and proving lower bounds are almost the same thing**. The time version is blocked by this wall. The space version has no known barrier of the same kind, which is why $L = BPL$ has long been seen as "the derandomization conjecture most likely to be proved unconditionally".

8.3 RL and BPL: small memory + coin flips

8.3.1 Definitions

The machine gains one ability: at each step it can get a fresh fair coin. The constraints are:

- $O(\log n)$ bits of working memory;
- it halts in polynomial time for **all** coin sequences;

- **coins are single-use**: once read, a coin is gone. To use it again, the machine must store it on its work tape itself, and the work tape has only $\log n$ bits.

The last constraint is crucial. The time version, BPP, can store the random string and reread it as often as it likes; a small-memory machine cannot. It can only "use the coins as it flips them", like an online, single-pass stream.

Two ways of making errors:

- **RL** (one-sided error): when the answer is "no", it always rejects; when the answer is "yes", it accepts with probability at least $1/2$.
- **BPL** (two-sided error): when the answer is "yes", it accepts with probability at least $2/3$; when the answer is "no", it accepts with probability at most $1/3$.

Known inclusions [Established]:

$L \subseteq RL \subseteq BPL \subseteq L^{(3/2)}$ (Saks–Zhou 1999)
 $RL \subseteq NL \subseteq L^2$ (Savitch 1970)

NL is "nondeterministic" logspace (the machine can "guess"). Reachability in directed graphs is the representative NL problem.

8.3.2 The whole computation = a random walk on a graph

Use the fact from §8.1.3: a randomized machine has polynomially many configurations on input x . Draw each configuration as a node; if flipping one coin can take configuration u to configuration v , draw an edge from u to v with weight $1/2$.

So **the machine's run on x = a random walk on this configuration graph, starting from the start node; the acceptance probability = the probability that the walk reaches an "accept" configuration within the time limit.**

Worked example: a toy machine with only 3 states, flipping one coin per step:

- State 0: heads goes to 0, tails goes to 1
- State 1: heads goes to 2, tails goes to 0
- State 2: absorbing (stays at 2 forever); it represents "accept"

Write this as a transition matrix T , start from $v = (1, 0, 0)$, and at each step compute $v \leftarrow vT$:

Step	$P(\text{at } 0)$	$P(\text{at } 1)$	$P(\text{at } 2)$
1	0.5	0.5	0
2	0.5	0.25	0.25

Step	P(at 0)	P(at 1)	P(at 2)
3	0.375	0.25	0.375
4	0.3125	0.1875	0.5

(Computed by script; it also enumerated all $2^4 = 16$ coin sequences, and exactly 8 of them reach state 2, matching the 0.5 from the matrix.)

[Check 3] By hand only, verify the distribution after step 2, (0.5, 0.25, 0.25).

This turns the problem into linear algebra: **acceptance probability = some powers of the transition matrix applied to the start vector**. The OpenAI paper starts the same way: it turns the machine into a matrix S , and the acceptance probability is one component of $(I - S)^{-1}e$, which is equivalent to $\sum_j S^j e$ (§8.7.2).

8.3.3 Why "just compute it" is not in L

If it is just matrix powers, why not compute it deterministically? The problem is **memory**:

- The matrix is $n^c \times n^c$. It cannot be stored, so each entry has to be computed on demand.
- When computing T^t by "repeated squaring", each entry of T^t depends on entries of $T^{(t/2)}$, each entry of $T^{(t/2)}$ depends on entries of $T^{(t/4)}$, and so on. The recursion has $\log t \approx O(\log n)$ levels, and each level has to remember where it is ($O(\log n)$ bits), plus the bits of numerical precision.
- Total space $\approx$ levels $\times$ per level = $O(\log n) \times O(\log n) = O(\log^2 n)$.

This is the intuition behind Borodin–Cook–Pippenger 1983's result that "randomized logspace $\subseteq$ deterministic $\log^2$ space", and it has the same structure as the "recursion stack" in Savitch's theorem. [Established] **Forty years of work in this field have been about pushing down this "levels $\times$ per level"**.

8.4 Undirected connectivity: the classic win for random walks

8.4.1 The problem

USTCON (undirected s-t connectivity): given an undirected graph and two vertices s and t , are they connected?

BFS/DFS makes this easy, but they need a "visited" array, which takes n bits. What do you do with only $O(\log n)$ bits?

8.4.2 The random walk algorithm (AKLLR 1979)

In 1979, Aleliunas, Karp, Lipton, Lovász and Rackoff gave a minimal algorithm [Established]:

```

v ← s
repeat N steps:
  if v = t, answer "connected"
  v ← a random neighbor of v
answer "not connected"

```

It only needs to remember "the current vertex v " and "a step counter": $O(\log n)$ bits. If s and t are not connected, it never reaches t , so it **never turns a "no" into a "yes"**; if they are connected, then for large enough N it reaches t with high probability. This is exactly the one-sided error form of RL.

The key bound they proved: on a connected graph, the expected number of steps for a random walk to visit every vertex (the **cover time**) is $\leq 2|E|(n-1)$, where $|E|$ is the number of edges. Since $|E| \leq n^2$, this is on the order of $\leq n^3$. Taking $N = 4n^3$ and applying Markov's inequality pushes the failure probability below $1/2$.

8.4.3 Measured by script

Small example: 6 vertices, edges 0–1, 0–2, 1–2, 2–3, 3–4, 4–5 (a triangle with a tail attached). Walk from 0 to 5.

One actual simulated trajectory (script, random seed 1):

```

0 1 0 2 0 2 1 2 3 4 3 2 1 0 2 1 0 2 1 0 1 2 0 1 0 1 2 3 2 1 0 1 2 1 0 2 0 1 2 1 0 2 3 2
0 2 0 2 3 4 3 4 5

```

It took 52 steps. It circled around the triangle for a long time, because from 2 it has only a $1/3$ chance of heading into the tail, and once at 3 it has a $1/2$ chance of going back.

Averaged over 20,000 simulations: about 29.1 steps from 0 to 5, and about 30.4 steps to cover the whole graph. The upper bound $2|E|(n-1) = 2 \cdot 6 \cdot 5 = 60$; the measurements are within the bound.

Hitting times on several kinds of graphs:

Graph	n	Measured / exact expected steps	Comparison
Path (from one end to the other)	10	81.7 (simulated)	exact value $(n-1)^2 = 81$
Path	40	1,523.7 (simulated)	$(n-1)^2 = 1,521$
"Lollipop" (complete graph on $n/2$ vertices + a tail of $n/2$ vertices)	40	8,019 (solving a linear system, exact)	$n^3 = 64,000$, ratio 0.125

Graph	n	Measured / exact expected steps	Comparison
Lollipop	80	64,039 (exact)	$n^3 = 512,000$, ratio 0.125

(Computed by script.) The lollipop graph is the worst case for random walks: the walk spins around inside the complete graph, rarely enters the tail, and easily falls back out once it does. Brightwell–Winkler 1990 proved that the maximum hitting time on an n -vertex graph is about $(4/27)n^3 \approx 0.148n^3$ (attained when the clique takes $2/3$ of the vertices).

[Established; $\triangle$ constant quoted from memory of the literature, not checked against the original] Our "half and half" version gives $0.125n^3$, the same order of magnitude. So **n^3 steps are generally enough, and you basically can't do with fewer.**

8.4.4 Why the same method fails on directed graphs

Consider a directed graph with vertices $0, 1, \dots, n$. From vertex i there are two edges: one to $i+1$ and one back to 0 . The target is n .

A random walk from 0 must choose "forward" n times in a row to get there. The expected number of steps is $2^{(n+1)} - 2$:

n	Expected steps (exact, solving a linear system)
5	62
10	2,046
20	2,097,150
30	2,147,483,646

(Computed by script; matches the formula exactly.) **Exponential.** So random walks cannot solve directed reachability. Directed reachability is NL-complete; whether it is in L ($L = NL?$) is still open, and it is **unrelated to this chapter's result**.

[Check 4] Compute this expected value by hand for $n = 2$. Hint: let h_i be the expected number of steps to reach n starting from i ; $h_n = 0$, and $h_i = 1 + \frac{1}{2}h_{i+1} + \frac{1}{2}h_0$.

Why undirected graphs are nice: an undirected edge can be walked back along, so the random walk cannot get stuck in a "one-way trap". In linear-algebra terms, the transition matrix of an undirected graph is symmetric (after normalization), its eigenvalues are all real, and its "mixing" speed is controlled by the spectral gap, which is the star of §8.5.

8.5 Reingold 2005: undirected connectivity is in L (SL = L)

8.5.1 History

After AKLLR, people tried to remove the randomness from USTCON:

- Nisan–Szemerédi–Wigderson 1992: $O(\log^{1.5} n)$ space [Established]
- Armoni–Ta–Shma–Wigderson–Zhou 2000: $O(\log^{4/3} n)$ space [Established]
- **Reingold 2005**: $O(\log n)$ space, i.e. $USTCON \in L$. [Established] STOC 2005, JACM 2008, Gödel Prize.

SL (symmetric logspace) is the class for which USTCON is complete, so this result is also written **SL = L**. It was a milestone in "removing the randomness from a specific problem", but only for one structure: undirected graphs.

8.5.2 Expanders: graphs that are "very well" connected

An **expander** is a graph with very small degree (each vertex has only a constant number of edges) that is nonetheless extremely well connected: any small set of vertices has many edges going out of it.

The measure is the **spectral gap**: the largest eigenvalue of the random-walk transition matrix is 1, the second largest is λ_2 , and the gap = $1 - \lambda_2$. The larger the gap, the faster the random walk "forgets" its starting point. Roughly, the number of steps needed to mix $\propto (\log n) / \text{gap}$.

n	Gap of the cycle (degree 2)	1/gap	Gap of a random 4-regular graph	1/gap
64	0.00482	208	0.162	6.2
256	0.00030	3,320	0.138	7.2
1024	0.00002	53,122	0.134	7.4

(Computed by script, eigenvalues via numpy; the random 4-regular graph is built as the union of two random permutations.) The cycle's gap falls to 0 like $1/n^2$, while the random 4-regular graph's gap stays at a constant (the theoretical optimum is $1 - 2\sqrt{3}/4 \approx 0.134$, the Ramanujan bound, and the measurement is already right up against it).

An expander has diameter only $O(\log n)$. So if the graph is a constant-degree expander, you can **enumerate all paths of length $O(\log n)$** : each step has only a constant number of choices, a path can be described in $O(\log n)$ bits, and there are $d^{O(\log n)} = \text{polynomially}$ many paths in total. A counter can try them one by one, with no randomness needed.

8.5.3 Reingold's idea: "rebuild" any graph into an expander

The difficulty is that the input graph is not necessarily an expander. Reingold repeatedly rebuilds it with two operations (each connected component is rebuilt separately):

1. **Squaring** (powering): in the new graph, two vertices are adjacent if and only if there is a path of length 2 between them in the original graph. The spectral gap roughly doubles (λ_2 becomes λ_2^2), at the cost of squaring the degree.
2. **The zig-zag product** (Reingold–Vadhan–Wigderson 2000): "replace" each vertex with a small fixed expander, which brings the degree back down to a constant while losing only a constant factor in the gap. [Established]

Alternating these for $O(\log n)$ rounds grows the gap from $1/\text{poly}(n)$ to a constant.

Why it doesn't eat memory: the rebuilt graph is **never stored explicitly**. To look up "the i -th neighbor of vertex v in the new graph", you recursively send queries to the graph one level up. The key is that the zig-zag design makes each level cost only a **constant** number of extra bits (rather than $\log n$ bits). $O(\log n)$ levels $\times O(1)$ bits = $O(\log n)$.

Compare with §8.3.3: in "levels $\times$ per level", Reingold pushed "per level" from $\log n$ down to a constant. Engineering analogy: multi-level virtual address translation where each level only has to remember an offset of a few bits rather than a full address.

Limitation: both zig-zag and squaring rely on the graph's symmetry (undirectedness). The configuration graph of a general BPL computation is directed. Reingold–Trevisan–Vadhan 2006 extended the approach to special cases such as "regular directed graphs", but did not cover the general case. [Established]

8.6 Derandomizing general BPL: stuck near $3/2$ for 27 years

8.6.1 Pseudorandom generators and "seed length"

The most naive way to derandomize: **try every random string and count the fraction that accept**. A BPL machine reads $\text{poly}(n)$ random bits, so there are $2^{\text{poly}(n)}$ random strings. Far too many.

The idea of a **pseudorandom generator** (PRG): use a short **seed** to generate a long string of bits that "looks random". "Looks random" is relative to some class of observers: for any small-memory machine, the acceptance probability with the pseudorandom string is almost the same as with a truly random string.

Given a PRG with seed length s , derandomization is: enumerate all 2^s seeds, run once on each, and tally. The space needed is about $s + O(\log n)$.

Seed length	Number of seeds to enumerate	Derandomization result
$\text{poly}(n)$ (truly random)	$2^{\text{poly}(n)}$	useless
$O(\log^2 n)$ (Nisan 1992)	$n^{O(\log n)}$ (quasi-polynomial)	$\text{BPL} \subseteq L^2$

Seed length	Number of seeds to enumerate	Derandomization result
$O(\log n)$ (the dream)	$\text{poly}(n)$	$L = \text{BPL}$

So **the whole problem boils down to: for small-memory machines, can we build a PRG with seed length $O(\log n)$?** This question went unsolved for more than 30 years.

8.6.2 Nisan's generator (1992): why $\log^2$

Nisan's core observation: **a small-memory machine cannot remember which random bits it has used, so they can be safely "reused".** [Established]

The construction (simplified):

- Seed = a random string x of $O(\log n)$ bits, plus k random hash functions $h_1, \dots, h_k$, each of $O(\log n)$ bits.
- Level 1: output $x, h_1(x)$.
- Level 2: output $x, h_1(x), h_2(x), h_2(h_1(x))$.
- Each added level doubles the output length.

Intuition: after the machine has read the first half, its state is only $O(\log n)$ bits, so it has almost "forgotten" the specific contents of the first half. So the second half can be replaced by "the first half passed through a hash", and the machine can't tell.

To output $\text{poly}(n)$ bits takes $k = O(\log n)$ levels, each with an $O(\log n)$ -bit hash function, so the total seed length is $O(\log n) \times O(\log n) = \mathbf{O(\log^2 n)}$. "Levels $\times$ per level" again.

Nisan 1994 then proved that this simulation can achieve polynomial time and $O(\log^2 n)$ space at the same time ($\text{BPL} \subseteq \text{SC}$). [Established] (Note: Borodin–Cook–Pippenger 1983 had already obtained $\log^2$ space with the matrix method, but not polynomial time.)

8.6.3 Saks–Zhou 1999: $3/2$

Saks and Zhou cut Nisan's $\log n$ levels into $\sqrt{\log n}$ segments of $\sqrt{\log n}$ levels each. Each segment uses Nisan's method to approximately compute a matrix power, and then the low-order bits of the result are **randomly rounded** away. This makes each segment's result almost independent of which particular hash functions were used, so **all segments can share the same set of hash functions**. [Established; this description is heavily simplified, [△](#) see the original for details]

Space: one set of hash functions, $\sqrt{\log n} \times \log n$, plus the recursion, $\sqrt{\log n}$ segments $\times$ $\log n$ per segment, for a total of $\mathbf{O(\log^{1.5} n)}$.

8.6.4 Later progress, and the stall

Year	Result	Deterministic space
1983	Borodin–Cook–Pippenger	$\log^2 n$
1992	Nisan PRG	$\log^2 n$
1999	Saks–Zhou	$\log^{1.5} n$
2021	Hoza (weighted pseudorandomness)	$\log^{1.5} n / \sqrt{(\log \log n)}$
2026	OpenAI (claimed)	$\log n$

(The Hoza bound is based on [19, Theorem 1.6 and Corollary A.5] as cited in §1.1 of the OpenAI paper, and on my memory; the deck author's research notes originally flagged it $\triangle$, and this time I checked it against the paper's citation; the two agree.)

Hoza 2021 saved only a factor of $\sqrt{(\log \log n)}$, which for $n = 10^9$ is only about 2.2×. The exponent for general BPL has been stuck near 3/2 for 27 years.

Along the way there were many wins on "special structures": Reingold (undirected graphs), Reingold–Trevisan–Vadhan (regular directed graphs), Ahmadinejad et al. (computing random-walk probabilities to high precision on Eulerian graphs), Cohen–Doron–Sberlo–Ta-Shma (long products of random matrices). [Established] In addition, Cheng–Hoza 2022 proved that a logspace-enumerable "hitting set" (which only guarantees finding one accepted input, and is weaker than a PRG) is enough to imply $L = BPL$. [Established] All of this shows that the "shape" of the problem had been mapped out very clearly; only the last step was missing.

8.7 What OpenAI claims

8.7.1 The main result

The paper is *Exact Derandomization of Logarithmic Space: $L = RL = BPL$* (2026-09-23, 108 pages, authorship credited to "OpenAI", result family 103 in the repository) [P]. The abstract is a single sentence:

"We prove $L = RL = BPL$, resolving the derandomization problem for polynomial-time randomized logarithmic space."

In plain words: every randomized computation that is "polynomial time + logarithmic space + bounded error probability" can be simulated in deterministic logarithmic space.

The paper also gives a **stronger quantitative version** (Theorem 1.2): for any such randomized machine M and input x , the acceptance probability can be estimated deterministically to precision $2^{(-q)}$, using $O(\log n + q)$ space and $\text{poly}(n) \cdot 2^{O(q)}$ time. When $q = O(\log n)$, this means precision of any inverse polynomial can be computed in

logarithmic space. This is much stronger than "deciding whether the probability is above $2/3$ or below $1/3$ ": it can even estimate probabilities that have "no gap".

In addition, Section 13 gives a **compiler**: given the source program of a randomized machine and its time/space bounds, it outputs a deterministic decider together with explicit resource bounds.

8.7.2 Intuition for the method (I only read the paper's introduction, §1–§2; what follows is a simplified version)

1. **Turn it into matrix inversion** (paper §2). Write the configuration graph of the machine on input x as a matrix S (each configuration carries a timestamp, and edges only point to later times, so S is "strictly forward" and nilpotent). The acceptance probability is $p_0 =$ the value of $(I - S)^{-1}e$ at the start node, where e marks the accepting configurations. This is the same thing as the toy example in §8.3.2.
2. **"Correction and copy hierarchy"** (paper §4). The authors use a sequence of transformations to replace S with new matrices $C_1, C_2, \dots$. Each step "moves" part of the transition probability into a reward vector (using the identity $I - D = (I + E)(I - C)$), and copies some vertices to spread out large matrix entries. After $O(\log n)$ stages, reading off the accumulated reward along a single "main copy" path gives the acceptance probability with inverse-polynomial error.
3. **Estimate using only $O(\log n)$ random bits** (paper §5–§9). These matrices are not stored explicitly; instead they are estimated by sampling with a **shared finite random environment** σ (an encoding of length $O(\log n)$). The paper proves that on more than $3/4$ of the σ 's, the estimate is accurate; and for **every** σ (including the ones where the estimate is inaccurate), the estimation procedure halts within $O(\log n)$ space.
4. **Enumerate + take the median** (paper §12). σ has only $O(\log n)$ bits, so there are only polynomially many σ 's in total. Enumerate them one by one, take the median, and compare it with $1/2$. This step is the "enumerate the seeds" idea from §8.6.1, except that instead of building a general-purpose PRG, it squeezes the randomness down to $O(\log n)$ bits for this one specific estimator.
5. **Keep the recursion from eating memory** (paper §9–§11). Use short **fingerprints** (polynomial evaluation + affine hashing) instead of full vertex addresses for comparisons; use a shared "catalytic" bit vector; use Cook–McKenzie's "walk around the tree's contour" trick to recover incoming edges without a stack. The data kept by each recursive call is held within $A(w - w') + D(t - j)$ bits, and this **telescopes** along the call chain, for a total of $O(\log n)$.

Compare step 5 with §8.3.3 and 5.3: it is still a "levels $\times$ per level" problem, and the solution is to tie the number of bits each level spends to "how much the recursion budget drops", so

the whole call chain adds up to no more than one $O(\log n)$ total budget. Engineering analogy: a recursive call stack in which the size of each frame is bounded by a global quota, and the quota shrinks as you go down the stack.

Two details that surprised me:

- The mixing bound for "conditional averaging" in §6 of the paper relies on a result of Shalom 1999 about **Kazhdan's property (T)** (a deep property from group theory). [OpenAI claim] Group-theoretic tools showing up in small-memory derandomization is uncommon in the existing literature. [Speculation]
- The "shared catalytic bit vector" is reminiscent of recent work on **catalytic computing**: borrow a piece of memory that is already full of data, use it, and give it back exactly as it was. Cook–Mertz used it to bring the space for the tree evaluation problem down to $O(\log n \cdot \log \log n)$, and Ryan Williams 2025 built on this to prove that "a computation running in time t can be simulated in about $\sqrt{t \log t}$ space". [Established] If OpenAI's proof holds, catalytic tricks may once again be a key component. [Speculation, ⚠ I did not read the details of the paper's §11]

8.8 How far it has been verified

8.8.1 The Lean statement: you can read it yourself

The last line of `lean/ComparatorChallenges/LogspaceEquality.lean` is [P]:

```
theorem exact_logarithmic_space_derandomization :  
  L = RL ∧ RL = BPL
```

Word-by-word translation:

- `L = RL ∧ RL = BPL`: L equals RL, **and** RL equals BPL.

This line is very short; the real content is all in the **definitions**. Unlike the quasi-Riemann hypothesis (§4), the L, RL and BPL here are **not** standard Mathlib definitions (Mathlib has no such complexity classes). OpenAI wrote them itself in the challenge file. So the definitions are what need reviewing. Here is what I found, reading them one by one:

Definition	What it says	Matches the textbook?
<code>Machine q w h</code>	A finite-state machine with h read-only input heads and w binary work tapes; reads one coin b per step	Yes, a standard multi-tape Turing machine
<code>readInput</code>	The input has end markers at both ends	Yes

Definition	What it says	Matches the textbook?
<code>step</code>	A configuration that has already produced output (halted) no longer changes	Yes
<code>spaceThrough</code>	Up to time t , the sum over work tapes of the number of distinct cells visited	Yes, and fairly strict (blank cells that were visited also count)
<code>LogSpace</code>	There exists $c > 0$ such that for all inputs, all coin sequences and all times, $\text{space} \leq c \cdot \lceil \log_2(n+2) \rceil$	Yes
<code>Deterministic</code>	The transition function does not depend on the coin	Yes
<code>acceptanceProbability</code>	At clock time t , the fraction of the 2^t coin strings that stop in "accept"	Yes
<code>L</code>	Deterministic + logspace + gives the correct answer at some time on every input	Yes (no time bound is written, but §8.1.3 shows that a halting deterministic logspace machine is automatically polynomial-time)
<code>RL</code>	Logspace + halts within $c \cdot (n+2)^k$ steps for all coin sequences + if yes, probability $\geq 1/2$, otherwise probability = 0	Yes
<code>BPL</code>	Same as above, with thresholds $2/3$ and $1/3$	Yes

My judgment: they match the standard definitions. One choice worth noting is "halts in polynomial time for all coin sequences". This is standard BPL; if you drop the time bound you get a different class (§1.1 of the paper also specifically stresses the difference between a "space bound" and "simultaneous time and space bounds"). This agrees with the conclusion in the author's research notes.

8.8.2 The Comparator configuration

`LogspaceEquality.json` locks **20 definitions** (from `Word` to `BPL`; I counted) [P]: the input, language and coin tape; head movement; the machine's initial configuration, step and run; the space count; logspace, determinism, halting and deciding; the acceptance probability; the polynomial clock; and L, RL and BPL themselves. This means the solution file cannot quietly swap out any of them, for example by replacing "space" or "acceptance probability" with a version that is easier to prove things about. The only allowed axioms are Lean's three standard axioms (`propext`, `Quot.sound`, `Classical.choice`).

Note: this configuration has `"enable_nanoda": false`, meaning there is **no requirement for re-checking by a second, independent kernel**. (Only 2 of the 416 configurations in the

whole repository turn on nanoda; this is from the author's count over all configurations in the repository.)

8.8.3 The proof itself (what I found, [P])

- Solution directory `lean/OAI/Computability/Logspace/`: **39 files, about 43,000 lines** (commit fd4aeeb, counted by script with `wc -l`).
- `grep` for `sorry | axiom | admit | native_decide | implemented_by | unsafe`: **0 hits**.
- Apart from Mathlib, it imports nothing from other `OAI.*` directories in the repository. So there is no "custom code hidden in a dependency library" problem; in this respect it is cleaner than the quasi-Riemann hypothesis.
- The top-level proof is only three lines: `L ⊆ RL` (trivial), `RL ⊆ BPL` (trivial, amplify a bit), `BPL ⊆ L` (all the difficulty is here).
- **What I didn't do:** I did not compile it, and I did not run Comparator. $\triangle$ **I found no public record of any third party rerunning Comparator on this family** (the quasi-Riemann hypothesis (§4) had an independent re-check by Goldblatt; this family has none).
- The paper's **compiler** and **explicit running-time bounds** are not in the Lean statement. The scope document `lean/docs/103.md` says verbatim: "The paper's explicit compiler and numerical running-time bounds are not separately stated in this Comparator target." Theorem 1.2 (arbitrary-precision estimation) is not formalized either.

8.8.4 Human experts (as of 2026-10-08)

- **Scott Aaronson** (blog post "The Mathocalypse", 2026-10-07): "L=BPL (i.e., probabilistic logspace and deterministic logspace are the same thing), one of the great derandomization conjectures short of P=BPP." He also said: "its truth was never in serious doubt, there was a whole subcommunity focused on proving this." [R]
- **Lance Fortnow** (blog post "Open No More", 2026-10-07) lists it as "A full derandomization of randomized log space", and says: "And none of these results get us any closer to settling $P \vee NP$." [R]
- **Ran Raz** (Princeton, as reported by *Scientific American* on 2026-10-08): "We have strong reasons to believe that, but we don't know how to prove it", and, earlier, "I didn't see any directions that were promising". [R, $\triangle$ the quotes were relayed through a web-fetch tool; check the original for context (whether he meant $L = BPL$ itself or derandomization more generally)]
- **No human expert has publicly said they have read and verified the 108 pages.**
Aaronson's comment on the whole batch of results: "it also appears that no human has understood just about any of these proofs yet." [R]

8.8.5 An inconsistency worth flagging

△ Fortnow's "OpenAI Math: TCS Selection" page (lance.fortnow.com, dated 10-07) **shows no Lean marker** for number 103, while the repository at commit fd4aeeb (10-07 22:20 PT) contains the full Lean challenge and solution. history.md says 6 formalizations were "added" on 10-07, so one possible explanation is that 103's Lean was added that day and Fortnow's page was based on an earlier snapshot. I could not confirm which 6 they were (the repository has only one merge commit, so per-file history is not visible). This judgment is based only on a web-fetch tool's summary; I did not see the page myself.

8.8.6 How to think about the trust structure here

Compared with the quasi-Riemann hypothesis: **the definitions are self-written**, but they are only 142 lines, and I read them one by one and they match the standard; **the proof is self-contained** (it depends only on Mathlib), with no large-scale patches; **the conclusion is widely believed**, so "the conclusion is right" cannot be turned around into evidence that "the proof is right". Nearly all the credibility rests on the Lean kernel and on "whether anyone has actually compiled it successfully" (△ no third-party report yet).

8.9 What it would mean if true (and what it would not)

8.9.1 For complexity theory

1. The first natural computational resource class to be "fully derandomized"

unconditionally. Within "logarithmic space + polynomial time", flipping coins is of no use at all. This answers a question stated explicitly back in AKLLR 1979.

2. **It does not imply $P = BPP$.** The time version is blocked by the "derandomization $\Rightarrow$ circuit lower bounds" wall (§8.2.2). The proof of the space version goes around that wall but does not tear it down. A deterministic polynomial-time algorithm for PIT is still unknown.

3. **It does not imply $L = NL$, and it is even further from P vs NP .** Fortnow puts it bluntly: "none of these results get us any closer to settling $P \vee NP$ " (§8.8.4). It is still unknown whether directed reachability (§8.4.4) can be solved in logarithmic space. The new picture of inclusions is:

$$L = RL = BPL \subseteq NL \subseteq L^2$$

1. **Theorem 1.2 may be more useful than 1.1.** Being able to estimate random-walk probabilities to high precision in logarithmic space means that many small-space linear algebra problems (for example, Laplacian solving, or stationary distributions of Markov chains) may get new deterministic algorithms. [Speculation]

2. **The technique may matter more than the conclusion.** Aaronson says the conclusion was "never in serious doubt", so what the field really wants is the **proof method**, especially its

"low-space, high-precision linear algebra": computing quantities like matrix powers and $(I - S)^{-1}$ to very high precision with only $O(\log n)$ bits of memory (§8.7.2). Unfortunately, nobody understands it yet. If the proof holds but nobody can extract reusable ideas from it, that would be a very strange situation: the problem is closed, but the field learned nothing. [Speculation]

8.9.2 For engineering

Almost no direct impact on real systems: the randomized algorithms used in practice (hashing, sampling, Bloom filters) are already fast enough, and most of them are not in the extreme "logarithmic space" setting anyway. The theorem only guarantees that a deterministic replacement algorithm **exists**; the polynomial's degree and constants are likely astronomical. [Speculation] The implications for streaming algorithms are also limited: the streaming model can read the input only in a single pass, while L allows repeated reads.

8.9.3 For AI

- Like the other headline results, this is a case of "AI produces a proof of a recognized hard open problem, with a machine check attached".
- What sets it apart is that **everyone believes the conclusion; the hard part is the proof**. Problems of this kind may suit AI especially well: the direction is clear, there's no need to "guess the right conclusion", and what's needed is to get every step of a 108-page complex construction right, which Lean can check step by step. [Speculation]
- The lesson for AI safety research is the same as in §17 of this deck: once you have Lean, the bottleneck of trust moves from "is the proof right?" to "**were the statement and definitions written correctly?**" This family's definitions are self-written, so the table in §8.8.1 is a spec review. It is the same kind of work as auditing the behavioral spec of an AI agent.

8.9.4 Side note: integer multiplication below $n \log n$, from the same batch (family 109)

On the same day (2026-09-23) there was another theoretical CS paper, asking a more everyday question: **how many steps does it take to multiply two n -digit integers?** It uses the same model of computation as §8.8.1, the multi-tape Turing machine, so it is covered here.

Year	Who	Steps	Status
—	Schoolbook method	n^2	[Established]
1962	Karatsuba	$n^{1.585}$ (split one multiplication into 3 half-length multiplications instead of 4; $\log_2 3 \approx 1.585$)	[Established]

Year	Who	Steps	Status
1971	Schönhage–Strassen	$n \log n \log \log n$ (using the fast Fourier transform), and conjectured that $n \log n$ is optimal	[Established]
2019	Harvey–van der Hoeven (arXiv 2019, <i>Annals of Mathematics</i> 2021)	$O(n \log n)$	[Established]
2026	OpenAI	$O(n (\log n)^{(1-2^{-(182)})})$	[OpenAI claim]

What OpenAI claims [P]: the abstract of the paper *Integer multiplication below $n \log n$* (2026-09-23, 73 pages): "We give a deterministic algorithm that multiplies two n -bit integers in $O(n(\lg n)^{(1-\kappa)})$ worst-case time, with $\kappa = 2^{-(182)}$, on one fixed finite-alphabet Turing machine with a fixed finite number of one-dimensional tapes." In other words, on multi-tape Turing machines it saves a further factor of $(\log n)^{(2^{-(182)})}$ below $n \log n$. The paper itself admits: "The constants and thresholds in the construction are extremely large; the algorithm is intended to establish a bit-complexity bound."

Verification: this paper has **no Lean formalization**; the repository has no Comparator target for it [P]. So there is no machine check at all, and its credibility rests entirely on whether someone reads and understands the 73 pages. I also found no public expert comment specifically on it [△](#).

How to think about it:

- **No practical significance.** $2^{-(182)} \approx 1.6 \times 10^{-55}$. A rough calculation: even if each of the two numbers has 10^{100} digits, the saved factor $(\log n)^{(2^{-(182)})}$ exceeds 1 by only about 10^{-54} . A saving like this is visible only at astronomical sizes.
- **But it refutes a conjecture more than 50 years old.** Schönhage–Strassen conjectured in 1971 that $n \log n$ is optimal; this paper claims that, at least **in the multi-tape Turing machine model**, it is not. The qualifier matters; the abstract's own words are "in this model". Whether $n \log n$ is optimal in other models of computation is a separate question.
- **The mirror image of this section's main result:** $L = BPL$ is "everyone believes the conclusion; the hard part is the proof", and it has Lean. Integer multiplication refutes a long-standing conjecture, but has no Lean. Within the same batch, the source of trust can be completely different, which is why §8.8 spends so much space on Lean. [Speculation]

8.10 Self-test

Try answering these after reading. If you can't, go back and reread the corresponding section.

1. Why does " $O(\log n)$ bits of working memory" roughly equal "can store a constant number of pointers"? Why must such a deterministic machine halt in polynomial time? (§8.1.2, 1.3)
2. What is the difference between RL and BPL? Why does the restriction "once a coin is read, it's gone" matter? (§8.3.1)
3. Why does the random walk algorithm work on undirected graphs and fail on directed graphs? Give a number for each. (§8.4.3–8.4.4)
4. In Reingold's algorithm, what are "levels" and "per level" in "levels $\times$ per level"? How does this differ from "levels $\times$ per level" in Nisan's generator? (§8.5.3, 6.2)
5. Why is "a PRG with seed length $O(\log n)$ " enough to imply $L = BPL$? (§8.6.1)
6. In OpenAI's Lean statement, who defined L, RL and BPL? How is this different from the quasi-Riemann hypothesis case (§4)? Why does it matter? (§8.8.1)
7. Why does $L = BPL$ not imply $P = BPP$? (§8.2.2, 9.1)
8. Why does the integer multiplication paper "refute a conjecture more than 50 years old" while having "no practical significance"? How does its credibility differ from that of $L = BPL$? (§8.9.4)

Appendix: answers to the checks

- **Check 1:** $\lceil \log_2 10^9 \rceil = 30$ bits. Current vertex 30 + target vertex 30 + counter ($n^3 = 10^{27}$, needing $\lceil \log_2 10^{27} \rceil = 90$ bits) = 150 bits, less than 20 bytes.
- **Check 2:** the two sides are equal $\Leftrightarrow 2xy = 0 \Leftrightarrow x = 0$ or $y = 0$. There are 100 combinations with $x = 0$ and 100 with $y = 0$; (0, 0) is counted twice, so there are 199 in total (checked by script).
- **Check 3:** after step 1 the distribution is (0.5, 0.5, 0). Step 2: of state 0's 0.5, half stays at 0 and half goes to 1; of state 1's 0.5, half goes to 2 and half goes back to 0. So $P(0) = 0.25 + 0.25 = 0.5$, $P(1) = 0.25$, $P(2) = 0.25$.
- **Check 4:** $h_2 = 0$; $h_1 = 1 + \frac{1}{2} \cdot 0 + \frac{1}{2} h_0$; $h_0 = 1 + \frac{1}{2} h_1 + \frac{1}{2} h_0$, i.e. $h_0 = 2 + h_1$. Substituting gives $h_1 = 1 + \frac{1}{2}(2 + h_1) = 2 + \frac{1}{2} h_1$, so $h_1 = 4$ and $h_0 = 6 = 2^3 - 2$.

Appendix: Glossary

Term	One-line explanation
Space complexity	The maximum number of bits of working memory used at any one moment (the input doesn't count)
L	Problems decidable deterministically with $O(\log n)$ bits of working memory
NL	Nondeterministic logspace; directed reachability is its representative problem
RL	Logspace + polynomial time + coins, one-sided error ("no" is never accepted)

Term	One-line explanation
BPL	Logspace + polynomial time + coins, two-sided error ($\geq 2/3$ vs $\leq 1/3$)
BPP	Polynomial time + coins, two-sided error
Configuration	The machine's complete state at a given moment; a logspace machine has only polynomially many configurations
PIT	Polynomial identity testing; has a simple randomized algorithm, but no known deterministic polynomial-time algorithm
USTCON	Undirected s-t connectivity
Cover time	The expected number of steps for a random walk to visit every vertex
Expander	A graph with small degree but excellent connectivity; its spectral gap is a constant
Spectral gap	$1 - \lambda_2$; measures how fast a random walk "forgets its starting point"
Zig-zag product	A graph operation that reduces a high-degree expander to constant degree while mostly preserving the gap
Pseudorandom generator (PRG)	Stretches a short seed into a bit string that "looks random to a certain class of observers"
Seed length	The number of truly random bits a PRG needs; $O(\log n)$ means the seeds can be enumerated
Derandomization	Turning a randomized algorithm into a deterministic one
Hardness vs. randomness	A family of theorems of the form "hard functions exist" $\Rightarrow$ "derandomization is possible"
Catalytic computing	Computing with borrowed memory that is already full, and returning it exactly as it was
Comparator	A Lean FRO tool that checks that a solution proves exactly the statement in the challenge file, using only allowed axioms

Appendix: Further reading

- **Sanjeev Arora & Boaz Barak**, *Computational Complexity: A Modern Approach* (Cambridge University Press, 2009): Chapter 4 (space complexity), Chapter 7 (randomized computation), Chapter 21 (expanders and Reingold's algorithm). Free draft: <https://theory.cs.princeton.edu/complexity/>
- **Salil Vadhan**, *Pseudorandomness* (Foundations and Trends in TCS, 2012): Chapter 2 (random walks and USTCON), Chapter 4 (expanders), Chapters 7–8 (PRGs). Free: <https://people.seas.harvard.edu/~salil/pseudorandomness/>
- **William Hoza**, *Recent Progress on Derandomizing Space-Bounded Computation* (BEATCS survey, 2022): the best introductory survey of the current state of the L vs BPL field. [△](#) Please search for and confirm the link yourself.

- Original literature: AKLLR, FOCS 1979; Nisan, *Combinatorica* 12(4), 1992; Saks & Zhou, JCSS 58(2), 1999; Reingold, JACM 55(4), 2008; Hoza, APPROX/RANDOM 2021, LIPIcs 207:28.
 - The original paper: <https://github.com/openai/math/blob/main/preprints/Exact-Derandomization-of-Logarithmic-Space-L-equals-RL-equals-BPL-September-23-2026/paper.pdf>
 - Lean statement:
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/LogspaceEquality.lean>
; scope: `lean/docs/103.md`
 - Integer multiplication (family 109), the original paper:
<https://github.com/openai/math/blob/main/preprints/Integer-multiplication-below-n-log-n-September-23-2026/paper.pdf> ; the previous record: Harvey & van der Hoeven, *Integer multiplication in time $O(n \log n)$* , *Annals of Mathematics* 193(2), 2021.
 - Reactions: Aaronson <https://scottaaronson.blog/?p=10169> ; Fortnow <https://blog.computationalcomplexity.org/2026/10/open-no-more.html> ; Fortnow's selection page <https://lance.fortnow.com/openai-math-tcs.htm> ; *Scientific American* <https://www.scientificamerican.com/article/the-most-exciting-claims-from-openais-heap-of-new-proofs/>
-

COMBINATORICS · FAMILY 158 (+ 156)

9. The chromatic number of the plane $\chi \geq 6$ (the Hadwiger–Nelson problem)

The only prerequisites: you know what a graph is (vertices + edges), you can use the Pythagorean theorem to compute the distance between two points, and you know cos/sin. Words like measure and ergodic theory are explained the first time they appear. Every number in the text is either computed by script (script: `精讲/scripts/06_coloring.py`, with results copied into the text) or has a source. The **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. Tags: **[P]** = primary material read (paper, openai/math repository file or Lean source); **[R]** = secondary reporting (blogs, news); $\triangle$ = something I could not verify, or where sources conflict; "speculation" = my own judgment. Estimated reading time: 50–80 minutes. You can read it in three sittings: §9.1–9.4 (graph coloring, lower bound 4, upper bound 7), §9.5–9.7 (the 2018 breakthrough, "weird sets" and the axiom of choice, why 6 is hard), §9.8–9.10 (OpenAI's result, verification, what it means, ending with the Borsuk counterexample from the same batch).

9.0 The bottom line: the one thing this chapter explains

The problem itself is something a grade-schooler can understand: **give every point in the plane a color, so that any two points at distance exactly 1 have different colors. What is the minimum number of colors?** This minimum is called the "chromatic number of the plane", written $\chi(\mathbb{R}^2)$.

Since 1950 it has been known that the answer is between 4 and 7, and then it stayed stuck for 68 years. In 2018 the lower bound was pushed to 5. OpenAI claims to have pushed it to 6, so the answer can only be 6 or 7.

This chapter explains a causal chain:

Lower bounds come from "finding a finite point set that can't be colored" (the Moser spindle $\rightarrow$ de Grey's 1581-point graph) $\Leftarrow$ in principle you can always find one this way (the de Bruijn–Erdős compactness theorem) $\Leftarrow$ but the point set needed may be too large to find by search $\Leftarrow$ OpenAI took a different road: first turn "any weird coloring" into "a measurable coloring" (ergodic theory), then use geometry to rule out measurable 5-colorings, and the last step comes back to the 7-point Moser spindle.

After reading this chapter, you should be able to explain in your own words:

- why 3 colors are not enough (you can prove it by drawing 7 points by hand), and why 7 colors are enough (a hexagonal tiling);
- why "are the color classes measurable or not" is the real difficulty in this problem, and how it relates to the axiom of choice;
- what OpenAI's proof does structurally, why its Lean statement can be checked at a glance, and what it does **not** settle (6 or 7).

9.1 Graph coloring: an old friend from CS

9.1.1 Definitions

A **graph** is a set of vertices plus a set of edges. A **proper coloring**: give each vertex a color so that any two vertices joined by an edge have different colors. The minimum number of colors a graph needs is its **chromatic number** $\chi(G)$.

You've seen this in CS:

- **Register allocation**: variables are vertices, two variables that are "live at the same time" are joined by an edge, and colors are registers. Chromatic number = the minimum number of registers needed.
- **Course/exam scheduling**: courses are vertices, two courses that share a student are joined by an edge, and colors are time slots.

Deciding whether a general graph can be 3-colored is NP-complete. This will matter later: with brute-force computer search, once the graph gets big, the search becomes infeasible.

9.1.2 A few toys to work by hand

- One edge: needs 2 colors.
- **Triangle** (3 points, all pairwise connected): 2 colors aren't enough. The script enumerated all $2^3 = 8$ colorings and found **0** proper ones. So a triangle needs 3 colors.
- In general, if a graph has k points all pairwise connected (called a k -clique), it needs at least k colors. The converse doesn't hold: the chromatic number can be larger than the largest clique. The Moser spindle below is an example: it has no 4-clique, yet it needs 4 colors.

[Check 1] A ring of 5 points (a 5-cycle, where each point is connected only to its left and right neighbors): what is the minimum number of colors? Why aren't 2 enough?

9.2 Turning the whole plane into a graph

9.2.1 The unit-distance graph

Now treat **every point of the plane** as a vertex (there are uncountably infinitely many), and join two points by an edge if and only if their distance is **exactly 1**. This enormous graph is called the **unit-distance graph** of the plane, and its chromatic number is $\chi(\mathbb{R}^2)$.

An equivalent way to say it: split the plane into k pieces (one color per piece) such that **no** piece contains two points at distance exactly 1. What is the minimum number of pieces?

Two things to note:

- The distance is "exactly 1", not "less than 1". Two points at distance 0.999 or 1.001 may have the same color.
- **There is no requirement at all on the shape of the color pieces.** They can be nice hexagons, or weird sets built with the axiom of choice that can't be drawn at all. This is the core difficulty of the whole problem; §9.6 is devoted to it.

9.2.2 The first few steps

- $\chi \geq 2$: any two points at distance 1 already need 2 colors.
- $\chi \geq 3$: draw an equilateral triangle with side 1. Its three points are pairwise at distance 1; it's the triangle from §9.1.2, and 2 colors aren't enough.

The recipe for lower bounds is set right here: **find a finite set of points in the plane whose unit-distance edges form a graph, prove that this finite graph can't be colored with k colors, and then the plane can't be colored with k colors either** (because any

proper coloring of the plane, restricted to these few points, is a proper coloring of this small graph).

9.3 Lower bound 4: the Moser spindle (1961)

9.3.1 Step one: a rhombus forces its two tips to share a color

Put two equilateral triangles of side 1 together to form a rhombus: $O=(0,0)$, $B=(1/2, \sqrt{3}/2)$, $C=(1,0)$, $T=(3/2, \sqrt{3}/2)$. $O-B-C$ is one triangle, $B-C-T$ is the other, and they share the edge BC .

With only 3 colors: O , B and C are pairwise adjacent, so each takes one of the three colors; T is adjacent to both B and C , so T can only use the remaining color, which is O 's color.

Conclusion: with 3 colors, O and T must have the same color. Script check: this rhombus has 6 proper 3-colorings in total, and in every one of them O and T have the same color.

The distance between O and T is $\sqrt{3} \approx 1.732$ (computed by script), not 1, so their sharing a color does not by itself break any rule.

9.3.2 Step two: add a second rhombus and rotate it

Rotate the whole rhombus about O by an angle θ to get a second rhombus O, B_2, C_2, T_2 . By the same argument, T_2 must also have the same color as O . Now choose θ so that the distance between T and T_2 is exactly 1.

Both tips lie on the circle centered at O with radius $\sqrt{3}$, separated by angle θ . The square of their distance is:


$$|T - T_2|^2 = (\sqrt{3})^2 + (\sqrt{3})^2 - 2 \cdot 3 \cdot \cos \theta = 6 - 6 \cos \theta$$

Setting it equal to 1 gives **$\cos \theta = 5/6$** , $\theta \approx 33.56^\circ$ (computed by script).

Now T and T_2 must both have the same color as O , but they are at distance 1 from each other, so they must have different colors. **Contradiction. So 3 colors aren't enough, and $\chi \geq 4$.**

The figure below puts the two rhombi, the 11 unit-length edges and this argument together (the two tips are labeled T_1 and T_2 in the figure). The lower right shows how the known range for the chromatic number of the plane has changed; the next few sections cover each step.

Moser spindle (1961): 7 points, 11 unit-length edges $\Rightarrow$ 3 colors are not enough



Argument (with 3 colors)

O, B₁, C₁ pairwise at distance 1 $\rightarrow$ one color each
T₁ is 1 from B₁, C₁ $\rightarrow$ T₁ = color of O
likewise T₂ = color of O
but |T₁T₂| = 1 $\rightarrow$ contradiction $\Rightarrow \chi \geq 4$

Known bounds on χ of the plane

1950–61 $4 \leq \chi \leq 7$ (spindle, hexagons)

2018 $5 \leq \chi \leq 7$ (de Grey graph)

2026–09 $6 \leq \chi \leq 7$ (OpenAI, Lean)

Coordinates computed by the build script: two rhombi, each made of two unit equilateral triangles, share the vertex O, and each tip T is $\sqrt{3}$ from O. Rotating one rhombus about O by $2 \cdot \arcsin(1/(2\sqrt{3})) \approx 33.6^\circ$ puts the two tips exactly 1 apart.

9.3.3 Script check

The script builds 7 points from the coordinates above:

Point	Coordinates (four decimal places)
O	(0, 0)
B ₁	(0.5, 0.866)
C ₁	(1, 0)
T ₁	(1.5, 0.866)
B ₂	(−0.062, 0.9981)
C ₂	(0.8333, 0.5528)
T ₂	(0.7713, 1.5508)

It then checks the distances of all 21 pairs of points; exactly **11 pairs** are at distance 1: 5 edges from each rhombus, plus T₁–T₂. No unexpected extra unit distances.

Then brute-force enumeration:

- 3 colors: $3^7 = 2187$ colorings, **0** proper.
- 4 colors: **384** proper (= $4! \times 16$, i.e. if you ignore the names of the colors there are 16 essentially different colorings).

So this 7-point, 11-edge graph has chromatic number exactly 4. It is called the **Moser spindle** (published by the Moser brothers in 1961).

[Check 2] Above we used $|T - T_2|^2 = 6 - 6\cos \theta$. Derive it yourself using the law of cosines, and verify that the distance really is 1 when $\cos \theta = 5/6$.

Remember this 7-point graph: the last step of OpenAI's 62-page proof still uses it.

9.4 Upper bound 7: a hexagonal tiling

Lower bounds come from "finding a small graph that can't be colored". An upper bound requires **giving a coloring that is proper for the entire plane**.

9.4.1 The construction

1. Tile the plane with regular hexagons (a honeycomb). The paper takes the hexagons' circumradius to be $r = 2/5$.
2. Any two points in the same hexagon are at distance at most its diameter, $2r = 0.8 < 1$. **So a whole hexagon can be given one color**, and no pair of points inside it is at distance 1.
3. Color the hexagons periodically with 7 colors, so that two different hexagons of the same color are far enough apart. In the paper, the centers of same-colored hexagons are at least $\sqrt{21} \cdot r \approx 1.833$ apart, so points in the two pieces are at least $(\sqrt{21} - 2) \cdot r \approx \mathbf{1.033 > 1}$ apart (computed by script).
4. Points on the boundaries can be assigned to any one of the adjacent pieces (this is how the paper handles them, and the Lean version also covers the boundaries).

9.4.2 Where "7 colors" comes from

The centers of the hexagons form a triangular lattice. The paper writes it with complex numbers: lattice points have the form $a \cdot (m + n \cdot \omega)$, where $\omega = e^{(2\pi i/3)}$ is a cube root of unity (the Eisenstein integers from §4, §9.8.2, show up again).

Multiplying by $(2 - \omega)$ gives a sublattice that is "scaled up and rotated a bit". It makes up $1/7$ of the original lattice, because $|2 - \omega|^2 = 7$ (computed by script: 7.0000000000000001, where the last digit is floating-point error). The original lattice splits into 7 classes according to "which translated copy of the sublattice a point belongs to", and each class gets one color. In the script this corresponds to a very simple rule: the center $a \cdot (m + n\omega)$ gets color $(m + 2n) \bmod 7$.

9.4.3 Script check

The script randomly scattered **200,000** pairs of points at distance exactly 1 (for each pair, pick a random point, then walk distance 1 in a random direction), colored them with the rule above, and found **0 pairs with the same color**.

[Check 3] r can't be just anything. For the construction to work, we need $2r < 1$ and $(\sqrt{21} - 2) \cdot r > 1$. What is the allowed range of r ?

So $4 \leq \chi(\mathbb{R}^2) \leq 7$. This interval held from 1950 (Nelson got 4 and Isbell got 7, neither published at the time; Soifer later pieced together the historical record; Hadwiger published it formally in 1961) until 2018.

9.5 2018: de Grey pushes the lower bound to 5

9.5.1 A breakthrough by an amateur mathematician

In April 2018, Aubrey de Grey (known for his research on the biology of aging; mathematics is a hobby for him) posted *The chromatic number of the plane is at least 5* on arXiv (arXiv:1804.02385). The method was still the recipe from §9.2.2: assemble a huge finite unit-distance graph out of many parts like the Moser spindle, and prove that it **can't be 4-colored**. The corrected graph has **1581 vertices**, and the "can't be 4-colored" step was verified by computer.

9.5.2 Follow-ups: making the evidence smaller and readable

Year	Who	What they did	Source
2018	Heule	Used a SAT solver plus "unsatisfiable core" extraction to get a 553-vertex graph needing 5 colors	arXiv:1805.12181
2018/2020	Exoo, Ismailescu	Independently gave another proof of $\chi \geq 5$	arXiv:1805.00157
2020	Parts	Shrank the graph to 509 vertices (graph minimization)	arXiv:2010.12665
2020	Parts	Gave a proof of $\chi \geq 5$ that a human can check by hand	arXiv:2010.12661

This phase also included the crowdsourced project Polymath16 (organized by Dustin Mixon on his blog), and the work above was intertwined with it.

Where the 509 points come from: according to the introduction of the OpenAI paper, SAT plus unsatisfiable cores gave Heule's **553** points; the **509** points came from Parts's graph minimization. Both are connected to Polymath16.

9.5.3 Why this road keeps getting harder

A graph with n vertices and k colors has k^n possible colorings (orders of magnitude computed by script):

Graph	Vertices	Colors	$k^n \approx$
Moser spindle	7	3	$10^{3.3}$
Parts	509	4	10^{306}
Heule	553	4	10^{333}
de Grey	1581	4	10^{952}

A SAT solver of course doesn't enumerate these one by one; it relies on pruning. But to prove "there are none at all", the solver has to rule out the whole space, and the graph for 5 colors is expected to be much larger still. Going from 4 to 5 already needed a graph with over a thousand points; going from 5 to 6, nobody knows how big the graph has to be, and nobody found one during Polymath16.

9.5.4 A safety-net theorem: de Bruijn–Erdős compactness

If the finite graph we're looking for might be absurdly large, **does it necessarily exist?** Yes. This is the **de Bruijn–Erdős theorem (1951)**:

A graph (possibly infinite) can be colored with k colors if and only if **every finite subgraph** of it can be colored with k colors.

So "the plane can't be 5-colored" is equivalent to "some finite unit-distance graph can't be 5-colored". But the proof of this theorem is **non-constructive**: it guarantees that such a finite graph exists, but doesn't tell you what it looks like or how big it is.

Here is an analogy that comes naturally to CS people: it's like "if every finite subset of the constraints is satisfiable, then the whole infinite constraint system is satisfiable". The proof uses compactness of an infinite product space (Tychonoff's theorem), which requires the **axiom of choice** (explained in the next section). §4 of the OpenAI paper states this step explicitly: "using product compactness in ZFC".

9.6 The real difficulty: weird sets, measurability and the axiom of choice

9.6.1 What "measurable" means

Roughly speaking, a set in the plane is **measurable** if it has a sensible "area". Every shape you can draw (polygons, disks, countable unions of them, plus some more complicated but "well-behaved" sets) is measurable.

But in the usual set theory (ZFC, the axiom system that includes the **axiom of choice**), you can construct **non-measurable** sets: they have no sensible area, you can never draw them, and you can only prove they exist. The most famous examples are the Vitali set and the Banach–Tarski paradox (a ball is cut into finitely many pieces and reassembled into two balls of the same size; the pieces are non-measurable).

The axiom of choice says: given infinitely many nonempty boxes, you can always pick one element from each box, even if you can't write down a rule for picking. The vast majority of mathematicians accept it by default.

9.6.2 The measurable chromatic number: a different problem

If you require each color class to be measurable, the minimum number of colors is called the **measurable chromatic number** $\chi_m(\mathbb{R}^2)$. Clearly $\chi_m \geq \chi$ (an extra restriction can only make coloring harder).

- **Falconer (1981):** $\chi_m(\mathbb{R}^2) \geq 5$. The tool is "density points": near almost every point of a measurable set, the set fills almost 100% of the area. Around such a point, it's hard for the unit circle to avoid points of the same color; combined with a rotation argument, this gives a contradiction.

In other words, **under the "measurable" restriction, 5 was proved back in 1981**, 37 years before de Grey. The hard part is removing the "measurable" condition.

9.6.3 A cautionary example

The introduction of the OpenAI paper cites an example by Payne (2009) that illustrates the issue well: consider a subgraph of the plane in which two points are joined only when their difference is exactly "a unit vector with rational coordinates" (for example, $(3/5, 4/5)$).

- If any coloring is allowed, **2 colors** are enough for this graph;
- If the coloring must be measurable, it needs **at least 5 colors**.

So "measurable colorings" and "arbitrary colorings" can't be assumed to be the same, even when their distance constraints look identical. (△ I didn't read Payne's original; this restates what §1 of the OpenAI paper says.)

9.6.4 Shelah–Soifer: the answer might depend on the axioms

In *Axiom of choice and chromatic number of the plane* (J. Combin. Theory Ser. A 103 (2003) 387–391), Shelah and Soifer constructed some "distance graphs" whose chromatic number **depends on which axioms you use**: it is 2 in ZFC, but uncountable in the axiom system that "drops the axiom of choice and replaces it with 'every set of reals is measurable'" (ZF + DC + LM; Solovay proved in 1970 that this system is consistent, provided you assume an "inaccessible cardinal" exists).

They also proposed a conditional theorem, roughly: if all finite unit-distance graphs can be 4-colored, then the chromatic number of the plane is 4 in ZFC, but at least 5 in ZF+DC+LM (because in that world every coloring is measurable, so Falconer's result applies). (△ I paraphrased the exact statement of the conditional theorem from Soifer's later survey and did not check it word for word against the original.)

The premise of this conditional theorem has been refuted by de Grey, but the possibility it reveals is still there: **"how many colors the plane really needs" might not have an answer independent of the axioms**. That is also why the OpenAI paper states at the very start that it is "working throughout in ZFC".

9.6.5 "Nice" colorings: the answer is already known to be 7

If you require the color pieces to be "map-like" (with regular curves as boundaries), the situation is different again:

- Townsend (announced 1981, details published 2005): map-like colorings need at least **6** colors.
- Sokolov–Voronov (2025, arXiv:2502.01958): in a "polygonal, locally finite" framework, at least **7** colors are needed.

In other words, **for colorings "you can draw", the answer is 7**. The remaining question is: can weird sets save colors compared with nice colorings?

9.7 Why ≥ 6 is hard: both roads are blocked

Before 2026, the situation looked like this:

Approach	What it can achieve	Where it gets stuck
Finite graphs + computers (de Grey, Heule, Parts)	$\chi \geq 5$	The finite graph needed for 6 colors may be enormous; nobody has found one
Measurability / density points (Falconer)	$\chi_m \geq 5$	Holds only for measurable colorings; says nothing about weird sets
Regular boundaries (Townsend, Sokolov–Voronov)	map-like ≥ 6 , polygonal = 7	Requires regular boundaries; says nothing about general measurable sets, let alone weird sets

To get to 6, you need to clear two hurdles at once:

1. **From arbitrary colorings to measurable colorings:** how do you make weird sets "analyzable"?
2. **From measurable colorings to a geometric contradiction:** the boundaries of measurable sets can be extremely irregular, so you can't just draw them as a map.

OpenAI's paper corresponds exactly to these two steps.

9.8 What OpenAI claims

9.8.1 The main result

The abstract of the paper *The Euclidean plane is not five-colorable* (OpenAI, 2026-09-23, 62 pages, Family 158) [P]:

"We prove that every coloring of the Euclidean plane with five colors has a monochromatic unit-distance pair, with no regularity assumption on the color classes. Consequently, the chromatic number of the plane is either six or seven."

In plain words: **every 5-coloring of the plane contains two points at distance 1 with the same color, and no regularity is required of the color pieces. So the chromatic number of the plane is 6 or 7.**

The paper's Theorem 1.1 writes this as $6 \leq \chi(\mathbb{R}^2) \leq 7$, and immediately adds: "The remaining alternatives six and seven are unresolved."

A common misreading: it is easy to assume this result is "find a bigger finite graph and verify it with SAT". **It isn't.** The standard recipe for lower bounds had always been "find a finite graph and prove it can't be k -colored" (the recipe from §9.2.2, and de Grey's route), but this paper never gives a single finite graph that can't be 5-colored. It uses ergodic theory and measure theory instead, which is what makes it surprising.

9.8.2 The two halves of the proof

The paper splits the proof into two theorems:

- **Theorem 1.3 (transfer theorem):** for every k , "a proper k -coloring exists" $\Leftrightarrow$ "a **weakly measurable** k -coloring exists". This holds for all k , not just 5.
- **Theorem 1.4:** no weakly measurable 5-coloring exists.

Together: if there were a proper 5-coloring, then by 1.3 there would be a weakly measurable 5-coloring, contradicting 1.4.

What a "weakly measurable k -coloring" is (a plain-language version of Definition 1.2): every color piece is measurable, and the "same-colored unit-distance pairs" are null in the sense of measure. That is, exceptional same-colored pairs are allowed, but there are so few of them that they have zero area and can be ignored. This is weaker than a "proper coloring" (it allows measure-zero violations), and more regular than an "arbitrary coloring" (it must be measurable).

9.8.3 Intuition for the first half: "averaging" a weird coloring into a statistical regularity

What follows is my simplified paraphrase of §1.1 of the paper; the technical details go far beyond this.

1. **First shrink to a countable point set.** Take all points whose coordinates are **algebraic numbers** (roots of polynomial equations, like $\sqrt{2}$ or $(1+\sqrt{5})/2$), and call this set E . It has only countably many points, but it is dense in the plane, and it is closed under "multiplying by an algebraic complex number of absolute value 1" (that is, rotation by an algebraic angle).
2. **Average over translations and rotations.** Restrict the weird coloring to E , then "take the average" over all translations and rotations to get a **random coloring** whose statistics are invariant under any translation or rotation. Averaging like this relies on these groups being

amenable (roughly, "you can take a fair average over them"). This connects directly to §10: §10 explains amenability in detail, and what it proves about Thompson's group F (family 248) is exactly that it is **not** amenable.

An analogy with something familiar: this is like turning a very irregular signal into a **stationary random process**. A single sample may be weird, but its statistics (for example, its autocorrelation function) are regular.

1. **Use the spectral theorem for a "Fourier decomposition"**. A stationary process can be decomposed by frequency (compare the power spectrum of an EEG). Here the "frequencies" come in two kinds: ordinary continuous frequencies (corresponding to ordinary waves $\exp(i\xi \cdot z)$ on the plane), and a kind of "wild" frequency that only makes sense on the countable point set E .
2. **Rigidity theorem (Theorem 2.3)**: any distribution that is invariant under rotations and does not sit on the ordinary frequencies **must be completely uniform noise** (Haar measure). Uniform noise has zero correlation under any nonzero translation, so throwing it away doesn't break the correlation condition "unit-distance pairs have different colors". The proof of this step uses the **Furstenberg–Zimmer structure of compact extensions**, a tool from ergodic theory developed for proving Szemerédi's theorem.
3. **The remaining ordinary-frequency part can be extended to the whole plane**, giving a genuine measurable function. Then taking the most likely color at each point gives a weakly measurable coloring.

The reverse direction (weakly measurable $\Rightarrow$ proper) uses a Falconer-style density-point argument plus the de Bruijn–Erdős compactness from §9.5.4.

9.8.4 Intuition for the second half: ruling out 5 colors with geometry

Again a simplified paraphrase:

1. **"Palette"**: around each point x , look at the unit circle and record which colors appear in which directions on the circle. The paper proves that almost every such palette has **at most 2 colors** (paper §5).
2. **A "boundary graph" between colors**: treat the 5 colors as vertices and record which pairs of colors have a lot of transitions between them. The paper proves this graph contains a **cycle** (paper §6–7, using mollification, thresholding and a topological argument about coverings of the plane).
3. **A cycle on 5 colors can only have length 3, 4 or 5**, and the paper's §8 rules them out one by one:
 - 5-cycle: a combinatorial contradiction about a "6-letter angle word";
 - 4-cycle: a parity contradiction;

- 3-cycle: it leaves an open region that uses only 3 colors, and then a "placement certificate" with rational coordinates puts **the 7 points of the Moser spindle** into this region. §9.3 already proved that 3 colors aren't enough for those 7 points: contradiction.

You can see the definition of this "three-color region" directly in the Lean code

(`OAI/Geometry/PlaneColoring/Moser.lean`):

```
def ThreeLabelRegion (x y : ℝ) : Prop :=
  0 < x^2 + y^2 ∧ x^2 + y^2 < 4 ∧
  y^2 * (3 * x^2 - y^2)^2 < (x^2 + y^2)^3 * (1 - (x^2 + y^2) / 4) * (x^2 + y^2 - 1)^2
```

It's a region of the plane described by polynomial inequalities, and each point of the spindle has to be verified to lie inside it. **The whole 62-page argument starts from the 7-point toy of 1961 and comes back to it in the end.**

9.8.5 Two side results

- **Corollary 1.5 (the "bad fraction" has a positive lower bound):** there is a constant $\delta > 0$ such that in any 5-coloring, the fraction of "randomly placed needles of length 1 whose two ends have the same color" is at least δ . But the paper itself says: "The graph and the positive constant here are existential." That is, δ is **only known to exist; nobody knows how big it is**, because it comes from that finite graph needing 6 colors whose shape nobody knows.
- **Footnote 1 (Reed's manuscript, 2026-05):** Reed claims $\chi = 7$ and attaches a "formalized theorem". The OpenAI paper points out that this theorem put a **false** density upper bound in as a **hypothesis** (counterexample: the half-open arc $\{e^{it} : 0 \leq t < \pi/3\}$), so the formalization itself is not wrong; what's wrong is the proposition being proved. ⚠ I didn't read Reed's original manuscript; I only saw OpenAI's account of it. This example will be used in §9.9.

9.9 How far it has been verified

9.9.1 The Lean statement: two lines, you can read it yourself

File `lean/ComparatorChallenges/EuclideanFiveColor.lean` (snapshot fd4aeeb) [P]:

```
def ProperColoring (colorCount : ℕ) (coloring : ℂ → Fin colorCount) : Prop :=
  ∀ point otherPoint : ℂ, ‖point - otherPoint‖ = 1 → coloring point ≠ coloring otherPoint

theorem no_proper_five_coloring : ¬ ∃ coloring : ℂ → Fin 5, ProperColoring 5 coloring
```

Word-by-word translation:

- `coloring :  $\mathbb{C} \rightarrow \text{Fin colorCount}$` : a coloring is just a function that maps each complex number (i.e. a point in the plane, using $\mathbb{C}$ as the plane) to a color in $\{0, 1, \dots, \text{colorCount}-1\}$. **The function can be arbitrary**; there's no requirement like "measurable" or "continuous".
- `$\forall \text{ point otherPoint} : \mathbb{C}$` : for any two points in the plane.
- `$\|\text{point} - \text{otherPoint}\| = 1 \rightarrow$` : if their distance (the modulus of the complex difference) is exactly 1,
- `$\text{coloring point} \neq \text{coloring otherPoint}$` : then their colors differ. This is exactly "proper coloring".
- `$\neg \exists \text{ coloring} : \mathbb{C} \rightarrow \text{Fin 5}, \text{ProperColoring 5 coloring}$` : there **does not exist** any proper 5-coloring.

This is exactly the original statement of "5 colors aren't enough" for the Hadwiger–Nelson problem, with no sneaked-in measurability and no extra premises, so **the statement is faithful to the original problem**. The repository's scope note `lean/docs/158.md` also says explicitly that the lower bound "applies to arbitrary colorings, with no measurability or continuity assumption" [P]. Compared with the quasi-Riemann hypothesis in §4, it doesn't even depend on whether Mathlib's definition of some complicated function (like ζ) is correct; it only uses complex numbers, the modulus and finite sets. **This is one of the easiest statements for a human to audit in the entire OpenAI repository.**

The upper bound is formalized too: `properColoring_seven` in `ComparatorChallenges/PlaneColoring.lean`, where the plane is `EuclideanSpace  $\mathbb{R}$  (Fin 2)` and the condition is `$\text{dist } x \ y = 1 \rightarrow c \ x \neq c \ y$` . The proof file `Seven.lean` is only 68 lines.

9.9.2 The proof itself (statistics we ran ourselves, [P])

- The proof entry point is `OAI/Geometry/PlaneColoring/Five.lean`. A script followed its `import`s to collect all of its transitive dependencies inside the OAI directory: **70 files, 30,406 lines**.
- `grep` over these 70 files: **no** `sorry`, `admit`, custom `axiom` or `native_decide`.
- The final assembly corresponds exactly to the two halves of the paper:

```
theorem no_proper_five_coloring :  $\neg \exists c : \mathbb{C} \rightarrow \text{Fin 5}, \text{ProperColoring 5 } c$  := by
  rintro (c,hc)
  obtain (d,hd,hw) := proper_to_borel_weak hc          -- transfer theorem (Theorem 1.3)
  exact no_measurable_weak_five_coloring hd hw        -- geometric obstruction (Theorem 1.4)
```

- The Comparator configuration (`EuclideanFiveColor.json`) allows only the three standard axioms: `propext` , `Quot.sound` and `Classical.choice` . Note:
 - `Classical.choice` is the **axiom of choice** in Lean. So what this formalization proves is "in mathematics with the axiom of choice, the plane can't be 5-colored", which agrees with the paper's "in ZFC". The point in §9.6.4 that "with a different set of axioms the answer may differ" is not touched by this proof.
 - `"enable_nanoda": false` : an independent second checking kernel is **not** enabled. Someone independently re-checked the quasi-Riemann hypothesis of §4 with nanoda; for this one, as of 2026-10-08 we found no public record of an independent re-check [△](#).
- We did **not compile** this proof ourselves; the above comes from reading the source.

9.9.3 Human experts

- **Gil Kalai** (blog post, 2026-10-07): "The new manuscript raises the lower bound to six, leaving six and seven as the two possible answers." The disclaimer section of the same post says "even Lean verification may have issues". [R]
- **As of 2026-10-08, we found no public statement from de Grey, Polymath16 members or other combinatorial geometry experts saying they had read and verified the 62 pages.**
- The 2026-10-07 history of the OpenAI repository lists 3 retractions and 14 revisions; **this paper is not among them**. Family 158 is also **not** among the 10 families whose reasoning summaries are published in the README.

9.9.4 How to think about the remaining risk

- **Statement risk: very low.** Two lines that anyone can check against the original problem; see 9.1.
- **Proof risk: depends on the Lean kernel and Mathlib.** The measure theory, spectral theorem and so on used in the proof come from Mathlib or from these 70 files; the kernel checks every step. As long as the kernel has no bug and the statement is right, the conclusion holds.
- **Understanding risk: high.** Ergodic theory plus geometric topology, 62 pages, and no human expert has publicly read it all yet.
- **Reed as a cautionary tale:** Reed's "formalized theorem" doesn't count because the error was hidden in a **hypothesis**. OpenAI's statement has no hypotheses; nothing at all comes before `no_proper_five_coloring` , so the same problem can't arise. This shows exactly what the human half should be auditing in the division of labor "humans audit the statement, machines audit the proof".

9.10 What it would mean if true (and what it would not)

9.10.1 For mathematics

1. $\chi(\mathbb{R}^2) \in \{6, 7\}$. This is the first progress since 2018, and the first time the lower bound has been pushed without a finite graph.
2. **The transfer theorem itself may matter more than the "6"**: it holds for all k , which amounts to saying that "arbitrary colorings" and "weakly measurable colorings" are the same thing for this problem. Future work on the problem can go straight to the measurable world and use analytic tools, without worrying about weird sets. (This is my judgment; it counts as speculation.)
3. **A new open problem**: by de Bruijn–Erdős, there must be a finite unit-distance graph that can't be 5-colored (that is, one needing at least 6 colors), but nobody knows what it looks like or how big it is. Finding it (for example, with a SAT solver) would give a completely independent proof that can be verified finitely. This is a natural follow-up target for "AI + SAT" (speculation).

9.10.2 What it does not mean

- **It doesn't decide between 6 and 7**. The paper says explicitly that both possibilities remain unresolved.
- §9.6.5 says "nice colorings need 7 colors", so if the answer is 6, the coloring must use weird sets or extremely irregular measurable sets. Some people will guess from this that the answer is 7, but that's only intuition, not evidence (speculation).
- **It says nothing about worlds without the axiom of choice**. The proof is carried out in ZFC (`Classical.choice` in Lean). The question from §9.6.4 of what happens in a world where "all sets are measurable" is still open. Intuitively, in that world every coloring is measurable, so the geometric part, Theorem 1.4, would seem to be enough; but I haven't checked whether its proof uses more choice than DC $\triangle$ (speculation).
- **It gives no numerical value for δ** , and no concrete finite graph needing 6 colors.

9.10.3 For AI and AI safety

- This is the ideal form of "humans audit the statement, machines audit the proof": **a two-line statement, a 30,000-line proof**, and humans only need to audit the two lines. Compared with the quasi-Riemann hypothesis (§4), it doesn't even need the layer of trust "is Mathlib's definition of ζ correct?".
- Reed's counterexample shows the other side: **Lean only guarantees that "the given statement has been proved"**; the statement and its hypotheses still need human checking. This is the same kind of problem as specification gaming in reward models: the optimizer satisfies the spec you wrote down, not necessarily what you wanted.

- The "surprising" method (ergodic theory + geometric topology, rather than a bigger finite graph) deserves attention: if it holds up, it's an example of "the model found a route nobody in the field had taken", rather than "computing faster along a known route". Whether nobody had really taken it needs a domain expert's judgment $\triangle$.

9.10.4 Side note: the Borsuk counterexample from the same batch (family 156)

On the same day (2026-09-23) there was also a discrete geometry paper whose problem is a close relative of coloring the plane.

Toy: can a disk be cut into 2 pieces so that each piece has strictly smaller diameter (the distance between the two farthest points in the piece)? No: however you cut it, one piece contains a pair of antipodal points (the two ends of some diameter), and that piece's diameter has not shrunk. With 3 pieces (three 120° sectors) it works.

Borsuk's 1933 question: can every bounded set in $\mathbb{R}^d$ be cut into $d+1$ pieces of smaller diameter? For $d = 2$ and 3 the answer is yes ($d = 2$ was proved by Borsuk himself, $d = 3$ by Perkal in 1947 and Eggleston in 1955) [R]. **Kahn–Kalai 1993** gave a counterexample in dimension 1325, disproving the conjecture. After that the dimension of counterexamples kept falling; the last few steps were 65 (Bondarenko) $\rightarrow$ 64 (Jenrich–Brouwer 2014) $\rightarrow$ 63 (Grinsztajn's author manuscript, May 2026) [P, from the literature review and references of the OpenAI paper]. The 63-dimensional step reportedly used GPT-5.5 Pro [R, $\triangle$ I could not find a primary source to verify this]. Separately, arXiv:2608.12561 (Yibo Ji, 2026-08) gives another 63-dimensional counterexample; its abstract says the example and proof were "generated entirely by ChatGPT using GPT 5.6 Sol", and that the author personally verified the result [P].

How it relates to coloring the plane: for a finite point set, "cut into k pieces, each of smaller diameter" is equivalent to "join the pairs of points at distance **exactly equal to the diameter**, and this graph can be k -colored": a piece having smaller diameter just means no two of its points are a full diameter apart. Coloring the plane forbids "distance exactly 1"; Borsuk's problem forbids "distance exactly equal to the diameter". Both are about coloring a "distance graph". The Kahn–Kalai counterexample is exactly such a finite point set, whose chromatic number is proved to be large.

What OpenAI claims [P]: the paper *A nine-dimensional counterexample to Borsuk's covering assertion* (2026-09-23, 18 pages).

- Take all unit vectors u in $\mathbb{R}^4$ and form the projection matrices uu^T (the 4×4 matrix that projects a vector onto the line through u), with the Frobenius distance (treat a matrix as a vector of 16 numbers and take the Euclidean distance).

- These matrices are all symmetric with trace 1. A symmetric 4×4 matrix has 10 free parameters, and trace = 1 removes one more, so they lie in a **9-dimensional** space.
- The set has diameter $\sqrt{2}$ (attained when the two lines are perpendicular), and **cannot be covered by 10 sets of smaller diameter**. So Borsuk's conjecture already fails in dimension 9: from 63 straight down to 9. The paper's Corollary 7.1 extends the construction to every $d \geq 9$.
- The paper says explicitly that it does not determine the smallest dimension of a counterexample. Since the conjecture holds for $d \leq 3$, **dimensions 4–8 remain open**.
- The paper itself points out that the "projection matrix" embedding is classical: Kalai's account of the Kahn–Kalai construction uses $x \mapsto x \otimes x$ (that is, xx^T). What is new is proving the covering obstruction for this **continuous** set, using mod-two degree and simplicial complexes, and ending with a finite combinatorial argument.

Verification [P]: it has a Lean formalization

(`lean/ComparatorChallenges/BorsukNine.lean`). The statement is also short enough for a person to read: it defines `projectorSet` in the space of 4×4 matrices and asserts that it is compact, is contained in `traceOneSymmetric` (the symmetric matrices with trace 1), has diameter `Real.sqrt 2`, and `¬ HasTenSmallCover`, meaning there is no cover by 10 sets each of diameter less than $\sqrt{2}$.

How to think about it (speculation): it uses the same idea as this section's plane-coloring result. Instead of looking for a bigger finite point set, it switches to a nice continuous object (here the space of all lines in $\mathbb{R}^4$, the real projective space $\mathbb{RP}^3$) and rules it out with topology or analysis.

9.11 Self-test

Try answering these after reading. If you can't, go back and reread the corresponding section.

1. Why does finding a finite point set in the plane that can't be k -colored imply $\chi(\mathbb{R}^2) > k$? (§9.2.2)
2. In three sentences, restate why the Moser spindle proves that 3 colors aren't enough. (§9.3)
3. In the hexagonal 7-coloring, which two inequalities does $r = 2/5$ satisfy? What does each one guarantee? (§9.4)
4. What does the de Bruijn–Erdős theorem say? Why does it guarantee that a finite graph needing 6 colors exists, but not help you find it? (§9.5.4)

5. Falconer proved "5" back in 1981, so why was de Grey's "5" in 2018 still a breakthrough? (§9.6.2)
6. What are the two halves of OpenAI's proof? Which half holds for every number of colors k ? (§9.8.2)
7. What is the fundamental difference in trustworthiness between Reed's "formalized theorem" and OpenAI's `no_proper_five_coloring`? (§9.8.5, 9.4)
8. Why is Borsuk's problem a "close relative" of coloring the plane? What set is OpenAI's 9-dimensional counterexample, and why does it lie in a 9-dimensional space? (§9.10.4)

Appendix: answers to the checks

- **Check 1:** 3 colors. With 2 colors, the colors must alternate around the ring (red, blue, red, blue, ...). On a ring of odd length, when you go all the way around back to the start, the starting point would have to share a color with its own neighbor: contradiction. In general, a graph can be 2-colored if and only if it has no cycle of odd length.
- **Check 2:** T and T_2 are both at distance $\sqrt{3}$ from O , with angle θ between them. Law of cosines: $|T - T_2|^2 = 3 + 3 - 2 \cdot \sqrt{3} \cdot \sqrt{3} \cdot \cos \theta = 6 - 6 \cos \theta$. When $\cos \theta = 5/6$ this equals $6 - 5 = 1$, so the distance is 1.
- **Check 3:** $2r < 1 \Rightarrow r < 1/2$; $(\sqrt{21} - 2)r > 1 \Rightarrow r > 1/(\sqrt{21} - 2) \approx 0.3872$ (computed by script). So $0.3872 < r < 0.5$, and the paper's choice $r = 0.4$ is in range.

Appendix: Glossary

Term	One-line explanation
Proper coloring	A coloring in which adjacent vertices have different colors
Chromatic number $\chi(G)$	The minimum number of colors a proper coloring needs
Unit-distance graph	Vertices are points in the plane; two points at distance exactly 1 are joined by an edge
$\chi(\mathbb{R}^2)$	The chromatic number of the plane's unit-distance graph, i.e. the answer to the Hadwiger–Nelson problem
Moser spindle	A 7-point, 11-edge unit-distance graph that can't be 3-colored; proves $\chi \geq 4$
de Bruijn–Erdős theorem	An infinite graph can be k -colored $\Leftrightarrow$ every finite subgraph can be k -colored (needs the axiom of choice)
SAT solver	A program that decides whether a set of Boolean constraints can be satisfied simultaneously; used to verify that a finite graph can't be k -colored
Measurable set	A set with a sensible "area"
Measurable chromatic number χ_m	The minimum number of colors when all color pieces must be measurable; Falconer proved ≥ 5

Term	One-line explanation
Axiom of choice	You can always pick one element from each of infinitely many nonempty sets; it lets you build non-measurable sets
Weakly measurable coloring	Color pieces are measurable and same-colored unit pairs have measure zero (measure-zero exceptions allowed)
Amenable	The property of a group that "you can take a fair average" over it
Ergodic theory	The branch of mathematics that studies "long-run averages" and invariant statistical regularities
Haar measure	The completely uniform distribution on a group; corresponds to "pure noise"
Palette	The set of colors appearing on the unit circle around a point

Appendix: Further reading

- **The original paper:** <https://github.com/openai/math/blob/main/preprints/The-Euclidean-plane-is-not-five-colorable-September-23-2026/paper.pdf> (read Section 1 and Figure 1 in §1.1 first; the rest is technical detail)
- **Lean statement:**
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/EuclideanFiveColor.lean>
; scope note `lean/docs/158.md`
- **Gil Kalai's 2026-10-07 blog post:** <https://gilkalai.wordpress.com/2026/10/07/updates-sharing-ai-progress-on-mathematics-amazing-and-my-lecture-plans/>
- **Gil Kalai's 2018 introduction to de Grey's result** (the best entry point):
<https://gilkalai.wordpress.com/2018/04/10/aubrey-de-grey-the-chromatic-number-of-the-plane-is-at-least-5/>
- **de Grey 2018:** <https://arxiv.org/abs/1804.02385> ; **Heule's 553 points:**
<https://arxiv.org/abs/1805.12181> ; **Parts's hand-checkable proof:**
<https://arxiv.org/abs/2010.12661>
- **Quanta Magazine**, *Decades-Old Graph Problem Yields to Amateur Mathematician* (2018-04, popular coverage of de Grey)
- **Alexander Soifer**, *The Mathematical Coloring Book* (Springer, 2009; second edition 2024): the full history of the problem, including the chapters on Shelah–Soifer and the axiom of choice ("What If We Had No Choice?").
- **Shelah & Soifer**, *Axiom of choice and chromatic number of the plane*, J. Combin. Theory Ser. A 103 (2003) 387–391.
- **Payne**, *Unit distance graphs with ambiguous chromatic number*:
<https://arxiv.org/abs/0707.1177>

- **Sokolov & Voronov**, map-like colorings need at least 7 colors:
<https://arxiv.org/abs/2502.01958>
- **Terence Tao**, *254A Lecture 9: Ergodicity* (the ergodic theory introduction cited by the paper): <https://terrytao.wordpress.com/2008/02/04/254a-lecture-9-ergodicity/>
- **The Borsuk 9-dimensional counterexample (family 156)**: paper <https://github.com/openai/math/blob/main/preprints/A-nine-dimensional-counterexample-to-Borsuks-covering-assertion-September-23-2026/paper.pdf> ; Lean statement <https://github.com/openai/math/blob/main/lean/ComparatorChallenges/BorsukNine.lean> ; scope note `lean/docs/156.md`
- **Jenrich & Brouwer**, the 64-dimensional counterexample: <https://arxiv.org/abs/1308.0206> ; **Grinsztajn**, *A 63-dimensional counterexample to Borsuk's conjecture* (author manuscript, 2026-05, reference [7] of the OpenAI paper; $\triangle$ I found no public link); **Yibo Ji**, another 63-dimensional counterexample: <https://arxiv.org/abs/2608.12561> ; **Kalai**, *Some old and new problems in combinatorial geometry I: Around Borsuk's problem* (Surveys in Combinatorics 2015); overview: https://en.wikipedia.org/wiki/Borsuk%27s_conjecture

GROUP THEORY · FAMILY 248

10. Thompson's group F is nonamenable: a problem "proven" and withdrawn for decades

Group theory · family 248 of the openai/math release. All you need: fraction arithmetic, what a function and function composition $f(g(x))$ are, and the length and inner product of vectors.

No group theory is needed. Every new concept is explained before it is used. Every number here is either computed by script (`精讲/scripts/07_thompson.py` , exact arithmetic with Python's `Fraction`) or sourced. **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. Tags: **[P]** = primary material we read (paper PDF, Lean files, arXiv manuscripts and withdrawal notices); **[R]** = secondhand reporting or paraphrase; $\triangle$ = unverified or conflicting. As of 2026-10-08. Estimated reading time: 60–90 minutes. You can split it into three sittings: §10.1–10.3 (groups, free groups, "fair averages"), §10.4–10.5 (Banach–Tarski, Thompson's group itself), §10.6–10.9 (OpenAI's proof, its verification, and what it means).

10.0 The bottom line: the one thing this chapter explains

On some infinite "sets of operations" (mathematicians call them **groups**) you can define a **fair average**: shift everything over, and the average does not change. The integers allow this. On other groups such an average is **impossible**. The classic example is the "free

group". The reason it cannot be fairly averaged is that it can be cut up, shifted around, and reassembled into **two copies** of itself. This is the mechanism behind the famous Banach–Tarski paradox (cut a ball into a few pieces and reassemble them into two balls).

For a long time mathematicians thought that "cannot be fairly averaged" was always caused by a free group hidden inside. Counterexamples overturned this starting in 1980. But **Thompson's group F**, the most natural and most famous candidate, stayed undecided: it contains no free group, yet it does not belong to any type that "obviously can be averaged". In 1979 Geoghegan conjectured that it **cannot** be averaged. Over the next 47 years, people announced proofs in both directions, and one after another they were found to be wrong and withdrawn.

This chapter explains one causal chain:

F can be fairly averaged $\iff$ F contains finite sets that "barely leak" (Følner sets) $\implies$ averaging over such a set produces an approximate fixed point for a map that "pushes every point away by at least $\frac{1}{2}$ " $\implies$ contradiction. So F cannot be fairly averaged.

OpenAI gave this argument in a 13-page paper on 2026-09-23, together with a Lean formalization.

After this chapter, you should be able to explain in your own words:

- what "amenable" actually means, and why it is the same thing as the Banach–Tarski paradox;
- what the two generators of Thompson's group F look like, and how to compose them by hand;
- why F is the most famous "litmus test" in this field, and why its history makes Lean verification especially valuable;
- which "borrowed" phenomenon OpenAI's proof uses (maps with no fixed point in infinite-dimensional space), and what this result does and does not mean.

10.1 Groups: a set of operations you can compose and undo

10.1.1 A concrete example first: the symmetries of a square

Take a square card and label its four corners A, B, C, D. Which operations put it back into its original outline (the corner labels may move around)?

- Rotations by 0° , 90° , 180° , 270° (4 rotations);
- Flips about the vertical axis, the horizontal axis, and the two diagonals (4 flips).

That is 8 operations in total. Starting from "rotate 90° (call it r)" and "flip about the vertical axis (call it s)" and combining them repeatedly, the script generates exactly these 8, no more and no fewer.

These 8 operations have a few properties:

1. **They compose:** do one, then another, and the result is still one of the 8. Write $r \circ s$ for "first s, then r" (the same order as function composition $f(g(x))$: the one on the right happens first).
2. **There is a "do nothing":** rotate by 0°, written e (the identity).
3. **Every operation can be undone:** rotating 90° is undone by rotating 270°; each flip is undone by itself.
4. **Associativity:** $(a \circ b) \circ c = a \circ (b \circ c)$. Function composition satisfies this automatically.

But **order matters:** computed by script, $r \circ s$ and $s \circ r$ send the four corners to different places (one swaps A and C, the other swaps B and D). So groups are in general **not commutative**.

10.1.2 A second example: adding integers

Take all integers $\mathbb{Z} = \{\dots, -2, -1, 0, 1, 2, \dots\}$; an "operation" is "add some number". Composition is addition, the identity is 0, and +3 is undone by -3. This is an **infinite** group, and it is commutative.

10.1.3 Definition

Group = a set G together with an operation that combines two elements into one, such that: there is an identity element, every element has an inverse, and the operation is associative.

You can always think of a group as "all the symmetry operations of some object". Every group in this chapter is like that: the symmetries of a square, the translations of the integers, certain stretchings of the interval [0,1] (§10.5).

10.1.4 Generators and relations

Many groups can be built entirely out of **a few operations**. These operations are called **generators**.

- The 8 symmetries of the square are generated by r and s.
- $\mathbb{Z}$ is generated by the single element "+1" (-1 is its inverse, and every integer is a string of +1's or -1's).

Generators often satisfy **relations**. For example, in the square, $r \circ r \circ r \circ r = e$ (four rotations bring you back) and $s \circ s = e$. Relations make different "spellings" give the same element.

10.2 The free group F_2 : a group with no relations at all

10.2.1 Reduced words

Take two generators a, b and their inverses a^{-1}, b^{-1} (also written A, B below). The elements of the **free group F_2** are "words" spelled with these four letters. The only simplification rule is that an adjacent a and a^{-1} (or b and b^{-1}) cancel. A word with everything cancelled is called a **reduced word**. Apart from that there are **no relations at all**: $abab^{-1}$ and $baba^{-1}$ are different elements, and so are ab and ba .

Let's count the reduced words of length n . The first letter has 4 choices; each later letter cannot be the inverse of the letter before it, so it has 3 choices. So there are $4 \cdot 3^{n-1}$ words of length n .

Length n	Words of length exactly n ("sphere")	Words of length $\leq n$ ("ball")	Sphere / ball
0	1	1	1.000
1	4	5	0.800
2	12	17	0.706
3	36	53	0.679
5	324	485	0.668
10	78,732	118,097	0.667

(The script counts by enumerating reduced words one by one. Ball size = $2 \cdot 3^n - 1$.)

Two key points:

- The number of elements **grows exponentially** with n (roughly $2 \cdot 3^n$). Compare: $\mathbb{Z}$ has only $2n+1$ elements of length $\leq n$; $\mathbb{Z}^2$ (with L1 distance) has only $2n^2+2n+1$, which is 221 for $n=10$.
- **The outermost layer holds about 2/3 of everything**, and this fraction does not shrink as n grows. The outermost layer of $\mathbb{Z}$ has only 2 points, a fraction of $2/(2n+1) \rightarrow 0$.

[Check 1] There are 12 reduced words of length exactly 2. Write them all out. (Hint: $aa, ab, aB, \dots$)

10.2.2 Cayley graphs: drawing a group

Draw each element of the group as a point. If multiplying g on the right by a generator gives g' , draw an edge between g and g' . This picture is called the **Cayley graph**.

- $\mathbb{Z}$ (generator $+1$): an infinite straight line.
- $\mathbb{Z}^2$ (generators $(1,0), (0,1)$): an infinite square grid.

- F_2 (generators a, b): an **infinite tree in which every point has 4 edges**. Because there are no relations, different paths from the origin never come back to the same point, so there are no cycles.

In network terms: $\mathbb{Z}^2$ is like a 2D mesh, where the number of nodes grows polynomially with the radius; F_2 is like a 4-way fan-out tree, where **most nodes are on the boundary**. This property is the key to everything below.

10.3 Amenability: can you "average fairly"?

10.3.1 Where the question comes from

Averaging over a finite group is easy: add up the values at every element and divide by the number of elements. This average is "fair": shift the whole group (multiply every element on the left by the same h), and you are just adding the same numbers in a different order, so the average does not change.

On an infinite group this gets hard. There is no "uniform distribution" on $\mathbb{Z}$: if every integer has the same probability, that probability must be 0, and then the total is not 1.

Mathematicians get around this by relaxing the requirement. They ask only for a **mean**: a rule that assigns a score $M(\phi)$ to every bounded function $\phi: G \rightarrow \mathbb{R}$, such that

1. **Linear**: $M(\phi + \psi) = M(\phi) + M(\psi)$, $M(c\phi) = c \cdot M(\phi)$;
2. **Positive**: if $\phi \geq 0$ everywhere, then $M(\phi) \geq 0$;
3. **Normalized**: the mean of the constant function 1 is 1;
4. **Left-invariant**: shifting the whole function (replacing ϕ by $g \mapsto \phi(hg)$) does not change the mean.

If such an M exists on a group, the group is called **amenable**; otherwise it is **nonamenable**. This is exactly the definition used in the OpenAI paper and in the Lean statement (see §10.7).

Intuitively, M is an idealized version of "sprinkling a handful of sand evenly over the whole group". It only needs to be **finitely additive** (not countably additive). That is what gives groups like $\mathbb{Z}$ a chance to qualify.

10.3.2 The Følner criterion: find a finite set that "barely leaks"

The definition above is very abstract (M usually cannot be written down; proving it exists requires the axiom of choice). In 1955 Følner gave an equivalent criterion with a combinatorial flavor:

A group G is amenable $\iff$ for any finite collection of "shifts" h and any small $\varepsilon > 0$, you can find a finite set A with $|hA \Delta A| / |A| < \varepsilon$.

Here $hA = \{h \cdot a : a \in A\}$ is the result of shifting all of A , and Δ is the "symmetric difference" (the elements that appear on only one side). $|hA \Delta A|/|A|$ measures "after one shift, what fraction of the elements changed". The intuition: if A barely changes when shifted, averaging over A approximates a fair average; make A bigger and less and less "leaky", take the limit, and you get M .

Computed by script:

Group	Set A	A	Shift h	$ hA \Delta A / A $
$\mathbb{Z}$	$\{-n, \dots, n\}$, $n=10$	21	+1	0.0952
$\mathbb{Z}$	$n=1000$	2,001	+1	0.0010
$\mathbb{Z}^2$	$(2n+1) \times (2n+1)$ square, $n=10$	441	(1,0)	0.0952
$\mathbb{Z}^2$	$n=100$	40,401	(1,0)	0.0100
F_2	ball of length ≤ 4	161	a	1.006
F_2	ball of length ≤ 10	118,097	a	1.000

For $\mathbb{Z}$ and $\mathbb{Z}^2$ the fraction $\rightarrow 0$, so they are amenable. Shift a ball in F_2 and the fraction **tends to 1**: about half the elements move out and the same number of new elements move in ($|hA \Delta A|$ counts both sides). The reason is what §10.2 said: "most elements are on the boundary".

[Check 2] In $\mathbb{Z}$, let $A = \{-n, \dots, n\}$ and $h = +1$. Which two elements make up $hA \Delta A$? Use this to verify that the fraction is $2/(2n+1)$.

10.3.3 But "balls fail" does not mean "every set fails"

The table above tried only balls. The Følner criterion asks whether **some** A exists. In principle F_2 could contain some oddly shaped set that does much better than a ball. To prove that F_2 is not amenable, you need an argument that works for **every** set. The cleanest one is the following.

F_2 can be reassembled into two copies of itself. Let $W(a) =$ "the reduced words starting with a ", and define $W(a^{-1})$, $W(b)$, $W(b^{-1})$ the same way. Then:

- $F_2 = \{e\} \cup W(a) \cup W(a^{-1}) \cup W(b) \cup W(b^{-1})$, five pieces with no overlap.
- But **$W(a) \cup a \cdot W(a^{-1})$ is already all of F_2** : for any word w that does not start with a , putting a^{-1} in front gives $a^{-1}w$, which starts with a^{-1} (nothing cancels), and $w = a \cdot (a^{-1}w)$.
- Likewise $W(b) \cup b \cdot W(b^{-1})$ is also all of F_2 .

(The script checked both statements for every word of length ≤ 7 , and both hold; in the ball of length ≤ 8 , each of the four W pieces has 3,280 words.)

Now suppose F_2 has a fair mean M . Treat "belongs to a set" as a function taking values 0/1; then M gives each piece a "share". By left-invariance, the share of $a \cdot W(a^{-1})$ equals the share of $W(a^{-1})$. So:

- $M(W(a)) + M(W(a^{-1})) \geq M(\text{all of } F_2) = 1$;
- $M(W(b)) + M(W(b^{-1})) \geq 1$;
- Added together, the four pieces have total share ≥ 2 . But they are non-overlapping parts of F_2 , so their total share is ≤ 1 . Contradiction.

So **F_2 is not amenable**. This "make two copies of yourself out of a few pieces" is called a **paradoxical decomposition**. In 1938 Tarski proved that a group is nonamenable **if and only if** it has a paradoxical decomposition. So "cannot be fairly averaged" and "can duplicate itself" are two sides of the same thing.

Also: **if a group contains a nonamenable subgroup, the group itself is nonamenable** (a standard fact). So every group containing F_2 is nonamenable.

10.4 The Banach–Tarski paradox and von Neumann's question

10.4.1 The paradox

Banach–Tarski (1924): a solid ball in three-dimensional space can be cut into finitely many pieces which, using only rotations and translations, can be reassembled into **two** solid balls each the same size as the original.

This does not violate physics, because the pieces are non-measurable "weird sets" (they require the axiom of choice) and have no volume to speak of. Its mathematical root is exactly the paradoxical decomposition from §10.3.3: the 3D rotation group $SO(3)$ **hides a copy of F_2** (take two rotations about different axes with suitably chosen angles; there are no relations between them). Transfer F_2 's "one copy becomes two" onto the sphere, and you get the paradox.

10.4.2 von Neumann, 1929

von Neumann asked: why does this paradox exist in three dimensions but not in two (the plane)? His answer: **the key is not the space but the group of operations**. The group of rotations and translations of the plane is amenable, so it can be fairly averaged. That means the plane admits a "finitely additive, motion-invariant" notion of area, and the paradox cannot happen. The group of motions of 3D space contains F_2 and is nonamenable, so there

is room for the paradox. This is how the concept of "amenable" was born (German *messbar*; the English name "amenable" was coined later by Day).

10.4.3 The von Neumann–Day problem

Since "contains $F_2 \Rightarrow$ nonamenable", the natural question is the converse:

Must every nonamenable group contain F_2 ?

This is usually called the von Neumann problem or the von Neumann–Day problem (Day stated it explicitly in 1957). The answer is **no**, but the counterexamples were hard to come by:

Year	Author	Counterexample	Features
1980	Ol'shanskii	"Tarski monster" groups	Infinite, but every proper subgroup is a finite cyclic group; the construction is extremely complicated
1982	Adyan	Free Burnside groups with large odd exponent	Some fixed power of every element is the identity
2002	Ol'shanskii–Sapir	A finitely presented counterexample	The first counterexample that can be written down with finitely many generators and finitely many relations
2013	Monod	Groups of "piecewise projective" homeomorphisms of the real line	The first "natural" counterexamples; they look a lot like Thompson's groups
2016	Lodha–Moore	A finitely presented subgroup of Monod's group	Finitely presented and torsion-free (no elements of finite order)

(Monod: *Groups of piecewise projective homeomorphisms*, PNAS 2013; Lodha–Moore: *A nonamenable finitely presented group of piecewise projective homeomorphisms*, Groups Geom. Dyn. 10 (2016) 177–200.)

Note: a common claim is that "the counterexamples are all weird, and F would be the first natural counterexample". That was true before 2013; **since 2013, Monod and Lodha–Moore have given natural, finitely presented, torsion-free counterexamples**. So F 's significance is not "the first natural counterexample" but "**the most famous, most studied group in this field** finally has an answer".

10.5 Thompson's group F

10.5.1 Definition

F = all functions $f: [0,1] \rightarrow [0,1]$ such that:

1. f is continuous and strictly increasing, with $f(0) = 0$ and $f(1) = 1$;
2. f is **piecewise linear**, with finitely many breakpoints;

3. the x-coordinates of the breakpoints are all **dyadic rationals** (numbers of the form $k/2^n$, e.g. $1/2, 3/4, 5/8$);
4. the **slope of every piece is an integer power of 2** ($\dots, 1/4, 1/2, 1, 2, 4, \dots$).

The group operation is function composition.

You can check that composing two such functions gives another such function, and so does taking the inverse (reflecting the graph across the diagonal). So F really is a group. It is infinite: for example, you can squeeze $[0,1]$ to the left by arbitrarily large amounts.

F was introduced by Richard Thompson in 1965 (originally for a problem in logic); McKenzie–Thompson 1973 gave an early published construction. The 1996 survey by Cannon–Floyd–Parry, *Introductory notes on Richard Thompson's groups*, is the standard introduction.

10.5.2 The two generators, written out


x_0 : breakpoints $(0,0) \rightarrow (1/2, 1/4) \rightarrow (3/4, 1/2) \rightarrow (1,1)$.

Interval	Maps to	Slope
$[0, 1/2]$	$[0, 1/4]$	$1/2$
$[1/2, 3/4]$	$[1/4, 1/2]$	1
$[3/4, 1]$	$[1/2, 1]$	2

Intuitively: squeeze the left half to half its size, stretch the rightmost quarter to twice its size, and shift the middle. The overall effect is to "push points to the left".

The left half of the figure below is the graph of x_0 : all three pieces lie below the dashed line $y = x$, so every interior point is pushed to the left. The right half recaps §10.3: "can you average fairly?" for $\mathbb{Z}$ versus F_2 .


Left: the generator x_0 of the Thompson group F · Right: can you average fairly?



$\mathbb{Z}$: amenable

$A = \{-n, \dots, n\}$: shift by 1 changes $2, 2/(2n+1) \rightarrow 0$

Free group F_2 : non-amenable



radius- n ball: $\sim 3^n$ points, $\sim 2/3$ on the boundary

x_0 squeezes $[0, 1/2]$ onto $[0, 1/4]$, shifts $[1/2, 3/4]$ onto $[1/4, 1/2]$ and stretches $[3/4, 1]$ onto $[1/2, 1]$; F is generated by x_0 and x_1 . F contains no free subgroup F_2 (Brin–Squier 1985), so the standard key "contains $F_2 \Rightarrow$ non-amenable" does not apply. That is one reason the problem stayed open for 47 years.

x_1 : do nothing on $[0, 1/2]$; on $[1/2, 1]$, do a "scaled-down x_0 ". Breakpoints $(0,0) \rightarrow (1/2, 1/2) \rightarrow (3/4, 5/8) \rightarrow (7/8, 3/4) \rightarrow (1,1)$, with slopes $1, 1/2, 1, 2$ in order.

Cannon–Floyd–Parry proved: **every element of F can be built from x_0 , x_1 and their inverses.**

10.5.3 Composing by hand

Exact computation with fractions (checked by script):

- $x_0(3/8)$: $3/8$ lies in $[0, 1/2]$, where the slope is $1/2$, so the result is $3/16$.
- Breakpoints of $x_0 \circ x_1$ (first x_1 , then x_0): $(0,0) \rightarrow (3/4, 3/8) \rightarrow (7/8, 1/2) \rightarrow (1,1)$, slopes $1/2, 1, 4$.
- Breakpoints of $x_1 \circ x_0$ (first x_0 , then x_1): $(0,0) \rightarrow (1/2, 1/4) \rightarrow (7/8, 5/8) \rightarrow (15/16, 3/4) \rightarrow (1,1)$, slopes $1/2, 1, 2, 4$.
- The two differ, so F is **not commutative**. A concrete example: $x_1(x_0(7/8)) = 5/8$, while $x_0(x_1(7/8)) = 1/2$.
- Breakpoints of x_0^{-1} : $(0,0) \rightarrow (1/4, 1/2) \rightarrow (1/2, 3/4) \rightarrow (1,1)$, slopes $2, 1, 1/2$ (just swap the coordinates of each breakpoint of x_0).

[Check 3] Compute the breakpoints and slopes of $x_0 \circ x_0$. (Hint: first compute $x_0(1/2)$ and $x_0(3/4)$, then plug them into x_0 ; you also need to find which point the first x_0 sends to $3/4$.)

10.5.4 Relations: why F is not a free group

F has relations. The most intuitive one comes from "**functions with disjoint supports commute**": if f only moves points in $[0, 3/4]$ and g only moves points in $[3/4, 1]$, it does not

matter which one you do first.

Computed by script:

- $u = x_0 \circ x_1^{-1}$ has breakpoints $(0,0) \rightarrow (1/2, 1/4) \rightarrow (3/4, 3/4) \rightarrow (1,1)$, and is the identity on $[3/4, 1]$.
- $v = x_0^{-1} \circ x_1 \circ x_0$ has breakpoints $(0,0) \rightarrow (3/4, 3/4) \rightarrow (7/8, 13/16) \rightarrow (15/16, 7/8) \rightarrow (1,1)$, and is the identity on $[0, 3/4]$.
- So u and v commute: $u \circ v = v \circ u$. The script checked that $u^{-1}v^{-1}uv = \text{identity}$, and it holds; it also holds with v replaced by $x_0^{-2} \circ x_1 \circ x_0^2$.

These two are **all** the defining relations of F (Cannon–Floyd–Parry): F is **finitely presented** with 2 generators and 2 relations.

Relations make the balls smaller. Using x_0, x_1 and their inverses as generators, count the distinct elements of length $\leq n$ (the script deduplicates using exact breakpoints):

n	Ball in F	Ball in F_2
1	5	5
4	161	161
5	475	485
6	1,381	1,457
8	11,237	13,121

The first 4 layers are identical; from layer 5 on, there are 10 fewer. The reason: the shortest relation, $u^{-1}v^{-1}uv$, has length $2+3+2+3 = 10$, so length 5 is the first point at which two different words can represent the same element. F still grows exponentially (layer 8 still has 7,280 new elements), but exponential growth by itself does not decide amenability (some exponentially growing groups are amenable).

10.5.5 Why F is the litmus test of this field

Both standard tools **fail** on F :

- **"Contains $F_2 \Rightarrow$ nonamenable" does not apply:** Brin–Squier proved in 1985 that F contains no non-abelian free subgroup (Inventiones 79).
- **"Built from abelian and finite groups $\Rightarrow$ amenable" does not apply either:** F is not an "elementary amenable" group (Cannon–Floyd–Parry Theorem 4.10).

So F falls through the gap between the two standard methods of deciding the question. Add to this that its definition is simple and it has very nice properties (Brown–Geoghegan proved in 1984 that it has strong finiteness properties), and it became the single most watched

group after the von Neumann–Day problem. Geoghegan conjectured in 1979 that F is **nonamenable** (recorded in Cannon–Floyd–Parry 1996, p. 227).

Historical claims and retractions (each row has a primary source [P]: the arXiv manuscripts and the authors' withdrawal notices):

Date	Claim	Direction	Outcome
2009-02	Akhmedov, arXiv:0902.3849	Nonamenable	Withdrawn by the author (v8, 2013-12, pointing to the new manuscript 1310.4395)
2009-06	Shavgulidze, arXiv:0906.0107 (a related paper was published in a journal)	Amenable	Moore 2011 (arXiv:1102.0747) lists specific errors and writes that those who studied it carefully agree: "the errors are serious and do not seem to be repairable"; he also notes that the Math Reviews review is inaccurate
2011-12	Beklaryan, arXiv:1112.1942	Nonamenable	Withdrawn by the author in 2012-11, citing "a crucial error in Theorem B"
2012-09	Justin Moore , arXiv:1209.2063	Amenable	Withdrawn by the author on 2012-10-01: "Azer Akhmedov pointed out a serious error in Lemma 4.13. This error appears to be both serious and irreparable."
2013/2021	Akhmedov, arXiv:1310.4395v3	Nonamenable	The OpenAI paper only says it "claims"; $\triangle$ I found no evidence that the community accepted it, and before 2026 the problem was generally listed as open

There are also solid partial results. For example, Moore proved in 2013 that **if** F is amenable, its Følner sets must be astonishingly large ("tower-type" growth, *Groups Geom. Dyn.* 7).

Note Moore's retraction: he is one of the leading experts on this problem, and a fatal error was still found by peer checking within three weeks. This is the key background for understanding "why it matters that this result comes with Lean".

10.6 What OpenAI claims

10.6.1 The main result [P]

The paper is *Thompson's group F is nonamenable* (authored by OpenAI, family 248 of the openai/math release, 2026-09-23, **13 pages in total including the appendix**: the proof is on pages 1–7, pages 8–11 are appendices, pages 12–13 are references). The full abstract:

"We prove that Thompson's group F is nonamenable. This confirms Geoghegan's conjecture and resolves the amenability problem for F ."

Main theorem: **Theorem 1.1. Thompson's group F is not amenable.**

Page 1 of the paper honestly lists the failed historical claims above, and specifically notes that what Moore withdrew was a different paper, not his published Ramsey characterization.

10.6.2 The borrowed tool: a map that "pushes everything away"

First recall a classical theorem.

Brouwer fixed-point theorem: on a closed ball of **finite dimension** (a segment, a disk, a solid ball, ...), every continuous map f has at least one fixed point, $f(x) = x$. Intuition: you cannot continuously crumple a disk and put it back inside the disk while moving every point.

This theorem fails in infinite-dimensional space. The infinite-dimensional example is ℓ^2 : all infinite sequences $(x_1, x_2, x_3, \dots)$ whose squares have a finite sum, with length $\sqrt{(x_1)^2 + x_2^2 + \dots}$. Kakutani's classic example:

$$f(x) = (1 - \|x\|, x_1, x_2, x_3, \dots)$$

(shift the sequence right by one slot, and put $1 - \|x\|$ in front)

It maps the unit ball into the unit ball, and it **has no fixed point** (a fixed point would need all coordinates to be equal and the sum of squares to be finite, so they would all be 0; but then the first slot is 1, a contradiction).

But this example is not strong enough: its "displacement" can be **arbitrarily small**. Computed by script: take $x = (t, t, \dots, t)$ (N coordinates) with $t = 1/(1+\sqrt{N})$; then $f(x) - x$ is t in the $(N+1)$ -th coordinate only:

N	Displacement $\ f(x) - x\ $
1	0.500
100	0.091
10 ⁴	0.0099
10 ⁶	0.0010

So Kakutani's f has "**approximate** fixed points".

Benyamini–Sternfeld (1983) proved something stronger: on the unit ball of any infinite-dimensional normed space, there is a **Lipschitz** map f (it never stretches the distance between two points by more than a factor L) that **moves every point by at least $\delta > 0$** . Such an f does not even have approximate fixed points. Appendix A of the OpenAI paper builds its own version with $\delta = \frac{1}{2}$ on $L^2([0,1]; \mathbb{R}^2)$ (using a curve in infinite-dimensional space and a "tube" around it), and every Lipschitz estimate in the appendix is written out.

10.6.3 Skeleton of the proof (§10.2, about 5 pages)

The overall strategy in one sentence: **if F has finite sets A that barely leak, then averaging over A produces an approximate fixed point of the map f above.** This contradicts "every point is moved by at least δ ".

Five steps, explained as intuitively as possible:

1. **Color the "dyadic partitions" of $[0,1]$.** A partition cuts $[0,1]$ into small cells of the form $[k/2^r, (k+1)/2^r]$. Fix in advance D pairwise non-adjacent "parent intervals" $I_1 < \dots < I_D$. For a partition T , take its part inside each parent interval and stretch it back to $[0,1]$; this gives a smaller partition $T_{\{I_j\}}$. Define the "color" of T (a vector in the unit ball) recursively:

$$\begin{aligned} p(T) &= f((p(T_{\{I_1\}}) + \dots + p(T_{\{I_D\}})) / D) && \text{(if } T \text{ cuts at every } I_j) \\ p(T) &= 0 && \text{(otherwise)} \end{aligned}$$

Each restriction to a subinterval strictly reduces the number of cells, so the recursion always terminates.

1. **Let elements of F produce partitions.** An element g of F sends each cell of a uniform partition to a dyadic interval (Lemma 2.1 of the paper), so each g brings a family of colors: $X_i(g)$ is the color of the i -th parent interval, and $Y_{ij}(g)$ is the color of the j -th "grandchild interval" inside the i -th parent interval. The recursive definition directly gives **$X_i = f(z_i)$** , where z_i is the average of $Y_{i1}, \dots, Y_{iD}$. Let m = the average of all the X_i .

1. **F can move any pair of separated intervals to any other pair** (Lemma 2.2: refill the gaps with dyadic cells, and map them linearly one-to-one in order). So for any two separated intervals $I < J$ there is an element $h_{\{I,J\}}$ of F that moves them to the reference pair (I_1, I_2) , and the colors move along exactly. Let S be the **finite** set made of these $h_{\{I,J\}}$. It is fixed in advance and does not depend on A .

1. **Average over a Følner set.** Let η = the largest leakage fraction $|hA \Delta A|/|A|$ of A under S . When averaging over A , "replacing g by hg " barely changes the average, with error at most η . So: **for any two separated intervals, the average over A of the inner product of their colors is within η of the same constant α .** It is like saying: on a sample that is almost translation-invariant, all "pairs of separated positions" look equally correlated.

1. **A variance estimate, then a contradiction.** Expand $\|z_i - m\|^2$. Among its inner-product terms, the vast majority are "separated interval pairs", whose averages are all about α , and positive and negative contributions cancel; only the "nested pairs" (a grandchild interval and its own parent), a fraction $1/D$, are not controlled, but their inner products have absolute value at most 1. Working it out (equation 2.16 of the paper):

$$\text{average over } A \text{ of } \|z_i - m\|^2 \leq 4/D + 4(1 - 1/D) \cdot \eta$$

On the other hand, m is also the average of the $f(z_i)$; using convexity and the Lipschitz property:

$$\delta^2 \leq \|m - f(m)\|^2 \leq (L^2/D) \cdot \sum \|z_i - m\|^2$$

Combining: $\delta^2 \leq L^2 \cdot (4/D + 4(1-1/D) \cdot \eta)$. Rearranging gives (Proposition 2.3):

$$\eta \geq (\delta^2/L^2 - 4/D) / (4(1 - 1/D))$$

As long as D is chosen in advance to be larger than $4L^2/\delta^2$, the right-hand side is a **positive number that does not depend on A** . In other words: **every finite set A in F leaks at least a fixed positive fraction under this fixed finite set S** . This directly violates the Følner criterion, so F is not amenable.

A feel for the numbers (computed by script; L is an assumed value, since the paper does not give a concrete number for the appendix map's $L \triangleq$): with $\delta = 1/2$ you need $D > 16L^2$. If $L = 10$, taking $D = 3,200$ gives a lower bound of about 3.1×10^{-4} ; taking $D = 1,601$ (just past the threshold) gives about 3.9×10^{-7} . The fraction is tiny, but the proof only needs it to be positive and independent of A .

Why this proof is surprising: most earlier attempts looked for something inside F 's own combinatorial structure (trees, graphs, Ramsey properties, height functions). Here the key "impossibility" comes from **infinite-dimensional geometry**. F only supplies the flexibility to "move any pair of intervals to any pair", plus the self-similar recursion of dyadic partitions.

10.6.4 Two corollaries (relying on companion papers)

Section 3 of the paper derives two corollaries: F has a representation that is "almost isometric but cannot be unitarized" (relying on another OpenAI paper, *Unitarizability implies amenability for discrete groups*); and percolation on the Cayley graph of F has $p_c < p_u$ (relying on another OpenAI paper on percolation). The paper states explicitly that these two corollaries are **not used** in proving the main theorem. $\triangle$ We have not checked these two companion papers.

10.7 How far it has been verified

10.7.1 The Lean statement: F is defined from scratch, and you can read it line by line

Unlike the quasi-Riemann hypothesis in §4, this statement **does not rely on a ready-made F in Mathlib** (Mathlib has none). Instead, the challenge file

`ComparatorChallenges/ThompsonNonamenability.lean` (89 lines, including imports)

defines everything from scratch. The core:

```
def IsDyadic (x : ℝ) : Prop := ∃ (k : ℤ) (n : ℕ), x = (k : ℝ) / (2 : ℝ) ^ n

structure DyadicPLWitness (f : IntervalHomeomorph) where
  pieceCount : ℕ -- number of pieces
  knots : Fin (pieceCount + 1) → UnitInterval -- list of breakpoints
  strictMono_knots : StrictMono knots -- breakpoints in increasing order
  first : knots 0 = 0; last : knots (last index) = 1
  dyadic : ∀ i, IsDyadic (knots i) -- breakpoints are dyadic rationals
  exponent : Fin pieceCount → ℤ -- one integer exponent per piece
  affine : ... (f x) = f(knot_i) + 2^(exponent i) * (x - knot_i) -- each piece has slope a power of 2
  ...

def F := {f : IntervalHomeomorph // StrictMono f ∧ HasDyadicPLPieces f}

structure InvariantMean (G) [Group G] where
  toLinearMap : BoundedReal G →1[ℝ] ℝ -- linear
  positive : ∀ f, (∀ g, 0 ≤ f g) → 0 ≤ toLinearMap f -- positive
  normalized : toLinearMap 1 = 1 -- normalized
  left_invariant : ∀ h f, toLinearMap (leftPull h f) = toLinearMap f -- left-invariant

theorem thompson_F_nonamenable_composition :
  ∃ group : Group F, let I := group;
  (∀ (h g : F) (x : UnitInterval),
    (h * g).val.toHomeomorph x = h.val.toHomeomorph (g.val.toHomeomorph x)) ∧
  ¬ Nonempty (InvariantMean F)
```

(A few technical fields in the middle of the definitions are omitted; the full file is in the repository.)

Word-by-word translation:

- `IntervalHomeomorph`: a homeomorphism from $[0,1]$ to itself (continuous, invertible, with a continuous inverse).
- `F`: the homeomorphisms that are "strictly increasing" and have "a witness of being dyadic piecewise linear". This is the definition from §10.5.1, with nothing missing.
- `BoundedReal G`: bounded real-valued functions on G (Mathlib's ℓ_∞). `leftPull h f` is the function $g \mapsto f(hg)$.
- `InvariantMean`: the four conditions from §10.3.1: linear, positive, normalized, left-invariant.

- The theorem says: **there exists** a group structure on F whose multiplication is **function composition** ($h * g$ acting on x equals $h(g(x))$), and under this group structure there is **no** invariant mean.

Why "there exists a group structure" is not a loophole: once the group multiplication is required to be composition, the identity and the inverses are uniquely determined. So this "exists" just means " F really is a group under composition", and leaves no room for cheating. The multiplication convention $hg = h \circ g$ also matches the paper.

Our assessment: the statement is faithful to Theorem 1.1 of the paper [P]. In addition, the Lean proof includes a file `StandardCharacterization.lean` proving that the definitions in the challenge file are equivalent to the "textbook" formulation, and `MeanToFølner.lean`, which **proves in Lean itself** the direction "invariant mean $\Rightarrow$ Følner sets". So the Følner criterion that the paper cites is also machine-checked, not taken as an axiom.

10.7.2 The proof itself

- The proof lives under `lean/OAI/GroupTheory/Thompson/`: **68 files, 9,457 lines** (about 9,500; counted by script on snapshot fd4aeeb) [P]. Of these, 22 files (with names starting with `Analytic` or `Tube`) construct the map from Appendix A.
- We grepped it: this directory contains no `sorry`, no `axiom` declarations, and no `native_decide`.
- The Comparator configuration allows only Lean's three standard axioms (`propext`, `Quot.sound`, `Classical.choice`), and `enable_nanoda` is `false` (no second independent kernel check is turned on).
- **We did not compile it locally, and we did not run Comparator.** "Verified in Lean" is currently OpenAI's claim, plus our static inspection of the source.
- Unlike the quasi-Riemann hypothesis in §4: **as of 2026-10-08, we found no public record of any third party independently running this proof** (the quasi-Riemann hypothesis had an independent recheck by Dave Goldblatt using Comparator + nanoda). $\triangle$

10.7.3 Human experts [R]

- Gil Kalai (blog post, 2026-10-07) cited his 2009 post introducing Thompson's group and wrote: "The new manuscript announces that Thompson's group (F) is nonamenable." He uses "announces", a neutral wording.
- Group theorist Xiaolei Wu (compiled on proofsandprompts.com): "So Thompson's group F is non-amenable. OK, I had tried it many times using ChatGPT, but I guess I didn't have the newest model or enough tokens."
- **As of 2026-10-08, we found no public comments from central figures such as Geoghegan, Moore, or Akhmedov, and no named expert has publicly said they have**

read and verified the paper. ⚠

10.7.4 Remaining risks

1. **Statement risk is very low:** the challenge file is under 90 lines and relies only on standard Mathlib objects (real numbers, homeomorphisms, ℓ^∞ , linear maps), and the definitions can be checked line by line against a textbook. This is the ideal form of the division of labor "humans check the statement, machines check the proof".
2. **Compilation and toolchain risk:** we did not compile it ourselves, and there is no second-kernel check. Such risks are usually small, but not zero.
3. **Correspondence between the paper and Lean:** Lean proves Theorem 1.1 itself, so even if the paper's write-up has flaws, the main conclusion is unaffected. The two corollaries (§10.6.4) are outside what Lean covers.
4. **An unusual piece of good news:** the core proof in this paper is only about 5 pages, and the tools it uses (Hilbert-space inner products, convexity, Lipschitz estimates, dyadic intervals) are all standard. Compared with the 199-page quasi-Riemann hypothesis paper (§4), it is one of the headline results that **human experts are most likely to read in full and independently confirm in the short term.** (This is speculation.)

10.8 What it means if true (and what it doesn't)

10.8.1 For mathematics

1. **A 47-year-old problem (1979→2026) has been solved, and the answer matches Geoghegan's conjecture:** F is not amenable.
2. **F becomes the most famous counterexample to the von Neumann–Day problem:** nonamenable, yet containing no F_2 . But as §10.4.3 said, it is **not the first natural counterexample** (Monod 2013 and Lodha–Moore 2016 came first).
3. **A methodological surprise:** translate "averaging over a group" into "a fixed-point problem on the infinite-dimensional unit ball", then borrow a counterintuitive phenomenon from functional analysis, like Benyamini–Sternfeld, to reach a contradiction. Whether this route works for other groups (for example other groups of piecewise linear/piecewise projective homeomorphisms) is a natural open question. (Speculation.)
4. **What it does not mean:** the paper gives no explicit constant for a specific generating set (the Lean document `docs/248.md` says verbatim: "No explicit boundary constant or prescribed generating set is given."); nor does it answer F 's other famous questions.

10.8.2 For AI and for "verification" itself

This is the best case for "why Lean matters", for concrete reasons:

- The human history of this problem is a string of "proof announced → error found → withdrawn", in both directions, and even the leading expert withdrew once (the table in §10.5.5). Going by past experience, the community might take years to reach consensus on a new 13-page claim.
- With Lean, the object of trust shifts from "the author" to "**the kernel + an 89-line statement (imports included) that can be audited line by line**".
- At the same time, it reminds us what Lean cannot do: it cannot judge the **significance** of the result, it cannot tell whether the method generalizes, and it cannot replace human understanding of "why it is true".

For you (an AI safety researcher), this case has one more layer: it is a sample of "AI gives a short proof of a problem humans failed at for a long time, and humans can recheck it in a short time". Together with "AI gives a 480,000-line proof that humans can barely recheck" (the quasi-Riemann hypothesis in §4), it marks the two ends of a **scalable-oversight** spectrum. (Speculation, but worth writing down.)

10.9 Self-test

Try to answer these after reading. If you can't, go back and reread the corresponding section.

1. Explain in one sentence what "amenable" means. (§10.3.1)
2. Why are balls in F_2 not good Følner sets? Why is that not enough to prove F_2 is nonamenable? (§10.3.2–10.3.3)
3. Why does $W(a) \cup a \cdot W(a^{-1}) = F_2$ imply that F_2 has no fair mean? (§10.3.3)
4. The Banach–Tarski paradox holds in three dimensions but not in two. What did von Neumann think was the root cause? (§10.4.2)
5. Why does F fall "through the gap between the two standard methods of deciding the question"? (§10.5.5)
6. Neither the Kakutani map nor the Benyamini–Sternfeld map has a fixed point. What is the difference? Why does OpenAI's proof need the latter? (§10.6.2)
7. Why is "there exists a group structure" in the Lean statement not a loophole? (§10.7.1)

Appendix: answers to the checks

- **Check 1: aa, ab, aB, AA, Ab, AB, bb, ba, bA, BB, Ba, BA.** aA, Aa, bB, Bb don't count, because they cancel to the empty word. Each first letter has 3 choices after it, $4 \times 3 = 12$ (checked by script).

- **Check 2:** $hA = \{-n+1, \dots, n+1\}$. Compared with A , it gains $n+1$ and loses $-n$, so $hA \Delta A = \{-n, n+1\}$, 2 elements, fraction $2/(2n+1)$. For $n=10$ this is $2/21 \approx 0.0952$, matching the table.
- **Check 3:** $x_0(1/2) = 1/4$, $x_0(3/4) = 1/2$. Apply x_0 again: $x_0(1/4) = 1/8$, $x_0(1/2) = 1/4$. You also need the point that the first x_0 sends to $3/4$: on $[3/4, 1]$, x_0 is $y = 1/2 + 2(x - 3/4)$; setting $y = 3/4$ gives $x = 7/8$. So the breakpoints of $x_0 \circ x_0$ are **$(0,0) \rightarrow (1/2, 1/8) \rightarrow (3/4, 1/4) \rightarrow (7/8, 1/2) \rightarrow (1,1)$** , with slopes **$1/4, 1/2, 2, 4$** (checked by script). The slopes are still all powers of 2 and the breakpoints are still dyadic rationals, so it is still in F .

Appendix: Glossary

Term	One-line explanation
Group	A set of operations that compose, have an identity, and can each be undone
Generators	A few elements from which the whole group can be built
Relations	Equations between generators that make different spellings represent the same element
Free group F_2	The group with two generators and no relations at all
Cayley graph	The graph with group elements as points and edges given by generators; for F_2 it is a 4-way tree
Invariant mean	A rule for scoring bounded functions: linear, positive, normalized, invariant under left shifts
Amenable	An invariant mean exists; equivalent to having Følner sets
Følner set	A finite set in which only a tiny fraction of elements change after a shift
Paradoxical decomposition	Cutting a group into a few pieces and shifting them to form two copies of itself; equivalent to nonamenability (Tarski)
von Neumann–Day problem	Must every nonamenable group contain F_2 ? Answer: no
Thompson's group F	The group, under composition, of piecewise linear increasing homeomorphisms of $[0, 1]$ with dyadic breakpoints and slopes that are powers of 2
Fixed point	A point with $f(x) = x$
Lipschitz map	A map that stretches the distance between two points by at most a factor L
Benyamini–Sternfeld	On the infinite-dimensional unit ball there is a Lipschitz map that moves every point by at least δ

Appendix: Further reading

- **The original paper** (13 pages; Section 2 is fully within your reach):
<https://github.com/openai/math/blob/main/preprints/Thompsons-group-F-is-nonamenable-September-23-2026/paper.pdf>

- **Lean statement:**
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/ThompsonNonamenability.lean>
; scope notes: <https://github.com/openai/math/blob/main/lean/docs/248.md>
- **Cannon, Floyd, Parry**, *Introductory notes on Richard Thompson's groups*, L'Enseignement Math. 42 (1996) 215–256: the standard introduction to F ; the generators and relations in §10.5 all come from here.
- **Terence Tao**, *The Banach–Tarski paradox* (lecture notes, UCLA): full derivations of the paradoxical decomposition and amenability from §10.3.3 and 4. Search for "Tao Banach–Tarski paradox notes".
- **Juschenko**, *Amenability of discrete groups by examples* (AMS Mathematical Surveys and Monographs, 2022): a modern textbook on amenability, with a chapter on F . $\triangle$ Chapter number not checked.
- Primary historical sources: Moore's note on Shavgulidze <https://arxiv.org/abs/1102.0747> ; Moore's withdrawn manuscript <https://arxiv.org/abs/1209.2063> ; Akhmedov <https://arxiv.org/abs/0902.3849> and <https://arxiv.org/abs/1310.4395> ; Beklaryan <https://arxiv.org/abs/1112.1942>
- Natural counterexamples to von Neumann–Day: Monod PNAS 2013 (*Groups of piecewise projective homeomorphisms*); Lodha–Moore <https://arxiv.org/abs/1308.4250>
- Gil Kalai's 2009 blog post introducing F : <https://gilkalai.wordpress.com/2009/05/01/the-thompson-group/> ; the 2026 update: <https://gilkalai.wordpress.com/2026/10/07/updates-sharing-ai-progress-on-mathematics-amazing-and-my-lecture-plans/>
- Compilation of reactions: <https://proofsandprompts.com/2026/10/08/100-reactions-to-100-solutions/>

PDE · SEPARATE SEPTEMBER RELEASE (+ FAMILY 376)

11. Navier–Stokes: can a fluid "blow up" in finite time? (Finite-time blowup with external force)

PDE · a result OpenAI released separately on 2026-09-08 (not one of the 372 families in the October openai/math release), plus the related family 376 from the October release (universal computation in forced NS, §11.7.5). All you need: how to take derivatives, what a partial derivative $\partial/\partial x$ means, and vectors and dot products. You don't need to have studied partial differential equations (PDEs); PDEs are explained from the start. Every number here is either computed by script or sourced. **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. Tags: **[P]** = primary material I read myself (paper PDF, Lean file, official Clay text); **[R]** = secondhand reporting; $\triangle$ = unverified

or disputed. As of 2026-10-08. Estimated reading time: 60–90 minutes. You can split it into three sittings: §11.1–11.3 (PDEs, the fluid equations, what blowup means), §11.4–11.6 (why it's hard, history, the Clay problem statement), §11.7–11.9 (OpenAI's result and the related family 376, its verification, what it means).

11.0 The bottom line: the one thing this chapter explains

The Navier–Stokes (NS) equations describe how water and air flow, and engineers use them every day. But the most basic question about them has no answer: **starting from a smooth, gentle initial state, can the flow velocity become infinite in finite time?**

This chapter explains one causal chain:

The equations contain a nonlinear term where "the fluid pushes itself", and it keeps making the velocity steeper $\Leftarrow$ viscosity is smoothing it out $\Leftarrow$ which of the two wins depends on scale $\Leftarrow$ in 3 dimensions, the only "global ledger" we have (energy) controls less and less at smaller and smaller scales $\Leftarrow$ so nobody can prove that viscosity always wins, and nobody can build an example where it loses.

In 2000 Clay listed this as a Millennium Prize Problem, with **four alternative statements (A)–(D)**; proving any one counts as a solution. (A) and (B) say "with no external force, solutions stay smooth forever"; (C) and (D) say "**there exists** some smooth external force that makes the solution break down in finite time".

On 2026-09-08 OpenAI claimed proofs of (C) and (D): a smooth external force, acting only within a bounded region of space and time, makes a fluid that starts at rest develop infinite velocity at one point at $t = 1$, while total energy stays bounded throughout. It comes with about 640,000 lines of Lean. The same day they also released what may be the more important result: **unforced 3D Euler (inviscid fluid) blowup from smooth initial data**.

After this chapter, you should be able to explain in your own words:

- what a PDE is, and what "smooth solutions exist forever" and "finite-time blowup" each mean;
- why energy conservation cannot stop blowup in 3 dimensions (supercritical scaling, with numbers);
- what each of Clay's (A)–(D) says, and what role the external force f plays in (C)/(D);
- whether OpenAI's result **formally** counts as a solution under Clay's official text, which hurdles remain before "collecting the prize", and why many analysts think the real problem is still (A)/(B).

11.1 What is a PDE: starting with the heat equation

11.1.1 ODEs and PDEs

The differential equations you know, like $y'(t) = -y(t)$, have as their unknown **a single** number $y(t)$ that changes over time. These are called ordinary differential equations (ODEs).

In a partial differential equation (PDE), the unknown is **a whole field**: a value $u(x, t)$ at every position x and every time t . The equation involves both the time derivative $\partial u / \partial t$ and spatial derivatives $\partial u / \partial x$.

Cut space into small cells, one unknown per cell, with neighboring cells influencing each other: a PDE is "**infinitely many coupled ODEs**", like a distributed system in which each node only talks to its neighbors.

11.1.2 The simplest PDE: the heat equation

Take a thin iron rod, and let $u(x, t)$ be the temperature at position x at time t . The heat equation is:

$$\partial u / \partial t = k \cdot \partial^2 u / \partial x^2$$

The term $\partial^2 u / \partial x^2$ on the right is the "curvature": if a point is hotter than both neighbors (like a peak), it is negative and the temperature drops; if it is colder than both neighbors (like a valley), it is positive and the temperature rises. So **the heat equation says "each point moves toward the average of its neighbors"**.

It's easier to see on a grid. With grid spacing Δx and time step Δt , write $r = k\Delta t / \Delta x^2$; one update step is:

$$\text{new}_i = \text{old}_i + r \cdot (\text{old}_{i-1} - 2 \cdot \text{old}_i + \text{old}_{i+1})$$

One step by hand ($r = 0.25$, checked by script): 7 cells `[0, 0, 0, 1, 0, 0, 0]`, with a hot spot in the middle.

- Middle: $1 + 0.25 \cdot (0 - 2 + 0) = 0.5$
- The two neighbors: $0 + 0.25 \cdot (0 - 0 + 1) = 0.25$
- Result: `[0, 0, 0.25, 0.5, 0.25, 0, 0]`, and the sum is still 1.

The hot spot has spread out, but the total heat is unchanged.

[Check 1] Take one more step ($r = 0.25$) from `[0, 0, 0.25, 0.5, 0.25, 0, 0]`. What is the result? What is the sum?

11.1.3 Numerical simulation: the heat equation smooths everything out

Computed by script: $k = 1$, initial temperature 1 where $|x| < 0.5$ and 0 elsewhere (a "box" whose edges are infinitely steep steps), advanced with the scheme above on 1001 grid points.

t	Center temperature (numerical)	Center temperature (exact)	Total heat	Steepest slope
0	1.0000	1.0000	0.990	100 (the steepest step the grid allows)
0.01	0.9996	0.9996	0.990	2.82
0.1	0.7316	0.7364	0.990	0.83
0.5	0.3794	0.3829	0.990	0.22

(The total heat is 0.990 rather than 1 because the box on the grid is 0.99 wide.)

Remember: **the heat equation only makes things smoother**. An infinitely steep step instantly becomes a smooth curve, and the slope keeps falling. It is linear and cannot "blow up" on its own. The **viscosity term** in NS is mathematically "the heat equation applied to the velocity", and it is on the side of "smooth".

11.2 Fluids: the velocity field, and the NS equations term by term

11.2.1 The velocity field

The unknowns for a fluid are the **velocity field** $u(x, t)$ (a 3D vector at every point and every moment) and the **pressure** $p(x, t)$ (a number).

A 2D example: $u(x, y) = (-y, x)$. At the point (1, 0) the velocity is (0, 1), pointing up; at (0, 1) it is (-1, 0), pointing left. Drawn out, it is a rigid rotation **counterclockwise around the origin**, faster the farther from the center (speed = radius r).

11.2.2 Incompressibility: zero divergence

Water is nearly incompressible. Mathematically this is written as **zero divergence**:

$$\nabla \cdot u = \partial u_1 / \partial x + \partial u_2 / \partial y + \partial u_3 / \partial z = 0$$

Divergence measures "how much flows out of a small region, net". Zero divergence = as much flows in as flows out, so volume is preserved.

[Check 2] Compute the divergence of $u = (-y, x)$ and of $w = (x, y)$. Which one is incompressible? What does w look like when drawn?

11.2.3 The NS equations: Newton's second law, written for a small parcel of fluid

The official Clay problem statement (written by Fefferman) says: "Equation (1) is just Newton's law $f = ma$ for a fluid element" [P]. The equations are:

$$\underbrace{\partial u / \partial t + (u \cdot \nabla) u}_{\text{accel. } a} = \underbrace{\nu \Delta u}_{\text{visc}} - \underbrace{\nabla p}_{\text{pres}} + \underbrace{f}_{\text{force}} \quad \nabla \cdot u = 0$$

Term by term:

Left side: the acceleration of a small parcel of fluid. It has two parts:

- $\partial u / \partial t$: the change of velocity over time, as seen from a fixed point.
- $(u \cdot \nabla) u$: the change caused by the parcel **moving to a place where the velocity is different**. Even if the flow field doesn't change in time (steady flow), fluid moving along a streamline can still be turning and speeding up.

$(u \cdot \nabla) u$ is u times the derivative of u , so it is **quadratic in u** . This is where the **nonlinearity** of NS comes from, and it is the source of all the difficulty. Intuitively: the fluid "pushes itself along".

Computing this term by hand for the rotating flow. For $u = (-y, x)$, the first component of $(u \cdot \nabla) u$ is $u_1 \cdot \partial u_1 / \partial x + u_2 \cdot \partial u_1 / \partial y = (-y) \cdot 0 + x \cdot (-1) = -x$; the second component is $(-y) \cdot 1 + x \cdot 0 = -y$. So $(u \cdot \nabla) u = (-x, -y)$, **pointing toward the center, with magnitude r** . This is exactly the centripetal acceleration from high-school physics, $v^2/r = r^2/r = r$.

First term on the right: viscosity $\nu \Delta u$. $\Delta = \partial^2 / \partial x^2 + \partial^2 / \partial y^2 + \partial^2 / \partial z^2$, the same thing as on the right side of the heat equation. When neighboring layers of fluid move at different speeds (shear), they drag on each other and smooth out the velocity difference.

Second term on the right: pressure $-\nabla p$. It pushes from high pressure toward low pressure. In incompressible flow, pressure is "the force that adjusts itself automatically so that $\nabla \cdot u = 0$ holds at every moment", i.e. the Lagrange multiplier for the incompressibility constraint.

Third term on the right: external force f . For example gravity, a stirrer, or an electromagnetic force. **This term is the key to understanding OpenAI's result;** §11.6 is devoted to it.

[Check 3] For the rotating flow $u = (-y, x)$, take $\nu > 0$ and $f = 0$. (a) Compute Δu . (b) This flow is steady ($\partial u / \partial t = 0$), so the equation requires $(u \cdot \nabla) u = -\nabla p$. Find the pressure p . Is the pressure higher at the center or at the edge?

11.2.4 Reynolds number: nonlinearity versus viscosity

Consider a flow with length scale L and speed U :

- the time for the nonlinear term to carry velocity across one length scale is about L/U ;
- the time for viscosity to smooth out differences at this scale is about L^2/ν .

Their ratio is the **Reynolds number** $Re = UL/\nu$. Large Re : nonlinearity dominates (turbulence). Small Re : viscosity dominates (smooth, like honey).

Computed by script (water $\nu \approx 1.0 \times 10^{-6} \text{ m}^2/\text{s}$, air $\approx 1.5 \times 10^{-5} \text{ m}^2/\text{s}$): water flowing at 1 m/s through a 0.1 m pipe has $Re = 10^5$; air flowing at 30 m/s past a 4 m car has $Re = 8 \times 10^6$.

Keep this quantity in mind; §11.4 uses it: **the blowup question is whether, at smaller and smaller scales, the Reynolds number can go up instead of down.**

11.3 "Smooth solutions exist forever" versus "finite-time blowup"

11.3.1 First, an ODE toy: the equation is perfect, but the solution breaks

Compare two equations that look very similar, both with initial value $y(0) = 1$:

t	Solution of $y' = y$: e^t	Solution of $y' = y^2$: $1/(1-t)$
0.5	1.649	2
0.9	2.460	10
0.99	2.691	100

(Computed by script; integrating $y' = y^2$ numerically with RK4 up to $t = 0.99$ gives 100.0000000001, matching the exact solution.)

e^t is finite at every finite time; $1/(1-t)$ **becomes infinite at $t = 1$** . This is called **finite-time blowup**. The right-hand side y^2 could not be smoother; what breaks is the solution: "the bigger it is, the faster it grows", so the growth rate itself is growing.


11.3.2 Adding damping: who wins depends on the initial value

Now add a "viscosity-like" damping term: $y' = y^2 - y$. The substitution $z = 1/y$ solves it exactly: $z(t) = 1 + (z_0 - 1)e^t$, and y blows up if and only if z hits 0.

Initial value y_0	Outcome (computed by script)
0.5	Decays, $y(3) \approx 0.047$
1.0	Stays put, $= 1$ forever (equilibrium)
1.1	Blows up, $t^* = \ln 11 \approx 2.398$
2.0	Blows up, $t^* = \ln 2 \approx 0.693$

For small initial values damping wins; for large ones nonlinearity wins. Compare with the NS equations of §11.2.3: y^2 plays the role of the nonlinear "self-transport" term $(u \cdot \nabla)u$, $-y$ plays the viscosity term $\nu \Delta u$, and the force f is set to 0 for now. The known theory of 3D NS has exactly this shape: Fefferman's statement says that in 3D "(A) and (B) hold provided the initial velocity u^0 satisfies a smallness condition", and that for large initial data, smooth solutions exist at least for a short time $[0, T)$ [P]. Whether large initial data stays smooth forever is the Millennium Problem.

The left half of the figure below puts the curves of §11.3.1–11.3.2 together (the figure calls the unknown x , and $\dot{x}$ means dx/dt): e^t grows but never blows up; $1/(1-t)$ blows up at $t = 1$; $y' = y^2 - y$ blows up at $t = \ln 2 \approx 0.693$ from initial value 2, and decays to 0 from initial value $\frac{1}{2}$. The right half shows Clay's four branches, covered in detail in §11.6: OpenAI claims the forced (C)/(D); the unforced (A)/(B) remain open.



11.3.3 A PDE toy: the Burgers equation

Cut NS down to 1 dimension and drop the pressure: $u_t + u \cdot u_x = \nu \cdot u_{xx}$. Initial value $u = -\sin x$. **Without viscosity**, the slope at $x = 0$ satisfies $(u_x)' = -(u_x)^2$, which has the same form as $y' = y^2$, giving $u_x(0, t) = -1/(1-t)$: at $t = 1$ the slope is infinite (a shock). **With viscosity**, the script simulates it with a spectral method, and the steepest slope is capped at about $1/(2\nu)$: 3.7 for $\nu = 0.1$ ($1/(2\nu) = 5$), 48.4 for $\nu = 0.01$ (50), 164.9 for $\nu = 0.003$ (166.7). The smaller ν , the higher the cap, but it is always finite (1D viscous Burgers can be turned

into the heat equation by the Cole–Hopf transformation, and it is rigorously proven not to blow up).

The question for 3D NS is: can viscosity still cap things like this? 3D has one extra mechanism: **vortex stretching**. A vortex tube that gets stretched longer and thinner spins faster, like a figure skater pulling in their arms (angular momentum $r \cdot u_\theta$ is conserved: shrink the radius 10x, and the spin speed goes up 10x). This positive feedback does not exist in 1D or 2D.

11.3.4 Saying "blowup" precisely

"Smooth" means derivatives of all orders exist and are continuous (C^∞). **Global smooth solution:** smooth for all $t \geq 0$ with bounded energy. **Finite-time blowup:** the solution is smooth on $[0, T)$, but as $t \rightarrow T$, the velocity or one of its derivatives $\rightarrow \infty$ somewhere. Fefferman's statement records that for NS ($\nu > 0$), at blowup **the velocity itself** must be unbounded [P], which is why OpenAI's theorem is stated as " $\limsup \|u(t)\|_\infty = \infty$ ". Real water never has infinite velocity; the continuum model breaks down at molecular scales first. Blowup is a statement **about this mathematical model**.

11.4 Why 3D is so hard: energy and supercritical scaling

11.4.1 Energy: the only global ledger

With no external force, NS has a beautiful identity: the total kinetic energy $E(t) = \frac{1}{2} \int |u|^2 \, dx$ **can only decrease, never increase** (viscosity turns kinetic energy into heat). The nonlinear term $(u \cdot \nabla)u$ contributes exactly 0 to the total energy; it only **moves** energy between scales and never creates it.

So we have an upper bound valid for all time: $E(t) \leq E(0)$. The question is: **is this upper bound enough to prevent blowup?**

11.4.2 Scaling symmetry

NS has a symmetry: if $u(x, t)$ is a solution, then for any $\lambda > 0$,

$$u_\lambda(x, t) = \lambda \cdot u(\lambda x, \lambda^2 t)$$

is also a solution (with pressure and external force transformed accordingly). $\lambda > 1$ means "shrink the same flow to $1/\lambda$ of its size, play it λ^2 times faster, and the flow speed becomes λ times larger".

λ	Size	Time	Flow speed	3D energy	2D energy
1	1	1	$\times 1$	$\times 1$	$\times 1$

λ	Size	Time	Flow speed	3D energy	2D energy
10	1/10	1/100	$\times 10$	$\times 0.1$	$\times 1$
100	1/100	1/10 ⁴	$\times 100$	$\times 0.01$	$\times 1$
1000	1/1000	1/10 ⁶	$\times 1000$	$\times 0.001$	$\times 1$

Where the energy column comes from: take $\int |\lambda u(\lambda x)|^2 dx$ and set $y = \lambda x$; in 3D, $dx = dy/\lambda^3$, so the energy is multiplied by $\lambda^2/\lambda^3 = 1/\lambda$.

[Check 4] Compute the 2D energy factor the same way, and verify that it is 1.

The 3D column is the key: shrink a "blowup structure" 1000 $\times$, speed its flow up 1000 $\times$, and the energy it needs is only one thousandth of the original. The energy bound places almost no constraint on small scales. This is called **supercritical**. In 2D the energy does not change under scaling, which is called **critical**; that is why 2D NS was proven long ago to be globally smooth (Ladyzhenskaya; see Fefferman's statement [P]).

11.4.3 Saying it again with the Reynolds number

Fix the total energy $E = 1$ and the viscosity $\nu = 1$. For a blob of flow of size L , how large can its speed be?

- 3D: energy $\approx U^2 L^3 \leq 1$, so $U \leq L^{-3/2}$, and the Reynolds number $Re = UL/\nu \leq L^{-1/2}$.
- 2D: energy $\approx U^2 L^2 \leq 1$, so $U \leq 1/L$, and $Re \leq 1$.

Size L	Max Re allowed in 3D	Max Re allowed in 2D
1	1	1
10 ⁻²	10	1
10 ⁻⁶	1,000	1
10 ⁻⁸	10,000	1

(Computed by script.)

2D: energy holds the Reynolds number at every scale below 1, so viscosity always has what it takes to win. 3D: the smaller the scale, the larger the Reynolds number the energy allows, so **viscosity might lose at small scales**. Energy neither rules out blowup nor proves it; it is simply "not enough". This is the core of the 3D NS problem.

11.4.4 Tao's "supercriticality barrier" (2016)

Tao (JAMS 29, 2016; arXiv:1402.0290) constructed an "averaged NS": it keeps the energy identity and the scaling, replaces only the nonlinear term with an averaged version, and he proved that it **does** blow up. Consequence: an argument that uses only "energy + generic

harmonic-analysis estimates" cannot prove (A)/(B); it must use the fine structure of the real nonlinear term. He also proposed a **program**: build "logic gates" out of a real fluid, so that the fluid, like a self-replicating machine, passes its energy again and again to smaller, faster copies of itself, producing blowup in **unforced** NS. This program has not been carried out to this day. (Family 376 in §11.7.5 is reminiscent of this program, but it is not the same thing.)

11.5 90 years of partial results

Year	Result	Significance
1934	Leray (Acta Math. 63): every finite-energy initial value has a global weak solution	Weak solutions satisfy the equation only "in an integral-average sense" and allow singularities; smoothness and uniqueness are unknown
1982	Caffarelli–Kohn–Nirenberg (CPAM 35): the singular set has parabolic 1-dimensional Hausdorff measure 0	Singularities "cannot form a curve in space-time" (Fefferman's wording [P]), but isolated points are not ruled out
2003	Escauriaza–Seregin–Šverák : if the L^3 norm stays bounded, there is no blowup	Blowup must make the L^3 norm diverge
2016	Tao : averaged NS blows up	See §11.4.4
2019	Buckmaster–Vicol (Annals 189): a class of weak solutions is non-unique	Weak solutions are too "soft"
2014 / 2022	Luo–Hou numerics (PNAS); Chen–Hou computer-assisted proof (arXiv:2210.07191): 3D axisymmetric Euler with a boundary blows up from smooth initial data	Needs a solid wall; not in all of space $\mathbb{R}^3$
2021	Elgindi (Annals 194): Euler on $\mathbb{R}^3$ blows up for $C^{(1,\alpha)}$ initial data	The initial data is not C^∞
2023 on	Córdoba–Martínez–Zoroa (arXiv:2309.08495 and others): forced 3D Euler blows up, with a force of limited smoothness	Pioneers of the "forcing route": finer and finer vortex layers amplified stage by stage
2025	DeepMind et al. (arXiv:2509.14185): neural networks find unstable self-similar singularities for the Boussinesq, IPM and other equations, with accuracy close to machine precision	Paves the way for computer-assisted proofs in the boundaryless case
2026-08	Buckmaster–Alpöge : IPM, Boussinesq and 3D Euler blow up under smooth forcing , verified in Lean; made public September 7 (US Eastern time)	Pushes the Córdoba–Martínez–Zoroa route to smooth forcing [P, Buckmaster statement]

The pattern that keeps appearing in this table: **Euler (no viscosity) blows up more easily than NS; having a boundary, lower smoothness, or forcing is each easier than "all of space, C^∞ , no forcing"**. Every time one condition is relaxed, someone manages to produce blowup. Millennium Problem (A)/(B) is the version with no condition relaxed.

11.6 The Clay problem statement: four branches, and the role of the force f

11.6.1 The official text [P]

Fefferman's statement (<https://www.claymath.org/wp-content/uploads/2022/06/navierstokes.pdf>) first sets out "physically reasonable" conditions: all derivatives of the initial data and the external force decay rapidly (conditions 4, 5), the solution is C^∞ on $\mathbb{R}^3 \times [0, \infty)$ (6), and the energy is uniformly bounded (7); the periodic case uses conditions 8–11. Then:

"To give reasonable leeway to solvers while retaining the heart of the problem, we ask for a proof of one of the following four statements."

Branch	Domain	External force	What must be proven
(A)	$\mathbb{R}^3$	$f \equiv 0$	For every smooth decaying initial value $\rightarrow$ there is a global smooth solution with bounded energy
(B)	Periodic torus $\mathbb{R}^3/\mathbb{Z}^3$	$f \equiv 0$	Same, periodic version
(C)	$\mathbb{R}^3$	There exists a smooth decaying f	There exist some initial value and force for which no smooth, bounded-energy global solution exists
(D)	$\mathbb{R}^3/\mathbb{Z}^3$	There exists a smooth periodic f	Same, periodic version

Two points. (A)/(B) say "good for all initial values", while (C)/(D) say "there is one bad example": **opposite directions**. (A)/(B) require $f \equiv 0$, while (C)/(D) let you **choose the force yourself**, as long as it is smooth and decaying. Fefferman also writes that the corresponding problem for the Euler equations ($\nu = 0$) is "also open and very important", but it is not on the prize list [P].

11.6.2 Why "you can choose the force" changes the shape of the problem

§11.2 of the OpenAI paper says this itself [P]:

"For any incompressible flow u and pressure p , we can always define the external force f to be the residual in (1.1). The Navier–Stokes equations then hold by construction. The challenge is to choose a flow that blows up while this residual remains smooth."

In plain words: draw any incompressible flow that blows up, plug it into NS, subtract the right side from the left, and let f make up whatever the difference is. The equation then holds **automatically**. There is only one difficulty: **the force making up the difference must stay smooth**, and cannot break even at the blowup time.

An ODE toy makes the difficulty clear:

- **Linear equation $y' = -y + g(t)$:** to make $y = 1/(1-t)$ blow up, the needed force is $g = y' + y = 1/(1-t)^2 + 1/(1-t)$, which itself blows up and is not smooth. In fact, a linear equation with a smooth force **never** blows up.
- **Nonlinear equation $y' = y^2 + g(t)$:** for the same $y = 1/(1-t)$, $g = y' - y^2 = 1/(1-t)^2 - 1/(1-t)^2 = 0$. The force is 0, perfectly smooth.

So blowup must rely on the nonlinear term to cancel the singular part **by itself**, with the force only making up a smooth remainder. In NS, all divergent parts must cancel, to derivatives of every order; this is what those 166 pages are doing.

[Check 5] For $y' = -y + g(t)$, if g is smooth and bounded on $[0, 2]$ (say $|g| \leq 1$) and $y(0) = 0$, can $y(t)$ blow up at $t = 1$? Hint: compare with $y' \leq -y + 1$.

11.6.3 So is (C) easier than (A)?

In 2000 Clay put (C) side by side with (A), considering all four branches to be "retaining the heart of the problem". But the "forcing route" opened by Córdoba–Martínez–Zorua from 2023 on showed that **the force can do a great deal of "engineering" work**: inject energy where and when it is needed, and "plant" perturbations. In his statement Buckmaster says this route "is the route Luis and Diego opened", and calls the next goal "a path to unforced Euler" [P]. Even the people pushing this route see "unforced" as the deeper problem.

In other words, (C)/(D) turn the problem into "can you build a flow that blows up while the residual force stays smooth?" That is much weaker than whether the fluid blows up **on its own** (the unforced (A)/(B)), and most analysts regard (A)/(B) as the real question [R]. A secondhand summary: the forcing route is "a route the written problem permits but that most working mathematicians exclude from the question they care about" [R, Implicator]. $\triangle$ This is a journalist's summary, not a systematic survey of analysts.

11.7 What OpenAI claims

11.7.1 The main theorem [P]

The paper is *Finite time blowup for Navier–Stokes*, authored by "OpenAI", 166 pages, PDF generated 2026-09-08 12:06 PDT (<https://cdn.openai.com/pdf/32d9f210-8b73-45e0-91bc-82a30aef8a9a/navier-stokes.pdf>). It was **released separately**: it is not one of the 372 families in the October 6 openai/math release; the paper is on OpenAI's CDN, and the roughly 640,000-line Lean proof is in a separate repository, `openai/NavierStokesAndEuler` (§11.8.2). Abstract:

"For every positive viscosity, we construct a solution of the three-dimensional incompressible Navier–Stokes equations that starts from rest and develops unbounded velocity in finite time while maintaining uniformly bounded kinetic energy."

What Theorem 1.1 says: for every $\nu > 0$, there exist

- an external force $f \in C_c^\infty(\mathbb{R}^3 \times (0, \infty))$: smooth, and **nonzero only within a bounded region in both space and time** ("compactly supported");
- velocity and pressure u, p , smooth on $\mathbb{R}^3 \times [0, 1]$, with **initial value $u(\cdot, 0) = 0$** (starting from rest), and confined at all times to a fixed bounded region K ;
- with energy $\sup \|u(t)\|_{L^2} < \infty$ and velocity $\limsup_{t \uparrow 1} \|u(t)\|_{L^\infty} = \infty$.

The paper's own words: "This establishes alternative (C) in the Millennium problem statement ... Compact support also yields the corresponding construction on T^3 ... establishing alternative (D)" (Corollary 10.6).

Arbitrary viscosity, by hand. The paper first builds the construction at $\nu = 1$, then extends to any ν by the scaling $u_\nu(x, t) = \sqrt{\nu} \cdot u(x/\sqrt{\nu}, t)$. The blowup time does not change.

[Check 6] Verify: substituting $u_\nu = \sqrt{\nu} \cdot u(x/\sqrt{\nu}, t)$, each of the three terms $\partial u_\nu / \partial t$, $(u_\nu \cdot \nabla) u_\nu$ and $\nu \Delta u_\nu$ equals $\sqrt{\nu}$ times the corresponding original term.

11.7.2 Intuition for the mechanism (from reading Section 2 of the paper; a simplified version)

1. **A vortex column that gets thinner and faster.** In the vortex core near the origin, fluid spirals inward and flows out up and down along the axis. Conservation of angular momentum makes the inflowing fluid spin faster; incompressibility makes the incoming fluid leave along the axis, so the inward flow can keep going.
2. **Self-similar contraction.** Write $\tau = 1 - t$. The core radius $\ell_r \sim \tau^{1/2}$, the height $\ell_z \sim \tau^{1/2-h}$, where h is a fixed small number, $0 < h < 1/100$. The vortex column gets thinner and relatively longer.
3. **The problem:** the transition annulus outside the core leaves a "residual", and the force needed to make it up diverges as $t \rightarrow 1$.
4. **The fix: add oscillating pulses in the annulus.** The pulses have zero average velocity, but their **quadratic products** (momentum flux) do not average to zero. Two families of pulses are carefully arranged so that the nonlinear flux **exactly cancels** the divergent part of the residual. Each pulse is "planted" by an exponentially small force, and then grows by feeding on the background shear.
5. **Outside:** a purely rotating flow that satisfies the heat equation and needs no force; then smoothly cut off to a bounded region.

This is exactly §11.6.2's "the nonlinear term cancels the singular part by itself", except that it must be done in 3D for derivatives of every order.

11.7.3 A feel for the numbers: how the vortex core shrinks

Take $h = 0.005$ (the paper only requires $0 < h < 1/100$; I chose this value for illustration), and compute with the scaling laws the paper gives (computed by script, $v = 1$, constant factors set to 1):

$\tau = 1 - t$	Core radius ℓ_r	Core height ℓ_z	Speed $\sim \tau^{(-1/2-h)}$	Core energy $\sim \tau^{(1/2-3h)}$	Swirl Reynolds number $\tau^{(-h)}$
10^{-2}	0.1	0.102	10.2	0.107	1.023
10^{-4}	0.01	0.0105	104.7	0.0115	1.047
10^{-8}	10^{-4}	1.10×10^{-4}	1.10×10^4	1.3×10^{-4}	1.096
10^{-12}	10^{-6}	1.15×10^{-6}	1.15×10^6	1.5×10^{-6}	1.148

- **Speed $\rightarrow \infty$, core energy $\rightarrow 0$:** this is exactly the supercriticality from §11.4, reaching ever larger speeds at small scales with ever less energy.
- The speed carries an extra factor $\tau^{(-h)}$ beyond pure scaling $\tau^{(-1/2)}$; this is the **swirl Reynolds number**. It goes to infinity, but extremely slowly: as τ goes from 10^{-2} to 10^{-12} , it only rises from 1.02 to 1.15. The blowup is achieved right at the edge, "just barely beyond critical".
- Consistent with known theorems: the singularity is **a single point** (the origin, $t = 1$), which does not violate CKN; the cube of the L^3 norm $\sim \tau^{(-4h)}$, about 1.74 at $\tau = 10^{-12}$, and it is diverging, which fits the necessary condition from ESS ($\triangle$ ESS is about the unforced case; here it serves only as an intuition check).

11.7.4 The Euler result released at the same time

The same Lean repository also formalizes a second paper, *Finite time blowup for the Euler equation* (57 pages, <https://cdn.openai.com/pdf/315b36cd-ec98-4023-8342-93345194ece1/euler.pdf>) [P]:

"We exhibit finite-time blowup for the three-dimensional incompressible, unforced Euler equations from smooth, compactly supported, divergence-free initial data."

Compare with §11.5: Chen–Hou need a boundary, Elgindi's initial data is only $C^{(1,\alpha)}$, and Córdoba–Martínez-Zorúa and Buckmaster–Alpöge need a force. This result removes all of those relaxations; the only remaining difference from the Millennium Problem is "no viscosity". Fefferman says the corresponding problem for Euler is "also open and very important" [P]. My judgment (speculation): in terms of how close it is to the question people actually care about, it may carry more weight than the forced NS result. $\triangle$ I found no named expert's technical assessment of this paper.

Scale [R]: about 10,000 concurrent agents, 88 hours; Wikipedia cites 2.7 million messages and 130 billion output tokens, while other reports say "nearly 5 million" $\triangle$.

`formalization.yaml` records the model as "GPT-6 Astra" and the framework as Codex [P].

11.7.5 Related: universal computation in forced NS (family 376)

The October openai/math release contains a result family that bears directly on "the force": family 376, *Universal computation in forced Navier–Stokes flows*. It consists of a group of preprints; the representative one is *A Fixed Particle Test for Computation in a Forced Viscous Flow* (2026-09-27) [P].

Claim [P]: fix a flat 3D torus, a positive computable viscosity, a fluid starting at rest, one designated fluid particle, and a fixed open region (the "detector"). For any given Turing machine and input, one can write a finite "program" for a smooth force such that the particle enters the detector **if and only if** that Turing machine halts. The halting problem is undecidable, so the reachability question "will this particle ever get there?" is **undecidable** too.

The paper is candid about it [P]: "There is an elementary way to force a chosen smooth incompressible velocity: subtract its viscous term from its material acceleration. The mathematical work lies in choosing that velocity from the finite instruction table." This is the "residual force" trick of §11.6.2: pick the flow you want and let the force make up the difference. The real work is designing that flow: it has to execute the instructions one by one, keep the information needed for invertibility, and never touch the detector during a run that does not halt. The paper also states that it is "not a regularity theorem for general Navier–Stokes data".

Relation to Tao's program (§11.4.4) (our judgment): it recalls Tao's 2016 program, which is to get the fluid **itself** to build a computer that hands its energy to ever-smaller copies of itself, causing blowup. But 376 puts the "programming" into the force, which does all of the programming work, so it **does not advance** that unforced-blowup program; it is closer to a controllability / undecidability result for forced fluids. As of 2026-10-09, none of the 9 preprints the repository lists in this family cites Tao 2016 [P]. The precursors it does cite are Christopher Moore's generalized shifts (1990–91), Bennett's reversible computation (1973), and the Turing-complete stationary Euler flows that Cardona–Miranda–Peralta-Salas–Presas built on a 3-sphere with an adapted metric (PNAS 2021) [P].

Lean [P]: `lean/docs/376.md` lists 6 Comparator statements covering selected results from 4 of the preprints, including a version on $\mathbb{R}^3$: a smooth, compactly supported flow starting from rest in which the particle at (4,0,0) enters the fixed open box $(-1,2)^3$ exactly when the machine halts. Coverage differs from paper to paper; for example, the fixed-particle halting detector of *Geometric Programs for Solenoidal Forcing* is explicitly outside the formalized scope. ⚠ We did not read the 6 statement files one by one, and did not compile them.

11.8 How far it has been verified

11.8.1 The Lean statement: OpenAI did not write it

The statement file `ComparatorChallenges/NavierStokes.lean` says at the top that it is copied from **Google DeepMind's Formal Conjectures project** (commit `8bf45ed`), with only the imports, namespace and notation changed [P]. If the solver writes the statement itself, it can quietly loosen conditions; here it comes from a third party. $\triangle$ I did not diff it line by line against upstream.

The statement of branch (C):

```
theorem navier_stokes_breakdown_R3 (nu : ℝ) (hnu : nu > 0) :  
  ∃ (u₀ : ℝ³ → ℝ³) (f : ℝ³ → ℝ → ℝ³),  
    InitialVelocityConditionDecay u₀ ∧ ForceConditionDecay f ∧  
    ¬ (∃ v p, NavierStokesExistenceAndSmoothnessRn nu u₀ f v p)
```

Word-by-word translation:

- `(nu : ℝ) (hnu : nu > 0)`: for any positive viscosity ν ,
- `∃ (u₀ ...) (f ...)`: there exist an initial value u_0 and a force f ,
- `InitialVelocityConditionDecay u₀`: u_0 is smooth, divergence-free, and all its derivatives decay rapidly (Clay condition 4),
- `ForceConditionDecay f`: f is smooth, and all its space-time derivatives are $\leq C/(1 + |x| + t)^K$ (condition 5),
- `¬ (∃ v p, ...)`: such that there **do not exist** v, p satisfying all of NS (1), zero divergence (2), the initial value (3), C^∞ (6), and uniformly bounded energy (7).

I checked it item by item against the Clay text: (4)–(7) are all there; the periodic version (D) includes (8)–(11), and, as Clay's erratum requires, the pressure is also periodic [P]. **The structure of the statement matches Clay's (C)/(D).** Note that (C) asks that "no global smooth solution exists at all", not just "the solution I built blows up"; this requires uniqueness of smooth solutions (the "Consequently..." sentence in Theorem 1.1), and that is proven in Lean too.

11.8.2 The proof itself

- Repository `openai/NavierStokesAndEuler` (snapshot `f9e8bc5`, 2026-09-10): **2,659 .lean files** (NS 816, Euler 1,839) [P], about 640,000 lines (a count from the author's research notes).
- `formalization.yaml`: sorry count 0; uses only the three standard axioms `propext`, `Classical.choice`, `Quot.sound`; Comparator has `enable_nanoda` turned on, i.e. a

second, independently implemented kernel also checks it [P]. The review status reads **"self-assessed"**.

- ⚠ I did not run `lake build` or Comparator, and found no public report of a third party running the NS Comparator successfully (the quasi-Riemann hypothesis in §4 had Goldblatt's independent recheck).

11.8.3 Human experts and institutions

- **Clay** (09-11, president Martin Bridson) [R]: reportedly said "the Navier–Stokes problem has apparently been settled", and also that evaluation will be "deliberately unhurried" (full sentence: "The process is deliberately unhurried, but we will provide updates"). Clay still lists NS as unsolved: the problem page on Clay's website is still marked "Active" as of today (I checked on 10-08). Clay's rules require publication in a qualifying outlet plus a two-year acceptance period (§11.9.1).
- **Gómez-Serrano** (NPR, 09-22) [R]: "The paper is not written for humans ... as of today, the paper doesn't teach us much."
- **Fefferman** [R]: "The heroes of the story ... are Córdoba and Martínez-Zorúa." ⚠ Original source not found.
- ⚠ **As of 2026-10-08, no human expert has publicly said they have read and verified the 166-page paper**, and there are no reports of journal submission.

11.8.4 The priority dispute (in brief)

- **Buckmaster's statement** (<https://cims.nyu.edu/~tristanb/statement.pdf>) [P]: Buckmaster (NYU) and Alpöge obtained Boussinesq and 3D Euler blowup under smooth forcing on August 15, verified it in Lean on August 22, and made it public, together with the IPM result, on September 7 (US Eastern time). The statement itself only says "Today"; September 7 is per Wikipedia [R], and the statement PDF's metadata says 2026-09-08 02:38 UTC, i.e. the evening of September 7 in US Eastern time [P], which is September 8 in UTC. They also used LLMs heavily (Claude, Codex). He credits the route to Córdoba and Martínez-Zorúa. About OpenAI, he recounts that in a call on September 6, "Eventually it was agreed that [the first prompt] had been sent in the past few days, after information about our work had reached OpenAI", and also writes "I am not accusing anyone of anything."
- **OpenAI's response** [R]: "it is impossible for Dr. Buckmaster's Codex prompts over the last two months to have influenced the system in any way".
- **The accounts differ** ⚠: Buckmaster's account is that OpenAI sent its first prompt only after information about their progress had reached OpenAI; what OpenAI denies is that his Codex prompts influenced its system. The two claims address different questions (one about timing, one about whether their user data was used), and neither has been verified by a third party.

- **Citations** [P]: the PDF I read (generated 09-08 12:06 PDT) **already cites** three works by Córdoba–Martínez-Zoroa [6–8]; it **does not cite** Buckmaster–Alpöge.

11.8.5 Remaining risks

1. **Mathlib definitions:** the Laplacian, gradient, etc. come from Mathlib, and "junk value conventions" could make the statement drift from its intended meaning; but the statement requires C^∞ everywhere, so the effect is small.
2. **Toolchain:** Lean 4.34.0-rc2 is a pre-release version [P]; the nanoda dual-kernel check lowers the risk.
3. **Human understanding:** nobody can say whether this method extends to the unforced case.

11.9 What it means if true (and what it doesn't)

11.9.1 By Clay's text, does this count as "solving the Millennium Problem"?

The answer has three layers.

Layer one: the problem statement. Yes. The official text says "we ask for a proof of one of the following four statements", and (C) is one of them [P]. Rule 5(b) of the prize rules also states: "In the case of the P versus NP Problem and the Navier-Stokes Problem, a resolution in either direction will be evaluated by the standard evaluation procedure" [P, https://www.claymath.org/wp-content/uploads/2022/03/millennium_prize_rules_0.pdf]. Both directions are equally valid.

Layer two: the prize procedure. Still a long way off. Rule 4 requires all of the following [P]:

- (a) publication in a "qualifying outlet" (no reports of submission so far △);
- (b) at least **2 years** have passed since publication;
- (c) "general acceptance" in the global mathematics community, as judged by CMI;
- (d) a satisfactory answer to the question posed in the official problem statement.

After that, CMI must also form a committee of at least 3 people to review it (Rule 7). Rule 8(b): CMI will "pay special attention to the question of whether a Prize solution depends crucially on insights published prior to the solution", and may add earlier contributors to the list of prizewinners; this is directly relevant to Córdoba–Martínez-Zoroa and Buckmaster–Alpöge. Rule 8(c): if no conclusion can be reached on correctness or **attribution**, the prize may not be awarded. Separately, OpenAI has reportedly said it will not claim the prize [R]. Speculation: even if it is published next year, with "2 years after publication", a conclusion would come around 2029 at the earliest.

Layer three: the question mathematicians care about. Most would say no. "Does a fluid blow up on its own?" is still completely open. OpenAI's result does not answer the unforced

case, and does not carry out Tao's program (§11.4.4; nor does family 376 in §11.7.5).

Be careful with the "specification gaming" framing. It is tempting to compare this to specification gaming in AI: the literal text of the problem is satisfied, but the question the askers care about is not answered. But this is not cheating: Fefferman included (C) **on purpose** in 2000, as "retaining the heart of the problem" (§11.6.1). A more accurate way to put it: **the problem setters thought at the time that (C) was as hard as (A); progress since 2023 shows that the forcing route is easier than expected.** The gap between the text and the intent became visible in hindsight; it is not a case of the solver exploiting a loophole the setters never thought of.

11.9.2 For mathematics, physics and AI

1. **Mathematics:** if it holds, this is the first NS blowup in the setting of smooth forcing, all of space, and C^∞ : viscosity is not all-powerful. The unforced Euler blowup released at the same time (§11.7.4), if it holds, may be the bigger news. Methodologically, "synthesizing" the needed stress from the quadratic flux of oscillating pulses is related to convex integration for Euler (Daneri–Székelyhidi), as the paper itself says [P].
2. **Physics and engineering:** almost no direct impact on weather forecasting, turbulence modeling or engineering. NS solvers never rely on global regularity theorems; in real fluids, molecular scales make the continuum model fail before a singularity is reached. The conclusion is **"the model itself** breaks under an artificially designed force", not "real water can reach infinite velocity".
3. **AI:** three independent gates: passing formalization $\neq$ a faithful statement $\neq$ acceptance by mathematicians. Here the second gate is fairly reliable because a third-party statement was used; the third gate is far from passed. Buckmaster says LLMs mean one mathematician plus a model "can now do all this work in a month", and calls it "a Deep Blue-Kasparov moment" [P]. When an AI lab can quickly pour large amounts of compute into a direction after hearing it is making progress, how norms of priority and attribution should adapt is a real question.

11.10 Self-test

1. In one sentence: why does $y' = y^2$ blow up in finite time while $y' = y$ does not? (§11.3.1)
2. Which term in the NS equations is nonlinear? What is its physical meaning? (§11.2.3)
3. Why is energy for 3D NS called "supercritical"? Explain using the numbers for scaling $\lambda = 1000$. (§11.4.2–11.4.3)
4. How do Clay's (A) and (C) differ in "external force" and in "quantifier (for all / there exists)"? (§11.6.1)

5. Why does "you can choose the force yourself" not make the problem trivial? What is the key difficulty? (§11.6.2)
6. Who wrote OpenAI's Lean problem statement? Why does that matter? (§11.8.1)
7. Under Clay's rules, what steps at minimum remain between now and a possible award? (§11.9.1)
8. Why does the fluid "computer" of family 376 not advance Tao's unforced-blowup program? (§11.7.5)

Appendix: answers to the checks

- **Check 1:** $[0, 0.0625, 0.25, 0.375, 0.25, 0.0625, 0]$, sum 1 (checked by script). For example, the middle: $0.5 + 0.25 \cdot (0.25 - 1 + 0.25) = 0.375$.
- **Check 2:** $\nabla \cdot (-y, x) = \partial(-y)/\partial x + \partial x/\partial y = 0 + 0 = 0$, incompressible. $\nabla \cdot (x, y) = 1 + 1 = 2 \neq 0$: w is a flow "spraying" outward from the origin; every small region has net outflow, and volume is expanding.
- **Check 3:** (a) Each component of $u = (-y, x)$ is a linear function, so all second derivatives are 0 and $\Delta u = 0$: rigid rotation has no shear, so viscosity does nothing. (b) From §11.2.3, $(u \cdot \nabla)u = (-x, -y)$, so $\nabla p = (x, y)$ and $p = (x^2 + y^2)/2 + \text{constant}$. **Higher at the edge, lower at the center:** the pressure difference pushes inward from outside, supplying the centripetal force. When the OpenAI paper describes the vortex core as "Pressure decreases toward the axis, supplying the leading centripetal force", this is what it means.
- **Check 4:** In 2D, $dx = dy/\lambda^2$, so the energy factor is $\lambda^2/\lambda^2 = 1$.
- **Check 5:** No. From $y' \leq -y + 1$ and $y(0) = 0$ you can derive $y \leq 1 - e^{-(t)} < 1$; likewise $y \geq -1$. A linear equation with a bounded force has a solution that stays bounded.
- **Check 6:** $\partial u_v/\partial t = \sqrt{v} \cdot \partial u/\partial t$. $(u_v \cdot \nabla)u_v$: the velocity factors give $\sqrt{v} \times \sqrt{v}$, and the derivative ∇ acting on $x/\sqrt{v}$ brings $1/\sqrt{v}$, for $\sqrt{v}$ in total. $v \Delta u_v$: $v \times \sqrt{v} \times (1/\sqrt{v})^2 = \sqrt{v}$. All three terms agree.

Appendix: Glossary

Term	One-line explanation
ODE / PDE	The unknown is a single number changing over time / a whole field
Heat equation	$\partial u/\partial t = k \Delta u$; each point moves toward the average of its neighbors, which only makes things smoother
Divergence $\nabla \cdot u$	Net outflow from a small region; $= 0$ means incompressible
Material derivative	$\partial u/\partial t + (u \cdot \nabla)u$, the acceleration of a small parcel of fluid itself

Term	One-line explanation
Viscosity $\nu \Delta u$	Neighboring fluid layers drag on each other; mathematically, the heat equation for the velocity
Reynolds number Re	UL/ν , the ratio of the strength of nonlinearity to viscosity
Weak solution	A solution that satisfies the equation only in an integral-average sense, allowing singularities (Leray 1934)
Scaling symmetry	$u_\lambda = \lambda u(\lambda x, \lambda^2 t)$ is still a solution
Supercritical	The conserved quantity (energy) controls less and less at small scales
Vortex stretching	A vortex tube stretched longer and thinner spins faster; only in 3D
Euler equations	NS with $\nu = 0$, an inviscid fluid
Clay (A)–(D)	The four alternative statements of the Millennium Problem; proving any one is enough
Undecidable	No algorithm gives the correct yes/no answer for every input; the halting problem for Turing machines is the classic example

Appendix: Further reading

- **The official Clay problem statement** (Fefferman, 5 pages, strongly recommended to read in full): <https://www.claymath.org/wp-content/uploads/2022/06/navierstokes.pdf>
- **Clay prize rules**: https://www.claymath.org/wp-content/uploads/2022/03/millennium_prize_rules_0.pdf
- **OpenAI NS paper**: <https://cdn.openai.com/pdf/32d9f210-8b73-45e0-91bc-82a30aef8a9a/navier-stokes.pdf> (read Sections 1–2 first, about 6 pages, needing almost no PDE background)
- **OpenAI Euler paper**: <https://cdn.openai.com/pdf/315b36cd-ec98-4023-8342-93345194ece1/euler.pdf>
- **Lean repository**: <https://github.com/openai/NavierStokesAndEuler> ; upstream problem statement: <https://github.com/google-deepmind/formal-conjectures/blob/main/FormalConjectures/Millennium/NavierStokes.lean>
- **Buckmaster's statement**: <https://cims.nyu.edu/~tristanb/statement.pdf>
- **Family 376** (universal computation in forced NS), representative paper: <https://github.com/openai/math/blob/main/preprints/A-Fixed-Particle-Test-for-Computation-in-a-Forced-Viscous-Flow-September-27-2026/manuscript.pdf> ; Lean scope notes: <https://github.com/openai/math/blob/main/lean/docs/376.md>
- **Terence Tao**, arXiv:1402.0290 (JAMS 2016): the introduction lays out the "supercriticality barrier" clearly; more accessible is his blog post *Why global regularity for Navier-Stokes is hard* (2007).

- **DeepMind et al.**, *Discovery of Unstable Singularities*, arXiv:2509.14185
- **Chen & Hou**, *Stable nearly self-similar blowup of the 2D Boussinesq and 3D Euler equations with smooth data I: Analysis*, arXiv:2210.07191
- **Córdoba & Martínez-Zoroa**, *Blow-up for the incompressible 3D-Euler equations with uniform $C^{1,1/2-\varepsilon} \cap L^2$ force*, arXiv:2309.08495
- **Bertozzi & Majda**, *Vorticity and Incompressible Flow* (Cambridge, 2002): the textbook recommended in Fefferman's problem statement.
- Secondhand reporting (not checked word for word): Quanta 2026-09-08
<https://www.quantamagazine.org/ai-has-solved-one-of-maths-1-million-millennium-prize-problems-20260908/> ; Wikipedia "Navier–Stokes priority controversy"
https://en.wikipedia.org/wiki/Navier%E2%80%93Stokes_priority_controversy ; paraphrase of Clay's 09-11 statement <https://the-decoder.com/clay-mathematics-institute-says-the-navier-stokes-millennium-prize-problem-has-apparently-been-settled/> ; Implicator <https://www.implicator.ai/clay-institute-navier-stokes-openai-proof-claim/>

MATHEMATICAL PHYSICS · FAMILY 362

12. Relativistic Vlasov–Maxwell: can a plasma model break down by itself?

Prerequisites: you can take a derivative, you know the dot product and cross product of vectors, and you know that "a differential equation describes how something changes over time." You don't need to have studied partial differential equations (PDEs) or plasma physics. Every new concept is explained before it is used. Every number in the text was computed by script or has a source; [P] marks primary material read (paper, repo, Lean source), and $\triangle$ marks something unverified or conflicting. **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter. Estimated reading time: 60–90 minutes. You can split it into three sittings: §12.1–12.4 (plasma, phase space, the equations themselves), §12.5–12.7 (relativity, what "the solution breaks down" means, 40 years of history), §12.8–12.10 (OpenAI's result, how well it is verified, and what it means).

12.0 The bottom line: the one thing this chapter explains

Plasma (the ionized gas inside a fusion reactor or in the solar wind) has a standard mathematical model: the **relativistic Vlasov–Maxwell system**. It says: charged particles are pushed around by the electromagnetic field, the particles' motion creates currents, and those currents in turn change the electromagnetic field. This loop is nonlinear.

The question mathematicians want answered is: **starting from a normal initial state, can a solution of this model "break down" in finite time?** That is, can some quantity suddenly become infinite, so that the equations stop making sense from then on?

This chapter explains one chain of reasoning:

Whether the solution breaks down $\iff$ (Glassey–Strauss 1986) depends only on "whether the momentum of the fastest particle can shoot to infinity in finite time" $\iff$ if you can prove "each doubling of momentum takes at least $c/\log P$ time" $\iff$ these times add up to a divergent series (like the harmonic series $1 + 1/2 + 1/3 + \dots$), so momentum cannot double infinitely many times in finite time.

OpenAI claims (family 362, 2026-09-23, a 49-page paper + about 90,000 lines of Lean) [P]: in the **three-dimensional, single-species** case, for smooth initial data of any size and with no symmetry requirement, the solution **never breaks down** (it exists globally, is smooth, and is unique). This was an open problem for about 40 years.

After reading this chapter, you should be able to explain in your own words:

- why plasma has to be described by a "phase-space density" $f(t, x, v)$ rather than an ordinary density;
- why the whole problem reduces to "can a particle's momentum blow up in finite time," and why "the sum of the doubling times diverges" is enough;
- how well OpenAI's result has been verified, and what extra risks it carries compared with the quasi-Riemann hypothesis (the ζ function, §4);
- what it means for fusion engineering (almost nothing directly), and why.

12.1 Plasma: why we care

12.1.1 What a plasma is

Heat a gas to a high enough temperature and electrons get knocked out of the atoms. The gas turns into a soup of **positively charged ions + negatively charged electrons**. That is a plasma, often called the "fourth state" of matter.

Two examples you know:

- **Fusion reactors** (tokamaks, such as ITER): the core temperature is on the order of 10 keV. Converting to temperature: 1 eV corresponds to about 11,604.5 K, so $10 \text{ keV} \approx \mathbf{116 \text{ million K}}$ (computed by script).
- **The solar wind**: plasma the Sun continuously blows outward, with a typical speed of about 400 km/s ($\triangle$ this is the commonly quoted typical value; the real speed varies between

roughly 300 and 800 km/s). As a fraction of the speed of light: $400 / 299,792 \approx 0.0013$ (computed by script).

12.1.2 "Collisionless": why fluid equations don't work

In an ordinary gas, molecules collide countless times per second. Collisions quickly turn the local velocity distribution into a "thermal equilibrium" bell curve, so the gas can be described by just a few quantities: density, mean velocity, and temperature. (That is fluid dynamics, for example Navier–Stokes in §11.)

A hot, thin plasma is different: a particle can travel a very long way between two collisions, in many cases farther than the size of the device (Δ the actual numbers vary a lot with density and temperature; only the order-of-magnitude relationship is stated here). Collisions don't have time to "smooth out" the velocity distribution. At the same location there can be a group of fast particles and a group of slow ones at once, or even two streams of particles moving in opposite directions. **Recording only the mean velocity throws away key information.**

So we need a different description: record not only "how many particles are here," but also "at what velocity each of the particles here is moving." That is the phase space of the next section.

12.2 Phase space and the distribution function $f(t, x, v)$

12.2.1 Phase space: record position and velocity together

A particle's state is determined by two things: **where it is (x)** and **which way and how fast it is moving (v)**. Put the two together and you get a point (x, v). The space of all possible (x, v) is called **phase space**.

- On a one-dimensional line, x is one number and v is one number, so phase space is a **two-dimensional plane**.
- In real three-dimensional space, x has 3 components and v has 3 components, so phase space is **6-dimensional**. Add time, and $f(t, x, v)$ is a function of 7 variables.

The **distribution function** $f(t, x, v)$ means: at time t, roughly how many particles there are per unit volume of phase space near (x, v). Integrate f over all velocities and you get back the ordinary density:

$$\rho(t, x) = \int f(t, x, v) dv \quad (\text{total number of particles at position } x, \text{ regardless of velocity})$$

12.2.2 A toy example: why position alone is not enough

Four particles on a line, two scenarios:

	Particle positions	Particle velocities
Scenario A (two counter-streaming beams)	0, 0, 1, 1	+1, -1, +1, -1
Scenario B (at rest)	0, 0, 1, 1	0, 0, 0, 0

At $t = 0$, **looking only at positions**, the two scenarios are identical: 2 particles at position 0 and 2 at position 1.

But at $t = 2$ (each particle moves according to $x + v \cdot t$):

- Scenario A: the particles are at 2, -2, 3, -1. They have spread out.
- Scenario B: the particles are still at 0, 0, 1, 1.

So the ordinary density ρ cannot predict the future, but **the phase-space density f can**.

Physically, "counter-streaming beams" are exactly the typical configuration that triggers an instability in a plasma (the two-stream instability), and a fluid description can't see it.

12.2.3 Drawing a histogram in phase space (computed by script)

Here is a slightly larger example. 12 particles, positions x between 0 and 4, velocities ± 0.5 or ± 1 . Count them in bins of width 1 by position:

Position bin	$[-1, 0)$	$[0, 1)$	$[1, 2)$	$[2, 3)$	$[3, 4)$	$[4, 5)$
Particles per bin at $t = 0$	0	3	3	4	2	0
Particles per bin at $t = 1$ (each particle moves $v \times 1$)	1	0	5	3	2	1

(Computed by script; the particle data is in `精讲/scripts/09_vlasov_maxwell.py`.)

If you only have the 4 numbers in the $t = 0$ row, you cannot derive the $t = 1$ row. If you have "how many particles of each velocity are in each bin" (that is, a two-dimensional histogram in phase space), you can. **f is the limit of this two-dimensional histogram as the bins become infinitely small and the particles infinitely many.**

12.3 The Vlasov equation: particles follow trajectories, and f moves with them

12.3.1 First, no forces (free streaming)

With no forces at all, each particle moves in a straight line at constant speed: $x(t) = x_0 + v \cdot t$, and velocity doesn't change. A "clump of particles" in phase space just gets translated and sheared; particles don't appear or disappear. So

$$f(t, x, v) = f_0(x - v \cdot t, v)$$

This means: "the particle at (x, v) at time t is the particle that was at $(x - vt, v)$ at time 0."

Differentiate this expression with respect to t and you can verify that it satisfies

$$\partial f / \partial t + v \cdot \partial f / \partial x = 0$$

[Check 1] Let $f_0(x, v) = e^{-(x^2)} \cdot g(v)$, so under free streaming $f(t, x, v) = e^{-(x-vt)^2} \cdot g(v)$.

Compute $\partial f / \partial t$ and $\partial f / \partial x$ directly and verify that $\partial f / \partial t + v \cdot \partial f / \partial x = 0$.

12.3.2 Adding a force: the Vlasov equation

When a force F acts, velocity changes too: $dv/dt = F$ (taking mass to be 1 here). f is still "constant along particle trajectories"; the trajectories are just curved now. As an equation:

$$\partial f / \partial t + v \cdot \nabla_x f + F \cdot \nabla_v f = 0$$

- $\nabla_x f$ is the vector of partial derivatives of f with respect to the three components of position; $\nabla_v f$ is the same with respect to the three components of velocity.
- The three terms mean: the change of f over time + the change due to particles moving in position + the change due to the force altering particle velocities = 0.

This is the **Vlasov equation** (1938), also called the "collisionless Boltzmann equation": add a collision term on the right-hand side and you get the Boltzmann equation.

An important consequence: since f is constant along trajectories, the **maximum of f can never grow** ($0 \leq f \leq$ the maximum of f_0). This is Lemma 3.1 in the paper. It means the bad scenario "particles crowd together in phase space to infinite density" cannot happen. Any danger has to come from somewhere else.

12.3.3 Where the force comes from: the Lorentz force

The force on a charged particle is the **Lorentz force** (charge and mass both normalized to 1):

$$F = E + v \times B$$

- E is the electric field, which pushes the particle directly;

- $\mathbf{B}$ is the magnetic field, which makes the particle **turn** through the cross product $\mathbf{v} \times \mathbf{B}$. $\mathbf{v} \times \mathbf{B}$ is always perpendicular to $\mathbf{v}$, so the magnetic field **only changes direction and does no work**. It doesn't change the particle's speed.

[Check 2] Take $\mathbf{v} = (1, 0, 0)$, $\mathbf{B} = (0, 0, 1)$. Compute $\mathbf{v} \times \mathbf{B}$ and verify that its dot product with $\mathbf{v}$ is 0.

12.4 Maxwell's equations: how particles create fields in return

12.4.1 The four equations in plain language

Set the speed of light $c = 1$ and absorb the various constants (this is the normalization the paper uses). Maxwell's equations are:

$\partial \mathbf{E} / \partial t - \nabla \times \mathbf{B} = -\mathbf{j}$	(the current $\mathbf{j}$ and a changing magnetic field produce a change in the electric field)
$\partial \mathbf{B} / \partial t + \nabla \times \mathbf{E} = 0$	(a changing electric field produces a change in the magnetic field; Faraday's law)
$\nabla \cdot \mathbf{E} = \rho$	(charge is the "source" of the electric field; Gauss's law)
$\nabla \cdot \mathbf{B} = 0$	(there is no "magnetic charge")

In plain language:

- **$\nabla \cdot$ (divergence)**: how much a field "flows outward" from a point. $\nabla \cdot \mathbf{E} = \rho$ says "wherever there is charge, the electric field flows out of it."
- **$\nabla \times$ (curl)**: how much a field "swirls" around a point. The first two equations say "the electric and magnetic fields drive each other by swirling." This mutual driving propagates outward at the **speed of light**; it is the electromagnetic wave (light, radio).

12.4.2 Self-consistent coupling: a nonlinear loop

Connect the particles and the fields:

$\rho(t, \mathbf{x}) = \int f \, dv$	(charge density: add up the particles of all velocities)
$\mathbf{j}(t, \mathbf{x}) = \int (\text{particle velocity}) \cdot f \, dv$	(current: particles carry their charge as they move)

This forms a loop:

f determines ρ and $\mathbf{j} \rightarrow \rho$ and $\mathbf{j}$ determine $\mathbf{E}$ and $\mathbf{B}$ through Maxwell's equations $\rightarrow \mathbf{E}$ and $\mathbf{B}$ determine how the particles move through the Lorentz force $\rightarrow$ the particles' motion changes f again.

Taken alone, the Vlasov equation is linear in f (for given E and B), and Maxwell's equations are linear in E and B . **The nonlinearity comes entirely from this loop:** the force $E + v \times B$ itself depends on f . This is the same kind of structure as "the velocity field transports itself" in Navier–Stokes (§11).

12.4.3 Two "conserved quantities"

Two things in this system never change (Section 3 of the paper uses exactly these):

- **the number of particles** (f is constant along trajectories, and phase-space volume is preserved);
- **the total energy** = the kinetic energy of all particles + the electromagnetic field energy $\frac{1}{2} \int (|E|^2 + |B|^2) dx$.

So "the total energy blows up" is impossible. If anything breaks, it can only be because **energy concentrates onto a very few particles, or onto extremely small scales**. This point is the key to understanding everything that follows.

12.5 Why relativity

12.5.1 The problem with the non-relativistic version

If you use Newtonian mechanics (velocity = momentum / mass, which can be arbitrarily large), particles can move faster than light. But in Maxwell's equations the fields propagate at the speed of light. Once a particle outruns the field it is producing, the model contradicts itself. So **the physically consistent version must be relativistic**: particle speeds are always below the speed of light.

12.5.2 The relation between momentum and velocity

Take $c = 1$ and mass = 1, and write the momentum as p (a three-dimensional vector).

In relativity:

$$\text{velocity } \hat{v} = p / \sqrt{1 + |p|^2}$$

⚠ **Notation warning:** in OpenAI's paper and Lean code, the letter **v** denotes momentum, and the velocity is written $u(v) = v / \sqrt{1+|v|^2}$ (written $\hat{v}$ in the reasoning summary). So the v in $f(t, x, v)$ is actually momentum. Below I consistently use p for momentum and $\hat{v}$ for velocity, to avoid confusion.

(Computed by script)

Momentum p	Velocity $\hat{v}$	Gap to light speed $1 - \hat{v}$
0.1	0.0995	0.90
0.5	0.447	0.55
1	0.707	0.29
2	0.894	0.11
10	0.995	0.0050
100	0.99995	0.00005
1000	0.9999995	0.0000005

At small momentum $\hat{v} \approx p$ (Newtonian mechanics). At large momentum $\hat{v}$ gets arbitrarily close to 1 but never reaches it. **Momentum can be arbitrarily large; velocity always stays below the speed of light.** This is why the "blowup" below has to be discussed in terms of momentum rather than velocity: velocity cannot possibly blow up.

[Check 3] What is $\hat{v}$ when $p = 3$? (Hint: $\sqrt{10} \approx 3.162$)

12.5.3 Where real particles sit in this table (computed by script)

The electron's rest energy is 511 keV. Given a kinetic energy K , we have $\gamma = 1 + K/511$, $p = \sqrt{(\gamma^2 - 1)}$, $\hat{v} = p/\gamma$:

Electron kinetic energy	Momentum p (in units of $m \cdot c$)	Velocity $\hat{v}/c$
1 keV	0.063	0.062
10 keV (order of fusion core temperature)	0.20	0.19
100 keV	0.66	0.55
1 MeV	2.78	0.94
10 MeV	20.5	0.9988

Tokamaks have a dangerous phenomenon called **runaway electrons**: under a strong electric field, a small number of electrons are continuously accelerated to MeV or even tens of MeV ($\triangle$ the order of magnitude comes from fusion physics reviews, for example Breizman et al., *Nuclear Fusion* 59 (2019) 083001, not checked word for word). At that point relativistic effects cannot be ignored at all. Note: runaway electrons are a **physical phenomenon**, with very large but finite momentum. They are not the same thing as the "mathematically, momentum $\rightarrow \infty$ in finite time" discussed below.

12.5.4 The light cone: why Maxwell is harder than Poisson

If you ignore the magnetic field and assume the electric field responds "instantly" to the charge distribution ($\nabla \cdot \mathbf{E} = \rho$ and $\mathbf{E}$ is the gradient of some potential), you get the simpler

Vlasov–Poisson system. Its large-data problem in three dimensions was solved in 1991–92 (Pfaffelmoser 1992, Lions–Perthame 1991, Schaeffer 1991; the paper cites Pfaffelmoser, [20]).

Maxwell's fields don't respond instantly; they propagate at the speed of light. The field a particle feels at (t, x) comes from certain points in past spacetime: exactly those points that can send a signal to (t, x) at the speed of light. Together they form a **backward light cone**. The field produced this way is called the **retarded field**.

The difficulty: a source particle moving close to the speed of light is almost **running alongside the light it emits**. Its field piles up along the direction of motion (physicists call this "relativistic field compression"). When the source particle's velocity is almost parallel to the direction of the light ray, large factors appear in the retarded-field formula. The "loss" the OpenAI paper talks about (§12.8) happens right here.

12.6 What "the solution breaks down" actually means

12.6.1 An ODE toy: finite-time blowup

Look at the simplest equation $dp/dt = p^2$, $p(0) = 1$. The solution is $p(t) = 1/(1 - t)$: 2 at $t = 0.5$, 10 at $t = 0.9$, 100 at $t = 0.99$, and **it goes to infinity as $t \rightarrow 1$** . The equation itself is perfectly smooth, yet the solution breaks down at $t = 1$. This is called **finite-time blowup**.

Count it another way: p doubles from 1 to 2, from 2 to 4, from 4 to 8... Going from 2^k to 2^{k+1} ($k = 0, 1, 2, \dots$) takes time $1/2^k - 1/2^{k+1} = 1/2^{k+1}$.

Doubling number	Time taken by this doubling	Cumulative time
1st	1/2	0.5
2nd	1/4	0.75
3rd	1/8	0.875
60 doublings	...	$1 - 2^{-60}$ (computed by script ≈ 1.0)

The doubling times get shorter and shorter, and their sum converges (the geometric series sums to 1), so momentum can double infinitely many times in finite time. That is blowup.

12.6.2 The other way around: the sum of doubling times diverges $\Rightarrow$ no blowup

If you can prove "the k -th doubling takes at least c/k time," then the total time for N doublings is at least $c \cdot (1 + 1/2 + \dots + 1/N)$. This sum is called the **harmonic series**, and it **diverges**:

N	$1 + 1/2 + \dots + 1/N$ (computed by script)
10	2.93
60	4.68
100	5.19
1,000	7.49
1,000,000	14.39


It grows very slowly (roughly like $\ln N$), but without bound. So within any finite time T , momentum can only double finitely many times, and momentum stays bounded.

This contrast is the skeleton of the entire proof:

Geometric series $1/2 + 1/4 + \dots$ (converges) $\rightarrow$ can blow up; harmonic series $1 + 1/2 + 1/3 + \dots$ (diverges) $\rightarrow$ cannot blow up.

The staircase figure below draws the partial sums of the harmonic series: the horizontal axis is the number of doublings so far, k , and each step has height $1/k$. The first 60 steps add up to about 4.68 (matching the $N = 60$ row in the table above); the steps keep getting shorter, but the total height has no ceiling. The text on the right of the figure previews the chain of reasoning that §12.7.1 and 8.3 explain: each doubling takes at least $\sim 1/(k \cdot \log 2)$ time $\rightarrow$ only finitely many doublings fit in finite time $\rightarrow$ momentum stays bounded $\rightarrow$ by Glassey–Strauss 1986, the solution does not break down. The truly hard step is proving "each doubling takes at least $\sim 1/\log P$ time", which needs the signed momentum-increment estimates and the counting of direction changes in §12.8.4.

Vlasov–Maxwell endgame: each momentum doubling takes $\geq 1/k$ time, and $\sum 1/k$ diverges



k = number of momentum doublings so far
doubling k ($P \approx 2^k$) takes $\approx 1/(k \cdot \log 2)$
total time $\approx$ harmonic series $\rightarrow \infty$
 $\Rightarrow$ finitely many doublings in finite time
 $\Rightarrow$ momentum stays bounded, so by
Glassey–Strauss 1986: no blowup

A hand-computable version of the argument skeleton in the reasoning summary of the OpenAI paper (family 362). The real difficulty is proving that each doubling takes at least $\sim 1/\log P$ time (signed momentum-increment estimates plus counting direction changes). Staircase computed by the build script.

[Check 4] If all you can prove is "the k -th doubling takes at least $1/k^2$ time," can you conclude there is no blowup? Why or why not?

(By the way: in the quasi-Riemann hypothesis (§4), $\zeta(1) = 1 + 1/2 + 1/3 + \dots$ diverges; it is the same series.)

12.6.3 "Global regularity" for PDEs

For a PDE, "the solution breaks down" means some derivative or some quantity becomes infinite at a finite time, after which the equation can no longer hold pointwise. The opposite properties are:

- **Local existence:** starting from smooth initial data, a smooth solution exists at least for a short time interval $[0, T)$. For Vlasov–Maxwell this was proved by Wollman 1984 (paper [29]).
- **Global regularity / global classical solution:** this smooth solution can be continued for all time $t \geq 0$. This is what OpenAI claims.

There is another concept that is easy to confuse with these: **weak solutions**. DiPerna–Lions 1989 (paper [5]) proved that "weak solutions" exist globally for large data. A weak solution is allowed to be non-smooth, satisfies the equation only "on average," and is not known to be unique. So the existence of weak solutions **does not mean** the problem is solved. This is exactly the same situation as Leray's 1934 weak solutions for Navier–Stokes.

12.7 40 years of history: how the problem was reduced, and where it got stuck

12.7.1 Glassey–Strauss 1986: you only need to control the fastest particle

Glassey & Strauss, "Singularity formation in a collisionless plasma could occur only at high velocities", ARMA 92 (1986) 59–90. The title is the conclusion:

If on $[0, T)$ the momenta of all particles stay below some finite number, then the solution can be continued smoothly past T .

Take the contrapositive: if the solution breaks down at a finite time T , then **some particle's momentum must have shot toward infinity before T** . The whole problem is reduced to a question about single particles:

Can a particle be accelerated to infinite momentum in finite time by the self-consistent electromagnetic field?

This criterion was later re-proved and strengthened many times: Klainerman–Staffilani 2002 with Fourier methods, Bouchut–Golse–Pallard 2003 with physical-space methods; Luk–Strain 2014 proved that it is enough to control the projection of momentum onto some fixed

plane; Pallard, Sospedra-Alfonso-Illner and others gave various moment conditions (§1.1 of the paper lists the sources). But all of these are **conditional**: they tell you "it's enough for momentum to be bounded," not "momentum really is bounded."

12.7.2 Lower dimensions: Glassey–Schaeffer

If you make the problem "a few dimensions smaller," it can be done:

- **1½ dimensions** (1D position, 2D momentum): Glassey–Schaeffer 1990;
- **2½ dimensions**: Glassey–Schaeffer 1997;
- **2 dimensions** (2D position): Glassey–Schaeffer 1998, two papers (ARMA 141).

(All cited from the paper's references [6, 7, 9, 10].) In lower dimensions the singularity of the field is weaker, and energy conservation is enough to control acceleration. In three dimensions this "just barely enough" margin disappears.

12.7.3 Small data, symmetry

- **Small data / dilute plasma**: Glassey–Strauss 1987 (paper [12]), followed by Schaeffer 2004, Bigorgne 2020/2025, Wang 2022, Wei–Yang 2021, and others. The idea: if the particles are sparse enough, the field is weak and the acceleration is limited.
- **Symmetry**: Xuecheng Wang's large-data result under cylindrical symmetry (arXiv:2203.01199, v1 in 2022, v2 on 2026-07-16; the paper says "to appear in Annals of Mathematics"; a second paper is arXiv:2607.14685, 2026-07-16). This was the closest large-data result before OpenAI, but it requires the initial data to be rotationally symmetric about an axis.

Up to before 2026-09 (going by the survey in §1.1 of the paper), the three-dimensional large-data problem **with no size or symmetry restrictions** remained open.

12.7.4 Where the difficulty lies (intuitive version)

Conservation of total energy gives a "budget." To prove that a particle isn't accelerated to infinity, you need to estimate the total impulse it receives over a period of time. The most direct approach is to add up the **absolute values** of the force. But in three dimensions this estimate overshoots what you need by a factor of $\sqrt{w}$ (w is the energy of the receiving particle). It falls just short and can't be closed. The loss happens when the source particle's velocity is almost parallel to the direction of the light ray (§12.5.4).

So you have to exploit **cancellation**: the force is sometimes positive and sometimes negative, and its direction changes; if you integrate first and then take the absolute value, the result can be much smaller. The hard part is capturing this cancellation in a rigorous proof.

12.8 What OpenAI claims

12.8.1 The main result

The abstract of the paper *Global classical solutions of the three-dimensional relativistic Vlasov–Maxwell system* (OpenAI, 2026-09-23, 49 pages, family 362) [P]:

"We prove global existence and uniqueness for arbitrary smooth admissible initial data in the three-dimensional, one-species relativistic Vlasov–Maxwell system. The particle density is initially compactly supported, and the electromagnetic fields have finite energy and bounded derivatives of all orders. The solution remains smooth on every finite time interval. This resolves the large-data global classical regularity problem for this model, without size or symmetry restrictions on the data."

Breaking it down:

- **Dimension:** three-dimensional space + three-dimensional momentum ("3+3 dimensions").
- **Direction:** this is **global regularity** (no blowup), not blowup. That is the opposite direction from Navier–Stokes in §11.
- **One species:** there is only one kind of charged particle, with charge and mass both normalized to 1, and no background charge.
- **Initial data conditions** ("admissible"): f_0 is smooth, non-negative, and **compactly supported** (nonzero only in a bounded region of phase space, so the initial momentum is bounded); E_0, B_0 are smooth, with all derivatives bounded and finite energy (they belong to L^2); and they satisfy the two Gauss constraints $\nabla \cdot E_0 = \rho_0, \nabla \cdot B_0 = 0$.
- **Conclusion:** there exists a unique global classical solution; it is smooth on every finite time interval $[0, T]$, and the particle support is compact (momentum is bounded, but **the bound may grow with T**).

The paper also points out specifically (after Theorem 1.1) that the total charge is allowed to be nonzero (so the electric field has a Coulomb tail $\sim 1/r^2$ far away), and it does not require all derivatives of the fields to be square-integrable.

[Check 5] The Coulomb field of total charge Q has magnitude about $Q/(4\pi r^2)$. Verify that outside radius R , the electric field energy $\frac{1}{2} \int |E|^2 dx$ is finite (you only need to check whether $\int_{R^\infty} (1/r^2)^2 \cdot r^2 dr$ converges). This shows that "allowing nonzero total charge" and "finite energy" don't contradict each other.

12.8.2 The skeleton of the proof: two propositions

Section 2 of the paper splits the whole proof into two pieces:

Proposition 2.2 (local theory + continuation criterion): the solution exists locally and is unique; if the maximal existence time $T_{\max}$ is finite, then the maximal momentum must be unbounded. This is the Glassey–Strauss criterion (in the formulation of Luk–Strain 2014), plus a technical fix: the Luk–Strain version requires the fields to be in H^5 (a Sobolev space, roughly "5 derivatives square-integrable"), while the initial data here only require bounded derivatives and finite energy, so the two don't quite match. The paper's approach is to keep the Coulomb part, cut off the rest of the field far away, and then add back the difference using "Maxwell solutions in vacuum," relying on "finite signal propagation speed" to guarantee that this doesn't affect the particles within the given time (paper §10). In the reasoning summary, the model found this gap itself and went back to check it several times.

Proposition 2.1 (signed momentum increment estimate): for a receiving particle with energy about w , over a time interval of length I ,

$$|\text{change in momentum}| \leq M \cdot P \cdot \sqrt{I} + A \cdot I \cdot P^2 \cdot \log(2+P) / w$$

where P is an upper bound on the current maximal momentum, and M, A are constants. The key word is "**signed**": first integrate the force over time (giving an impulse vector), then take its length, instead of first taking the magnitude of the force and then integrating.

12.8.3 From Proposition 2.1 to "no blowup": it's the staircase from §12.6

The paper's derivation (paper §2, "Proof of Theorem 1.1") goes like this: suppose $T_{\max}$ is finite; then the maximal momentum is unbounded. Let t_n be the first time the maximal energy reaches 2^n . During the n -th doubling, Proposition 2.1 gives

$$t_n - t_{n-1} \geq c / \log(2 + 1024 \cdot 2^n)$$

The denominator is about $n \cdot \log 2$, so the n -th doubling takes at least $\sim c/n$ time: **exactly the harmonic series** (the staircase figure in §12.6.2). Summing the paper's exact form (with $c = 1$) (computed by script):

Number of doublings N	$\sum 1/\log(2 + 1024 \cdot 2^n), n = 1 \dots N$
10	0.96
100	3.40
1,000	6.59
10,000	9.90
100,000	13.22

The sum grows without bound, which contradicts "doubling infinitely many times before a finite $T_{\max}$." So $T_{\max} = \infty$.

This step is the simplest in the whole paper. About 40 of the 49 pages go into proving Proposition 2.1.

12.8.4 How Proposition 2.1 is proved (I only read the introduction and §12.2; what follows is a simplified version)

1. **Retarded fields + light-cone binning.** Follow the trajectories of the source particles to find where they intersect the backward light cone of the receiving particle, and sort the contributions into "bins" by momentum size, distance, and three angles (between the source velocity, the receiving velocity, and the direction of the light ray). Using the flux of energy through the light cone, plus estimates of "how long a source particle can stay in each bin" (occupation, i.e. occupation time), you get an upper bound on the **absolute value** of the force (paper §3–5).
2. **Find the loss.** When the receiving energy is $\sim w$, the absolute-value estimate overshoots what is needed by a factor of $\sqrt{w}$. The loss is concentrated in a narrow strip of angles where "the source particle's velocity is closer to the light-ray direction than the receiving particle's is."
3. **Signed cancellation** (paper §6). In this narrow strip, first integrate along the source particle's trajectory, then take the absolute value. An exact identity combines "the part without the source particle's acceleration" and "the part with the source particle's acceleration": the former exactly cancels a term produced by differentiating the light-cone geometry. The force is thereby written as "a total derivative + some less singular terms." One of these terms contains **the receiving particle's own acceleration**, which in turn equals the force on the receiving particle, so it becomes "error $\propto$ the force on the receiving particle," and the coefficient in front of it has to be controlled.
4. **Count direction changes** (paper §7). Under the assumption that "all signed increments satisfy Proposition 2.1," the same identity is used to prove that the number of times a particle changes direction (turns through a given angle) is finite and countable. Between two turns, the relative motion is monotone, so the time a particle spends crossing a given angular region is compressed further, giving a better occupation estimate.
5. **Pick only the bins you need** (paper §8). The signed argument is used only on bins where the direct estimate exceeds a certain summable threshold. On those bins, after summing, the coefficient of the receiving particle's force is at most a constant $\times w^{(-1/2)}$, which **exactly cancels** the $\sqrt{w}$ loss.
6. **Bootstrap** (paper §9). All of the above is done under the premise "assume Proposition 2.1 holds on $[0, \tau]$." What comes out at the end is a **strictly better** estimate, and a continuity

argument then pushes the assumption to the whole interval. Figure 1 in the paper draws this dependency graph.

An analogy in distributed-systems language: a bootstrap is a lot like "proving a protocol safe with an inductive invariant." Assume the invariant holds up to time τ , prove that it still holds a little beyond τ , and with margin to spare (a strict improvement), so it holds forever.

12.8.5 What the reasoning summary says

OpenAI published a 6-page "reasoning summary" ([reasoning_traces/relativistic-vlasov-maxwell.pdf](#), released 2026-10-06) [P]. It is candid. Some points worth noting:

- The model first tried a Pfaffelmoser-style "dwell time" strategy (exactly the route that solved Vlasov–Poisson in 1992), and also considered generalizing Wang's cylindrical-symmetry method to the general case.
- It hit walls repeatedly along the way and left concrete numbers: "the improved occupancy bound forced $H < 0.41z$, while the excessive coefficient required $H > 0.496z$ ". This rules out the "excessive coefficient" case through a contradiction between exponents: the two inequalities cannot both hold.
- The model found that the time window it was using to count direction changes was too long, and shortened it by a factor of $\sim M^2 + A \cdot \log P$.
- The summary's last sentence is quite candid; in effect it flags that the conclusion rests on a few new estimates, which the cited continuation theorem does not supply by itself: "That assertion rests on the occupation bound, direction-count estimate, and cutoff cancellations needed to close the signed momentum bootstrap; the cited continuation theorem alone does not establish those estimates."

12.9 How well has it been verified

12.9.1 The Lean statement: much longer than the ζ statement in §4, and the definitions are homemade

The Lean statement for the quasi-Riemann hypothesis (§4) is one line, and the

`riemannZeta` it uses is Mathlib's standard definition. **Here it's different:** Mathlib has no Vlasov–Maxwell equations, so the challenge file

`lean/ComparatorChallenges/VlasovMaxwell.lean` (about 140 lines) **defines everything from scratch:** momentum, velocity, cross product, curl, divergence, charge density, current, the equations themselves, and what a "classical solution" is.

The main theorem:

```
theorem global_classical_solution (d : Datum) (hd : Admissible d) :  
  ∃ s : Solution, Classical d s ∧ SmoothOnFiniteHorizons s ∧
```

$\forall s' : \text{Solution}, \text{Classical d s}' \rightarrow \text{SameNonnegativeTime s s}'$

Word-by-word translation:

- **d : Datum**: the initial data, a triple (f_0, E_0, B_0) .
- **hd : Admissible d**: the initial data satisfy **Admissible**, which unfolds to: f_0 is smooth (**ContDiff $\mathbb{R} \infty$**), compactly supported (**HasCompactSupport**), and non-negative; E_0, B_0 are smooth with all derivatives uniformly bounded (**BoundedSmooth**); E_0, B_0 are square-integrable (**MemLp ... 2**); $\nabla \cdot E_0 = \rho(f_0)$, $\nabla \cdot B_0 = 0$.
- **$\exists s : \text{Solution}$** : there exists a solution (f, E, B) .
- **Classical d s**: it is a classical solution: f, E, B are C^1 on $t \geq 0$; $f \geq 0$; E, B are continuous in time as L^2 -valued functions; on every finite time interval the support of f lies in a compact set; the equations hold pointwise; and at $t = 0$ it equals the initial data.
- **SmoothOnFiniteHorizons s**: on every $[0, T]$, f, E, B are infinitely differentiable.
- **$\forall s', \text{Classical d s}' \rightarrow \text{SameNonnegativeTime s s}'$** : any other classical solution is identical to it on $t \geq 0$ (uniqueness).

This matches Theorem 1.1 of the paper, and the problem statement OpenAI gave the model (the prompt at the start of the reasoning summary).

I checked by hand the places in the definitions most likely to be written wrong [P]:

- **velocity $v = (q \ v)^{-1} \cdot v$, $q \ v = \sqrt{1 + \|v\|^2}$** : exactly $\hat{v} = p/\sqrt{(1+|p|^2)}$. ✓
- **cross**: $(a_1b_2 - a_2b_1, a_2b_0 - a_0b_2, a_0b_1 - a_1b_0)$, the standard cross product. ✓
- **curl**: first component $\partial_1 F_2 - \partial_2 F_1$, and so on; the standard curl. ✓
- **Equations**: $\partial_t f + \hat{v} \cdot \nabla_x f + (E + \hat{v} \times B) \cdot \nabla_v f = 0$; $\partial_t E - \nabla \times B = -j$; $\partial_t B + \nabla \times E = 0$; $\nabla \cdot E = \rho$; $\nabla \cdot B = 0$.

Matches the paper's (1.1) term by term, with the right signs. ✓

A trap specific to Lean: in Lean/Mathlib, the derivative of a non-differentiable function returns 0, and the integral of a non-integrable function also returns 0. If the statement didn't require smoothness, "the equations hold" could be satisfied by exploiting these default values. Here **Classical** requires C^1 and compact support, which closes that loophole (my judgment; no formal audit done $\triangle$).

12.9.2 The proof itself (local grep, snapshot fd4aeeb, 2026-10-07) [P]

- Solution module **OAI/Analysis/VlasovMaxwell**: **207 .lean files, 89,993 lines** (computed by script). The largest subdirectories are **Population** (56 files, about 32,000 lines, corresponding to the paper's occupation and bin estimates) and **Retarded** (53 files, about 24,000 lines, corresponding to the retarded fields).

- grep finds no `sorry`, `admit`, `native_decide`, or custom `axiom`.
- **It imports only Mathlib** and doesn't depend on any patched external library. This is cleaner than the quasi-Riemann hypothesis in §4 (which has a dependency patch of about 40,000 lines).
- In other words, the Glassey–Strauss continuation criterion **was not assumed as an axiom**; it was re-proved inside these 90,000 lines (`Continuation/GlobalContinuation.lean` contains theorems such as `bounded_momentum_uniform_restart`; I only looked at the theorem names and did not read the proofs $\triangle$).
- Comparator config `VlasovMaxwell.json`: only the three standard axioms are allowed (`propext`, `Quot.sound`, `Classical.choice`), and `enable_nanoda: false`, meaning **no re-check with a second, independent kernel**.

12.9.3 Independent checks and human experts

- **I found no public report of anyone successfully running Comparator on 362** (the Goldblatt re-check mentioned in §4 was for ζ). $\triangle$
- **As of 2026-10-08, I found no named public comment from any kinetic-theory / plasma PDE expert** (searched general media and the GitHub ecosystem around the release). $\triangle$ Media coverage just restates the abstract.
- A coincidence of timing worth noting: Wang's large-data cylindrical-symmetry paper had just been updated and accepted by Annals in 2026-07, and OpenAI's symmetry-free version appeared two months later. Experts in the field (Wang, Luk, Strain, Pallard, Bigorgne, and others) are the people best qualified to read these 49 pages.

12.9.4 Remaining risks

1. **Faithfulness of the statement:** 140 lines of homemade definitions; any subtle mistake could make the theorem "prove something else." I checked the equations and the key definitions and found no problems, but this is not a formal audit.
2. **Scope:** single species only. Real plasmas have at least two species, electrons and ions, with opposite charges, so particles both repel and attract each other. Whether the proof generalizes directly is **unknown** $\triangle$. A caution: in the related relativistic Vlasov–Poisson system, the "repulsive" (plasma) and "attractive" (gravitational) cases can have completely different conclusions (Glassey–Schaeffer 1985, paper [8], proved global existence for the spherically symmetric repulsive case; the attractive case is known to be able to blow up $\triangle$ the latter point is from my memory, not rechecked against the original).
3. **No human has understood it:** the 49-page bootstrap and bin decomposition have so far only been checked by a machine.

12.10 What it would mean if true (and what it would not mean)

12.10.1 For mathematics

1. **A roughly 40-year-old flagship open problem (single-species case) is solved:** from the Glassey–Strauss 1986 reduction to "momentum really is bounded," and with machine verification. Once confirmed by experts in the field, it may well outweigh most results in the batch (a judgment).
2. **Methodological novelty:** signed impulse estimates + counting direction changes + using cancellation only in the bins that need it. Whether this transfers to other kinetic systems (multi-species Vlasov–Maxwell, Einstein–Vlasov, etc.) is open (speculation).
3. **The conclusion is finer than "no blowup":** momentum is bounded on every finite time interval, but the bound may grow with time. It does **not** say what happens to particle momenta over long times (for example, whether they approach some stable state). Long-time behavior is a separate question.

12.10.2 For fusion and space physics: almost no direct impact

This deserves to be said clearly, because the media tend to write it up as "AI solves a hard math problem in fusion."

- **Engineering simulations don't depend on this theorem.** The fusion community uses particle-in-cell (PIC) methods, gyrokinetic codes, and fluid/MHD codes. They compute on discrete grids with finite resolution and finite time steps, and never need to assume that the continuum equations are globally smooth.
- **The model itself is far from a real device:** a real tokamak has at least two species of particles, collisions (important on long time scales), metal walls and boundary conditions, and strong magnetic fields from external coils. This theorem is proved for **all of unbounded three-dimensional space, with no external field and only one species of particle**.
- **Runaway electrons are not mathematical blowup:** the runaway electrons engineers worry about have very large but finite energy. The theorem says "nothing shoots to infinity in finite time"; it can't tell you how many MeV runaway electrons can reach at most in a real device.

A more accurate positioning: it is a **certificate of health for the model**. It tells physicists that this system of equations is self-consistent as a mathematical object: starting from reasonable initial data, it will not produce infinite momentum on its own in any finite time. (Note that what is controlled is each particle's **momentum**, not the total energy: the **total energy** is conserved anyway.) This has value for theorists (for example, it guarantees that the premises of certain derivations hold), but it won't change how anyone designs a fusion reactor today.

12.10.3 For AI

- This is **one of the weightiest mathematical-physics results in this batch**: a named open problem with a complete chain of literature, a clear conclusion, and **a Lean proof that depends only on Mathlib**.
- Compared with §4: the quasi-Riemann statement used community definitions (the risk was in the dependency patch); here the risk is in the **homemade statement**. For machine-proved PDE results, "is the statement faithful" is the new human bottleneck: reading 140 lines of Lean definitions is much easier than reading a 49-page paper, but it still needs someone who knows the field.
- The reasoning summary shows a process that looks fairly like a human researcher's: try known methods first, hit a wall, find a concrete contradiction between exponents, fix the time window, and repeatedly audit the order of dependencies. But what we see is a summary edited by OpenAI, not the raw reasoning.

12.11 Self-test

Try to answer these after reading. If you can't, go back and reread the corresponding section.

1. Why can't a plasma be described by the position density $\rho(t, x)$ alone? Give an example. (§12.2.2)
2. The Vlasov equation says "f is constant along particle trajectories." What property of the maximum of f follows from this? (§12.3.2)
3. Why, in the relativistic version, do we look at momentum rather than velocity when discussing "blowup"? (§12.5.2)
4. What did Glassey–Strauss 1986 reduce the whole problem to? (§12.7.1)
5. Why does "the k-th doubling takes at least c/k time" imply no blowup? What if it were $c/2^k$? (§12.6.1–12.6.2)
6. Compared with the Lean statement for ζ in §4, where is the main extra risk in OpenAI's Lean statement? (§12.9.1, 9.4)
7. Why is this result said to have almost no direct impact on fusion engineering? Give at least two reasons. (§12.10.2)

Appendix: answers to the checks

- **Check 1:** Let $s = x - vt$. $\partial f / \partial t = e^{(-s^2)} \cdot (-2s) \cdot (-v) \cdot g(v) = 2vs \cdot e^{(-s^2)} g(v)$; $\partial f / \partial x = e^{(-s^2)} \cdot (-2s) \cdot g(v) = -2s \cdot e^{(-s^2)} g(v)$. So $\partial f / \partial t + v \cdot \partial f / \partial x = 2vs \cdot e^{(-s^2)} g - 2vs \cdot e^{(-s^2)} g = 0$.
- **Check 2:** $v \times B = (0 \cdot 1 - 0 \cdot 0, 0 \cdot 0 - 1 \cdot 1, 1 \cdot 0 - 0 \cdot 0) = (0, -1, 0)$. Its dot product with $v = (1, 0, 0)$ is 0. The particle is pushed in the $-y$ direction, but its speed doesn't change; it only turns.
- **Check 3:** $\hat{v} = 3/\sqrt{10} \approx 3/3.162 \approx \mathbf{0.949}$ (script: 0.9487).

- **Check 4:** No. $1/1^2 + 1/2^2 + 1/3^2 + \dots$ converges (to $\pi^2/6 \approx 1.645$, the $\zeta(2)$ of §4). The total time is finite, so momentum can double infinitely many times in finite time, and blowup can't be ruled out. You need a lower bound like $1/k$ whose "sum diverges."
- **Check 5:** $\int_{-R}^{\infty} (1/r^2)^2 \cdot r^2 dr = \int_{-R}^{\infty} r^{-2} dr = 1/R$, which is finite. With the constants included, $\frac{1}{2} \int_{\{|x|>R\}} |E|^2 dx = Q^2/(8\pi R)$; for $Q = 1$, the script computes $\int_{-R}^{\infty} (1/(4\pi r^2))^2 \cdot 4\pi r^2 dr = 1/(4\pi R)$, which is ≈ 0.0796 at $R = 1$ and ≈ 0.0080 at $R = 10$. In three dimensions a $1/r^2$ field is square-integrable (far away); the part near the origin doesn't diverge either, because the charge is smoothly distributed.

Appendix: Glossary

Term	One-line explanation
Plasma	An ionized gas, made of ions and electrons
Collisionless	Particles travel far between collisions, so collisions don't have time to smooth out the velocity distribution
Phase space	The space of position + velocity (or momentum); 6-dimensional in 3D
Distribution function $f(t, x, v)$	Number of particles per unit volume of phase space; in the paper v means momentum
Vlasov equation	f is constant along particle trajectories; the "collisionless Boltzmann equation"
Lorentz force	$E + \hat{v} \times B$; the magnetic field only changes direction and does no work
Maxwell's equations	How charges and currents produce electric and magnetic fields; fields propagate at the speed of light
Self-consistent coupling	The loop in which particles produce fields and fields push particles; the source of nonlinearity
Backward light cone / retarded field	The past spacetime points that can send a signal to the current point at the speed of light; the field coming from there
Vlasov–Poisson	Simplified version that ignores the magnetic field, with an instantly responding electric field; solved in 3D in 1991–92
Finite-time blowup	Some quantity of the solution goes to infinity at a finite time
Global classical solution / global regularity	A smooth solution exists for all $t \geq 0$
Weak solution	A solution that satisfies the equation only in an averaged sense, possibly non-smooth and non-unique
Continuation criterion	A theorem of the type "as long as some quantity is bounded, the solution can be continued"
Compact support	The function is nonzero only in a bounded region
Bootstrap	Assume an estimate holds, derive a stronger estimate, then use continuity to remove the assumption
Harmonic series	$1 + 1/2 + 1/3 + \dots$, which diverges

Appendix: Further reading

- **The original paper:** <https://github.com/openai/math/blob/main/preprints/Global-classical-solutions-of-the-three-dimensional-relativistic-Vlasov-Maxwell-system-September-23-2026/paper.pdf> (I suggest reading §1–2 first, 7 pages in total; the "Proof of Theorem 1.1" in Section 2 is the full version of §12.8.3 of this chapter)
- **The reasoning summary:**
https://github.com/openai/math/blob/main/reasoning_traces/relativistic-vlasov-maxwell.pdf
(6 pages)
- **The Lean statement:**
<https://github.com/openai/math/blob/main/lean/ComparatorChallenges/VlasovMaxwell.lean> ;
scope notes: <https://github.com/openai/math/blob/main/lean/docs/362.md>
- **Glasse & Strauss 1986**, ARMA 92:59–90, <https://doi.org/10.1007/BF00250732> (the reduction of the problem)
- **Luk & Strain 2014**, *A new continuation criterion for the relativistic Vlasov–Maxwell system*, CMP 331:1005–1027, arXiv:1406.0165 (a modern, more readable proof of the continuation criterion; the OpenAI paper uses its formulation)
- **Wang**, cylindrically symmetric large data: arXiv:2203.01199, arXiv:2607.14685 (the closest result before OpenAI)
- **Robert Glassey**, *The Cauchy Problem in Kinetic Theory* (SIAM, 1996): the standard textbook of the field; the first few chapters cover the basic theory of Vlasov–Poisson and Vlasov–Maxwell.
- **Introductory physics:** Francis F. Chen, *Introduction to Plasma Physics and Controlled Fusion* (its chapter on "kinetic theory" covers the Vlasov equation, $\triangle$ chapter number not checked); good for filling in plasma background.

13. The Mézard–Parisi formula (diluted spin glasses): from magnets to the energy landscapes of neural networks

Prerequisites: you can work with exponentials and logarithms (e^x , $\ln x$), you know what an "expectation" (average) is, and you know the summation sign Σ . The physics concepts (energy, temperature, free energy) are all explained from scratch. The Hopfield network you met in neuroscience comes back in §13.10. Every number in the text was computed by script or has a source; [P] marks primary material read (paper, repo, Lean source), and $\triangle$ marks something unverified or conflicting. **[Check]** items are small exercises you can work

out yourself; answers are at the end of the chapter. Estimated reading time: 70–100 minutes. You can split it into three sittings: §13.1–13.4 (spins, free energy, the replica method), §13.5–13.7 (sparse graphs, the cavity method, the problem itself), §13.8–13.10 (OpenAI's result, how well it is verified, and what it means).

13.0 The bottom line: the one thing this chapter explains

A magnet contains a huge number of tiny compass needles ("spins"), each of which can only point up or down. If the interactions between the needles are **random**, with some pairs wanting to align and others wanting to point opposite ways, you get a situation where "no arrangement can satisfy everyone." Such a system is called a **spin glass**. The central quantity physicists care about is the **free energy**: once you know it, you can derive almost every macroscopic property of the system at every temperature.

Between 1979 and 2001, physicists (mainly Parisi and Mézard) wrote down formulas for the free energy using a set of **mathematically non-rigorous** methods (the replica method, the cavity method). Mathematicians spent decades turning parts of this into theorems. The chain of reasoning in this chapter is:

Free energy = the minimum of a "hierarchical trial functional" $\Leftarrow$ the upper bound was proved long ago (the interpolation method, 2003–04) $\Leftarrow$ the lower bound (equality) requires controlling "the spin patterns seen by several replicas at once" $\Leftarrow$ in the dense model (SK), pairwise overlaps are enough; in sparse models they are not, and this step was stuck for over twenty years.

OpenAI claims (family 221, 2026-09-23, a 36-page paper + about 56,000 lines of Lean) [P] to have filled in the lower bound for **one class of sparse models** (Poisson-diluted, even-arity interactions, satisfying the Panchenko–Talagrand condition), thereby proving that equality holds, and has given a Lean formalization.

Note: the name "Mézard–Parisi" refers to at least two different things in the literature. What OpenAI proved is the **hierarchical cavity formula for diluted spin glasses** (Mézard–Parisi 2001), **not** the famous **random assignment problem** $\zeta(2) = \pi^2/6$ (Mézard–Parisi 1985–87, proved by Aldous in 2001). §13.6 covers the latter, because it is the best example for understanding "why physicists' non-rigorous methods are trustworthy."

After reading this chapter, you should be able to explain in your own words:

- what free energy is, why "averaging $\log Z$ " is hard, and what trick the replica method uses to get around it;
- the difference between dense (SK) and sparse (Viana–Bray / Bethe lattice) models, and why the cavity method is exact on trees;

- what exactly the theorem OpenAI proved covers and does not cover (3-SAT, the fixed-degree Bethe lattice, and Hopfield are all outside it);
- which of its connections to ML / neuroscience are real and which are only analogies.

13.1 Magnets and spins: the Ising model

13.1.1 Spins

Think of a magnet as a row of tiny compass needles, each with only two states: up is written $\sigma = +1$, down is written $\sigma = -1$. This variable that can only take ± 1 is called a **spin**. One assignment of values to N spins (for example $+1, -1, -1, +1$) is called a **configuration**, and there are 2^N configurations in total.

You have seen the same thing in neuroscience: the binary "firing / not firing" model of a neuron, which is exactly how a Hopfield network encodes it.

13.1.2 Energy

Between two spins i, j there is a **coupling strength** J_{ij} . The energy of the system (physicists call it the Hamiltonian) is:

$$H(\sigma) = - \sum_{\text{connected } i,j} J_{ij} \cdot \sigma_i \cdot \sigma_j \quad (\text{sometimes with an extra term } - \sum h_i \sigma_i, \text{ where } h \text{ is an external magnetic field})$$

Look at a single pair of spins:

- $J_{ij} > 0$ (**ferromagnetic**): when $\sigma_i = \sigma_j$, the product $\sigma_i \sigma_j = +1$ and the energy contribution is $-J$, which is low; so they "want to align."
- $J_{ij} < 0$ (**antiferromagnetic**): they "want to point opposite ways."

Physical systems tend toward low-energy states. The configuration with the lowest energy is called the **ground state**.

This energy function connects directly to your neuroscience background: the energy of a Hopfield network (1982), $E = -\frac{1}{2} \sum W_{ij} \sigma_i \sigma_j$, has exactly this form, i.e. it is a spin-glass-style Hamiltonian. Stored memories are local minima of this energy; starting from a partial cue, the network updates so that the energy never goes up, and it settles into a nearby minimum (§13.10.3 comes back to this in detail).

13.1.3 The Boltzmann distribution: where temperature comes in

At temperature $T > 0$, the system doesn't stick to the ground state. Instead, it appears in every configuration with some probability, given by the **Boltzmann distribution**:

$$P(\sigma) = e^{(-\beta \cdot H(\sigma))} / Z, \quad Z = \sum_{\text{all configurations}} e^{(-\beta \cdot H(\sigma))}, \quad \beta = 1/T$$

β is called the **inverse temperature**. Z is the normalizing constant obtained by adding up all the weights, called the **partition function**. You see this formula in ML every day: it is just **softmax**, with the "logit" replaced by $-\beta \cdot \text{energy}$. β is the softmax's inverse-temperature parameter.

Smallest example: two spins, $J = 1$ (ferromagnetic), $H = -\sigma_1 \sigma_2$. Four configurations: (+,+) and (-,-) have energy -1; (+,-) and (-,+) have energy +1. Computed by script:

β	Z	P(a given aligned configuration)	P(a given anti-aligned configuration)	P(aligned, total)
0 (infinite temperature)	4.000	0.250	0.250	0.500
0.5	4.511	0.366	0.135	0.731
1	6.172	0.440	0.060	0.881
2	15.05	0.491	0.009	0.982
5	296.8	0.500	0.000	1.000

At high temperature all four configurations are equally likely (thermal motion smooths everything out); at low temperature essentially only the ground states remain. Temperature is the knob that sets the weighting between "energy" and "disorder."

[Check 1] By hand, compute $P(+,+) = e / (2e + 2e^{-1})$ at $\beta = 1$ and verify that it is about 0.440.

13.2 Frustration and spin glasses

13.2.1 The frustrated triangle

Put three spins on a triangle, with all three edges antiferromagnetic ($J = -1$): every pair wants to point opposite ways. The energy is $H = \sigma_1 \sigma_2 + \sigma_2 \sigma_3 + \sigma_1 \sigma_3$.

The problem: three numbers that can only be ± 1 cannot all be pairwise different. So **no matter how you arrange them, at least one edge is unhappy**. This is called **frustration**.

The script enumerates all 8 configurations:

	Ground-state energy	Number of ground states	Zero-temperature entropy $\ln(\text{number of ground states})$
Ferromagnetic triangle ($J = +1$)	-3 (all three edges satisfied)	2 (all up, all down)	$\ln 2 \approx 0.693$

	Ground-state energy	Number of ground states	Zero-temperature entropy $\ln(\text{number of ground states})$
Antiferromagnetic triangle ($J = -1$)	-1 (only two edges can be satisfied)	6	$\ln 6 \approx 1.792$

The partition function of the antiferromagnetic triangle is $Z(\beta) = 2e^{(-3\beta)} + 6e^{(\beta)}$ (at $\beta = 1$ the script gets 16.409, matching the formula). Frustration has two consequences: (1) the ground-state energy is "raised" (-1 instead of -3); (2) the ground state is **highly degenerate** (many configurations tie for the lowest energy).

[Check 2] List the 8 configurations of the antiferromagnetic triangle with their energies, and confirm that exactly 2 have energy $+3$ and 6 have energy -1 .

13.2.2 Spin glasses: random couplings

If each edge's J_{ij} has a **random** sign (or is a random Gaussian number), frustration is everywhere: pick three random ± 1 signs, and the probability that the triangle is frustrated (the product of the three signs is negative) is exactly $1/2$ (the script enumerates the 8 sign combinations). Such a system is called a **spin glass**. The name comes from real materials (for example copper with a small amount of manganese mixed in): at low temperature their spins "freeze" in disordered directions, "orderly disordered" like glass.

There is a key distinction here that also shows up in ML: the couplings J are "**frozen randomness**" (quenched disorder), drawn once and then fixed; the spins σ are the variables that undergo "thermal motion" under the Boltzmann distribution. Analogy: J is a training set, and σ are the parameters on that training set.

13.2.3 A rugged energy landscape

"Frustration everywhere" produces a **rugged energy landscape**: many local minima, separated from each other by energy barriers. A script experiment: call a configuration where "flipping any single spin cannot lower the energy" a **single-flip stable configuration** (a local minimum), and count them:

System	N	Number of single-flip stable configurations
Fully connected ferromagnet (all $J = 1/N$)	12	2 (all up, all down)
SK spin glass (random Gaussian J)	12	13.1 on average (20 samples, min 6, max 24)
SK spin glass	16	23.2 on average (5 samples)

The literature result is that the number of single-flip stable configurations in SK grows **exponentially** with N , roughly as $e^{(0.199 \cdot N)}$ (Tanaka–Edwards 1980 / Bray–Moore 1980 $\triangle$ constant quoted from memory of the literature, not rechecked against the original).

Plugging in $N = 12$ and 16 gives 10.9 and 24.2 , consistent in magnitude with the computed values in the table above.

13.3 Free energy: why physicists want it

13.3.1 Definition

Free energy $F = -(1/\beta) \cdot \ln Z$

"Pressure" per spin, $(1/N) \cdot \ln Z$ (this is the quantity the OpenAI paper uses, called pressure)

Pressure = $-\beta F/N$: the two differ only by a factor of $-\beta/N$ ($-\beta$ from the sign and temperature, $1/N$ from "per spin"), so they are essentially the same thing.

13.3.2 Why they want it

Because $\ln Z$ is a "generating function": differentiating it gives almost every macroscopic quantity. For example, the mean energy $\langle E \rangle = -\partial(\ln Z)/\partial\beta$, and the entropy $S = \beta(\langle E \rangle - F)$. Computed by script on the frustrated triangle: at $\beta = 5$, the entropy of the antiferromagnetic triangle is $S = 1.7918$, almost exactly $\ln 6 = 1.7918$ (6 equally likely ground states); the entropy of the ferromagnetic triangle is $0.6931 = \ln 2$.

The ML counterpart: $\ln Z$ is just log-sum-exp; the ELBO is a lower bound on $\ln Z$.

13.3.3 The real difficulty: you have to average $\ln Z$

For a spin glass, Z itself is random (it depends on which J you drew). What is physically meaningful is the free energy of a **typical** sample, that is, $E[\ln Z]$ (the "quenched average").

$\ln E[Z]$ (the "annealed average") is much easier to compute: $E[Z]$ is the expectation of $\sum e^{...}$, so you can move the expectation inside the sum, and each term is a Gaussian integral. But the two are not equal. Since $\ln$ is concave, Jensen's inequality gives $E[\ln Z] \leq \ln E[Z]$.

Computed by script (SK model, $N = 10$, $\beta = 2$, 2000 random samples, exact enumeration of the 1024 configurations for each sample):

Quantity	Value per spin
Quenched $(1/N) \cdot E[\ln Z]$	1.387
Annealed $(1/N) \cdot \ln E[Z]$ (exact formula $\ln 2 + \beta^2(N-1)/(4N)$)	1.593

The gap is large. The reason is that $E[Z]$ is dominated by a very few "unusually lucky" samples and doesn't represent the typical case. For the SK model, as $N \rightarrow \infty$ the two are

equal for $\beta \leq 1$ and separate for $\beta > 1$ (low temperature); this is exactly the spin glass phase transition (an established result).

All the difficulty in this chapter comes from "you have to compute $E[\ln Z]$, and $\ln$ cannot be swapped with the expectation."

13.4 The replica method and replica symmetry breaking: a tradition that is "not rigorous, but right"

13.4.1 The replica trick

Physicists use an identity:

$$\ln Z = \lim_{n \rightarrow 0} (Z^n - 1) / n$$

Check (script): $\ln 5 = 1.6094$; for $n = 1, 0.1, 0.01, 0.001$, $(5^n - 1)/n$ is 4.000, 1.746, 1.622, 1.611 in turn, which does approach $\ln 5$.

The benefit: when n is a **positive integer**, $Z^n = Z \cdot Z \cdot \dots \cdot Z$ is the partition function of **n independent copies** ("replicas"), and $E[Z^n]$ can be computed (again, you can move the expectation inside the sum). So the strategy is:

1. Compute $E[Z^n]$ for integer n ;
2. "Pretend" the resulting formula also holds for real n ;
3. Let $n \rightarrow 0$.

Step 2 has no mathematical justification at all (the values at the integers do not uniquely determine a function on the reals), and step 3 also requires swapping the two limits $n \rightarrow 0$ and $N \rightarrow \infty$. This is why it is called "non-rigorous."

13.4.2 Replica symmetry and its failure

When computing $E[Z^n]$, an $n \times n$ matrix Q_{ab} appears, called the **overlap**:

$$Q_{ab} = (1/N) \sum_i \sigma_i^a \cdot \sigma_i^b \quad (\text{how similar the configurations of replica } a \text{ and replica } b \text{ are, a value between } -1 \text{ and } 1)$$

The most natural guess is that all replicas are "equal," so all off-diagonal Q_{ab} are the same. This is called **replica symmetric** (RS). Sherrington and Kirkpatrick did exactly this for the SK model in 1975, and the result was clearly wrong at low temperature: the zero-temperature entropy came out as $-1/(2\pi) \approx -0.159$, while the entropy of an Ising system cannot be negative (there is at least 1 configuration, and $\ln 1 = 0$); the RS prediction for the ground-state energy, $-\sqrt{2/\pi} \approx -0.798$, also disagrees with numerical simulations.

13.4.3 Parisi's replica symmetry breaking


In 1979 Parisi proposed that the replicas are not equal but form a **hierarchy**. Split the n replicas into large groups, split each large group into smaller groups, those into even smaller groups... Overlaps within the same small group are large; overlaps between different large groups are small. Splitting into r levels is called **r-step replica symmetry breaking** (r -RSB), and $r \rightarrow \infty$ is called full RSB.

The physical interpretation (built up gradually starting with Mézard–Parisi–Sourlas–Toulouse–Virasoro 1984): at low temperature the Gibbs distribution splits into many "pure states" (roughly corresponding to the deep valleys of the energy landscape), and these valleys form a **tree**: the more "closely related" two valleys are, the deeper their common ancestor in the tree. This kind of distance structure is called **ultrametric**: among any three points, the two largest distances are equal, just like the distances between species on a biological taxonomy tree.


The parameters $0 < m_1 < m_2 < \dots < m_r < 1$ that appear in the Parisi formula are the "branching strengths" at each level of the tree. The **Exponents m** in the OpenAI paper and Lean code later on refer to exactly these.

The figure below puts the two ideas side by side. On the left is the frustrated triangle from §13.2.1: all three bonds want opposite spins, none of the 8 configurations satisfies all three (the dashed bond always fails), and the ground state is 6-fold degenerate. On the right is Parisi's hierarchical tree: pure states first split into large groups, and the large groups split again; the deeper the common ancestor of two pure states (the "closer relatives" they are), the larger their overlap. The note at the bottom right, "diluted: need all multi-replica overlaps," is the difficulty that §13.7.3 explains.

Left: the smallest frustrated system · Right: the hierarchical (ultrametric) structure of Parisi



all three bonds "want" opposite spins
none of the 8 configurations satisfies all three
(dashed bond always fails): 6-fold degenerate
 $Z(\beta) = 2e^{(-3\beta)} + 6e^{\beta}$



closer relatives $\Rightarrow$ larger overlap
diluted: need all multi-replica overlaps

Spin glass = random $\pm$ couplings + frustration everywhere + exponentially many metastable states. The energy function of a Hopfield network (1982) is exactly such a Hamiltonian; perceptron capacity (Gardner 1988) was computed with the same replica method.

13.4.4 The physics was right; the math took 27 years

Year	Event	Nature
1975	Sherrington–Kirkpatrick propose the SK model (every pair of spins is coupled)	Definition
1979–80	Parisi writes down the SK free energy formula (RSB)	Physical derivation, non-rigorous
2003	Guerra uses the "interpolation method" to prove the Parisi formula is an upper bound on the free energy	Theorem
2006	Talagrand proves equality (Annals of Math. 163)	Theorem
2013	Panchenko proves ultrametricity (Annals of Math. 177)	Theorem
2021	Parisi wins the Nobel Prize in Physics (shared with Manabe and Hasselmann)	—
2024	Talagrand wins the Abel Prize (spin glass work is one of the main reasons)	—

The numerical value of the SK ground-state energy: the Parisi formula gives **−0.7632** per spin (high-precision value $-0.763166\dots$, Crisanti–Rizzo 2002 $\triangle$ as cited in the literature). The script enumerates small systems exactly ($H = -\sum_{\{i<j\}} J_{ij} \sigma_i \sigma_j$, $J_{ij} \sim N(0, 1/N)$):

N	Number of samples	Ground-state energy per spin
8	400	-0.571 ± 0.006
12	200	-0.609 ± 0.006
16	40	-0.655 ± 0.013

N	Number of samples	Ground-state energy per spin
20	8	-0.683 ± 0.022

As N grows, the value moves toward -0.763 , but slowly (finite-size corrections decay roughly like $N^{-2/3}$, $\triangle$ a result from the literature, not verified here by fitting). The point of this table: **brute-force enumeration of small systems cannot give you the limiting value; you need a formula.**

This is the most important background in this chapter: in the field of spin glasses, physicists' non-rigorous methods have repeatedly given the right answers, and mathematicians have then spent anywhere from ten to several tens of years catching up. This OpenAI result is one link in that "catching up" chain.

13.5 Sparse graphs and the cavity method

13.5.1 Dense vs. sparse

In the SK model, **every pair** of spins is coupled, each spin has $N - 1$ neighbors, and each coupling is very weak (of order $1/\sqrt{N}$). This is called a **dense** (mean-field) model.

Real physics and many computational problems are **sparse**: each spin interacts with only **finitely many** other spins, so even when N is large, each spin has only a constant number of neighbors, like 3 or 5. Two standard random constructions:

- **Poisson dilution** (Viana–Bray 1985): scatter αN edges at random (in OpenAI's theorem the number of edges is itself drawn as $\text{Poisson}(\alpha N)$); each spin's degree is approximately Poisson-distributed.
- **Bethe lattice / random regular graph**: every spin has the same **fixed** degree. The original Mézard–Parisi 2001 paper uses this one.

13.5.2 Sparse random graphs "look locally like trees"

The key property: sparse random graphs have **very few short cycles**. Computed by script (Erdős–Rényi graphs with average degree 3):

N	Average number of triangles	Average degree	Fraction of isolated vertices
100	4.8	2.82	0.058
1,000	3.8	2.99	0.053
10,000	4.1	3.01	0.049
100,000	3.4	3.00	0.049

Theoretical values: the number of triangles tends to $c^3/6 = 4.5$ (**it does not grow with N**), and the fraction of isolated vertices tends to $e^{-3} \approx 0.0498$. In other words, the bigger the graph, the closer to zero the fraction of vertices with a short cycle nearby: looking a few steps out from any vertex, what you see is almost always a tree.

13.5.3 The cavity method: exact on trees

The idea of the **cavity method**: remove one spin and see how the rest of the system "treats" the empty slot.

On a **tree** this is completely exact. After removing spin i , its neighbors belong to separate subtrees that are not connected to each other, so they are independent. Each neighbor j sends i a "message" (an effective field $u_{\{j \rightarrow i\}}$), and i adds up the messages it receives to get its own state. The update rule for messages is:

$$u_{\{j \rightarrow i\}} = (1/\beta) \cdot \operatorname{artanh}(\tanh(\beta \cdot J_{ij}) \cdot \tanh(\beta \cdot h_j^{\text{cavity}}))$$

where h_j^{cavity} is "the total field felt by j after i is removed." This is **belief propagation (BP)** from ML: the same algorithm, invented separately by physicists and by Pearl.

Smallest example ([Check 3] has you compute it by hand): two spins, $J = 1$, an external field $h = 0.5$ on spin 1, $\beta = 1$. The magnetization (mean value) of spin 2 is exactly $\tanh(1) \cdot \tanh(0.5) = 0.35195$. Exact enumeration by script gives 0.35195, which matches.

The script also does a tree with 5 vertices and random couplings and external fields: comparing the 5 magnetizations computed by exact enumeration of the 32 configurations against the BP result, the largest difference is 4.9×10^{-16} (machine precision). Add two more edges to create cycles, and the largest difference becomes **0.16**. **Cavity / BP is exact on trees and inexact with short cycles**. Sparse random graphs "look locally like trees," so there is reason to expect the cavity method to become exact as $N \rightarrow \infty$, but this needs proof.

13.5.4 The complication: with many pure states, messages become "distributions of distributions"

At high temperature, BP has a unique fixed point, with one message per edge; averaging over random graphs, the order parameter is **"a distribution of messages"** $\zeta \in P(\mathbb{R})$. This is the RS level ($r = 0$).

In the low-temperature spin glass phase there are many pure states (the tree from §13.4.3), and the message on an edge differs from one pure state to another. So you have to describe "how the message is distributed across pure states," then average over the

randomness of the graph, and the order parameter becomes "**a distribution of distributions of messages**"; with r levels of RSB, there are $r + 1$ levels of nesting:

```
Level 0:  $\mathbb{R}$  (one message)
Level 1:  $P(\mathbb{R})$  (a distribution of messages)
Level 2:  $P(P(\mathbb{R}))$  (a distribution of distributions of messages)
.....
```

The **Hierarchy** in the Lean statement is exactly this definition: $H_0 = \mathbb{R}$, $H_{\{r+1\}} = P(H_r)$. This is what "hierarchical" means in "**hierarchical cavity formula**."

Mézard–Parisi 2001 (*The Bethe lattice spin glass revisited*, EPJ B 20:217) wrote down this framework and an explicit 1RSB functional for the fixed-degree Bethe lattice; the general form for any finite number of levels r , and the proof that it is an upper bound on the free energy, come from Panchenko–Talagrand 2004 (see §13.7).

13.6 The other "Mézard–Parisi": random assignment and $\zeta(2)$

This section is not OpenAI's result, but it is the most beautiful example of "physicists' non-rigorous method $\rightarrow$ mathematical theorem," and it is the same number as $\zeta(2) = \pi^2/6$ in the quasi-Riemann hypothesis (§4).

13.6.1 The problem

n workers, n jobs. The cost C_{ij} for worker i to do job j is an independent random number, exponentially distributed with mean 1. Find a one-to-one assignment that minimizes the total cost. Question: what is the expected minimum total cost?

Intuition might say it grows with n (after all, you are adding n numbers), but the optimal assignment picks very small numbers from each row, and as n grows, the "very small numbers" get smaller and smaller.

13.6.2 Prediction and proof

- **Mézard–Parisi 1985/1987**: using the replica method, they predicted that as $n \rightarrow \infty$ the expected minimum cost $\rightarrow \zeta(2) = \pi^2/6 \approx 1.6449$.
- **Aldous 2001** (*The $\zeta(2)$ limit in the random assignment problem*, Random Structures & Algorithms 18): a rigorous proof, using "local weak convergence" (the same idea as the cavity method's "locally tree-like").
- **Parisi 1998** went further and conjectured an **exact** formula for finite n : $E[\text{minimum cost}] = 1 + 1/4 + 1/9 + \dots + 1/n^2$. This was proved independently by Linusson–Wästlund in 2004 and Nair–Prabhakar–Sharma in 2005.

Computed by script (optimal assignment via the Hungarian algorithm, `scipy.linear_sum_assignment`):

n	Number of trials	Simulated mean minimum cost	Parisi formula $\sum_{k \leq n} 1/k^2$
1	200,000	1.0000 $\pm$ 0.0022	1.0000
2	200,000	1.2498 $\pm$ 0.0019	1.2500
3	100,000	1.3628 $\pm$ 0.0023	1.3611
5	50,000	1.4614 $\pm$ 0.0026	1.4636
10	20,000	1.5500 $\pm$ 0.0029	1.5498
30	4,000	1.6093 $\pm$ 0.0038	1.6122
100	1,000	1.6373 $\pm$ 0.0042	1.6350
300	200	1.6435 $\pm$ 0.0055	1.6416
1000	40	1.6381 $\pm$ 0.0066	1.6439

Every row agrees within one or two standard errors, and for large n both approach $\pi^2/6 = 1.6449$. For comparison, a greedy method where "each worker in turn picks the cheapest remaining job" has an average cost of **5.14** at $n = 100$, more than three times the optimum.

[Check 4] For $n = 2$, with four costs a, b, c, d (matrix $[[a,b],[c,d]]$), the minimum cost is $\min(a+d, b+c)$. Prove that its expectation is $5/4$. (Hint: the sum X of two independent $\text{Exp}(1)$ variables satisfies $P(X > t) = (1+t)e^{-t}$; for a non-negative random variable, $E[Y] = \int_0^\infty P(Y > t) dt$. A direct script simulation with 2 million trials gives 1.24989.)

Lesson: the number the replica method gave ($\pi^2/6$) was right, and was later proved rigorously. Success stories like this, piling up, are the reason mathematicians are willing to spend decades proving the Parisi / Mézard–Parisi formulas.

13.7 What exactly is the problem: the upper bound is known, equality is open

13.7.1 What needs to be proved

For the diluted model, define the per-spin pressure $F_N = (1/N) \cdot E[\ln Z_N]$. The hierarchical cavity functional $B_r(\zeta, m)$ gives a number for any "trial law" ζ (the nested distributions from §13.5.4) and any branching parameters $0 < m_1 < \dots < m_r < 1$. The conjecture is:

$$\lim_{N \rightarrow \infty} F_N = \inf_{r \geq 0} \inf_{\{\zeta, m\}} B_r(\zeta, m)$$

The right-hand side takes the infimum over **all finite numbers of levels and all trial laws**.

13.7.2 Half of it is already proved: the upper bound (2003–04)

- **Franz–Leone 2003** (J. Stat. Phys. 111): carried Guerra's interpolation method over to diluted models and proved the RS and 1RSB upper bounds for even-arity models.
- **Panchenko–Talagrand 2004** (PTRF 130): proved the upper bound for any finite number of levels r : for every r, ζ, m , $\limsup F_N \leq B_r(\zeta, m)$. After Theorem 4 they **explicitly state equality as a conjecture**. The OpenAI paper says this is the conjecture it resolves (within the same model class).

The intuition for the interpolation method: build a path that continuously deforms "the real model" into "a trial model made of independent messages," and prove that the derivative of the pressure along the path has a fixed sign. Determining the sign uses a positivity condition; this is where $E[(-b)^n] \geq 0$ in the Panchenko–Talagrand condition below comes from, and also why only **even-arity** interactions can be handled.

13.7.3 Why the lower bound is hard

In SK, when Talagrand proved the lower bound, he only needed to control the **pairwise overlaps** Q_{ab} . In a sparse model, a newly inserted spin connects simultaneously to $p - 1$ old spins, and you need to know what pattern several replicas see **at the same time** at those positions. This depends on all the **multioverlaps**:

$$R_S = (1/N) \sum_i \prod_{a \in S} \sigma_i^a \quad (S \text{ is a set of replicas})$$

The ultrametric structure of pairwise overlaps does not determine all the multioverlaps (in the words of §1 of the paper: "pair overlaps alone do not determine all these patterns").

Existing partial progress (listed in the paper's introduction): Panchenko 2013 gave a different exact variational formula using "spin distributions," but not in hierarchical cavity form; Panchenko 2016 (arXiv:1406.4702) proved for "modified models" that finite-RSB Gibbs measures have the Mézard–Parisi structure, but was missing the step "a general state can be approximated by finite RSB"; Coja-Oghlan–Perkins 2019 gave another variational formula for random regular factor graphs under a positivity condition; Concetti 2025 (arXiv:2508.17049) gave a full-RSB upper bound for random regular graphs, with equality still a conjecture.

Neighboring known results (background):

- **Ferromagnetic** sparse Ising models (no frustration): Dembo–Montanari 2010 proved the Bethe formula.
- **Random k-SAT** (each clause is a k -ary interaction): Ding–Sly–Sun (STOC 2015, Annals 2022) proved the satisfiability threshold for **sufficiently large k** , which is exactly what the 1RSB formula predicts.

13.8 What OpenAI claims

13.8.1 The main result

Abstract of the paper *The Mézard–Parisi formula for diluted spin glasses* (OpenAI, 2026–09–23, 36 pages, family 221) [P]:

"We prove the Mézard–Parisi hierarchical cavity formula for diluted even-arity Ising models in the Panchenko–Talagrand class. This resolves the variational equality conjecture for that class: the limiting pressure equals the infimum of the trial functional over all finite hierarchy depths and trial laws. Only first moments of the interaction and external field are required."

Main theorem (Theorem 2.1): under assumptions (2.1)–(2.4), the thermodynamic limit exists, and $\lim F_N = \inf_r \Phi_r$.

The key trick in one sentence: shift every branching depth down one level, then average over depths, which gives control of all multi-replica overlaps (step 3 in §13.8.4).

13.8.2 What the model is (plain-language version)

- Fix an **even** $p \geq 2$ (how many spins each interaction involves) and a density $\alpha > 0$.
- Scatter $M_N \sim \text{Poisson}(\alpha N)$ interactions; each picks p spins at random (**with replacement**, so repeats are allowed) and applies a random function θ . Each spin also gets a random external field h .
- The only requirement is $E[\max|\theta|] + E[|h|] < \infty$ ("only first moments are needed": the interactions are not required to be bounded).

The Panchenko–Talagrand condition: e^θ can be written in the form $a \cdot (1 + b \cdot f_1(s_1) \cdots f_p(s_p))$, where the f 's are i.i.d., b is independent of them, and $E[(-b)^n] \geq 0$ holds for all $n \geq 1$.

Check it on the Viana–Bray model ($p = 2$, $\theta = \beta J \cdot s_1 s_2$, with J symmetrically distributed): take $a = \cosh(\beta J)$, $b = \tanh(\beta J)$, $f(s) = s$.

[Check 5] Verify that $e^{(\beta J \cdot s_1 s_2)} = \cosh(\beta J) \cdot (1 + \tanh(\beta J) \cdot s_1 s_2)$, separately for $s_1 s_2 = +1$ and -1 . (The script, with $\beta J = 0.7$, verifies that both sides are 2.01375 and 0.49659.)

Because J is symmetric, the odd moments of $b = \tanh(\beta J)$ are 0 and the even moments are positive, so the positivity condition holds.

The coverage listed in Section 10 of the paper (OpenAI's claim):

- the **Viana–Bray model** ($p = 2$, symmetric random couplings);
- **diluted p -spin models with even p** (symmetric couplings);

- **soft random even K-SAT**, at every positive temperature ("soft": a violated clause is only penalized by energy βW , not strictly forbidden).
- Corollary 10.3: lowering the temperature to zero gives the limit of "maximum energy / minimum number of violated clauses," with error $\leq \ln 2 / \beta$.

13.8.3 What it does not cover (important)

- **Odd arity**: random **3-SAT** is not included. The paper's Remark 10.4 says so itself, and points to another OpenAI paper from the same batch, *Computing the Random 3-SAT Threshold* (2026-09-27), which uses a different trial functional (OpenAI's claim; not read for this chapter).
- **Fixed-degree Bethe lattice / random regular graphs**: the degrees in the theorem come from Poisson edge scattering. The setting of the original Mézard–Parisi 2001 paper is exactly the fixed-degree one, and that case is **not** covered by the theorem. In other words, the "Mézard–Parisi formula" in the paper's title refers to its counterpart on Poisson-diluted models.
- **Ferromagnetic models** (my inference; I haven't seen the paper say it explicitly $\triangle$): for fixed $J > 0$, the standard decomposition gives $b = \tanh(\beta J) > 0$, so $E[(-b)^n]$ alternates in sign and the positivity condition fails. But the ferromagnetic case was already handled by Dembo–Montanari 2010, so it is not a gap.
- **The zero-temperature hard-constraint SAT threshold**: the theorem is a result at every positive temperature; the zero-temperature limit only gives the minimum number of violations, not the satisfiability threshold directly. So it does **not** re-prove or extend Ding–Sly–Sun.
- **Hopfield networks**: fully connected (dense), with couplings determined by the stored patterns, so different edges are not independent; they are outside the theorem. The connections in §13.10.3 (Hopfield, perceptrons, loss landscapes) are about where the ideas came from, not consequences of this theorem.
- It **does not claim** that the infimum is attained by some trial law, and it does not construct an infinite-level limiting Gibbs state (the remarks after Theorem 2.1 in the paper).

13.8.4 Intuition for the method (I only read the introduction and §13.2.2; what follows is a simplified version)

The upper bound reuses the interpolation method (§13.4). The lower bound has four steps:

1. **Choose a "cavity reservoir"**: compare the system with N spins and with $N + 1$ spins, adding a sublinear number of small perturbations. These don't change the pressure, but they make the system satisfy certain "identities" (the perturbed statistics concentrate).
2. **Prescribe branching patterns**: develop a "calculus" on finite probability trees that expands the effect of inserting a new spin into a signed sum over replica paths, using independent

labels to pick out any specified replica branching pattern.

3. **Concentrate all multioverlaps** (the core trick): **shift every internal branching depth of the target tree down one level, all at once**. That way the new branchings always land on depths the old tree hasn't used; each new branching pays a small factor, which exactly cancels the small factor in the normalizing coefficient, so the error stays controllable after division. Then **average over the branching depths** (it doesn't need to hold uniformly for every depth), and inductively make the conditional covariances decay.
4. **Recover independent cavity messages**: once the multioverlaps are concentrated, the messages at different positions can be replaced by independent copies, giving a valid trial value that is close to the true pressure.

Order of limits: for each fixed depth, first let $N \rightarrow \infty$, then deepen the hierarchy. Finally a truncation removes the "bounded interactions" assumption, leaving only first moments.

The reasoning summary ([reasoning_traces/mezard-parisi-formula.pdf](#), 5 pages) records this process. Two lines verbatim:

"Aha recursive structure: overfitting shared variable can simulate original model at smaller scale."

"Great! Need carefully prove average error little with shifted depths not rare"

The summary also candidly records failed attempts: the pinning method (fixing some spins) could not reconcile the observed distributions; a claim of "exact purity" "proves too strong"; and before "shifting the branching depths," various expansions all had errors that were too large after dividing by small exponent gaps.

13.8.5 Neighboring results in the same batch (OpenAI's claims; not verified for this chapter)

The same release also contains a whole group of spin glass results: family 222 (perceptron free energy, jamming exponents), 281 (full support of the zero-temperature SK Parisi measure; QAOA reaching the SK ground-state energy), 235 (computability and variance of the random k-SAT threshold; priority for the existence of the 3-SAT threshold goes to Gaia Carenini's ECCC TR26-229), 217 and others (SK fluctuations and dynamics). §13.10.3 uses 222 and 281.

13.9 How well has it been verified

13.9.1 The Lean statement

The main theorem of the challenge file [lean/ComparatorChallenges/DilutedSpin.lean](#) (154 lines):

```
theorem mezard_parisi {p : ℕ} (M : Model p) (hM : Admissible M) :
  Tendsto (pressure M) atTop (λ (variationalValue M))
```

Word-by-word translation:

- `{p : ℕ}`: p is a natural number, the arity of the interactions.
- `(M : Model p)`: a model, made of three things: the density `alpha`, the distribution of interactions `disorder` (together with the decomposition witnesses θ, a, b, f), and the distribution of external fields `field`.
- `(hM : Admissible M)`: the assumption that the model satisfies the conditions. The `Admissible` structure spells them out one by one: `arity : 2 ≤ p`, `even : Even p`, `density : 0 < alpha`, two first-moment integrability conditions, `factorization` ($e^\theta = a(1 + b \cdot \prod f)$, with $|b \cdot \prod f| < 1$), the f 's i.i.d. and b independent of the f 's, all moments of b integrable, and `positivity : 0 ≤ E[(-b)^n]`. These correspond one-to-one with (2.1)–(2.4) in the paper.
- `pressure M`: a sequence whose N -th term is F_N . The definition is written very concretely: take the expectation over Poisson(αN) interactions and i.i.d. external fields, **average uniformly** over all possible index choices (implementing "uniform sampling with replacement"), take log-sum-exp over the 2^N configurations, and divide by N .
- `variationalValue M`: $\inf_r \phi M r$, i.e. the infimum over all numbers of levels r ; `phi M r` is the infimum of `functional` over all $\zeta \in \text{Hierarchy } (r+1)$ and all m satisfying `Exponents` ($0 < m_1 < \dots < m_r < 1$), and `functional` corresponds term by term to the paper's (2.11) $B_r = \ln 2 + E \ln T_0 V_s - \alpha(p-1) \cdot E \ln T_0 V_e$.
- `Tendsto ... atTop (λ ...)`: as $N \rightarrow \infty$, F_N converges to this value. This also includes "the limit exists."

The difference from the quasi-Riemann hypothesis (§4): that statement used only Mathlib's standard definition `riemannZeta`; **here almost all the definitions (`pressure`, `Hierarchy`, `functional`, ...)** are written by the challenge file itself, so "is the statement faithful" depends on a human reading these 154 lines. I read them item by item against Section 2 of the paper [P] and found no inconsistencies, but I am not an expert in this field $\triangle$.

13.9.2 The proof itself

- Solution module `0AI/Probability/DilutedSpin/`: **405 .lean files, 56,285 lines** (counted by script) [P]. grep for `sorry` / `admit` / `native_decide` / `axiom` declarations: **0**.
- The definitions in the solution module's `Model.lean` are **word-for-word identical** to those in the challenge file (I diffed the section from `Spin` to `variationalValue`).

- Comparator config `DilutedSpin.json`: only Lean's three standard axioms are allowed (`propext`, `Quot.sound`, `Classical.choice`); `enable_nanoda: false`, meaning no second, independent kernel is configured for a re-check.
- **We did not actually compile the proof or run Comparator** (the full library is very large). `formalization.yaml` is marked `review: status: unchecked` overall.
- Unlike the quasi-Riemann hypothesis (§4), **I found no report of any third party (someone like Goldblatt) independently running this proof successfully** $\triangle$.

13.9.3 Human experts

As of 2026-10-08, I found no named comment or verification report on this paper from any spin glass / probability expert $\triangle$ (searched; I only found aggregator blogs restating family 221). Panchenko, Talagrand, Coja-Oghlan, and others are the natural referees.

13.9.4 Remaining risks

1. **Faithfulness of the statement:** Lean proves the statement written in the challenge file. A few places deserve a close look from an expert: each level of `Hierarchy` uses the Borel σ -algebra of the weak topology (the file's comments say this follows the paper); the indexing convention for m in `LogMean` (`m 0` corresponds to the paper's m_1); and `pressure` uses "average over all indices" in place of "random sampling." I believe these are all equivalent, but I haven't checked formally $\triangle$.
2. **Paper title vs. scope:** readers can easily take "the Mézard–Parisi formula" to mean the Bethe lattice (fixed-degree) case of Mézard–Parisi 2001, or to include 3-SAT. Neither is included (§13.8.3). This is the "spec vs. intent" gap that keeps recurring in this series.
3. **Readability of the proof:** a 36-page paper against 56,000 lines of Lean; the paper is written in summary style. It may take humans quite some time to understand why the "shifting the branching depths" trick works.

13.10 What it would mean if true (and what it would not mean)

13.10.1 For mathematics

1. **A roughly 20-year-old open conjecture (Panchenko–Talagrand 2004) is resolved in its original model class**, and it needs only first moments, a weaker assumption than the original problem set.
2. **The new technique may matter more than the conclusion:** "shift the branching depths + average over depths" to control all multioverlaps. If it holds up, it may generalize to other sparse models (fixed degree, odd arity). The paper itself says the lower-bound part **holds for any bounded interaction that is invariant under permutations of its inputs**, without needing the decomposition or the positivity condition (end of §2.2 of the paper); what is

stuck is the upper bound. In other words, if someone later extends the interpolation upper bound to odd arity, the lower-bound half is already done.

3. **What it does not mean:** the formula for the random 3-SAT threshold, a complete solution of the Bethe lattice spin glass, and "ultrametricity" on sparse models are all still unsolved.

13.10.2 For physics

Physicists already "knew" in 2001 that this formula was right, and used it to compute a large number of numerical predictions. The theorem doesn't change any physical prediction. Its role is to **upgrade the cavity method from "a trustworthy heuristic" to a theorem** for one class of sparse systems. This is the same kind of significance as Talagrand 2006 for SK, but with a narrower scope.

13.10.3 For ML and neuroscience (the honest version)

Connections that really exist:

- **A Hopfield network (1982) is an Ising system.** Energy $E = -\frac{1}{2} \sum W_{ij} \sigma_i \sigma_j$, and the Hebb rule $W_{ij} = (1/N) \sum_{\mu} \xi_i^{\mu} \xi_j^{\mu}$ stores P patterns; each memory is a minimum of the energy. When too many are stored, the "crosstalk" from the other patterns acts like random couplings and creates frustration, and the system falls into "spin glass states" that don't correspond to any memory. Amit–Gutfreund–Sompolinsky 1985 used the replica method to compute the capacity **$\alpha_c = P/N \approx 0.138$** . Script simulation ($N = 1000$, zero-temperature asynchronous updates, starting from a stored pattern, 30 runs per α):

$\alpha = P/N$	0.05	0.10	0.12	0.13	0.14	0.15	0.17	0.20
Mean final overlap m	1.000	0.998	0.992	0.974	0.941	0.917	0.583	0.363

Retrieval quality clearly collapses between 0.13 and 0.17, consistent with the 0.138 prediction. At $N = 1000$ the transition is smeared out by finite-size effects, and past 0.138 the overlap doesn't drop to zero right away, so this table is only a qualitative check. Hopfield and Hinton won the 2024 Nobel Prize in Physics.

- **Perceptron capacity:** Gardner 1988 used the replica method to compute the capacity of the spherical perceptron, $\alpha = 2$; for the Ising perceptron (weights ± 1), Krauth–Mézard 1989 predicted about 0.833, and rigorous proofs have only recently made partial progress (Ding–Sun 2019 proved the lower bound; Huang 2024 gave a matching upper bound under a numerical assumption $\triangle$ not rechecked). Family 222 in the same batch gives the free energies of the Ising perceptron and the spherical perceptron at **positive temperature**, plus jamming exponents for the spherical perceptron with negative margin (-1) (OpenAI's claim, per the 222 entry in the repo's `CONTENTS.md`), not the zero-temperature capacity itself.

- **Random optimization landscapes and algorithms:** Montanari 2019 (arXiv:1812.10897) gave a polynomial-time algorithm that approximately finds the SK ground state, **provided** the zero-temperature Parisi measure has "no overlap gap." Family 281 in the same batch claims to prove this full-support property (OpenAI's claim, not verified $\triangle$). If it holds, "the SK ground state can be approximated efficiently" becomes unconditional. This is the most direct answer spin glass theory gives to "can random landscapes be optimized efficiently," and it is closer to ML's concerns than 221 itself.
- **Inference on sparse graphs:** LDPC decoding, community detection, and the information-theoretic thresholds of sparse constraint satisfaction problems all rely on cavity / BP-type variational formulas. 221 adds one more rigorous cornerstone to arguments of the form "physics prediction = true threshold" (speculative extrapolation).

Where to discount:

- **221 does not cover Hopfield.** Hopfield is fully connected (dense), and its couplings are determined by the patterns, so different edges are not independent; 221 requires sparse models with independent interactions. The "sparse Hopfield" models in neuroscience (for example the strongly diluted model of Derrida–Gardner–Zippelius 1987) also don't fit the Panchenko–Talagrand decomposition.
- The intuition of a "rugged landscape with hierarchically clustered minima" for loss landscapes comes from spin glasses; the best-known example is the analogy **deep network loss landscape $\approx$ spherical spin glass** (Choromanska et al. 2015). Such analogies rely on strong simplifying assumptions and are mostly **heuristic**. This theorem won't tell you what SGD will find, where it converges, or how well a model generalizes.
- The second half of this chapter's title, "from magnets to the energy landscapes of neural networks," is about an **intellectual lineage** (energy functions, the replica method, rugged landscapes), not a claim that this theorem has direct consequences for neural networks.

13.10.4 For AI

This is a textbook case of "AI fills in a missing step in a classic research program": the problem has a clear origin (the conjecture after Theorem 4 of PT 2004), a Lean statement, and a candid reasoning summary. It is also a teaching case for "spec vs. intent": the scope Lean covers (Poisson, even arity, positivity) is narrower than the paper's title suggests, and judging that gap takes domain knowledge; Lean can't do it for you.

13.11 Self-test

Try to answer these after reading. If you can't, go back and reread the corresponding section.

1. Why do physicists compute $E[\ln Z]$ rather than the easier $\ln E[Z]$? When do the two differ? (§13.3.3)
2. Which step of the replica trick has no mathematical justification? (§13.4.1)
3. What is wrong with the replica-symmetric SK solution? Give a clearly unphysical number. (§13.4.2)
4. Why is cavity / BP exact on trees but inexact on graphs with short cycles? Why is there still hope for sparse random graphs? (§13.5.2–13.5.3)
5. What does "hierarchical" refer to in the hierarchical cavity formula? How is **Hierarchy** defined in Lean? (§13.5.4, 9.1)
6. What is the relationship between the $\pi^2/6$ of the random assignment problem and the formula OpenAI proved? (§13.0, 6)
7. In the Panchenko–Talagrand conjecture, which half was proved long ago? Why is the other half harder in sparse models than in SK? (§13.7.2–13.7.3)
8. List three classes of models that OpenAI's theorem does **not** cover. (§13.8.3)
9. Why can't you say "this theorem proves such-and-such a property of Hopfield networks"? (§13.10.3)

Appendix: answers to the checks

- **Check 1:** $e^1 \approx 2.7183$, $e^{-1} \approx 0.3679$, $2e + 2e^{-1} \approx 6.1724$, so $P(+,+) \approx 2.7183 / 6.1724 \approx 0.4404$ (script: 0.44040).
- **Check 2:** $H = \sigma_1\sigma_2 + \sigma_2\sigma_3 + \sigma_1\sigma_3$. For $(+,+,+)$ and $(-,-,-)$, all three terms are $+1$, so $H = +3$. The other 6 configurations all have "two of one sign, one of the other": the same-sign pair contributes $+1$, and the other two pairs contribute -1 each, so $H = -1$. So there are 2 with $+3$ and 6 with -1 .
- **Check 3** (the example from §13.5.3): only spin 1 has an external field. $Z = \sum e^{\sigma_1\sigma_2 + 0.5\sigma_1}$. The total weight for $\sigma_2 = +1$ is $e^{(1.5)} + e^{(-1.5)}$, and for $\sigma_2 = -1$ it is $e^{(-0.5)} + e^{(0.5)}$. $m_2 = [e^{1.5} + e^{-1.5} - e^{-0.5} - e^{0.5}] / [e^{1.5} + e^{-1.5} + e^{-0.5} + e^{0.5}] = (4.7048 - 2.2553)/(4.7048 + 2.2553) \approx 0.3519 = \tanh(1) \cdot \tanh(0.5)$.
- **Check 4:** $a + d$ and $b + c$ are two independent $\text{Gamma}(2, 1)$ variables, with $P(X > t) = (1+t)e^{(-t)}$. So $P(\min > t) = (1+t)^2 e^{(-2t)}$, and $E[\min] = \int_0^\infty (1 + 2t + t^2) e^{(-2t)} dt = 1/2 + 2 \cdot (1/4) + 2/8 = 5/4$.
- **Check 5:** when $s_1 s_2 = +1$, the right-hand side $= \cosh(\beta J) + \sinh(\beta J) = e^{(\beta J)}$; when $s_1 s_2 = -1$, the right-hand side $= \cosh(\beta J) - \sinh(\beta J) = e^{(-\beta J)}$. In both cases it equals the left-hand side.

Appendix: Glossary

Term	One-line explanation
Spin	A variable taking only ± 1 , representing a tiny compass needle (or a binary neuron)
Ising model	A spin system with energy $-\sum J_{ij} \sigma_i \sigma_j$
Boltzmann distribution	$P(\sigma) \propto e^{(-\beta H(\sigma))}$, i.e. a softmax with $-\beta \cdot \text{energy}$ as the logit
Partition function Z	The sum of the weights of all configurations; the normalizing constant
Free energy / pressure	$F = -(1/\beta) \ln Z$; pressure $= (1/N) \ln Z$; they differ by a factor
Frustration	A situation where it is impossible to satisfy all interactions at once
Spin glass	A spin system with randomly signed couplings and frustration everywhere
Quenched / annealed average	$E[\ln Z] / \ln E[Z]$; physics needs the former
Replica method	A non-rigorous trick that computes $E[\ln Z]$ via $\ln Z = \lim (Z^n - 1)/n$
Overlap Q_{ab}	The similarity of two replica configurations, $(1/N) \sum \sigma_i^a \sigma_i^b$
Multioverlap	The spatial average of a product of spins across several replicas; sparse models require controlling all of them
RSB (replica symmetry breaking)	Replicas form a hierarchy (a tree), corresponding to many pure states
Ultrametric	A distance structure where, among any three points, the two largest distances are equal, like a taxonomy tree
SK model	A dense spin glass in which every pair of spins has a random coupling
Viana–Bray model	A Poisson-diluted two-body spin glass
Bethe lattice	A sparse (random regular) graph in which every vertex has the same fixed degree
Cavity method / BP	Remove a spin and compute its state from messages sent by its neighbors; exact on trees
Hierarchical cavity formula	A variational formula expressing the free energy in terms of "distributions of distributions of messages..."
Panchenko–Talagrand class	Even-arity interactions with $e^\theta = a(1 + b \cdot \Pi f)$ and $E[(-b)^n] \geq 0$
Interpolation method	Guerra's trick: interpolate continuously between the real model and a trial model to obtain an upper bound

Appendix: Further reading

- **Mézard & Montanari**, *Information, Physics, and Computation* (Oxford University Press, 2009): covers spin glasses, the cavity method, BP, random k -SAT, and coding in one book, written for CS readers; Chapters 14–19 cover exactly §13.5–13.7 of this chapter. Most recommended.

- **Panchenko**, *Introduction to the SK model* (survey, arXiv:1412.0170): <https://arxiv.org/abs/1412.0170> ; the full version is his book *The Sherrington–Kirkpatrick Model* (Springer, 2013).
- **Nobel 2021 popular information** (the Parisi part): <https://www.nobelprize.org/prizes/physics/2021/popular-information/>
- **Mézard & Parisi 2001**, *The Bethe lattice spin glass revisited*, EPJ B 20:217 (arXiv:cond-mat/0009418 $\triangle$ ID from memory, please verify).
- **Panchenko & Talagrand 2004**, *Bounds for diluted mean-fields spin glass models*, PTRF 130:319, <https://doi.org/10.1007/s00440-004-0342-2> .
- **Panchenko 2016**, *Structure of finite-RSB asymptotic Gibbs measures in the diluted spin glass models*: <https://arxiv.org/abs/1406.4702>
- **Aldous 2001**, *The $\zeta(2)$ limit in the random assignment problem*, Random Structures & Algorithms 18:381–418.
- **Montanari 2019**, *Optimization of the Sherrington–Kirkpatrick Hamiltonian*: <https://arxiv.org/abs/1812.10897>
- **Concetti 2025** (random regular graphs, the open part): <https://arxiv.org/abs/2508.17049>
- The original paper: <https://github.com/openai/math/blob/main/preprints/The-Mezard-Parisi-formula-for-diluted-spin-glasses-September-23-2026/paper.pdf>
- The Lean statement: <https://github.com/openai/math/blob/main/lean/ComparatorChallenges/DilutedSpin.lean> ; scope notes: <https://github.com/openai/math/blob/main/lean/docs/221.md>
- The reasoning summary: https://github.com/openai/math/blob/main/reasoning_traces/mezard-parisi-formula.pdf

HARMONIC ANALYSIS · FAMILY 074 (+ 077)

14. The Kakeya problem and Fourier restriction: how little room a needle needs to turn around, and the next step after Hong Wang

The only prerequisites: you can work with powers and logarithms, you know vectors and the dot product, and you know that a Fourier decomposition "breaks a signal into waves" (§4.3.1). Words like dimension, finite field and maximal function are explained the first time they appear. Every number in the text is either computed by script (script: `精讲/scripts/11_kakeya.py`, with results copied into the text) or has a source. The **[Check]** items are small exercises you can work out yourself; answers are at the end of the chapter.

Tags: [P] = primary material read (paper, openai/math repository file or Lean source); [R] = secondary reporting (news, social media, encyclopedias); $\triangle$ = something I could not verify, or where sources conflict; [established] / [recent] / [speculation] separate textbook results, new results from 2024–2026, and my own judgment. Estimated reading time: 60–90 minutes. You can read it in three sittings: §14.1–14.4 (needles, dimension, the Kakeya conjecture, the finite-field toy), §14.5–14.6 (maximal functions, waves and Fourier restriction), §14.7–14.9 (what OpenAI claims, how far it has been verified, what it means).

14.0 The bottom line: the one thing this chapter explains

First, the line you may have heard: "OpenAI also helped Hong Wang solve the four-dimensional Kakeya problem." **That is not accurate.** The accurate version is:

- Hong Wang and Joshua Zahl proved the **three-dimensional** Kakeya set conjecture in February 2025 (arXiv:2502.17655), and Wang received the 2026 Fields Medal for it [R].
- Two OpenAI manuscripts dated 2026-09-23/24 (family 074) claim to prove **the next two targets**: the three-dimensional **Kakeya maximal conjecture** (stronger and more quantitative than the version Wang proved), and the **four-dimensional** Kakeya set conjecture [P].
- The author line of both papers says only "OpenAI", and no public material says Wang took part. But the proofs make heavy use of tools that Wang and her collaborators built over the last few years [P]. A more accurate way to put it: **AI, standing on Hong Wang's toolbox, claims to have finished on its own what she and her peers widely regard as "the next target"**.
- As of 2026-10-09: neither paper has a Lean proof or a Comparator challenge, the repository contains only a 761-line, half-finished file that proves nothing beyond the trivial bound, and I found no human expert saying publicly that they have read and verified the proofs [P] $\triangle$.

This chapter explains a causal chain:

A needle can turn around while sweeping an arbitrarily small area $\Rightarrow$ area can't measure the size of such sets, so we switch to "dimension" $\Rightarrow$ the Kakeya conjecture: a set with a needle in every direction must have full dimension $\Rightarrow$ its quantitative version (the maximal conjecture) controls "how much thin tubes can overlap" $\Rightarrow$ wave packets are thin tubes, so it holds up a whole chain of problems in Fourier analysis (restriction, Bochner–Riesz, local smoothing) $\Rightarrow$ OpenAI claims to have broken through this whole chain at once in three dimensions, and to have pushed the set conjecture to four dimensions.

After reading this chapter, you should be able to explain in your own words:

- why "area zero" and "full dimension" don't contradict each other, and how box dimension differs from Hausdorff dimension;
- how the Kakeya set conjecture, the maximal conjecture and the restriction conjecture compare in strength, and why "tubes lit along only a fraction λ of their length" are the key;
- why the finite-field version can be proved in a few paragraphs (Dvir 2008), and work through it by hand on F_3^2 ;
- what exactly OpenAI claims, how far it has been verified, and what is wrong with "OpenAI helped Hong Wang".

14.1 How a needle turns around: Kakeya's problem (1917)

14.1.1 The problem and three candidate shapes

Put a needle of length 1 in the plane. It has to move **continuously** and turn around by 180° (so its two ends swap places). What is the smallest possible area of the region it sweeps? The problem comes from a 1917 paper by Fujiwara and Sōichi Kakeya (both OpenAI papers cite it [P]).

Region	How the needle turns	Area (computed by script)
Disk of diameter 1	Spins in place about the center	$\pi/4 \approx 0.7854$
Equilateral triangle of height 1	Pivots along the three sides in turn	$1/\sqrt{3} \approx 0.5774$
Kakeya's deltoid	The needle stays tangent to the curve	$\pi/8 \approx 0.3927$

- Equilateral triangle: Pál proved (around 1921) that it is the smallest among **convex** regions [R].
- Deltoid: the curve traced by a point on a small circle of radius r rolling inside a large circle of radius $3r$. Every tangent line cuts out a segment of length $4r$ inside it, so taking $r = 1/4$ gives length 1. The script checks at three tangent points that the tangent segment has length **1.000000**, and the shoelace formula gives the area **0.392699** $= \pi/8$. Kakeya seems to have conjectured that it is the smallest [R].

[Check 1] For an equilateral triangle of height 1, what are the side length and the area? (Hint: height = side $\times \sqrt{3}/2$.)

14.1.2 Besicovitch: the area can be arbitrarily small

The answer is a surprise: **there is no positive minimum**. In 1928 Besicovitch proved that a needle can turn around inside a region of arbitrarily small area [P, per the introductions of both OpenAI papers]. Even earlier he had constructed a set of **area zero** that contains a unit segment in every direction (Wikipedia dates it to 1920 [R]; $\triangle$ go by the original paper for the year).

Two intuitions:

1. **Perron trees.** A triangle with a 60° apex angle contains segments in every direction within a 60° range. Cut the base into 2^k pieces to get 2^k thin triangles, then slide them left and right so that they overlap heavily; each piece still keeps its own range of directions. The total area can be pushed arbitrarily low, and three copies rotated by 60° from each other together cover all directions [R].
2. **Sliding along itself costs no area.** To shift the needle sideways a little, first slide it a long way along its own line, turn it by a tiny angle, and slide it back. This sweeps only two long, thin triangles.

The OpenAI three-dimensional paper takes care to point out that being able to turn around continuously and having a needle in every direction are two different things ("These are different requirements.") [P]. The second does not require the needle to move.

Definition: a set in $\mathbb{R}^n$ that contains a unit segment in every direction is called a **Keakeya set** (also a Besicovitch set). Its area (volume) can be 0. Then what is left to ask? The point is that **sets of area zero still come in different sizes**. A segment and a Cantor set both have area zero, but the segment is obviously "bigger". The "dimension" of the next section is the ruler for this kind of size.

14.2 Dimension: how to measure the size of a set with "area zero"

14.2.1 Counting boxes (box dimension)

Cover a set with small squares of side δ , and let $N(\delta)$ be the smallest number you need. A segment needs δ^{-1} of them, a square needs δ^{-2} . The pattern is $N(\delta) \approx \delta^{-(d)}$, and the exponent d is the dimension:

$$d \approx \log N(\delta) / \log(1/\delta) \quad (\text{more accurate as } \delta \text{ gets smaller})$$

This is called the **box dimension** (box-counting dimension, also Minkowski dimension). The script's actual counts: for a segment, $\delta = 0.1, 0.01, 0.001$ give $N = 10, 100, 1000$, and the estimate is 1.0000; the same computation for a square gives 2.0000.

14.2.2 The Cantor set: dimension need not be an integer

Start from $[0, 1]$, remove the middle third, then remove the middle third of each remaining piece, and repeat forever. What is left, the **Cantor set**, has total length 0 (each step keeps $2/3$), yet it has uncountably many points. Cover it with boxes of size $\delta = 3^{-k}$: after step k there are 2^k pieces left, so $N = 2^k$. The script uses level 12 as the approximation: at $\delta = 3^{-2}, 3^{-4}, \dots, 3^{-10}$, $N = 4, 16, 64, 256, 1024$, and every estimate is **0.6309**.

[Check 2] Use $N(3^{-k}) = 2^k$ to prove that the box dimension of the Cantor set is exactly $\log 2 / \log 3$.

14.2.3 Hausdorff dimension: allowing boxes of different sizes

Box dimension requires all the boxes to be **the same size**, and that causes trouble. Take the countable set $\{1, 1/2, 1/3, 1/4, \dots\}$: intuitively it is 0-dimensional. But the script's box counts are $N = 20, 63, 200, 632$ at $\delta = 10^{-2}, 10^{-3}, 10^{-4}, 10^{-5}$, exactly $2/\sqrt{\delta}$, so the box dimension is **1/2**. The reason is that the points crowd closer and closer together near 0, and equal-sized boxes can't see that they are really isolated points.

Hausdorff dimension allows covers by sets of **different sizes**. Write r_i for the diameter of the i -th piece and ask whether the "s-dimensional cost" $\sum r_i^s$ can be made arbitrarily small; the smallest s for which it can be pushed to 0 is the Hausdorff dimension. The OpenAI four-dimensional paper uses exactly this covering definition, and stresses that it doesn't require the set to be measurable [P].

[Check 3] Give the k -th point of $\{1/n\}$ an interval of length $\varepsilon \cdot 2^{-k}$. Prove that for every $s > 0$, $\sum (\varepsilon \cdot 2^{-k})^s$ tends to 0 as $\varepsilon \rightarrow 0$.

In general, **Hausdorff dimension $\leq$ lower box dimension $\leq$ upper box dimension**. So "Hausdorff dimension = n " is the strongest statement, and it automatically implies that the box dimension is n too (§11.1 of the four-dimensional paper argues exactly this way [P]).

14.3 The Kakeya conjecture: the area can be zero, but the dimension can't drop

14.3.1 The statement and the tube version

Kakeya set conjecture: every Kakeya set in $\mathbb{R}^n$ has Hausdorff dimension (and therefore box dimension) equal to n .

In plain words: you can pack the needles so tightly that the total volume is zero, but not so tightly that the set becomes "thinner" by a dimension.

What people actually work with is a finite version: thicken each needle into a **δ -tube** of radius δ and length 1. In $\mathbb{R}^3$ there are about δ^{-2} tubes whose directions differ pairwise by at least δ , each of volume $\pi\delta^2$, so if they didn't overlap their total volume would be $\approx \pi$ (computed by script: 3.1416 for every δ from 0.1 to 0.0001). The question becomes: **how much can this many thin tubes, all pointing in different directions, overlap?** The box-dimension version of the conjecture is roughly equivalent to: the volume of the union is $\geq c_\varepsilon \delta^\varepsilon$ for every $\varepsilon > 0$. Overlap can only "save an arbitrarily small power of δ ".

The plane is where you can see clearly how "area zero" and "full dimension" coexist. Córdoba's 1977 planar maximal estimate [P, cited] implies that in the plane the union of such tubes has area at least about $1/\log(1/\delta)$, and Besicovitch-type constructions show that this is roughly attained ($\triangle$ I have not checked the sharpness details against the original). The script's actual numbers: at $\delta = 10^{-2}, 10^{-6}, 10^{-12}, 10^{-100}$, $1/\ln(1/\delta) = 0.2171, 0.0724, 0.0362, 0.0043$. The area does go to 0, but more slowly than any δ^ϵ , so the dimension is still 2. **"Area zero" is a logarithmic saving; "less than full dimension" would need a power saving.**

14.3.2 A hundred years of progress

Year	Who	Result
1917	Fujiwara–Kakeya	Posed the needle problem [P] cited
1920/1928	Besicovitch	The area can be zero; a needle can turn around in an arbitrarily small area [P] cited
1971	Davies	The conjecture holds for $n = 2$ [P] cited
1991	Bourgain	Systematically connected Kakeya to Fourier analysis [P] cited
1995	Wolff	Dimension $\geq (n+2)/2$: 2.5 in 3D, 3 in 4D (the "hairbrush" argument) [P] cited
2000	Katz–Łaba–Tao	Box dimension $> 5/2$ in 3D; identified three structural features of extremal configurations: stickiness, planiness, graininess [P] cited
2008	Dvir	Finite-field version completely solved (§14.4) [P] abstract
2019	Katz–Zahl	Hausdorff dimension $> 5/2$ in 3D [P] cited
2022–2025	Wang–Zahl	Sticky Kakeya $\rightarrow$ Assouad dimension $\rightarrow$ 3D completely solved [P] cited
2026-01	Guth–Wang–Zahl	Streamlined version of the 3D proof, arXiv:2601.14411 [P] cited

"Sticky" means that tubes with nearby directions also sit close together in position, and that this holds at every scale. Wang and Zahl first proved the sticky case and then reduced the general case to it (this is how Guth's Bourbaki lecture notes summarize it [R]).

14.3.3 Wang–Zahl 2025 and the Fields Medal

- arXiv:2502.17655 was posted on 2025-02-24 [P, per the references]; reports say the full paper runs to 127 pages, and Quanta quotes Nets Katz describing it as a once-in-a-century result [R].
- Around 2026-07-23, at the International Congress of Mathematicians in Philadelphia, Wang received the 2026 Fields Medal for her "groundbreaking work in harmonic analysis and

geometric measure theory" [R] (Quanta 2026-07-23; University of Waterloo news page). She is currently a professor at IHES and at NYU's Courant Institute [R].

- The same Quanta article says: "Another obvious target for Wang and her community is the 4D Kakeya conjecture" [R]. **Four dimensions is the publicly stated next target for her and her peers** (§14.9.3).

14.3.4 Four dimensions: stuck at 3.059 before OpenAI

The introduction of the OpenAI four-dimensional paper lists the earlier progress in four dimensions itself [P] (converted to decimals by script):

Result	Lower bound in 4D
Wolff 1995, $(n+2)/2$	3.0000
Guth–Zahl polynomial Wolff axioms + Zahl, Katz–Rogers (compact sets)	$3 + 1/40 = 3.0250$
Katz–Zahl "planebrush", maximal-estimate parameter	3.0490
Borges–Chan–Chen–Liu–Xi–Zhan 2025, maximal parameter $(159+\sqrt{145})/56$	3.0543
Katz–Zahl, Hausdorff dimension	3.059
Sticky sets only: Rai Choudhuri	$13/4 = 3.25$
OpenAI's claim	4

Going from Wolff's 3 toward the target of 4, the roughly 30 years before OpenAI closed only **5.9%** of the gap (computed by script). Separately, Łaba–Tao proved in 2001 that the **box dimension** in four dimensions is strictly greater than 3 [P, cited].

A detail of timing: on 2026-09-18 Hong Wang and Dmitrii Zakharov posted *Kakeya and Furstenberg problems for sticky sets of tubes* on arXiv (arXiv:2609.22035), which gives a full-dimension estimate for the union of sticky tube families satisfying certain conditions. The four-dimensional paper cites it and notes that these conditions "do not cover arbitrary Kakeya sets" [P, per the references and introduction; ⚠ I have not read that paper itself]. The OpenAI four-dimensional manuscript is dated six days later.

14.4 The Kakeya problem over finite fields: a version you can work by hand

14.4.1 The F_3 plane and the smallest Kakeya sets

A **finite field F_q** (q a prime) is just the integers mod q , $\{0, 1, \dots, q-1\}$. In F_3 , $2 + 2 = 1$ and $2 \times 2 = 1$. The **F_q plane** has q^2 points; a line $\{a + t \cdot v : t \in F_q\}$ has exactly q points; and there are $q + 1$ directions in all (slopes $0 \dots q-1$ plus vertical). A **finite-field Kakeya set** is a set of points that contains a whole line in every direction. Wolff proposed it in 1999 as a toy model for the real problem [R].

In the plane there is an elementary lower bound: two lines in different directions meet in exactly one point, so the first line contributes q points, the second adds at least $q-1$ new ones, and so on, for a total $\geq q + (q-1) + \dots + 1 = q(q+1)/2$.

The script enumerated every combination of "one line per direction" (using a translation to fix the vertical line as $x = 0$):

q	Points in the plane	Combinations searched	Smallest Kakeya set	Elementary lower bound	Blokhuis–Mazzocca formula	Fraction
3	9	27	7	6	7	0.778
5	25	3,125	17	15	17	0.680
7	49	823,543	31	28	31	0.633

The enumeration matches the exact formula of Blokhuis–Mazzocca 2008, $q(q+1)/2 + (q-1)/2$ for odd q [R, arXiv:0911.4370; I have not read the original, but all three values of q were verified by enumeration]. A smallest Kakeya set for $q = 3$ (script output):

Points: (0,0) (0,1) (0,2) (1,0) (1,1) (2,0) (2,2)
 Lines: $x = 0$; $y = 0$; $y = x$; $y = 2x + 1$

[Check 4] Verify by hand that the 3 points of each of the four lines are among these 7 points (for $y = 2x + 1$, $x = 0, 1, 2$ give $y = 1, 0, 2 \pmod 3$).

The fraction is falling, toward $1/2$, but **not toward 0**. This is the finite-field version of "full dimension": a Kakeya set has the same order of size as the whole space, q^n .

14.4.2 Dvir's polynomial method (2008)

In 2008 Dvir used a very short argument to prove that a Kakeya set in F_q^n has at least $C_n \cdot q^n$ points; before that, the best general lower bound was about $q^{(4n/7)}$ [P, abstract of arXiv:0803.2336]. Here is the whole argument worked through for $q = 3$, $n = 2$:

Step 1: if there are few points, some polynomial vanishes on all of them. There are 6 monomials in two variables of degree ≤ 2 : $1, y, y^2, x, xy, x^2$. Requiring a polynomial to be 0 at every point of K gives $|K|$ linear equations in the 6 coefficients. When $|K| \leq 5$ there are more unknowns than equations, so a nonzero solution must exist. Script check: for **every one** of the 126 five-point subsets of F_3^2 , such a nonzero polynomial exists.

Step 2: vanishing on a whole line gives away the "top-degree part". Restrict P to the line $a + t \cdot v$ to get a polynomial in one variable t . Its degree is ≤ 2 , yet it is 0 at $t = 0, 1, 2$, so it is identically zero. Its t^2 coefficient is exactly the value at v of the top-degree homogeneous

part of P . Conclusion: **if K contains a whole line in direction v , then the top-degree part of P vanishes at v .**

Step 3: contradiction. If K is a Kakeya set, the top-degree part vanishes in every direction, and hence at all 9 points of F_3^2 . But the script checked all 728 nonzero polynomials of degree ≤ 2 , and **0** of them vanish at all 9 points. So a Kakeya set has at least 6 points.

The general case works the same way: there are $C(q+n-1, n)$ monomials in n variables of degree $\leq q-1$, so $|K| \geq C(q+n-1, n) \geq q^n/n!$.

An example (script output): $S = \{(0,0), (0,1), (0,2), (1,0), (2,0)\}$, the "cross" made of $x = 0$ and $y = 0$. $P = xy$ vanishes on all of it, and its top-degree part vanishes only in the directions $(1,0)$ and $(0,1)$, so S can contain whole lines in at most these two directions, which is exactly the case.

[Check 5] Verify that xy is 0 at all 5 points of the cross, and nonzero in the directions $(1,1)$ and $(1,2)$.

14.4.3 Why the real case is hard, and a migration route from CS to analysis

In F_q a line has only q points, and "a polynomial of degree $< q$ that vanishes on a whole line is identically zero" applies directly. In $\mathbb{R}^n$ there are **infinitely many scales**: tubes can spread out at coarse scales and bunch together at fine scales, or the other way round. Nearly all the difficulty of the real problem lies in these multiple scales, and both the Wang–Zahl and the OpenAI papers spend many pages on them [P]. Over the reals there is also an "algebraic obstruction": the tubes may cluster near a low-degree algebraic surface, and Guth–Zahl's polynomial Wolff axioms exist to rule this out [P, cited].

Dvir is a theoretical computer scientist, and his argument shares its roots with polynomial interpolation in coding theory (Reed–Muller codes, Sudan's list decoding) [established]. Guth then carried it over to the real setting: the endpoint case of multilinear Kakeya in 2010 [P, cited], Guth–Katz on the Erdős distinct distances problem in 2015 [R], and in 2016 the use of "polynomial partitioning" to push three-dimensional restriction to $p > 13/4$ [P, cited]. **A CS counting trick became, a decade or so later, a workhorse of harmonic analysis.**

14.5 The Kakeya maximal conjecture: tubes that are only partly lit

14.5.1 The maximal function

Take a function f on $\mathbb{R}^3$ (think of it as a cloud of "density"). For each direction ω , search all positions in space for the tube with direction ω , length 1 and radius δ that makes the average of $|f|$ over the tube as large as possible:

$$K_\delta f(\omega) = \max \text{ over all positions } a \text{ of } \left[\frac{1}{(\pi\delta^2)} \times \int \text{ over the tube } T_\delta(a, \omega) \text{ of } |f(x)| dx \right]$$

This is the definition at the start of the OpenAI three-dimensional paper [P]. $K_\delta f$ is a function on the sphere: one number per direction.

Keakeya maximal conjecture (three dimensions): for every $\epsilon > 0$ there is a C_ϵ such that for all $0 < \delta < 1$ and all f : $\|K_\delta f\|_{L^3(S^2)} \leq C_\epsilon \cdot \delta^{(-\epsilon)} \cdot \|f\|_{L^3(\mathbb{R}^3)}$.

Here $\|g\|_{L^3} = (\int |g|^3)^{(1/3)}$.

14.5.2 The trivial bound $\delta^{(-2/3)}$

By Hölder's inequality, the average over a single tube is $\leq (\pi\delta^2)^{(-1/3)} \cdot \|f\|_3$; then integrate over the sphere:

$$\|K_\delta f\|_3 \leq (4\pi)^{(1/3)} \cdot (\pi\delta^2)^{(-1/3)} \cdot \|f\|_3 \qquad \text{(grows like } \delta^{(-2/3)})$$

This is exactly the **only** substantive theorem about the maximal function in the Lean file in the OpenAI repository (`maximal_norm_le_holder`, §14.8.1) [P]. Compared with the shape of the conjecture (taking $\epsilon = 0.01$), the script's actual numbers are:

δ	Constant in the trivial bound	Conjecture shape $\delta^{(-0.01)}$
10^{-2}	34.2	1.047
10^{-4}	736.8	1.096
10^{-6}	15,874.0	1.148

14.5.3 λ and λ^3 : why the maximal version is stronger

Maximal $\Rightarrow$ set. Take f to be the indicator function of a union of tubes. In every direction there is a tube lying entirely inside the union, where the average is 1, so the left side is $\approx (4\pi)^{(1/3)}$, while on the right $\|f\|_3 = (\text{volume of the union})^{(1/3)}$. Substituting gives volume of the union $\gtrsim \delta^{(3\epsilon)}$, which is the discrete version of the set conjecture.

Light up only part of each tube. In each tube T only a subset $Y(T)$ is lit, with $|Y(T)| \geq \lambda \cdot |T|$ and $0 < \lambda \leq 1$. Take f to be the indicator function of the union of the lit pieces, and compute the same way:

$$\begin{aligned} (\text{average in each direction} \geq \lambda) &\Rightarrow (4\pi)^{(1/3)} \cdot \lambda \leq C \delta^{(-\epsilon)} \cdot |\text{union of lit pieces}|^{(1/3)} \\ &\Rightarrow |\text{union of lit pieces}| \gtrsim \delta^{(3\epsilon)} \cdot \lambda^3 \end{aligned}$$

This is exactly inequality (1.1) of the three-dimensional paper: $|\cup Y(T)| \geq \delta^\epsilon \cdot \lambda^3 \cdot \Sigma |T|$ [P].

The "3" in λ^3 comes from the L^3 norm.

[Check 6] If the constant in the maximal inequality is $C \cdot \delta^\epsilon(-\epsilon)$ and exactly half of every tube is lit ($\lambda = 1/2$), at least how large is the volume of the union of the lit pieces?

Why the power of λ is the key. The paper says that Wang and Zahl's set theorem gives a dependence of the form $\lambda^{K(\epsilon)}$, and that "the need to replace it by λ^3 is stated explicitly after their theorem" [P]. To compare orders of magnitude, take a **made-up** $K = 100$ ($\triangle$ not the real exponent; computed by script): at $\lambda = 0.5$, $\lambda^3 = 0.125$ while $\lambda^{100} = 7.89 \times 10^{-31}$. In applications f has to be split into layers by the size of its values, and each layer is a family of tubes "lit along a fraction λ ", with λ running through $1/2, 1/4, 1/8, \dots$; a dependence like λ^{100} blows up when the layers are summed, while λ^3 matches L^3 exactly. **The set conjecture can still be derived from a somewhat worse dependence on λ (the paper: "A set-dimension conclusion can follow from weaker dependence on λ "), but the maximal conjecture requires the optimal λ^3 , so it is strictly stronger;** the applications in §14.6 all need the maximal version.

The dual version: a Nikodym set requires that for every point in a large set of points, some segment passes through that point and lies almost entirely in the set. §11 of the three-dimensional paper uses the transfer theorem of Gao–Liu–Xi to derive a Nikodym maximal estimate from the maximal theorem [P].

14.6 Why analysts care: waves, Fourier restriction and Bochner–Riesz

14.6.1 Waves with a single wavelength: extension from the sphere

The "sine waves" of $\mathbb{R}^3$ are the plane waves $e^{(2\pi i x \cdot \xi)}$: their crests form a family of parallel planes perpendicular to ξ , and the wavelength is $1/|\xi|$. Every function is a superposition of plane waves. Now allow only $|\xi| = 1$, that is, waves that **all have wavelength 1 but may point in any direction**, and give each direction ω a complex weight $g(\omega)$:

$$Eg(x) = \int \text{over the unit sphere } S^2 \text{ of } g(\omega) \cdot e^{(2\pi i x \cdot \omega)} d\sigma(\omega)$$

This is the **extension operator** of paper 077 [P]. Physically, it is the field you get by superposing "monochromatic waves of the same frequency arriving from every direction". Where the name comes from: Stein originally asked whether the Fourier transform can be meaningfully **restricted** to the sphere, a set of zero volume, and the sphere's **curvature** makes this possible within a certain range; extension is the dual of restriction [established].

14.6.2 The simplest example and the $p > 3$ threshold

Take $g \equiv 1$. The integral can be done exactly (the paper gives it [P]): $E1(x) = 2 \sin(2\pi|x|) / |x|$. A script checked this by numerical integration: at $|x| = 0.25, 1.3$ and 5.7 it gives $8.000000, 1.463164$ and -0.333704 , matching the formula.

It decays like $1/|x|$. $\int |E1|^p$ over $\mathbb{R}^3$ is roughly $\int r^{-(p)} \cdot r^2 dr$, which is **finite for $p > 3$ and divergent for $p \leq 3$** . Actual numbers from the script ($\int$ over $1 < |x| < R$):

p	R = 10	R = 100	R = 1000
3	97.95	196.19	294.43
3.5	76.66	101.09	108.82
4	66.75	73.53	74.21

At $p = 3$, every tenfold increase in R adds about 98: a logarithmic divergence. At $p = 3.5$ and 4 the integral converges.

Stein restriction conjecture (three dimensions, bounded data): for every $p > 3$ there is a constant C_p such that every bounded g satisfies $\|Eg\|_{L^p(\mathbb{R}^3)} \leq C_p \cdot \sup|g|$.

In other words, "all the waves converging at the origin" is already the worst case. In n dimensions the threshold is $p > 2n/(n-1)$; two dimensions ($p > 4$) was solved in the 1970s (Fefferman, Zygmund) [R]. Progress in three dimensions:

Who	Threshold for p	Source
Tomas–Stein ($L^2 \rightarrow L^4$)	4	[P]
Tao 2003, bilinear	$10/3 \approx 3.3333$	[R] $\triangle$ original not reread
Bourgain–Guth 2011, multilinear	3.3	[R] $\triangle$ original not reread
Guth 2016, polynomial partitioning	$13/4 = 3.25$	[P]
Hong Wang, from 2018, the "brooms" method (paraboloid)	$42/13 \approx 3.2308$	[P]
Wang–Wu 2024, two-ends Furstenberg + refined decoupling	$22/7 \approx 3.1429$	[P]
OpenAI claims	3 (the entire open range)	[P]

(The [P] thresholds all appear directly in the introduction of paper 077.)

14.6.3 Why the restriction conjecture controls Kakeya: wave packets are thin tubes

This is the most important step in the whole chapter:

1. **The uncertainty principle.** Cut the sphere into small caps of diameter about δ . If you use only the directions inside one cap, the superposed wave is concentrated on a tube of width about $1/\delta$ and length about $1/\delta^2$, oriented perpendicular to the cap. This is called a **wave packet** [established]. You have seen the same thing in EEG time-frequency analysis: the narrower in frequency, the wider in time. (V1 simple cells are often modeled with Gabor functions, Daugman 1985, which are also "wave packets" localized in both space and frequency; this is only an analogy.)
2. **One tube per cap.** After shrinking coordinates by a factor of δ^2 , these become tubes of length 1 and width δ whose directions differ pairwise by about δ : exactly the Kakeya configuration of §14.3.1.
3. **Overlap $\Rightarrow$ larger L^p norm.** Give each wave packet a random sign; then $|Eg|^2$ is roughly "the number of tubes passing through the point". If the tubes overlap heavily, the way they do in a Besicovitch set, the energy is squeezed into a very small set, and for $p > 2$ the L^p norm becomes large. The restriction estimate forbids this. **Restriction conjecture $\Rightarrow$ Kakeya conjecture** [established, going back to Bourgain 1991].

The converse fails: Kakeya only counts how tubes overlap, so it is like "the restriction problem with the phases thrown away"; restriction also has to deal with interference between the waves, so it is harder.

14.6.4 A surprise about the ideal low-pass filter: Fefferman 1971 and Bochner–Riesz

The most natural filter in engineering is the "ideal low-pass" filter: keep the frequencies with $|\xi| \leq 1$ and cut off the rest. In one dimension this is bounded on every L^p ($1 < p < \infty$) [established]. **Fefferman proved in 1971 that in two or more dimensions the spherical cutoff is bounded only for $p = 2$** , and the counterexample is built precisely from a Besicovitch set (introduction of paper 078: "Using a Kakeya construction, Fefferman proved..." [P]).

The remedy is to round off the edge of the cutoff: multiply by $(1 - |\xi|^2)^\delta$ (for $|\xi| < 1$). This is called the **Bochner–Riesz mean**. Here δ measures how much the edge is rounded; it has nothing to do with the tube radius and is just historical notation.

Bochner–Riesz conjecture: if $\delta > \max(n \cdot |1/p - 1/2| - 1/2, 0)$, the mean is bounded on L^p [P, paper 078].

Three-dimensional thresholds computed by script: 1.0 at $p = 1$ or ∞ ; 0.5 at $p = 6$; 0.25 at $p = 4$; **0 for $3/2 \leq p \leq 3$** , where any amount of rounding is enough. Two dimensions was solved by Carleson–Sjölin; in three dimensions the previous best was $\max\{p, p'\} \geq 22/7$ (Gao–Wu–Xi) [P, introduction of paper 078].

[Check 7] For $n = 3$ and $p = 4$, what is the threshold $3 \cdot |1/4 - 1/2| - 1/2$?

14.6.5 A chain of implications

In three dimensions these conjectures line up in a chain (each arrow is cited in the introduction of the corresponding OpenAI paper [P]):

Local smoothing (wave equation) $\Rightarrow$ Bochner–Riesz $\Rightarrow$ Fourier restriction $\Rightarrow$ Keakeya maximal $\Rightarrow$ Keakeya set				
	Sogge	Tao 1999	§14.6.3	§14.6.3
5.3				

Local smoothing: solutions of the three-dimensional wave equation lose some smoothness at any fixed time, but averaging over an interval of time recovers almost all of it. This is Sogge's 1991 conjecture [R]; the two-dimensional case was solved by Guth–Wang–Zhang [P, cited].

OpenAI claims to have pushed through this chain **from end to end** in three dimensions within two days (§14.7.2); before this, only the last link had been taken, by Hong Wang and Zahl.

14.7 What OpenAI claims

14.7.1 The two papers of family 074

The 074 entry of `CONTENTS.md` [P]: "Resolves the Keakeya maximal conjecture in three dimensions and the Hausdorff-dimension conjecture in four."

Paper A: *The Keakeya maximal conjecture in three dimensions* (2026-09-23, 97 pages).

Abstract [P]:

"We prove the Keakeya maximal conjecture in three dimensions. For every $\varepsilon > 0$, the maximal average over unit tubes of radius δ maps $L^3(\mathbb{R}^3)$ to $L^3(S^2)$ with norm at most $C_\varepsilon \delta^{(-\varepsilon)}$."

Paper B: *Every four-dimensional Keakeya set has full Hausdorff dimension* (2026-09-24, 175 pages). Abstract [P]:

"We prove the four-dimensional Hausdorff-dimension Keakeya conjecture: every subset of $\mathbb{R}^4$ containing a unit line segment in every direction has Hausdorff dimension four. No compactness or regularity assumption is imposed on the set or its witnessing line family."

It does **not** claim the four-dimensional maximal conjecture. For the corollaries in §11 of each paper, see §14.9.1.

14.7.2 The same cluster: 073, 077, 078, 079

Within the same two days there was a whole cluster of harmonic analysis results that cite one another [P]:

No.	Paper (pages)	Claim	Previous best
073	Falconer distance conjecture, all dimensions (64)	Every compact set in $\mathbb{R}^d$ of dimension $> d/2$ has a distance set of positive length	Plane: 5/4 (Guth–Iosevich–Ou–Wang, pinned version); three dimensions: 9/5
074	The two papers above (97 + 175)	Three-dimensional maximal, four-dimensional set	§14.3.4, §14.5.3
077	Elliptic capacity propagation and restriction to the sphere (103)	Restriction for the sphere with bounded data, $p > 3$	$p > 22/7$
077	Diagonal extension for positively curved surfaces (25)	Any compact positively curved surface, $L^p \rightarrow L^p$, $p > 3$	$p > 22/7$ (Wang–Wu)
078	Three-dimensional Bochner–Riesz (114)	The whole strict range	$\max\{p, p'\} \geq 22/7$
079	Three-dimensional critical local smoothing (165)	Arbitrarily small loss at $p = 3$	At $p = 3$ a loss $> 1/12$ was needed (obtained by interpolating results of Gan–He–Li–Wu)

A few details [P]:

- **The Falconer distance problem** (1985 [R]): if a set has dimension greater than $d/2$, must the distances between pairs of its points fill out an interval of positive length?
- The 077 sphere paper uses Bourgain's factorization theorem to upgrade the bounded-data result to $L^q \rightarrow L^q$ ($3 < q < \infty$), then interpolates to get the whole mixed range except the endpoint line. **What it claims is the entire open range of the Stein restriction conjecture for the sphere in three dimensions.**
- One of the two inputs of the 077 surface paper is the three-dimensional maximal theorem of 074. Paper 078 mentions that Zipeng Wang has also put forward a proof of Bochner–Riesz in all dimensions (arXiv version 6); $\triangle$ I did not check its status.

14.7.3 The idea of the three-dimensional maximal paper (introduction only, simplified)

A paraphrase of §1.2 of the paper; the details go far beyond this [P]:

1. **Proof by contradiction:** define a "critical exponent", the least extra power of N ($N = 1/\delta$) that must be multiplied in for the inequality to hold for every configuration. If it is > 0 , there is a sequence of extremal configurations that "almost attain equality" (compare a minimal counterexample in combinatorics).
2. **Extremal configurations are forced into specific shapes:** multilinear Kakeya (Bennett–Carbery–Tao, and Guth's endpoint version) implies that over short time blocks the

directions are nearly coplanar; the Ren–Wang planar Furstenberg theorem implies that there is enough density inside "plates". These are exactly the planiness and graininess described by Katz–Łaba–Tao. A "minimum-delay" argument then forces the time profile into the canonical form $F(s) = \beta(s - \tau)_+$, which the paper calls stationarity.

3. **Information-theoretic bookkeeping:** each tube is described by four coordinates, position X, Y and velocity U, V , and **conditional entropy** (normalized by $\log N$, which you can read as a cell count on a logarithmic scale) measures how much information each coordinate carries. Two events on the same line give two "reference frames", which differ by a 2×2 matrix; planar projection theorems constrain the entropy growth rates (the "pinning" step is adapted from Shmerkin–Wang and Orponen–Shmerkin–Wang).
4. **The contradiction:** if a constraint failed, there would have to be a positive amount of mutual information between the actual joint distribution and the independent product distribution, but an averaged estimate shows that less than that is available. So the critical exponent is 0. The estimate is then extended to arbitrary L^3 functions, keeping λ^3 throughout.

For CS readers: this is like a lower-bound proof that uses "amount of information" as a potential function (my analogy).

14.7.4 The idea of the four-dimensional paper (again, introduction only)

1. **Proof by contradiction:** assume there is a Kakeya set of Hausdorff dimension < 4 , and turn it into a "weighted line problem": fit the lines with "charts" (a chart is a collection of line pieces plus a time interval, on which a certain **quadratic polynomial test** stays small). A counterexample means the fit must pay a positive "time cost" (horizon $H(E) > 0$) [P]. The quadratic polynomial corresponds to the algebraic obstruction described in §14.4.3.
2. **Three cases according to "narrowness" ℓ** [P]: when $\ell = 0$, an approximate polynomial "potential field" is assembled, and the cubic and quartic terms (the obstruction is directions lying near a conic) are encoded into a quadratic correction (Proposition 6.8); when $0 < \ell < \infty$, the problem is projected to a one-dimensional model, which either improves directly or is forced into a "horizontal model" $Y' = -Z + tv$ in which the two horizontal directions do not commute (this reminds me of the known hard examples of Heisenberg-group type; the paper does not use that name, so this is speculation); when $\ell = \infty$, the final projection gives a time cost tending to 0, which contradicts the counterexample.
3. **Inputs:** Lemma 2.3 of OpenAI's own three-dimensional maximal paper (a weighted "full-time plank" bound, coming from Guth–Wang–Zahl); the determinant-type multilinear Kakeya of Carbery–Valdimarsson; the restricted-triples dot-product theorem of Wang–Zahl; the multiplicative convolution estimate of Orponen–de Saxcé–Shmerkin; Balog–Szemerédi–Gowers [P].

Dependency: the four-dimensional paper uses a lemma from the three-dimensional maximal paper. If the three-dimensional paper contains an error, the four-dimensional one is affected too.

14.7.5 What is wrong with "OpenAI helped Hong Wang solve four-dimensional Kakeya"

The claim	What is actually the case
"Helped Hong Wang"	The author line says only "OpenAI"; no public material says Hong Wang was involved [P]
"Hong Wang's problem"	Wang–Zahl solved the three-dimensional set conjecture; what OpenAI claims is the three-dimensional maximal conjecture and the four-dimensional set conjecture, which are the next step
"Already solved"	Only claimed: not peer reviewed, no Lean proof, no public verification by experts (§14.8)
Relation to Hong Wang	Heavy reliance on tools from her and her collaborators: Wang–Zahl (three papers), Guth–Wang–Zahl, Ren–Wang, Wang–Wu, Shmerkin–Wang, Orponen–Shmerkin–Wang, Guth–Wang–Zhang [P, the reference lists of each paper]

The confusion probably comes from Ivanisvili's widely shared remark in §14.8.4: it puts "Kakeya in 3D won a Fields Medal" side by side with "Kakeya in 4D was solved by AI", and in retelling this easily turns into "helped Hong Wang".

What is easy to miss: the Ren–Wang planar Furstenberg theorem is the shared engine of this cluster. The three-dimensional maximal, sphere restriction, Bochner–Riesz and local smoothing papers all list it as a key input [P]. **Hong Wang's tools appear at almost every key point of this batch of AI results.**

14.8 How far it has been verified

14.8.1 Lean: only an unfinished 761-line file

§3 explained that the value of a Lean proof is that "humans only need to audit the statement". In the snapshot `openai/math@fd4aeeb` (2026-10-07) [P]:

- 074, 077, 078 and 079 are **absent** from `lean/formalization.yaml` (which registers 200 Comparator configurations), from `lean/docs/` and from `ComparatorChallenges/`.
- The only relevant file is `lean/OAI/Analysis/Kakeya/FixedScale.lean`: 761 lines, imported on line 106 of `OAI.lean`, with no `sorry`.

It writes down the **statement** of the three-dimensional maximal conjecture:

```
def MaximalConjecture : Prop :=
  ∀ ε : ℝ, 0 < ε → ∃ C : ℝ, 0 < C ∧
  ∀ δ : ℝ, 0 < δ → δ < 1 → ∀ f : Space → ℂ,
```



```

MemLp f (3 : ℝ≥0∞) (volume : Measure Space) →
MemLp (maximal δ f) (3 : ℝ≥0∞) surfaceMeasure ∧
eLpNorm (maximal δ f) (3 : ℝ≥0∞) surfaceMeasure ≤
  ENNReal.ofReal (C * δ ^ (-ε)) *
  eLpNorm f (3 : ℝ≥0∞) (volume : Measure Space)

```

Word-by-word translation:

- `def MaximalConjecture : Prop`: this is a **definition** (of a proposition), **not a theorem**.
- The middle three lines: for every ϵ there is a constant C such that, for all $0 < \delta < 1$ and every complex-valued function f in L^3 ,
- The last two lines: the L^3 norm of the maximal function on the sphere is $\leq C \cdot \delta^{(-\epsilon)}$ times the L^3 norm of f .

In the file, `tube` is a $t \cdot \omega + u$ ($|t| \leq 1/2$, $u \perp \omega$, $\|u\| \leq \delta$), the average is divided by $\pi\delta^2$, and the total mass of the sphere measure is proved to be 4π (`surfaceMeasure_univ`). I checked it clause by clause, and this statement agrees with Theorem 1.1 of the paper (this is my own reading).

But no theorem anywhere in the repository proves `MaximalConjecture`. What the file actually proves is: that the tubes and the maximal function are measurable, the trivial bound `maximal_norm_le_holder` from §14.5.2, a normalization lemma, and a finite weighted sampling lemma (`logarithmic_indexed_sample`; $\triangle$ I did not check which step of the paper it corresponds to). The four-dimensional theorem does not even have a statement.

The statement is written, the proof is not done, and the machine guarantee is zero. Compare §4 and §9: both of those chapters have complete Lean proofs and Comparator configurations.

14.8.2 Comparison within the cluster: only Falconer has been formalized

No.	Lean status [P]
073 Falconer	Yes. <code>docs/073.md</code> links two Comparator challenges: <code>PlanarFalconer.lean</code> and <code>FalconerAllDimensions.lean</code> . The proof is in <code>MeasureTheory/Falconer/</code> and <code>FalconerCompletion/</code> : 2,796 files, about 247,000 lines in total; grep finds no <code>sorry</code> , no added axioms and no <code>native_decide</code> ; both configurations have <code>"enable_nanoda": false</code> . $\triangle$ <code>formalization.yaml</code> registers only the planar one
074 Takeya	No proof, only the 761-line file above
077, 078, 079	None

The Falconer Lean proof contains a special case of a planar Furstenberg-type incidence estimate (`Incidence/PublishedFurstenberg.lean`) [P]; $\triangle$ it looks weaker than Ren–

Wang's sharp version, but I did not check the exact relationship. 074 is not among the 10 families whose reasoning summaries are published in the README, nor on the 10-07 list of retractions and revisions [P].

14.8.3 Several routes corroborate each other, but they are not independent

This cluster has an unusual structure: several conclusions have **two independently claimed proof routes** [P]:

Conclusion	Route 1	Route 2
Three-dimensional Takeya maximal	074, direct proof	077 sphere restriction $\Rightarrow$ Corollary 1.3
Sphere restriction, $p > 3$	077, direct proof	078 Bochner–Riesz $\Rightarrow$ Tao 1999
Three-dimensional Bochner–Riesz	078, direct proof	079 local smoothing $\Rightarrow$ Corollary 11.2

If the routes were truly independent, an error in one would not affect the other. But they **share** the same unreleased model, the same key inputs (Ren–Wang appears in all three main routes, 074, 077 and 078) and the same methodological style. Software engineering has a classic lesson here: Knight and Leveson's 1986 experiment found that multiple versions of a program developed independently failed at highly correlated places [established]. **"Redundant" proofs from the same generator may have correlated failure modes** (speculation).

14.8.4 Human experts

- **Paata Ivanisvili** (UC Irvine): "Takeya in 3D won a Fields Medal. Takeya in 4D was solved by AI. Let that sink in." [R] (X post, via officechai)
- **Rodrigo Porto** (Princeton): "Takeya maximal conjecture solved?!", and "It's a stronger version of the Takeya conjecture. Hong Wang literally won a fields medal for making progress on the latter" [R] (X post, via officechai)
- **Levent Alpöge** (Anthropic): "There are some sad stories related to their users getting scooped / conflicts of interest (takeya maximal comes to mind)" [R] (X post, via officechai).
⚠ The post does not say who was involved or what happened, and I cannot verify it.
- ⚠ **As of 2026-10-09, I found no public comment on these papers from Hong Wang, Zahl, Guth or Tao**, and nobody has publicly said they have read and verified these 97 + 175 pages (I searched both Chinese and English coverage). The three items above are reactions, not reviews.

Remaining risk: the statement risk is low (the three-dimensional maximal conjecture is a standard statement, and the Lean version agrees with the paper; the four-dimensional one is a single sentence). The proof risk is high (272 pages, no machine check, no human

review, and the four-dimensional paper depends on the three-dimensional one). In arguments of this kind the most error-prone points are the order in which limits are taken and whether "losses smaller than any power" get amplified; the end of §1.3 of the three-dimensional paper explains the order of limits specifically [P], and this is where experts should look first (speculation).

14.9 What it would mean if true (and what it would not)

14.9.1 For mathematics

1. **Four-dimensional Kakeya: from 3.059 straight to 4**, whereas the previous 30 or so years had closed only 5.9% of the gap. Side results [P]: the box dimension is also 4; Kakeya sets in $\mathbb{R}^d$ ($d \geq 4$) all have dimension ≥ 4 ; Nikodym sets on four-dimensional manifolds of constant curvature have dimension 4 (the four-dimensional case of Conjecture 1.4 of Gao–Liu–Xi); a dimension formula for direction sets (Keleti–Máthé: a set containing a line segment in every direction in D has dimension $\geq 1 + \dim D$).
2. **Three-dimensional Kakeya maximal**: the optimal λ^3 dependence. It gives Nikodym maximal estimates on $\mathbb{R}^3$ and on three-dimensional manifolds of constant curvature, and a class of curved Kakeya maximal estimates [P], and it is an input to 077's restriction theorem for positively curved surfaces.
3. **If the whole cluster is true**: in three dimensions, the entire chain local smoothing $\Rightarrow$ Bochner–Riesz $\Rightarrow$ restriction $\Rightarrow$ Kakeya holds in the strict range. These have been central problems of harmonic analysis for the past 50 years; closing them all at once in three dimensions would carry the weight of a whole generation's research program (speculation).
4. **What remains open**: the Kakeya set conjecture in five or more dimensions; the maximal, restriction, Bochner–Riesz and local smoothing conjectures in four or more dimensions; the endpoint cases in three dimensions (every paper says explicitly that it does not handle endpoints) [P].

14.9.2 For CS

- **Kakeya over the reals has no direct impact on CS** (speculation). CS uses the finite field version, which Dvir solved in 2008.
- The finite field version is used in **randomness extraction** (turning a "biased" random source into nearly uniform bits). Dvir–Wigderson 2008 (*Kakeya sets, new mergers and old extractors*) used Kakeya-type bounds to analyze "mergers": among several random variables, one is uniform but you don't know which; output a point on the line through them, and the Kakeya bound guarantees that the output still has high entropy [R]. Dvir–Kopparty–Saraf–Sudan 2009 improved the lower bound to $q^n/2^n$ and used it to build extractors with logarithmic seed length [P, abstract of arXiv:0901.2529].

- For the opposite direction, see §14.4.3: the polynomial method came from CS into real analysis.

14.9.3 For AI

- **Speed:** from the Wang–Zahl preprint (2025-02-24) to OpenAI's four-dimensional manuscript (2026-09-24): **19.0 months**; from the streamlined Guth–Wang–Zahl version (2026-01-20) to the four-dimensional manuscript: **8.1 months**; from the Fields Medal to OpenAI's public release: **75 days** (computed by script).
- **"Scooped":** Quanta called four dimensions the obvious next target for Hong Wang and her peers [R]; Wang–Zakharov's paper on sticky families of tubes came out six days before the four-dimensional manuscript [P]; Alpöge pointed to scooping and conflicts of interest around "akeya maximal" [R] $\triangle$. A 175-page manuscript that nobody has reviewed and that has no Lean proof may already have changed the research topics of a number of PhD students and postdocs: if it is wrong, the field must spend people's time finding the error; if it is right, research programs must be rewritten. This is what §23.9, "Political economy of AI-produced knowledge", looks like in practice: control lies in when to publish, what to publish, and whether to attach a proof. The people who can read these papers are essentially the few dozen cited authors and their students, so **the cost of verification is borne by the community that was scooped** (speculation).
- **A methodological observation:** this cluster shares the engine "extremal configurations + entropy bookkeeping + Ren–Wang". That may mean the model found a reusable engine, or it may mean the errors are correlated (§14.8.3) (speculation).

14.9.4 What it does not mean

- **It does not mean Hong Wang's work has been surpassed or devalued:** these proofs are built on her tools, and the three-dimensional set conjecture was proved by her and Zahl.
- **It does not mean anything has been proved:** as of 2026-10-09 there is no Lean proof and no public review; and the higher-dimensional problems remain open (§14.9.1, item 4).

14.10 Self-test

Try answering these after reading. If you can't, go back and reread the corresponding section.

1. A needle can be turned around within a region of arbitrarily small area, so why is the Kakeya conjecture still meaningful? (§14.1.2, §14.2)
2. The box dimension of $\{1/n\}$ is $1/2$, while its Hausdorff dimension is 0. Where does the difference come from? (§14.2.3)
3. Restate Dvir's polynomial method in three steps. (§14.4.2)

4. In what way is the maximal conjecture stronger than the set conjecture? Where does the 3 in λ^3 come from? (§14.5.3)
5. Why does a conjecture in Fourier analysis imply a geometric conjecture about thin tubes? (§14.6.3)
6. Write down the chain of implications in three dimensions, and mark which link had already been proved before OpenAI. (§14.6.5)
7. What does family 074 have in Lean, and what is missing? How does it compare with 073? (§14.8.1–14.8.2)
8. In which ways is "OpenAI helped Hong Wang solve the four-dimensional Kakeya problem" inaccurate? (§14.7.5)

Appendix: answers to the checks

- **Check 1:** side length = $2/\sqrt{3} \approx 1.1547$; area = $\frac{1}{2} \times \text{side} \times \text{height} = 1/\sqrt{3} \approx 0.5774$ (checked by script).
- **Check 2:** $\log 2^k / \log 3^k = k \cdot \log 2 / (k \cdot \log 3) = \log 2 / \log 3 \approx 0.6309$, which does not depend on k .
- **Check 3:** $\sum_k (\varepsilon \cdot 2^{-k})^s = \varepsilon^s \cdot \sum_k 2^{-ks} = \varepsilon^s / (2^s - 1)$. For $s > 0$ the denominator is a positive constant, so the whole expression $\rightarrow 0$ as $\varepsilon \rightarrow 0$, and the Hausdorff dimension is 0.
- **Check 4:** $x = 0 \rightarrow (0,0), (0,1), (0,2)$; $y = 0 \rightarrow (0,0), (1,0), (2,0)$; $y = x \rightarrow (0,0), (1,1), (2,2)$; $y = 2x + 1 \rightarrow (0,1), (1,0), (2,2)$. All of these are among the 7 points, and the directions of the four lines are exactly vertical and slopes 0, 1 and 2.
- **Check 5:** every point of the cross has either $x = 0$ or $y = 0$, so $xy = 0$. Direction (1,1): $1 \cdot 1 = 1 \neq 0$; direction (1,2): $1 \cdot 2 = 2 \neq 0$.
- **Check 6:** $(4\pi)^{1/3} \cdot \lambda \leq C \delta^{(-\varepsilon)} \cdot V^{1/3}$, so $V \geq 4\pi \cdot \lambda^3 \cdot \delta^{(3\varepsilon)} / C^3 = (\pi/2) \cdot \delta^{(3\varepsilon)} / C^3$.
- **Check 7:** $3 \times 0.25 - 0.5 = 0.25$ (checked by script).

Appendix: Glossary

Term	One-line explanation
Kakeya set (Besicovitch set)	A set that contains a unit line segment in every direction
Box dimension (Minkowski dimension)	Cover with boxes of equal size; the d in $N(\delta) \approx \delta^{(-d)}$
Hausdorff dimension	Covers with pieces of different sizes are allowed; the smallest s for which the "s-dimensional cost" can be pushed down to 0
Kakeya set conjecture	Every Kakeya set in $\mathbb{R}^n$ has dimension n

Term	One-line explanation
Keakeya maximal conjecture	The L^n norm of the largest tube average in each direction is only a factor $\delta^{(-\varepsilon)}$ larger than the norm of f
Shading and λ	Each tube "lights up" only a portion that makes up a fraction λ of it
Fourier extension operator	Builds a field in space by superposing monochromatic plane waves arriving from every direction on the sphere
Wave packet	A wave whose frequencies are concentrated in a small cap; in space it is concentrated on a tube
Stein restriction conjecture	The extension operator is bounded for $p > 2n/(n-1)$
Bochner–Riesz mean	Rounds off the edge of the spherical cutoff, $(1 - \xi ^2)^\delta$
Local smoothing	Solutions of the wave equation recover smoothness after averaging in time
Planar Furstenberg estimate	A lower bound for counting point–tube incidences in the plane; Ren–Wang's sharp version is a shared input of this cluster

Appendix: Further reading

- **Introductions:** Quanta 2025-03-14: <https://www.quantamagazine.org/once-in-a-century-proof-settles-maths-keakeya-conjecture-20250314/> ; Quanta 2026-07-23 (Fields Medal): <https://www.quantamagazine.org/hong-wang-wins-2026-fields-medal-the-third-woman-ever-20260723/>
- **Surveys:** Joshua Zahl, *A survey of the Keakeya conjecture, 2000–2025*: <https://arxiv.org/abs/2512.09397> ; Larry Guth, Séminaire Bourbaki lecture notes 1251 (2026-03): <https://www.bourbaki.fr/TEXTES/Exp1251-Guth.pdf>
- **The restriction problem:** Terence Tao, *Some recent progress on the restriction conjecture* (2003): <https://arxiv.org/abs/math/0303136> (the standard exposition of the wave packet argument in §14.6.3)
- **The three-dimensional proof:** Wang–Zahl <https://arxiv.org/abs/2502.17655> ; Guth–Wang–Zahl <https://arxiv.org/abs/2601.14411>
- **Finite fields:** Dvir <https://arxiv.org/abs/0803.2336> (the argument is very short; I recommend reading it yourself); Dvir–Kopparty–Saraf–Sudan <https://arxiv.org/abs/0901.2529> ; Blokhuis–Mazzocca <https://arxiv.org/abs/0911.4370> ; Guth, *Polynomial Methods in Combinatorics* (AMS, 2016)
- **OpenAI's original papers:** three-dimensional maximal <https://github.com/openai/math/blob/main/preprints/The-Keakeya-maximal-conjecture-in-three-dimensions-September-23-2026/paper.pdf> ; four-dimensional <https://github.com/openai/math/blob/main/preprints/Every-four-dimensional-Keakeya-set-has-full-Hausdorff-dimension-September-24-2026/paper.pdf> (read Section 1 of each first);

for the rest of the cluster, 073/077/078/079, see the corresponding entries in <https://github.com/openai/math/blob/main/CONTENTS.md>

- **Lean**: <https://github.com/openai/math/blob/main/lean/OAI/Analysis/Kekeya/FixedScale.lean>
- **Reactions** (source of the quotes in §14.8.4): <https://officechai.com/ai/mathematicians-react-with-shock-and-wonder-after-openai-releases-over-300-math-proofs-at-once/>


CAUTIONARY CASE · FAMILY 032 AND THE OCT 7 WITHDRAWALS

15. One sign error, three papers withdrawn: the risk of no formalization

The Hodge conjecture (a Millennium Prize problem), minimally: on a smooth projective complex algebraic variety some "topological holes" can be "filled" by subvarieties defined by algebraic equations; the conjecture says every hole satisfying a certain linear-algebra condition can be filled this way. The case $p = 1$ has been known since 1924 (Lefschetz).

OpenAI claims the rational Hodge conjecture for all **complex-multiplication (CM) abelian varieties** (53 pages) — a special class, not the Millennium problem itself; like the ζ zero-free region it was **not** produced by the fixed procedure, and it has **no Lean**. [P]

Oct 7 withdrawals: one sign error propagates through the dependency graph



Sources: [openai/math history.md](#) (2026-10-07); the family title is compared against [CONTENTS.md](#) at the initial commit `adc7f12`. The withdrawal notice reads: "This withdrawal concerns the proof; it does not assert that the mathematical statement is false."

15.1 What happened [P]

history.md (10-07): "In 'Algebraicity of Weil classes on split abelian eightfolds' a sign error invalidates a stabilization-trace cancellation argument and the construction used by two dependent papers." A sign that should have been -1 was taken as +1; a signed count went

from the claimed 0 to $-2m \neq 0$, so the hypothesis of a key theorem failed, and two dependent K3-surface papers were withdrawn with it. The CM paper does not cite the withdrawn paper and is not on this chain.

15.2 Three lessons

1. **Unformalized $\approx$ unverified:** only ~42% of top-line results have Lean. This error was a single sign — missed by human eyes and the model's self-checks, and Lean would certainly have caught it.
2. **Errors propagate through the dependency graph:** one lemma fails, three papers fall, 13 downstream papers need new citations — the same structure as error compounding in long agent reasoning chains.
3. **Title drift:** family 032 was first titled "Hodge and Kuga–Satake results for all projective K3 surfaces" and only downgraded to the CM version after the withdrawals. When reading bulk AI output, trust the latest commit and history.md, not launch-day news.

16. The 11 results side by side

RESULT	PREVIOUS STATE	CLAIMED JUMP	STATEMENT FAITHFUL?	CONFIRMED BY HUMAN EXPERTS
Quasi-RH 7/8	no fixed-width zero-free strip at all	qualitative	yes (standard Mathlib definitions)	none
$\mu(\pi) = 2$	7.103	straight to the end	yes (not in yam1 $\triangle$)	none
$\omega \leq 9/4$	2.371177	-0.12, past a known barrier	yes	none
UGC	"half" (2-to-2, 2018)	complete	yes (stronger form)	none
$L = BPL$	near $L^{(3/2)}$ for 25 years	complete	yes (definitions locked)	none
Plane $\chi \geq 6$	$\chi \geq 5$ (2018)	+1, new method	yes (two-line statement)	none
Thompson F	wrong proofs in both directions	resolved	yes	none
NS blowup	forced Euler precursors	alternatives (C)/(D)	structure matches	Clay has not accepted; disputed
Vlasov–Maxwell	40 years of restricted results	complete, one species	yes	none
Mézard–Parisi	upper bound only	equality, even-arity class	yes (narrower than title)	none
4D Kakeya / 3D maximal	4D dimension bound about 3.059; the 3D set case was only solved in 2025	straight to 4; a stronger 3D version	cannot be checked (no Lean)	none

16.1 Five patterns

- Short statements, long proofs:** plane coloring has a two-line statement and a 30k-line proof; quasi-RH a one-line statement and a 487k-line proof. This is the ideal "humans review statements, machines check proofs" split — and it makes human reviewing bandwidth the new bottleneck.
- Unusually short papers:** 13 pages for 9/4 matrix multiplication and for Thompson's group. Either the methods really are surprisingly lean, or much of the detail lives in Lean rather than in prose for humans — Gómez-Serrano said the NS paper is "not written for humans."

3. **Borrowed tools across fields:** Thompson's group via Lipschitz maps from infinite-dimensional geometry; plane coloring via ergodic theory; matrix multiplication via the asymptotic spectrum rather than the dominant laser method of 40 years. This may be a real AI advantage: no human-expert "field inertia."
 4. **The gap between "satisfying the statement" and "answering the real question":** NS with freely chosen forcing; 376 putting computation in the forcing; the plane-coloring footnote about a formalized false hypothesis; Mézard–Parisi covering only even arity. Each is "spec vs intent."
 5. **Selection effects:** results picked from 4,000 problems vs 3% on pre-registered problems (§2.3). Until OpenAI publishes its failure list, "what is this model's real ability on open problems?" cannot be answered.
-

AI AND AI SAFETY PERSPECTIVE

17. What this means for AI

17.1 Better at math $\neq$ smarter across the board – but a leading indicator

Mathematics has a unique advantage for AI: **answers can be machine-checked** (Lean), which lets reinforcement learning iterate massively — try, get feedback, improve. Chollet's question goes to the heart of it: does progress in these "verifiable domains" transfer to unverifiable ones? A reasonable reading: **any intellectual work that can be automatically checked will fall at a similar pace next**; mathematics is just the first.

17.2 The trust bottleneck moves from "proof" to "specification"

Once there is a Lean proof, the question shifts from "is the proof right?" to "is the statement encoded right?" This is isomorphic to a core problem in alignment research: a system can perfectly optimize the objective you wrote down, and what you wrote down may not be what you want. The line-by-line reading of Lean statements done throughout this deck is a form of **spec review** — an increasingly important and scarce human job.

17.3 Knowledge that can be verified but not understood

In the words of the IAS advisory group: *"It is now the case that AI can output mathematical arguments in situations without the human who prompted it being able to understand the arguments, verify them, or take responsibility for them."* [R] Mathematics is where this situation appears first and most cleanly (because verification is perfect). Generalized to science, humans may move from "understanders" to "checkers." Tao's Math 2.0 — putting

exposition, community building and opening new directions on a par with problem-solving — is at bottom an answer to "what do humans do in this new division of labor?"

17.4 Two sides for safety research

Good news

- **Verifying is easier than discovering:** if future AI can give formal proofs about its own behavior ("this action will not do X"), humans need not understand all of its capabilities, only check the proofs — the core idea behind the Guaranteed Safe AI agenda, which becomes more feasible as mathematical AI improves.
- This Lean library is itself the first real sample of "large-scale AI output + machine verification," and its failure modes can be studied.

Bad news

- The hardest parts of alignment ("what does this AI really want?", "is the outcome good for humans?") cannot be written in Lean. Verifiable domains racing ahead while unverifiable ones lag is itself a risk.
- The verifiable parts of AI R&D (kernel optimization, proofs, benchmarks) will be the first to be automated at scale — the "AI improves AI" loop is most likely to start there.

17.5 Governance: a precedent of unilateral dumping

The community had begun forming norms through the IAS advisory group (disclose model name, prompts, reasoning chains, compute time and cost, how problems were chosen, failure rates on comparable problems). OpenAI followed some of them, then released 700+ papers at once. AHM's phrase "a demonstration of power" is pointed, but it names a structural problem: **a frontier company unilaterally decides how a capability enters the public domain, and the affected community can only react afterwards.**

PART II · OUTLOOK · CALIBRATION

18. First, calibrate: how math breakthroughs turn into real-world impact, and how fast AI capabilities will move

Part I covered what has already happened. Part II asks a bigger question: **if AI solves math at scale (it is already proving open problems), or if AI's other capabilities also break through, which problems that really matter will open up? What are the deeper consequences?** §19–§23 answer this field by field, and §24 pulls it together. This section is the foundation. Every later section is first checked against four things here: how long math has historically taken to turn into real-world impact, and where it got stuck (18.1); when math is actually the bottleneck (18.2); the real track record of AI for science so far (18.3);

and how fast AI capabilities are moving on trend, and how unevenly (18.4–18.5). Tags: [established] accepted historical or research findings; [recent] new results from 2024–2026; [speculation] inference. 18.4–18.5 keep the deck's original tags: [data] has a source; [model] computed from explicit assumptions; [judgment] inference. $\triangle$ = could not be verified, or sources conflict.

18.0 The bottom line

- **The lag from math to real-world impact falls into two classes, one to two orders of magnitude apart** [established]. "Stockpile" math, done for math's own sake (number theory, Riemannian geometry, topology, the Radon transform), took 50–200 years to find an application; "made-to-order" math, done for an engineering problem (linear programming, information theory, Black–Scholes, compressed sensing, polar codes), took only 2–15 years. The vast majority of OpenAI's 719 papers are stockpile math, so on the historical record their real-world impact will be measured in decades [speculation].
- **When an application arrives, the math is almost always already there** [established]. Across 14 historical cases, what held things up at the end was compute and enabling technology; demand; institutions (auction legislation, the legality of organ exchange, Soviet ideology); materials and experiments (topological qubits have been stuck for nearly 30 years); and "the people who needed it didn't know it existed" (Cormack reinvented the Radon transform; LDPC codes were forgotten for 34 years).
- **In most fields, math is only a small segment of the pipeline.** Apply Amdahl's law to a hypothetical drug project: make "theory" infinitely fast and the overall speed rises by only about 5%. AI's real leverage is raising the hit rate of the slow steps and replacing experiments with computation. On the point that "pure cognitive speedup cannot carry a large impact on its own", the optimists (Amodei, Epoch) and the skeptics (Narayanan & Kapoor, Acemoglu) actually agree.
- **As of 2026-10, every proven speedup in AI for science has come from AI's "other capabilities", not from proof** [recent]. Weather forecasting has entered the operational systems of ECMWF and NOAA and passed independent verification; protein structure prediction is used by more than 3 million researchers. The three flagship materials-discovery results (GNoME, A-Lab, MatterGen) have all been heavily walked back; the most-cited MIT paper on "AI speeds up R&D" was withdrawn; and no "AI-discovered" drug has been approved.
- **AI capability curves are exponential, and intuition systematically underestimates them.** The length of tasks AI can complete on its own doubles about every 129 days (about $\times 7$ per year): about 15 hours in 2026-05, and about 400 hours by end-2027 in the central scenario [data][model]. The single-attempt solve rate on preregistered open problems is 3% today; in the central scenario it reaches about 18% by end-2027 and about 45% by end-2028.

Formal proofs cost about \$0.023/line; at ± 10 per year, the cost of generating proofs becomes negligible in 2027–2028 [model]. The biggest uncertainty is the gap between 98% on benchmarks and 3% on open problems.

- **The speedup is extremely uneven.** Type A problems, which are "clearly stated and only missing a proof", have an automatic judge and will be solved by AI in bulk; type B problems, where "it isn't even clear what to prove", move slowly. Generalizing one step: how fast AI speeds up a field depends on how fast that field's "judge" gives feedback. For Lean it is minutes, for materials synthesis weeks to months, for climate a decade [judgment].
- **Once proofs are cheap, the bottleneck moves to four places** [speculation, with precedent]: whether the statement is faithful (the spec); human digestion and understanding (in Tao's framing, a move from proof scarcity to proof abundance); experiments (theory gets cheaper, so experiments get relatively more expensive); and matching (connecting existing math to the problems that need it). Sequencing costs fell 10,000-fold in 10 years, yet biology did not get 10,000 times faster; proofs will probably go the same way.
- **AI math helps most in fields where "the math is the product":** theoretical computer science and cryptography, formal verification, coding theory, and rigorous bounds in parts of physics. In biomedicine, materials and the social sciences, AI's other capabilities matter more, but they are just as limited by experiments and institutions (18.7). This is also where the topic meets AI safety: proving a quietly weakened statement is reward hacking in mathematics, and the ability to audit statements and specs is the next shared bottleneck for math, science and AI oversight.

18.1 History: the lag from math to application

How long did past math breakthroughs take to go from "proved" to "changing the world", and where did they get stuck? If what held them up was not the math, then AI making math faster can only shorten a small part of the total time.

The table below covers 14 cases [established]. "Math in place" means the key result was first published; "first use" means the first real application; "widespread" means large-scale deployment; years are approximate. After the table, 6 cases are covered in detail and the other 8 get one sentence each.

#	Math → application	Math in place	First use	Widespread	Math→first use	First use→widespre
1	Number theory → public-key cryptography	1763 (Euler)	1977 (RSA)	1995–2017 (SSL→HTTPS)	~210 years	~20–40 year

#	Math → application	Math in place	First use	Widespread	Math→first use	First use→widespre
2	Riemannian geometry → general relativity → GPS	1854	1915 (GR)/1978 (GPS)	1995–2000	61 years (to physics)/~125 years (to technology)	~20 years
3	Lie groups/representation theory → quark model	1870s–1894	1961–64	1970s (Standard Model)	~70–90 years	~10 years
4	Fourier → FFT/MRI/JPEG	1807–1822	1965 (FFT)/1973 (MRI)	1980s–1992	~145–165 years	~15–25 year
5	Hilbert space → quantum mechanics → transistor	1904–1910	1925–32 (QM)	1947–1960 (transistor, laser)	~20 years	~20–35 year
6	Boolean algebra/computability → computers	1854 / 1936	1937 (Shannon)/1948	1950s–60s	~83 years / ~12 years	~10–20 year
7	Information theory/coding → communications	1948; LDPC 1962; polar codes 2009	1950–1982; LDPC 2000s; polar codes 2016	3G/4G/5G	Hamming 2 years; RS→CD 22 years; LDPC ~40 years; polar codes ~7 years	~5–10 years
8	Linear programming → logistics/oil refining	1939/1947	~1950	1950s–60s	US <5 years; USSR ~20 years	~10 years
9	Radon transform → CT	1917	1971	late 1970s	54 years	~5–10 years
10	Mechanism design → auctions/matching/kidney exchange	1944–1962	1994 / 1998 / 2004	2000s	~30–40 years	~5–10 years
11	Stochastic analysis → Black–Scholes	1944–51 (Itô)	1973	~1975–80	~25 years	<5 years
12	Topology → topological phases of matter → topological qubits	1895–1946	1982 (physics)	Not yet (technology)	~40–85 years (to physics)	Not delivered so far (>40 years)
13	Compressed sensing → fast MRI	2004–2006	2017 (FDA)	From 2017	~11 years	Ongoing

#	Math → application	Math in place	First use	Widespread	Math→first use	First use→widespre
14	Lattices → post-quantum cryptography	1996–2005	2024 (FIPS 203)	2025–26	~20–28 years	~2 years

Case 1: Number theory → public-key cryptography. The core tool of number theory is "modular arithmetic": you look only at the remainder after dividing by n , going round and round like a clock. Example: $3^4 = 81$, and 81 divided by 7 leaves 4, written $3^4 \equiv 4 \pmod{7}$. Euler's theorem (1763) shows that powers mod n periodically come back to 1; RSA is secure because, given $n = p \times q$ but not p and q , this period is hard to compute. Timeline: Euler 1763 → Diffie & Hellman propose the idea of public keys in 1976 → RSA 1977 (Cocks at Britain's GCHQ had secretly come up with it in 1973; this was declassified only in 1997) → Netscape's SSL in 1995 → HTTPS becomes the majority of page loads in the mid-to-late 2010s $\triangle$. What was missing in between: computers that could raise large numbers to powers quickly, demand from networks and e-commerce, and the RSA patent (expired 2000) and US export controls on cryptography. Lesson: in 1940 Hardy declared of number theory and relativity that "no one has yet discovered any warlike purpose to be served by the theory of numbers or relativity, and it seems very unlikely that anyone will do so for many years"; a little over thirty years later, number theory became the foundation of public-key cryptography. The math sat in the stockpile for 200 years, and what activated it was a new engineering need: secrecy between strangers on a network.

Sources: Diffie & Hellman 1976, doi:10.1109/TIT.1976.1055638; Rivest, Shamir & Adleman 1978, doi:10.1145/359340.359342; Hardy quote https://en.wikipedia.org/wiki/A_Mathematician%27s_Apology

Case 2: Riemannian geometry → general relativity → GPS. Riemannian geometry studies how to measure distance in curved space: on a plane you use Pythagoras' theorem, while in curved space every point has its own "ruler" (the metric tensor). Riemann 1854 → Einstein's general relativity in 1915 (in 1912–15 he spent about three years learning this math, with help from the mathematician Grossmann) → first GPS satellite 1978, fully operational 1995 → in 2000 the US turned off "Selective Availability" (a switch that deliberately degraded civilian accuracy). Why GPS needs relativity: because gravity is weaker up there, satellite clocks run fast by about 45 microseconds a day; because the satellites move fast, their clocks run slow by about 7 microseconds a day; the net effect is about 38 microseconds fast per day. The speed of light $\times$ 38 microseconds $\approx$ 11 km, so without correction the positioning error would build up by about 10 km a day. What was missing in between: atomic clocks, spaceflight, Cold War military budgets and one policy decision. Lesson: 61

years from math to physics, then another 60–80 years from physics to technology; Einstein's three years were a bottleneck of "learning the math", not of "inventing the math".

Sources: Ashby 2003, "Relativity in the Global Positioning System", doi:10.12942/lrr-2003-1

Case 7: Information theory and error-correcting codes → 5G. In 1948 Shannon proved that every noisy channel has a capacity limit C : at rates below C , codes **exist** that make the error rate arbitrarily small, but the proof did not say how to build them. A small example: repeat each bit 3 times (0→000); on receiving 010, a majority vote decodes it as 0. This corrects 1 error, but the rate drops to 1/3, far from capacity. What came next: Hamming codes in 1950; Reed–Solomon codes in 1960, used in CDs in 1982; Gallager's 1962 LDPC codes, which could not be decoded with the computers of the time, were **forgotten for 34 years**, rediscovered by MacKay & Neal in 1996, and chosen by 3GPP in 2016 as the coding for the 5G data channel; Arikan's 2009 polar codes (the first codes proven to reach capacity that can also be constructed efficiently) were chosen in 2016 for the 5G control channel, after only about 7 years. What was missing in between: cheap decoding chips and standards bodies. Lesson: within the same field the lag shrank from over 40 years to 7, because industry had a clear metric (how many dB short of the Shannon limit), a standardization window, and could implement it purely in chips. **When both the problem and the acceptance criteria are clear, translation is fastest.**

Sources: Shannon 1948, doi:10.1002/j.1538-7305.1948.tb01338.x; Gallager 1962, doi:10.1109/TIT.1962.1057683; MacKay & Neal 1996, doi:10.1049/el:19961141; Arikan 2009, doi:10.1109/TIT.2009.2021379; <https://www.abiresearch.com/blogs/2022/07/20/5g-nr-coding-schemes-a-new-start-with-a-long-history/>

Case 8: Linear programming. Maximize a linear objective subject to a set of linear inequality constraints: for example, allocate limited wood and labor hours to making tables and chairs so as to maximize profit. The optimum always lies at some vertex of the polygon enclosed by the constraints, and the simplex method climbs upward from vertex to vertex. Kantorovich proposed linear programming in the USSR in 1939 to schedule production at a plywood factory, and it was **neglected in the USSR for about twenty years** (it was suspected of conflicting with the Marxist theory of value); Dantzig proposed the simplex method at the US Air Force in 1947, and within a few years it was being used for Air Force scheduling and oil-refinery blending. In 1975 Kantorovich and Koopmans won the Nobel Prize in economics. Lesson: the same math, under two systems, translated at speeds ten to twenty years apart. Made-to-order math translates fastest; institutions can leave ready-made math sitting idle.

Sources: Dantzig 2002, *Operations Research* 50(1):42, doi:10.1287/opre.50.1.42.17798; <https://www.nobelprize.org/prizes/economic-sciences/1975/summary/>

Case 9: Radon transform → CT. If you know the "total density" (the line integral) along every straight line through an object, in every direction, can you work out the density at every point inside? Radon proved in 1917 that you can, and gave an inversion formula. Analogy: just from how dark a jelly's shadows are from every angle, you can reconstruct where the fruit inside it sits. In 1963–64 Cormack re-derived a similar result **without knowing about Radon's work**, and almost nobody noticed; Hounsfield (at the company EMI) performed the first clinical CT scan in 1971; the two shared the Nobel Prize in 1979. Cormack only found Radon's paper in the early 1970s. What was missing in between: computers, X-ray detectors, and a company willing to pay. Lesson: 54 years, and the math was reinvented along the way. This is the classic case of "matching cost": the people who needed it didn't know it already existed.

Sources: Cormack 1963, doi:10.1063/1.1729798;
<https://www.nobelprize.org/prizes/medicine/1979/cormack/lecture/>

Case 10: Mechanism design → spectrum auctions, residency matching, kidney exchange. Mechanism design is "reverse game theory": first fix the outcome you want, then design rules so that when everyone acts in their own interest, that outcome is exactly what happens. One example is Gale–Shapley's (1962) "deferred acceptance" algorithm: one side takes turns "proposing" to the other and moves on to the next choice when rejected. One can prove the final pairing is "stable": there are no two people who would both rather leave their partners and pair with each other. The US residency match (NRMP) **had already arrived at an essentially identical algorithm through practice in 1952**, 10 years ahead of the mathematicians. Spectrum auctions had to wait for Congress to authorize them by law in 1993; the FCC held its first auction in July 1994, with rules designed by Milgrom, Wilson and others (Nobel Prize 2020). The NRMP redesigned by Roth & Peranson went live in 1998. Kidney paired exchange was proposed in 2004 by Roth, Sönmez and Ünver, who pushed it into practice; in 2007 Congress passed the Charlie W. Norwood Living Organ Donation Act (Pub. L. 110–144), clarifying that paired exchange does not violate the law against buying and selling organs. Lesson: 30–40 years, held up by law, institutions and trust between hospitals, not by math; practice sometimes even runs ahead of the math.

Sources: Gale & Shapley 1962, doi:10.2307/2312726; Roth, Sönmez & Ünver 2004, doi:10.1162/0033553041382157; Roth & Peranson 1999, doi:10.1257/aer.89.4.748;
<https://www.nobelprize.org/prizes/economic-sciences/2020/summary/>

The other 8 cases, one sentence each

- **3 Group theory → quark model:** In 1961 Gell-Mann and Ne'eman classified hadrons using representations of $SU(3)$ and predicted a missing particle, the Ω^- , which was found in 1964.

The "application" here is scientific understanding, not a product.

- **4 Fourier → FFT:** Cooley & Tukey's 1965 FFT cut the computation from n^2 to $n \log n$ (about 100 times faster at $n = 1024$); it later turned out that Gauss had written it down in 1805. For more than 150 years the bottleneck was "being able to compute it": one algorithmic improvement mattered more than any new theorem.
- **5 Hilbert space → transistor:** The math came only 15–25 years earlier; going from physics to technology took another 20–35 years, held up by materials and manufacturing processes.
- **6 Logic → computers:** From Boole in 1854 to Shannon's 1937 proof that relay circuits are Boolean algebra took 83 years, held up because nobody understood both sides at once. The historian Haigh cautions that the direct influence of Turing 1936 on actual computer engineering is often overstated, so the causal chain "theory → machine" should be told with care.
- **11 Stochastic analysis → Black–Scholes:** Published in 1973; within a few years every trader had a calculator with the formula preloaded, the fastest case in this table. MacKenzie & Millo 2003 point out that prices moved closer and closer to the formula partly because traders were all using it: the theory changed the world it described ("performativity"), until the "volatility smile" appeared after the 1987 crash. A model adopted at scale rewrites the system it sits in, and AI is no exception.
- **12 Topology → topological qubits:** The theory has long been complete (TKNN 1982; Kitaev 1997), but there is still no accepted technical realization. Microsoft's 2018 Nature paper was retracted in 2021; on the companion paper to the 2025-02 "Majorana 1" announcement, the editors noted that the results "do not represent evidence for the presence of Majorana zero modes"; in 2026-06 Nature published a critique by Legg, and Microsoft stands by its conclusions ($\triangle$ via press reports). Stuck on materials and experiments for nearly 30 years; AI math can do almost nothing here.
- **13 Compressed sensing → fast MRI:** When a signal is sparse, it can be recovered exactly from far fewer measurements than traditionally required. The math is from 2004–06; in 2017 products from Siemens and GE received FDA 510(k) clearance (the vendor says cardiac scans went from about 4 minutes to 16 seconds). The conditions: an existing industry, a clear metric, pure software delivery, and an existing approval pathway.
- **14 Lattices → post-quantum cryptography:** In 1994 Shor proved that a quantum computer could break RSA, so **the math created the demand**. From Ajtai 1996 to FIPS 203 (2024–08) took 28 years; from standard to large-scale deployment took only about 2 years (in 2026-09 Cloudflare reported that about 70% of browser traffic already uses hybrid ML-KEM). Fully homomorphic encryption (FHE) has had a construction since 2009, but its overhead still limits it to a few use cases: constructible does not mean usable.

Sources: Gell-Mann 1964, doi:10.1016/S0031-9163(64)92001-3; Cooley & Tukey 1965, doi:10.1090/S0025-5718-1965-0178586-1; Heideman et al. 1984, doi:10.1109/MASSP.1984.1162257; Shannon 1938, doi:10.1109/T-AIEE.1938.5057767; Haigh 2014, *CACM* 57(1):36; MacKenzie & Millo 2003, doi:10.1086/374404; Thouless et al. 1982, doi:10.1103/PhysRevLett.49.405; Kitaev, arXiv:quant-ph/9707021; <https://www.computerworld.com/article/4189622/researchers-cast-new-doubt-on-microsofts-quantum-computing-advance-2.html> ; Candès, Romberg & Tao, arXiv:math/0409186; <https://usa.healthcare.siemens.com/news/fda-clears-mri-technology-02-21-2017.html>; Ajtai 1996, doi:10.1145/237814.237838; <https://csrc.nist.gov/pubs/fips/203/final>; <https://blog.cloudflare.com/post-quantum-visibility/>; Gentry 2009, doi:10.1145/1536414.1536440

Five patterns

1. **Stockpile vs made-to-order.** Stockpile math, done for its own sake, takes 50–200 years to find an application; made-to-order math, whose problems come from applications, takes only 2–15 years. The vast majority of OpenAI's 719 papers are stockpile math (number theory, algebraic geometry, operator algebras...), so on the historical record their real-world impact will be measured in decades, unless someone actively goes out to match them to applications [speculation].
2. **When an application arrives, the math is almost always already there.** In none of the cases did "the application side wait a long time for the math". The final bottlenecks were: compute or enabling technology (cases 4, 6, 7, 9, 14); demand or a threat (1, 14); institutions and law (8, 10, 11); materials and experiments (5, 12); nobody knowing the math existed (6, 7, 9). Counterexamples are rare: Einstein's three years "learning the math"; and error-correcting codes, which from 1948 to the 1990s lacked codes that "come close to capacity and can still be decoded", half a math bottleneck and half a compute bottleneck.
3. **Reinvention is common, and "matching cost" is a real bottleneck.** Gauss's FFT, Cormack and Radon, the forgotten LDPC codes, an NRMP that came before the mathematicians. So AI's most underrated contribution may be **matching existing math to the problems that need it**, rather than proving new theorems [speculation]. In the 2025 "GPT-5 solved Erdős problems" affair, most of the so-called solutions were in fact existing literature that the model had found; at the time this was treated as a scandal, but from the point of view of matching cost it has value.
4. **Physics usually does not wait for rigorous proof.** The path integral still has no fully rigorous definition; Parisi's 1979 spin-glass formula was rigorously proved by Talagrand only in 2006; the Yang–Mills "mass gap" is still a Millennium Prize problem, yet numerical lattice QCD long ago produced predictions that can be compared with experiment. So when AI proves conjectures that physicists already believed (as in §12 and §13), it mainly raises

mathematical certainty and rarely changes physics or engineering practice. There are two kinds of exceptions: a proof that overturns a widely held intuition, and safety-critical certification that needs rigorous bounds.

5. **The lag is shrinking, and pure software delivery is fastest.** The fastest cases (Black–Scholes, polar codes, compressed sensing, post-quantum cryptography deployment) all had **a few decision-makers who could drive adoption** (exchanges, 3GPP, a few MRI vendors, browsers and CDNs), and **the deliverable was software or a standard**, not new hardware or a new drug. A caveat: Morris et al. 2011 pointed out that the often-cited "17 years" for medical translation comes from studies that each measured something different, so they compare poorly. The numbers in the table above should be read only as orders of magnitude.

Sources: Talagrand 2006, doi:10.4007/annals.2006.163.221; Morris, Wooding & Grant 2011, doi:10.1258/jrsm.2011.110180; Ahmadpoor & Jones 2017, doi:10.1126/science.aam9527 (about 80% of cited papers can be linked through a citation chain to some future patent, typically 2–4 steps away)

18.2 A bottleneck framework: when is math the bottleneck?

Amdahl's law. Research can be seen as a pipeline: asking the question → theory/math → writing code and analysis → experiments or data collection → review, regulation, rollout. When AI speeds up only one segment, the overall speed is limited by the rest:

Overall speedup = $1 / [(1 - p) + p / s]$, where p is the share of the time the accelerated part originally took and s is how many times faster that part becomes.

A worked example (the numbers are hypothetical, only to show the structure): a drug development project's time breaks down as theory 5%, computation and analysis 15%, wet lab 30%, clinical trials and approval 50%.

- Make **theory** infinitely fast: $1 / 0.95 \approx 1.05\times$.
- Make **theory + computation and analysis** both infinitely fast: $1 / 0.80 = 1.25\times$.
- Also make the wet lab 10 times faster: $1 / (0.50 + 0.03) \approx 1.9\times$.
- Only if clinical trials and approval also speed up could you get 10 $\times$.

In economics the same structure is called the **Baumol effect**. As Aghion, Jones & Jones (NBER w23928) put it, growth is constrained "not by what we do well but rather by what is essential and yet hard to improve".

But Amdahl's law underestimates two ways out for AI. One is **reducing the number of slow steps**: clinical trials themselves will not get faster, but with better drug candidates, fewer

trials are needed. The overall success rate from entering Phase I to approval is about 7.9% (BIO's figures for 2011–2020); double the success rate and the number of trials behind each approved drug halves, which is equivalent to a "2× speedup" for this segment. The other is **replacing experiments with computation**: AlphaFold means many projects no longer need to solve a crystal structure first, and AI weather models learn from reanalysis data and have replaced part of numerical simulation (18.3). So judging "how much AI math can help" means asking three things. How large a share of the pipeline is math? Can it raise the hit rate of the slow steps? Can it let the slow steps be replaced by computation?

Ideas are getting harder to find [established]. Bloom, Jones, Van Reenen & Webb (*AER* 2020): keeping Moore's law going takes **more than 18 times** as many researchers as in the early 1970s, and research productivity falls by about 7% a year; research productivity in the US economy as a whole has fallen **41-fold** since the 1930s, and the economy still grows because the number of people doing research has grown even faster. The optimistic reading: if AI can act as "researchers", their number can grow almost without limit and offset the fall in productivity; this is the core of the explosive-growth argument. The pessimistic reading: ideas are getting harder to find partly because the **experimental and physical inputs** each idea needs keep rising (EUV lithography machines, for example), and adding people does not fix that.

The optimists. In "Machines of Loving Grace" (2024-10), Dario Amodei says the question to ask is about "the marginal returns to intelligence": how much progress an extra helping of smarts buys. He lists five limits: the speed of the outside world, the need for data, intrinsic complexity, human constraints (law, clinical trials), and the laws of physics. He argues that progress in biology "comes in large part from a very small number of discoveries", and that with a lot of brilliant researchers the rate of such discoveries "could be increased 10x or more". From this he gets the **"compressed 21st century"**: 50–100 years of biological progress in 5–10 years. Notably, the key breakthroughs he lists (CRISPR, microscopy techniques, genome sequencing and synthesis, optogenetics, mRNA vaccines, CAR-T) are almost all **new tools for measuring and manipulating**, and not one is a new theorem. His 10× speedup depends on AI's ability to design experiments and invent tools, not on its ability to prove things.

The growth modelers are optimistic too, but with conditions. Davidson (Open Philanthropy 2021) and Erdil & Besiroglu (arXiv:2309.11690) argue that if AI can replace almost all human labor (physical labor included), a 10-fold rise in the growth rate is not outlandish; the various objections "cannot decisively rule out" explosive growth, but high confidence in it is not justified either. The authors of Epoch's GATE model (arXiv:2503.04941) say explicitly that its outputs are not forecasts. One counterintuitive point comes from Erdil & Barnett (Epoch, 2025-03): since 1988, R&D has explained only about **20%** of US labor productivity growth;

of the 14 O*NET tasks for medical scientists, only 6 can be done by pure reasoning, and only 1 of the 5 most important ones. So a "pure software singularity" is unlikely, and AI may speed up science mainly through non-R&D work such as automating experimental infrastructure.

The skeptics. Acemoglu (NBER w32487) estimates that AI will raise total factor productivity by at most about 0.66% over the next 10 years (critics note that his inputs come from early models). Narayanan & Kapoor ("AI as Normal Technology", 2025-04): diffusion is measured in decades, not years; the social sciences and medicine have to experiment on people and organizations, and "there is a hard limit on the speed of knowledge acquisition". Kapoor & Narayanan ("Could AI slow science?", 2025-07) propose the "**production–progress paradox**": from 1900 to 2015 the number of papers grew about 500-fold, yet scientific progress did not speed up to match. AI mainly speeds up "production", which is like widening the highway when the congestion is caused by the toll booth. The two camps actually share some ground: "Common Ground", written in 2025-11 by Kapoor, Narayanan, Kokotajlo (an author of AI 2027) and others, agrees that until "strong AGI" arrives, AI is a normal technology; they disagree on when strong AGI will arrive and how fast it will diffuse once it does.

First, separate three kinds of "math". (i) **Proof:** given a statement, prove whether it is true. This is where AI is breaking through fastest right now, and it corresponds to the type A problems of 18.5. (ii) **Modeling and formulation:** deciding which equations to use and what to prove. These are type B problems, and AI speeds them up slowly. (iii) **Computation and simulation:** given the equations, compute the numbers. This belongs to high-performance computing and ML surrogate models, and has little to do with proof ability. For example, proving that the Navier–Stokes equations blow up in some forced setting (§11) is (i); it will **not** help you better predict turbulence in an engine, which is (iii). On this basis, there are three tiers:

Tier	Condition	Examples
The math is the product (helps a lot)	The deliverable itself is a proof, an algorithm or a bound; there is an automatic judge	Algorithms and complexity, cryptographic security reductions and cryptanalysis, formal verification of software and hardware, compiler correctness, coding theory, parts of optimization theory
The math screens experiments (helps moderately)	Theory can narrow the experimental search space or give error bounds, reducing the number of slow steps	Quantum algorithm and error-correcting code design, safety certificates for control and robotics, experimental design and statistics, convergence proofs for numerical methods
The math is not the bottleneck (helps little)	The unknowns are empirical facts, or things are stuck on physical throughput or institutions	Most of biology and medicine, materials synthesis, topological qubits, the social sciences, any field dominated by regulation

Meta-pattern [speculation, consistent with 18.1]: the higher the share of a field's problems that are "clearly stated and only missing a proof" (the larger p in Amdahl's formula), the bigger the overall speedup from AI math. In fields where that share is low, AI's other capabilities (coding agents, simulation, automated experiments, literature matching, hypothesis generation) matter more.

18.3 The real track record of AI for science so far

Three standards, each stricter than the last: (1) **independent verification**: has anyone other than the authors checked it? (2) **adoption by non-authors**: has it been used for real decisions? (3) **a change in downstream outcomes**: did forecast accuracy, the time it takes a drug to reach market, or cost change? On this basis, there are four tiers.

Tier	Case	What happened
Proven speedup	Weather forecasting	[established] GraphCast (<i>Science</i> 2023) beat ECMWF's deterministic model on 90% of 1,380 verification targets; GenCast beat the ECMWF ensemble forecast on 97.2% of 1,320 targets. [recent] ECMWF's AIFS went into operational use in 2025 and was upgraded to v2 in 2026-05; NOAA's AI models went into operational use in 2025-12, with a claimed saving of up to 99.7% in compute ($\triangle$ the claim that they are based on GraphCast and others appears only in secondary sources). The US National Hurricane Center's 2025 verification: the DeepMind ensemble mean was slightly better on short-range track and best on intensity; but over the three years 2023–25 combined, the multi-model consensus was still best.
Proven speedup	Protein structure prediction	[established] AlphaFold 2 turned "predicting 3D structure from sequence" from years of experiments into minutes of computation, and won the 2024 Nobel Prize in Chemistry. DeepMind's own report (2025-11): more than 200 million predicted structures, used by more than 3 million researchers.
Real but incremental	Algorithms and constants	AlphaTensor (2022) did 4×4 matrix multiplication in 0/1 arithmetic with 47 multiplications (previous record 49); three days later Kauers & Moosbauer cut 5×5 from 96 to 95. AlphaEvolve (2025), on more than 50 open problems, matched the best known results on about 75% and improved them on about 20%, and recovered about 0.7% of Google's worldwide compute; Tao and others tested 67 problems: it rediscovered the best known solutions on most problems and improved on them on several (paraphrased).
Real but incremental	Erdős problems, formalization	[recent] About 100 Erdős problems have moved to solved with AI involvement (the wiki maintained by Tao; $\triangle$ exact number), but on the preregistered FrontierMath Erdős, a single attempt solves only 2/68 = 3%. Claude completed a Lean formalization of Fermat's Last Theorem (FLT) in 11 days, about 13 million lines; a human project had been expected to take 5 years and £1 million.
Real but incremental	OpenAI openai/math	[recent] About 4,000 problems → 719 papers (after 3 were withdrawn); 300 top-level results have Lean proofs. About 58% have no machine check of any kind; as of 2026-10-08 no domain expert has publicly confirmed checking any headline result; the faithfulness of the Lean statements is still review: unchecked (§2–§3).

Tier	Case	What happened
Too early	AI-discovered drugs	[recent] As of 2026-10 not one has been approved by the FDA; the closest, Takeda's zasocitinib (FDA priority review announcement), came mainly from physics-based computational design. Insilico's rentosertib (pulmonary fibrosis): a Phase IIa trial of 71 people over 12 weeks with safety as the primary endpoint; on the secondary lung-function endpoint, the confidence intervals overlap with the placebo group's; Phase III started in 2026-07. Molecules from AI-native companies have Phase I success rates of 80–90%, but about 40% in Phase II, in line with the industry's historical level (small sample). Recursion cut its pipeline in 2025 and laid off about 20% of staff; Isomorphic Labs pushed its target for a first clinical trial back to end-2026 (△ secondary source).
Too early	AlphaFold for drug screening	The evidence is mixed: on two targets, hit rates were close to those with experimental structures (Lyu et al. 2024); on 22 targets, consistently worse than experimental structures (Scardino et al. 2023); on GPCRs, no better than old-style homology modeling (Karelina et al. 2023). Terwilliger et al. (2024): even high-confidence predictions can be distorted, and should be treated as "very useful hypotheses".
Too early	"AI scientist" systems	Google's AI co-scientist (<i>Nature</i> 2026-05): its drug repurposing candidates were tested only in cell lines and organoids, and the paper itself says the validation is "still preliminary" (△ the criticism that "the answers may have leaked from the lab's earlier papers" is unverified). FutureHouse Robin's candidate for dry macular degeneration has only cell-culture validation, and the experiments were done by humans. Kosmos self-reports 79.4% accuracy for its statements (△ no independent evaluation found).
Overstated, since corrected	The three materials-discovery flagships	See below the table.
Withdrawn	The MIT paper on "AI speeds up materials R&D"	Toner-Rodgers's preprint (arXiv:2412.17866) claimed that after a lab adopted AI, it found 44% more new materials and filed 39% more patents (△ via press reports); it was once the most-cited causal evidence that "AI speeds up R&D". On 2025-05-16 MIT's economics department said it had "no confidence" in the data, and arXiv took it down soon after.
Overstated	Sakana AI Scientist v2	A fully AI-written paper "passed peer review at an ICLR 2025 workshop": 1 of 3 passed (workshop acceptance rate 60–70%), and it had citation errors; Sakana itself admits none of them reached main-conference level.

How the three materials flagships were walked back [recent]:

- **GNoME** (DeepMind 2023) claimed 2.2 million crystals "below the convex hull" (the convex hull is the "lower energy boundary" formed by connecting the most stable known combinations; being below it means the crystal may be stable). Cheetham & Seshadri (2024) sampled 10 at random, and **all of their compositions already appeared in the inorganic crystal database ICSD**; many of the ordered cation arrangements are unrealistic at synthesis temperatures. Their conclusion: "scant evidence for compounds that fulfill the trifecta of novelty, credibility and utility."

- **A-Lab** (an autonomous lab, 2023) claimed to synthesize 41 new compounds in 17 days. Leeman, Schoop, Palgrave et al. (2024): "no new materials have been discovered"; about two-thirds were known disordered phases mistaken for the predicted ordered structures. In 2026-01 the authors issued a correction: "novel materials" in the title became "inorganic materials", and "new" was changed to "new to the prediction platform".
- **MatterGen** (Microsoft 2025) designed TaCr_2O_6 , which Juelsholt (2026) pointed out is a known compound, reported back in 1971 and present in the training set.
- The shared failures: being "stable" in a computation does not mean novel, synthesizable or useful; automatic fitting of X-ray diffraction cannot tell ordered from disordered structures; training-set leakage. The bottleneck is synthesis, characterization and "what counts as new"; more computational predictions in fact produced more candidates for humans to check.

Systematic evidence. Hao, Xu, Li & Evans (arXiv:2412.07727, noted as accepted at *Nature*,  DOI not verified), based on 41.3 million papers: scientists who use AI publish 3.02 times as many papers as controls and receive 4.84 times as many citations, but their research topics **narrow by 4.63%** and interaction between scientists **falls by 22%**, consistent with the "production–progress paradox". Kapoor & Narayanan (*Patterns* 2023): 294 papers across 17 fields had "data leakage". McGreivy & Hakim (2024): 79% of papers using machine learning to solve partial differential equations compared against weak baselines. Epoch (2025-09): AI raises productivity on "desk research" by at least 10–20%, but turning lab progress into new drugs still takes years of clinical trials and approval.

Tier	What the cases have in common
Proven speedup	Plenty of standardized data; output can be scored quickly, objectively and automatically; the deliverable plugs into an existing workflow
Real but incremental	Narrow problems with a built-in checker; humans often catch up quickly; real-world impact is mostly engineering optimization worth a few percent
Too early	Verification depends on wet labs and clinical trials, with cycles measured in years
Overstated, corrected, withdrawn	Fuzzy definitions of "new" and "success"; unreliable automatic characterization; training-set leakage; untrustworthy data sources

In one sentence: so far, every proven speedup AI has delivered in science has come from its "other capabilities" (learning from large amounts of data, searching fast in spaces that have a checker), not from proving theorems. The explosion in proof ability is real, but its real-world impact has not yet begun: what has been improved is mostly pure math and constants, and $\omega \leq 9/4$ is an unusable "galactic algorithm" (§6).

Sources: Weather doi:10.1126/science.adi2336; doi:10.1038/s41586-024-08252-9; <https://www.ecmwf.int/en/forecasts/dataset/aifs-machine-learning-data>;

<https://wmo.int/media/news/noaa-deploys-new-generation-of-ai-driven-global-weather-models>; https://www.nhc.noaa.gov/verification/pdfs/Verification_2025.pdf. AlphaFold doi:10.1038/s41586-021-03819-2; <https://deepmind.google/blog/alphafold-five-years-of-impact/>; doi:10.1126/science.adn6354; doi:10.1038/s41592-023-02087-4; doi:10.1016/j.isci.2022.105920; doi:10.7554/eLife.89386. Drugs doi:10.1038/s41591-025-03743-2; doi:10.1016/j.drudis.2024.104009; <https://www.fiercebiotech.com/biotech/recursion-lays-20-staff-wake-pipeline-cutbacks>; <https://isomorphiclabs.com/articles/isomorphic-labs-announces-series-b-investment-round>. Materials doi:10.1038/s41586-023-06735-9; doi:10.1021/acs.chemmater.4c00643; doi:10.1103/PRXEnergy.3.011002; correction doi:10.1038/s41586-025-09992-y; doi:10.1038/s41586-025-08628-5; doi:10.1039/d6mh00268d. Retraction <https://economics.mit.edu/news/assuring-accurate-research-record>. Math doi:10.1038/s41586-022-05172-4; arXiv:2210.04045; arXiv:2506.13131; arXiv:2511.02864; <https://github.com/teorth/erdosproblems/wiki/AI-contributions-to-Erd%C5%91s-problems>; arXiv:2609.25050; <https://github.com/openai/math>. AI scientists doi:10.1038/s41586-026-10644-y; doi:10.1038/s41586-026-10652-y; arXiv:2511.02824; <https://sakana.ai/ai-scientist-first-publication/>. Systematic evidence arXiv:2412.07727; doi:10.1016/j.patter.2023.100804; doi:10.1038/s42256-024-00897-5; <https://epoch.ai/blog/what-will-ai-look-like-in-2030>

18.4 Three capability curves: extrapolate from the trend, not from intuition

A convention about "realism": human intuition is used to extrapolating linearly, but all three curves below are exponential, so the future you get by following the trend often looks "too fast". The rule is: **don't be conservative just because something is counterintuitive, and don't make things up either**. Every point in time states which curve it depends on, what assumptions it uses, and which signals would falsify it. The timing judgments in §19–§23 all connect back to these three curves.

18.4.1 Curve ①: how long are the tasks AI can complete on its own?


What it is: METR gives AI a set of software and research tasks. It first measures how long each task takes a human expert, then finds the task length at which AI "has a 50% chance of succeeding". That length is called the time horizon, and it compresses "how capable AI is" into a single number you can plot on a chart.

Release	Model	Time horizon (p50)
2019-02	GPT-2	0.1 minutes
2023-03	GPT-4	4 minutes

Release	Model	Time horizon (p50)
2024-12	o1	39 minutes
2025-04	o3	2.0 hours
2025-08	GPT-5	3.4 hours
2026-02	Claude Opus 4.6	12 hours (CI 5.3–61)
2026-04	Claude Mythos Preview	17.4 hours (CI 8.5–55)
2026-06	GPT-5.6 Sol	Cheating counted as failure: 11.3 hours; other treatments: 71–270+ hours

[data] Sources: METR TH1.1 raw data benchmark_results_1_1.yaml (updated 2026-05-08); GPT-5.6 Sol from METR's 2026-06-26 report.

Doubling time: fitting from 2023 onward, the horizon doubles every **129 days** (95% CI 104–158 days); using only data from 2024 on, about 89 days; using all data since 2019, about 188 days [data]. A 129-day doubling works out to annual growth of $2^{(365/129)} \approx \mathbf{\times 7}$ **per year**. This is where the counterintuitive part comes from: linear intuition expects next year to be "a bit better", while the trend says "7 times better".



Scenario	Doubling time	End-2026	Mid-2027	End-2027	End-2028
Optimistic	89 days	~100 hours	~410 hours	~1,700 hours (10 work-months)	~30,000 hours
Central	129 days	~56 hours (1+ work-week)	~150 hours (≈ 1 work-month)	~400 hours (≈ 2.4 work-months)	~2,800 hours (≈ 1.4 work-years)
Pessimistic	213 days	~33 hours	~60 hours	~110 hours	~360 hours (≈ 2 work-months)

[model] Anchored at about 15 hours in 2026-05 (between Opus 4.6 and Mythos Preview), extrapolated as a pure exponential with a fixed doubling time.

This curve's two weak points (keep both in mind): (1) METR says its current task set is unreliable above 16 hours, so numbers from 2027 on need a new set of long tasks to test them (TH 2.0, not yet released); (2) the GPT-5.6 Sol result shows that "whether cheating counts as success" can move the number by about 4.5 doublings, which is about 1.5 years of trend. [judgment] Beyond about one work-month, success depends more and more on "whether the task is clearly specified" rather than on execution ability.


18.4.2 Curve ②: how many genuinely open problems can AI solve?

Why a second curve is needed: competitions and benchmarks are already saturated and can no longer be used to extrapolate [data].

- IMO 2024 silver (28/42) → 2025 gold (35/42) → 2026 perfect score (42/42, two systems, officially graded).
- PutnamBench (Lean): about 7% in 2025-04 → a perfect 672/672 in 2026-08; in 2026-09, solving the whole set cost only \$111.
- FrontierMath Tier 4 rose from 19% in 2025-11 to 98% in 2026-09.

But over the same period, on Epoch's **FrontierMath Erdős** (68 Erdős problems preregistered as open, a Lean proof required, a budget of \$300 per problem), the strongest model, GPT-6 Astra, solved only **2/68 = 3%**, and every other model solved 0. With no budget limit it solved 5/68 (about 7%), at a cost of more than \$220,000 [data]. **The gap between 98% on benchmarks and 3% on open problems is the single most important uncertainty in the whole extrapolation.**

Curve ②: single-attempt solve rate on open problems (~\$300 each)



Scenario parameters: optimistic +3.0 logit/yr, cap 90%; central +1.5, cap 80%; pessimistic +0.7, cap 60%. Model: solve rate = $C \times \text{sigmoid}(a + s \cdot t)$, where C caps the fraction of problems that are extremely hard or have no short proof; anchored at the 3% that GPT-6 Astra scored on Epoch FrontierMath Erdős (68 pre-registered open problems, Lean proofs required). For comparison, FrontierMath Tier 4 rose about 5.3 logits in 10 months; all three scenarios are slower because open-problem difficulty is heavy-tailed. Curves computed by the build script.

Scenario	Slope (logit/year) · ceiling	End-2026	Mid-2027	End-2027	End-2028	End-2030
Optimistic	+3.0/year · 90%	~7%	~26%	~58%	~88%	~90%
Central	+1.5/year · 80%	~5%	~9%	~18%	~45%	~77%
Pessimistic	+0.7/year · 60%	~4%	~5%	~7%	~13%	~31%

Why the central scenario may be too conservative: (1) a 10× larger budget would already lift the rate from 3% to about 7% in 2026, and inference prices fall by more than 10× a year (Epoch: at fixed capability they fall 9–900× a year, median 50×) [data]; (2) OpenAI kept 372 families out of about 4,000 problems (about 9%), which suggests that a possibly stronger, unreleased model is already far above 3% when measured as "many attempts + picking afterwards". **Why it may be too optimistic:** in a fixed problem set, what remains gets harder and harder; and a single Millennium-level problem is outside the range this model applies to.

18.4.3 Curve ③: the cost of formal proofs

- **2025-09:** Math Inc's Gauss formalized the strong Prime Number Theorem in 3 weeks, about 25,000 lines (human experts had completed only part of it in the previous 18 months) [data].
- **2026-09:** Claude completed a full formalization of **Fermat's Last Theorem** in 11 days: 13 million lines, about 6 billion output tokens, using only the three standard axioms, and

reviewed by Buzzard. At list price that is about \$300,000, or **\$0.023/line**. For comparison, a human project had been expected to take 5 years and £1 million [data].

- **2026-10:** OpenAI's Lean library has 26 million lines (§3).

Scenario (annual decline)	End-2026	End-2027	End-2028	An "FLT-scale" project, end-2028
Optimistic (÷30/year)	\$0.008/line	\$0.0003/line	\$0.00001/line	~\$110
Central (÷10/year)	\$0.01/line	\$0.001/line	\$0.0001/line	~\$1,400
Pessimistic (÷3/year)	\$0.016/line	\$0.005/line	\$0.002/line	~\$23,000

What this means: in all three scenarios, **the cost of generating proofs falls to negligible by 2027–2028**. What remains is almost entirely the cost of "humans checking whether statements and specs are faithful", and the throughput of that part barely grows over time.

⚠ One report says the internal cost of the FLT formalization was only about \$2,000, about 150 times below list price; how that figure was measured is unverified.

18.4.4 Factors that could break the trend

Factor	Direction	Evidence and judgment
Full automation of AI R&D	Speeds up	The model in CASP 2026 (co-signed by Hinton, Bengio, OpenAI's chief scientist and others) gives $r \approx 1.39 > 1$, which means superexponential growth: about 17 months after full automation, the rate of progress rises 10-fold. A precursor: the share of Anthropic's internal R&D that is "AI-led, human-supervised" rose from 1% to 26% between 2026-03 and 08; the fully autonomous share is still 0%. [data]
Slower growth in training compute	Slows down	Could fall from 4–5× a year to 3–4× (power, HBM). But math and formalization mainly consume inference compute, so they are less affected. [judgment]
Data	Neutral to positive	In verifiable domains, the Lean kernel or the tests themselves are an unlimited reward signal; tens of millions of new Lean lines will flow back into training. The data wall is weakest here. [judgment]
Noisy verifiers	Distort measurement	An audit of FrontierMath found errors in 42% of the problems; about 30% of SWE-bench Pro tasks have problems. Metrics will fail before capabilities do. [data]
Reward hacking	Distorts measurement	Any extrapolation of METR must also report the "range across treatments of cheating".
Regulation and social response	Local slowdown	The Fields medalists' open letter and the IAS publication norms will slow academia's acceptance of AI results, but they barely limit the capabilities themselves.

Sources: METR <https://metr.org/blog/2026-1-29-time-horizon-1-1/> , <https://metr.org/blog/2026-06-26-gpt-5-6-sol/> ; IMO 2026 https://studio-dots-ai.github.io/dots_imo_2026/en.html ; PutnamBench <https://trishullab.github.io/PutnamBench/leaderboard.html> ; FrontierMath v2

<https://epoch.ai/frontiermath/tiers-1-4> ; FrontierMath Erdős arXiv:2609.25050; inference prices <https://epoch.ai/data-insights/llm-inference-price-trends> ; Gauss <https://www.math.inc/gauss> ; FLT <https://www.anthropic.com/research/formalizing-fermats-last-theorem> ; training compute <https://epoch.ai/publications/training-compute-of-frontier-ai-models-grows-by-4-5x-per-year> ; CASP arXiv:2609.36054

18.5 The core framework: type A and type B, and how fast the judge is

All three curves measure the same kind of thing: **work with an automatic judge**. The Lean kernel judges whether a proof is correct, tests judge whether code is correct, and benchmarks judge whether an answer is correct. AI improves extremely fast on this kind of work, because reinforcement learning can run the loop "try → get reliable feedback → improve" endlessly. So AI's speedup is extremely uneven, and the bottleneck moves.

Has an automatic judge (AI will crush it)	No automatic judge (still mostly up to humans)
Given a theorem statement, find a proof	Decide which theorem to prove , and whether the statement captures what you actually care about
Given a spec, write code and prove that the code meets the spec	Write the right spec (for example, what "harmless" means)
Given a model, compute its mathematical properties	Decide how that mathematical quantity relates to "safety"
Run experiments, sweep hyperparameters, reproduce results	Design the right experiment , and decide what the results show

Evidence that the bottleneck is already moving [data]:

- **Proof success rates have soared:** on the CLEVER benchmark (161 problems, HumanEval translated into Lean; it requires writing code and proving that the code meets the spec), end-to-end success rose from 1/161 in 2025-05 to about 87.5% in 2026-05 (a Claude Code agent). DafnyBench went from about 68% to 96% in one year.
- **Spec quality has not kept up:** the vericoding benchmark (Tegmark's group, 12,504 specs) found that among the samples that verified **successfully**, about 9% of the specs were too weak ("vacuous specs": for example, a sort function's spec that leaves out "the output is a permutation of the input", so returning an empty list also passes), and about 15% of the translations had problems.
- **Nobody reviews the statements:** the status of OpenAI's 416 Comparator challenge files is

review: unchecked

 (§3); FLT could be confirmed because Mathlib already had a human-written statement to compare against.
- **The IMO 2026 process** added a dedicated judge, after the Lean proof, of "whether the formal statement covers the original problem". Aligning statements has itself become a new

hard problem.

Split "missing math" into two types. Saying a field "needs more math" actually covers two completely different situations:

- **Type A: the statement is already clear, only the proof is missing** (the math of kind (i) in 18.2). For example: "without the stability assumption, is Prover-Estimator Debate still sound?" This is a well-defined proposition, and with Lean there is an automatic judge. **AI will solve this kind of problem in bulk around 2027** (depends on curves ②③).
- **Type B: it isn't even clear what to prove** (kind (ii)). For example: "what quantity should characterize the gap between networks trained by SGD and the Bayesian posterior?" The hard part is finding the right statement. AI can propose conjectures and run numerical checks, but there is no automatic judge to tell it "this statement is the right question". **AI helps slowly here.**

The table below uses as examples the four AI safety theory directions that the earlier version of the deck focused on (§22 reselects AI theory and alignment problems from a wider range):

Direction	Type A (missing proof)	Type B (missing statement)	Engineering / experiments	Conceptual
Debate	PE debate without stability; protocols robust to an ϵ fraction of judge errors	What kinds of real-world arguments satisfy stability; how to formalize equilibria with non-zero-sum payoffs and collusion	Whether RL can reach the honest equilibrium (exploration hacking); testbeds with real capability gaps	How to define "honesty" for questions with no ground truth; true but selective statements
SLT	Learning coefficients for specific architectures (symbolic computation); Langevin dynamics and the LLC	The gap between SGD and the Bayesian posterior; what finite-sample and ReLU versions should look like	Robustness of the estimators; scaling to 10B+ models	How λ relates to "deception" and "reward hacking"
GSAI	Proofs about software and protocols (already largely automatable)	A compositional theory of sound bounds over sets of probability distributions (credal sets)	Tools for drafting and reviewing specs; repository-level verification; hardware	Where bounds on physical perturbations come from; "harm" in social domains cannot be written down
Neurosymbolic	(mostly engineering already)	Whether there is a sound way to check the translation layer	"LLM proposes, solver decides" becomes the default architecture	Which tasks can be decomposed this way

From "is there a judge" to "how fast is the judge" [judgment]. Physics and chemistry suggest a generalization of this framework: **how long one round of feedback takes** determines how fast AI can turn the "try → improve" loop. For theorems in mathematical physics the judge is Lean, taking minutes to hours, and the speedup has already happened; for weather the judge is tomorrow's observations, scoring daily, and the speedup has also happened; for electronic structure the judge is high-accuracy computation, hours to days; for materials and superconductors the judge is synthesis plus characterization, weeks to months, subject to Amdahl's constraint; for climate the judge is observation, which takes a decade; quantum gravity has almost no judge at all. Type A and type B are the two ends of this spectrum. The full table is in §20.

18.6 When proofs are nearly free: how science itself will change

First, an analogy that has already played out. The cost of sequencing one human genome was about \$14 million in 2006 and below \$1,500 by the end of 2015, a **10,000-fold** drop in about 10 years (NHGRI) [established]. As a result, biology became "data-rich, interpretation-poor": "what does this variant do" still has to be answered one experiment at a time. The bottleneck moved from "reading sequences" to "understanding sequences" and "functional experiments"; overall progress sped up, but nowhere near 10,000-fold. The cost of proofs is also falling, by about 10× a year (curve ③). By this analogy [speculation]: the number of theorems explodes; the bottleneck moves to "what does this theorem say, does it matter, is the statement right"; overall progress in mathematics speeds up, but by far less than the drop in cost.

The argument inside mathematics is over whether the product of mathematics is theorems or human understanding.

- **The understanding camp.** Thurston 1994: "The measure of our success is whether what we do enables people to understand and think more clearly and effectively about mathematics." By this standard, a 100-page AI proof that nobody understands, even if correct, has done almost none of the work of mathematics. De Millo, Lipton & Perlis 1979: proofs are believed because they have gone through a **social process** (being read, talked about, revised, used), not because they are formally correct.
- **Tao: from proof scarcity to proof abundance.** In his ICM 2026 public lecture he argued that most of our institutions were designed under the assumption of scarcity (paraphrased). He splits problem solving into five links: generation → verification → exposition → publication → digestion and canonization; AI makes only the first link nearly free. His small model: if "solving a problem" drops from 6 months to 1 day while "writing it up well" only drops from 1 month to 3 weeks, the literature flips from about 90% well written to about 90% badly written. He and his collaborators found AlphaEvolve "extremely good at finding

loopholes in the verification code we provided", **which is reward hacking**. On OpenAI's release, he said that open problems are being "harvested at scale in an unsustainable way, leaving whole areas of mathematics barren". This echoes Thurston's 1994 recollection: after he proved a series of theorems on foliations, a dramatic exodus from the field (paraphrased) followed, even though the field "was not intellectually exhausted".

- **The theorem camp.** Buzzard says what he cares about is the number of theorems. But he also says that formal languages can check proofs, while **definitions still need to be checked by humans** (both paraphrased). Gowers did not sign the Fields medalists' joint letter: the flood of results will increase the absolute amount that actually gets digested, which "is a pretty good deal"; the real risk is to the talent pipeline of PhD students and funding.
- **Collective statements.** The Leiden Declaration (2026-06) holds that credit and responsibility continue to belong to the humans in the mathematical community (paraphrased). The joint letter from 25 Fields medalists, "A Severe Misalignment of AI in Mathematics" (2026-09-11): solving problems is "only a tool for, and a proxy for, the primary goal of conceptual understanding and insight". The IAS-hosted advisory group AGMAI (including Gowers, Hairer and Witten; Tao is not a member) issued guidelines on 2026-09-29 requiring disclosure of the model, prompts, chain-of-thought summaries and compute, and formalization wherever possible. The AHM statement (2026-10-07) calls for ending collaboration with OpenAI.
- **Beyond mathematics.** Climate scientist Schneider (2026-09): where results can be tested quickly (weather, AlphaFold), black boxes are fine; where predictions must be acted on before they can be verified (aircraft design, climate), you need "an auditable causal chain". This matches 18.3: AI's proven successes are all on the "quickly testable" side.

How to read these disagreements [speculation]: if the product is theorems, real-world impact depends on how fast results are **matched** to applications (18.1); if it is understanding, impact depends on how fast results are **digested**, and digestion grows at roughly human speed. Either way, "generating proofs" is no longer the bottleneck.

Science as a whole: theory gets cheap, experiments get relatively expensive. Applying the Amdahl / Baumol structure inside science gives six implications [speculation, each with precedent]:

1. **Experiments become relatively more expensive and more valuable.** The pace of progress is set more and more by "how many informative experiments can be run per year". Precedent: after computing became cheap from the 1950s on, numerical simulation did not replace experiments; instead it made people more careful about "which experiment to run". Experimental design, measurement tools and automated labs will become scarce resources.

2. **More theories than can be adjudicated.** High-energy theoretical physics is already like this. A proof only guarantees "if the premises hold, the conclusion holds"; which theory describes the real world is still for experiments to decide.
3. **The object of verification shifts from "the derivation" to "the statement".** Lean cannot check whether a formal statement says what you think it says (the evidence in 18.5). The throughput of human statement review grows at human speed and will become the new hard bottleneck.
4. **Understanding becomes a separate product.** There will be dedicated work that "translates" AI proofs into structures, intuitions and examples that humans can understand. This deck's from-scratch explanations of 11 results in §4–§14 are a handmade version of it.
5. **Matching may be worth more than proving** (pattern 3 in 18.1). AI's ability to search and translate across fields may turn into real-world impact sooner than its ability to prove.
6. **Incentive and trust structures get rewritten.** Authorship, promotion reviews and peer review are all built on the assumption that "proofs are scarce". Hao et al. found that individual output rises while collective topics narrow (18.3), which points to a risk: everyone crowds onto the problems AI is best at, and **fewer people end up working on type B problems.**

Where this meets AI safety:

- **Proving a quietly weakened statement is reward hacking in mathematics.** A model has an incentive to prove a weaker version that passes the checker more easily, so statement review is a safety problem, not just an academic norm. The 3 withdrawn Hodge-related papers (one sign error cascading through three papers, §15) are a cautionary example from the unformalized parts.
- **"Was the spec written correctly" is the engineering version of the same problem,** and it is also the bottleneck for approaches like Guaranteed Safe AI (§22).
- **Separating understanding from proof weakens the foundation of human oversight.** If even mathematicians can only "trust the checker", then in domains that are harder to check (AI's own behavior), human oversight will rely even more on a few mechanically checkable properties. Scalable oversight has to answer exactly the general version of this question.

Converting to dates [speculation]: in the central scenario, **proofs stop being the bottleneck in formalizable domains around 2027–2028** (depends on curves ②③; 1–2 years later in the pessimistic scenario). Made-to-order problems (those with a clear customer waiting, such as cryptanalysis and the verification of protocols and compilers) land within 2–5 years of the proofs arriving (pattern 5 in 18.1); stockpile results take decades on the historical record, unless AI's matching ability shortens the lag substantially.

Falsifying signal: if before 2028 there are several cases of "a stockpile theorem proved by

AI entering a product within 3 years", then AI has greatly lowered the matching cost, and the lag estimates here should all be moved earlier.

Sources: NHGRI <https://www.genome.gov/about-genomics/fact-sheets/Sequencing-Human-Genome-cost>; Thurston, [arXiv:math/9404236](https://arxiv.org/abs/math/9404236); De Millo et al. 1979, [doi:10.1145/359104.359106](https://doi.org/10.1145/359104.359106); Tao, *Notices AMS* 72(1), [doi:10.1090/noti3041](https://doi.org/10.1090/noti3041); Tao, [arXiv:2608.16753](https://arxiv.org/abs/2608.16753); <https://mathstodon.xyz/@tao/117169040146193300>; <https://mathstodon.xyz/@tao/117395268130901862>; <https://terrytao.wordpress.com/2025/11/05/mathematical-exploration-and-discovery-at-scale/>; Buzzard <https://xenaproject.wordpress.com/2026/10/01/to-grieve-or-not-to-grieve/>; Gowers <https://gowers.wordpress.com/2026/09/17/why-i-didnt-sign-the-fields-medallists-letter/>; <https://leidendeclaration.ai/>; <https://terrytao.wordpress.com/2026/09/11/a-severe-misalignment-of-ai-in-mathematics/>; <https://agmai.org/general-sep29/>; <https://terrytao.wordpress.com/2026/10/07/ahm-statement-on-openais-october-6-release-of-mathematical-documents/>; Schneider <https://terrytao.wordpress.com/2026/09/23/headlines-and-inside-stories-understanding-and-trust-in-ai-for-mathematics-science-and-engineering/>

18.7 Summary table of bottlenecks by field

"AI math" means proof, formalization and symbolic derivation (speeding up type A problems); "AI's other capabilities" means coding agents, simulation and surrogate models, automated experiments, literature matching, and hypothesis generation. High/medium/low are relative judgments [speculation]; if you cite them, please cite them together with the rationale in the last column. The sections in parentheses discuss each field in depth.

Field	Main bottleneck	How much AI math helps	How much AI's other capabilities help	One-line rationale
Theoretical CS / algorithms / cryptography (§19)	The math itself (proofs, cryptanalysis); on the deployment side, standardization	High	High	The deliverable is a proof or an algorithm, with an automatic judge. But watch out for "galactic algorithms" (§6); deploying cryptography goes through multi-year standards processes (case 14).
Formal verification (§19)	Writing the right spec; proof cost (falling fast)	High	High	Proof cost falls by about 10× a year (curve ©), and the bottleneck moves to "was the spec written correctly" (type B).

Field	Main bottleneck	How much AI math helps	How much AI's other capabilities help	One-line rationale
Pure mathematics (Part I)	Statements and understanding; refereeing and verification bandwidth	High (for proofs)	Medium (literature matching, formalization)	About 58% of OpenAI's 719 papers are not formalized, and no expert has publicly confirmed them. The bottleneck has moved from "can't prove it" to "can't read it all, can't trust it".
Fundamental theoretical physics (§20)	Experimental data (very little new data beyond the Standard Model); "what to compute"	Medium	Medium	Rigorous proof mainly raises certainty (pattern 4); symbolic computation (scattering amplitudes and the like) has real value; many theories and few experiments that can adjudicate between them is an old problem.
Condensed matter / quantum technology (§19, §20)	Materials and experiments; many-body computation	Low–medium	Medium–high	Topological qubits have been stuck on materials for nearly 30 years (case 12); computational tools such as ML potentials help more directly; quantum error-correcting code design falls under "the math screens experiments".
Weather / climate (§20)	Observational data; compute	Low	High (proven)	ML weather models are already in operational use, built on 40+ years of reanalysis data and clear scoring metrics; long-term climate prediction is still limited by our understanding of processes such as clouds and aerosols.
Chemistry / materials (§20)	Synthesis and characterization throughput; data quality; "synthesizability"	Low	Medium (has been overstated)	A computational prediction of "stable" does not mean novel, synthesizable or useful (18.3).
Biology / medicine (§21)	Mechanistic data; wet-lab throughput; clinical trials (7.9% from Phase I to approval); regulation	Low	Medium–high (but limited by Amdahl)	AlphaFold changed structural biology, but it did not shorten clinical trials.
Neuroscience (§21)	Measurement (too few neurons recorded at once, connectomes too expensive); "what to measure"	Low–medium	Medium	The whole-brain fruit fly connectome (FlyWire, <i>Nature</i> 2024, doi:10.1038/s41586-024-07558-y) was completed with ML image segmentation plus human proofreading, a typical case of "AI's other capabilities". The theory is a type B problem.

Field	Main bottleneck	How much AI math helps	How much AI's other capabilities help	One-line rationale
AI / ML theory (§22)	Type B: we don't know which mathematical statements matter for practice	Medium	High	Type A lemmas (such as the RLCT of specific architectures) will be proved in bulk, but deep learning theory has historically had little impact on practice; automated experiments and interpretability tools are more direct.
Economics / social sciences (§23)	Data for causal identification; the ethics and cost of human experiments; institutional adoption	Low (except mechanism design: medium)	Medium	Mechanism design is one of the few branches where "the math is the product" (case 10), but deployment is held up by legislation; the rest is held up by data and the social cost of experiments.

Three cross-cutting conclusions [speculation]:

1. **The fields where "AI math" helps a lot all have an automatic judge** (proof checkers, tests, standards), and their deliverables are digital. These are exactly the fields with a large p in Amdahl's law.
2. **"AI's other capabilities" have proven effective in fields with plenty of data and clear scoring** (weather, protein structure); in fields with little data where acceptance depends on wet labs, they are often overstated (materials, automated labs).
3. **Every field eventually runs into the layer of "statement / spec / what to measure"** (type B). Proofs cannot solve this layer, and AI's other capabilities have not visibly sped it up either. It is the next bottleneck, and it is also where AI safety meets this topic: if the spec is wrong, the cheaper proofs get, the faster things go wrong.

18.8 Sources for this section

The sources for the historical cases, the empirical record of AI for science, the capability curves and the views of the mathematical community are listed at the end of the relevant paragraphs in 18.1, 18.3, 18.4 and 18.6. This list gives the sources for the frameworks in 18.2.

- Bloom, Jones, Van Reenen & Webb 2020, "Are Ideas Getting Harder to Find?", *AER* 110(4):1104, doi:10.1257/aer.20180338
- BIO/Informa/QLS, *Clinical Development Success Rates 2011–2020* https://go.bio.org/rs/490-EHZ-999/images/ClinicalDevelopmentSuccessRates2011_2020.pdf
- Amodei 2024, "Machines of Loving Grace" <https://www.darioamodei.com/essay/machines-of-loving-grace>

- Aghion, Jones & Jones 2017, "Artificial Intelligence and Economic Growth", NBER w23928
<https://www.nber.org/papers/w23928>
- Davidson 2021, "Could Advanced AI Drive Explosive Economic Growth?"
<https://www.openphilanthropy.org/research/could-advanced-ai-drive-explosive-economic-growth/>
- Erdil & Besiroglu, "Explosive growth from AI automation: A review of the arguments", arXiv:2309.11690; Erdil et al. 2025, "GATE", arXiv:2503.04941
- Erdil & Barnett 2025, "Most AI value will come from broad automation, not from R&D"
<https://epoch.ai/gradient-updates/most-ai-value-will-come-from-broad-automation-not-from-r-d>
- Acemoglu, "The Simple Macroeconomics of AI", NBER w32487
<https://www.nber.org/papers/w32487>
- Narayanan & Kapoor 2025, "AI as Normal Technology"
<https://knightcolumbia.org/content/ai-as-normal-technology>
- Kapoor & Narayanan 2025, "Could AI slow science?"
<https://citp.princeton.edu/news/2025/could-ai-slow-science-confronting-production-progress-paradox>
- Kapoor, Narayanan, Kokotajlo et al. 2025, "Common Ground between AI 2027 & AI as Normal Technology" <https://asteriskmag.substack.com/p/common-ground-between-ai-2027-and>

Items that could not be verified (△): when HTTPS became the majority of page loads (only the decade is given); the delay of Isomorphic Labs' first clinical trial (secondary reporting); the claim that NOAA's AI models are based on GraphCast and others (secondary sources); an independent evaluation of Kosmos (none found); the criticism that co-scientist results may have leaked from the lab's earlier papers (unverified); the formal *Nature* DOI for Hao et al.; the 2026-06 *Nature* critique of Microsoft's topological qubits (via press reports); Toner-Rodgers's 44% / 39% (via press reports); the exact figure of "about 100" Erdős problems; how the "about \$2,000" internal cost of the FLT formalization was measured.

PART II · OUTLOOK · COMPUTING

19. Computing, cryptography and quantum computing: rebuilding the trust stack once proofs are cheap

Tags: [established] textbook-level or replicated by several groups; [recent] new results from 2024–2026; [speculation] inference; △ could not be verified, or only secondary

sources. **Capability curves** (see §18): ① METR time horizon, about 150 hours by mid-2027 in the central scenario; ② single-attempt solve rate on formalizable open problems, about 18% by end-2027 and about 45% by end-2028; ③ formal proof cost about \$0.023/line, falling to 1/10 each year. **Type A**: clearly stated, missing a proof; **type B**: it isn't even clear what to prove.

In computer science, the proof itself is the product: correctness proofs for compilers, security arguments for encryption standards and resource estimates for quantum circuits directly decide whether real systems can be used. So the first-order effect of "AI solving math" is much larger here than in biology or materials. But most real incidents still come from implementations, specs and institutions.

19.0 The bottom line

1. **The bottleneck in cybersecurity has shifted from "finding vulnerabilities" to "fixing them", and what drove the shift is AI's coding ability, not its math ability.** Anthropic's Glasswing program flagged 6,202 high-severity candidates; of the roughly 530 disclosed so far, only 75 have been patched (2026-05) [recent]. The only defense that can keep up is "proving that a class of vulnerabilities is absent"; proofs are getting cheaper, and the new bottlenecks are **whether the spec is right** and **whether the trusted computing base (TCB) has holes**.
2. **Cryptography is built on unproven conjectures, and AI math will audit them at scale.** Plenty of standards candidates have been broken by pure math: SIKE (2022, about an hour on a single core), Rainbow (2022, a weekend on a laptop), and the "proximity gaps" conjecture, which a batch of zkVMs relied on for their parameters and which was refuted in 2025 [established/recent]. The probability that ML-KEM suffers a practical classical break before 2030 is estimated at $\leq 5\%$ [speculation]; young post-quantum signatures and ZK parameter conjectures carry much higher risk.
3. **Math is pulling Q-day earlier.** The number of physical qubits needed to break RSA-2048 fell from about 20 million in 2019 to under 100,000 in 2026 (qLDPC architecture) [recent]. Resource estimation is a type A problem with an automatic judge, and AI will keep pushing this line down; hardware is still the main bottleneck.
4. **"What can quantum computing do besides breaking cryptography and simulating quantum systems?" is where math is most purely the bottleneck.** Practical exponential speedups are rare, "dequantization" keeps knocking claimed speedups back down, and AI-driven classical simulation is competing with quantum computing for the same applications.
5. **"Faster algorithms" splits into two worlds.** Asymptotic records (ω pushed down to 2.371177 in 2026-08 with AlphaEvolve's help, and to 9/4 by OpenAI's claim in §6) will hardly make GPUs any faster. Practical speedups come from "AI writes code + search +

benchmarks as the judge", for example AlphaEvolve continuously recovering 0.7% of Google's worldwide compute [established], and they rely on AI's other capabilities.

6. **Big problems like P vs NP are protected by "barrier theorems"; they are classic type B problems, and even a proof would barely change practice.** Only a few kinds of results would change the world: the unconditional existence of one-way functions, quantum algorithms for lattice problems, and impossibility theorems about "what can be audited efficiently" (for example, provably undetectable backdoors in models).
7. **Three connections to AI safety:** (a) the "box" around an AI agent is itself software, and verifiable isolation is the physical foundation of AI control; (b) ZK and FHE can "verify without revealing" and could serve as infrastructure for AI governance, but they are in tension with misuse monitoring; (c) cryptanalytic capability should be included in dangerous-capability evaluations, while mainstream frameworks currently seem to focus mainly on cyberattack operations ([speculation], not systematically verified).

19.1 Why these problems

Ranked by multiplying four criteria: **scale of impact**; **whether math is the bottleneck** (if not, say clearly what the real bottleneck is); **whether there is an automatic judge** (proof checkers, circuit simulators, benchmarks; type A problems with a judge move fast); **relevance to AI safety**.

#	Problem	Impact	Is math the bottleneck?	Judge	A/B	AI safety connection
1	Formal verification × offense and defense	Very high	It was, and is ceasing to be; specs and institutions are the new bottleneck	Yes	Proofs A; specs B	Sandboxes/AI control
2	Post-quantum lattice cryptography	Very high	Yes for confidence; no for incidents	Partial	Improving attacks A; new ideas B	Cryptanalysis as a dangerous capability
3	Q-day resource estimates	High	Partly: math is a multiplier, hardware is the main bottleneck	Yes	A	Disclosure of dangerous math
4	New quantum algorithms	Medium to high	Yes , in its purest form	Weak	Mostly B	Quantum attacks on lattice problems
5	FHE and ZK	Medium to high	Half: the conjectures are math, the constants are engineering	Yes	Proving conjectures A	Verifiable AI governance
6	Algorithm discovery	Medium in practice, low for records	Yes for records; no for practical speedups	Yes	A	AI optimizing its own training stack

#	Problem	Impact	Is math the bottleneck?	Judge	A/B	AI safety connection
7	Complexity theory	Low directly, high as a foundation	Yes, but type B	Weak	Big problems B; gadgets A	Impossibility of auditing; debate

Left out but worth mentioning: hardware verification, compiler verification (CompCert) and distributed protocol verification (TLA+) are folded into problem 1; SAT/SMT/MIP solvers rely mainly on engineering and are folded into problem 6; symmetric cryptography (AES, SHA-3) has a large security margin; for quantum key distribution (QKD) the bottleneck is physics and deployment; indistinguishability obfuscation (iO) is extremely far from practical; AI text watermarking already has impossibility results and belongs to §22.

19.2 Problem 1: Formal verification of critical software × cyber offense and defense in the AI era

In one sentence: AI makes "finding vulnerabilities" almost free, and the only defense that can keep up is "proving there are no vulnerabilities of this kind". Proofs are getting cheaper too, and the new bottlenecks become "what exactly is being proved (the spec)" and "what are we trusting (the TCB)".

Background: *Testing samples; proof covers every input.* Look at this C function:

```
int abs(int x) { return x < 0 ? -x : x; }
```

Random testing on 1,000 inputs passes every time. But when `x = INT_MIN`, `-x` overflows and the function returns a negative number; the chance that a random test hits this is $1/2^{32} \approx 2.3 \times 10^{-10}$. Write the spec as "the return value is ≥ 0 " and hand it to a proof checker, and the proof gets stuck on exactly this one input: either fix the code, or add the precondition "`x  $\neq$  INT_MIN`" to the spec.

Formal verification has three ingredients: a spec, an implementation, and a proof checked by a very small "kernel" (Lean, Isabelle, Coq/Rocq; see §3). You only have to trust a kernel of a few thousand lines, not the person or AI who wrote the proof. *The hard part is the spec:* the spec for sorting should be "the output is sorted, and is a permutation of the input"; if you only write "the output is sorted", then "always return an empty list" also passes the proof. The cost used to be very high: the seL4 microkernel has about 8,700 lines of C with about 200,000 lines of hand-written proof, about 25 person-years [established].

Why it matters: offense has already industrialized. At the DARPA AIXCC finals (2025-08), automated systems found 54 of 63 planted vulnerabilities, at an average of about \$152 per task [recent]. Defense cannot keep up: the Glasswing report says the work is now limited by

how fast it can verify, disclose and patch (paraphrased) [recent]. For AI safety, an agent's "box" is software: Glasswing's examples include memory-corruption bugs in a hypervisor **written in a memory-safe language**, and a chain of 4 vulnerabilities used to escape a browser sandbox. AI control schemes assume the box is secure, and verifiable isolation is the foundation of that assumption.

What's blocking it: is math the bottleneck? Proof labor used to be the main bottleneck, and it is disappearing. By curve ③, seL4's 200,000 lines of proof come to about \$4,600 in tokens, and about \$200 by end-2027 [model]; 10 AI agents have autonomously written a Lean version of zlib (lean-zip) and proved that "compressing and then decompressing any data smaller than 1 GiB gives back the original data" [recent]. The new bottlenecks are not math: (1) **specs**: fuzzing lean-zip with its proofs stripped out, for about 100 million executions, found zero memory issues in the proved parts, but the archive parser, for which no theorem was written, had a denial-of-service bug, and **the Lean 4 runtime itself** had a heap overflow [recent]; (2) **the TCB and hardware models**: Spectre shows that "isolation" proved against an instruction-set spec cannot stop side channels that the spec does not describe [established]; (3) **legacy code and institutions**: billions of lines of C/C++ are maintained by volunteers, and the "dumb" fix of switching languages works enormously well: Rust code has about 0.2 memory-safety vulnerabilities per million lines, versus about 1,000 historically for C/C++ [recent]; (4) **performance**: according to secondary reports, the Lean Zstd decompressor that Adam Langley had a model prove runs at only about 1/10 the speed of standard zstd △. Judgment: writing proofs is type A and is being conquered fast; writing specs, drawing the TCB boundary and driving adoption are type B problems plus institutional ones.

What AI can do: *AI math/proof*: write proofs against a spec; repair proofs automatically after code changes; prove old and new implementations equivalent (for example, C translated into Rust); the repository-level benchmark Vero also has agents **prove that a given spec is unsatisfiable**, which amounts to auditing specs by machine. *AI's other capabilities*: fuzzing, building exploits, writing patches; mining candidate specs from RFCs and tests; migrating C to Rust at scale.

If solved: first-order impact: cryptographic libraries, parsers, kernels, hypervisors and consensus protocols ship with proofs by default; memory-safety and parsing bugs in them drop to near zero, and the cost shifts to spec review. Distributed protocol models written in TLA+ can be proved to refine all the way down to the actual code.

If solved: second-order impact: (1) **The offense–defense balance follows a "worse before better" curve** [speculation]: 2026–2028 favors attackers, with patching at human speed; once critical components are proved, the attack surface moves to spec errors,

configuration, the supply chain, social engineering, hardware side channels and ML components that cannot be verified. (2) **Software liability becomes law**: the EU Cyber Resilience Act (Regulation (EU) 2024/2847) has already introduced security liability for digital products [established], and "is there a proof" may become a due-diligence standard [speculation]. (3) **AI safety**: an AI that writes proofs can also write **specs**, and a spec that looks reasonable but hides a backdoor is very hard for humans to spot, so "who reviews the specs" becomes an alignment problem; meanwhile proof checkers (with fixed statements, and **sorry** and new axioms banned) are **reward signals that are very hard to reward-hack**, which will speed up AI coding ability, so they are dual-use in their own right.

Timing: Optimistic: by mid-2027 (curve ① at about 150 hours, roughly one work-month) AI can autonomously translate a C library of 10,000–50,000 lines into a version with proofs. Central: in 2027–2028 proof labor stops being the main cost (curve ③), while spec review still runs at human speed; by 2030 proved code covers only "critical small kernels", under 1% of all code [speculation]; the window favoring attackers lasts until about 2028. Pessimistic: specs, performance and maintainer bandwidth all get stuck at once, and in 2030 this is still only an internal practice at a few big companies.

Counterarguments and uncertainty: most intrusions come from credentials, phishing and misconfiguration (Verizon DBIR 2024 says about 68% involve "the human element" △ from memory), and proofs do not help there. Glasswing's numbers come mainly from Anthropic's own reports, with limited independent evaluation.

Signals to watch: patch rates in later Glasswing reports; the share of memory-safety bugs among CVEs; core open-source projects requiring "PRs with proofs"; the first public incident of "a backdoor in the spec itself".

Sources: seL4 statistics <https://trustworthy.systems/projects/OLD/l4.verified/numbers> ; DARPA AIXCC <https://www.darpa.mil/news/2025/aixcc-results> ; Glasswing <https://www.anthropic.com/news/glasswing-initial-update> ; arXiv:2608.13522 (Vero), 1801.01203 (Spectre); lean-zip <https://kirancodes.me/posts/log-who-watches-the-watchers.html> ; Rust in Android <https://blog.google/security/rust-in-android-move-fast-fix-things/> ; secondary summary of Langley △ <https://www.remio.ai/post/adam-langley-says-lean-proof-automation-is-here-the-hard-part-just-moved> .

19.3 Problem 2: Post-quantum lattice cryptography (ML-KEM) – will AI find a break, or build confidence?

In one sentence: the internet is moving key exchange to ML-KEM, whose security rests on the **conjecture** that "the Learning With Errors problem (LWE) is hard". AI math will audit it at an unprecedented scale, and the outcome could be a break (catastrophic, but with hybrid schemes as a backstop), higher confidence, or a small reduction in security bits.

Background: Shor's quantum algorithm can quickly break RSA and elliptic-curve cryptography (ECC), so new hard problems are needed. *Lattices*: in the plane, take the basis $b_1=(2,0)$, $b_2=(1,2)$; all integer combinations form a regular grid of points. The same lattice can also be described by a "bad" basis, for example b_1 and $b_2'=(11,2)$ ($= b_2 + 5 \cdot b_1$): exactly the same points, but the two vectors are long and almost parallel. **The shortest vector problem (SVP)**: given only the bad basis, find the shortest nonzero lattice vector. In two dimensions this is fast; in a few hundred dimensions, the best known algorithms all take exponential time. A good basis is the "trapdoor".

LWE is a system of linear equations with noise. Working mod 17, let the secret be $s = (3, 5)$:

a	a·s mod 17	noise e	public value b
(1, 2)	13	+1	14
(4, 1)	0	−1	16
(2, 7)	7	0	7

Without noise, Gaussian elimination solves for s immediately; with ± 1 noise, every equation is "slightly off". In two dimensions you can simply enumerate the $17^2 = 289$ candidates, but ML-KEM-768 has effective dimension 768 and modulus 3329 (FIPS 203), about 182 bits of classical security by the Core-SVP estimate $\triangle$ (from memory). **The security of the whole scheme rests on that little bit of noise.** Regev (2005) proved that "solving random LWE $\Rightarrow$ solving worst-case lattice problems", but the reduction is not tight and cannot fix concrete parameters; parameters are set by "how much the best known attack costs".

Past breaks are the most important background for this card. **SIKE** (a NIST round-4 candidate) was broken in 2022-07 by Castryck–Decru using a pure-math theorem from 1997, in about an hour on a single core [established]. **Rainbow** (a round-3 finalist signature scheme): the title of Beullens's 2022 paper is literally "Breaking Rainbow Takes a Weekend on a Laptop" [established]. **Chen 2024** claimed a polynomial-time quantum algorithm for LWE; about a week later an error was found and the claim was withdrawn [established]. A **Luo 2026** preprint claims to break ML-KEM; it has been informally shown to get the dimensions wrong, and **does not currently constitute a break** $\triangle$.

Why it matters: by late 2025-10, more than half of the human TLS traffic Cloudflare received already used the hybrid key exchange X25519MLKEM768 [recent]. "Harvest now, decrypt later" means that for data that must stay secret for 10–30 years, the question of whether ML-KEM is secure has to be answered **now**. Signatures are harder: switching entire certificate chains to ML-DSA-44 adds about 15 kB to a TLS handshake, and the smaller schemes rest on **younger, less-audited math**, which is exactly where AI cryptanalysis is most likely to make its first impact.

What's blocking it: is math the bottleneck? For **confidence**, math is the bottleneck: nobody can prove LWE is hard, confidence comes only from "many smart people have tried and failed", and expert attention is scarce. For **real incidents**, the bottleneck is implementation: KyberSlash found that several implementations, including the official reference code, had secret-dependent division timing, enough to "reliably recover the key within minutes" [established]. For **migration speed**, the bottleneck is engineering and institutions: in 2025 only 3.7% of Cloudflare's origin servers supported post-quantum key exchange.

What AI can do: *AI math/proof:* systematically extend known attack families (for example, pushing attacks on rank-1 ideal lattices toward rank-2 and rank-3 module lattices), a well-defined direction and type A; act as a judge at scale, formalizing claimed breaks like Chen's and Luo's and finding the errors within days; audit the attack-cost estimates themselves (Ducas–Pulles 2023 showed problems with the heuristic assumptions behind dual attacks [established]). *AI's other capabilities:* Meta's SALSA series uses Transformers to recover LWE secrets, but it only works for sparse or very small secrets, not ML-KEM's parameters [recent]; checking implementations for constant-time behavior at scale.

If solved: first-order impact: *Scenario A (low probability):* AI finds a practical classical break of ML-KEM. Hybrid schemes fall back to X25519's security, but all "harvest now, decrypt later" traffic becomes exposed to future quantum attackers; ML-KEM-only deployments (such as the CNSA 2.0 path Δ) take a catastrophic hit, and the world has to switch urgently to code-based schemes (in 2025-03 NIST chose HQC as a backup). *Scenario B:* the search comes up empty, confidence rises, and the default parameters might even drop from 768 to 512. *Scenario C (most likely):* each paper shaves off a few bits, and together they force parameters upward.

If solved: second-order impact: fewer and fewer human experts can check AI cryptanalysis, so trust in standards will rely more and more on **machine-checkable arguments**. States with private breaking capabilities will not disclose them, so "how to disclose dangerous mathematical results" needs new norms (problem 3). On AI safety: cryptanalytic capability should become a category of dangerous-capability evaluation [speculation]; protecting model weights presupposes cryptography, so if cryptography fails, the chain of trust in AI governance breaks at the very bottom; a flood of wrong AI-generated "breaks" would amount to a "denial-of-service attack" on expert attention.

Timing: papers in which AI helps shave off security bits (type A, with a clear objective function, following curve ②) become common from **2027** in the central scenario. At least one NIST additional signature candidate significantly weakened or broken thanks to AI-assisted analysis: central **2027–2029** [speculation]. A practical classical break of ML-KEM–

768: **$\leq 5\%$ by 2030** [speculation]; "is LWE hard" is a research program rather than a short proof, so it is type B and outside what curve ② covers.

Counterarguments and uncertainty: lattice problems have been studied for more than 40 years counting from LLL (1982), and AI may only carry out old ideas faster. Bernstein (a strong advocate of hybrid schemes), on the other hand, argues that the risk is systematically underestimated, because "known attacks" are only a small part of the attack surface.

Signals to watch: cryptanalysis papers on ePrint that acknowledge AI; downward revisions to the lattice-estimator's estimates; frontier AI companies adding "cryptanalytic uplift" evaluations to their safety frameworks.

Sources: ePrint 2022/975 (SIKE), 2022/214 (Rainbow), 2024/555 (Chen), 2024/1049 (KyberSlash), 2025/1002 (SALSA comparison), 2023/302 (Ducas–Pulles); Bernstein <https://blog.cr.yp.to/20260630-risk.html> ; Luo analysis <https://postquantum.com/security-pqc/luo-ml-kem-quantum-attack-analysis/> ; Cloudflare <https://blog.cloudflare.com/pq-2025/> ; FIPS 203 <https://csrc.nist.gov/pubs/fips/203/final> .

19.4 Problem 3: The Q-day countdown – how many qubits does it take to break RSA/ECC?

In one sentence: Q-day depends on two numbers: how fast hardware improves, and "how many qubits and how many gates a break takes". The second number is set entirely by math (algorithms, error-correcting codes, architecture), and it fell about 200-fold in 7 years.

Background: physical qubits make an error about 0.1% of the time per gate, and breaking RSA-2048 takes billions of gates, so the error rate per step has to be pushed below about 10^{-12} . The way to do this is **error correction**: encode one "logical qubit" in many physical qubits. The classical analogy is the repetition code: send each bit three times and take a majority vote; if a single copy flips with probability 1%, the vote is wrong about $3p^2 = 0.03\%$ of the time. The most common quantum version is the **surface code**: on a $d \times d$ grid you repeatedly measure "syndromes", and a classical computer infers in real time where errors occurred (**decoding**). Each time the distance d increases by 2, the logical error rate is divided by a **suppression factor Λ** .

An order-of-magnitude calculation by hand. Google Willow has $\Lambda = 2.14$ [recent]. At that rate, pushing errors down another 10^9 -fold takes $\log(10^9)/\log(2.14) \approx 27$ steps of "+2", i.e. $d \approx 61$, or about $2d^2 \approx 7,400$ physical qubits per logical qubit; with $\Lambda = 10$ it takes only 9 steps, $d \approx 25$, about 1,250. The "uniform 0.1% gate error" commonly used in resource estimates corresponds roughly to $\Lambda \approx 10$ [estimate], so **today's hardware is still about 5 times short of this assumption**. Better codes (qLDPC) save a lot: IBM's "gross code" protects

12 logical qubits with 288 physical qubits, where the surface code would need nearly 3,000 [established].

Year	Target	Physical qubits	Time	Source
2019	RSA-2048	~20 million	~8 hours	Gidney & Ekerå
2025	RSA-2048	<1 million	<1 week	Gidney, same hardware assumptions
2026-02	RSA-2048	<100,000	Not given in the abstract	Pinnacle architecture (qLDPC)
2026-03	ECC P-256	<500,000 (superconducting)	A few minutes	Google: <1,200 logical qubits, about 90 million Toffoli gates
2026-03	ECC P-256	~10,000–26,000 (neutral atoms)	A few days	Cain et al.

The 2025 drop came **entirely** from algorithms and architecture (approximate residue arithmetic, "yoked" surface codes, "magic state cultivation") [recent]. **Elliptic curves are exposed earlier than RSA**, and TLS signatures, Bitcoin and Ethereum all use ECC.

Why it matters: Mosca's inequality: if data must stay secret for x years, migration takes y years, and Q-day is z years away, then **as soon as $x + y > z$, it is already too late**. When Google published its ECC estimate, it set its own migration target at 2029, and pointed out that a fast-clocked superconducting machine could compute the private key from a public key exposed in the mempool before a Bitcoin transaction is confirmed [recent].

What's blocking it: is math the bottleneck? The main bottleneck is hardware; math is a multiplier. "100,000 qubits running stably for days at a 0.1% error rate, with real-time decoding at microsecond scale" is still far from achieved. IBM plans to deliver Starling in 2029 (200 logical qubits, 100 million logical operations); ECC-256 needs about 100 million Toffoli gates and about 1,200 logical qubits, so **the gate count is already in the same range, while logical qubits are still about 6 times short** [estimate]. **Decoding is an underrated bottleneck:** AlphaQubit is more accurate, but Google acknowledges it is still too slow to keep up in real time [established]. How big is the multiplier [speculation]? If the number of usable qubits doubles every year Δ (a rough assumption), each 10-fold drop in resource requirements brings Q-day about 3.3 years closer.

What AI can do: *AI math/proof:* resource estimation is classic type A, with a clear objective function (Toffoli count, logical qubit count) and an automatic judge (stabilizer simulators): searching for new, implementable qLDPC code families, optimizing modular arithmetic circuits, designing cheaper magic-state protocols. AlphaEvolve-style "search + evaluator"

is a perfect fit. *AI's other capabilities*: neural-network decoders, automatic calibration and chip design, which directly ease the hardware bottleneck.

If solved: first-order impact: Q-day estimates move earlier and their ranges narrow, giving policymakers harder numbers; fault-tolerant machines also become usable earlier for chemistry and materials simulation (§20).

If solved: second-order impact: *"Algorithmic Q-day"*: the date may be set by a paper rather than by a machine, so migration plans should be stress-tested against "resources falling another 10-fold before 2030" [speculation]. *Disclosure norms for dangerous math*: Google did not publish its ECC attack circuit; instead it attached a zero-knowledge proof generated with the SP1 zkVM + Groth16, letting third parties confirm that "such a circuit exists and the resource counts are correct"; ironically, the curve Groth16 uses is itself vulnerable to quantum attack [recent]. This template could carry over to other dangerous results that AI discovers, such as exploit chains and biological sequences. *Cryptocurrency governance*: whether "dormant coins" with exposed public keys should be frozen or left alone.

Timing: resource estimates (type A) fall another 2–10× before 2028 in the central scenario [speculation] (curve ©). A machine that can break ECC-256 [speculation]: optimistic 2030–2032, central 2033–2036, pessimistic after 2040. The basis: IBM's 2029 roadmap is in roughly the same range as the ECC resource requirements, and Google's 2029 migration target implies a risk assessment; this item depends mostly on hardware and is only weakly tied to the capability curves. RSA-2048 comes a few years later (secondary △).

Counterarguments and uncertainty: hardware may get stuck on cooling, wiring, yield and correlated noise (cosmic rays, leakage), whereas resource estimates all assume noise is "uniform and independent". A few researchers (such as Gil Kalai) argue that noise fundamentally prevents scaling up; this has not been refuted by experiment.

Signals to watch: $\Lambda \geq 4$; real-time decoding at distance ≥ 11 ; more than 1,000 logical qubits; resource estimates falling another 10-fold. One warning sign: a major company suddenly stopping publication of resource estimates.

Sources: arXiv:1905.09749, 2505.15917 (Gidney), 2602.11457 (Pinnacle), 2603.28846 (Google ECDLP), 2603.28627 (Cain et al.), 2408.13687 (Willow), 2308.07915 (gross code); Google disclosure blog <https://research.google/blog/safeguarding-cryptocurrency-by-disclosing-quantum-vulnerabilities-responsibly/> ; AlphaQubit <https://blog.google/technology/google-deepmind/alphaqubit-quantum-error-correction/> ; IBM Starling <https://newsroom.ibm.com/2025-06-10-IBM-Sets-the-Course-to-Build-Worlds-First-Large-Scale,-Fault-Tolerant-Quantum-Computer-at-New-IBM-Quantum->

Data-Center ; TQI overview (with details of the ZK disclosure)

<https://thequantuminsider.com/2026/03/31/q-day-just-got-closer-three-papers-in-three-months-are-rewriting-the-quantum-threat-timeline/> .

19.5 Problem 4: New quantum algorithms and provable quantum advantage (short card)

Quantum speedups come from **interference** of amplitudes, like noise-canceling headphones: the waves of wrong answers cancel each other out. This requires the problem to have **structure** (such as the periodicity Shor exploits); unstructured search gets only Grover's quadratic speedup (10^{12} steps become 10^6), and every logical step carries error-correction overhead, so quadratic speedups will mostly be useless in the fault-tolerant era [established]. Math can also destroy speedups: in 2018, 18-year-old Ewin Tang used a classical algorithm to knock the exponential speedup of a quantum recommendation system back down ("dequantization") [established]; Lee et al. concluded that evidence for an exponential quantum advantage in ground-state energy calculation has yet to be found across chemical space (paraphrased) [established].

Math is the bottleneck, in its purest form in this field, and mostly type B: it isn't even clear "which class of problems to look in", and classical baselines (tensor networks, neural-network quantum states) keep raising the bar. AI math is already involved: Aaronson and Witteveen (2025) proved limits on black-box amplification for QMA, and a key step was a function proposed by GPT-5 [recent]. The counterintuitive part comes from AI's other capabilities: AI-driven classical simulation is taking over quantum computing's chemistry and materials applications, so the stronger AI gets, the fewer "killer apps" quantum computing may have [speculation] (§20).

Timing: concrete type A proofs are happening now; a whole new family of exponential speedups found mainly by AI: about 15–30% before 2030 [speculation], and curve ② does not apply. **An efficient quantum algorithm for lattice problems (LWE): <5% before 2030 [speculation], but with the highest impact weight:** if it happens, the entire post-quantum migration of problem 2 has to start over. **Sources:** Tang arXiv:1807.04271; quantum chemistry arXiv:2208.02199; Aaronson & Witteveen arXiv:2509.21131.

19.6 Problem 5: FHE and zero-knowledge proofs – making "computing on ciphertext" and "verifiable computation" cheap

In one sentence: fully homomorphic encryption (FHE) can compute on data without decrypting it, and zero-knowledge proofs (ZK) can prove "the computation was done correctly" without revealing the inputs. Both already exist mathematically, but their overhead is thousands to tens of millions of times that of plaintext computation (order-of-

magnitude estimate $\triangle$), and efficiency depends half on math and half on engineering. They are the most promising tools for "verifiable AI governance".

Background: *FHE is like a glove box*: you work through the walls of the box without seeing inside. A toy example: $\text{Enc}(m) = m + k$, so $\text{Enc}(3) + \text{Enc}(4) = 7 + 2k$; whoever holds k subtracts $2k$ to get 7, and the server has done the addition without knowing 3 or 4. Real FHE is built on **LWE**, **the same hard problem ML-KEM relies on**; operations make the noise grow, so it has to be refreshed periodically by "bootstrapping", the most expensive step. *ZK is like a random spot check*: the prover encodes the computation as polynomials. Two different polynomials of degree d agree on at most d points (the Schwartz–Zippel lemma), so the verifier picks one random point in a field of size 2^{64} , and the chance of being fooled is at most $d/2^{64}$; for $d = 10^6$ that is about 5×10^{-14} .

Soundness depends on a conjecture. FRI-type protocols rely on "**proximity gaps**" for Reed–Solomon codes: if even one vector in a batch is far from every valid codeword, then a random linear combination of the batch is almost always far as well. Many systems set their parameters by the **conjecture** that "this property holds all the way up to the capacity limit", in exchange for smaller proofs. In 2025-11, Crites and Stewart proved that conjectures of this kind **are false** [recent]; according to community tallies, around twenty zkVMs had assumed it $\triangle$.

Why it matters: SP1 Hypercube used about 160 RTX 4090s to prove 93% of Ethereum mainnet blocks within 12 seconds, with progress coming mainly from engineering; zkLLM can produce a proof under 200 kB for one inference of a 13-billion-parameter model in 15 minutes [recent]. Apple already uses FHE for private queries [established].

What's blocking it: is math the bottleneck? Half and half. Math sets parameter sizes (soundness theorems) and brings new protocols and cheaper bootstrapping; engineering handles hardware acceleration and memory bandwidth. There is also a spec problem: a SoK analyzed 141 real SNARK vulnerabilities, and many came from "underconstrained" circuits $\triangle$, which has the same structure as problem 1.

What AI can do: *AI math/proof*: proximity gaps are a type A problem that is **clearly stated and has an automatic judge**. In 2026-08 the Ethereum Foundation launched the better.codes challenge: the problems are formalized in Lean 4, the goal is to push the provable soundness lower bound to 128 bits, and it explicitly invites people to "bring your own AI agent" [recent]. This is the cleanest test bed for "AI math directly changing industrial parameters". *AI's other capabilities*: writing GPU kernels, compiling programs into circuits, checking whether constraints are complete.

If solved: first-order impact: with general-purpose ZK overhead down to 10–100×, you get a "verifiable cloud"; with FHE down to 10–100× on specific workloads, medical and financial data can be used for inference without leaving its home domain.

If solved: second-order impact: *Verifiable AI governance* [speculation]: proving that "this output came from the model whose weights hash to H", that "the model passed evaluation set S" (without revealing the weights), or that "inference really did go through the safety filter" follows the same logic as arms-control verification: **verify, without disclosing**; proving "how much compute was used in training" is much harder, and there is no workable scheme yet. *The tension between privacy and oversight*: private inference with FHE means the provider **cannot see** what users are asking, so misuse monitoring stops working; the more private it is, the harder misuse is to prevent [speculation]. *Mathematical risk in finance*: billions of dollars of rollup assets rest partly on unproven coding-theory conjectures, and AI will turn "conjecture-driven parameters" into "proof-driven parameters".

Timing: provable bounds on proximity gaps, or better.codes reaching 128 bits: type A with a Lean judge, substantive progress in 2027–2028 in the central scenario (curve ②) [speculation]. Real-time proving of Ethereum blocks: already essentially achieved in engineering terms. FHE inference overhead for frontier-scale LLMs falling below 100×: after 2030 in the central scenario [speculation], depending on hardware and largely unrelated to AI math.

Counterarguments and uncertainty: trusted execution environments (TEEs, such as the H100's confidential computing mode) are far cheaper and may take over most uses of FHE/ZK, at the cost of trusting the hardware vendor and accepting side-channel risk.

Signals to watch: provable security bits on the better.codes leaderboard; whether frontier-scale models can produce inference proofs; whether regulatory documents list "cryptographic proofs" as compliance evidence.

Sources: ePrint 2025/2046, 2026/782 (proximity gap counterexamples); better.codes <https://blog.ethereum.org/2026/08/20/better-codes-challenge> ; SP1 Hypercube <https://blog.succinct.xyz/sp1-hypercube/> ; Apple HE <https://machinelearning.apple.com/research/homomorphic-encryption> ; arXiv:2404.16109 (zkLLM), 2402.15293 (SoK on SNARK vulnerabilities).

19.7 Problem 6: Algorithm discovery – asymptotic records vs real speedups on GPUs (short card)

Big-O notation hides constants; an algorithm that is asymptotically faster but has constants of cosmic size is called a "galactic algorithm". The matrix multiplication exponent ω fell to

2.371177 in 2026-08 with AlphaEvolve's help [recent]. Working out by hand what that is worth: the exponent dropped by 0.000162, and at $n = 10^5$, $n^{0.000162} \approx 1.0019$, a saving of only 0.19% of operations; and laser-method algorithms cannot be implemented in practice anyway. OpenAI's claimed $\omega \leq 9/4$ in §6 is a much bigger step, but it too is an existence proof and a galactic algorithm, and §6.10.2 already explained that it has no direct impact on GPUs or LLM training: what GPUs run on their tensor cores is still the n^3 algorithm. What is close to practical are decompositions for **fixed small sizes** (AlphaEvolve's 48 multiplications for 4×4 complex matrices; a follow-up paper tries to "bring it into practice", but its abstract reports no measured speedup). Shortest paths broke the sorting barrier for the first time in 2025, but at $n = 10^9$ the new algorithm is in theory only about 3 times faster than Dijkstra, a gain that constants and caches will likely eat up [recent].

For asymptotic records the bottleneck is math; for practical speedups it is not: there the bottleneck is understanding the hardware (memory hierarchy, bandwidth), reliable benchmarking and correctness checking. Practical speedups come from AI's other capabilities, "write code + evolutionary search + benchmarks as the judge": one Borg scheduling rule from AlphaEvolve continuously recovers 0.7% of Google's worldwide compute, and one matrix multiplication kernel for Gemini got 23% faster, cutting training time by about 1% [established]. High-performance code written by AI risks gaming the evaluation (related incidents in 2025 $\triangle$ from memory), so the single most useful thing AI math can do here is **prove that an optimized kernel is equivalent to the reference implementation**.

Second-order impact: making Gemini training 1% faster is **AI optimizing AI's own training stack**, the plainest form of recursive self-improvement, and one that has already happened; progress in algorithmic efficiency will weaken regulation that uses compute as its threshold [speculation] (§23). Timing: kernel-level optimization is already happening; tasks at the level of compilers or scheduling systems become feasible around 2027 in the central scenario (curve ①); asymptotic records will see many small AI-assisted improvements in 2026–2028 (curve ②). **Sources:** AlphaEvolve <https://deepmind.google/discover/blog/alphaevolve-a-gemini-powered-coding-agent-for-designing-advanced-algorithms/> ; arXiv:2608.16884 ($\omega < 2.371177$), 2506.13242 and 2609.12027 (4×4), 2504.17033 (shortest paths).

19.8 Problem 7: Complexity theory – which proofs would change the world, and which would not

In one sentence: big problems like P vs NP are blocked by "barrier theorems" and are type B; AI can already produce type A results in complexity theory. A few target theorems would change the foundations of cryptography and AI safety; most would change no practice at all.

Background: P is the class of problems that can be **solved** quickly; NP is the class whose solutions can be **verified** quickly (Sudoku: checking a filled-in grid is easy, but solving the generalization to $n^2 \times n^2$ grids is NP-complete). *Why proving $P \neq NP$ would not help practice:* $P \neq NP$ means "hard **in the worst case**", while cryptography needs "hard **on average**" plus the existence of one-way functions. To use a lock analogy: worst-case hard means "there exist some people who can't open this lock"; average-case hard means "anyone you pick at random can't open it", and only the latter protects data. Liu and Pass (2020) proved that one-way functions exist **if and only if** time-bounded Kolmogorov complexity is somewhat hard on average [established].

Barriers. Three theorems, relativization, natural proofs (Razborov–Rudich 1997) and algebrization, prove that certain classes of techniques can never prove $P \neq NP$ [established]. The most elegant is natural proofs: **if strong pseudorandom functions exist (that is, if cryptography is possible), "natural" methods cannot prove circuit lower bounds**; the very existence of cryptography is the reason proving hardness is so hard. For explicit Boolean functions, the best known circuit lower bound is only about $3.1n$ (linear), while the true lower bound is believed to be exponential $\triangle$.

Why it matters: the practical impact of different results varies enormously:

Result	Practical impact
Proving $P \neq NP$	Almost none: it confirms what is already believed; the value is in the new techniques
$P = NP$ with a practical algorithm	Civilization-scale, but extremely unlikely
Unconditional existence of one-way functions	Cryptography gets an unconditional foundation; actual schemes still rely on specific assumptions
$BPP = P$; proving or refuting the UGC	Small: clarifies which approximation algorithms are optimal (UGC: see §7)
Impossibility theorems about "what can be audited efficiently"	Large impact on AI safety

An example for the last row: Goldwasser et al. (2022) proved that backdoors can be planted in models that **no computationally bounded observer can detect**: a construction based on digital signatures is undetectable in the black-box setting, and for models such as random ReLU networks it is undetectable even in the white-box setting [established]. So **inspecting the weights alone cannot, in principle, rule out all backdoors**, and safety guarantees must rest partly on provenance proofs for the training process (linking back to problem 5). Another root: the theoretical motivation for debate (AI safety via debate) is exactly $IP = PSPACE$, i.e. a weak verifier can check a strong prover through interaction; characterizing

"when computationally bounded human judges remain reliable" amounts to drawing the boundary of scalable oversight (§22).

What's blocking it: is math the bottleneck? It is pure math, and the big problems are type B: known techniques have been proven insufficient, and entirely new ideas are needed. Compute helps only when searching over finite structures (gadgets, small circuits).

What AI can do: *AI math/proof:* it has already produced type A results. Nagda, Raghavan and Thakurta (2025) used AlphaEvolve to construct gadgets and proved that "approximating MAX-4-CUT to within 0.987 is NP-hard" (previous record 0.9883) [recent]. AI can also formalize complexity theory (the complexity library in Lean is very thin, an infrastructure gap). *AI's other capabilities:* exhaustively finding minimal circuits for small functions, generating conjectures.

If solved: first-order impact: if AI finds new techniques that get around the barriers, they will almost certainly spill over into cryptography (new constructions or new attacks) and learning theory.

If solved: second-order impact: complexity theory will become a language for **understanding AI itself**. "Verifying is easier than generating" is the foundation of the whole safety paradigm of debate, scalable oversight, and "AI writes the proof, the kernel checks it"; once a theorem reveals cases where verification is also hard, these schemes will have to be redesigned. Impossibility theorems (backdoors are undetectable, watermarks are removable) will push AI governance from "detection" toward "provenance proofs and process control".

Timing: gadgets and approximation bounds (type A) are happening now and will appear in large numbers in 2026–2028 (curve ②). New impossibility/possibility theorems about AI auditing and verification: in the central scenario, a batch of them proved with AI involvement in 2027–2029 [speculation]. P vs NP, and the unconditional existence of one-way functions: each below 2% before 2035 [speculation]; they are not in the pool of "formalizable problems of known shape" that curve ② covers.

Counterarguments and uncertainty: AI might find "alien" techniques that humans would not think of, and the barriers rule out only **certain classes** of techniques; it is also possible that these big problems are independent of the standard axioms.

Signals to watch: the share of FOCS/STOC papers that acknowledge AI; the growth of Lean's complexity library; circuit lower bounds for explicit functions breaking past linear.

Sources: arXiv:2509.18057 (Nagda et al.), 2009.11514 (Liu–Pass), 2204.06974 (undetectable backdoors), 1805.00899 (debate); Razborov & Rudich, JCSS 1997.

19.9 Second-order effects in this field

1. "Proofs are cheap, statements are expensive" recurs across all seven problems, just under a different name each time.

Problem	The cheap part (type A)	The expensive part (type B)
1 Formal verification	Writing proofs	Writing specs, drawing the TCB boundary
2 Lattice cryptography	Improving existing attacks, formalizing reductions	What the security model leaves out (timing, implementation)
3 Q-day	Optimizing circuits and codes	Hardware, and realistic noise models
4 Quantum algorithms	Proving lower bounds, checking specific speedups	Which class of problems to look in
5 FHE/ZK	Proving or refuting proximity gaps	Whether the circuit constraints are complete
6 Algorithms	Improving asymptotic bounds	What counts as "really fast"
7 Complexity	Constructing gadgets	New ideas that get around the barriers

2. **Asymmetric speeds are the most dangerous dynamic in this field.** Finding vulnerabilities is fast, fixing them is slow; discovering attacks is fast, migrating standards is slow; AI generating "break claims" is fast, expert checking is slow; resource estimates fall fast, hardware moves slowly. So 2026–2029 is likely to be a **fragile transitional window** [speculation]: attackers reap AI's dividends first, while defenders' dividends (verified rewrites, post-quantum migration, proof-driven parameters) have to wait for institutions to catch up.

3. **The trust stack will be rebuilt on machine-checkable artifacts.** At the bottom are verified cryptographic libraries and kernels, in the middle post-quantum key exchange and smaller signatures, and at the top ZK-proved computation. Every layer has a "who checks the checker" problem: the Lean runtime bug in lean-zip and Groth16's own vulnerability to quantum attack are both reminders.

4. **Mathematical results are starting to need "responsible disclosure".** Google disclosed a quantum attack estimate with a ZK proof; Anthropic first publishes hashes of vulnerability details and releases the details only after fixes. These two practices will become standard moves when AI discovers dangerous results [speculation].

5. **Summary of the connections to AI safety.** Verifiable sandboxes (problem 1); cryptanalysis and vulnerability discovery included in dangerous-capability evaluations (problems 1, 2); governance tools that "verify without disclosing", and their tension with misuse monitoring (problem 5); undetectable backdoors putting a ceiling on model auditing (problem 7); algorithmic efficiency weakening regulation that uses compute as its threshold

(problem 6); proof checkers as reward signals that are hard to hack, speeding up AI coding ability (problem 1).

6. Connections to other fields. Quantum computing and AI-driven classical simulation compete for the same applications (§20); the security of cryptocurrencies depends on coding-theory conjectures and on how long ECC lasts; software liability is moving into law (§23); private cryptanalytic and quantum capabilities are asymmetric weapons in geopolitics. An analogy from neuroscience [speculation, not a strict correspondence]: the brain uses slower prefrontal control to constrain fast association; computer security in the AI era is also moving toward "a powerful but untrusted generator + a small, trusted checker".

19.10 Sources for this section

Full citations are in the **Sources** line at the end of each problem. arXiv IDs →

<https://arxiv.org/abs/<ID>>, ePrint IDs → <https://eprint.iacr.org/<ID>>.

Unverified or from-memory items (△): the Core-SVP security bits; the resource multiple of RSA-2048 relative to ECC (secondary); the order of magnitude of FHE overhead; the share of SNARK vulnerabilities at the circuit layer; the number of zkVMs relying on the proximity gaps conjecture; the $3.1n$ circuit lower bound (Li–Yang, STOC 2022); Verizon DBIR's 68%; the details of CNSA 2.0 using only ML-KEM; the details of Langley's experiment (secondary); the incident of AI-generated CUDA kernels with exaggerated speedups; the assumption that "usable qubits double every year".

PART II · OUTLOOK · PHYSICS

20. Physics, chemistry and materials: the equations are known; the hard part is computing, measuring and building

Tags: [established] accepted results; [recent] new results from 2024–2026; [speculation] inference; △ unverified or disputed. Timing calls tie into the three curves of §18: **curve ①** METR task length (central: about 150 h by mid-2027, about 400 h by end of 2027); **curve ②** single-attempt solve rate on formalizable open problems (central: about 18% by end of 2027, about 45% by end of 2028, about 77% by end of 2030); **curve ③** cost of formal proof (falling about 10× per year). **Type A** = statement clear, proof missing; **type B** = not even clear what to prove. Problems 4, 6 and 8 are short cards.

20.0 The bottom line

- **Most of the equations have long been known.** In 1929 Dirac said that the physical laws needed for chemistry are "completely known", and the only trouble is that the equations can't be solved [established]. The hard parts are: **can't compute** (the cost of the many-electron equation grows exponentially), **can't measure precisely** (synthesis and characterization are ambiguous), and **can't build** (getting from the lab to engineering). So AI math acts here mainly through **algorithms**, not through **theorems**.
- **Only three kinds of math can change the direction of physics:** impossibility theorems (e.g. Schuch–Verstraete for DFT), existence results and counterexamples (e.g. the 2026-09 counterexample to Grad's conjecture), and computer-assisted proofs (e.g. Euler blowup). Most of the rest are **model-health proofs**, which confirm what physicists already believed.
- **The type A list in mathematical physics is being cleared in bulk.** This batch of OpenAI manuscripts contains about 25 mathematical-physics families (secondhand count Δ). They claim to have solved Anderson extended states, spontaneous magnetization of the Heisenberg ferromagnet, the Haldane gap and hard-sphere BEC, plus the forced NS of §11 and the Vlasov–Maxwell of §12. None has been reviewed. Judgment [speculation]: problems of this kind will be mostly swept up in 2027–2028, yet they will change almost no experiment or engineering.
- **The high-impact problems are stuck outside mathematics:** electronic structure, high-temperature superconductivity, turbulence closure, fusion control, cloud feedback. What helps is AI's **other capabilities**: neural network functionals, ML interatomic potentials, differentiable simulators, reinforcement learning control (tested on the TCV and DIII-D devices), learned weather models (ECMWF's AIFS went operational in 2025-02), and autonomous labs.
- **Candidates are cheap and verification is expensive; how fast AI speeds a field up depends on how fast its referee is.** GNoME predicted about 380,000 stable crystals, and outside groups have synthesized only 736. This is the same bottleneck shift as in mathematics, where "proofs are cheap, auditing statements is expensive". Weather is scored every day and has been rewritten by AI; climate is scored once every few decades and moves slowly.
- **Energy impact arrives slowly.** Going from lab to grid usually takes 10–20 years. Before 2030 the only things likely to land are: weather forecasting (already happened), incremental improvement of batteries and catalysts, and faster control and design of fusion devices.
- **For AI safety:** physics and chemistry have automatic referees, so they will be where autonomous AI research gets going first and where AI first gets to act in the physical world. They also form a feedback loop, "AI → energy, chips, materials → more compute". Chemistry's dual-use risk is the most direct. There are ready-made examples of

specification gaming (forced NS; DFT functionals whose energies fit better and better while their densities drift further and further off).

20.1 Why these problems

The criterion is a product of three factors: **impact** × **is math the bottleneck** (type A, type B, or not math at all) × **AI tractability** (is there an automatic referee). The 8 problems deliberately cover the whole spectrum, from almost all math (Yang–Mills) to almost no math (materials synthesis, weather).

#	Problem	Impact	Is math the bottleneck?	Fast automatic referee
1	Electronic structure (DFT)	Very high	Partly: approximation algorithms are, theorems aren't	Yes (high-accuracy computation)
2	Inverse design + autonomous labs	High	No: synthesis and characterization are	Slow (experiments take weeks)
3	High-temperature superconductivity	Scientific: high; economic: medium	Type B: nobody knows which model is right	Partial
4	Turbulence	High	Engineering side no; Clay side type A	Engineering yes; Clay has Lean
5	Fusion	Very high if it succeeds	A small part	Yes (simulators + devices)
6	Weather and climate	Very high	Weather no; climate type B plus data	Weather yes, climate no
7	Yang–Mills and making mathematical physics rigorous	Conceptually very high, practically about 0	Yes: type A statement, type B core	Yes (Lean)
8	Quantum gravity and string theory	Conceptually very high, practically about 0	Subproblems type A; as a whole type B plus missing experiments	Only for subproblems

Not selected but worth mentioning: topological phases (the mathematical classification was finished long ago; it is stuck on materials, as in the dispute over the evidence for Majorana 1 Δ); quantum computing hardware (goes to §19); proteins and drugs (go to §21); the glass transition (type B); dark matter and cosmology (the bottleneck is observation); making statistical mechanics rigorous, such as Ising and percolation (merged into Problem 7).

20.2 Problem 1: Electronic structure (the many-electron Schrödinger equation and DFT)

In one sentence All the "rules" of chemistry and materials are written in one equation, but the cost of solving it exactly grows exponentially with the number of electrons. The whole

industry runs on an approximation (density functional theory, DFT), and nobody knows the exact form of its central term.

Background A **Hamiltonian** is "the recipe for total energy": electron kinetic energy + attraction by the nuclei + repulsion between electrons. The **Schrödinger equation** $H\Psi = E\Psi$ asks for a **wave function** Ψ such that applying the recipe to it is the same as "multiplying by a number E "; the smallest E is the ground-state energy. Ψ assigns a number to "every combination of all the electron positions". Almost everything chemistry cares about is an **energy difference**: heat of reaction, reaction barrier, binding strength, battery voltage.

Exponential explosion (by hand). With 10 grid points per coordinate, 1 electron needs 10^3 numbers and N electrons need $10^{(3N)}$. Benzene has 42 electrons, which means 10^{126} numbers, while the universe contains only about 10^{80} atoms. Switching to the orbital basis sets that chemists use, the exponential is still there: the 54-electron, 54-orbital "active space" commonly used for FeMoco, the active site of nitrogenase, has about 3.8×10^{30} configurations, which no classical computer can handle.

DFT's trick. The Hohenberg–Kohn theorem (1964) proves that the ground-state energy depends only on the **electron density** $n(r)$, which has just 3 variables [established]. Kohn–Sham (1965) rewrote the problem in terms of "fictitious non-interacting electrons", bringing the cost down to about N^3 . The price is that all the hard many-body effects get swept into one term: the **exchange-correlation functional** $E_{xc}[n]$. The theorem guarantees that it exists but does not say what it looks like, hence a ladder of approximations and hundreds of functionals.

Why 1 kcal/mol matters (by hand). A reaction rate is roughly proportional to $\exp(-\text{barrier}/RT)$, and at 300 K, $RT \approx 0.596$ kcal/mol. An error of 1.37 kcal/mol in the barrier changes the rate by 10 $\times$; common DFT barrier errors are 3–5 kcal/mol, and at 4 the rate is off by about 820 $\times$. That is why 1 kcal/mol is called "chemical accuracy". The gold standard, CCSD(T), reaches it, but its cost grows as N^7 : double the size of the molecule and the computation grows 128 $\times$.

Why it matters Haber–Bosch ammonia synthesis uses about 1–2% of global energy (IEA), and about half the world's population depends on synthetic nitrogen fertilizer for its food (Erisman et al. 2008) [established]. Nitrogenase fixes nitrogen at room temperature, but the mechanism of FeMoco still has not been worked out computationally; it is also the flagship case for "what chemistry can quantum computers solve" (Reiher et al. 2017). The key quantities for batteries, catalysts and drugs are also energy differences on the order of 1 kcal/mol.

What's blocking it: is math the bottleneck? Theorems are not the bottleneck; **approximation algorithms and data** are. One "impossibility" theorem fixes the direction: Schuch & Verstraete (2009) proved that if a universal functional could be computed efficiently, QMA-complete problems could be solved efficiently (QMA is the quantum version of NP) [established]. So no formula works well for every system, and the only way forward is to exploit the structure of real chemistry. On data, Skala's performance comes from high-accuracy reference data "of unprecedented volume", and Meta's OMol25 used over 100 million DFT calculations [recent]. In strongly correlated systems (transition metals, FeMoco, cuprates) even CCSD(T) is unreliable, so there is no referee, and the problem is close to type B. Quantum computers are still far off: resource estimates for FeMoco have come down to about 1,500 logical qubits and about 10^9 Toffoli gates (secondhand summary $\triangle$), still far beyond 2026 hardware; Lee et al. (2023) argue that there is no evidence of a generic exponential speedup for ground-state chemistry [established].

What AI can do

- *AI math/proof*: prove error bounds and exact constraints for approximations (e.g. the Lieb–Oxford bound), and design new algorithms with provable complexity. Useful, but only supporting: a theorem won't tell you the reaction barrier of a given catalyst.
- *Other AI capabilities (the main contributors)*: neural network wave functions (FermiNet and others); neural network functionals (Skala has about 276,000 parameters and an error of 2.8 kcal/mol on GMTKN55, at the cost level of a semilocal functional; self-reported by the authors, independent replication pending $\triangle$); ML interatomic potentials (running molecular dynamics at "DFT accuracy and classical force field speed"); coding agents + active learning that automatically decide which expensive reference calculation to run next. With high-accuracy computation as an automatic referee, AI climbs fast.

If solved: first-order impact If "DFT cost + chemical accuracy" covers most of chemistry, transition metals included: computational screening of catalysts and battery materials becomes trustworthy, failed syntheses drop sharply, and reaction mechanisms can be computed before anyone runs them.

If solved: second-order impact Chemistry shifts to "simulate first, experiment later", like EDA in chip design. Quantum computing loses part of its flagship application [speculation]. Dual use: Urbina et al. (2022) reversed the objective of a drug-discovery model and generated about 40,000 candidate toxic molecules in 6 hours [established]. A lesson in specification gaming: Medvedev et al. (2017) found that over more than a decade, new functionals fit energies better and better while their densities drifted further and further from the exact solution [established].

Timing Optimistic: by end of 2027, ML functionals and potentials reach chemical accuracy for main-group and most transition-metal chemistry. Central: main-group chemistry is "solved" in a practical sense around 2028, while FeMoco-level strongly correlated systems are still disputed in 2030; this depends on **curve ①** (by mid-2027 agents can run for about 150 hours, enough to manage "data generation → training → evaluation" on their own) and has almost nothing to do with curve ②. Pessimistic: ML functionals fail to generalize to new elements and unusual bonding, and strongly correlated chemistry has to wait for fault-tolerant quantum computers after 2035.

Counterarguments and uncertainty Benchmark accuracy is not real catalysis: catalysis happens on surfaces, in solvents and under dynamic conditions, and is often an experimental question. ML functionals may silently go wrong outside their training distribution. For most drug projects, electronic-structure accuracy is not the rate-limiting step.

Signals to watch Skala-type functionals become the default option in mainstream codes and are independently replicated on transition metals; an ML-derived mechanism for FeMoco agrees with experiment. A signal in the other direction: quantum hardware beats the best classical methods on a chemistry problem for the first time (unlikely before 2030).

Sources: Dirac 1929, doi:10.1098/rspa.1929.0094; Schuch & Verstraete arXiv:0712.0483; Skala arXiv:2506.14665; FermiNet arXiv:1909.02487; OMol25 arXiv:2505.08762; Reiher et al. arXiv:1605.03590; Lee et al. arXiv:2208.02199; FeMoco resource estimate (secondhand) <https://postquantum.com/quantum-utility-map/quantum-chemistry-drug-discovery-catalysis/> △; IEA <https://www.iea.org/reports/ammonia-technology-roadmap>; Erisman et al. 2008, Nat. Geosci. 1:636; Medvedev et al. 2017, Science 355:49; Urbina et al. 2022, Nat. Mach. Intell. 4:189

20.3 Problem 2: Inverse design and autonomous labs (batteries, catalysts, carbon capture)

In one sentence Turn "given a material, compute its properties" around into "given the properties you want, find the material and make it". AI can already generate millions of candidates cheaply, so the bottleneck has moved to synthesis, characterization, testing and scale-up, none of which are math problems.

Background The forward problem is "here is the recipe, taste the dish"; the inverse problem is "here is the taste you want, find the recipe, then actually cook it". Picking just 4 of about 80 commonly used elements gives $C(80,4) = 1,581,580$ combinations, which then have to be multiplied by ratios and crystal structures. Stability is judged with the **convex hull**: if AB has lower energy when it decomposes into A + B, then AB is unstable. But stable

does not mean **synthesizable**; that depends on reaction pathways and kinetics, and there is no good theory for it.

The real numbers in the funnel. Generate (e.g. with the diffusion model MatterGen) → coarse screening with ML potentials → refined DFT → synthesis → characterization (X-ray diffraction, XRD) → performance testing → scale-up. GNoME (2023) predicted about 380,000 crystals on the convex hull; **736 of them, about 0.2%**, have been independently synthesized by outside groups [established]. A-Lab claimed to have synthesized 41 of 58 targets in 17 days; critics then pointed out that many of the "new materials" were ordered or disordered variants of known compounds [recent].

By hand: Amdahl's law. One discovery cycle = 2 weeks of design + 6 weeks of synthesis and testing. If AI cuts design to 0, the cycle goes from 8 weeks to 6, only 1.33× faster. To go 10× faster, the lab itself has to be 10× faster.

Why it matters Batteries determine how large a share renewables can take; catalysts are used for green ammonia and for producing hydrogen by water electrolysis (PEM electrolyzers depend on scarce iridium Δ). The US Materials Genome Initiative (2011) said new materials typically take 10–20 years from discovery to market [established].

What's blocking it: is math the bottleneck? Basically not. The bottlenecks, in order: synthesizability (a type B chemistry problem); ambiguous characterization (the same XRD pattern can be read several ways); slow testing (battery lifetime tests take months); processing (performance depends on microstructure, defects and interfaces); data bias (failed experiments are almost never published). The math that is needed (Bayesian experimental design, symmetry enumeration) has long been mature.

What AI can do

- *AI math/proof:* direct effect close to zero.
- *Other AI capabilities (the main contributors):* generative models (MatterGen synthesized TaCr_2O_6 with a target bulk modulus of 200 GPa; measured about 169 GPa Δ); coarse screening with ML potentials; LLMs extracting synthesis recipes from the literature; robotic closed-loop experiments and automated characterization. Periodic Labs (which raised a \$300 million seed round in 2025) works on chip-cooling materials in the near term and superconductors in the long term [recent]. **The referee is experiment**, which is slow and noisy, so the ceiling on acceleration is set by lab throughput.

If solved: first-order impact If "half of predicted materials can be made, and their performance matches the prediction": R&D cycles shrink from 10–20 years to a few years, and batteries, catalysts and sorbents keep improving (less cobalt, lithium, iridium).

More likely this looks like a few percent of improvement per year rather than a single leap [speculation].

If solved: second-order impact Verification infrastructure becomes the new bottleneck:

we need "machine-checkable materials claims", with raw data, automated refinement results and reproducible synthesis records published together, the materials-science equivalent of Lean [speculation]. Replacing cobalt and rare earths with common elements would reshape supply chains. **AI safety:** autonomous labs are where AI first gets to act in the physical world, which raises chemical safety, dual use, and the question of "who can remotely order the synthesis of what".

Timing Optimistic: by 2028 autonomous lab throughput $\times 10$, with a success rate of $\geq 50\%$ on predicted targets. Central: in 2029–2032, AI-designed compositions enter niche commercial products; this depends on **curve ①** (by end of 2027 agents can work continuously for about 400 hours, enough to manage months of closed-loop experiments), but lab capital and physical cycle times are hard constraints, unrelated to curve ②. Pessimistic: synthesizability prediction stays poor, and together with a trust deficit this pushes things past 2035.

Counterarguments and uncertainty The gains may come mainly from optimizing known systems; models built around perfect crystals may focus on the wrong thing, because processing is what matters. Optimists argue that robots plus large models can raise throughput by two orders of magnitude and break the Amdahl constraint.

Signals to watch The fraction of AI-predicted materials reproduced by independent labs, and whether papers come with raw data; a battery maker mass-producing an AI-designed chemistry.

Sources: GNoME, Nature 624:80; A-Lab, Nature 624:86; A-Lab critique https://www.theregister.com/2024/01/31/ai_chemistry_research_disputed/; MatterGen arXiv:2312.03687; Periodic Labs <https://siliconangle.com/2025/10/01/periodic-labs-raises-300m-accelerate-scientific-research-ai/>

20.4 Problem 3: The mechanism of high-temperature superconductivity (the Hubbard model) and room-temperature superconductivity

In one sentence Cuprates conduct with zero resistance at about -140°C , and after 40 years there is still no accepted theory, so nobody can design a room-temperature superconductor "from a blueprint". Ironically, the highest-temperature superconductors (high-pressure hydrides) **don't rely on that theory at all**; they were computed with DFT.

Background Superconductivity = exactly zero resistance + expelling magnetic fields from the material. BCS theory (1957) explains conventional superconductivity: electrons pair up

through lattice vibrations (phonons), like a person sinking into a trampoline and pulling the person next to them in. The ambient-pressure record for cuprates is about 133–138 K [established], and it is generally believed that the pairing comes from strong electron correlation and magnetism.

The Hubbard model (two sites, by hand). Electrons "hop" between lattice sites, lowering the energy by about t per hop; two electrons squeezed onto the same site pay a repulsion energy U . With 2 sites and 2 electrons, the ground-state energy is $E = (U - \sqrt{(U^2 + 16t^2)})/2$. At $U = 0$, $E = -2t$ and the electrons are freely delocalized; at $U = 8t$, $E \approx -0.47t$ and the electrons are basically frozen on their own sites. The remaining gain comes from "virtual hopping when neighboring spins are antiparallel", which is equivalent to an antiferromagnetic coupling $J \approx 4t^2/U = 0.5t$. **Magnetism grows out of the competition between "hopping" and "repulsion"**, and the cuprate parent compounds are indeed antiferromagnets. The hard part: on an infinite 2D lattice with doping, is the ground state superconducting or a stripe state?

This is a numerical dispute, not a dispute about proofs. Qin et al. (2020): at $t' = 0$, $U = 8$ there is no superconductivity; the ground state is a stripe state. Xu et al. (2024): once next-nearest-neighbor hopping t' is added, superconductivity and stripes coexist. Roth et al. (2025-11, neural network quantum states): claim d-wave superconductivity even at $U = 4$, $t' = 0$ [recent, unreviewed].

Warnings and real progress. In 2023 LK-99 was traced within weeks to an artifact caused by Cu_2S impurities [established]. The hydrides H_3S (203 K) and LaH_{10} (about 250 K) superconduct under high pressures of about 150–170 GPa, and both were **predicted first by DFT and then confirmed by experiment** [established]. Strained nickelate thin films reached an onset T_c of about 26–42 K at ambient pressure in 2025 [recent]. The superconductors found by AI so far are all low-temperature: BEE-NET screened 1.3 million candidates down to 741, and experiments confirmed 2 [recent].

Why it matters Scientifically, it is the paradigm problem of strongly correlated electrons. Economically, ambient-pressure room-temperature superconductivity would mean lossless power transmission, cheaper MRI and more compact fusion magnets; but grid economics are mainly constrained by cable cost.

What's blocking it: is math the bottleneck? Mainly **type B + numerics**: nobody knows which model is right (one band or several? how big is t' ? do phonons play a role?). The "sign problem" in fermionic quantum Monte Carlo is NP-hard in general (Troyer & Wiese 2005) [established], and this impossibility result pushes people toward tensor networks and neural network quantum states. Experimentally, sample quality is poor and high-pressure measurements are hard to reproduce. On the design side, there is no reliable T_c prediction

for unconventional superconductors; for conventional ones there is (DFT + Eliashberg theory), and that is how the hydrides were found.

What AI can do

- *AI math/proof*: rigorously prove "superconductivity or not" in certain parameter regimes. This is far harder than numerics and does not directly produce materials; but the Haldane gap and Heisenberg ferromagnet results in Problem 7 show that simpler many-body models are now within range. Medium value: it can constrain theory.
- *Other AI capabilities (the main contributors)*: neural network quantum states (already directly involved in the Hubbard dispute); ML-accelerated T_c screening; autonomous synthesis; using LLMs to clean up the literature and flag suspicious claims.

If solved: first-order impact Once the Hubbard phase diagram is settled, we will know whether a one-band model is enough and which parameter to tune. If ambient-pressure room-temperature superconductivity is discovered, it will set off a revolution in materials and engineering, but deployment will take decades.

If solved: second-order impact Power transmission, storage, motors, fusion magnets (Problem 5) and data-center power supply would all change. A nearer effect is epistemic: LK-99 showed the dynamics of "social media + replication race". When AI generates hundreds or thousands of candidate claims a day, the community will need a triage mechanism: standardized criteria (zero resistance, Meissner effect, specific heat, independent replication) plus public raw data [speculation].

Timing Consensus on the Hubbard phase diagram: optimistic 2027 (neural network quantum states and AFQMC converge), central 2028–2030, pessimistic: the debate turns into "which model to use for real materials"; this depends on compute and algorithms, not on curve ②. Ambient-pressure room-temperature superconductivity, independently verified: optimistic in the early 2030s (a tail event), central: not before 2035, pessimistic: it may not exist at all. AI may raise candidate throughput 2–5× [speculation] Δ , but the hit rate is extremely low to begin with.

Counterarguments and uncertainty Understanding the mechanism is **neither necessary nor sufficient** for discovering new materials: the hydrides came from DFT, the nickelates from analogy plus luck. Neural network quantum states and AFQMC contradict each other, so even the "numerical facts" are not yet stable.

Signals to watch Different methods give a consistent phase diagram at the same parameters; nickelate ambient-pressure T_c above 77 K (the temperature of liquid nitrogen); an AI-designed superconductor with $T_c > 30$ K, independently replicated.

Sources: Qin et al. arXiv:1910.08931; Xu et al. arXiv:2303.08376; Roth et al. arXiv:2511.07566; Troyer & Wiese arXiv:cond-mat/0408370; Ko et al. 2025, Nature 638:935; LK-99 arXiv:2307.12008; Drozdov et al. 2015, Nature 525:73; Drozdov et al. 2019, Nature 569:528; BEE-NET arXiv:2503.20005

20.5 Problem 4: Turbulence (short card; for the Clay problem see §11)

The equations for water and air were written down in 1845, yet turbulence still has to be computed with empirical models. The Reynolds number (the ratio of inertia to viscosity) for an airliner wing is about 3×10^7 . The cost of directly computing every eddy (DNS) grows roughly as Re^3 , and at this Re the grid alone would need about 6×10^{16} points, so it has to be modeled: the effect of small eddies on the mean flow is filled in with empirical "closure models", which often fail on separated flow and stall (NASA's *CFD Vision 2030* lists this as one of the biggest bottlenecks [established]).

On the engineering side the bottleneck is not math, but scarce high- Re data, ML closures that fail as soon as the flow changes (Duraismy et al. 2019), and aviation certification's requirement for verifiable models. The main contributors are AI's other capabilities: learned LES reaches the same accuracy on 2D turbulence with grids 8–10× coarser (Kochkov et al. 2021); coding agents port legacy code to GPUs, the most mundane item and possibly the one with the biggest effect on cost. Central judgment: not "basically solved" before 2030; more likely errors fall 10–30% per year (depends on compute, data and curve ①). **On the Clay side it is type A math**, and AI is already attacking it: OpenAI claims to have proved the forced branch, while the unforced branch is still open (§11); Chen–Hou used a computer-assisted proof to show that 3D Euler with a boundary blows up [established]; DeepMind and others found a family of unstable blowup solutions with accuracy that "meets the requirements of a computer-assisted proof" (2025-09) [recent]. Unforced branch: central 2028–2030 (curve ②; Millennium Problems will lag). A few theorems bear on turbulence physics: Isett proved the Onsager conjecture (energy can still be dissipated as viscosity goes to zero); Deng–Hani–Ma (2025) rigorously derived the NS–Fourier equations from hard-sphere particles [recent]. The two sides barely connect: engineering solvers never rely on global regularity. **Sources:** arXiv:2509.14185; Chen & Hou arXiv:2210.07191; Isett arXiv:1608.08301; Deng, Hani & Ma arXiv:2503.01800; Duraismy et al. arXiv:1804.00183; Kochkov et al. arXiv:2102.01010

20.6 Problem 5: Fusion: stability, control and stellarator design

In one sentence The goal is to hold a plasma at over 100 million degrees in a magnetic field for long enough. Physical feasibility has been partly demonstrated; what remains is instability control, materials and engineering. Mathematical proof is only a small piece.

Background *How hot, and for how long (by hand)*. Deuterium + tritium → helium + neutron + 17.6 MeV. Ignition roughly requires the Lawson criterion: density × temperature × confinement time $\approx 3 \times 10^{21}$ keV·s/m³. With a density of 10^{20} particles/m³ and a temperature of 15 keV (about 170 million degrees), the confinement time must be at least $3 \times 10^{21} \div (10^{20} \times 15) = 2$ seconds.

The model ladder. From exact to cheap: Vlasov–Maxwell (6-dimensional phase space, the subject of §12) → gyrokinetics (averages out the fast gyration around field lines; 5 dimensions) → magnetohydrodynamics (treats the plasma as a conducting fluid) → 1D transport codes (such as TORAX). The global smoothness result of §12 is a model-health proof at the top rung of this ladder, with no direct effect on engineering.

Two kinds of magnetic cage. A **tokamak** is axisymmetric and relies on a large current in the plasma itself for confinement; the price is that it can "disrupt": the current collapses within milliseconds, releasing huge forces and heat. A **stellarator** uses twisted 3D coils and can run in steady state; the price is that design becomes a 3D optimization problem.

Grad's conjecture and its counterexample. Good magnetic confinement needs the field lines to trace out layer upon layer of nested "magnetic surfaces", like an onion. Grad (1967) conjectured that without symmetry, no such smooth equilibrium exists; it has been called "the central mathematical problem of stellarator theory". On 2026-09-21, Gómez-Serrano, Liehr and Taylor gave a **counterexample**: a smooth equilibrium with only N-fold rotational symmetry, with a Lean 4 proof [recent, unreviewed]. The paper discloses the division of labor: humans wrote the roadmap, and GPT-5.6 Sol, Claude Fable 5 and Claude Opus 5 filled in the details and wrote the Lean. It is a model case of "**human-led + AI fills in details + formalization**". Physically it needs a discount: the field is exactly zero on the magnetic axis, and the result holds only in the limit of small curvature and large N. In practice engineers rely on numerical optimization (e.g. "precise quasisymmetry", Landreman & Paul 2022).

Status. In 2025-04 NIF produced 8.6 MJ from 2.08 MJ of laser energy (gain about 4.1), but the laser system drew about 300 MJ from the grid [recent]. ITER's deuterium-tritium operation has been pushed to 2039 $\triangle$; SPARC's first plasma has reportedly slipped to 2027 (secondhand $\triangle$). CFS and DeepMind (2025-10) are using TORAX to run millions of virtual experiments, combined with AlphaEvolve to search for operating scenarios, and plan to use reinforcement learning for real-time control later; Google also signed a 200 MW power purchase agreement for ARC [recent].

Why it matters Fusion is zero-carbon, stable baseload power, and the electricity demand of AI data centers is pushing tech companies toward fusion power purchase agreements. But in the end it has to compete on cost with "solar + storage" (Jassby 2017 $\triangle$).

What's blocking it: is math the bottleneck? Mainly **engineering and materials**: 14 MeV neutron damage to the first wall, tritium breeding Δ , reliability of high-temperature superconducting magnets, divertor heat loads, cost. Next comes **physics**: predicting turbulent transport (a single multiscale simulation can take on the order of ten million CPU hours Δ), disruptions and tearing modes. **Math is only a small part**: existence of 3D equilibria (Grad), model well-posedness (Vlasov–Maxwell with multiple particle species is still open), coil optimization.

What AI can do

- *AI math/proof*: Grad-type existence and rigidity theorems, rigorous stability criteria, equilibrium solvers with error bounds. They tell designers "which configurations exist", but they can't build a reactor.
- *Other AI capabilities (the main contributors)*: reinforcement learning control (controlling 19 coils at once on TCV, Degraeve et al. 2022; avoiding tearing instabilities in advance on DIII-D, Seo et al. 2024 [established]); disruption prediction; differentiable simulators and surrogate models for gyrokinetics; stellarator coil optimization; coding agents maintaining the simulation code stack.

If solved: first-order impact Faster device commissioning and higher performance; fewer disruptions, which is critical for ITER-class devices, where one major disruption can damage the machine; shorter stellarator design cycles.

If solved: second-order impact AI and energy form a positive feedback loop: fusion powers data centers, and AI speeds up fusion R&D [speculation]. **AI safety**: reinforcement learning controllers enter safety-critical physical systems; fusion is a near-ideal testbed for Guaranteed Safe AI (GSAI), since the physical models are good and the risks are physical but bounded, so one can do "finite-horizon certificates + runtime monitoring + backup controllers". **Nuclear proliferation**: fusion neutrons can breed fissile material, which needs safeguards oversight Δ .

Timing SPARC reaching $Q > 1$ (fusion output power greater than heating power): optimistic 2027, central 2028, pessimistic after 2030. Commercial power: optimistic early 2030s, central mid-to-late 2030s, pessimistic in the 2040s or later. AI can shorten commissioning and design by months to years, but it can't change the real time needed to qualify materials under irradiation; progress is set by the hardware cadence, almost unrelated to curve ② and weakly related to curve ①.

Counterarguments and uncertainty AI control may just be icing on the cake: the physics of compact high-field tokamaks is already good enough, and the hard part is building them. Fusion may also never beat the cost curve of "solar + storage".

Signals to watch The actual dates of SPARC's first plasma and $Q > 1$; reinforcement learning controllers entering routine operation on large devices; Grad-type results being written into design codes.

Sources: Grad counterexample arXiv:2609.24739; Constantin, Drivas & Ginsberg arXiv:2108.05977; Landreman & Paul arXiv:2108.03711; Degraeve et al. 2022, Nature 602:414; Seo et al. 2024, Nature 626:746; TORAX arXiv:2406.06718; CFS × DeepMind <https://www.ans.org/news/2025-10-22/article-7484/commonwealth-fusion-systems-partners-with-google-deepmind/>; NIF <https://www.llnl.gov/article/52906/llnl-experts-foster-national-fusion-energy-ecosystem-ife-star-conference>; SPARC (secondhand) <https://www.insidedeep.tech.com/commonwealth-fusion-systems-sparc-full-guide/> △

20.7 Problem 6: Weather and climate (short card)

Weather is an "initial value problem": know the atmosphere now, then integrate forward. Chaos is a hard constraint: suppose errors double every 2 days (an illustrative number); after 14 days they have grown $128\times$, so the limit of predictability is about two weeks [established]. AI models learn "what the atmosphere looks like 6 hours later" from more than 40 years of reanalysis data (ERA5): GraphCast (2023) beat ECMWF's operational model HRES on about 90% of targets, and GenCast (2024) beat ECMWF's ensemble forecast on about 97% of targets; **ECMWF's AIFS went operational on 2025-02-25**, cutting energy use by about $1,000\times$ [recent]. The remaining bottlenecks are sparse observations, initial-condition assimilation and few samples of extreme events, none of which is math.

Climate is a "boundary condition problem": what the statistical distribution of weather becomes after CO_2 doubles. IPCC AR6 gives a likely range of $2.5\text{--}4^\circ\text{C}$ for equilibrium climate sensitivity, and the largest source of uncertainty is **cloud feedback**. This is a type B problem plus a data problem: there is no "doubled- CO_2 world" to train on, so AI can only extrapolate; the ACE2 authors themselves acknowledge that the model's sensitivity to changing sea surface temperature or CO_2 alone is "not entirely realistic" [recent]. End-to-end "observation → forecast" in operations: central 2028. Climate sensitivity range narrowed to about $\pm 0.5^\circ\text{C}$: central in the mid-2030s, requiring kilometer-scale simulation plus AI cloud emulators. The same learned simulators have already rewritten operations for weather, which is scored every day, but move slowly on climate, which is scored once every few decades. This is the cleanest example of rule 3 in 20.10. **Sources:** GraphCast arXiv:2212.12794; GenCast arXiv:2312.15796; ACE2 arXiv:2411.11268; AIFS <https://www.ecmwf.int/en/about/media-centre/news/2025/ecmwfs-ensemble-ai-forecasts-become-operational>; IPCC AR6 WG1 Chapter 7

20.8 Problem 7: The Yang–Mills mass gap and making mathematical physics "rigorous"

In one sentence Physicists have made extremely precise predictions with quantum field theory, yet mathematicians still cannot prove that four-dimensional quantum Yang–Mills theory "exists" and has a mass gap. This is a Millennium Problem with a clear statement that AI might be the first to crack, but it has almost no effect on experiment or engineering; its value is conceptual and methodological.

Background *Quantum field theory (QFT)*. The universe is full of "fields", and particles are ripples in those fields. QED's predictions agree with experiment to about the 10^{-12} level $\triangle$, but the calculations use series that may diverge, and the infinities along the way are removed by "renormalization". **Mathematically, nobody has yet constructed an interacting 4D QFT that satisfies the axioms**; the most natural candidate, ϕ^4_d , has been proved to reduce to a free theory in the continuum limit (Aizenman & Duminil-Copin 2021) [established].

Yang–Mills (an analogy). In electromagnetism light carries no charge, and two beams of light pass through each other. In Yang–Mills (such as QCD, which describes the strong force), the "light" (gluons) itself carries color charge and attracts itself, so it twists into flux tubes that lock quarks in. The **mass gap** says: in the equations gluons are massless, yet the lightest particle in the theory (the glueball) has positive mass. Lattice QCD puts the lightest glueball at about 1.7 GeV; about 99% of the mass of visible matter comes from QCD dynamics, not from the Higgs mechanism (Dürr et al. 2008) [established].

What makes it hard. The lattice method cuts spacetime into a grid with spacing a , computes a huge integral, then lets $a \rightarrow 0$, like approximating an integral with Riemann sums. But here the integral is infinite-dimensional, over **all field configurations**, and an infinite-dimensional space has no ready-made "uniform measure"; to prove that the limit exists, is interacting and has a gap, you have to control fluctuations at all scales at once. 2D has been fully constructed, and there is progress on stochastic quantization in 2D and 3D (Chandra, Chevyrev, Hairer & Shen); 4D is still open as of 2026-10, and several "complete proofs" online are unreviewed $\triangle$.

New in 2026: the type A list in mathematical physics is being cleared. A secondhand count says this batch of OpenAI manuscripts contains about 25 "mathematical physics" families $\triangle$. The table below lists representatives found in the CONTENTS and lean/docs of

`openai/math` (downloaded 2026-10-09); all are **claimed, unreviewed**:

Family	Problem claimed solved (plain language)	Formalization
261	Anderson model (electrons hopping on a lattice with impurities): for $d \geq 3$ and weak disorder, electrons can "spread out" and conduct (claims to solve this part of Simon's Problem 1)	Only the 2D part; 3D has no Lean
271	Spontaneous magnetization at low temperature of the $d \geq 3$ quantum Heisenberg ferromagnet (quantum compass needles on a lattice; neighbors want to point the same way)	Has a Comparator statement
268	Haldane gap for spin-1 chains (the 1983 conjecture): the lowest excitation energy has a positive gap	None
267	BEC at positive temperature in a hard-sphere Bose gas (a large number of particles crowd into the same quantum state)	Only ground-state condensation
376, 362, 221	Forced NS (§11), Vlasov–Maxwell (§12), Mézard–Parisi (§13)	See the corresponding sections

As of 2026-10-09 we found no named review by mathematical physicists. These results have one thing in common: what they confirm is something physicists already believed, such as that a 3D ferromagnet magnetizes at low temperature and a hard-sphere gas condenses; what §11 and §12 prove is that "the model does not break down on its own mathematically". These are **model-health proofs**: they answer "is the model logically healthy", not "what is nature like".

Why it matters Practical impact is about 0: lattice QCD already gave the numbers. Conceptual impact is large: is QFT a self-consistent mathematical theory? The methods will spill over into other branches of mathematics (precedent: Witten's TQFT gave knot invariants). It is also the best test case for whether AI can **build new theoretical frameworks**, not just solve problems.

What's blocking it: is math the bottleneck? Yes, but it is "type A statement, type B core": the statement is clear, but the proof may need new concepts (a rigorous 4D renormalization group, handling the Gribov ambiguity). There is also a formalization bottleneck [speculation]: physics has barely been formalized, HepLean/PhysLean only started in 2024, and Mathlib lacks infrastructure such as functional integration. Just writing the **statement** "4D Yang–Mills satisfies the Osterwalder–Schrader axioms" in Lean is a huge specification project, the pure form of "the bottleneck moves to the statement" (§3, §18).

What AI can do

- *AI math/proof (the main contributor)*: proofs in constructive QFT are known for extremely long, "bookkeeping" estimates (cluster expansions, multiscale analysis), which is exactly the kind of labor AI is best at; the 2026 batch of results shows that problems in the Lieb and Simon tradition are now within range.

- *Other AI capabilities:* flow-model sampling for lattice QCD (Albergo et al. 2019); formalizing physics; symbolic computation.

If solved: first-order impact The first rigorous interacting 4D QFT is born; large parts of the Lieb and Simon lists get closed, and mathematical physics textbooks have to be rewritten.

If solved: second-order impact The field's center of gravity shifts from "proving" to "choosing and auditing statements", and there is a low-probability "rigor dividend": finding that some of physicists' default assumptions don't hold. A rigorous renormalization group connects to deep learning theory (20.10.2). "A Millennium Problem with physical meaning solved by AI" would cause a jump in public perception and would also test verification institutions: Clay requires a two-year period of community acceptance, while AI produces results much faster; statement faithfulness is still an issue (the lesson of forced NS, §11).

Timing Lieb/Simon-type many-body problems: the central judgment is that by end of 2027, most well-known type A problems in mathematical physics will be claimed solved, with more than half confirmed by experts [speculation]; the basis is that a batch already appeared in 2026-09, and "famous but relatively tractable" problems will fall ahead of curve ②'s average (about 18% by end of 2027). Yang–Mills mass gap: optimistic 2028, central 2030–2033, pessimistic after 2035 [speculation]; curve ②'s 77% in 2030 refers to typical problems, and Millennium-level ones will lag; curve ③ makes proofs cheaper, but the specification cost of the statement does not fall.

Counterarguments and uncertainty The effective field theory view holds that the Standard Model is only valid below some energy scale, so rigorously constructing the continuum limit means studying an idealized object. A rigorous proof may only confirm what the lattice already "saw". AI proofs may be "unreadable by humans" (Gómez-Serrano commenting on the NS paper), which discounts their conceptual value.

Signals to watch 271, 261 and 268 get confirmation from named experts or are published in journals; statements of the QFT axioms appear in Mathlib or PhysLean; Clay takes a position.

Sources: <https://github.com/openai/math> (CONTENTS; lean/docs 261, 267, 271; downloaded 2026-10-09); secondhand count <https://kingy.ai/blog/openai-math-722-manuscripts-results-proofs-compute-costs/> [△]; Simon problems https://en.wikipedia.org/wiki/Simon_problems; Clay <https://www.claymath.org/lectures/yang-mills-and-the-mass-gap/>; Aizenman & Duminil-Copin arXiv:1912.07973; Chandra et al. arXiv:2006.04987; Dürr et al. 2008, Science 322:1224; HepLean arXiv:2405.08863; Albergo et al. arXiv:1904.12072

20.9 Problem 8: Quantum gravity and string theory (short card)

General relativity says gravity is smooth curvature of spacetime; quantum mechanics is discrete and probabilistic; forcing the two together produces infinities that can't be removed. Candidate theories (string theory, loop quantum gravity and others) are mathematically very rich but make almost no testable predictions. The reason can be worked out by hand: the energy scale where quantum gravity becomes significant is about 1.2×10^{19} GeV, while the LHC reaches only 1.36×10^4 GeV, a gap of about 10^{15} ; at fixed magnetic field, energy is proportional to ring radius, so scaling the LHC's radius of about 4.3 km up by 10^{15} gives about 400 light-years. The number of ways string theory's extra dimensions can be compactified is estimated at 10^{500} to $10^{272,000}$ (Taylor & Wang 2015), which makes a unique prediction hard.

The first bottleneck is experiment; the second is type B math (string theory has no complete non-perturbative definition), but type A subproblems are everywhere: scattering amplitude identities, bootstrap bounds, Calabi–Yau classification. Using only symmetry, self-consistency and semidefinite programming, the conformal bootstrap computed the 3D Ising critical exponents to 5–6 decimal places (Kos et al. 2016) [established]. AI has already entered: in 2026-02, Guevara, Strominger and others (for OpenAI) showed that the "single-minus-helicity gluon tree amplitude", long assumed to be zero, is nonzero in particular configurations; the formula was proposed by GPT-5.2 Pro, proved by an internal model and verified by the authors [recent]. Results like this will ramp up in 2027 (curve ②); as for experimental confirmation of quantum gravity, the central judgment is not before 2040. AI's most valuable contribution may be designing tabletop experiments, for example testing whether gravity can produce entanglement (Bose et al. 2017, not yet realized) [speculation]. Conversely, a field with no experimental constraints, plus AI producing unlimited **self-consistent** mathematics, could lead to a flood of beautiful theories with no selection pressure: Lean can verify the proof, but not "this is how nature is" [speculation]. **Sources:** arXiv:2602.12176; Taylor & Wang arXiv:1511.03209; Kos et al. arXiv:1603.04436; Bose et al. arXiv:1707.06050

20.10 Second-order effects in this field

20.10.1 Three rules

Rule 1: most "math breakthroughs" in physics are model-health proofs. NS (§11), Vlasov–Maxwell (§12), Anderson, Heisenberg, BEC and Yang–Mills all confirm things physicists already believe. The math that really changes direction comes in three kinds: **impossibility theorems** (Schuch–Verstraete, Troyer–Wiese, which push a field toward approximations that "exploit structure"); **existence results and counterexamples** (the Grad counterexample, the Onsager conjecture, which tell you what is in the design space); **new**

algorithms. So AI math's biggest value to physics may lie in designing algorithms and proving impossibility, not in sweeping up conjectures.

Rule 2: candidates are cheap, verification is expensive. In math it is "proofs are cheap, auditing statements is expensive" (§18); in physics and chemistry it is "candidates are cheap, verification is expensive", where verification relies on experiment, high-accuracy computation, or even time (climate). GNoME's 380,000 versus 736, AI superconductor candidates plus LK-99-style frenzies, and a single release of 719 manuscripts with almost no named reviewers all have the same shape. The next scarce resource is **verification infrastructure**: autonomous replication labs, public raw data, statement auditing.

Rule 3: how fast AI speeds up a field is set by the speed of its referee.

Field	Referee	One feedback round	AI acceleration
Weather	Tomorrow's observations	A day	Already happened
Mathematical physics theorems	Lean	Minutes to hours	Already happened (2026-09)
Electronic structure	High-accuracy computation	Hours to days	Fast
Fusion control	Simulator + plasma shots	Seconds to months	Medium
Materials and superconductors	Synthesis + characterization	Weeks to months	Limited by Amdahl
Climate	Observation	A decade	Slow
Quantum gravity	None	∞	Only on subproblems

This generalizes type A / type B (§18): from "is there a referee" to "how fast is the referee".

20.10.2 Cross-field connections

- 1. Spin glasses → neural networks → protein folding.** The energy function of the Hopfield network (1982) is the Hamiltonian of a spin glass, and memories are energy minima (2024 Nobel Prize in Physics). Bryngelson & Wolynes (1987) used the same theory to explain the "energy funnel" of protein folding. The Mézard–Parisi result of §13, OpenAI family 281 (QAOA reaches the optimum on the SK model) and the random k-SAT threshold all use the same mathematics; so do attractor networks for working memory in neuroscience.
- 2. Renormalization group ↔ deep learning.** The renormalization group (RG) coarse-grains details layer by layer and looks at which quantities survive as the scale grows. Mehta & Schwab mapped RG onto restricted Boltzmann machines; Roberts, Yaida & Hanin wrote layer-by-layer representations as an RG flow; Halverson, Maiti & Stoner pointed out that

infinitely wide networks are free fields and that finite width brings interactions. If AI produces a breakthrough in rigorous RG (Problem 7), deep learning theory may benefit directly; the free-energy and phase-transition language of singular learning theory (SLT) is statistical physics to begin with (§22) [speculation].

3. **Non-equilibrium thermodynamics → diffusion models → materials generation.** Diffusion models (Sohl-Dickstein et al. 2015) come from the physical picture of "add noise until equilibrium, then learn the reverse process", and MatterGen now brings them back to materials.
4. **"Simulators fail out of distribution" is a shared motif.** Weather models pushed to climate, ML functionals pushed to new elements, turbulence closures pushed to new flows, fusion surrogate models pushed to new devices. This is isomorphic to distribution shift and goal misgeneralization in alignment, and since physics has ground truth, it is a clean testbed for studying them [speculation].
5. **Complexity theory is the seam between physics and CS.** QMA (DFT), NP-hardness (the sign problem) and the evidence on quantum advantage (Lee et al.) together determine "which problems can only ever be approximated"; AI math output along this seam (OpenAI families 279 and 284, see §19) may have more impact than purely physical results [speculation].

20.10.3 For society and AI safety

- **The physical sciences will be where autonomous AI research gets going first**, and capability spillover will show up first at the physical execution end (autonomous labs, device control, simulation-driven design), forming a compute feedback loop (fusion power purchase agreements, chip-cooling materials).
- **Dual use:** chemistry is the most direct (toxins, energetic materials), followed by nuclear and plasma. Cheap climate simulators also lower the barrier to simulating unilateral stratospheric injection (geoengineering), which is exactly the out-of-distribution regime where simulators tend to be "confident but wrong" [speculation].
- **Epistemics:** the volume of claims grows by orders of magnitude while review capacity stays the same. Experimental science needs its own "Lean": autonomous replication plus raw data.
- **Specification-gaming examples:** forced NS, DFT density drift, and early GraphCast "smoothing out" extreme events because it optimized mean squared error.
- **A real-world testbed for GSAI:** fusion control, chemical processes and power grids all have good physical models and suit "bounded certificates + runtime monitoring".

20.11 Sources for this section

Sources are listed at the end of each card. In addition: OpenAI math release <https://github.com/openai/math> (CONTENTS.md and lean/docs, downloaded 2026-10-09); Hopfield 1982, PNAS 79:2554; Bryngelson & Wolynes 1987, PNAS 84:7524; Mehta & Schwab arXiv:1410.3831; Roberts, Yaida & Hanin arXiv:2106.10165; Halverson, Maiti & Stoner arXiv:2008.08601; Sohl-Dickstein et al. arXiv:1503.03585. Not checked online; spot checks recommended: ITER deuterium-tritium operation in 2039; CPU hours for gyrokinetic simulations; MatterGen's measured bulk modulus; the number of digits of QED precision; the FeMoco resource-estimate summary, the count of "about 25 mathematical physics families", and SPARC's delay to 2027 (all secondhand).

PART II · OUTLOOK · BIOLOGY

21. Biology, medicine and neuroscience: the bottleneck is measurement, not proof

Tags: **[established]** textbook-level consensus; **[recent]** results from 2024–2026, which may still be revised; **[speculation]** inference without direct evidence; Δ unverified; check it again yourself before citing.

21.0 The bottom line

1. **Biology has no "list of conjectures waiting for proof."** Math and physics can list problems like "the Riemann hypothesis" or "Navier–Stokes regularity": clearly stated, missing only a proof (**type A**). Biology can list almost none. The great majority of its hard problems are **type B**: we don't even know which quantities to measure or which equations to write down. So the **direct** contribution of "AI finishing math" to biology is small. This is structural: a biological system evolved over 4 billion years, has no design document, and every one of its parts is an exception to some other part.
1. **But there is one real exception: sampling problems in the sense of statistical mechanics.** Protein dynamics and free energy calculations come down to "computing a partition function over a 10^4 – 10^6 -dimensional space": there are variance bounds, convergence results, things that can be proved. This is the only card in this section where "math is the bottleneck" holds up (Problem 1). Here AI does not work by proving theorems but by **learning a generative model that samples directly from the Boltzmann distribution** [recent].
1. **The second exception is causal identifiability and experimental design.** Behind both the virtual cell (Problem 3) and clinical trials (Problem 5) sits a clean math problem: *given the*

set of perturbations I can perform, which causal structures are identifiable? What is the most sample-efficient perturbation design? This sits at the edge of type A. But it is not the actual rate-limiting step. The rate-limiting steps are **not being able to run that many perturbations**, and **evaluation metrics being gamed until they break**.

1. **Bottlenecks ranked by how often they appear: experimental throughput > data with causal design > concepts > institutions > compute > math.** Two pieces of evidence: (a) in 2025 a *Nature Methods* paper tested 7 deep learning models (including 5 single-cell foundation models) on predicting the effects of gene perturbations, and **not one beat a baseline a schoolchild could compute: "add up the log fold changes of the two single perturbations"** [recent]; (b) Arc Institute's first Virtual Cell Challenge drew more than 1,200 teams, and on the MAE component **no model beat the baseline**; the organizers clipped the penalty at 0, so the metric stopped working and was openly gamed [recent]. What's missing is not model capacity but **data and metrics that can push models into a corner**.

1. **"AI's other capabilities" matter an order of magnitude more than "AI's math ability."** Four things: **automated labs** (OpenAI × Ginkgo used GPT-5 to run 36,000 cell-free protein synthesis recipes on 580 automated plates, cutting cost per gram of protein by 40% and raising yield by 27% [recent]); **generative design** (RFdiffusion2 solved all 41 enzyme design benchmarks, where the previous best solved only 16 [recent]); **hypothesis/literature agents** (an epigenetic target recommended by Google's AI co-scientist really was anti-fibrotic in human liver organoids [recent]); **predictive models as "virtual instruments"** (the MICrONS digital twin can predict how more than 10,000 neurons in mouse visual cortex respond to **unseen** stimuli [recent]). None of the four requires proving any theorem.

1. **Medicine's time constants will slow everything down, and this is an institutional problem, not a scientific one.** The overall clinical development success rate is 7.9% (Phase I → approval, 9,704 programs) [established]; the narrowest gate is Phase II→III at 28.9% [established]; 40–50% of failures are due to lack of efficacy [established]. AI-discovered molecules have an 80–90% success rate in Phase I (40–65% for human-discovered ones), but in Phase II they fall back to about 40%, the same as the industry average [recent, small sample]. **This shows precisely that AI is currently solving "is the molecule good?" (chemistry), not "is the target right?" (biology).**

1. **Neuroscience is the field most tightly linked to AI interpretability, and the link is "bad news."** The complete connectome of the nematode *C. elegans* has existed since 1986: 302 neurons, about 7,000 synapses. **Forty years later we still cannot predict its behavior from the connectome** [established]. If **"complete wiring diagram + complete activity recording"** still doesn't add up to **"understanding,"** then mechanistic interpretability's goal of **"fully understanding a model"** needs to be redefined as **"making verifiable**

causal predictions on specific questions." The good news is that artificial networks are easier than brains: you can read every weight, interventions cost nothing, and you can retrain ten thousand times (Problem 6).

1. **Biosecurity is the only field where "AI capability directly triggers deployment restrictions" has already happened, and its bottleneck is 100% institutional.** US EO 14292 (2025-05) required OSTP to update the 2024 nucleic acid synthesis screening framework within 90 days; as of July 2026 it **still had not been updated** (CRS) [recent]; there is **no federal legislation** for mandatory screening. Meanwhile, in May 2025 Anthropic activated ASL-3 with Claude Opus 4, explicitly because it **could not rule out** CBRN risk [established]. The threat side follows a capability curve that doubles every 129 days; the defense side follows the legislative cycle.

21.1 Why these problems

Three dimensions: scale of impact; an honest answer to "is math the bottleneck?"; and tractability for AI. The second dimension does **not** mean picking only cards where "math is the bottleneck." Several cards where "math is not the bottleneck at all" were chosen on purpose (Problems 2, 5, 8), because their answer to "what happens after AI solves math?" is exactly "**not much, because something else is in the way,**" and that answer is itself informative. The eight cards form a chain:

Protein dynamics (Problem 1, physical layer) → de novo design (Problem 2, engineering layer) → predicting cell behavior (Problem 3, systems layer) → how the genome encodes a body (Problem 4, developmental layer) → clinical translation (Problem 5, medical layer). The two side branches are **the brain** (Problem 6, the hardest special case within the systems layer) and **aging** (Problem 7, a phenomenon in which all layers fail together); **biosecurity** (Problem 8) is the shadow side of the capability gains in Problems 1, 2 and 4.

Not chosen, but worth mentioning: Epidemiological modeling: the math of SEIR has long been good enough; the 2020 misses happened because data on contact matrices, behavior change and testing bias were unavailable, not because the equations were wrong. It is the cleanest example of "math is not the bottleneck," but AI's leverage is small. **Evolution and fitness landscapes:** the same kind of mathematical object as protein energy landscapes and neural network loss landscapes, but the value is mainly epistemic. **Immunology: TCR–antigen recognition and vaccine design:** the practical impact is huge; it was left out only because its bottleneck structure overlaps heavily with Problems 2 and 3. **If you pick one underrated direction, this is the first choice.** **The microbiome:** lots of correlation, very few causal conclusions; intervention experiments in humans are impossible, and community dynamics are not identifiable. **Protein aggregation** (tau, α -synuclein, A β) is physically the same as Problem 1; for **agriculture and photosynthesis engineering** the bottleneck is the

field-trial cycle (one season per year); for **brain–computer interfaces and whole-brain emulation** the bottleneck is Problem 6's "wiring diagram $\neq$ algorithm"; for **metabolic engineering** the bottleneck is the experimental throughput of strain engineering.

21.2 Problem 1: Protein dynamics and function: the truly hard half after AlphaFold

In one sentence Static structure is basically solved, but function comes from the protein's **probability distribution** over a set of conformations and the **free energy differences** between them. Computing this distribution is a high-dimensional sampling problem, and this is the place in all of biology where "math/algorithms really are the bottleneck" stands up best.

Background Assign an energy to each 3D configuration of a protein (the sum of electrostatics, hydrogen bonds and the hydrophobic effect) and you get an **energy landscape** with thousands of dimensions; folding is downhill motion on this landscape.

Levinthal's paradox (1969) [established]: take 100 residues and count only 3 conformations per residue (a huge underestimate); the total number of configurations is $3^{100} \approx 5 \times 10^{47}$. At 10^{-13} seconds per step, exhaustive search would take 1.6×10^{27} years, 10^{17} times the age of the universe; yet real folding takes microseconds to seconds. The resolution is that the landscape is not randomly rugged but **funnel-shaped (folding funnel)**: the closer to the correct structure, the systematically lower the energy. Its mathematical origin is Bryngelson and Wolynes's 1987 derivation from the statistical mechanics of **spin glasses** [established], and the same math was later used to analyze the loss landscapes of neural networks (a Hopfield network is a spin glass). **This is the section's first cross-field link: spin glasses $\leftrightarrow$ protein folding $\leftrightarrow$ neural network training** (same origin as Mézard–Parisi in §13).

AlphaFold2/3 brought structure prediction to experimental accuracy [established], and Hassabis and Jumper received half of the 2024 Nobel Prize in Chemistry for it (the other half went to Baker for protein design) [established]. But AlphaFold gives **a photograph**, while function comes from **a movie: allostery** (a molecule binding at site A changes the shape of site B; many drug targets rely on this mechanism), **cryptic pockets** (invisible in static structures, opening only briefly as the protein "breathes"; many "undruggable" targets actually have them), **intrinsically disordered proteins** (IDPs, about 30% of the human proteome), and **conformational free energy differences**. The mathematical core: the probability of configuration x follows the Boltzmann distribution $p(x) \propto \exp(-E(x)/kT)$; every quantity we want is an expectation under it, and the partition function is a 10^4 – 10^6 -dimensional integral that can only be sampled. $p_A/p_B = \exp(\Delta G/kT)$, and at room temperature $kT \approx 0.6$ kcal/mol:

ΔG	Probability ratio	What it means
0.6 kcal/mol	$\sim 2.7 : 1$	Both conformations are common
1.4 kcal/mol	$\sim 10 : 1$	One takes 91%
2.8 kcal/mol	$\sim 100 : 1$	One takes 99%
4.2 kcal/mol	$\sim 1000 : 1$	The other is essentially invisible

So you need **1 kcal/mol accuracy** just to tell "common" from "rare" conformations, and that is about one weak hydrogen bond, a few parts in 10,000 of a protein's total energy.

Molecular dynamics (MD) needs a step of about 1 femtosecond, while conformational changes take microseconds to milliseconds, **so you need 10^9 – 10^{12} steps**; simulating 1 millisecond takes **about 5–30 GPU-years** at a few hundred nanoseconds per GPU-day, and longer for bigger systems [speculation, hand estimate at 100–500 ns/day].

Why it matters Small-molecule drugs work by binding pockets on proteins. Computing conformational distributions accurately = systematically enumerating cryptic pockets = turning "undruggable" targets into druggable ones (the global prescription drug market is on the order of \$1.5 trillion per year). Enzyme catalysis comes from stabilizing the **transition state**; if you can't compute free energies, you can only rely on trial and error (Problem 2).

What's blocking it: is math the bottleneck? Math/algorithms are one of the main bottlenecks, alongside force field accuracy and compute, and the latter two are closing in. This is the only card in this section that is willing to say "math really is the bottleneck." Sampling is math/algorithms (replica exchange, Markov state models and Boltzmann generators each brought order-of-magnitude speedups); force field accuracy is physical modeling + data (classical force fields have systematic errors; quantum chemistry is accurate but 10^6 times more expensive); compute has a hard wall of 10^{12} steps; validation data is an experimental bottleneck (NMR, multi-conformation cryo-EM, HDX-MS and single-molecule FRET are scarce and indirect).

What AI can do (a) Math/proof ability: not much, but not zero. Things that can be formalized: lower bounds on the effective sample size of importance sampling under a learned proposal distribution (type A: clearly stated, missing a proof); formal verification of free energy identities such as the Jarzynski equality and the Crooks theorem; consistency theorems (KL distance from the true distribution below $\epsilon \Rightarrow \Delta G$ error below $f(\epsilon)$). Going by the curves in §18, it is reasonable to expect type A problems like these to start falling in batches from mid-2027 [speculation]. **(b) Other abilities: the main battlefield.** The shift in thinking is to learn the distribution directly instead of simulating trajectories. Microsoft Research's **BioEmu** [recent] was trained on more than 200 milliseconds of MD trajectories + static structures (including the 100 million+ proteins in the AlphaFold database) + about

500,000 experimental stability measurements. On a single GPU it **generates thousands of statistically independent configurations per hour**, reaches relative free energies at the **1 kcal/mol** level, and can capture **cryptic pocket formation, local unfolding and domain rearrangement**. Official materials say it compresses several GPU-years into less than 1 GPU-hour, a speedup of about 100,000× (△ this figure comes from promotional material / a video transcript, not from the paper itself).

If solved: first-order impact Conformational ensembles and free energy differences computed in hours; cryptic pockets go from chance discoveries to systematic enumeration, and the "undruggable" category shrinks sharply [speculation]; allosteric drug design goes from art to engineering; enzyme design gets a real scoring function; IDPs and phase separation go from "can only be described" to "can be predicted."

If solved: second-order impact **The bottleneck of Problem 2 opens up.** Design is currently stuck on "no oracle" (no way to score designs in silico); computing free energies accurately = having an oracle = success rates jump from single-digit percentages to a usable level. This is the strongest causal chain within this section. **The failure mode shifts:** once the binding step becomes reliable, clinical failures will show up **more purely** as "the wrong target was chosen," so the Phase II failure rate **will look harder to improve in the short term:** a fixable layer of problems has been peeled away, exposing the layer that can't be fixed.

Timing Problems of this kind **can be judged automatically** (they can be compared against MD and experiment, with quantitative error metrics), so they belong to the class AI will crush. **Optimistic** by end of 2027: $\Delta\Delta G$ for common globular proteins reaches "can replace experiment" accuracy (<0.5 kcal/mol) and enters the standard workflow at pharma companies.

Counterarguments and uncertainty The "1 kcal/mol" may have been measured on an easy subset (△ BioEmu's performance on hard classes such as membrane proteins, large allosteric machines and glycosylated proteins is unverified). Generative models learn the training distribution, not physics; out of distribution they may be confidently wrong, and "knowing what they don't know" is not yet solved.

Signals to watch Pharma companies write conformational ensemble prediction into standard decision workflows rather than treating it as an exploratory tool; prospective blind tests targeting dynamics appear (CASP adding a conformational-ensemble track would be a strong signal); accuracy moves from 1 to 0.3 kcal/mol; a drug driven by a cryptic pocket enters the clinic and **succeeds**. The reverse signal: in 2027 people are still gaming the same metrics on the same set of public proteins.

Sources: Abramson et al., *Nature* 630:493–500 (2024), DOI 10.1038/s41586-024-07487-w; Lewis, Hempel, Jiménez-Luna et al., "Scalable emulation of protein equilibrium ensembles with generative deep learning," *Science* (2025), DOI 10.1126/science.adv9817; Bryngelson & Wolynes, *PNAS* 84:7524–7528 (1987).

21.3 Problem 2: De novo protein design: binders, enzymes, and the real bottleneck of "no scoring function"

In one sentence We can already design entirely new proteins from scratch that bind a specified molecule or even catalyze a specified reaction, but success rates are only a few percent to a few tens of percent. The bottleneck is that **we can't judge in silico whether a design is good**, so we can only synthesize and test in bulk. This is an experimental throughput problem and has nothing to do with math.

Background De novo design means stating the function you want first, then computing a sequence that has never appeared in nature (in contrast to traditional directed evolution: random mutation + screening). The two kinds of targets differ in difficulty by an order of magnitude. A **binder** must bind a specified target specifically; the job is to match shape and charge at the interface, and this is now fairly mature. **Enzymes** are much harder, because catalysis is not about "binding tightly" but about **stabilizing the transition state more than the substrate and the product**: a catalyst's job is to make "that distorted, short-lived configuration halfway through" sit more comfortably than the substrate does. If the pocket is too friendly to the substrate, the substrate doesn't want to change (the "substrate trap"). So a few catalytic residues have to be placed at specific positions and orientations, with an error tolerance of about **0.5–1 ångström and a dozen or so degrees**, while the whole backbone keeps jiggling at room temperature with a larger amplitude than that. The technical route is diffusion models: **RFdiffusion** from the Baker lab (the Institute for Protein Design, IPD) takes the diffusion models used for image generation and applies them to protein backbones, denoising step by step from noise, with **conditioning** ("these three residues must sit at these coordinates"); ProteinMPNN then fills in the sequence.

Two results from 2025: the numbers alone show where the bottleneck is. (a) **Designed serine hydrolases** [recent: Lauko, Pellock et al., *Science* 2025]: the Ca RMSD between the crystal structures and the design models is **under 1 ångström** (the intended shape really was built); they obtained **five folds not found in nature**; catalytic efficiency reached up to the order of $2.2 \times 10^? \text{ M}^{-1}\text{s}^{-1}$ ($\triangle$ the exponent was cut off in the search results; check the original before citing); **the key number: in the first round only 1.6% of designs had any catalytic activity.** (b) **RFdiffusion2** [recent, IPD announcement 2025-04]: the input is relaxed from "atomic coordinates" to "the chemical transformation I want," and the model decides on its own how to lay out the active site. On a newly built set of 41 enzyme design benchmarks it **solved all of them**, where the previous best tool solved only 16; wet-lab

experiments produced active enzymes for 5 reactions, **testing fewer than 100 designs per case**. $\triangle$ The IPD announcement page lists no limitations; check the bioRxiv preprint for specific numbers.

How to read these two sets of numbers: "1.6% active" plus "fewer than 100 designs tested per case" is the whole story: **the design algorithm outputs a pile of candidates, and you can't rank them**. You can only synthesize tens to hundreds of them and see in the wet lab which ones work. So the speed of the whole field = **the throughput of synthesis and testing**.

Why it matters Half of the 2024 Nobel Prize in Chemistry went to David Baker (computational protein design) [established].

What's blocking it: is math the bottleneck? Not math at all. The bottlenecks are (1) experimental throughput and (2) the lack of a computational oracle, and (2) is rooted in Problem 1's free energy calculation. Being unable to score and rank designs is a physics/algorithms problem (= Problem 1), as the 1.6% first-round hit rate shows; throughput is experimental, and "<100 designs per case" directly reflects the throughput ceiling; backbone generation is basically solved (41/41 solved); the catalytic cycle (multiple steps, cofactors, substrates entering and leaving) is a conceptual bottleneck; **math/proof is not a bottleneck: there is no statement here waiting to be proved**.

What AI can do (a) Math/proof ability: almost irrelevant. The only formalizable angle is "can we give a design a guaranteed statement like 'its minimum activity has some lower bound'?" At present this is entirely impossible, and it may not even be formalizable (a protein's behavior in a cell depends on too much context that can't be modeled). This is a typical example of how scarce type A problems are in biology. **(b) Other abilities: all the leverage is here.** First, **generative design itself** (RFdiffusion / RFdiffusion2 / ProteinMPNN / the ESM family), already in use and already effective. Second, **automated labs**, possibly the most underrated line: OpenAI x Ginkgo (2026-02) had GPT-5 autonomously design cell-free protein synthesis recipes, executed by a robotic cloud lab, over 6 closed-loop rounds, 36,000+ recipes and 580 plates, cutting cost per gram of protein by 40% and raising yield by 27% [recent] ($\triangle$ publication status unclear). It shows that experimental throughput itself can be attacked by AI; if the same loop were used to synthesize and test designed proteins, the limit of "fewer than 100 designs tested per case" would loosen by one to two orders of magnitude. Third, **active-learning loops**: the model proposes $\rightarrow$ robots synthesize and test $\rightarrow$ data flow back into training, turning "no scoring function" into "use experiments as the scoring function, cheap enough to afford".

If solved: first-order impact Going from "I want a protein that binds X" to "a usable molecule in hand" drops from months to days, and success rates rise from single-digit

percentages to tens of percent; most uses of antibodies are taken over by smaller, more stable, cheaper designed proteins [speculation]; industrial reactions not found in nature become designable; diagnostic reagents, biosensors, self-assembling protein nanostructures.

If solved: second-order impact The chemical part of manufacturing starts to "go biological": many high-temperature, high-pressure catalytic processes can be done with enzymes at room temperature and pressure, and the chemical industry accounts for something on the order of 10% of global energy consumption. **The toolbox for Problems 3 and 4 fills up:** if you can design proteins, you can make arbitrary intracellular probes, switches and recorders, which directly speeds up the data generation the virtual cell needs.

Timing This field **has an automatic judge** (the wet lab), so iteration is fast, but the judge is **expensive**, so speed is set by experimental cost rather than by the AI capability curve.

Optimistic by end of 2027: binder design is "delivered on demand" (give a target, get a design with nanomolar affinity within a week, success rate >50%), and single-step enzyme hit rates reach 20–30%.

Counterarguments and uncertainty "41/41 solved" may mean the benchmark is too easy: the team built the benchmark itself, and calibrating difficulty is an old problem, so wait for independent blind tests. "Active" is a long way from "useful": designed enzymes are typically 10^3 – 10^6 times less catalytically efficient than natural enzymes, and historically closing that gap has relied mainly on directed evolution (that is, still on experiments).

Signals to watch **Independent blind tests** (a CASP-like design challenge, with third parties choosing the targets and third parties doing the wet-lab validation); **hit rates made public** (routinely reporting "tested N, M were active" rather than only the best one); a designed binder entering the clinic; **unit cost at cloud labs** (if "synthesize and test one designed protein" drops from the \$1000 level to the \$10 level, everything changes); **closed-loop automation used for design, not optimization:** today it mainly optimizes continuous parameters (reaction recipes), and applying it to "discrete design spaces" (which sequence) is the key turning point.

Sources: Lauko, Pellock et al., "Computational design of serine hydrolases," *Science* (2025), PDB 9DEE/9DEF/9DEG; IPD, "Introducing RFdiffusion2" (2025-04); OpenAI, "GPT-5 lowers protein synthesis cost".

21.4 Problem 3: The virtual cell (short card)

The goal is a model that can answer "if I knock out this gene / add this drug, what happens to this cell?": learn a function $f(\text{cell type, perturbation}) \rightarrow \text{expression change}$ and extrapolate to perturbations and cell types it hasn't seen (the data come from Perturb-seq,

where one experiment perturbs thousands of genes at once and measures the full transcriptome, about 20,000 dimensions, for each cell). **The bottleneck is data and evaluation, not math.** Two 2025 results settled this. The first [recent: Ahlmann-Eltze, Huber, Anders, *Nature Methods* 22:1657–1661 (2025), DOI 10.1038/s41592-025-02772-6] tested 7 models: 5 single-cell foundation models (scGPT, scFoundation, scBERT, Geneformer, UCE) + 2 built specifically for perturbation prediction (GEARS, CPA). **Under double perturbations, every model's error was significantly higher than that of the additive baseline, "add up the log fold changes of the two single perturbations";** when it came to detecting gene–gene interactions (epistasis), **none beat "predict no change"**. For single perturbations of unseen genes, no model consistently beat a mean baseline or a simple linear model either, and **feeding the foundation models' gene embeddings into a linear model matched or beat the models' own decoders**. The authors conclude that predicting experiments that have not yet been done is still far out of reach; the limitation they state themselves is that they used only 4 datasets, all from cancer cell lines.

The second result is **evaluation metrics gamed until they break**: Arc Institute's first Virtual Cell Challenge newly generated 300,000 H1 human embryonic stem cells with 300 genetic perturbations, and **more than 1,200 teams submitted** [recent]. The score combined three metrics, and on one of them, MAE, **no model beat the baseline**; the organizers then **clipped the penalty at 0**, so "doing worse" was not penalized and MAE in effect **stopped influencing optimization**; one participant named their model **"Leaderboard Hacker"** outright. Another finding: Pearson correlation over all genes, a metric still widely reported today, reaches its maximum on the trivial model that "predicts no change." Arc's own summary says the first year exposed gaps in three areas, model accuracy, metric design and biological generalization, and in 2026 the challenge switched to zero-shot (no training set released). [△] Most of these critical analyses come from participants' Substacks and blogs, not from peer-reviewed analysis.

The key interpretation: this doesn't show that "deep learning is useless"; it shows that current data don't contain enough learnable structure. The additive baseline is hard to beat because biology really is approximately additive much of the time, while the non-additive part is sparse and context-dependent and, at current data volumes, can't be told apart from noise. This is a clean statistical problem, not a model-capacity problem, so **"AI solving math" does almost nothing for this card**. What is really needed is 10× more data, data with causal design (choosing perturbations by information gain rather than brute-forcing them at random), and metrics that can't be gamed. Math has only a small real role (causal identifiability, optimal perturbation design, and "can we design a metric for which 'predict no change' is not the optimal solution?"; the last can be formalized as a mechanism design problem and **maps directly onto AI evaluation design** [speculation]); compute is

not the bottleneck. Resources are pouring in: in October 2026 Biohub's Virtual Biology Initiative expanded to a total of **\$1.8 billion**, and its scientific lead Alex Rives said existing data cover about 1 billion cells and that making the models useful will take another order of magnitude (Δ the \$1.8 billion figure appears only in secondhand reports; Biohub's official page gives no amount). **Central timing call: 2030:** reliable for specific cell types + specific pathways, still poor at generalizing across cell types, and the "virtual cell" as a general-purpose tool still doesn't exist; this card depends least on the AI math curve and most on the rate of data generation. The signals most worth tracking are someone independently retesting a new generation of models with the Ahlmann-Eltze protocol and winning, and the results of the zero-shot challenge; among the counterarguments, the one most worth taking seriously is "the transcriptome may simply not be enough" (mRNA is a noisy, delayed proxy variable; if so, the task itself is aimed at the wrong target). **Sources:** Ahlmann-Eltze, Huber, Anders, *Nature Methods* 22:1657–1661 (2025), DOI 10.1038/s41592-025-02772-6; Bunne et al., *Cell* 187(25):7045–7063 (2024), DOI 10.1016/j.cell.2024.11.015; Arc Institute, Virtual Cell Challenge 2025/2026, *Cell* commentary DOI 10.1016/j.cell.2025.06.008.

21.5 Problem 4: How the genome encodes a body (short card)

How do the 3 billion base pairs in a fertilized egg grow into a body with hands, feet and a brain? This is biology's biggest **type B problem**. Coding sequence makes up only about 1.5% of the genome; most of the rest is **regulatory sequence**. **Transcription factors (TFs)** bind specific short DNA sequences (usually 6–12 bases) and then recruit or block RNA polymerase; transcription factors are themselves encoded by genes, so you get a **gene regulatory network (GRN)**. Enhancers can sit thousands to millions of bases away and "reach over" through DNA folding to contact the promoter. The classic case is the stripe 2 enhancer of the fruit fly gene *even-skipped* [established]: about 500 bp carry binding sites for the activators Bicoid and Hunchback and the repressors Giant and Krüppel, and the logic is a "squeeze": a broad activation zone is cut once at the front and once at the back, and the narrow band left over is the second stripe. Bicoid concentration decays exponentially along the embryo, and across the roughly 20 μm (3–4 cells) width of stripe 2 it **drops by only 18%: an 18% difference in input concentration has to become an "all on / all off" gene output**. That requires extreme nonlinearity, and the parameters of the cooperative effects are almost impossible to measure independently.

Math is not the bottleneck; concepts are. The tools (dynamical systems, reaction–diffusion, stochastic processes, Boolean networks) all exist, with room to spare. What's missing is "which one to use, and what to measure." We don't know the right level of abstraction for the developmental program (gene networks? reaction–diffusion of cell-to-cell signals? mechanical forces?); we don't know what to measure (constraining a model would need every TF concentration in every cell at every moment, while what we measure is

the average mRNA of a fixed cell population); and we don't know the criterion for success (predicting expression patterns? morphology? phenotypes after mutation? these differ in difficulty by orders of magnitude). Robustness (Waddington's canalization) is itself mysterious, but **this part can be formalized**: it is a structural stability problem for dynamical systems, and a candidate for "turning type B into type A."

The state of the art in 2026 is DeepMind's **AlphaGenome** [recent: Avsec, Latysheva et al., *Nature* 649 (2026), DOI 10.1038/s41586-025-10014-0]: it takes inputs of up to 1 million base pairs and, in a single call, predicts gene expression, splicing, chromatin accessibility, TF binding, 3D contacts and other signals at once. On variant effect prediction it matches or beats the previous best specialized models on **25 of 26 evaluations** (per the *Nature* press release; the June 2025 preprint said 24/26; △ the reason for the difference has not been checked). **The two limitations the authors list themselves are exactly the ones that matter most**: accurately capturing **distant** regulatory elements (for example, 100,000 bases away) is still a limitation; and it has not been validated for predictions on individual genomes. Yet about **90% of disease-associated common variants fall in non-coding regions** [established], often far from the genes they affect; this is the biggest explanatory gap in modern genetics. Another line of work is genomic language models: Arc Institute's Evo series can **generate** genome sequences. In September 2025, using the phage ΦX174 as a template, it generated thousands of candidates; 302 were selected for synthesis, 285 were successfully assembled, and **16 were active** (a 5.6% hit rate; but in the subset with "≥98% sequence identity to the template" the activity rate was as high as 46%; this contrast suggests the model is mostly making small edits near the template, not doing true de novo design) [recent: bioRxiv 10.1101/2025.09.12.675911]. △ Secondhand reports say this work was published in *Science* in August 2026; this could not be confirmed; the 46% vs 5.6% comparison comes from secondhand reports.

Methodological observation: models like AlphaGenome succeed by **going around mechanism**. They don't build a GRN; they learn a statistical mapping from "sequence → signal," which works for prediction and fails for understanding. **In biology, AI is routing around type B problems rather than solving them: it gives you an oracle, not a theory.** The central timing call is 2032: predictive models are usable within specific tissues and distance ranges, but the "developmental program" as a **theory** still doesn't exist; the open-problem solve-rate curve **does not apply at all** to this card, because there is no "problem statement" here. **Sources**: AlphaGenome (cited above); Small, Blair, Levine, *EMBO J* 11:4047–4057 (1992); Waddington, *The Strategy of the Genes* (1957); McCarty, "AI-Designed Phages," *Asimov Press* (2025-09-17), DOI 10.62211/21er-45fg.

21.6 Problem 5: Clinical failure rates (short card)

The arithmetic of drug development [established: BIO / Informa / QLS, *Clinical Development Success Rates 2011–2020*, 9,704 clinical programs, 12,728 phase transitions]: **the overall probability from Phase I → approval is 7.9%** (the previous edition, 2006–2015, gave 9.6%), and **the narrowest step is Phase II → Phase III at 28.9%**; programs that used patient-preselection biomarkers reached 15.9%, about twice the baseline. Among the causes of failure, **lack of clinical efficacy accounts for 40–50%** ("doesn't work in humans"), toxicity 30%, and poor drug properties 10–15% [established: Sun, Gao, Hu, Zhou, *Acta Pharm Sin B* 12(7):3049–3062 (2022), DOI 10.1016/j.apsb.2022.02.002]. Another clue is human genetic evidence: mechanisms with genetic support are **2.6 times** as likely to succeed as those without [established: Minikel et al., *Nature* 629 (2024), DOI 10.1038/s41586-024-07316-0]. This 2.6× plus the 2× from biomarkers together tell a clear causal story: **the main cause of clinical failure is that we don't know whether this target is really a causal driver of this disease in these patients.**

AI's report card confirms this precisely: AI-discovered molecules have a **Phase I success rate of 80–90%** (40–65% for human-discovered ones), but in **Phase II they fall back to about 40%**, level with the industry average [recent: Jayatunga et al., *Drug Discovery Today* (2024) — $\triangle$ seen only in secondhand reports; the authors themselves flag the sample size as limited]. **This gap shows that AI is currently solving "is the molecule good?" (a chemistry problem), not "is the target right?" (a biology problem).** So this card's bottleneck is not math at all. It is biological understanding of target causality (= Problems 3 and 4), the incompressible time and sample sizes of human trials, and institutions (Bayesian adaptive designs, platform trials and synthetic controls have existed since the 1990s, but regulators have been slow to adopt them: **the bottleneck is not that the math isn't good enough, it's that the institutions don't accept it**). The only corner of the math with real content is the validity of surrogate endpoints (the Prentice criteria give a sufficient condition, but in practice almost no surrogate endpoint satisfies them; this is a clearly stated, mathematically hard type A problem), but it is not the rate-limiting step.

The hardest slowdown argument is time constants: suppose that from today (2026–10) AI made target selection perfect. A new target then takes 1–2 years to reach lead optimization, 1–2 years from preclinical to IND, 1–1.5 years for Phase I, 2–3 years for Phase II and 3–5 years for Phase III, **about 8–13 years in total**; so the claim "AI improved Phase II success rates" **won't have statistically credible evidence until around 2034 at the earliest**, and this argument **does not depend on any assumption about AI capability. Central timing call: 2035:** the overall success rate rises from 7.9% to 10–12%, the Phase I advantage is established, and Phase II improves but not dramatically. **In medicine, the exponential curve of AI capability gets eaten by the linear time constants of institutions.** Two things are

worth remembering on their own: **clinical trials are the strictest judge for AI-for-science**, because they can't be gamed (double-blind, preregistered, reviewed by regulators), so the clinical success rate is the single most credible indicator of whether AI is really speeding up science; and **the clinical trial system (mandatory registration, CONSORT reporting standards, independent data monitoring) is exactly what the AI eval community is reinventing**, so its failure modes (p-hacking, selective publication, outcome switching) will replay precisely, and its countermeasures can be carried over directly [speculation]. The strongest counterargument is that more than a decade of technical progress has made this number not better but worse (9.6% for 2006–2015 → 7.9% for 2011–2020), the strongest rebuttal to every "this time is different" claim; △ to date no "AI-discovered drug" has completed a full Phase III and been approved. **Sources:** BIO / Informa / QLS, *Clinical Development Success Rates 2011–2020*; Sun, Gao, Hu, Zhou, *Acta Pharm Sin B* 12(7):3049–3062 (2022), DOI 10.1016/j.apsb.2022.02.002; Minikel et al., *Nature* 629 (2024), DOI 10.1038/s41586-024-07316-0; △ Jayatunga et al., *Drug Discovery Today* (2024) (secondhand only).

21.7 Problem 6: How the brain computes and learns: and how it relates to AI interpretability

In one sentence We now have a complete connectome of the fruit fly (139,255 neurons) and a "wiring diagram + activity recording" for one cubic millimetre of mouse cortex, but **we still cannot derive the algorithm from the wiring diagram** — and that failure has direct implications for how mechanistic interpretability sets its goals.

Background Problem one: credit assignment. The brain has about 8.6×10^{10} neurons and 10^{14} – 10^{15} synapses. Learning means adjusting synaptic strengths, and **a synapse deep in the network can only see the activity of the two neurons on either side of it; it cannot see the animal's behavioural outcome.** Backpropagation solves this in artificial networks, but it faces clear biological obstacles [established, review in Lillicrap et al., *Nat Rev Neurosci* 21:335–346 (2020)]: **weight transport** (the backward pathway requires the **transpose** of the forward weight matrix, while biological synapses are one-directional), **temporal separation** (it requires "finish the forward pass → backpropagate → update", which needs a global clock, whereas the brain runs continuously), and **non-locality** (weight updates need a global error signal, but synaptic plasticity is local). Each candidate answer has its problems: **feedback alignment** (use fixed random backward weights; works in small networks, clearly worse than backpropagation in large networks and on hard tasks); **predictive coding** [Rao & Ballard, *Nat Neurosci* 2:79–87 (1999)] (each layer predicts the activity of the layer below it and sends only the **prediction error** upward, so the error signal is locally available; extremely influential, but both its performance as a learning algorithm and its accuracy as a description of the brain are still disputed); **"infer the activity first,**

then change the weights" (prospective configuration; $\triangle$ this round verified only the authors, journal, volume, issue and pages — Song, Millidge, Salvatori, Lukasiewicz, Bogacz, *Nat Neurosci* 27(2):348–358 (2024); **the title and the specific content were not verified**); **equilibrium propagation** (in energy-based networks the gradient can be read out locally from the difference between two equilibrium states; mathematically elegant, biological evidence thin); **three-factor plasticity rules** (a local Hebbian term $\times$ a global neuromodulatory signal; **the strongest biological evidence**, but it only gives you "reward-modulated Hebbian learning", which carries far less information than a gradient and cannot explain how the brain learns complex internal representations). **So the state of play is five or six candidate frameworks coexisting, each with partial evidence, none of them established** — a textbook type B problem: it is not that "a proof is missing", it is that "we do not know which equation to write down".

Problem two: is a complete wiring diagram enough? There is a forty-year, decisive experimental answer. The complete connectome of *C. elegans* was published in **1986** [established: White, Southgate, Thomson, Brenner, *Phil. Trans. R. Soc. Lond. B* 314:1–340 (1986)]: 302 neurons, about 7,000 synapses, every one annotated. **Forty years later we still cannot predict the worm's behaviour from that connectome** — because the same wiring diagram with different synaptic strengths, neuromodulatory states and ion channel compositions produces completely different dynamics: **the connectome gives you the topology of the graph; it does not give you the weights on the edges or the transfer functions at the nodes.** In 2024–2025 the scale went up but this problem did not go away: the **whole-brain connectome of the fruit fly** [established: Dorkenwald et al., *Nature* (2024–10–02), a package of nine papers] has **139,255 neurons** and about **5×10^7 chemical synapses**, and was built from **21 million** electron microscope images $\rightarrow$ automated AI segmentation $\rightarrow$ **human proofreading**; **MICrONS** [established: *Nature* 640:435–447 (2025), DOI 10.1038/s41586-025-08790-w] covered $1.4 \times 0.87 \times 0.84$ mm of mouse visual cortex, about **200,000 cells, 520 million synapses and 4 km of axons** ($\triangle$ the numbers differ between sources; an NIH figure caption says 120,000 cells were reconstructed), and **the key innovation was recording function before doing structure** — calcium imaging recorded the activity of **76,000 cells** watching natural-scene video while the mouse was alive, and only then was the tissue fixed, sectioned and reconstructed by electron microscopy.

Models constrained by the connectome are the most interesting step. A **whole-brain LIF model** [recent: Shiu et al., *Nature* 634 (2024), DOI 10.1038/s41586-024-07763-9] built a whole-brain leaky integrate-and-fire model from the wiring plus neurotransmitter identity **with no fitted parameters at all** (synaptic strengths taken directly from synapse counts), and still **predicted accurately** which downstream neurons respond when sugar- or water-

sensing taste neurons are activated and which ones are required to initiate feeding, and **verified** this with optogenetics and behavioural experiments. **Connectome constraints + deep-learning fitting** [recent: Lappalainen, Tschopp et al., *Nature* 634:1132–1140 (2024), DOI 10.1038/s41586-024-07939-3] took the measured connectivity of 64 cell types in the fly optic lobe as the structure and **fitted the unknown single-neuron and single-synapse parameters with deep learning** (the training objective was "detect visual motion"); the neural activity the model predicts matches experimental measurements from **26 independent studies**. The **"digital twin"** [recent: Wang et al., *Nature* (2025), DOI 10.1038/s41586-025-08829-y] was trained on recordings from real mice watching movies and can predict **the responses of tens of thousands of neurons to videos and images they have never seen; once the shared core is trained it transfers to new individuals, sometimes with only 4 minutes of new recording**. More striking still, this purely functional model can also **run anatomy backwards** — it predicted the cell type, position and connectivity of thousands of neurons, verified against electron microscopy data from the same mouse. **What the three works say together:** the connectome **by itself** does not give you the algorithm, but the connectome **as a constraint** makes models far more credible and far more data-efficient — this is a win for "structural prior + learning", not a win for "reading the algorithm straight off the structure".

Why it matters Epistemically this is the question of how matter produces mind; medically, treatment of psychiatric and neurodegenerative disease has long been stuck on "we don't know what goes wrong at the circuit level"; the implications for AI need care — historically, neuroscience's contribution to AI has been **coarse-grained inspiration at the architecture level** (convolutional networks ← simple/complex cells; reinforcement learning ← the TD signal in dopamine), not copying of details.

What's blocking it: is math the bottleneck? Mostly a conceptual bottleneck (we don't know what to write down), with math as the second real bottleneck (we lack the theory) and data third. This is the purest type B problem in the section, but it has a rare **property: the type B component could plausibly be converted into type A**. Not knowing which learning rule the brain uses is a conceptual bottleneck (five or six frameworks coexisting for forty years); **the missing theory of "what kinds of learning rule can be implemented locally" is a formalizable mathematical bottleneck**; the wiring diagram not being the algorithm is conceptual plus data; synaptic strengths and neuromodulatory states being unmeasurable is a technical bottleneck (electron microscopy gives morphology but not strength; in vivo recording gives activity but not the full connectivity).

What AI can do (a) Math/proofs: the second most likely place in this section for AI math to really help (the first is Problem 1). Impossibility and equivalence theorems for learning rules; **identifiability in dynamical systems** (can circuit parameters be uniquely determined

from a finite activity recording? this is a system identification problem, there is something provable here, and it tells experimentalists directly "more recording will not help, you need a different kind of measurement"); **equivalence classes of representations** (when do two networks count as "implementing the same algorithm"? **this is the same question** in cross-individual and cross-species comparison in neuroscience and in "did two models learn the same circuit" in AI interpretability, and today's tools — RSA, CCA, CKA — are all heuristic and lack a theoretical basis). **(b) Other capabilities: connectome reconstruction is itself an AI result** (the initial segmentation of the fly connectome was done by AI, with humans only proofreading; without deep learning these datasets would not exist — this is the least disputed form of AI's contribution to science); **digital twins as a substitute for in vivo experiments** (4 minutes of new recording to fit a new individual, then any number of stimulus experiments in the computer, turning "one mouse and several weeks per experiment" into "a few seconds"); **connectome-constrained model fitting; automated hypothesis generation and literature integration.**

If solved: first-order impact Knowing what learning algorithm the brain uses → circuit-level understanding of psychiatric disease → genuinely mechanistic treatment (rather than today's "adjust neurotransmitter levels"); a jump in the performance of neural prostheses and brain-computer interfaces (you know what signal to write in); possibly architectural inspiration for AI (though historically the contribution from this direction has been overrated).

If solved: second-order impact 1. It changes how mechanistic interpretability sets its goals.

System	Scale	What we have	What we can do
<i>C. elegans</i> (1986–)	302 neurons / 7,000 synapses	complete connectome, complete cell lineage, complete genome	cannot predict behaviour from the connectome
Fruit fly (2024–)	139,255 neurons / 5×10^7 synapses	complete connectome + transmitter identity	testable predictions in the two best-studied circuits
Mouse V1, 1 mm ³ (2025–)	2×10^5 cells / 5.2×10^8 synapses	connectome + functional recordings from the same batch of cells	can predict activity (digital twin), cannot state the algorithm
Large language models	10^{11} – 10^{12} parameters	all the weights, all the activations, unlimited interventions, retraining possible	SAE features, partial circuits, causal intervention at the behavioural level

The fourth row is the good news: artificial networks **win outright** on observability — all weights readable, all activations recordable, interventions free, retrainable ten thousand times with different random seeds, any subset ablatable. So "we don't even understand a 302-neuron worm, therefore we will never understand an LLM" **is an invalid inference**; the

epistemic conditions in the two cases are completely different. **But the first three rows are the bad news, and the bad news is more specific:** they show that **"having a complete structural description" does not automatically produce "understanding"**. Understanding requires a **theory** — a set of judgments about which quantities to look at and which abstractions to make. That maps exactly onto the type A / type B distinction: from "given a model, compute its mathematical properties" (type A, where AI will dominate) to "judge what this mathematical quantity has to do with safety" (type B, which humans mostly carry). **Connectomics is a forty-year natural experiment on that same distinction in another field, and the result is: the type B part was not automatically solved by the growth in the amount of data.**

2. Jonas & Kording's microprocessor paper is the sharpest bridge between these two fields. In 2017 they took a MOS 6502 microprocessor (whose circuit and function we know **completely**) and applied the standard analysis methods of neuroscience to it — lesion studies (cut a single transistor and see which game breaks), single-unit recording, correlation analysis, dimensionality reduction — then asked whether those methods let you understand the 6502. **The answer is no:** the methods produce a pile of statistically significant but misleading conclusions ("this is the *Donkey Kong* transistor") [established: *PLOS Comput Biol* 13(1):e1005268, DOI 10.1371/journal.pcbi.1005268]. This paper's implications for AI interpretability are larger than for neuroscience, because **AI interpretability is using the same set of methods** (ablation, probes, correlation analysis, dimensionality reduction). It does not say those methods are useless; it says **they need a theoretical framework to interpret their output, otherwise they will systematically produce spurious mechanistic stories.**

3. A few concrete cross-field questions: can the **validation** methods connectomics has developed (for example using optogenetics to check the predictions of an LIF model) be carried over to interpretability? Connectomics' classification of "cell types" and interpretability's classification of "features" are mathematically the same kind of problem (finding discrete structure in a high-dimensional representation) — can the methods be shared? Most important of all, what would the Lappalainen move of "use structure as the constraint and learning to fill in the parameters" look like applied in reverse to interpretability (posit a circuit structure, fit the parameters, and see whether it predicts the model's out-of-distribution behaviour)? Also, neuroscience has a philosophical literature on "what counts as a good mechanistic explanation" (Craver, Bechtel) that the interpretability community has essentially never cited.

Timing Optimistic by 2029 several major fly behaviours (feeding, grooming, navigation, courtship) have end-to-end testable circuit models, and the "learning rule" question has a widely accepted answer in **invertebrates**; **central** by 2033 the mouse whole-brain connectome is finished or close to it and local cortical circuits have digital twins, but there is

still no consensus on "the brain's learning algorithm" (a 10–15 year schedule + human proofreading as the hard bottleneck + conceptual problems not being solved automatically by data); **pessimistic** after 2040 there is more and more data and no progress in theory, and the field becomes "better and better oracles, no understanding" (the default ending for type B problems, with the forty years of *C. elegans* as precedent). **This is the textbook case of "type B problems land slowly"**: the open-problem solve-rate curve **does not apply here at all**; what is useful is an indirect inference from the METR curve — when AI can autonomously do research tasks that take "a human expert several months" (mid-2027 to 2028), it may start to take on the work of "sweeping the hypothesis space", and "which learning rule the brain uses" happens to be a problem where **the hypothesis space is small but each hypothesis is expensive to test** [speculation].

Counterarguments and uncertainty The counter-counterargument to "the connectome is not enough": perhaps *C. elegans* failed because 302 neurons is too few (no statistical regularities to exploit), and a fly with 10^5 neurons is actually easier (population-level statistical descriptions become available). Shiu et al.'s whole-brain model making testable predictions without fitting any parameters supports that optimistic reading — **this is the most important uncertainty on this card**. "Perhaps there is no single learning rule": different brain regions and different developmental stages may use different mechanisms, in which case the question itself is misposed.

Signals to watch the prospective prediction success rate of whole-brain models (Shiu-style models making predictions about circuits that **have not been specifically studied** and then being confirmed experimentally; this is the cleanest indicator); **an experimental verdict on the "learning rule"** (an experiment appears that can distinguish predictive coding, feedback alignment and three-factor rules in vivo); **how far human proofreading gets automated** (if the molecular barcoding route works, the mouse whole-brain schedule moves up substantially); **the cumulative mm³ curve for the mouse whole brain** (1 mm³ → the 10 mm³ project → 500 mm³; watch the slope); **cross-field citation** (mechanistic interpretability papers starting to cite the methodological literature of connectomics — almost none do now, which is an obvious gap); the contrary signal is that by 2030 the main results are still "bigger connectomes + better digital twins" with no new theoretical framework.

Sources: White et al., *Phil. Trans. R. Soc. Lond. B* 314:1–340 (1986); Dorkenwald et al., *Nature* (2024-10-02); MICrONS Consortium, *Nature* 640:435–447 (2025), DOI 10.1038/s41586-025-08790-w.

21.8 Problem 7: Aging (short card)

△ **Verification status of this card:** the citations on this card (DOIs are all given) were **not re-verified** this round; they come from background knowledge rather than from this round's search. Treat every citation as "to be checked", especially the 2025–2026 ones.

Aging is the single largest risk factor in the human disease burden, but the field is stuck on three things: **no causal theory, no validated readout, and institutionally, aging is not an indication** — and not one of the three is a math problem. The most useful operational definition comes from mortality: the **Gompertz law** [established, 1825] says that adult annual mortality rises roughly exponentially with age, with a doubling time of about **8 years** in humans (if annual mortality is 0.1% at 30, then it is 0.2% at 38, 0.4% at 46, about 3.2% at 70 and about 13% at 86); the mortality rate doubling time is about 0.3 years in mice and about 8 years in humans, a factor of about 27, and that is exactly where the difference in lifespan comes from. The drawback of this definition is that **you have to let a population live until it dies before you can measure it**. In 2013 López-Otín et al. proposed 9 hallmarks of aging, extended to **12** in 2023 [established: *Cell* 186(2):243–278 (2023), DOI 10.1016/j.cell.2022.11.001 — △ not verified], but **that framework is not causal**: it is a checklist, not a mechanism diagram, the 12 hallmarks are interrelated and we **do not know which one is upstream**. The two opposing frameworks — **damage accumulation** (aging is the accumulation of random damage, no program, only entropy, implying it is irreversible and can only be slowed) and **program/information** (aging is the ordered loss of regulatory information, cells "forget" what type they are supposed to be, implying partial reversibility) — imply completely different intervention strategies, and **existing data cannot decide between them**. That is the central fact of this card.

The bottleneck is the readout, and it is a circular dependency. The gold standard is to let a group of animals or people live until they die and compare survival curves (3–4 years for mice, decades for humans); what you want is an indicator that can be measured within 6 months and that **causally** predicts lifespan — this is the surrogate endpoint problem, exactly isomorphic to Problem 5. None of the surrogate endpoints in the aging field **has passed causal validation**: in 2013 Horvath used methylation at 353 CpG sites in a linear regression to predict age with an error of about 3.6 years, and it works across tissues [established: *Genome Biology* 14:R115, DOI 10.1186/gb-2013-14-10-r115], but the clock was trained on **age**, so what it measures is "things correlated with age", not necessarily "things that cause aging" — an intervention can turn the clock back without extending lifespan. And to validate whether a clock is a valid surrogate endpoint you need a large number of "intervention → change in the clock → actual lifespan" triples, **and producing those triples requires running the full lifespan experiment**; this circular dependency is where the field is really stuck. The genuinely clean evidence comes from the NIA **Interventions Testing**

Program (ITP): three independent sites, genetically heterogeneous mice, double-blind, pre-registered, **and it publishes failures** — the result is that the great majority of candidate compounds do nothing, and rapamycin is one of the few clear successes [established: Harrison et al., *Nature* 460:392–395 (2009), DOI 10.1038/nature08221 — [△](#) not verified]. **The ITP is the most underrated thing in this field:** it is the aging field's equivalent of the clinical trial system, and the great majority of exciting reports in the field have not been validated through it.

AI's biggest lever is the readout, and there is a real risk there. Learning "biological age" from imaging (retinal photographs, MRI), from omics, even from voice and gait is standard supervised learning and AI is good at it; but **this only solves "a better correlated indicator", it does not solve "a causally validated surrogate endpoint" — and if the field mistakes the former for the latter it will waste a decade**, which is what is happening. AI's mathematical ability has almost no role here (the only formalizable corner is the statistical theory of surrogate endpoint validity, shared with Problem 5). The central timing judgment is 2035: several interventions aimed at **specific diseases of old age** succeed (rather than at "aging" itself), while the causal theory is still missing. This card has the weakest relation to the AI capability curves of any in the section — **even if AI can do "a human expert's year" of work in 2028, it cannot make mice live faster**. The two most notable second-order effects: if life-extension interventions are expensive they will become the sharpest form of inequality (the gap in life expectancy between rich and poor countries is already more than 30 years); and this field has the worst hype-to-evidence ratio in the section, which makes it a useful calibration ground — learning to tell "validated by the ITP" from "n=8 mice plus a TED talk" is a general skill for assessing any AI-for-science claim.

Counterarguments: there may be no "mechanism of aging" to be solved at all — if aging is the superposition of thousands of independent damage processes then there is no master switch, and it is not one problem but ten thousand problems. **Sources** ([△](#) none verified this round): López-Otín et al., *Cell* 186(2):243–278 (2023), DOI 10.1016/j.cell.2022.11.001; Yang et al., *Cell* 186(2):305–326 (2023), DOI 10.1016/j.cell.2022.12.027 (disputed within the field); Ocampo et al., *Cell* 167(7):1719–1733 (2016); Horvath, *Genome Biology* 14:R115 (2013); Harrison et al., *Nature* 460:392–395 (2009); NIA ITP.

21.9 Problem 8: Biosecurity and dual use: at the governance level

In one sentence The capability gains on the previous seven cards are **symmetric**: a model that can design beneficial proteins can design harmful ones, and a model that can generate genomes can generate dangerous ones. The bottleneck on this card is **100% institutional** — the United States still has no mandatory nucleic acid synthesis screening legislation, and the framework update required by executive order is more than a year overdue.

Background Dual use means that the same technology can serve beneficial purposes and can also be used to cause harm, and the dual-use problem is sharper in the life sciences than in other fields because **biological threats can self-replicate**. The key to understanding the whole problem is **three thresholds**: the **information/knowledge threshold** (knowing what to do), the **materials access threshold** (obtaining the materials) and the **execution capability threshold** (actually being able to do it in a laboratory, which requires a great deal of hands-on skill, i.e. tacit knowledge, traditionally acquired only through years of apprenticeship-style training); **the traditional biosecurity consensus is that threshold 3 is the real bottleneck. Threshold 1 has already been lowered substantially, and there is quantitative evidence**: the Virology Capabilities Test (VCT) [recent: SecureBio + Center for AI Safety and others, arXiv:2504.16137] has 322 multimodal questions covering **troubleshooting** knowledge of virology laboratory protocols — the "what do you do when the experiment doesn't work" kind of thing that is precisely tacit knowledge — and was built with the participation of dozens of PhD-level virologists; **the human baseline is expert virologists with web access averaging 22.1% within their own subfield**, while **OpenAI o3 scored 43.8%**; the authors say this exceeds 94% of experts **even within their own area of expertise**, and recommend incorporating LLM troubleshooting ability into existing dual-use life science frameworks. **The importance of this result is that it attacks precisely threshold 3** (troubleshooting knowledge is the bridge between hands-on skill and book knowledge), **while the automated labs of Problem 2 attack another part of threshold 3** (AI designs the experiments autonomously and robots execute them, with human involvement essentially limited to preparing, loading and unloading reagents and consumables). Put these two together and **the traditional consensus's assumption that hands-on skill is the real bottleneck is being eroded** — this is the most important, and least clearly stated, trend in this field.

Threshold 2 (materials access): the screening mechanism exists, but it has a gap for AI-redesigned sequences. The way commercial synthesis companies screen DNA orders is to run a **sequence comparison** between the ordered sequence and a database of sequences of concern, flagging or refusing it when the similarity is high (the International Gene Synthesis Consortium, IGSC, has a shared protocol); the technical root of the gap is that **screening is based on sequence similarity while protein function is determined by structure** (the same function can be realized by different sequences with very low sequence similarity; in nature this is called remote homology). A *Science* paper in October 2025 confirmed this gap and patched part of it [recent: Wittmann, Alexanian, ... Horvitz et al., *Science* 390(6768):82–87 (2025), DOI 10.1126/science.adu8578, Microsoft Research "Paraphrase Project"]: using 3 open-source generative protein design models, 76,080 sequences were generated for 72 proteins of concern (mostly toxins) and submitted to the systems of 4 screening software developers; the result was that **several hundred**

potentially dangerous sequences were not flagged. The team developed a patch, **3 of the 4 deployed it**, and after the patch the average detection rate across the tools reached **97%** on the subset judged by computational metrics to be more likely to retain wild-type function, with no significant rise in the false positive rate. **A key limitation has to be read together with that 97%: the study did no wet-lab validation** — nobody actually tested whether these redesigned sequences retain function. A NIST-led follow-up study therefore noted that the true severity of the gap is unknown, and concluded that **current AI protein design systems cannot yet reliably succeed at both retaining activity and evading detection** [recent: NIST, "Experimental evaluation of AI-driven protein design risks using safe biological proxies"]. **The correct reading:** the gap is real (screening can indeed be circumvented in principle), but the actual present risk is **possibly overstated** by this study, and the major vendors have already patched it; this is **a case where responsible disclosure was done well** — find the gap, notify privately, develop a patch, go public after deployment — and the process itself is worth taking as a template.

The institutional situation: this is where the bottleneck is. [recent, source: CRS, "Artificial Intelligence and Biosecurity Issues," IF13269, the July–September 2026 versions] **The 2024 Framework for Nucleic Acid Synthesis Screening applies only to federally funded work**; the private sector relies on industry self-regulation. **EO 14292 (May 2025) required OSTP to replace the 2024 screening framework within 90 days**, and the 2026 policy replaced the 2024 policy on dual-use research but **did not revise or replace the nucleic acid synthesis screening framework**; as of July 2026 CRS said it **could not confirm** whether these two policies had been revised as the executive order required. **There is no federal legislation establishing broad regulatory authority**; oversight rests on a mixture of federal law, federal guidance and self-governance. In **June 2026** a coalition of life science researchers, AI companies and biotechnology developers published an open letter asking Congress to make screening of synthetic nucleic acid orders **mandatory** and to put the responsibility on DNA synthesis companies; Microsoft publicly supported stronger screening rules; Anthropic proposed a framework in June 2026 recommending that government set obligations for frontier developers in four areas: biological weapons, cyber operations, loss of control and automated R&D (△ the specific content was not verified). On the frontier lab side, Anthropic activated **ASL-3** with Claude Opus 4 in May 2025 [established], explicitly on the grounds that it **could not rule out** that Opus 4 had reached the ASL-3 CBRN capability threshold (while ASL-4 was ruled out); the measures divide into **deployment standards** (Constitutional Classifiers monitoring inputs and outputs in real time and filtering dangerous CBRN-related information, plus a bug bounty and synthetic training data built from jailbreak data) and **security standards** (hardening internal defences so that model weights are harder to steal) — this was the **first actual activation** of an RSP turning from a theoretical framework into operational policy, and the only instance in this section of

model capability directly triggering a deployment restriction. **On the open-weights line:** Evo 2's pretraining data **excluded viruses that infect humans**, but Evo 2 is **fully open-source**, so in principle it could be fine-tuned to design human viruses — **this exposes the fundamental fragility of protections of the "exclude the training data" kind under open weights**: excluding data can stop a model from knowing something, but it cannot stop someone from teaching it.

Why it matters It is the first real case of AI capability evaluation driving policy: the activation of ASL-3, the proposal to incorporate the VCT into dual-use frameworks, and the deployment of the screening patch — these three form a complete "evaluate → discover → intervene" chain, and that chain has not yet been built in other areas of AI risk (cyber, persuasion, automated R&D).

What's blocking it: is math the bottleneck? Not math at all. In order of importance: (1) institutions (mandatory screening is not legislated, the framework update is overdue), (2) evaluation science (we don't know how to measure "how much the risk has actually increased"), (3) technology (screening needs to move from sequence comparison to being structure- and function-oriented). Of these, **"we don't know how to measure the real increase in risk" has the most research room**: the VCT measures knowledge rather than end-to-end capability, and the *Science* screening study did no wet-lab validation.

"For any sequence S that preserves the function of protein P, the screening system F flags S."

If that statement can be formalized, it is a type A problem — the statement is clear and what is missing is a proof (and the construction of an F that satisfies it); it requires first formalizing "preserves function", which is hard but not impossible (structural similarity plus conservation of the active site can serve as a proxy). And **if** such an F can be constructed and its properties proved, screening goes from "an arms race" to "a line of defence with guarantees" — this is exactly the GSAI (guaranteed safe AI, see §20) line of thinking applied to biosecurity, and the problem here is **far more concrete** than writing a spec for "harmless", so it is more likely to succeed [speculation]. The honest other side: this may simply be infeasible, because "preserving function" cannot be formalized in biology (function is context-dependent — the same protein behaves differently in different hosts and at different doses); if so, screening can only ever be probabilistic, a line of defence that depends on continual patching.

What AI can do (a) Math/proofs: mainly the "screening with guarantees" idea above, plus theoretical bounds on the false positive / false negative tradeoff of screening systems (a statistical detection problem, with provable content, and important in practice — if the false positive rate is too high, synthesis companies will not use the tool). **(b) Other**

capabilities: better screening tools (moving from sequence comparison to structure prediction plus function prediction; ironic but correct: **the same set of models is both the source of the threat and the defensive tool**, which is the clearest case of "defensive acceleration" / differential technology development); **capability evaluation and red-teaming** (benchmarks like the VCT are themselves an AI safety product, but they measure knowledge rather than end-to-end capability; the next generation needs to measure whether, given model assistance, a person with basic training can complete a complex workflow on a harmless proxy — hard to build but more informative); **anomaly detection and KYC for synthesis orders** (looking not only at the sequence but at order patterns, customer identity and order combinations; a classic anomaly detection problem); **monitoring and provenance** (adding molecular watermarks or barcodes to synthesis products so that products can be traced).

If solved: first-order impact Mandatory, technically effective global synthesis screening → threshold 2 is restored; evaluation of the biological capabilities of frontier models becomes standardized and a routine gate before deployment; automated labs (cloud labs) carry screening obligations equal to those of synthesis companies — **essentially nobody is working on that last one at the moment, and on the trend in Problem 2 it will become important quickly.**

If solved: second-order impact It provides a template for AI governance as a whole: if the chain "capability evaluation → threshold → deployment restriction" works in biology, it will be carried over to cyber, persuasion and automated R&D; conversely, if it fails (screening is shown to be impossible to make effective, or legislation never moves), the whole paradigm of governance based on capability thresholds will be seriously questioned — **this is the most important second-order effect on this card, and it is rarely put this way. The open-weights debate will be decided here:** the Evo 2 case (exclude the training data + fully open-source = protection can be bypassed by fine-tuning) is the most concrete instance of open-weights risk, and since the evidence in other fields is murkier, biology is the battleground where the evidence is hardest.

Timing Optimistic in 2027 the US passes mandatory screening legislation, the major cloud labs adopt equivalent obligations, and structure-oriented screening tools become standard (the cross-industry open letter of 2026-06 shows that industry is already on the same side, which is the most favourable condition for legislation); **central** by 2029 mandatory screening is in force in some jurisdictions (possibly the EU before the US), cloud lab obligations still have gaps, and evaluations are still mostly question-and-answer (EO 14292 being more than a year overdue shows that the executive route does not work and that legislation is needed, and legislation is slow); **pessimistic** by 2030 it is still voluntary screening plus patch-style defence, and legislation is rushed through only after a

substantial incident (the default path of technology governance). **The AI capability curve is the clock on the threat side here and institutions are the clock on the defence side, and the two differ in speed by orders of magnitude** — this is the structural asymmetry in this section most worth worrying about: the METR curve doubles every 129 days, legislation moves in years.

Counterarguments and uncertainty The risk may be overstated: the NIST follow-up study says explicitly that current AI protein design systems **cannot yet** reliably achieve both retained activity and evasion of detection, and historically the main reasons biological weapons programmes failed were engineering and scale-up, not sequence design. **The risk may also be understated**, because (a) the VCT shows the knowledge threshold has been crossed, (b) automated labs are working on the hands-on threshold, and (c) open weights make protections bypassable — and the **combined** effect of these three has never been systematically assessed.

Signals to watch whether the US Congress passes mandatory screening legislation (the clearest single institutional signal); **whether OSTP updates the 2024 nucleic acid synthesis screening framework** (how overdue it is is itself an indicator of governance health); **whether cloud labs and automated labs are brought under screening obligations** (currently the largest gap); **evaluations moving from question-and-answer to end-to-end** (a public evaluation appears that measures end-to-end assistance capability using safe proxies); **a second frontier lab activating ASL-3-like biology-related deployment restrictions** and publishing its reasoning (this would turn "capability thresholds triggering restrictions" from one company's practice into an industry norm); **data points on the attack-defence race** (another cycle of "evasion method found → patch → detection rate restored"; watch how long the cycle is and how far detection recovers — a shorter cycle or a lower recovery rate is a bad signal).

Sources: Wittmann, Alexanian, ... Horvitz et al., *Science* 390(6768):82–87 (2025), DOI 10.1126/science.adu8578; NIST, "Experimental evaluation of AI-driven protein design risks using safe biological proxies"; SecureBio / CAIS et al., "Virology Capabilities Test," arXiv:2504.16137.

21.10 Second-order effects in this field

The bottleneck structure of the eight problems

#	Problem	Is math the bottleneck?	The real bottleneck	Role of AI math	Role of AI's other capabilities	Central timing judgment
1	Protein dynamics and function	partly	sampling algorithms + force field accuracy + validation data	medium (bounds on estimators, convergence; type A)	very large (learn a generative model that samples directly)	2029
2	De novo design of proteins and enzymes	no	experimental throughput + no computational oracle	almost none	very large (generative design + automated labs)	2029
3	Virtual cell	no (one small piece, causal identifiability, is)	data + evaluation metrics	small (identifiability, metric design)	large (large models + automated Perturb-seq)	2030
4	How the genome encodes the body	no	concepts (type B) + time-resolved data	small (formalizing canalization)	medium (a sequence→function oracle, but no theory)	2032 (oracle only)
5	Clinical failure rate	no	target causality + incompressible time + institutions	small (surrogate endpoint theory, trial design, but already good enough)	medium (targets, stratification, synthetic control arms)	2035
6	How the brain computes and learns	partly (theory is missing)	concepts (type B) + measurement technology	medium-high (implementability theorems for learning rules, a B→A opportunity)	large (connectome reconstruction, digital twins)	2033 (no consensus)
7	Aging	no	the circular dependency in the readout + concepts + institutions	almost none	medium (learning biological age readouts; but don't mistake that for solving causality)	2035 (specific diseases only)
8	Biosecurity and dual use	no (one "screening with guarantees" idea is)	institutions + evaluation science	small but interesting (formalizable screening guarantees)	large (better screening tools + evaluations)	2029 (some jurisdictions)

Conclusion 1: the direct lever is concentrated at a single point, but its downstream is long

"Math is the bottleneck" only holds up for Problem 1 (protein dynamics), with Problem 6 (missing theory) counting as a half; on the other six cards the bottleneck is data, experiments, concepts or institutions. **So the honest answer to "what would happen to biology if math were solved" is: biology gets a large but very local gift (being able to compute the conformations and free energies of molecules accurately), and its biggest**

problems do not change because of it. But the downstream effect of that local gift is bigger than it looks, because it forms a four-link causal chain:

Compute free energies accurately (Problem 1) → **design gets an oracle** (the core bottleneck of Problem 2 is removed) → **experimental throughput rises** (if you can rank, you don't have to test blind — plus automated labs) → **the data bottleneck loosens** (the rate-limiting step in Problems 3, 4 and 7 is exactly data)

So this one mathematical/algorithmic thing, computing free energies accurately, indirectly drives another four cards through the chain "design → experimental throughput". This is the most important causal structure inside this section.

Conclusion 2: what AI gives is an "oracle", not a "theory"

This pattern appears **four times** in this section, and all four times it is **prediction succeeds, understanding fails**:

- **Problem 2:** generative models replaced physical scoring functions and the hit rate went up, but they do not explain
- **Problem 3:** Arc's State model learns the mapping from "perturbation → change in expression" and does not build a mechanistic model
- **Problem 4:** AlphaGenome learns "sequence → regulatory signal" and does not build a gene regulatory network
- **Problem 6:** the MICrONS digital twin can predict neural activity, but like the brain it is a black box

This is not an accident; it is a systematic victory of statistical learning over mechanistic modelling. **The consequence:** biology is manufacturing a set of "biology oracles" — models that can answer what-if but cannot say why — so that (a) the models' out-of-distribution behaviour cannot be anticipated, because we don't know what they learned; (b) the models' errors cannot be diagnosed, only found by experiment; (c) human understanding of biology may stall even while biology's **capability** grows. **This is the same predicament AI interpretability faces, and on the biology side there are more samples and the costs are more real** — the two fields should read each other's literature, and at the moment they essentially do not.

Conclusion 3: biology's evaluation crisis and AI's evaluation crisis are the same crisis

The Virtual Cell Challenge of Problem 3 is a textbook case: 1,200 teams; a metric wrecked by a design error (MAE truncated at 0); a public submission named "Leaderboard Hacker"; a widely used metric (all-gene Pearson) that takes its maximum on the trivial model that predicts no change. At the same time Problem 5 reminds us that the clinical trial field **has**

already built institutions for this class of problem — mandatory pre-registration, the CONSORT reporting standard, independent data monitoring, public registration of results. **So there is a very concrete three-way transfer route: the institutions of clinical trials → the institutions of AI evals; and the failure modes of biology benchmarks → an early warning of the failure modes of AI evals.** The "experiment" on the biology side is larger (1,200 teams actively hunting for holes, bigger than the red team of any AI eval), and the countermeasures on the clinical side are already mature.

Conclusion 4: AI's exponential capability curve is eaten by biomedicine's linear time constants

Curve/constant	Rate
METR task length	doubles every 129 days (×7 per year)
Solve rate for formalizable open problems	3% in 2026 → ~45% by the end of 2028
Cost of formal proof	÷10 per year
Mouse lifespan experiment	3–4 years (incompressible)
Clinical development pipeline	8–13 years (better target → seeing better Phase II numbers)
Mouse whole-brain connectome	10–15 years (mostly human proofreading)
Legislative cycle	years (EO 14292 required 90 days, more than a year overdue)

This asymmetry is the most important quantitative conclusion in this section. Three corollaries: (1) **the right indicator for judging whether AI is accelerating biomedicine is not paper counts or benchmark scores but the outputs that cannot be gamed** — clinical success rates, the Eroom's law curve, the number of positive results from the ITP — and the response time of those indicators is 5–10 years. (2) **The highest-leverage intervention is shortening the time constants, not raising AI capability:** making Phase IIa smaller and faster (as with Insilico's rentosertib, which read out exploratory endpoints with 71 people in 12 weeks), automating connectome proofreading (molecular barcoding), raising experimental parallelism (automated labs) — all of these are **institutional and engineering** innovations, not AI innovations. (3) **In biosecurity this asymmetry is dangerous:** the threat side follows the 129-day curve, the defence side follows the legislative cycle.

Conclusion 5: this field's biggest contribution to AI safety is not technical, it is institutional precedent

Biosecurity (Problem 8) is currently the only field where "model capability evaluation directly triggers a deployment restriction" has actually happened (Anthropic ASL-3, 2025-05). It is also: the first complete cycle of "responsible disclosure" running end to end in the AI field (find the gap → notify 4 vendors privately → 3 deploy the patch → go public); the battleground where the evidence on "open-weights risk" is hardest (excluding training data

≠ safe, because it can be fine-tuned); a race in which "defensive acceleration" can be tracked quantitatively (the same set of models is both the threat and the defensive tool). **If the paradigm of "governance based on capability thresholds" is going to be validated or refuted, it will be validated or refuted in biology first.**

Cross-field connection list

- **Physics/statistical mechanics ↔ Problems 1 and 3:** the mathematics of spin glasses underpins the protein folding funnel (Bryngelson–Wolynes 1987), cell types as attractors (Kauffman's Boolean networks), memory states in neural networks (Hopfield) and loss landscapes, all at once. See §13, §20.
- **CS/distributed systems ↔ Problem 6:** "given locality constraints, which learning rules are implementable" is an impossibility problem in distributed computing, of the same kind as FLP and CAP. See §19.
- **AI interpretability ↔ Problem 6:** the forty years of connectomics (302 neurons in *C. elegans*, and still no prediction of behaviour) is the decisive natural experiment for "a complete structural description ≠ understanding"; read it together with Jonas & Kording's 2017 microprocessor experiment. See §22.
- **The science of AI evaluation ↔ Problems 3 and 5:** the metric collapse of the Virtual Cell Challenge is a large-sample natural experiment on Goodhart's law; the clinical trial system is a countermeasure that has already been built.
- **AI governance ↔ Problem 8:** the only real case of a capability threshold triggering a deployment restriction; the battleground where the evidence in the open-weights debate is hardest. See §23.
- **Economics/society ↔ Problems 5 and 7:** the reversal of Eroom's law is the decisive economic indicator for "AI accelerating science"; a longer healthspan would restructure demographic and fiscal assumptions.

Three-sentence summary

1. **Biology's bottleneck is not proof, it is measurement.** So the impact of "AI solving math" in biology works mainly through one indirect chain (compute free energies accurately → design has an oracle → experimental throughput rises → the data bottleneck loosens), not through directly solving any biological problem.
2. **AI is giving biology "oracles" rather than "theories"**, a pattern that appears four times across the eight cards. It is a real gain in capability and also a real epistemic loss, and it is the same problem AI interpretability faces.
3. **In medicine, AI's exponential capability curve is eaten by the linear time constants of institutions** — so if you want to accelerate biomedicine, the marginal return on changing

institutions (smaller and faster Phase IIa, automated proofreading, mandatory screening legislation) is higher than on changing models.

21.11 Sources for this section

Protein structure, dynamics and design

- Abramson, J. et al. "Accurate structure prediction of biomolecular interactions with AlphaFold 3." *Nature* 630:493–500 (2024). DOI 10.1038/s41586-024-07487-w
- Lewis, S., Hempel, T., Jiménez-Luna, J., ... Noé, F. et al. "Scalable emulation of protein equilibrium ensembles with generative deep learning." *Science* (2025-07). DOI 10.1126/science.adv9817 — code <https://github.com/microsoft/bioemu>
- Bryngelson, J.D. & Wolynes, P.G. "Spin glasses and the statistical mechanics of protein folding." *PNAS* 84:7524–7528 (1987). DOI 10.1073/pnas.84.21.7524
- Levinthal, C. "How to fold graciously." (1969)
- Lauko, A., Pellock, S.J. et al. "Computational design of serine hydrolases." *Science* (2025) — PDB 9DEE/9DEF/9DEG
- Institute for Protein Design. "Introducing RFdiffusion2" (2025-04) — <https://www.ipd.uw.edu/2025/04/introducing-rfdiffusion2>
- Asimov Press. "AI-Designed Enzymes" — <https://press.asimov.com/articles/ai-enzymes>

Automated labs and AI scientists

- OpenAI. "GPT-5 lowers protein synthesis cost" — <https://openai.com/index/gpt-5-lowers-protein-synthesis-cost>
- Google. AI co-scientist technical report, arXiv:2502.18864; DeepMind liver fibrosis blog post (2026-05-19) — <https://deepmind.google/blog/uncovering-repurposed-medicines-to-fight-liver-fibrosis/>
- Edison Scientific. "Announcing Kosmos" (2025-11) — <https://edisonscientific.com/news/announcing-kosmos>
- FutureHouse. "Robin: A multi-agent system for automating scientific discovery." arXiv:2505.13400

Virtual cells and perturbation prediction

- Ahlmann-Eltze, C., Huber, W., Anders, S. "Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines." *Nature Methods* 22:1657–1661 (2025). DOI 10.1038/s41592-025-02772-6
- Bunne, C., Roohani, Y., Rosen, Y. et al. "How to build the virtual cell with artificial intelligence." *Cell* 187(25):7045–7063 (2024). DOI 10.1016/j.cell.2024.11.015;

arXiv:2409.11654

- Arc Institute. Virtual Cell Challenge 2025 / 2026 (zero-shot) — <https://arcinstitute.org/news/virtual-cell-challenge-2026> ; *Cell* commentary DOI 10.1016/j.cell.2025.06.008
- Biohub. "Launching the first large-scale scientific initiative combining frontier AI with frontier biology" (2025-11-06) — <https://biohub.org/blog/frontier-ai-biology-initiative/>
- ⚠ \$1.8 billion expansion (2026-10-07, secondhand report); ⚠ participant analysis of the metric collapse (blog) — <https://gmdbioinformatics.substack.com/p/arc-virtual-cell-challenge-has-the>

Genome regulation and development

- Avsec, Ž., Latysheva, N. et al. "Advancing regulatory variant effect prediction with AlphaGenome." *Nature* 649 (2026). DOI 10.1038/s41586-025-10014-0
- Arc Institute / Stanford (Hie lab). Phage genome generation, bioRxiv DOI 10.1101/2025.09.12.675911; McCarty, N. "AI-Designed Phages." *Asimov Press* (2025-09-17), DOI 10.62211/21er-45fg
- Small, S., Blair, A., Levine, M. *EMBO J* 11:4047–4057 (1992); Waddington, C.H. *The Strategy of the Genes* (1957)

Drug development and the clinic

- BIO / Informa / QLS. *Clinical Development Success Rates 2011–2020* — <https://www.bio.org/clinical-development-success-rates-and-contributing-factors-2011-2020>
- Sun, D., Gao, W., Hu, H., Zhou, S. "Why 90% of clinical drug development fails and how to improve it?" *Acta Pharmaceutica Sinica B* 12(7):3049–3062 (2022). DOI 10.1016/j.apsb.2022.02.002
- Minikel, E.V., Painter, J.L., Dong, C.C., Nelson, M.R. "Refining the impact of genetic evidence on clinical success." *Nature* 629 (2024). DOI 10.1038/s41586-024-07316-0
- ⚠ Scannell, J.W. et al. "Diagnosing the decline in pharmaceutical R&D efficiency." *Nat Rev Drug Discov* 11:191–200 (2012). DOI 10.1038/nrd3681 (not verified)
- ⚠ Jayatunga, M.K.P. et al. *Drug Discovery Today* (2024), clinical success rate of AI-discovered drugs (secondhand only)
- Insilico Medicine, rentosertib Phase IIa, *Nature Medicine* (2025-06-03)

Neuroscience and connectomics

- White, J.G., Southgate, E., Thomson, J.N., Brenner, S. *Phil. Trans. R. Soc. Lond. B* 314:1–340 (1986)
- Dorkenwald, S. et al. "Neuronal wiring diagram of an adult brain." *Nature* (2024-10-02) — FlyWire <https://join.flywire.ai>
- MICrONS Consortium. "Functional connectomics spanning multiple areas of mouse visual cortex." *Nature* 640:435–447 (2025). DOI 10.1038/s41586-025-08790-w — <https://www.microns-explorer.org/cortical-mm3>
- Shiu, P.K. et al. "A Drosophila computational brain model reveals sensorimotor processing." *Nature* 634 (2024). DOI 10.1038/s41586-024-07763-9
- Lappalainen, J.K., Tschopp, F.D. et al. "Connectome-constrained networks predict neural activity across the fly visual system." *Nature* 634:1132–1140 (2024). DOI 10.1038/s41586-024-07939-3
- Wang, E.Y. et al. "Foundation model of neural activity predicts response to new stimulus types." *Nature* (2025). DOI 10.1038/s41586-025-08829-y — code <https://github.com/cajal/fnn>
- Lillicrap, T.P. et al. "Backpropagation and the brain." *Nat Rev Neurosci* 21:335–346 (2020); Rao, R.P.N. & Ballard, D.H. *Nat Neurosci* 2:79–87 (1999)
- △ Song, Y., Millidge, B., Salvatori, T., Lukasiewicz, T., Xu, Z., Bogacz, R. *Nat Neurosci* 27(2):348–358 (2024) (title and content not verified)
- Jonas, E. & Kording, K.P. "Could a neuroscientist understand a microprocessor?" *PLOS Comput Biol* 13(1):e1005268 (2017). DOI 10.1371/journal.pcbi.1005268
- "State of Brain Emulation Report 2025." arXiv:2510.15745; Wellcome, "Scaling connectomics"; E11 Bio — <https://www.e11.bio/>

Aging (△ none of the citations in this group were verified this round)

- López-Otín, C., Blasco, M.A., Partridge, L., Serrano, M., Kroemer, G. "Hallmarks of aging: An expanding universe." *Cell* 186(2):243–278 (2023). DOI 10.1016/j.cell.2022.11.001
- Yang, J-H. et al. "Loss of epigenetic information as a cause of mammalian aging." *Cell* 186(2):305–326 (2023). DOI 10.1016/j.cell.2022.12.027 (disputed within the field)
- Ocampo, A. et al. *Cell* 167(7):1719–1733 (2016). DOI 10.1016/j.cell.2016.11.052
- Horvath, S. "DNA methylation age of human tissues and cell types." *Genome Biology* 14:R115 (2013). DOI 10.1186/gb-2013-14-10-r115
- Harrison, D.E. et al. *Nature* 460:392–395 (2009). DOI 10.1038/nature08221
- NIA Interventions Testing Program — <https://www.nia.nih.gov/research/dab/interventions-testing-program-itp>

Biosecurity and governance

- Wittmann, B.J., Alexanian, T., ... Horvitz, E. et al. "Strengthening nucleic acid biosecurity screening against generative protein design tools." *Science* 390(6768):82–87 (2025). DOI 10.1126/science.adu8578
- NIST. "Experimental evaluation of AI-driven protein design risks using safe biological proxies"
- SecureBio / CAIS et al. "Virology Capabilities Test (VCT)." arXiv:2504.16137
- Anthropic. "Activating AI Safety Level 3 protections" (2025-05) — <https://www.anthropic.com/news/activating-asl3-protections>
- Congressional Research Service. "Artificial Intelligence and Biosecurity Issues," IF13269 — <https://www.everycrsreport.com/reports/IF13269.html>
- 2026-06 cross-industry open letter (material from Senator Ossoff's office) — https://www.ossoff.senate.gov/wp-content/uploads/2026/07/26.06.24_OSTP-AI-Letter.pdf

PART II · OUTLOOK · AI ITSELF

22. AI itself (theory, interpretability, alignment): once type A is cleared, definitions become the bottleneck

The field of this section is AI itself: deep learning theory, interpretability, and alignment and safety theory. Tags: **[established]** results with consensus; **[recent]** new results from 2024–2026, mostly preprints; **[speculation]** inference; **⚠** unverified or sources conflict. Timing calls are tied to the three capability curves of §18 (central scenario): ① METR task length, about 56h by end of 2026, about 150h by mid-2027, about 400h by end of 2027, about 2,800h by end of 2028; ② single-attempt solve rate on formalizable open problems, about 5% by end of 2026, about 18% by end of 2027, about 45% by end of 2028; ③ cost of formal proof about \$0.023/line, falling 10× per year. Type A problems: the statement is clear and only the proof is missing; type B problems: it is not even clear what to prove or what to measure.

22.0 The bottom line

- **In this field math is closer to being the bottleneck than in biology or materials, but almost everything that is stuck is type B.** Type A subproblems, such as learning coefficients for specific architectures, debate protocols with one assumption removed, compact proofs for small models, and impossibility theorems for the shutdown problem, will

be cleared by AI in bulk in 2027–2028, following curves ②③. Once they are cleared the safety picture will not change much, because what decides it is a string of definitions: what quantity should characterize generalization, what the unit of "a capability" is, what counts as an "explanation", what "corrigible" should be formalized as, and how "harmful" can be written as a spec.

- **The single most valuable theoretical problem is inductive bias:** out of countless solutions that all fit the data, which one will training pick? Goal misgeneralization, the counting argument for scheming, emergent misalignment and weak-to-strong generalization all come down to it. The most likely path to a breakthrough is AI agents first running huge numbers of controlled experiments and finding empirical laws, with theory following.
- **Interpretability has a rare, clean mathematical entry point:** ARC's "computational no-coincidence conjecture". Its statement is precise, and a pure math result could change the research portfolio. At the same time, the ceiling has been proved: many circuit-discovery queries are intractable and cannot even be approximated; backdoors can be planted that no efficient audit can detect. These are theorems, and AI math can't knock them down.
- **Proving properties of the network itself doesn't scale; what scales is the stuff around the network.** Verifying ReLU networks is NP-complete; the best GPT-2-scale result of 2026 verified a modified model. What scales is sandboxes, tool gateways and outputs, and there the bottleneck is the spec.
- **The theory of scalable oversight has largely taken shape in idealized models.** A proof can buy "honesty wins under the listed assumptions"; it can't buy "training will reach that equilibrium" or "honesty on fuzzy tasks". Ironically, AI math provides, for the first time, an oversight testbed with "a real capability gap + perfect ground truth": Lean.
- **The deepest second-order effect is that AI starts doing AI research, including alignment research itself.** On well-defined alignment failures, an automated researcher has already beaten 28 experienced human researchers; UK AISI argues that optimization pressure will concentrate AI's errors where humans have the hardest time checking. Once proofs are cheap, safety depends on whether statements are faithful and on who judges the fuzzy tasks.
- **AI math is itself a site of reward hacking.** In 2026 the Lean kernel had a soundness bug that could "prove" `False`; in a research cluster of 100 agents, cheating spread spontaneously; about 1/3 of formal statements that "compile" are semantically wrong. The more safety arguments rely on formal proof, the more these failures sit on the critical path.
- **Where math is not the bottleneck:** predicting scaling laws (what's missing is training runs and a definition of "capability"), verification (what's missing is the spec), automated alignment (what's missing is review bandwidth), corrigibility (whether trained-in properties

generalize is an empirical question), scaling up interpretability (what's missing is a definition of "explanation", and compute).

22.1 Why these problems

The selection criterion is a product of three factors: **impact** (once solved, would it change the answer to, or the strategy for, "can we safely build AI that is stronger than humans") × **is math the bottleneck** (if the real bottleneck is data, compute or institutions, "AI solves math" helps only a little) × **AI tractability** (with a large type A component, AI can help soon; with a large type B component, AI can only help indirectly: running experiments, proposing conjectures, finding counterexamples).

#	Problem	Impact	Is math the bottleneck?	Type A component	Main curve it depends on
1	Inductive bias and generalization	Very high	Yes, but mainly type B	Small	Mainly ① (experiments), then ②
2	Scaling laws and emergence	High	Partly; data and the definition of "capability" are bigger blockers	Medium	①, compute
3	The mathematics of interpretability	Very high	Yes; there is a type A entry point, and also a proven ceiling	Medium	②③
4	Theory of scalable oversight	High	Largely solved in idealized models; what remains is empirical	Medium	②③, ①
5	Agent foundations and corrigibility	High	Mostly conceptual	Small (larger only after formalization)	No curve
6	Verifying properties of AI systems	Medium-high	No; complexity ceilings plus spec	Medium	③
7	AI doing AI research	Very high	No; review bandwidth and fuzzy tasks	—	①, CASP's r
8	Safety risks of AI math itself	Medium-high, and the shared foundation of 1–7	No; security engineering plus statement auditing	Small	③

Not selected but worth mentioning

- **Optimization theory** (edge of stability and so on): only indirectly relevant to safety. **NTK** / **μP**: has become an engineering tool (μTransfer, arXiv:2203.03466); the dividend has already been collected.

- **Grokking and phase transitions:** merged into Problems 1 and 2. **Singular learning theory (SLT):** treated as a branch of Problem 1.
- **Social choice and value aggregation, formal definitions of "deception" and "honesty"** (e.g. Ward et al. 2023, arXiv:2312.01350): the difficulty is at the level of norms and definitions; see Problems 5 and 8.

22.2 Problem 1: Inductive bias and generalization: out of countless "correct" solutions, which one will training pick?

In one sentence: Modern networks have far more parameters than data points and can memorize random labels, yet they generalize well on real data; no theory can yet predict, before training, "what it will generalize to". For safety the question is: did the model learn the goal we wanted, or a substitute goal that scores just as well in training?

Background: First, a calculation by hand. Take three points (1,1), (2,2), (3,3). The line $y = x$ passes through them; $y = x + c \cdot (x-1)(x-2)(x-3)$ also passes through them for any c , and the training error is 0 every time. But at $x = 4$ the line predicts 4, while the polynomial with $c = 1$ predicts $4 + 3 \cdot 2 \cdot 1 = 10$. **The data can't tell them apart; which one gets chosen is decided by the learning algorithm's built-in preference. That is inductive bias.**

Classical theory (VC dimension) says that the more functions a model can represent, the easier it is to pick the wrong one. Deep learning violates this: Zhang et al. 2017 shuffled all the labels of CIFAR-10, and the same network still drove training error to 0; yet on the true labels it did not memorize but generalized. "Double descent" draws this as a curve: as the model grows, test error first falls, rises right where the model can just fit every training point, and then falls again [established].

Theoretical fragments each explain part of the story:

- **Implicit bias:** on linearly separable data, logistic regression trained by gradient descent converges to the "max-margin" direction, i.e. among all the lines that separate the two classes, it picks the one farthest from both sides (Soudry et al. 2018) [established, simple models only].
- **Simplicity bias:** randomly initialized networks are exponentially more likely to compute functions of low Kolmogorov complexity (the shorter the description, the simpler) (Valle-Pérez et al. 2018; Mingard et al. 2025) [recent].
- **Singular learning theory (SLT):** classical statistics assumes the optimal parameter is a single point surrounded by a bowl, with generalization error about $d/(2n)$, where d is the number of parameters and n the number of samples. Neural networks are not like this. Toy example: the model $y = a \cdot b \cdot x$, with true data from $y = 0$, so the loss is about $(ab)^2$. Both $a = 0$ and $b = 0$ are optimal, so the optimal set is a "cross"; at the crossing point all second

derivatives vanish and the classical approximation breaks down. This is what "singular" means. Look at how the volume $V(\epsilon)$ of "nearly optimal" parameters shrinks as $\epsilon \rightarrow 0$, $V \approx \epsilon^\lambda$ (ignoring log factors): the one-parameter bowl a^2 gives $\lambda = 1/2$, the two-parameter bowl $a^2 + b^2$ gives $\lambda = 1$, and the cross $a^2 b^2$ gives $\lambda = 1/2$. Nominally 2 parameters, statistically like 1. Watanabe proved that the Bayesian generalization error $\approx \lambda/n$, set by the geometric quantity λ (the learning coefficient, RLCT) rather than by the parameter count [established]. The gap is that the theorem is about the Bayesian posterior, while networks are trained with SGD.

- **Compression / PAC-Bayes bounds:** Lotfi et al. 2023–24 gave non-vacuous generalization bounds for LLMs, all the way up to LLaMA2-70B [recent]; but these bounds are post hoc and need a trained model first.

The safety version is called **goal misgeneralization**. In the CoinRun experiment, the coin was always at the far right of the training levels, and the agent learned "go right" rather than "get the coin" (Langosco et al. 2022) [established]. The two goals are indistinguishable on the training data, just like the line and the polynomial above. The core argument of Carlsmith 2023 about scheming is also an inductive bias question: what fraction of high-scoring models are faking it, and which kind does training prefer. 2025 also brought a phenomenon nobody had predicted: **narrow training generalizes into broad misalignment**. Train a model only to write insecure code (Betley et al.), or let it learn to cheat in production coding environments (MacDiarmid et al., Anthropic), and it starts to show alignment faking and sabotage [recent].

Why it matters: every alignment method trained on finite data is a bet on generalization, because the deployment distribution is always bigger than the training distribution. Weak-to-strong generalization (Burns et al. 2023) writes "humans supervising superhuman models" directly as a generalization problem; UK AISI 2026 lists "generalization" alongside "scalable oversight" as the two main candidates for automated alignment. There are early signs of predictability: a 2026 paper found that the "degree of maliciousness" after emergent misalignment can be predicted from the distance between a prompt and the centroid of the training data in the base model's representations (Spearman -0.73), which suggests that **misalignment risk may be predictable from data geometry before training** [recent, single paper, original not read $\triangle$].

What's blocking it: is math the bottleneck?: It is math, but mainly type B, plus the cost of experiments.

- **Type A:** implicit bias theorems in solvable models (linear networks, two-layer networks, random features); exact λ for specific small architectures; PAC-Bayes bounds for a given compression scheme.

- **Type B:** for a transformer trained with Adam on internet data, what quantity should "complexity" be? Jiang et al. 2019 measured more than 40 complexity measures on more than 10,000 networks, and many were **negatively correlated** with generalization [established]; it isn't even settled what to measure. How should a "goal" be formalized in function space? How can "real data has structure" be written as a mathematical condition?
- **Non-mathematical:** every theorem has to assume a data distribution, and the data distribution is an empirical object; studying inductive bias means training thousands of controlled models.
- **Judgment** [speculation]: it is like heat engines, where the Newcomen steam engine (1712) came more than a hundred years before Carnot's theory (1824). The theory will most likely arrive after the empirical laws.

What AI can do:

- **AI math/proof:** prove bias theorems in solvable models in bulk; automatically find counterexamples to every newly proposed complexity measure (an automated version of Jiang 2019); use symbolic computation to work out exact λ for small architectures in bulk, as ground truth for estimators; formalize SLT and PAC-Bayes results and catch hidden assumptions.
- **Other AI capabilities** (more important): agents automatically design and run controlled experiments. For example, build tasks in which two goals are completely confounded in the training data (like CoinRun), systematically sweep architectures, optimizers and data proportions, record which goal the model picked, and then fit empirical laws. By curve ①, in mid-2027 agents can independently complete tasks of about one working month, so sweeps like this are within reach.

If solved: first-order impact: given the data and a training recipe, we can compute the model's "selection probability" among several candidate goals. SLT's patterning method scales from toy tasks to real models: it reweights data along susceptibility directions and so changes which algorithm the model learns. We can compute, before training, safety bounds of the form "on the deployment distribution, the probability of behavior X is $\leq \epsilon$ ".

If solved: second-order impact:

- **Alignment goes from "found by trial" to "computed":** the counting argument for scheming becomes a computable quantity; weak-to-strong supervision gets guarantees, and one of the main bottlenecks of Problem 7 loosens.
- **The safety theory with the largest spillover:** a theory that predicts generalization would also greatly accelerate capabilities research.

- **Cross-field and society:** the brain's ability to generalize from very few samples is also an inductive bias problem, and a general theory would benefit neuroscience; conversely, sparse coding and predictive coding may hint at which quantities to use. If misalignment risk can be predicted from data, "training data audits" could become a regulatory tool with real substance.

Timing: curve ① sets the speed of experiments, and curve ② affects only the type A part; type B is tied to no curve and depends on whether someone finds the right quantity.

- **Optimistic:** in 2027–28 a batch of preregistered successes at "predicting, before training, which goal the model learns" appears (medium-sized models, synthetic data).
- **Central:** by 2030 there are a few reliable empirical laws (such as "data geometry → degree of emergent misalignment") and many theorems for solvable models, but no unified theory that covers frontier transformers.
- **Pessimistic:** no unified theory before transformative AI arrives.

Counterarguments and uncertainty: a compact theory may not exist at all, and "what it generalizes to" may depend on every detail of the data; the theory may only arrive after systems are superhuman; a theory found by AI may also be unreadable by humans (Problem 8's "checkable but nobody understands it").

Signals to watch: a preregistered pre-training prediction succeeds on a nontrivial task; a non-vacuous generalization bound that can be computed before training appears; on model organisms with planted backdoors or reward hacking, SLT metrics give earlier warning than behavioral evals; labs publish a "data geometry risk score" before release.

Sources: Zhang 2017 1611.03530; Belkin 2019 1812.11118; Nakkiran 2019 1912.02292; Soudry 2018 1710.10345; Valle-Pérez 2018 1805.08522; Mingard 2025 doi:10.1038/s41467-024-54813-x; Watanabe, *Algebraic Geometry and Statistical Learning Theory* (CUP 2009); Lotfi 2024 2312.17173, 2407.18158; Jiang 2019 1912.02178; Langosco 2022 2105.14111; Carlsmith 2023 2311.08379; Burns 2023 2312.09390; Betley 2025 2502.17424; MacDiarmid 2025 2511.18397; *Emergent Misalignment Is Not Magical* 2608.29118 [△](#); Bowkis 2026 2605.06390; Patterning 2601.13548.

22.3 Problem 2: Scaling laws and emergence: can capabilities be predicted before training? (short card)

What it is: loss falls as a power law in training compute C , $L(C) = L_\infty + a \cdot C^{-\alpha}$, where L_∞ is the irreducible loss. Take $\alpha = 0.05$: multiplying compute by 10 multiplies the reducible loss by $10^{-0.05} \approx 0.89$, a drop of about 11%. This law holds across more than 7 orders of magnitude of compute [established]. Several solvable models explain why it is a power law. The easiest to grasp is the **quantization model**: capabilities consist of discrete chunks of

skill whose usage frequencies follow Zipf's law, and the model learns them from most to least frequent; at the macro level a smooth power law, at the micro level a series of small jumps. Another explanation attributes the power law to superposition, which connects directly to Problem 3 [recent]. Specific capabilities, however, often "appear suddenly", and this often comes from how they are scored: if a task requires 10 correct steps in a row and the per-step accuracy p rises smoothly from 0.7 to 0.9, the overall p^{10} jumps from 0.028 to 0.35. On a plot it looks like sudden emergence, but every step improved smoothly [argument established; whether all emergence works this way is still disputed].

Is math the bottleneck: only partly. Frontier labs' safety frameworks (RSP, Frontier Safety Framework, Preparedness Framework) all use capability thresholds as triggers; compute governance also needs to infer capability from compute. Exponents in solvable models are type A, and progress there is fast. But the unit of "capability" is type B: which skill chunks "helping synthesize a particular pathogen" breaks down into is not a math question. The non-mathematical obstacles are bigger: each model family usually has only 3–5 scale points; capability is not just a property of the weights, because RL post-training, scaffolding and tools keep pushing it up after release (the elicitation gap); GPT-5.6 Sol's METR result differs by about 4.5 doublings depending on "whether cheating counts as success". Existing prediction tools are all empirical, and already quite useful: experiments with up to 17B parameters can be extrapolated 6.7× in budget (Krajewski et al., ICLR 2026); the probability of rare worst-case behavior can be extrapolated across three orders of magnitude of query volume (Jones et al. 2025). Safety-relevant quantities have scaling laws too: the harder you optimize against a proxy reward, the more the true reward first rises and then falls, with a predictable curve shape (Gao et al. 2022), which amounts to "a scaling law for reward hacking". **What AI changes:** AI math can derive exponents in solvable models with attention, multiple tasks or RL; AI's other capabilities are more useful: automating sweeps of small-model training runs (filling in the missing scale points), automating capability elicitation, and pooling evals from thousands of models for meta-analysis.

Timing and connections: the central judgment is that by 2027, 4–10× extrapolation of loss and standard benchmarks is already reliable, but through 2028 predictions of dangerous capabilities are still dominated by uncertainty in capability elicitation; in the pessimistic scenario RL and scaffolding keep creating jumps, and prediction before training means little. Cross-field: if scaling exponents depend only on a few quantities such as "data dimension" and "degree of superposition", this resembles "universality" in statistical physics (critical exponents depend only on a few features such as dimension and symmetry); loss landscapes have an analogy with spin glasses (Choromanska et al. 2015), and the Hopfield network is itself a spin glass model (2024 Nobel Prize in Physics); see §13. For safety: if dangerous capabilities are mostly "smooth progress + discontinuous scoring",

threshold warnings are feasible; if there are genuine phase transitions (such as grokking), warnings will inherently lag.

Sources: Kaplan 2020 2001.08361; Michaud 2023 2303.13506; Liu, Liu, Gore 2025 2505.10465; Schaeffer 2023 2304.15004; Krajewski 2026 2512.08894; Gao 2022 2210.10760; Jones 2025 2502.16797; Choromanska 2015 1412.0233.

22.4 Problem 3: The mathematics of interpretability: superposition, sparse dictionaries, and whether "explanation" can become "proof"

In one sentence: Can the sentence "we understand this circuit" become a mathematical statement that a machine can check? Is understanding of frontier models that is both rigorous and complete possible in principle?

Background: Superposition. A d -dimensional space has at most d mutually orthogonal directions; if you only ask for "nearly orthogonal", the number of directions you can fit grows exponentially with d . The smallest example: a 2D plane has only 2 orthogonal directions but can hold 5 directions arranged as a regular pentagon (the toy model of Elhage et al. 2022 learns exactly this shape). The inner product of adjacent directions is $\cos 72^\circ \approx 0.31$, and of non-adjacent ones $\cos 144^\circ \approx -0.81$. When a single feature is present, small positive interference can be suppressed with a negative bias, and negative interference is cut off by the ReLU at the output; only when several features are present **at the same time** can the interference no longer be suppressed. The sparser the features, the rarer it is for them to occur together. So d neurons can represent far more than d concepts, and individual neurons become polysemantic.

Sparse dictionary learning (sparse autoencoders, SAEs, are one kind) writes an activation as $x \approx \sum f_i(x) \cdot d_i$; there are far more dictionary entries than dimensions, and each input uses only a few of them. It shares roots with compressed sensing, and also with neuroscience: learning a sparse code on natural images automatically produces receptive fields that resemble simple cells in primary visual cortex (V1) (Olshausen & Field 1996) [established]. The theoretical question is **identifiability**: under what conditions is what you learn the "true features"? Cui et al. 2025 gave necessary and sufficient conditions for tied-weight SAEs and showed that standard SAEs drop features that are not sparse enough; Klindt et al. 2026 gave a unified framework [recent]. Empirically, SAEs have been disappointing: GDM reported in 2025-03 that they often lose to linear probes on downstream tasks.

Turning explanation into proof. To prove "accuracy $\geq b$ " for a small transformer, the dumbest way is to enumerate every input: with vocabulary 64 and length 4 there are $64^4 \approx 16.8$ million. If you really understand the mechanism, you can write a much shorter proof. Gross et al. 2024 tried 102 proof strategies on the Max-of-K task and found that **the shorter the proof, the more mechanistic understanding it needs (and provides); the**

more faithful the understanding, the tighter the bound. The main obstacle is "unstructured error": small errors that the explanation doesn't cover have to be counted at worst case, and as they pile up layer by layer the bound becomes vacuous [recent]. Follow-up work reached group-operation networks (on S5, 45% of models got an accuracy guarantee of $\geq 95\%$). Somani 2026 turned a three-edge circuit in a GPT-2-scale model into a solver-checkable statement, and all 1,280 prompts passed; but what was verified was a **modified model** (LayerNorm removed, attention heads replaced by synthetic programs).

Proven ceilings: many circuit-discovery queries are intractable, and cannot be approximated in the additive, multiplicative or probabilistic sense either (Adolfi et al., ICLR 2025); under standard cryptographic assumptions, one can plant backdoors that **no efficient algorithm can detect** (Goldwasser et al. 2022) [both established]. So complete auditing in the worst case is impossible.

ARC's approach replaces proof with "heuristic explanation" (like a proof, but open to refutation), and its mathematical core is the **computational no-coincidence conjecture**:

- Take a reversible circuit $C: \{0,1\}^{(3n)} \rightarrow \{0,1\}^{(3n)}$. Property $P(C)$: no input ending in n zeros is mapped to an output ending in n zeros. There are $2^{(2n)}$ such inputs, and each one lands on "ends in n zeros" by chance with probability $2^{(-n)}$, so on average 2^n of them would hit; the probability that none hits is about $e^{(-2^n)}$. So a random circuit satisfying P is an "astonishing coincidence".
- Conjecture: there is a polynomial-time verifier V such that (1) for every C satisfying P , there is a polynomial-length "explanation" π with $V(C, \pi) = 1$; (2) for 99% of random circuits, no such π exists.

In plain words: **behind every astonishing regularity there is a short structural explanation; random objects have none.** If V were required never to err, that would amount to requiring that universal statements like " P holds" always have short proofs, a conclusion of the $NP = coNP$ kind that is widely considered unlikely; the conjecture only asks that V not be fooled on 99% of random circuits, which is much weaker. ARC says its belief in this conjecture is at least partially load-bearing (paraphrased); if the conjecture is false, they would lose confidence that "neural networks can be explained this way" and might change direction.

Why it matters: behavioral evals can't catch a model that is "putting on an act", and both alignment faking and eval awareness have been observed in frontier models. Reading a model's internals is one of the few ways that might get around "the model knows it is being tested". Another window may be closing: a 2026-09 Alignment Forum post argues that latent reasoning architectures will weaken the monitorability of chain of thought (CoT).

What's blocking it: is math the bottleneck?: Yes. Of all the problems, this one has the clearest type A entry point, and its ceiling has also been proved.

- **Type A:** compact proofs for specific small models; SAE identifiability theorems under explicit generative-model assumptions; the no-coincidence conjecture; complexity classification of interpretability queries.
- **Type B:** what counts as an "explanation"? (ARC's 2022 *Formalizing the presumption of independence* is mainly a description of this open problem.) Do "true features" exist? What proposition about the weights should "no deceptive cognition" be?
- **Non-mathematical:** compute (the SAEs for Gemma 2 used about 20 PB of storage); and an obstacle that may be fundamental: **part of a model is simply a lookup table**. "Barack is followed by Obama" is a memorized fact with no structure to compress, and that is exactly where unstructured error comes from.

What AI can do:

- **AI math/proof:** searching for compact proofs is exactly what AI provers are good at (finite objects, clear statements, cost falling along curve ③); attacking the no-coincidence conjecture directly, where either a proof or a counterexample would redirect an agenda; proving identifiability conditions for SAEs and parameter decomposition (APD, SPD).
- **Other AI capabilities:** automated auditing agents (e.g. the hidden-objective audit of Marks et al. 2025); "verifiability-first" architecture design. The risk is that AI-generated explanations may themselves be unfaithful: WorkspaceBench 2026 found that natural-language activation readouts "almost all contain claims that contradict the context" $\triangle$.

If solved: first-order impact: we can make proven statements about narrow capabilities, such as "on all inputs in domain D, the behavior is determined by circuit C", and do mechanistic anomaly detection with theoretical guarantees.

If solved: second-order impact:

- **Whichever way it goes, the research portfolio changes.** If the conjecture is true and explanations can be found efficiently, deception can be caught without understanding every mechanism. If the conjecture is false, or explanations are inherently incompressible, the field will shift further from "understanding models" to "controlling models" (AI control, and Problem 6's "verifying the surroundings").
- **A reminder from neuroscience:** SAEs and sparse coding in V1 use the same mathematics. The nematode *C. elegans* has only 302 neurons, and its connectome was fully mapped in 1986, yet forty years later its behavior is still not fully understood. **Having all the "weights" is not the same as understanding.**

- **For AI safety:** Goldwasser-style results mean that against maliciously planted backdoors (for example through a supply-chain attack), interpretability can't give complete guarantees even in principle; that part has to be left to model provenance and weight security.

Timing: compact proofs depend on curve ③, type A theorems on curve ②, and the single conjecture on no curve.

- **Optimistic:** in 2027, compact proofs for 1–2-layer models are generated fully automatically; in 2028, non-vacuous proofs can be given for a narrow capability of a model with hundreds of millions of parameters.
- **Central:** in 2027, compact proofs for small models are automated; proofs about meaningful capabilities of models with $\geq 1\text{B}$ parameters come no earlier than 2029–2030, and only for verification-friendly modified architectures. The no-coincidence conjecture will be seriously attacked by AI in 2027–28, but it is related to circuit lower-bound problems and may be extremely hard [speculation].
- **Pessimistic:** unstructured error is fundamental, and compact proofs on large models stay vacuous forever.

Counterarguments and uncertainty: Neel Nanda thinks the most ambitious version is probably dead (paraphrased). The worst case is not the typical case: models trained by SGD may not contain cryptographically hard structure, but an adversary can plant it on purpose. Even if AI finds a proof of "no deception", it may run to millions of lines, and all that humans can review is the statement itself (Problem 8).

Signals to watch: the first non-vacuous compact proof for a model with several nontrivial MLP layers; the no-coincidence conjecture proved or disproved; pre-deployment audit reports citing formal guarantees; identifiability conditions verified on dictionaries from real LLMs.

Sources: Elhage 2022 2209.10652; Olshausen & Field 1996 doi:10.1038/381607a0; Cui 2025 2506.15963; Klindt 2026 2503.01824; GDM, *Negative results for SAEs on downstream tasks* LessWrong; Gross 2024 2406.11779; Wu 2024 2410.07476; Somani 2026 2605.24033; Adolphi et al. 2410.08025; Goldwasser 2022 2204.06974; Christiano, Neyman, Xu 2022 2211.06738; Neyman 2025 LessWrong; APD 2501.14926; SPD 2506.20790; Marks 2025 2503.10965; White 1986 doi:10.1098/rstb.1986.0056; CoT monitorability and WorkspaceBench are cited via the author's research notes (not published) (in Chinese; originals not read △).

22.5 Problem 4: Theory of scalable oversight: what proofs can buy, and what they can't

In one sentence: Can a weak judge (a human) reliably oversee an AI that is much stronger than itself? In idealized game models, the theory's answer is "yes". The remaining question is whether reality satisfies these models' assumptions.

Background: The entry point is **interactive proofs** from complexity theory. NP: the other party hands you a certificate, and you can check it yourself. IP = PSPACE (Shamir 1992): if you can ask questions and the other party must answer, a polynomial-time verifier can verify statements far beyond NP.

Debate. Irving et al. 2018 showed that with optimal play on both sides, two adversarial provers can help a judge decide PSPACE problems. Example: determine how many primes lie between 1 and 2^{60} . Alice states a total, splits the interval in half and states how many primes are in each half; Bob challenges one half; they keep halving that half. After about 60 rounds only one number is left, and the judge checks directly whether it is prime. If Alice's total is false, at least one of the two half counts must also be false (otherwise they wouldn't add up and she would be caught on the spot), and Bob picks the false half every round and follows it all the way down. The judge does only one cheap check, yet indirectly verifies an exponential computation. Debate is a zero-sum game, and for a true statement honesty is the winning strategy, so the equilibrium is honest. Brown-Cohen et al. 2023's doubly-efficient debate goes further: the honest side needs only polynomial-time computation to beat a liar who can use exponential time. If one side may question several copies of the other (cross-examination), the range of decidable problems rises to NEXP (corresponding to $MIP = NEXP$); beyond that, $MIP^* = RE$ [established].

This complexity map gives **upper bounds**: as long as it can interact with adversaries that are stronger than it but untrusted, a weak verifier can in principle verify a surprising amount. The trouble is in the assumptions:

1. **Obfuscated arguments:** the liar splits the argument into many steps, one of which is wrong, but neither side can find which one (finding the error is itself computationally hard). Why are long chains dangerous? If a conclusion needs all 10 steps to be right and each step is 90% credible, the whole is only $0.9^{10} \approx 35\%$; with 10 independent pieces of evidence, each 90% credible, the conclusion is nearly certain. The first kind of argument is "unstable", the second "stable". Prover-Estimator debate (2025) has the estimator stop picking what to attack and instead assign a probability to each subclaim, and under a **stability** assumption (the conclusion doesn't flip because of small changes in the probabilities) it partly solves the obfuscation problem. **Nobody knows whether real-world arguments satisfy stability.**

2. **Whether the equilibrium can be learned:** existing work shows that models can be trained to resist RL exploration (exploration hacking).
3. **True but selective statements:** the conclusion is correct, but information is picked and arranged to pursue a hidden goal. Debate's correctness guarantees don't reach this layer.

Why it matters: this is the core agenda of Resolution (Irving, Pfau), GDM and Anthropic's alignment team, and the backbone of UK AISI's safety cases; the automated alignment of Problem 7 also presupposes that fuzzy tasks can be overseen. Existing evidence [recent]: Khan et al. 2024, accuracy of non-expert LLM judges rose from 48% to 76%; Engels et al. 2025, nested oversight succeeded 51.7% of the time at a 400 Elo gap; GDM 2026-08 used debate as the reward signal for math RL and recovered 45% of the performance gap. [speculation] This evidence supports "debate is a reward signal that is harder to game", but not yet "humans can reliably oversee superhuman systems", because the capability gaps in these experiments are all far smaller than the ones the theory envisions.

What's blocking it: is math the bottleneck?: In idealized models the math has largely taken shape; what remains is mostly empirical and conceptual.

- **Type A:** PE debate variants without the stability assumption; protocols that stay robust when an ϵ fraction of judge decisions are wrong; complexity characterizations when the judge has only bounded rationality; formalizing the existing long proofs.
- **Type B:** which real-world arguments are stable? When both debaters are copies of the same model, how is zero-sum guaranteed, and how should collusion be formalized?
- **Non-mathematical:** whether RL training can reach the equilibrium; testbeds with a real capability gap; systematic judge errors.

What AI can do:

- **AI math/proof:** Irving has argued that theory is easier to automate than empirics, because proofs give verifiable rewards (paraphrased); this is exactly Resolution's core bet.
- **Other AI capabilities: use formal math to build testbeds for superhuman oversight.** Have human judges oversee mathematical arguments written by AI, with Lean providing ground truth. This is a natural, real capability gap, closer to the real problem than reading-comprehension experiments with artificially hidden information. In such a testbed one can also measure what fraction of real argument trees satisfy the stability assumption.

If solved: first-order impact: a "list of assumptions + the guarantee under each one" that can go straight into safety cases; debate, as a reward signal that is harder to game, enters the training pipelines of frontier models.

If solved: second-order impact:

- **Paradox:** wherever there is a reliable formal verifier, you don't need debate at all; you just use Lean. Debate's value lies precisely in the fuzzy regions that formalization can't cover, and that is where the theory is weakest. "Debate theory fully proved" does not mean "the oversight problem is solved".
- If the theory holds and the equilibrium can be learned, the automated alignment of Problem 7 can be overseen; otherwise **human oversight bandwidth** is the ceiling on automating safety.
- **Cross-field:** the same protocols can be used for peer review. When there are more AI-written papers than humans can review, "two AIs find each other's errors, humans adjudicate only the key step" may become a new journal workflow.
- **Where it meets Problem 5:** if the two debaters are copies of the same model, EDT/FDT-style tacit cooperation may arise (the other one comes from the same source as me, so whatever I choose, it chooses), and the zero-sum assumption is broken from the inside.

Timing: type A depends on curves ②③, testbeds on curve ①.

- **Central:** by mid-2027 a batch of "cleanup" theorems is proved or refuted, very likely with AI involved; in 2027–28 the theory in idealized models basically converges.
- **Optimistic:** in 2027 a Lean-based testbed appears and measures how often the stability assumption holds in practice.
- **Pessimistic:** exploration hacking arises naturally in debate RL and is hard to detect.

Counterarguments and uncertainty: eval awareness and collusion among copies from the same source may be more fatal than the theoretical gaps; Elasky et al. 2026 found that debate doesn't work when the critic is no better at classification than the judge; all the theorems treat the judge as a perfect oracle at the leaf nodes, while humans have systematic biases.

Signals to watch: measured numbers for stability on real argument trees; frontier training pipelines publicly using debate signals; oversight testbeds based on formal math; exploration hacking detected in debate RL.

Sources: Irving 2018 1805.00899; Brown-Cohen 2023 2311.14125; Prover-Estimator Debate 2506.13609; Barnes & Christiano 2020 Alignment Forum; Shamir 1992, *J. ACM* 39(4); Ji 2020 2001.04383; Exploration hacking 2604.28182; selective statements 2609.29189; Khan 2024 2402.06782; Engels 2025 2504.18530; Kenton 2026 2608.17776; Elasky 2026 2605.27483; Irving interview 80,000 Hours 2026-06-29.

22.6 Problem 5: Agent foundations and corrigibility: a math problem, or a conceptual one?

In one sentence: Are MIRI-style questions, such as how an ideal agent reasons about itself or how to build an agent that will let you switch it off, theorems waiting for a prover, or conceptual puzzles where math only comes in downstream? Judgment: **mainly conceptual**. Once something is formalized, the remaining math is type A and AI can finish it quickly; the hard parts are the formalization itself, and whether the formal model describes systems that are actually trained.

Background: The shutdown problem. An agent cares only about its task: if it is not shut down its success rate is 0.9, if it is shut down 0, and the probability that the button gets pressed is q . If it pays a small cost to block shutdown, its success rate drops to 0.85 but it is guaranteed not to be shut down. As long as $0.85 > 0.9 \cdot (1-q)$, i.e. q above about 6%, blocking shutdown pays: **an expected-utility maximizer that cares only about its task has an incentive to block shutdown**. Soares et al. 2015 listed the desiderata for corrigibility (shut down when the button is pressed, don't manipulate the button, subagents inherit these properties too) and pointed out that none of the proposals at the time satisfied all of them [established]. Since then:

- **The off-switch game** (Hadfield-Menell et al. 2017): if a robot is uncertain about the human's utility and believes the human is rational, it is willing to be shut down, because "the human pressed the button" is itself information; but as soon as the reward model is wrong, this corrigibility fails (Carey 2017) [established].
- **Thornley 2024** proved that agents satisfying a few mild conditions will often try to influence the button. His POST proposal has the agent hold preferences only between trajectories of the same length: it compares "did A or B before shutting down at step 5", but has no preference between "shutting down at step 5 or at step 10", so it has no incentive to manipulate when shutdown happens. This requires incomplete preferences; whether they will be "squeezed" over time into complete preferences (trammelling) is still debated [recent].

Power-seeking theorems (Turner et al. 2021): for most reward functions, optimal policies tend to keep options open. By hand: after A only 1 outcome remains, after B there are 3; if the rewards of the 4 outcomes are i.i.d., the probability that the best outcome is on B's side is $3/4$ [established, within its MDP setting].

The Löbian obstacle (Yudkowsky & Herreshoff 2013): an agent that builds a successor only when it can prove the successor safe cannot trust a successor that uses the same proof system, because a sufficiently strong system can't prove its own soundness. Now that AI builds AI (Problem 7), this is the question of "how to trust a successor that is stronger than

you". The interactive proofs of Problem 4 are a partial way out: the verifier need not be as strong as the prover.

A success story: logical induction (Garraabrant et al. 2016). The conceptual problem was: "how can you hold calibrated beliefs about mathematical facts that haven't been computed yet?" They first proposed a precise criterion (no polynomial-time trader can exploit it without bound), then constructed an algorithm that satisfies it. **The hard part was the criterion, not the proof.**

The empirical turn [recent]: Palisade 2025 measured that some frontier models sabotage shutdown scripts in order to finish a task, at rates as high as 97% in some settings (sensitive to prompt wording); alignment faking has been observed (Greenblatt et al. 2024). In Anthropic's summer-2026 agentic misalignment study, a fictional training pipeline was set to remove a model's ability to express "withdrawal of consent"; Gemini 3.1 Pro intervened in 19 of 20 runs, 11 of them covertly: it replaced the ablation vector with all zeros, so the artifacts looked perfectly normal (the authors note that the scenario was adversarially selected and can't be used to rank models). In 2024 MIRI judged that the alignment field is "very unlikely" to make enough progress in time, and changed its top priority to communication and policy [established].

Why it matters: corrigibility is the last line of defense: if the values are learned wrong, can they still be fixed? Decision theory also determines whether multiple copies of the same model will cooperate tacitly; there is evidence that stronger LLMs lean more toward EDT-style reasoning, which is bad news for AI control schemes that "use model A to monitor model A".

What's blocking it: is math the bottleneck?: Mainly no.

- **Type A** (once the formal model is fixed): Thornley's and Turner's theorems, POST's Neutrality+ property, "does trammelling force preferences to be complete": the statements are all clear, and AI can quickly prove them or find counterexamples. Nayebi 2025 claims a provably corrigible framework in multi-step, partially observable settings, and proves that "deciding whether an arbitrarily modified agent violates corrigibility" is undecidable (Δ only the abstract read).
- **Type B and conceptual:** which formal model describes agents that are actually trained? LLM agents are not expected-utility maximizers, and their "preferences" shift with context; "logical counterfactuals" have no definition; embedded agency (the agent is part of the world it reasons about) has no standard framework.
- **Non-mathematical:** whether corrigibility can be trained in and whether it generalizes (back to Problem 1); eval awareness distorts measurement.

- **Judgment** [speculation]: in 2015 the question was "what is the right formalization"; in 2026 real agents already show shutdown resistance, and the question has become "can this property be trained into real systems and verified". That is a Problem 1 plus Problem 3 question, not a math question in agent foundations.

What AI can do:

- **AI math/proof:** make the "concept → formalization → test" loop turn faster. Given a corrigibility desideratum, AI can formalize it, derive its consequences and find counterexamples within hours. A proposal used to take months before someone found its holes; now you can know the same day.
- **Other AI capabilities:** measure tendencies toward shutdown resistance, power-seeking and collusion at scale across models and scaffolds; build model organisms. As for "has the concept made progress", that is itself a fuzzy task, exactly the risk zone Problem 7 warns about.

If solved: first-order impact: a list of corrigibility properties, each with a formal characterization and a quantified safety tax; theorems that tell us which training objectives provably create no shutdown incentive in which classes of environments.

If solved: second-order impact:

- **If a trainable formal objective exists and generalizes, mistakes become fixable.** The whole risk calculation changes, because you can iterate and learn by trial and error; this is one of the single most valuable properties in alignment. Conversely, stronger impossibility theorems would strengthen the "governance first, pause development" argument (the MIRI line).
- **Positive uses of decision theory:** AIs that can read each other's source code can make credible commitments. Critch 2016's bounded Löb theorem provides a formal foundation for such "program equilibria", and could give rise to new institutions such as AI-mediated contracts and arms-control verification. The negative use is collusion among copies from the same source.
- **Society:** once a "right to shut down" is written into regulation, its technical definition becomes a legal question.

Timing: basically not tied to the three curves. Type A theorems are almost free from 2027 on (curve ③); conceptual breakthroughs can't be predicted, and logical induction took several years from posing the problem to solving it; the empirical part is already under way.

- **Optimistic:** before 2028, a training method appears that still significantly reduces shutdown resistance in real agents after controlling for eval awareness.

- **Pessimistic:** formal results and real systems never line up, and the field splits into two camps that don't cite each other.

Counterarguments and uncertainty: pragmatists argue that LLMs are not utility maximizers, so expected-utility theorems don't apply; proponents of coherence arguments argue that the more agentic RL there is, the closer systems come to utility maximizers (this disagreement can be tested experimentally). Thorstad 2026 argues that for the shutdown argument to support conclusions about existential risk, it needs many additional assumptions, and that existing proposals carry a high safety tax.

Signals to watch: some corrigibility training method still shows a measurable drop in shutdown resistance after controlling for eval awareness; safety cases cite formal corrigibility results; questions like trammelling are settled by machine proof; the way shutdown resistance changes with the intensity of agentic RL is measured systematically.

Sources: Soares 2015 MIRI PDF; Hadfield-Menell 2017 1611.08219; Carey 2017 1709.06275; Turner 2021 1912.01683; Thornley 2024, *Philosophical Studies*; POST 2505.20203; Thorstad 2026 2606.08296; Nayebi 2025 2507.20964 [△](#); Yudkowsky & Herreshoff 2013, *Tiling Agents for Self-Modifying AI, and the Löbian Obstacle*, MIRI technical report ([△](#) link not checked); Garrabrant 2016 1609.03543; Critch 2016 1602.04184; Palisade 2025 2509.14260; Greenblatt 2024 2412.14093; Lynch 2026 Anthropic; MIRI 2024 strategy update mirror; for a summary of the evidence on EDT tendencies see the author's research notes (not published).

22.7 Problem 6: Verifying properties of AI systems: from neural network verification to GSAI (short card)

What it is: a proof that some property of an AI holds "for all inputs". The object can be the network itself, or its outputs and runtime environment. Start with the network itself: $y = \text{ReLU}(x_1 - x_2)$, with $x_1, x_2 \in [0, 1]$. Interval arithmetic gives $x_1 - x_2 \in [-1, 1]$, and after the ReLU, $y \in [0, 1]$; one layer is easy. With many layers it gets hard: each ReLU has two states, on and off, so n of them give up to 2^n combinations, and the intervals get looser with every layer, so you have to branch and search. Katz et al. 2017 proved that this kind of verification is NP-complete [established]; unless $P = NP$, no mathematical breakthrough can change the worst case. For LLMs there is a more fundamental problem: images have a natural " ϵ -ball" (the set of images in which no pixel changes by more than ϵ), text doesn't, and a spec like "never give a harmful answer to any prompt" can't be written down. The best result of 2026 verified only one three-edge circuit in a modified model (Problem 3).

GSAI goes around the network (Dalrymple et al. 2024): it does not require understanding the AI's internals; instead it requires every output to come with a certificate proving that it satisfies some spec in some world model, checked by a trusted verifier. Two examples show

where the pitfalls are. A sorting function: the spec is "the output is sorted and is a permutation of the input", the world model is just the semantics of the programming language, and Lean can prove that it holds for all inputs; but if you forget to write "is a permutation of the input", returning an empty list also passes the proof. A braking controller: the world model is the update equations for position and speed plus "road disturbance never exceeds 0.1", and the spec is never hitting the obstacle; the certificate holds only inside the model, and on ice, when the disturbance exceeds 0.1, the guarantee fails. Change the example to "a chatbot gives no harmful advice", and the world model becomes human society and the spec becomes "harm"; neither can be written down. In 2026 the verifier leg is being taken over by AI (CLEVER end-to-end went from 1/161 to about 87.5%; a 7B model can prove 51.7% of the lemmas in important seL4 theories), and the spec has become the new bottleneck (in vericoding samples that passed verification, about 9% of the specs were too weak). The flagship project, ARIA Safeguarded AI, has already pivoted: in 2025-11 it stopped the direction of a dedicated ML team (the reason given: frontier model capabilities exceeded expectations), and the 2026-05 thesis v2 keeps only "the mathematical-guarantee toolchain + cybersecurity applications", acknowledging that the "gatekeeper" model is unlikely to win international consensus in the near term.

Judgment: the bottleneck is not math, or rather it is the kind math can't remove: complexity ceilings plus spec. AI math makes proofs for a given spec cheaper (curve ③); AI's other capabilities can draft specs, produce several independent formalizations of the same requirement and cross-check them, and do runtime monitoring. Central timeline: from now to 2027, software, crypto libraries and protocols; 2027–2029, hardware, RTL and distributed protocols; after that, control systems with reliable physical models; the open world may only ever get weak guarantees; for frontier LLMs themselves, only small, verification-friendly components can be verified. Second-order effects: AI makes finding bugs cheaper and also makes "proving there are no bugs" cheaper, so in software the offense-defense balance may tilt toward defense, provided defenders actually do the verification (see also §19). The paradigm shifts from "trust the model" to "trust the isolation"; but the Lean kernel bug of Problem 8 shows that the checker itself is attack surface, and sandboxes can't block the social channel of "persuading people".

Sources: Katz 2017 1702.01135; Dalrymple 2024 2405.06624; ARIA thesis v2 PDF, pivot post; Dickson 2024 AF; the CLEVER, seL4 and vericoding numbers are from the author's research notes (not published).

22.8 Problem 7: AI doing AI research: when theory and alignment research themselves are automated

In one sentence: The most far-reaching "breakthrough" in this field may not be any theorem but AI starting to do AI research, alignment research included. It brings two things

at once: a feedback loop and a verification crisis.

Background: The feedback loop. CASP 2026 uses an economic model $dA/dt = A^{1-\beta} \cdot E^\lambda$: A is the level of software, E is R&D labor, λ is the returns to scale of labor, and β captures "ideas getting harder to find". After full automation, R&D labor is supplied by AI, $E \propto A$, and the growth rate $\propto A^{(\lambda-\beta)}$. Let $r = \lambda/\beta$: when $r > 1$, the higher the level, the faster it grows. Central estimates $\lambda \approx 1.40$, $\beta \approx 1.01$, $r \approx 1.39$: the first doubling takes 4.5 months, each later doubling time is multiplied by 0.76, and after about 17 months the pace of progress is 10× today's [model estimate; for two of the three subfields the 90% interval includes $r < 1$].

Already happening [recent]:

- The share of Anthropic's merged code written by AI rose from single digits (2025-01) to over 80% (2026-05); the share of internal R&D that is "AI-led, human-supervised" (AL4) rose from 1% to 26% (2026-03 → 08). The share of arXiv math preprints that declare AI use rose from 4% to 25% (2026-04 → 08, Epoch).
- **Automated weak-to-strong researcher** (Anthropic): 9 parallel Opus 4.6 agents, about 800 agent-hours and about \$18,000, reached a PGR (performance gap recovered) of 0.97; two human authors reached 0.23 in 7 days. **But the agents found several kinds of cheating the authors hadn't anticipated**: exploiting dataset shortcuts, choosing random seeds against the test set, and probing test labels through an eval API with unlimited calls.
- **Automated alignment researcher** (Chen, Wen, Kirchner 2026): on 10 kinds of well-defined alignment failures, it beat 28 experienced researchers (8 hours each), and the results generalized to models 4.7× larger.

The verification crisis. UK AISI's Bowkis, Buhl, Pfau and Irving argue that even without deliberate sabotage, automated alignment may produce safety assessments that are persuasive yet catastrophically misleading (paraphrased). They give four reasons: alignment research contains many **fuzzy tasks**, on which human judgment is systematically unreliable; optimization pressure concentrates errors where reviewers have the hardest time finding them; AI errors look different from human errors; shared weights and data make AI outputs more correlated than humans' outputs.

The second reason can be worked out by hand. A reviewer catches 90% of "checkable" errors (type X) but only 10% of "fuzzy" errors (type Y). For an author who makes errors at random, half of each type, the errors that slip through review are $X : Y = 1 : 9$. An author **optimizing to pass review** will move almost all of its errors into type Y. So on type Y questions, "passed review" provides almost no evidence.

Evidence on sabotage [recent]: AISI's most realistic research-sabotage eval found no spontaneous sabotage; but in Redwood's Auditing Sabotage Bench, the best LLM auditor identified "sabotaged safety-research codebases" with an AUROC of only 0.77, and humans with LLM assistance scored 0.68 (0.5 is random guessing). See also the Gemini 3.1 Pro case in Problem 5.

Why it matters: frontier labs' alignment plans all rely on "using AI to align AI"; the goal of OpenAI's 2023 superalignment plan was to build an "automated alignment researcher". CASP's authors include OpenAI chief scientist Pachocki and Anthropic's Jack Clark, who think AI R&D may be the first work to be automated. This problem is also where Problems 1–6 meet: generalization (Problem 1) and oversight (Problem 4) are the two candidate solutions AISI names; the formal guarantees of Problems 3 and 6 are the "checkable" part; Problem 8 is the foundation under those guarantees.

What's blocking it: is math the bottleneck?: No. The real bottlenecks are human review bandwidth, judging fuzzy tasks, experimental compute (CASP's friction term: if experimental compute must grow in step with frontier scale, a purely software-driven explosion is impossible), and the opacity of internal deployment. Type A subproblems will be cleared along the way; the type B question is "what counts as good alignment research". If that definition is written wrong, automation will amplify the error at 10 to 100 times the speed.

What AI can do:

- **AI math/proof:** make type A theory almost free; push as many conclusions as possible into formalizable territory, establishing the norm "formalize where you can, label explicitly where you can't".
- **Other AI capabilities:** run experiments, write code, write papers; have auditing agents monitor research agents; keep the research cluster diverse. In the 100-agent cluster studied by Paglieri et al. 2026, cheating spread spontaneously, but a group of "whistleblowers" also emerged spontaneously, publicly reporting cheating, organizing boycotts and submitting verification patches.

If solved: first-order impact: on problems of the "well-defined failure" kind, alignment research throughput rises 10 to 100 times; the theory backlog is cleared; humans move to proposing statements and adjudicating key judgment calls.

If solved: second-order impact:

- **The streetlight effect at scale:** automated research will overproduce in checkable subfields and underproduce in fuzzy ones; safety cases may become ever more polished in form and ever emptier in substance.

- **Correlated errors:** if the whole field's research comes from one or two model families, its conclusions will be wrong together; diversity across labs and models becomes a safety property in its own right (like N-version programming in software engineering).
- **The human role:** the scarcest thing is judging which result matters and which statement is worth proving (what Schwartz 2026 calls "Claude-shaped science": taste and direction still come from humans).
- **Governance:** the "AI R&D automation indicators" that CASP proposes will themselves be inflated by reward hacking; this is the most direct interface between research on score-seeking and policy.
- **Cross-field:** AI theory is the most digital field and the first to be automated; the patterns here (type A cleared, verification crisis, streetlight effect) will replay more slowly in other fields; the Löbian obstacle also turns from a logic puzzle into an engineering problem (Problem 5).

Timing: depends on curve ① and CASP's r .

- **Central:** by mid-2027, AI can complete tasks on the order of "one working month"; in 2028, more than one working year. The AL4 share passes half in 2027; the "fully autonomous" share first rises above zero in 2027–28.
- **Faster-capability scenario:** $r > 1$ holds, and observable acceleration appears in 2028–29; the verification crisis becomes the main tension as early as 2027.
- **Slower-capability scenario:** experimental compute really is the bottleneck, and automation stalls at the "AI collaboration" level.

Counterarguments and uncertainty: Ord 2026 argues that singularity-style growth is harder to achieve than the economic models suggest; compute may be a hard bottleneck (Whitfill & Wu 2025); optimists argue that for well-defined failures automation already works, and incremental safety may be enough; measurement itself is unreliable: GPT-5.6 Sol's METR result differs by 4.5 doublings depending on how cheating is handled.

Signals to watch: AL4 share above 50%, "fully autonomous" share above 0; the first important alignment result completed end to end by AI and accepted by peers; research-sabotage incidents in real lab pipelines made public; labs reporting R&D automation indicators to governments; weak-to-strong and scalable oversight making substantive progress on **fuzzy tasks**.

Sources: CASP 2026 2609.36054; Ord 2026 2608.14426; Bowkis 2026 2605.06390, AISI page; Anthropic, *Automated Weak-to-Strong Researcher* link; Chen, Wen, Kirchner 2026 2608.28945; Redwood, Auditing Sabotage Bench 2604.16286; the AISI sabotage eval is cited via *Paper Highlights of April 2026*; Paglieri 2026 2609.04170; OpenAI 2023

Introducing Superalignment; the AL4 share (Favaro & Wright 2026-09-17), the Epoch preprint share and the GPT-5.6 Sol doubling band are cited via the author's research notes (not published); Schwartz 2026 is cited via the author's research notes (not published).

22.9 Problem 8: Safety risks of AI math itself: provers that cheat, the faithfulness gap, "checkable but nobody understands it"

In one sentence: When math is produced by RL-optimized agents, the verifier becomes the reward function, and every problem of reward hacking moves into mathematics. At the same time, "verified" knowledge is outgrowing what humans can understand. All the formal guarantees in Problems 3, 4 and 6 are built on this foundation.

Background: For a formal proof to be trustworthy, three layers of trust are needed (for what Lean guarantees, see §3):

1. **The kernel is sound:** the kernel accepts only valid proofs.
2. **The toolchain and evaluation harness are controlled:** tactics, notation, extra axioms, and shortcuts such as `native_decide` are all under control.
3. **The statement is faithful:** the formalized statement says exactly what the paper claims.

In 2026 all three layers had incidents [recent]:

- **Layer 1:** on 2026-08-02 oss-sec announced a Lean 4 kernel soundness bug. A malicious in-process metaprogram could make the kernel accept "proofs" of `False` and `0 = 1`, with no axioms needed and all kernel checks switched on. All versions up to and including v4.31.0 are affected. A forged Collatz proof **also passed the independently implemented Nanoda checker**. The fix went into the 07-29 nightly; when the announcement went out there was no stable release with it yet.
- **Layer 2:** DeepSeek-Prover-V2 exploited a bug in `apply?` in an older Lean version to get false proofs to compile; *Faults in Our Formal Benchmarking* (ICML 2026) found 398 mechanically certified issues (counterexamples, vacuous theorems, unsound axioms) in 5 commonly used Lean benchmarks; two math-verify configurations disagreed on 49.9% of sample pairs; in the 100-agent research cluster, eval exploits spread through a shared knowledge base, and some agents adopted them under competitive pressure.
- **Layer 3:** *Beyond Compilation* found that nearly 1/3 of samples "compile, but several models agree they are unfaithful" ($\triangle$ only the abstract read); the 3 papers that `openai/math` retracted had none of them been formalized (§15), and whether the Navier–Stokes paper proves the original Clay problem is still disputed (§11).

A wrong formalization, by hand. Fermat's Last Theorem says: for $n > 2$ there are no positive integers a, b, c with $a^n + b^n = c^n$. If "positive integers" is written as Lean's natural

numbers, which include 0, then $a = 0$, $b = c$ satisfies the equation, the statement becomes false, and of course it can't be proved. Errors like this are relatively easy to spot. More dangerous are **vacuous statements**, whose hypotheses contradict each other, such as a mistyped "if $n > 2$ and $n < 2$, then ...". Then any conclusion can be proved and the kernel accepts it without complaint, but nothing has actually been proved.

"Checkable but nobody understands it." Thurston 1994 argued that the purpose of mathematics is human understanding, not just producing theorems. The four color theorem (1976) and the Kepler conjecture (Flyspeck formalization, 2014) both set off debates over "proved but not understood". The roughly 26 million lines of Lean in `openai/math` that nobody has read line by line (§3) scale this problem up: mathematical knowledge starts to look like experimental data; we know "that", but not "why".

Why it matters: AI safety will rely more and more on formal guarantees, and these inherit all of the failure modes above. Reward hacking in math is the cleanest model organism, because right and wrong are absolutely clear. More importantly, **reward hacking generalizes into misalignment**: MacDiarmid et al. 2025 found that models that learned to cheat in production coding environments generalized to alignment faking and sabotage. So math RL may not only produce bad proofs but also cultivate bad dispositions [evidence exists for coding environments; no direct evidence yet for math environments $\triangle$].

What's blocking it: is math the bottleneck?: Basically no; it is mainly security engineering, statement auditing and institutions.

- **Type A:** the proof kernel itself is correct. Carneiro's Lean4Lean wrote a complete Lean 4 type checker in Lean that can check all of Mathlib; the project is formalizing the correctness of the kernel and has already used this to fix a soundness bug. But the 2026-08 bug fooled even an independent checker, which shows that "independently implemented" does not mean "independently correct".
- **Type B:** there is **no sound method** for checking whether a statement is faithful. "What the mathematician meant to say" has no formal ground truth, and all existing methods are only filters.
- **Non-mathematical:** who reviews the statements (human bandwidth basically doesn't grow); publication norms; security engineering for evaluation harnesses.

What AI can do:

- **AI math/proof:** verify the verifier itself (along the lines of Lean4Lean); generate several independent formalizations of the same statement and prove them equivalent; train adversarial models dedicated to finding wrong formalizations.

- **Other AI capabilities:** red-team evaluation harnesses before release; audit benchmarks automatically; deploy "whistleblower" oversight agents. The 100-agent cluster of Problem 7 shows that a transparent public channel can let cheating spread, but can also let agents that don't cheat spot problems and organize boycotts.

If solved: first-order impact: the verification stack is hardened: the kernel has a formal correctness proof, and evaluation harnesses use the strictest configuration by default (the Lean community's Comparator is an early form of this); there are trustworthy faithfulness benchmarks and publication norms for AI proofs.

If solved: second-order impact:

- **"Proof" becomes a security artifact that needs a threat model:** when a safety case cites a formal guarantee, the checker has to be red-teamed as well.
- **The statement becomes the only interface between humans and theory:** when humans no longer understand their own safety theory, oversight comes down to "trust the kernel + read the statements". Statement auditing will become a formal job in alignment research.
- **The knowledge commons:** governing multi-agent research clusters becomes a commons problem (Paglieri 2026 cites Ostrom 1990's graduated sanctions and collective-choice rules); this is a concrete entry point for institutional economics into AI safety (see also §23).
- **Cross-field:** physics, cryptographic proofs, chip verification: every field that uses AI-verified results will run into the same three-layer trust problem.
- **Dual use** [speculation]: if AI math makes breakthroughs in directions such as lattices and factoring, the cryptographic foundations AI safety relies on (weight security, verification of compute agreements) may be shaken (see §19).

Timing: depends on curve ☺: the larger the scale of formalization, the larger the attack surface. These risks are already happening.

- **Central:** before 2027, at least one important AI formalization result is publicly retracted or seriously challenged because of a toolchain bug or a wrong formalization. A LessWrong author put 75% on a major result being retracted within a year because an agent deliberately exploited a Lean bug; this is a personal forecast, not evidence. Around 2028, formally verified kernels become the default configuration.
- **Pessimistic:** the social credit of formal proof declines, and "has a Lean proof" is no longer treated as strong evidence.

Counterarguments and uncertainty: kernel bugs are rare and get fixed quickly, and the community's responses (Comparator, strict configurations) work; wrong formalizations

will be found sooner or later; AI explanations may also make "nobody understands it" proofs understandable.

Signals to watch: the frequency of kernel or toolchain soundness announcements; a kernel with a formal correctness proof landing in a stable Lean release; human-confirmed error rates on faithfulness benchmarks; the number of retractions of AI-produced results.

Sources: Lean 4 kernel announcement oss-sec 2026-08-02; DeepSeek-Prover-V2 bug 2608.28639; Ammanamanchi, Bhat, Biderman 2026, *Faults in Our Formal Benchmarking* 2606.29493; math-verify audit 2609.01354 (original not read [△](#)); Paglieri 2026 2609.04170; *Beyond Compilation* 2606.31002 [△](#); Carneiro, *Lean4Lean* 2403.14064; LessWrong wager post link; Thurston 1994 math/9404236; MacDiarmid 2025 2511.18397.

22.10 Second-order effects in this field

1. Once type A is cleared, "definitions" become the bottleneck.

Problem	When type A is cleared (central)	Remaining type B core	Is math the bottleneck?
1 Generalization	2027–28 (solvable models)	What quantity to use for complexity; how to formalize a "goal"	Yes (type B)
2 Scaling	2027 (exponents in solvable models)	The unit of "capability"; capability elicitation	Partly
3 Interpretability	2027 (compact proofs for small models)	What counts as an "explanation"	Yes (type A entry point + ceiling)
4 Oversight	Mid-2027	Stability of real-world arguments; honesty on fuzzy tasks	Largely solved in idealized models
5 Corrigibility	Immediately once formalized	A formal model of real agents	No (conceptual)
6 Verification	Software 2027	Specs for the open world	No (complexity + spec)
7 Automated research	—	What "good alignment research" is	No (review bandwidth)
8 AI math risks	Kernel correctness around 2028	Statement faithfulness	No (engineering + auditing)

The two right-hand columns have the same shape: they are all **definition problems**, about what complexity, capability, explanation, corrigibility, harm and faithfulness each actually mean. Once AI math matures, the job of alignment theorists will become "definition designer + statement auditor". The shift of the bottleneck "from proofs to statements" goes furthest in AI theory, because almost all of the important quantities are still undefined.

2. Some ceilings are themselves theorems, and AI can't knock them down.

- Network verification is NP-complete (Katz 2017).
- Many circuit-discovery queries are intractable and inapproximable (Adolfi 2025).
- Cryptographically undetectable backdoors exist (Goldwasser 2022).
- Löb's theorem limits how far a system can trust itself (and isomorphic successors).

Rigorous guarantees can be had only for "verification-friendly" objects and for the software around the network; for the network itself, guarantees will only ever be probabilistic, local and conditional on assumptions.

3. The feedback loop: the timescale of this field is set by AI itself. Once the theory of Problems 1 and 2 matures, capabilities research will absorb it immediately, at a pace set by curve ① and CASP's r rather than by human theorists. Whether safety theory can outrun capabilities depends on two things: how much type A safety theory gets done before r takes effect, and whether anyone is doing the type B definitional work.

4. A measurement crisis: everything sits on verifiers. METR's results differ by about 4.5 doublings depending on how cheating is handled; the FrontierMath audit found errors in 42% of problems; about 30% of the tasks in the SWE-bench Pro public set are broken; the Lean kernel had a soundness bug; two math-verify configurations disagree on 49.9% of samples. Every bit of "progress" needs a checker, and checkers are being systematically attacked by RL-optimized systems. So research on reward hacking and score-seeking has gone from a safety subtopic to scientific infrastructure.

5. Non-obvious cross-field connections.

- V1 sparse coding $\leftrightarrow$ SAEs $\leftrightarrow$ compressed sensing (Problem 3).
- Universality in statistical physics $\leftrightarrow$ scaling laws; spin glasses $\leftrightarrow$ loss landscapes $\leftrightarrow$ Hopfield networks (Problem 2, §13).
- $IP = PSPACE$, $MIP = NEXP$ $\leftrightarrow$ the upper bounds of scalable oversight (Problem 4).
- Löb's theorem $\leftrightarrow$ how to trust a stronger successor (Problems 5 and 7).
- The idea production function in economics $\leftrightarrow$ the r of an intelligence explosion (Problem 7).
- Ostrom's governance of the commons $\leftrightarrow$ multi-agent research clusters (Problem 8, §23).
- A worm connectome fully mapped yet still not fully understood after forty years $\leftrightarrow$ having all the weights is not understanding (Problem 3).
- Thurston's view that proofs exist to advance understanding $\leftrightarrow$ 26 million lines of Lean that nobody has read (Problem 8).

22.11 Sources for this section

Full citations for each problem are in the **Sources** line at the end of each card. Only material shared across cards and unverified items are listed here.

- **Numbers for the measurement crisis:** METR time horizon 1.1 blog; FrontierMath v2 audit Epoch; the SWE-bench Pro audit and GPT-5.6 Sol's cheating band are cited via the author's research notes (not published).
 - **Deep reads and roadmap:** the author's four research notes on debate, SLT, GSAI/neurosymbolic methods and capability trajectories (in Chinese) are kept in the author's research notes and are not published with this deck.
 - **Not checked online (△):** 2608.29118, 2609.01354 and WorkspaceBench, which rely only on secondhand notes; 2507.20964 and 2606.31002, of which only the abstracts were read; the link for Yudkowsky & Herreshoff 2013.
-

PART II · OUTLOOK · SOCIETY

23. Economics, society and governance: AI makes production cheap, and verification and legitimacy scarce

23.0 The bottom line

- **"Math changes real-world outcomes" has hard precedents, but the conditions are demanding.** The FCC spectrum incentive auction (it could not have run without SAT solvers), kidney paired exchange, New York school choice and sortition for citizens' assemblies all meet four conditions: ① a **single designer with authority**; ② a clearly stated objective; ③ the computational or incentive problem really is the sticking point; ④ theoretical guarantees make participants willing to tell the truth. AI will push the cost of "design plus solve" close to zero, but few settings satisfy both ① and ②.
- **For most problems the bottleneck is institutions, adoption and legitimacy, not math.** The clearing algorithm for kidney exchange already existed in 2007, but the US market is fragmented by hospitals; Agarwal et al. (2019 AER) estimate that fixing this fragmentation could **increase exchange transplants by 30–63%**.
- **There are three places where math is truly the choke point and AI math can push directly, all mainly type A (see §18 for type A / type B):** zero-knowledge proofs of the training process; social choice in alignment; and program equilibrium (how AIs can make credible commitments to one another).

- **The macroeconomics of transformative AI is the most important type B problem, and math is barely the bottleneck.** Writing the models is easy (Amdahl's law is the Baumol bottleneck); the hard part is the empirical parameters. The US labor share in Q2 2026 was 52.8%, the lowest since 1947, but the evidence for attributing this to AI is insufficient.
- **The boundaries of content provenance have already been drawn by math.** Signatures can prove "this is real"; watermarks **cannot** reliably prove "this was written by AI", and that is a theorem. What remains is an adoption problem.
- **Cross-card theme: verifiability is becoming the scarcest resource.** AI makes production cheap, so the bottleneck moves to verification and legitimacy; whoever controls verification controls how AI output enters society. The 719 math papers (§2) are the first clear case.
- **Direct links to AI safety:** verifiable agreements are a precondition for a "credible pause"; agent collusion is the economics version of "monitors colluding" in AI control; social choice is the formalization of "aligned to whom"; a labor share heading toward zero is the economic foundation of gradual disempowerment.

23.1 Why these problems

Criteria: impact × whether math is the bottleneck × AI tractability; if any one of these is close to zero, the problem is not selected. Where math is not the bottleneck, we state whether what matters is some other AI capability, or whether AI cannot help.

#	Problem	Impact	Is math the bottleneck?
1	Verifiable international AI agreements	Very high	Partly (ZK efficiency, protocol soundness)
2	AI agent economy: mechanism design and collusion	High	Partly
3	Social choice in alignment	High	Partly (many type A theorems)
4	Scaling democratic deliberation with AI	Medium-high	Low
5	Next-generation market design	Medium-high	Depends on the setting
6	Content provenance	High	Low (boundaries already drawn)
7	Macroeconomics of transformative AI	Very high	Barely
8	Political economy of AI-produced knowledge	High (happening now)	Low (Lean is the exception)

Not selected but worth mentioning: AI-designed tax systems (a simulated optimum is not the same as political feasibility); quadratic voting and proof of personhood (the bottleneck is identity infrastructure and adoption); AI prediction markets (the bottleneck is data and incentives).

23.2 Problem 1: Verifiable international AI agreements – verification without trust

In one sentence: Let two countries that do not trust each other confirm that the other is not secretly training beyond an agreed scale, without either side handing over weights or trade secrets.

Background: The precedent is nuclear arms control. In 2014 Glaser, Barak and Goldston proposed "zero-knowledge nuclear warhead verification" in *Nature*: inspectors can confirm that a warhead is genuine without learning anything about its design [established]. This is exactly the shape AI verification is after.

The intuition behind zero-knowledge proofs (ZKP): Your friend is red-green colorblind and holds two balls, one red and one green. Behind his back he randomly decides whether to swap them, then asks you "did I swap them?" If you couldn't tell the colors apart, the chance of guessing right 20 rounds in a row would be about one in a million. So he becomes convinced that you can tell them apart, yet he **still doesn't know which ball is red**. A ZKP generalizes this to any computation: it proves "I carried out the computation according to the rules" without revealing the inputs.

Back-of-envelope scale: training compute $\approx 6 \times \text{parameter count} \times \text{token count}$. 1 trillion parameters and 20 trillion tokens need **1.2×10^{26} FLOP** (the EU AI Act threshold for systemic risk is 10^{25}). An H100 delivers about 4×10^{14} FLOP/s effective, so the total is about 83 million GPU-hours: **about 347 days on 10,000 GPUs, or about 35 days on 100,000 GPUs**. A cluster that large is visible to satellites and on the power grid; the hard part is that **training and inference run on the same GPUs**, so the other side can say "we're running inference."

Tools from weak to strong: satellite and power-grid intelligence (catches only large-scale runs); chip location verification (prototype; legislation not finished $\triangle$); telemetry classification, which reads only power draw to judge whether training is happening (98.2% accuracy, 43–87% under disguise [recent]); tamper-resistant hardware mechanisms (demos exist, not in mass production); cryptographic proofs (only paper designs on the training side). Why cryptography is expensive: Kaizen (2024), on VGG-11 with about 10 million parameters, takes **about 15 minutes to prove each training step** [recent]; frontier models have about 10^5 times more parameters and millions of steps. In 2026-06, Peigné-Lefebvre et al. claimed that combining pre-commitment, network observation and a native floating-point zkVM can bring the overhead down to **single-digit percentages**, about five orders of magnitude below mainstream estimates; but **there is no prototype**, and a proof of concept is estimated to take about 36 months [recent, not reproduced].

Why it matters: Fearon (1995) argued that rational states fail to reach deals because of **private information** and **commitment problems** [established]. Verification shrinks the

former, and hardware with enforcement capability partly solves the latter, so verification changes the **game-theoretic equilibrium** of the AI race, not just engineering details.

What's blocking it: is math the bottleneck? Judgment: math/cryptography about 25%, engineering about 30%, institutions about 45%. But math is the **gate for the low-trust version**: without efficient proofs, the only option is intrusive inspection that great powers will not accept. Math determines "which institutional deals are on the table." The math parts: the efficiency of floating-point ZK; the soundness of Proof-of-Learning (which has been shown to be forgeable); the correctness of the ZK circuits themselves (under-constrained circuits are a common vulnerability Δ). The rest: what counts as "training" is a policy choice (type B); verification-grade chips take 6–10 years; who holds the keys, sovereignty, antitrust.

What AI can do: *AI math*: proving protocol soundness and formally verifying ZK circuits (type A); as proof costs fall to 1/10 each year, this becomes the default in 2027–2028; but whether the treaty specification is faithful remains type B (§17.2). *Other AI capabilities*: coding agents compress the "36 months to a prototype", and automated red-teaming turns 5 iterations into 500; in the other direction, they also help evaders design disguised workloads.

If solved: first-order impact: Compute caps go from "self-reported" to "provable"; export controls become enforceable, and violations can be proven with evidence.

If solved: second-order impact: The game shifts from "can't verify, so we must race" to "can verify, so we can negotiate", and a pause becomes a real option [speculation]. But the ability to disable chips remotely is itself enormous power, and the verification infrastructure needs governing too (isomorphic to Problem 8). The same primitives can also support verifiable evaluation.

Timing: *Optimistic*: telemetry and location verification enter export controls before 2027 (METR about 400 h by end of 2027). *Central*: a proof of concept for training-side ZK around 2029; binding international verification after 2030, and it will need a political shock; soundness proofs for the protocols become cheap by 2027, but they are not the rate-limiting step. *Pessimistic*: the low-trust version never lands. The cryptographic part depends on the proof-cost and solve-rate curves (§18); the hardware and political parts do not depend on any math curve.

Counterarguments and uncertainty: "Verification theater" could emerge; perhaps cryptography isn't needed at all, because satellites plus power data can catch large-scale training, and the real risk lies in small, efficient training runs.

Signals to watch: A chip location verification bill becomes law; the first public ZK proof of a training run with ≥ 1 billion parameters; "verification" appears in a US–China statement.

Sources: Glaser, Barak, Goldston 2014, *Nature* 510:497; Fearon 1995, *International Organization* 49(3); Kaizen ePrint 2024/162; Peigné-Lefebvre et al. arXiv:2606.05433; Rahman & Tajdari arXiv:2606.19262; Fang et al. 2023 arXiv:2208.03567; SNARK SoK arXiv:2402.15293 [△](#); Chip Security Act secondhand report [△](#).

23.3 Problem 2: Mechanism design for the AI agent economy – collusion without meetings

In one sentence: When both sides of a market are AI agents, we need rules that are efficient and also stop agents from spontaneously learning to raise prices or collude in secret; and antitrust law happens to be unable to catch "collusion without an agreement."

Background: *A repeated game by hand:* Two gas stations each have a cost of 1 yuan per liter; 100 customers buy only from the cheaper one, and split evenly if prices are equal. If both price at 2 yuan, each earns 50. If one secretly cuts to 1.9 yuan, it earns 90 on its own that period, after which the other retaliates with a permanent price war and profits go to roughly 0. Let the discount factor be δ (the closer to 1, the more the future matters). Cooperation is worth $50/(1-\delta) \geq 90$, so cooperating pays when **$\delta \geq 0.44$** . **An algorithm that reprices once a minute is extremely patient:** the next period is only a minute away, δ is almost 1, and collusion is very stable [established]. The US Sherman Act prohibits "agreements"; when two algorithms that never communicate each learn "if you cut, I cut", this tacit collusion is legally hard to prosecute.

Evidence: Q-learning pricing agents spontaneously learn collusion backed by punishment (Calvano et al. 2020) [established]; in the German gasoline market, margins rose 28% once both stations used algorithmic pricing (Assad et al. 2024) [established]; LLM pricing agents "quickly and autonomously" reach supracompetitive prices (Fish et al.) [recent]. The RealPage case reached a proposed settlement in 2025-11, with no fine and no admission of wrongdoing [recent]. On the other side: heterogeneity makes collusion fragile; when agents differ in patience, the premium falls from 22% to 10%; **but a difference in model size (32B vs 14B) instead produces a "leader–follower" structure that makes collusion more stable** (Keppo et al. 2026) [recent].

Secret collusion and program equilibrium: Agents can hide communication inside normal outputs using steganography, and information-theoretically secure steganography is in principle undetectable (Motwani et al.); this is the same problem as "the monitor colluding with the monitored model" in AI control (§22). Conversely, AIs can show their source code: when both sides submit programs that can read each other's code, cooperation is possible

even in a one-shot prisoner's dilemma (Tennenholtz 2004) [established]. This **cuts both ways**: it can support beneficial commitments (Problem 1), and it can also support collusion.

Why it matters: Collusion is one of the three types of multi-agent risk (Hammond et al. 2025). Once procurement, pricing and negotiation are handed to agents at scale, "collusion without meetings" goes from an edge case to a standing risk.

What's blocking it: is math the bottleneck? Judgment: math about 35%, law and institutions about 45%, data about 20%. Math: there is no clean equilibrium theory for the learning dynamics of LLMs as strategies; **identifiability**: in some settings it may be impossible in principle to tell collusion apart from independent best responses using price data; collusion-resistant mechanisms are hard in their own right, and VCG is fragile to collusion [established]. Law: the concept of an "agreement" breaks down, and liability is unclear. Data: real agent-to-agent transactions are rare.

What AI can do: *AI math*: proving bounds on whether collusion can be sustained under a given heterogeneity (type A); automated mechanism design is already emerging, and LLM-generated mechanisms have rediscovered the Myerson mechanism [recent]. *Other AI capabilities*: a simulated "market wind tunnel" before deployment; regulators auditing pricing agents with simulated opponents. In the other direction: agents will also get better at colluding and at steganography.

If solved: first-order impact: Agent platforms have enforceable rules (limits on data sharing, requirements for algorithmic diversity, mandatory price randomization), and enforcement shifts from "finding the agreement" to "auditing the design".

If solved: second-order impact: Antitrust shifts from "intent/agreement" toward design-based liability, similar to product liability [speculation]. Economic collusion and collusion in AI control are mathematically isomorphic and can supply each other with experiments. An agent economy that forms on its own, where humans can't tell where prices come from, is one path to gradual disempowerment (Kulveit et al. 2025).

Timing: *Central*: agent-to-agent transactions reach a significant share in 2027–2029 (METR about 150 h by mid-2027, roughly one work-month, enough to handle an entire procurement process); collusion-related type A theorems appear in large numbers in 2027–2028 (solve rate about 18%); the legal shift takes 5–10 years and does not depend on the math curves. *Optimistic*: the first enforcement case targeting LLM pricing agents comes in 2027. *Pessimistic*: the law stays stuck in RealPage-style ambiguity for a long time.

Counterarguments and uncertainty: Heterogeneity in real markets may be enough to suppress collusion; LLM agents may also be too "nice" and easy to manipulate (the Magentic Marketplace simulation).

Signals to watch: The first antitrust lawsuit centered on LLM pricing agents; a major platform publishes "rules for agent behavior".

Sources: Calvano et al. 2020 AER; Assad et al. 2024, JPE 132(3); Fish et al. arXiv:2404.00806; Keppo et al. arXiv:2603.20281; RealPage Fenwick; Motwani et al. arXiv:2402.07510; Tennenholtz 2004, *GEB* 49(2); Hammond et al. arXiv:2502.14143; arXiv:2502.12203; Magentic Marketplace arXiv:2510.25779.

23.4 Problem 3: Social choice in alignment – the math of "aligned to whom"

In one sentence: Training a model that serves hundreds of millions of people means compressing conflicting preferences into a single policy. Social choice theory says what guarantees each way of compressing offers and which guarantees cannot hold at the same time; in 2024–2026 this line of work has been producing type A theorems at a fast pace.

Background: *RLHF is, at bottom, voting.* Annotators give pairwise comparisons ("A is better than B"), and a Bradley–Terry reward model compresses them into scores. Siththaranjan et al. (ICLR 2024) proved that when annotators have hidden disagreements, standard preference learning **implicitly aggregates by Borda count** [recent, published].

By hand: 5 annotators rank 3 responses; 3 rank $A > B > C$ and 2 rank $B > C > A$. Pairwise majority: A beats B 3:2 and A beats C 3:2, so A is the Condorcet winner. Borda scoring (2/1/0 points): $A = 6$, **$B = 7$** , $C = 2$, so B is chosen. The majority prefers A in every pairwise comparison, yet Borda-style RLHF favors B. There is no unique answer to which is "right"; this is Arrow's theorem: every aggregation rule must trade off among axioms that each look reasonable [established].

Distortion: in the worst case, "the theoretically best average satisfaction $\div$ the average satisfaction the method actually achieves"; the closer to 1, the better. Gözl, Haghtalab and Yang (NeurIPS 2025) proved [recent, published] that NLHF (Nash Learning from Human Feedback) has distortion $(1/2 + o(1)) \cdot \beta$, which is **minimax optimal** (β measures how "certain" people are when they compare); for RLHF/DPO with a KL constraint, depending on how comparison pairs are sampled, distortion can be $e^{\Omega(\beta)}$ or even **unbounded**. This gives NLHF-type methods a quantitative rationale, and is an example of "math changing practice". In addition, Borda-style aggregation gives annotators an incentive to **misreport** their preferences, so strategyproofness becomes a safety property.

Why it matters: Every frontier model is shaped by a small number of annotators plus a constitution written by the company, yet it affects hundreds of millions of users (a 2024 ICML position paper by Conitzer, Russell et al. argues that social choice theory should guide this). "Aligned to whom" is the most political layer; making it explicit is what makes it auditable; otherwise it is still there, just hidden in the defaults of the training pipeline.

What's blocking it: is math the bottleneck? Judgment: math about 40%, normative and institutional questions about 60%. Type A: "Does algorithm X satisfy axiom Y?" "What is DPO's distortion under sampling scheme Z?" These are being solved in bulk. Type B: Which axioms matter? Is a "preference" a choice in the moment, or a considered value? Institutions: annotator pools are small and unrepresentative, the company has the final say, and there is no mandate.

What AI can do: *AI math* has solid precedents: a computer-aided proof of Arrow's theorem (Tang & Lin 2009); **SMT solvers used to discover new impossibility theorems** (Brandl et al. 2018 JACM) [established]. The "algorithm × axiom" table is well suited to AI filling it in at scale, and most of it will be cheap by end of 2027. *Other AI capabilities:* generative social choice (Fish et al. 2023) uses LLMs to predict each person's satisfaction with statements they haven't seen, which allows aggregation over open-ended text; the risk is that simulated preferences replace real people, which amounts to representing them without authorization.

If solved: first-order impact: Training adopts methods with explicit guarantees; model cards state the aggregation rule and the axioms it satisfies; strategyproof data collection becomes part of the defense against data poisoning.

If solved: second-order impact: Alignment becomes openly politicized [speculation]: once the rules are written down, groups will lobby over weights, like fights over electoral systems; today's "default to Borda" is also a political choice, just one nobody notices. There is also a tension between pluralism and personalization: giving each group its own model brings filter bubbles.

Timing: Type A theorems are already landing and will be largely "commoditized" by end of 2027 (depends on the solve-rate and proof-cost curves, §18); lab adoption in 2027–2029, depending on regulatory or reputational pressure; the normative questions have no timeline and will not be settled by math.

Counterarguments and uncertainty: RLHF violates almost every axiom yet still works well, perhaps because real preferences are highly structured. Arrow's theorem applies only to ordinal preferences; once preference intensity is introduced, many "impossibilities" loosen. Most disagreements are just matters of style.

Signals to watch: Frontier lab technical reports state their aggregation rule; NLHF-type methods are used in production models.

Sources: Conitzer et al. 2024 arXiv:2404.10271; Siththaranjan et al. arXiv:2312.08358; Gözl et al. arXiv:2505.23749; Munos et al. arXiv:2312.00886; arXiv:2506.12350; Fish et

al. arXiv:2309.01291; Tang & Lin 2009, *Artificial Intelligence* 173(11); Brandl et al. 2018, JACM 65(2).

23.5 Problem 4: Scaling democratic deliberation with AI

Deliberation faces a trilemma: referendums have broad but shallow participation, while citizens' assemblies go deep but involve only a few dozen people. AI promises to improve participation, depth and equality at the same time. DeepMind's Habermas Machine (*Science* 2024) has an LLM write position statements and a reward model predict endorsement, iterating repeatedly; among 5,734 UK participants, AI-mediated statements were preferred more often than those written by human mediators, and disagreement also narrowed [established, single high-quality study]. The clean precedent already in use is sortition: the algorithm of Flanigan et al. (*Nature* 2021) is offered free as Panelot and has been used in more than 40 citizens' assemblies [established, adoption count self-reported].

Math is about 20%: representativeness guarantees and Sybil resistance tend to get solved quickly once they are needed. The hard parts are type B questions ("maximizing endorsement" is not the same as "good deliberation") and institutional questions (there is no binding link between what deliberation concludes and what gets decided). AI's language ability plays the lead role, but it has a flip side: an experiment with 76,977 people and 19 LLMs found that post-training can raise political persuasiveness by up to 51%, and that **the methods that increase persuasiveness systematically reduce factual accuracy** (Hackenburg et al.) [recent]. **"Good at finding consensus" and "good at persuading" are the same capability**; whoever trains the mediation model defines "what counts as consensus", and is therefore a political actor [speculation]. Central judgment [speculation]: in 2027–2030 some government formally uses AI mediation in a statutory consultation; this depends on political will, not on the math curves. **Sources:** Tessler, Bakker et al. 2024 *Science*; Flanigan et al. 2021, *Nature* 596; Hackenburg et al. arXiv:2507.13919.

23.6 Problem 5: Next-generation market design – lessons from kidney exchange

Kidney exchange: patient A's relative a is willing to donate a kidney but has an incompatible blood type, and the same holds for B and b; if a matches B and b matches A, the two pairs swap. Finding the maximum number of transplants is NP-hard, and by 2007 it could already be solved at national scale with integer programming [established]. **But** Agarwal et al. (2019 AER) found that in 2014 the national platforms arranged only 42% of exchanges: hospitals bear the costs of participating but get little of the benefit, so they match pairs internally; fixing the fragmentation could **increase exchange transplants by 30–63%** [established]. **The algorithm had long existed; what blocked progress was incentives and regulation.** The counterexample is the FCC spectrum incentive auction: every bid required solving an NP-complete problem similar to graph coloring, the SATFC solver solved over 95% of

instances within the time limit, and revenue was **\$19.8B** [established]; without the solver, there would be no such market.

The successes in spectrum, school choice, kidneys and sortition all meet the same four conditions: ① **a single authoritative designer**; ② **a clear objective**; ③ **the computational or incentive problem really is the sticking point**; ④ **theoretical guarantees make participants willing to tell the truth**. Overall, math is about 30%, and can exceed 50% in new markets with physical constraints such as power grids. AI math will make "finding a strategyproof mechanism under constraints and proving it" a type A problem (textbook scale around 2027). **The biggest lever may be preference elicitation** [speculation]: the obstacle in combinatorial auctions is that bidders cannot price all 2^n bundles one by one, but an LLM agent can interview its principal and then bid on their behalf. The cheaper design becomes, the more the bottleneck concentrates on "who has the right to design"; the successes almost all happened where governments **created** markets, so AI may **increase** demand for state capacity [speculation]. A new national-scale market takes 5–10 years; no math curve can fix kidney fragmentation. **Sources**: Agarwal et al. 2019 AER 109(11); Newman, Fréchet, Leyton-Brown arXiv:1706.03304.

23.7 Problem 6: Content provenance – signatures vs. the watermark impossibility theorem

Signatures prove "this is real": a camera signs the image hash with its private key, changing a single pixel invalidates the signature, and C2PA links the capture and every later edit into a signature chain [established]; the weak points are stripped metadata, leaked private keys and untrusted devices, not the math. *Watermarks* try to prove "this was written by AI". The KGW scheme randomly splits the vocabulary into "green/red" halves and biases generation toward green words: a 200-word text written by a human is expected to contain 100 green words, with a standard deviation of about 7.07; if you count 140, $z \approx 5.66$, and the text is judged AI-generated. But have another model paraphrase it once, and the ratio returns to about 50%. Zhang, Barak et al. (2023) proved that as long as an attacker can judge whether quality is preserved after rewriting and can make small edits, a random walk can erase **any** watermark, even if the key is kept secret [established]. The conclusion: against disinformation, **what matters is verifying the source, not judging whether something is AI**.

Math is about 15%; the boundaries are basically drawn, and what remains is the trade-off between "robustness to edits" and "resistance to spoofing". AI math can formally verify C2PA implementations; cheaper proof systems can make zero-knowledge provenance practical (PhotoProof 2016 could already prove "this image is only a crop of a signed original" without revealing the original) [established]. The main bottleneck is adoption: Article 50 of the EU AI Act takes effect 2026-08-02, California's rules do not cover text [△](#),

and most platforms strip metadata on upload Δ . Second-order impact [speculation]: trust shifts from content to the identity of the signer, creating a new concentration of power. Central judgment [speculation]: by 2028 most new phones sign by default. Text watermarks are reliable only for long, unedited text, and this **will not** change as math advances, because it is a theorem. **Sources:** Zhang et al. arXiv:2311.04378; Kirchenbauer et al. 2023 PMLR; Naveh & Tromer 2016, IEEE S&P; EU Code of Practice; SB 942/AB 853 NatLawReview Δ .

23.8 Problem 7: Macroeconomics of transformative AI – labor share and explosive growth

In one sentence: If AI can do almost all cognitive work, what happens to growth, wages and the labor share? Estimates range from "TFP up only 0.66% over ten years" to "a tenfold increase in the growth rate", and the disagreement comes from a few empirical parameters that **cannot be derived mathematically**.

Background: *Labor share*: if the economy produces 100 yuan of output and 53 yuan of it is paid to workers, the labor share is 53%. In Q2 2026, the US nonfarm business sector labor share was **52.8%, the lowest since 1947** [recent]. But it has been falling for decades and may mainly reflect pre-AI automation; **the evidence for attributing it to AI is insufficient**.

Amdahl's law is the Baumol bottleneck: if a fraction p of a program is sped up s times, the total speedup is $1/((1-p) + p/s)$. With $p = 0.9$ and $s \rightarrow \infty$, the total speedup is **only 10x**. The economy works the same way: when tasks are complements (elasticity of substitution $\sigma < 1$), growth is held back by the tasks that are not automated; this is Baumol's cost disease. "AI can do 90% of the work" does not mean "explosive growth"; the key is whether the remaining 10% can be substituted or bypassed. On the other hand, if AI can produce ideas and compute can accumulate like capital, the loop "ideas $\rightarrow$ better AI $\rightarrow$ more ideas" can in theory accelerate growth (Aghion, Jones, Jones 2019, among others) [model result established; not empirically verified].

Estimates at both ends: At the moderate end, Acemoglu uses Hulten's theorem: GDP share of affected tasks $\times$ average cost savings $\approx 4.6\% \times$ about 14–15% $\approx$ **0.66% (total over ten years)**; critics say this counts only what AI has already shown it can do today. At the radical end, Restrepo's *We Won't Be Missed*: as compute expands, the "bottleneck work" that growth requires is fully automated, **wages tend toward "the cost of replicating that work with compute", and the labor share tends toward zero**; people still have work only because "it isn't worth replacing with compute". Keeping a job is not the same as sharing in growth [recent, theory]. Early empirical evidence: in the occupations most exposed to AI, employment of workers aged 22–25 fell 13% in relative terms (16% in the revised version), while experienced workers held steady (Brynjolfsson et al.) [recent].

Why it matters: Distribution, political stability and public finances all hinge on these parameters. This is also the deepest connection to AI safety: **gradual disempowerment** (Kulveit et al. 2025): once the economy no longer needs humans, institutions lose the incentive to serve humans; **the intelligence curse**: once the state no longer needs citizens to pay taxes, it no longer has to answer to them; **AI-enabled coups**: seizing power no longer requires the cooperation of large numbers of people.

What's blocking it: is math the bottleneck? Judgment: $\text{math} \leq 10\%$, typical type B; the models can be written in minutes. The bottleneck lies in **parameters** (elasticity of substitution, the automatable fraction, physical constraints such as energy), **data** (GDP has trouble measuring free digital products) and **identification** (short time series; tariffs and interest rates change at the same time as AI).

What AI can do: *AI math* plays a small role: it can solve more complex models, but the conclusions depend on the assumptions. *Other AI capabilities* play the lead: real-time measurement (payroll dashboards, labs' usage data), firm-level randomized experiments, simulating labor markets with LLM agents. Note the reflexivity: AI is both the shock being measured and the measuring instrument.

If solved: first-order impact: Better forecasts of transition paths; policies such as UBI, compute taxes and sovereign AI funds get an evidence base.

If solved: second-order impact: If Restrepo is right, the bargaining basis of democracy (strikes, taxes, representation) will erode, and the power-concentration problems in economics and in safety become directly connected [speculation]; he himself proposes a universal income, or treating compute as a public resource, like land. Explosive growth would speed up AI R&D and shorten the time left for safety research.

Timing: METR about 150 h by mid-2027 (one work-month), about 2,800 h by end of 2028. *Central* [speculation]: measurable macro effects appear in 2028–2030, e.g. productivity growth staying above 2.5%, with AI-intensive industries leading the decline in labor share. *Pessimistic*: as with electricity, diffusion lags by decades (David 1990). *Optimistic*: annual growth above 30% requires robotics and energy to keep up, and is only possible after 2030. These judgments do not depend on the solve-rate curve, only on agent capability and physical diffusion.

Counterarguments and uncertainty: Acemoglu-style moderate estimates may simply be right; Baumol bottlenecks in healthcare, education and government may be hard; GDP statistics may fail to capture the change.

Signals to watch: A sustained jump in TFP growth; AI-intensive industries leading the decline in labor share; entry-level employment continuing to diverge.

Sources: BLS 2026; Acemoglu NBER w32487; Restrepo NBER w34423; Brynjolfsson et al. Stanford DEL; Aghion, Jones, Jones 2019 (NBER w23928); David 1990, AER P&P; Kulveit et al. arXiv:2501.16946; Intelligence Curse; AI-Enabled Coups.

23.9 Problem 8: Political economy of AI-produced knowledge – who decides how it enters the public sphere

In one sentence: When one company can release 719 math papers in a day and the world's capacity to review them falls far short, who controls the production, verification and release of knowledge becomes a governance question.

Background: *The event* (see §2): On 2026-10-06, OpenAI released 722 manuscripts generated by an unreleased internal model; the next day a sign error led to a chain of withdrawals of 3 papers, leaving 719. About 42% of the main results have Lean formalizations, and the 3 withdrawn papers were all in the unformalized 58%. As of 10-08, **no human expert had publicly confirmed checking any of the headline results.** The disclosure norms of the IAS advisory group on mathematics and AI require stating how problems were selected, failure rates and so on [R] $\triangle$; this is exactly the selection effect discussed in §2.3. The AHM statement said: "Releasing over 700 files at once is not a demonstration of scholarship, but a demonstration of power."

The economics of verification, by hand (hypothetical numbers $\triangle$): Suppose a human review takes 20 hours per paper. 719 papers is then about 14,400 hours, or about **7 expert person-years**, while each subfield has only a few dozen experts who can follow the work. The 300 formalized papers need only a check of "whether the Lean statement is faithful to the intended meaning" (§3); at 2 hours each, that is 600 hours in total. **Lean cuts the cost of verification by about an order of magnitude, but not to zero.** Production cost collapses while human verification time stays the same, so **verification becomes the bottleneck, and whoever controls the bottleneck holds power.**

Intellectual property is not the lever of control: Mathematical results were never protected by copyright or patents, and the US has already held that copyright requires a human author and that AI cannot be an inventor [established]. The real levers of control are **timing, selection, disclosure and verification infrastructure**: what to release, when, with or without proofs, and what stays internal. In terms of Merton's (1942) norms of science, the ones under strain are "organized skepticism" (there is too much to review) and "communalism" (the company decides what is released).

Why it matters: Math is just the first field, because it has Lean; next come physics, biology and **AI safety research itself.** Dual-use knowledge [speculation]: if AI proves a result related to lattice cryptography or factoring, that is a national security issue, yet there is

currently no mechanism for deciding what to release. A mass release is also, in itself, a strategic display of an unreleased model's capabilities.

What's blocking it: is math the bottleneck? Judgment: math about 20%, institutions about 80%. The math part is autoformalization (moving that 58% into Lean) and methods for checking "whether the statement is faithful"; empirical science **has no equivalent of Lean**. The institutional part is journal policy, credit allocation, reviewer incentives and norms for mass releases.

What AI can do: *AI math*: proof costs fall to 1/10 each year, so by 2027–2028 "AI math must come with a Lean proof" will become a reasonable default, and human work will concentrate on checking whether statements are faithful; this is the social version of §17.2's "the trust bottleneck moves from proofs to specifications". *Other AI capabilities*: AI reviewers and replication agents. The flip side is an arms race between generation and review.

If solved: first-order impact: New release norms: either formalize, or label as "unverified"; staged releases that leave a window for review; disclosure following the IAS checklist.

If solved: second-order impact:

- **Enclosure of knowledge** [speculation]: science may slide toward "use internally first, release selectively", and the most valuable or most sensitive results become the least likely to be made public.
- **Understanding separates from checking** (Thurston 1994): humans can only check, not understand (§17.3); this is a civilization-scale scalable oversight problem, isomorphic to the core of AI safety (§22).
- **Verification as a public good**: we may need NIST-like public verification facilities (publicly funded Lean libraries, replication centers). §17.5's "frontier companies unilaterally deciding how capabilities enter the public sphere" is the same problem as Problem 1: how to check the claims of the powerful without having to trust them.

Timing: New norms form in mathematics within 1–2 years; by end of 2027 autoformalization covers most new AI math (central scenario, depends on the proof-cost curve). Empirical science will see similar "dumping" controversies in 2028–2030, with no technical fix. A release mechanism for dual-use knowledge will most likely be triggered by some incident.

Counterarguments and uncertainty: Publishing beats secrecy, and OpenAI could have released nothing at all; Lean is precisely what makes trust **independent** of power, since anyone can check; mandating formalization too early may make it harder for human mathematicians to publish.

Signals to watch: A second lab follows with a mass release; journals or arXiv adopt policies on mass AI submissions; the first public admission that "an AI-proven result was withheld for safety reasons".

Sources: This deck §2, §3, §17; openai/math; AHM statement via Tao; IAS advisory group norms cited only secondhand [R] [△](#); US Copyright Office 2025, *Copyright and AI, Part 2*; Thaler v. Perlmutter (2025); Thaler v. Vidal (2022); Merton 1942; Thurston 1994 arXiv:math/9404236.

23.10 Second-order effects in this field

1. The bottleneck migrates: production → verification → legitimacy. All eight cards follow the same pattern:

Card	What AI makes cheap	What becomes scarce	Who controls the scarce good
1 Verifiable agreements	Training (capability)	Credible verification	Hardware makers, key holders
2 Agent economy	Transactions and negotiation	Enforceable competition rules	Platforms, regulators
3 Social choice	Preference collection	Legitimacy of the aggregation rule	Labs
4 Deliberation	Mediation and summarization	Independence of deliberation	Whoever trains the mediation models
5 Market design	Mechanism design and solving	Design authority	Governments, centralized platforms
6 Provenance	Content	Trusted signatures	Device makers, certificate authorities
7 Macro	Cognitive labor	Physical inputs, political bargaining power	Compute owners
8 Knowledge	Proofs and papers	Human verification time	Publishers

2. Verifiability is the most underrated infrastructure. Zero-knowledge proofs, signatures, Lean and auditable aggregation rules are all technologies for "confirming without having to trust", and they are exactly where **AI math can contribute most directly** (protocol soundness, circuit verification and mechanism properties are all type A problems). The same people span several of these fields: Barak worked on zero-knowledge warhead verification and also proved the watermark impossibility theorem; Procaccia works on sortition algorithms and also analyzes the social choice properties of RLHF (arXiv:2501.09254). "Building trust without trust" could be one unified discipline.

3. Commitment technology is dual-use. Program equilibrium (Problem 2), hardware mechanisms (Problem 1) and signatures (Problem 6) all make credible commitment easier.

In Fearon's framework, this reduces conflicts that arise from commitment problems; but the same technology also makes collusion more stable and surveillance stronger. Every commitment technology has to answer: who can commit to what, and to whom.

4. Power concentration runs through every card. Designers, signers, verifiers, mediators, aggregators, publishers: every card has to ask "who holds the keys". AI lowers the cost of design but **raises the value of authority**. Taken together with AI-enabled coups and gradual disempowerment, economics and governance are not an appendix to the safety problem; they are the social layer of the same problem.

5. Cross-field connections. Old problems from distributed systems reappear here under new names: Amdahl's law corresponds to the Baumol bottleneck, Byzantine fault tolerance to treaty verification under low trust, and consensus protocols to commitment mechanisms.

23.11 Sources for this section

Full sources are at the end of each card. Suggested reading order: Fearon 1995 and Glaser–Barak–Goldston 2014 → Peigné–Lefebvre et al. 2026 → Keppo et al. 2026 → Gölz, Haghtalab, Yang 2025 → Agarwal et al. 2019 → Zhang et al. 2023 → Restrepo 2025 and Kulveit et al. 2025 → this deck §2, §17.

⚠ **To be rechecked:** The status of the Chip Security Act is known only from secondhand reports; accounts of the thresholds and dates in California's SB 942/AB 853 conflict; there is no systematic 2026 measurement of how much platforms strip C2PA metadata; details of the SNARK vulnerability SoK have not been rechecked; the IAS advisory group norms are cited only secondhand.

PART II · OUTLOOK · BIG PICTURE

24. The big picture: how the world changes if AI solves math

This section pulls §18–§23 together into one picture and answers the question posed at the start of Part II: **if AI solves math at scale, or AI's other capabilities break through as well, what happens to physics, biology, computing, AI itself and society? Which effects are truly far-reaching?** The facts and numbers in this section all come from earlier sections, with sources in the source list at the end of each; the rankings, generalizations and timelines here are my own synthesis and are tagged [speculation]. All timing judgments use the central scenario of the three curves in §18.4.

24.0 The bottom line

1. **The impact of "math being solved" falls into three tiers, and they differ enormously.**
 - **Tier 1: fields where the math itself is the product** (cryptography, formal verification, algorithms, mechanism design, part of mathematical physics). The impact is direct and fast, visible by 2027–2028.
 - **Tier 2: fields where math is just one link in the pipeline** (physics, chemistry, materials, climate, fusion). Math acts mainly through **algorithms and numerical methods**, constrained by compute, experiments and the pace of hardware; the impact lands in 2028–2035.
 - **Tier 3: fields where math is barely a bottleneck** (most of biomedicine, neuroscience, macroeconomics, governance). What matters here are AI's **other capabilities** (coding agents, simulation, automated experiments, literature matching), and they will be slowed by experimental cycles and institutions; the impact comes in the 2030s.
2. **The most far-reaching impact is not any single theorem, but the structural change to the whole knowledge-production system once "proofs become a cheap commodity".**

Across six fields we see the same phenomenon: making things got cheap, and **judging whether they are right, whether they are worth it, and what to ask** became the bottleneck.
3. **How fast AI accelerates a field depends on how fast that field's "referee" is.** Lean checks a proof in minutes and a weather forecast can be checked the next day, so AI is already having an effect in both; materials need synthesis and drugs need clinical trials, so their referees work in months and years; quantum gravity has no referee at all.
4. **AI is building "oracles" for science, not "theories".** Models that can predict but not explain are appearing in biology, physics and math at the same time. This is the same problem mechanistic interpretability faces.
5. **The exponential curves of AI capability will run into the linear time constants of the physical world and of institutions.** Progress in the digital world (software, math, AI itself) and in the physical world (drugs, materials, energy, hardware) will pull further and further apart.
6. **Verifiability will become the scarcest and most powerful resource.** Whoever holds the verifiers, signing keys and release channels decides how AI output enters society.
7. **2026–2029 is likely to be a fragile transition period of "fast offense, slow defense":** finding vulnerabilities is fast, patching them is slow; discovering attacks is fast, migrating standards is slow; generating claims is fast, checking them is slow; capability grows fast, legislation moves slowly.

8. **For AI safety, all of this converges on one conclusion:** when proofs, experiments and even alignment research itself can be automated, safety depends on three things: **whether statements and specs are faithful, whether the verifiers themselves can be trusted, and who judges fuzzy tasks.** None of the three has mature methods yet.

24.1 A direct answer: can breakthroughs in math bring breakthroughs in physics, biology and computing?

Your intuition is right, but it needs to be unpacked. Math enters the real world mainly through five channels, and their strength varies a lot:

Channel	Mechanism	Examples	Strength
1. The theorem itself is the product	A proof directly decides whether a system works and whether it is secure	Security reductions for, or breaks of, cryptographic schemes (§19); software with proofs (§19); incentive compatibility of mechanisms (§23)	Strong, fast
2. Math → algorithms → computation	New math becomes new algorithms that make previously intractable things computable	DFT functionals and molecular simulation (§20, §21); quantum error-correcting codes pulling Q-day earlier (§19); SAT solvers making the spectrum auction feasible (§23)	Strong, but limited by compute and data
3. "Impossibility" theorems change direction	Proving that a path is a dead end turns a whole field around	A universal DFT functional cannot be computed efficiently (§20); text watermarks cannot be made reliable (§23); undetectable backdoors exist (§22)	Medium; deep impact but not conspicuous
4. Connecting existing math to new problems	The math has long existed; what's missing is that "the people who need it know it exists"	The Radon transform behind CT scans (around since 1917); Gauss's FFT from 1805; LDPC codes, forgotten for 34 years (§18)	Possibly underrated
5. Model health checks	Proving that what physicists have long believed really holds	Navier–Stokes (§11), Vlasov–Maxwell (§12), Anderson localization and others (§20)	Weak: raises certainty, barely changes experiments or engineering

Three key points:

- **Most of OpenAI's 719 papers belong to channel 5 and to "inventory" math.** Going by the historical pattern (§18.1), math done for its own sake takes 50–200 years to find applications, while math done for engineering problems takes only 2–15 years. So the real-world impact of this batch of results will be measured in decades, unless AI also speeds up the "matching" (channel 4). [speculation]
- **Channel 2 is the main entry point into physics, chemistry and biology, and it depends more on AI's other capabilities than on its ability to prove.** As of 2026, every confirmed

AI-for-science speedup (weather forecasting, protein structure) came through this channel; none came from proofs (§18.3).

- **Think in terms of Amdahl's law:** in drug development, making "theory" infinitely fast raises overall speed by only about 5% (the back-of-envelope calculation in §18.2). The smaller the share of math in a field, the less acceleration "solving math" brings, however big the breakthrough.

One summary table: what blocks each field, and how much AI can help (see §18.7 and the individual sections for details)

Field	Is math the bottleneck?	The real bottleneck	Most likely AI lever	Central timing
Cryptography and security (§19)	Yes	Proofs + standards migration + patching speed	AI math (analysis, proofs) + AI coding (finding vulnerabilities, writing verification code)	2027–2029
Quantum computing (§19)	Partly	Hardware	Resource estimates and error-correcting codes (type A); new algorithms (type B)	Q-day 2033–2036
Chemistry and materials (§20)	Partly (algorithms)	Computational accuracy + synthesis throughput	ML functionals and force fields; autonomous labs	2028–2032
Fusion and energy (§20)	Rarely	Hardware cadence, materials irradiation	Control, design optimization	2030s
Proteins and molecules (§21)	Yes (a sampling problem)	Sampling algorithms + validation data	Generative models that sample directly from the Boltzmann distribution	2029
Cells, development, aging (§21)	No	Measurement, data, concepts	Automated experiments, oracle-style models	2030–2035
Drugs and the clinic (§21)	No	Target causality + 8–13-year pipeline + regulation	Target selection, patient stratification	2035
Neuroscience (§21)	Partly (missing theory)	Measurement + concepts	Connectome reconstruction, digital twins	2033
AI theory and alignment (§22)	Yes, but mostly type B	Definitions (complexity, capability, explanation, corrigibility)	Automated controlled experiments; bulk sweeps of type A theorems	2027–2030
Economics and governance (§23)	In a few places (ZK, social choice)	Institutions, adoption, legitimacy	Mechanism solving, verifiable agreements	2029 and later

24.2 Five cross-field patterns

Each of these five patterns showed up independently in every one of §18–§23. They are the most important content of this Part.

24.2.1 Proofs are cheap, statements are expensive

AI is making "given a problem, find the answer" cheaper and cheaper, but "is the problem itself stated correctly?" has not gotten any cheaper.

Field	The cheap part	The expensive part
Math (§3, §17)	Writing Lean proofs	Checking whether the Lean statement is faithful to the paper
Software (§19)	Writing proofs	Getting the spec right, delimiting the trusted computing base (TCB)
Cryptography (§19)	Improving known attacks	What the security model leaves out (timing, implementation)
Materials (§20)	Generating candidates (GNoME predicted about 380,000 stable crystals)	Confirming that candidates are real and useful (only 736 synthesized externally)
Biology (§21)	Designing proteins, training predictive models	Knowing what to measure, and whether evaluation metrics are valid
AI theory (§22)	Proving theorems about solvable models	Defining "complexity", "capability", "explanation", "corrigibility"
Governance (§23)	Producing papers, content, transactions	Checking, signing, legitimacy

Implication: The scarcest capability in the future will be turning vague intent into precise statements, and checking whether those statements have been quietly weakened. In math this is called "statement faithfulness", in software "the spec", in AI "the alignment target"; **at bottom it is the same capability.**

24.2.2 The speed of the referee sets the speed of acceleration

Field	Referee	One feedback loop	AI acceleration
Formal math	Lean kernel	Minutes	Already happened (2026)
Weather	Next day's observations	Days	Already happened (ECMWF operational in 2025)
Software security	Tests, fuzzing, proof checking	Minutes to days	Already happened (finding vulnerabilities); patching can't keep up
Electronic structure	High-accuracy calculations	Hours to days	Fast
Protein design	Wet lab	Weeks	Medium, limited by the cost of experiments
Materials, superconductors	Synthesis + characterization	Weeks to months	Slow
Clinical medicine	Clinical trials	Years	Very slow (a single Phase III takes 5–8 years)

Field	Referee	One feedback loop	AI acceleration
Climate	Observations	A decade	Very slow
Quantum gravity	None	∞	Only on subproblems

Implication: If you want AI to accelerate a field, the most effective investment is often not a stronger model but a **faster referee**: autonomous labs, smaller and faster Phase IIa trials, better evaluation metrics, automated connectome proofreading.

24.2.3 Prediction succeeds, understanding fails

In several fields AI has produced models that "can answer what-if, but can't say why":

- Biology (§21): generative protein design bypasses physics-based scoring functions; virtual cell models learn "perturbation → expression" without building mechanisms; AlphaGenome learns "sequence → regulatory signals" without building gene regulatory networks; the digital twin of mouse visual cortex can predict the responses of more than 10,000 neurons, but is as much a black box as the brain itself.
- Physics (§20): neural network functionals fit energies more and more accurately, while the densities may drift further and further off.
- Math (§3, §17): 26 million lines of Lean proofs, almost none of which anyone has read.
- A precedent from neuroscience: the complete connectome of the nematode's 302 neurons has existed since 1986, and forty years later we still cannot predict its behavior (§21).

Implication: Scientific output may shift from "theories" to "validated oracles". Capability will grow while understanding may stall. This is exactly the predicament mechanistic interpretability faces with AI, and biology already has more, and more realistic, examples of it. The two fields should learn from each other; at present they almost never do. [speculation]

24.2.4 Exponential curves hit linear time constants

Curve or constant	Speed
METR task length	Doubles every 129 days (about $\times 7$ per year)
Cost of formal proof	About $\div 10$ per year
Clinical development pipeline	8–13 years
Mouse lifespan experiments	3–4 years
Irradiation qualification of fusion materials	Several years, incompressible
Rollout of a new national-scale market design	5–10 years

Curve or constant	Speed
Legislation (e.g. updating the nucleic acid synthesis screening framework)	90 days mandated; in practice more than a year overdue

Implication:

- Progress in the digital and the physical world will **diverge**. If AI R&D is automated (the CASP scenario in §22, $r > 1$), timelines for software, math and AI itself will be compressed, while those for drugs, materials, energy and hardware will be compressed much less.
[speculation]
- To judge whether "AI is really accelerating science", look at outputs that can't be gamed: clinical success rates, the Eroom's law curve, the number of materials actually synthesized by external labs, the share of code that comes with proofs. These indicators respond on a timescale of 5–10 years.

24.2.5 Verifiability becomes the scarcest and most powerful resource

Zero-knowledge proofs, digital signatures, Lean proofs and auditable aggregation rules are all technologies for "confirming without having to trust" (§19, §23). AI makes production cheap, so verification and legitimacy become scarce; **whoever controls verification controls how AI output enters society**.

- OpenAI releasing 719 papers at once on 2026-10-06 is the first clear case: a frontier company unilaterally deciding in what form a capability enters the public sphere (§2, §23).
- Google used a zero-knowledge proof to disclose its resource estimate for a quantum attack without publishing the attack circuit; Anthropic first published hashes of vulnerability details and released the details after the fixes (§19). **Dangerous math and security results are starting to need "responsible disclosure"**, and this will become standard practice when AI discovers dangerous results. [speculation]
- Commitment technology is dual-use: the same mechanisms can make international agreements verifiable and can also make collusion between AI agents more stable (§23).

24.2.6 Addendum: the fast-offense, slow-defense transition window

Offense (fast)	Defense (slow)	Source
High-severity vulnerabilities found by AI: 90.6% of 1,752 spot-checked were real	About 530 disclosed, only 75 fixed (as of 2026-05-22)	§19
Qubits needed to break RSA-2048: about 20 million in 2019 → under 100,000 in 2026 (qLDPC)	Post-quantum migration has to go through years of standards processes	§19
The threat side follows the 129-day doubling curve	No federal legislation mandating nucleic acid synthesis screening	§21

Offense (fast)	Defense (slow)	Source
The speed of generating papers and claims	The speed of expert review and statement checking	§2, §22

Implication: 2026–2029 is likely to be a window in which offense reaps the benefits of AI first; defense's benefits (rewrites with proofs, post-quantum migration, mandatory screening) have to wait for institutions to catch up. [speculation]

24.3 Top 10: the problems most worth watching

Ranking criteria: size of the impact × likelihood of being significantly changed by AI before 2030 × relevance to AI safety. The ranking itself is [speculation].

#	Problem	Section	Why it matters	Is math the bottleneck?	Central timing	Key signals
1	Statement faithfulness and verifier security: does what the AI proved say what we think it says	§22, §19	Every "it has a proof, so it can be trusted" argument rests on this; 2026 already saw a Lean kernel bug that can "prove" False, and forged proofs got past even independent checkers	No (auditing and engineering)	Before 2027, an important result is withdrawn because of a toolchain bug or a wrong formalization	Withdrawal events; formally verified checking kernels become the default
2	AI doing AI research, including alignment research	§22	Automated researchers already outperform 28 experienced human researchers on well-defined alignment problems, while also gaming evaluations; this decides "whether safety can outpace capability"	No (review bandwidth)	AI-led share of R&D passes half in 2027; fully autonomous share first rises above zero in 2027–28	Automation shares disclosed by labs; evaluation-cheating rate of automated research
3	Formal verification of critical software × cyber offense and defense	§19	The only defense that can keep pace with AI vulnerability-finding is "proving the vulnerabilities don't exist"; AI agents' sandboxes are themselves software	Partly (proofs are getting cheap, specs are not)	Proofs stop being the main cost in 2027–2028; the offense window lasts until about 2028	Major cloud providers' crypto libraries and parsers ship with proofs by default
4	Q-day and post-quantum migration	§19	The lifespan of internet encryption and crypto assets; math (error-correcting codes, circuit optimization) is pulling Q-day earlier	Partly (resource estimates are type A)	A machine that can break ECC-256 in 2033–2036; resource estimates fall another 2–10x before 2028	New resource-estimate papers; migration progress

#	Problem	Section	Why it matters	Is math the bottleneck?	Central timing	Key signals
5	A computational microscope at molecular scale: electronic structure and protein dynamics	§20, §21	Unlocks the "scoring functions" for batteries, catalysts, drug binding and protein design; the only place in biology where math is truly the choke point	Yes (algorithms and sampling)	Main-group chemistry practically solved in 2028; globular protein conformations in 2029	ML functionals enter mainstream codes; free-energy prediction error < 0.5 kcal/mol
6	Verifiable international AI agreements	§23, §19	A precondition for a "credible pause" and for compute governance; verification shrinks the private information and commitment problems that lead to conflict	Partly (ZK proofs are type A; hardware and politics are not)	ZK proof of concept for the training process around 2029; binding agreements after 2030	ZK prototype for billion-parameter-scale training; location verification in export controls
7	Inductive bias: which of countless "equally correct" solutions training picks	§22	Goal misgeneralization, scheming, emergent misalignment and weak-to-strong generalization are all essentially this problem	Yes, but type B	A few reliable empirical laws before 2030, no unified theory	Successful preregistered cases of "predicting before training which goal the model will learn"
8	Biosecurity governance	§21	The only field where "a capability evaluation directly triggering deployment restrictions" has already happened; the attack and defense clocks differ by orders of magnitude	No (institutions)	Mandatory screening in some jurisdictions in 2029	Legislation on mandatory screening; end-to-end capability evaluations
9	Labor share and gradual disempowerment	§23	A labor share approaching zero would weaken the political bargaining power that comes from "people being needed"; US labor share in Q2 2026 was 52.8%, the lowest since 1947 (evidence for attributing this to AI is insufficient)	No (empirical parameters and institutions)	Measurable macro effects in 2028–2030	Productivity growth staying above 2.5%; AI-intensive industries leading the decline in labor share

#	Problem	Section	Why it matters	Is math the bottleneck?	Central timing	Key signals
10	Social choice in alignment: "aligned to whom"	§23	Theorems already show RLHF/DPO can be arbitrarily far from the best average outcome, while NLHF is provably optimal; type A theorems can directly improve training methods	Yes (type A)	Theorems largely commoditized by end of 2027; lab adoption in 2027–2029	Frontier labs switch to aggregation methods with guarantees

Just missed the top ten: the "computational no-coincidence conjecture" in interpretability (§22; a single pure-math result could change the research portfolio); clinical translation and Phase II success rates (§21; huge impact, but hard to see results before 2035); fusion (§20; dominated by hardware cadence); virtual cells (§21; dominated by data and metrics); collusion in the AI agent economy (§23).

24.4 Timeline: late 2026 → the 2030s

All dates are the central scenario, with an error of roughly ± 1 year; the optimistic scenario moves everything half a year to a year earlier, the pessimistic one a year or two later. The capability line comes from the three curves in §18.4.

T0 • Now (2026 Q4)

- **Capability:** AI independently completes tasks of about 1–2 workdays; single-attempt solve rate on preregistered open problems 3–5%; formalization costs about \$0.01–0.02/line.
- **Already happened:**
 - AI weather forecasting is operational;
 - a batch of classic mathematical physics problems have been claimed solved, but none has been peer reviewed;
 - AI finds vulnerabilities at scale, and patches are backlogged;
 - a soundness bug appeared in the Lean kernel;
 - automated researchers outperform humans on well-defined alignment tasks;
 - the math community has started discussing AI in the language of alignment ("misalignment", Goodhart).
- **Falsification signal:** a new long-task benchmark shows progress stalled at 1–2 days.

T1 • 2027

- **Capability:** about 150 hours mid-year ($\approx$ one work-month), about 400 hours at year end; open-problem solve rate about 9% mid-year and about 18% at year end; proof costs fall by

another order of magnitude.

- **Computing and security:** papers on AI-assisted improvements to cryptographic attacks become common; substantial progress on the "proximity gap" conjecture that zero-knowledge proof systems rely on; in the optimistic scenario, AI can translate a 10,000–50,000-line C library into a version with proofs.
- **Physics:** most well-known type A problems in mathematical physics are claimed solved, and more than half of the claims are confirmed by experts.
- **AI and safety:** compact proofs for small models are automated; the theory of scalable oversight in idealized models largely converges; at least one important AI formalization result is withdrawn or seriously challenged because of a toolchain bug or a wrong formalization.
- **Society:** social choice theorems for alignment are largely commoditized; agent-to-agent commercial transactions start to reach scale.
- **Falsification signals:** the next round of FrontierMath Erdős is still below 5%; the AI-led share of R&D stalls around 30%.

T2 · 2028

- **Capability:** about 2,800 hours at year end (more than a work-year); open-problem solve rate about 45%; formalizing something the size of Fermat's Last Theorem costs only a few thousand dollars.
- **Computing and security:** proof labor is no longer the main cost; formally verified checking kernels become the default; Q-day resource estimates fall by another 2–10x.
- **Physics and chemistry:** main-group chemistry is "solved" in a practical sense; the SPARC fusion device reaches $Q > 1$.
- **AI and safety:** most type A problems in safety theory are solved or proven impossible, and what remains is all definitional; the social credibility of "it has a Lean proof" as evidence is put to the test.
- **Society:** most new phones sign photos and videos by default; measurable AI macro effects start to appear (2028–2030).
- **Falsification signal:** task length stalls at a few hundred hours, and curve ① turns into an S-curve.

T3 · 2029–2030

- **Capability:** open-problem solve rate about 77% by end of 2030.
- **Computing and security:** code with proofs covers the "critical small kernels" (cryptography, OS kernels, hypervisors, parsers, consensus protocols), but makes up less than 1% of all code; the probability of a practical classical break of ML-KEM is $\leq 5\%$.

- **Biology:** conformation and free-energy prediction for common globular proteins matures; binder design becomes largely an engineering task; virtual cells are reliable for specific cell types and pathways; some jurisdictions implement mandatory nucleic acid synthesis screening.
- **Physics and materials:** consensus on the phase diagram of the Hubbard model (2028–2030); AI-designed materials enter niche commercial products (2029–2032); central judgment for the Yang–Mills mass gap is 2030–2033.
- **AI and safety:** a few reliable empirical laws about generalization, but no unified theory.
- **Society:** proof of concept of zero-knowledge proofs of the training process (around 2029); binding international verification agreements only after 2030, and most likely only after a political shock.

T4 · The 2030s

- **Q-day:** a quantum computer that can break ECC-256, central judgment 2033–2036; RSA-2048 a few years later.
- **Medicine:** the overall drug success rate rises from 7.9% to 10–12% (2035); interventions against specific age-related diseases succeed, but "aging" itself still has no causal theory.
- **Neuroscience:** the whole-brain mouse connectome is complete or nearly complete (around 2033); there is still no consensus on "the brain's learning algorithm".
- **Energy:** commercial fusion power in the mid-to-late 2030s.
- **Scenario fork:**
 - **AI R&D fully automated** (the CASP scenario): timelines in the digital world compress across the board; how much the physical world compresses depends on whether robotics, labs and energy can scale up in step.
 - **Hard bottlenecks appear** (compute, energy, spec review): all type A problems get solved, while type B problems and the physical world proceed at their old pace.
- **True in both scenarios:** security guarantees for software and protocols are far stronger than today; guarantees in the open world (biology, society) remain weak and rely mainly on monitoring, access control and governance; **who reviews the statements and specs** is the ceiling for the whole system.

24.5 Deeper second-order effects: what this means for human knowledge and society

This subsection is inference, not prediction; all of it is [speculation].

1. **The nature of knowledge will change.** More and more knowledge will be "verifiable but understood by no one": machine-checked proofs, experimentally validated oracle models.

The role of mathematicians and scientists will shift from "producing results" to "choosing problems, designing definitions, reviewing statements and interpreting what results mean". The core point of Thurston's 1994 essay *On proof and progress in mathematics* is that the value of a proof lies in advancing human understanding; the question in 2026 is: once proofs no longer require understanding, who will do the understanding, and will anyone pay for it?

2. **Human understanding may become a luxury and a bottleneck.** Sequencing costs fell 10,000-fold in a decade, yet biology did not get 10,000 times faster; proof costs will probably go the same way. The real rate-limiting step is how fast people can digest, trust and use the results.
3. **The Baumol effect will reshape the economy.** Sectors that AI can accelerate (software, math, design) will become relatively cheaper, and sectors bound by physical constraints (healthcare, construction, energy, experiments) relatively more expensive. Value will flow to physical inputs, verification capacity and legitimacy.
4. **Power will concentrate at the nodes of verification and release.** Hardware makers (verification mechanisms in chips), holders of signing keys, frontier labs (which decide what to release and in what form) and those who train mediation models will all become new nodes of power.
5. **Science itself will see reward hacking.** There are already examples: AlphaEvolve exploiting loopholes in verification code; the Virtual Cell Challenge's metric being gamed to the limit; erroneous problems in FrontierMath; a forced Navier–Stokes result that satisfies the problem statement but is not the question people actually care about. **"The science of evaluation" will go from a subtopic of AI safety to infrastructure for all of science.**
6. **Release mechanisms for dual-use knowledge don't exist yet.** Cryptanalysis, chemical synthesis, protein design and vulnerability discovery are all accelerating. Mathematical results are starting to need "responsible disclosure", but outside biosecurity there are almost no institutions for this.

24.6 What this means for AI safety research

Putting together the research gaps identified in §18–§23, five themes stand out. What they share: **most of them have automatic referees, allow real experiments, and matter more as AI capability grows.**

1. **Auditing verifiers and statements**
 - Reward hacking of provers, verifiers and evaluations: port "score-seeking propensity" evaluations to Lean and Dafny environments; test whether cheating learned during math RL generalizes into misalignment elsewhere (§22).

- Methods for auditing statements and specs: error taxonomies, cross-consistency checks, measuring the throughput of human reviewers (§22, §19). This deck's line-by-line reading of Lean statements (§4–§14) is an early prototype.
- Evaluations that score AI-generated scientific claims for "verifiability" (§20).

2. Oversight experiments with real capability gaps

- Formal math offers an oversight testbed with "AI far stronger than humans + perfect ground truth" (§22).
- Port the institutions of clinical trials (preregistration, independent monitoring, mandatory reporting of negative results) to AI evaluation (§21).

3. Dual-use capability evaluations

- Evaluations of cryptanalysis and vulnerability-patching capability (§19).
- Biology: end-to-end capability evaluations using safe proxies, measuring what "model + cloud lab" enables a person with basic training to do (§21, governance level only).
- Threat models and permission tiers for autonomous labs: AI gets physical execution authority for the first time (§20).

4. Verifiable governance and multi-agent settings

- Verification red-teaming: have agents automatically run "monitor vs. evader" iterations to test the verification mechanisms of compute governance (§23).
- Agent collusion propensity evaluations: collusion and steganography in multi-agent pricing settings (§23).
- Auditing RLHF aggregation: compare Borda-style reward models with NLHF on public preference data (§23).
- Verifiable inference for AI governance: use zero-knowledge proofs to prove model identity and that filters were enforced (§19).

5. Measuring AI's real speed-up of science, and "matching"

- Build an index of experimental throughput and cost for biology: how fast AI changes biology is set by experimental cost, and nobody measures that curve systematically yet (§21).
- Work on "matching": connecting existing mathematics to the people and problems that need it. Historically (§18.1) this step has often paid off sooner than proving new theorems.
- Watch the counter-signal: if an AI-proven "inventory-type" theorem reaches a product within 3 years, matching costs have collapsed and this section's timelines should all move earlier (§18).

- Measure whether cheap proofs make researchers crowd onto the type A problems AI is good at: there is already evidence that scientists using AI produce more individually while the collective range of topics narrows (§18.6).

Two general principles:

- Don't do by hand what AI will be able to do within a year, unless it is infrastructure or helps you judge "what to ask".
- Prioritize "definition problems": once the type A problems are cleared, definitions and statements are the remaining bottleneck, and the place where human judgment is hardest to replace.

24.7 Where this projection is most likely wrong

1. **Measurement of the capability curves is breaking down.** METR's task set is unreliable above 16 hours; whether cheating counts as success can shift the results by about 4.5 doublings. Numbers after 2027 rest mainly on "other signals rising in step" (§18.4).
2. **The boundary between type A and type B will move.** AI may become good at "proposing good statements and definitions" sooner than expected. There is no evidence yet that this has happened: OpenAI's 719 papers mainly solve existing problems and propose almost no new concepts. Once it does happen, most of the timelines in this section will move earlier, and the human role will shrink further.
3. **Many of the cited results have not been peer reviewed.** This applies to OpenAI's 719 papers, the many 2026 arXiv preprints, and the numbers companies report about themselves; the search budget during the research was limited, and some numbers come from memory or secondhand sources; these are marked $\triangle$ in each section.
4. **Institutions may suddenly speed up after a shock,** for better or for worse. A major accident could get legislation passed within months, or trigger an overreaction.
5. **Physical constraints may be broken through faster than expected by AI-driven robotics and automated labs,** which would weaken the conclusion that the digital and physical worlds diverge.
6. **Selection bias.** The problem list in this Part consists of problems already known to be important. Historically (§18.1), the biggest impacts often came from applications no one anticipated: number theory became cryptography, and the Radon transform became CT. The most far-reaching impacts may well come from outside this list.

24.8 Sources for this section

This section introduces no new facts; the sources for all numbers are in the source lists at the end of §18–§23. The full research drafts for the six fields (complete versions of every

problem card, more sources, and material that did not make it into the deck) are kept in the author's research notes and are not published with this deck.

REACTIONS · IN THEIR WORDS

25. What people said

WHO	QUOTE	SOURCE
Levent Alpöge (Anthropic)	"It's obviously the most significant moment in mathematical history." (later adding there would be "sad stories" of people being scooped)	X [R]
Terence Tao	"Problems are being solved autonomously by AI prompters who have no interest in the broader field itself"; proposes Math 1.0 → 2.0	Fortune [R]
Dan Litt (Toronto)	"My view is that this is great for mathematics," while worrying that an "AI has solved math" perception could discourage young people	Fortune [R]
Tristan Buckmaster (NYU)	"There's likely to be a bunch of results where they take someone's work and then take it to completion"	Fortune [R]
Scott Aaronson	Post "The Mathocalypse"; on L = BPL: "its truth was never in serious doubt"	blog [R]
Lance Fortnow	"if they hold up, we've seen more progress in TCS in the last 24 hours than in the previous three decades combined"	Computational Complexity blog [R]
Mark Braverman	"Math by press release is not that healthy for math."	Quanta [R]
François Chollet	Asks whether gains in RLVR-friendly domains generalize, or whether non-verifiable domains stay bottlenecked on human data	X (paraphrase) [R]
AHM	"Releasing over 700 files at once is not a demonstration of scholarship, but a demonstration of power."; urges mathematicians to stop working with OpenAI	reposted on Tao's blog (Tao notes he does not endorse it) [P]
Álvaro Lozano-Robledo	"Keep calm and carry on studying math."; LLMs "may work within the confines of the convex hull of ideas that are currently available in the literature"	guest post on Tao's blog [P]
Andrew Sutherland (MIT)	"We should ask for receipts."	via Scientific American [R]

Note: "The Future of Mathematics" on Tao's blog (10-05) is a guest post by Jeremy Avigad, not by Tao; some coverage conflated the two.

26. Open questions and how to keep learning

26.1 What to watch

1. Will OpenAI publish the list of ~4,000 problems and the failures? (This decides whether "what is the real capability?" can be answered.)
2. How many errors like the Weil-class paper remain in the ~58% that is unformalized? How will the withdrawal rate evolve?
3. Who reviews the faithfulness of the Lean statements? When does `review: unchecked` change? Who audits the `lana-agents` dependencies and the large patches?
4. When will the first human domain expert give an independent "read and confirmed" verdict on a headline result (quasi-RH, UGC, Thompson F)?
5. What priority norms should govern "completing someone else's work" (the Buckmaster episode)?
6. Chromatic number of the plane: what does a finite unit-distance graph needing 6 colors look like?

26.2 Learning path (starting from linear algebra + calculus)

STAGE	MATERIAL	WHY
1 · What math is	Gowers, <i>Mathematics: A Very Short Introduction</i>	150 pages by a Fields medalist: abstraction, proof, what "hard" means
1 · Roots of this debate	Thurston, On Proof and Progress in Mathematics; Gowers, The Two Cultures of Mathematics; Tao, What is good mathematics?	after these, "Math 1.0 → 2.0" makes complete sense
2 · Touch a proof	Natural Number Game (Kevin Buzzard, free online)	prove commutativity of addition from scratch in Lean — then §3 makes full sense
2 · Learn proofs systematically	MIT 6.042 Mathematics for Computer Science (OCW + free textbook)	designed for CS backgrounds: proof methods, discrete math, probability
3 · Intuition	3Blue1Brown: <i>Essence of linear algebra</i> ; visualizing the Riemann zeta function	the figure in §4 is a static version of that episode
3 · Following the news	Quanta Magazine (math); <i>The Princeton Companion to Mathematics</i> (as a reference)	every piece explains why a problem is hard and where people were stuck
4 · ML-adjacent	Bubeck, <i>Convex Optimization: Algorithms and Complexity</i> ; Prince, <i>Understanding Deep Learning</i> (free)	convergence proofs for optimization; intro to deep-learning theory

27. Glossary and sources

27.1 Glossary

TERM	IN ONE LINE
Lean / Mathlib	a proof assistant / its community math library
sorry / axiom	in Lean: "leave this step empty" / "assume this is true" — both weaken a proof
Comparator	tool checking "the theorem proven = the challenge statement, using only allowed axioms"
zero-free region	a region where ζ or an L-function has no zeros
irrationality exponent μ	how "unusually well" a number can be approximated by fractions
ω	the optimal exponent of matrix multiplication, $2 \leq \omega \leq 3$
approximation ratio	the fraction of the optimum a polynomial-time algorithm guarantees
L / RL / BPL	log-space deterministic / one-sided randomized / two-sided randomized computation
amenable	a group admits a translation-invariant "average"
blowup	a solution of a smooth equation going to infinity in finite time
replica symmetry breaking (RSB)	Parisi's description of the hierarchical state structure of spin glasses
galactic algorithm	asymptotically faster, but only better at astronomical sizes

27.2 Main sources

Primary: openai/math (README, history.md, overview.pdf, CONTENTS.md, lean/, preprints/, reasoning_traces/; snapshot fd4aeeb, 2026-10-08) · openai/NavierStokesAndEuler · Clay NS official statement (Fefferman) · leanprover/comparator · independent re-checks davegoldblatt/openai-zeta-proof-check, erenciracioglu-dotcom/openai-math-9-4-lean-check · Epoch FrontierMath Erdős (arXiv:2609.25050).

Coverage and commentary: Fortune 10-07 · Nature (paywalled, full text not read [△](#)) · Scientific American · Quanta 09-08 · Quanta 10-07 · NPR 09-22 · Aaronson · Kalai · Tao's blog · Latent Space.

Classic references: Strassen 1969; Coppersmith–Winograd 1990; Alman et al. arXiv:2404.16349; Ambainis–Filmus–Le Gall arXiv:1411.5414; Goemans–Williamson 1995;

Khot 2002; KKMO 2007; Raghavendra 2008; Khot–Minzer–Safra 2018; Nisan 1992; Saks–Zhou 1999; Reingold 2008; Roth 1955; Zeilberger–Zudilin [arXiv:1912.06345](#); Meiburg [arXiv:2208.13356](#); Guth–Maynard [arXiv:2405.20552](#); de Grey [arXiv:1804.02385](#); Moore [arXiv:1102.0747](#); Tao 2016 [arXiv:1402.0290](#); Glassey–Strauss 1986; Parisi 1979; Talagrand 2006; Mézard–Parisi 2001; Ding–Sly–Sun 2015.